CS101 Algorithms and Data Structures

Fall 2020

Homework 9

Due date: 23:59, November 23, 2020

1. Please write your solutions in English.

2. Submit your solutions to gradescope.com.

3. Set your FULL NAME to your Chinese name and your STUDENT ID correctly in Account Settings.

4. If you want to submit a handwritten version, scan it clearly. Camscanner is recommended.

5. When submitting, match your solutions to the according problem numbers correctly.

6. No late submission will be accepted.

7. Violations to any of the above may result in zero grade.

1: (4×2') Choice

The following questions are choice questions, each question may have **one** or **multiple** correct answers. Select all the correct answer, you will get half points if you choose a strict subset(excluding empty set) of the right answer.

*Note: You should write those answers **in the box** below.*

Question 1	Question 2	Question 3	Question 4
BC	A	D	AB

Question 1. *Which of the following statements are/is true?*

(A) *Prim's algorithm for computing a minimum spanning tree of a graph only works if the weights of the edges are non-negative.*

(B) *A minimum spanning tree for a connected graph has exactly $N - 1$ edges, where N is the number of vertices in the graph.*

(C) *A connected graph that has unique edge weights (there are **no** two edges with the same weight) has a unique minimum spanning tree.*

(D) *None of above is True*

Question 2. *Which of the following statements are/is true?*

(A) *For Prim's algorithm, it is possible to calculate the minimum spanning tree of some graph in $\Theta(E)$ time without using priority queue and adjacent list. E represents the number of edges in the graph.*

(B) *Kruskal's algorithm adds the least-weighted edge each time to the MST.*

(C) *The removal of some edge from an MST will **not** disconnect the graph.*

(D) *None of above is True.*

Question 3. *Consider the following algorithm that attempts to compute a minimum spanning tree of a connected undirected graph G with distinct edge costs. First, sort the edges in decreasing cost order (i.e., the opposite of Kruskal's algorithm). Initialize T to be all edges of G. Scan through the edges (in the sorted order), and remove the current edge from T if and only if it lies on a cycle of T.*
*Which of the following statements is true?*

(A) *The output of the algorithm will never have a cycle, but it might not be connected.*

(B) *The algorithm always outputs a spanning tree, but it might not be a minimum cost spanning tree..*

(C) *The output of the algorithm will always be connected, but it might have cycles.*

(D) *The algorithm always outputs a minimum spanning tree.*

Question 4. *Which of the following algorithms reflect(s) the idea of greedy algorithms?*

(A) *Kruskal*

(B) *Prim*

(C) *Quicksort*

(D) *Mergesort*

2: (8') Play with Prim and Kruskal

Here we give the adjacent list of a DAG(Directed Acyclic Graph). The letters on the left of the arrows are the source vertices of the edges. The destination vertex v and weights of the edge w is given in the form of $\boxed{v\ w}$ on the right of the arrows.

$$\boxed{A} \to \boxed{B\ 3} \to \boxed{C\ 5} \to \boxed{D\ 2}$$
$$\boxed{B} \to \boxed{C\ 4} \to \boxed{D\ 6} \to \boxed{E\ 4}$$
$$\boxed{C} \to \boxed{D\ 2}$$
$$\boxed{D} \to \boxed{E\ 10} \to \boxed{F\ 7}$$
$$\boxed{E} \to \boxed{F\ 3}$$
$$\boxed{F}$$

Suppose we start from vertex "A".

(1)(2') Draw the graph of the DAG with the adjacent list given above. Annotate the edges with their weight.

(2)(2') View the DAG above as an undirected graph and draw its **maximum** spanning tree.

(3)(2') Write down the sequence of edges added to the **maximum** spanning tree using Kruskal's algorithm.

(4)(2') Write down the sequence of edges added to the **maximum** spanning tree using Prim's algorithm.

**Solution:**

(1):

(2):

(3):
(D, E), (D, F), (D, B), (A, C), (B, C)

(4):
(A, C), (B, C), (B, D), (D, E), (D, F)

3: (12') Minimum Spanning Tree

We are given the following graph G:

For each question below, please give your answer firstly and then explain it briefly.

Note that questions below are independent to each other.

(1)(2') For what range of values of g are you **guaranteed** to include edge BE in the MST?

Solution: Construct a cut where one set is E and the other set is (A, B, C, D, F). In order to guarantee g is in the MST, it must be the lightest edge: $g < 8$.

(2)(2') For what range of values of k are you **guaranteed** to include edge CD in the MST?.

Solution: There is a cycle using edges AD, AC and CD. The heaviest edge in the cycle will not be used, so in order to make CD not the heaviest edge, we have: $k < 2$.

(3)(2') For what range of values of k are you **guaranteed to NOT** include edge CD in the MST?.

Solution: There is a cycle using edges AD, AC and CD. The heaviest edge in the cycle will not be used, so in order to make CD the heaviest edge, we have: $k > 2$.

(4)(3') Suppose an adversary can set g and k arbitrarily. What is the maximum cost of a MST that the adversary can force?

Solution: MST will only choose BE and CD if it is better than or equal to the cuts of the other edges, so the maximum cost is $1 + 2 + 3 + 4 + 8 = 18$.

(5)(3') Draw seven edges on the following graph with six vertices, and then mark six edges with cost=1 and one edge with cost=x such that an adversary can choose x such that the cost of the MST is arbitrarily large.

Solution:

4: (8') Building A Tower

We build up a tower with a pile of stones. Each layer of the tower is exactly one piece of stone. We index the layers from top to bottom, so layer 1 is the top layer. The stone weight of the i-th layer is s_i. The tower will fall if there exists a layer i s.t. $\Sigma_{j=1}^{i-1} s_j > s_i$. For example, if the tower is built with 5 stones and the stone weight from top to bottom is 2, 3, 5, 9, 100, then the tower would fall because $9 < 5 + 3 + 2 = 10$.

You are the designer of the tower and you have n stones. Each stone has different weight and the same height. Can you build the tower as high as possible? Give an $O(n\log n)$ algorithm and prove that the algorithm produces the highest tower. Show the time complexity of your algorithm. Your answer should contain at least the following parts:

1. Main idea of your algorithm (which can also include pseudo-code here). (2')
2. Proof of the correctness of your algorithm. (4')
3. Time complexity analysis. (2')

Solution:

1. Main idea: First sort the piles w.r.t. their weights. Second consider picking $i - th$ pile to our tower where i ranges from the lightest pile to the heaviest one. If $\sum_{j=1}^{i-1} s_j \leq s_i$, we add pile s_i to our tower. Else, we skip this pile and consider $(i+1) - th$. Keep doing this until the last pile, and the number of piles we have chosen represents the highest height.

2. Proof idea:

Denote our algorithm A, and assume there is an optimal algorithm OPT.

Let the piles A picks are (from light piles to heavy piles)

$$a_1, a_2, a_3, \cdots, a_k$$

and let the piles OPT picks are (also, from light piles to heavy piles)

$$b_1, b_2, b_3, \cdots, b_l$$

Assume the first $m - 1$ piles are the same of the 2 algorithms, which means

$$a_m \neq b_m$$

Based on our algorithm A, we guranteen that a_m is the lightest pile among all the rest piles (and it won't make the tower fall). Therefore we have

$$a_m < b_m$$

Then, let's replace the pile b_m by a_m in the answers return by OPT algorithm. After doing this, $b_1, \cdots, b_m$ become exactly the same as $a_1, \cdots a_m$. Besides,

$$\sum_{i=1}^{m} b_i > \sum_{i=1}^{m-1} b_i + a_m$$

, which means this operation won't affect $b_{m+1}, \cdots, b_l$.

Thus, we can keep doing this transformation and the answer return by OPT will be the same as answer return by our algorithm A. ∎

3. Time complexity: Sorting step takes $O(n\log n)$, and picking the piles from the lightest to heaviest one takes $O(n)$. In total, $O(n\log n)$.

5: (9') To be A Grandparent

Assume you are a grandparent and going to give your grandchildren some pieces of cake. However you cannot satisfy a child unless the size of cake piece he receives is no less than his expected size. Different children may have different expected sizes. Meanwhile, you can't give each child more than one piece. For example, if three children's expected sizes are 1, 3, 4 and the sizes of pieces of cake are 1, 2, then you could only satisfy the first child. Given the expected sizes of the children and the sizes of cake pieces that you have, how can you make the most children satisfied? Show that your algorithm is correct and analyze the time complexity. Same as Q4, your answer should contain at least the following parts:

1. Main idea of your algorithm (which can also include pseudo-code here). (3')
2. Proof of the correctness of your algorithm. (4')
3. Time complexity analysis. (2')

Solution

1. Algorithm:

Always trying to satisfy the child with the least expectation with the smallest piece of cake. More specifically, sort all children w.r.t. their expectation sizes and sort all pieces of cake w.r.t. their sizes. Say children after sorting: $a_1, a_2, \cdots, a_n$. Cake after sorting: $b_1, b_2, \cdots, b_m$. Give b_1 to a_1 if b_1 can satisfy a_1. If b_1 cannot satisfy a_1, we then consider $b_2, b_3, ...$ and so on. Keep doing this until: no cake is left or no children are left.

2. Proof idea:

Assume an optimal algorithm OPT gives child-cake assignments:

child	a_1	a_2	...	a_m	...	a_l
cake	b_1	b_2	...	b_m	...	b_l

Assume our algorithm A gives child-cake assignments:

child	a'_1	a'_2	...	a'_m	...	a'_k
cake	b'_1	b'_2	...	b'_m	...	b'_k

(We might suppose that in both assignments cakes are arranged in ascending order and this assumption will not affect the proof.)

Let m be the maximum value for

$$a_1 = a'_1, b_1 = b'_1; a_2 = a'_2, b_2 = b'_2; \cdots ; a_{m-1} = a'_{m-1}, b_{m-1} = b'_{m-1}$$

i.e. the first $(m-1)$ assignments from A and OPT are the same. Then it means that there is a disagreement on the $m-th$ pair: (a_m, b_m) and (a'_m, b_m') are not the same.

There are 3 cases for the $m-th$ pair being different:

1) $a_m = a'_m$ but $b_m \neq b'_m$ (Same child $a_m(a'_m)$ gets different size of cake)
2) $a_m \neq a'_m$ but $b_m = b'_m$ (Different children get the same size of cake $b_m(b'_m)$)
3) $a_m \neq a'_m$ and $b_m \neq b'_m$ (Different children get the different sizes of cake)

For case 1, we have a claim that $b_m > b'_m$ because, according to our algorithm, b'_m is the smallest piece of cake among all the rest and b'_m will also satisfy the $m-th$ child. Also, we have a claim that b'_m is not in $\{b_1, b_2, \cdots, b_l\}$ because we have supposed that b_i's are sorted and $b'_m < b_m$. Thus, we can simply replace b_m in the OPT solution by b'_m without any conflicts.

For case 2, in OPT solution, child a'_m must be somewhere behind a_m or he/she does not even get a piece of cake. If in OPT solution child a'_m does not get a piece of cake, we can replace a_m by a'_m in OPT solution without conflicts. Else if a'_m is somewhere behind a_m, we can also switch these 2 children and they are still satisfied respectively. The reasons are as follows. For a_m he/she will get the piece of cake b that WAS assigned to a'_m which is bigger than his/her origin cake(Of course he/she will still be satisfied!). For a'_m, b_m will still satisfy him/her because in case 2, $b_m = b'_m$.

For case 3, similar to case 2, in OPT solution child a'_m must be somewhere behind a_m or a'_m does not get a piece of cake. We just switch a'_m and a_m without losing optimality. Then it becomes case 1.

Therefore, we can always transform the optimal solution to our solution without losing optimality. Our algorithm is also optimal. ∎

3. Time complexity:

Say number of children and number of pieces of cake are both $\theta(n)$. We need to do sorting twice and this takes $O(nlogn)$. Then "always trying to satisfy child with least expectation with the smallest cake" takes $O(n)$. So in total, we need $O(nlogn)$