

# CS101 Algorithms and Data Structures

Graph traversal  
Textbook Ch 22.2/3/5

# Outline

- Graph traversal
  - Breadth-first
  - Depth-first
- Applications
  - Connectedness
  - Unweighted path length
  - Identifying bipartite graphs

# Graph Traversal

## Traversals of a graph

- A means of visiting all the vertices in a graph
- Also called *searches*

Similar to tree traversal, we have breadth-first and depth-first traversals on graphs

- Breadth-first requires a queue
- Depth-first requires a stack

# Graph Traversal

Different from tree traversal: there may be multiple paths between two vertices.

To avoid visiting a vertex for multiple times, we have to track which vertices have already been visited

- We may have an indicator variable in each vertex
- We may use a hash table or a bit array
- Requiring  $\Theta(|V|)$  memory

The time complexity of graph traversal cannot be better than and should not be worse than  $\Theta(|V| + |E|)$

- Connected graphs simplify this to  $\Theta(|E|)$
- Worst case:  $\Theta(|V|^2)$

# Breadth-first traversal

Breadth-first traversal on a graph:

- Choose any vertex, mark it as visited and push it onto queue
- While the queue is not empty:
  - Pop the top vertex  $v$  from the queue
  - For each vertex adjacent to  $v$  that has not been visited:
    - Mark it visited, and
    - Push it onto the queue

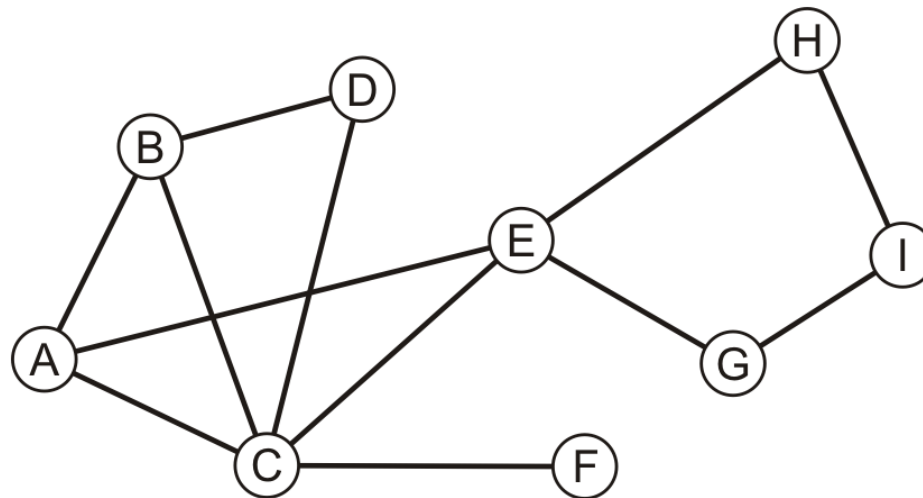
This continues until the queue is empty

- If there are no unvisited vertices, the graph is connected

The size of the queue is  $O(|V|)$

# Example


Consider this graph



# Example

Performing a breadth-first traversal

- Push the first vertex onto the queue




# Example

Performing a breadth-first traversal

- Pop A and push B, C and E

A






# Example

Performing a breadth-first traversal:

- Pop B and push D

A, B




# Example

Performing a breadth-first traversal:

- Pop C and push F

A, B, C




# Example

Performing a breadth-first traversal:

- Pop E and push G and H

A, B, C, E

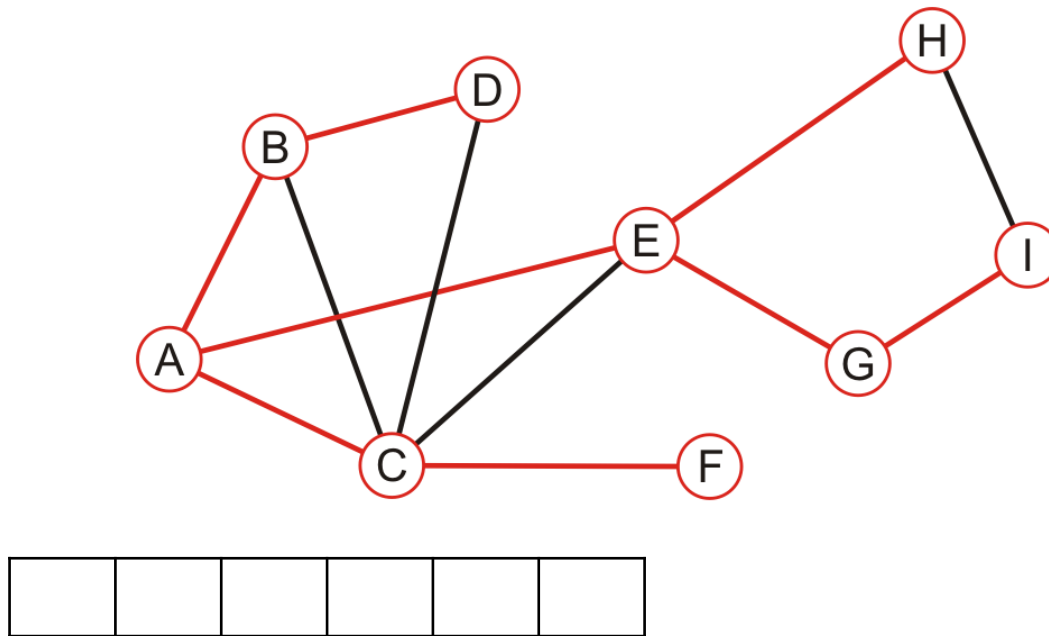


# Example

Performing a breadth-first traversal:

- The queue is empty: we are finished

A, B, C, E, D, F, G, H, I



# Iterative breadth-first traversal

An implementation can use a queue

```
void Graph::depth_first_traversal( Vertex *first ) const {
    unordered_map<Vertex *, int> hash;
    hash.insert( first );
    std::queue<Vertex *> queue;
    queue.push( first );

    while ( !queue.empty() ) {
        Vertex *v = queue.front();
        queue.pop();
        // Perform an operation on v

        for ( Vertex *w : v->adjacent_vertices() ) {
            if ( !hash.member( w ) ) {
                hash.insert( w );
                queue.push( w );
            }
        }
    }
}
```

# Depth-first traversal

Depth-first traversal on a graph:

- Choose any vertex, mark it as visited
- From that vertex:
  - If there is another adjacent vertex not yet visited, go to it
  - Otherwise, go back to the previous vertex
- Continue until no visited vertices have unvisited adjacent vertices

Two implementations:

- Recursive
- Use a stack

# Depth-first traversal

A recursive implementation:

```
void Vertex::depth_first_traversal() const {  
    for ( Vertex *v : adjacent_vertices() ) {  
        if ( !v->visited() ) {  
            v->mark_visited();  
            v->depth_first_traversal();  
        }  
    }  
}
```

# Depth-first traversal

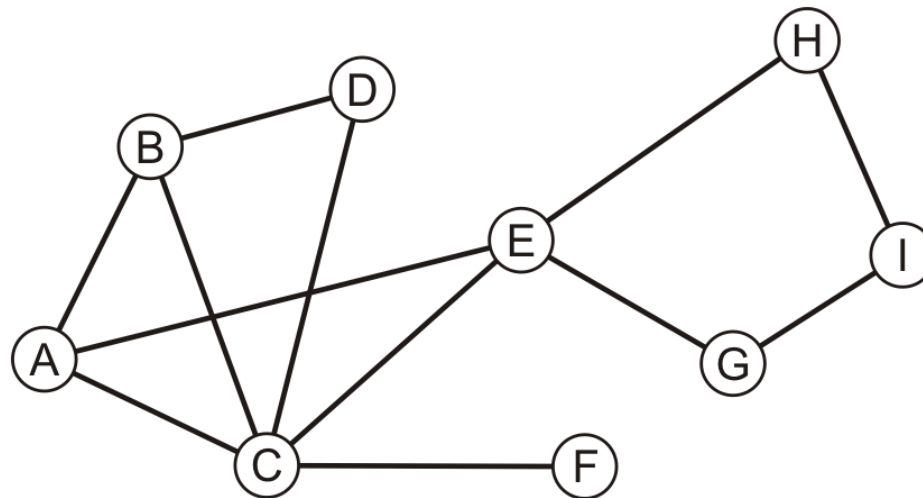
Use a stack:

- Choose any vertex
  - Mark it as visited
  - Place it onto an empty stack
- While the stack is not empty:
  - If the vertex on the top of the stack has an unvisited adjacent vertex  $v$ ,
    - Mark  $v$  as visited
    - Place  $v$  onto the top of the stack
  - Otherwise, pop the top of the stack



# Example

Perform a recursive depth-first traversal on this same graph

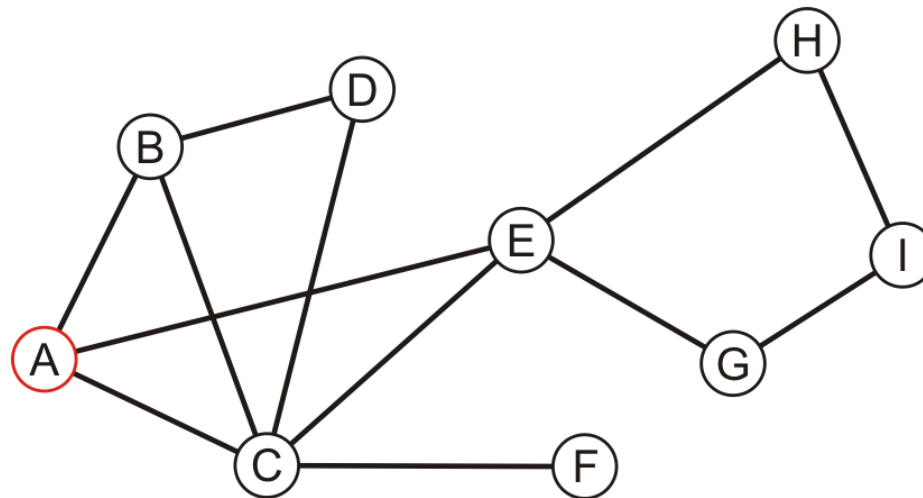


# Example

Performing a recursive depth-first traversal:

- Visit the first node

A

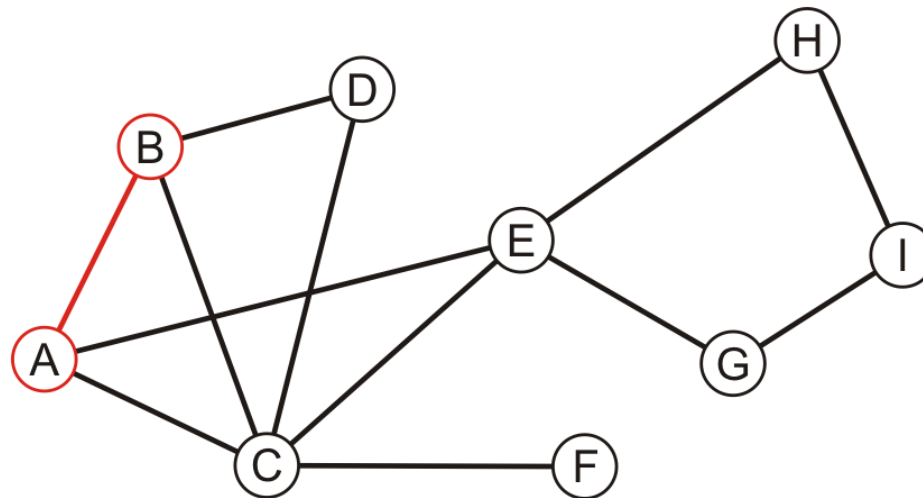


# Example

Performing a recursive depth-first traversal:

- A has an unvisited neighbor

A, B

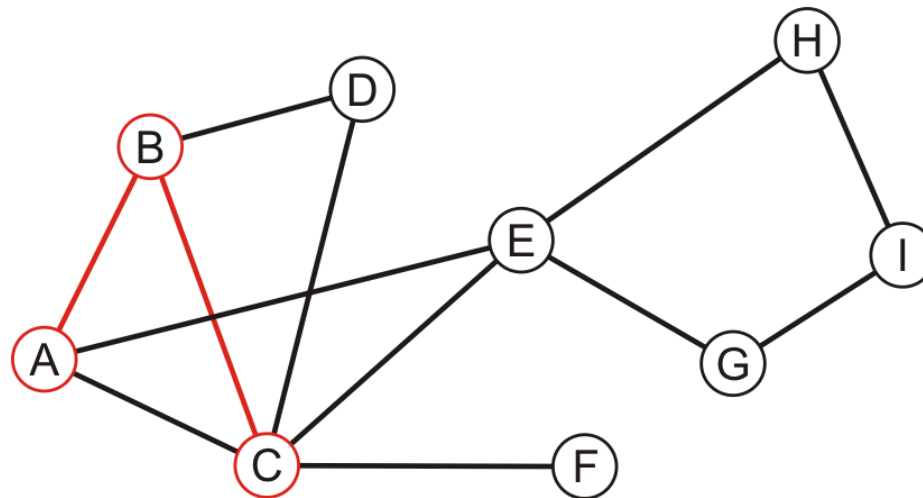


# Example

Performing a recursive depth-first traversal:

- B has an unvisited neighbor

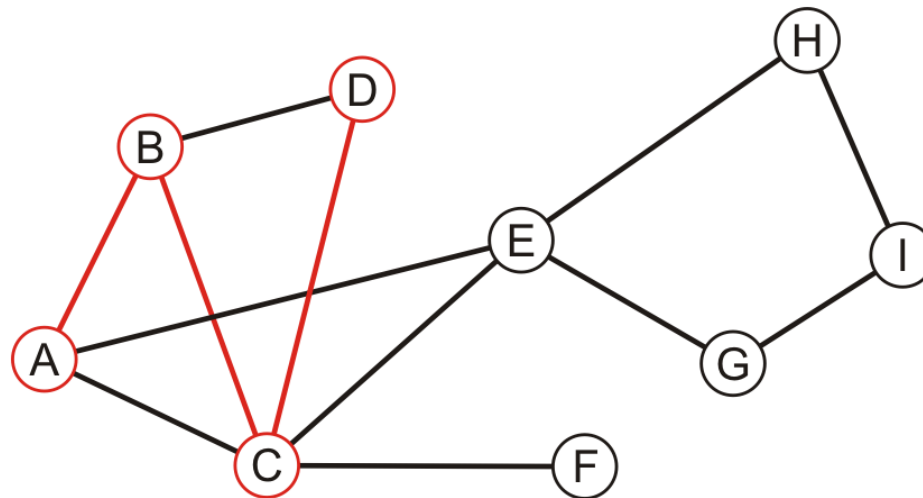
A, B, C



# Example

Performing a recursive depth-first traversal:

- C has an unvisited neighbor  
A, B, C, D

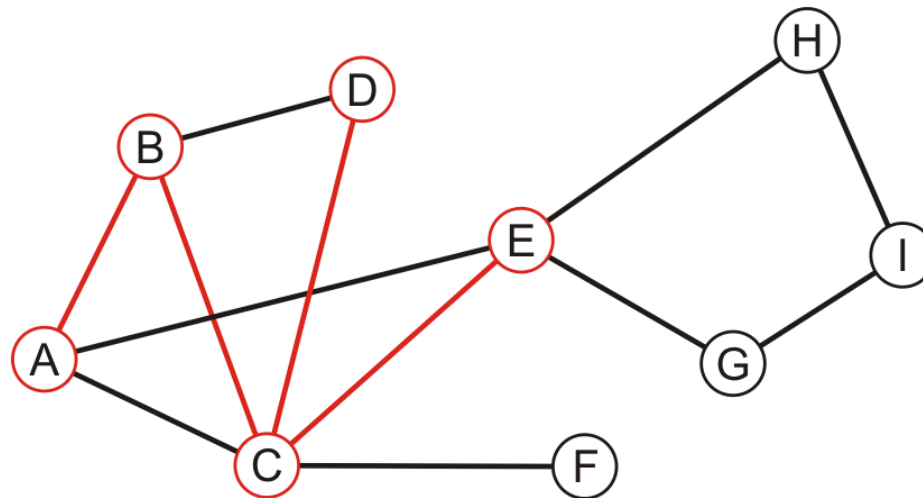


# Example

Performing a recursive depth-first traversal:

- D has no unvisited neighbors, so we return to C

A, B, C, D, E

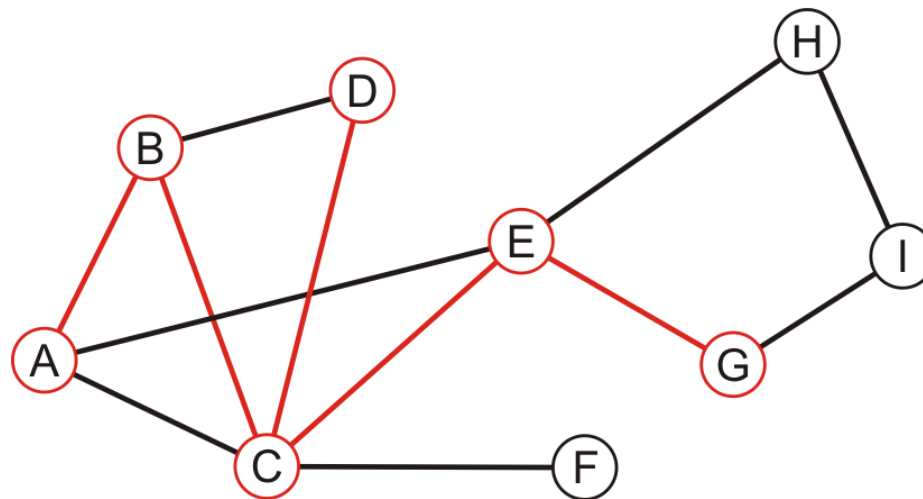


# Example

Performing a recursive depth-first traversal:

- E has an unvisited neighbor

A, B, C, D, E, G

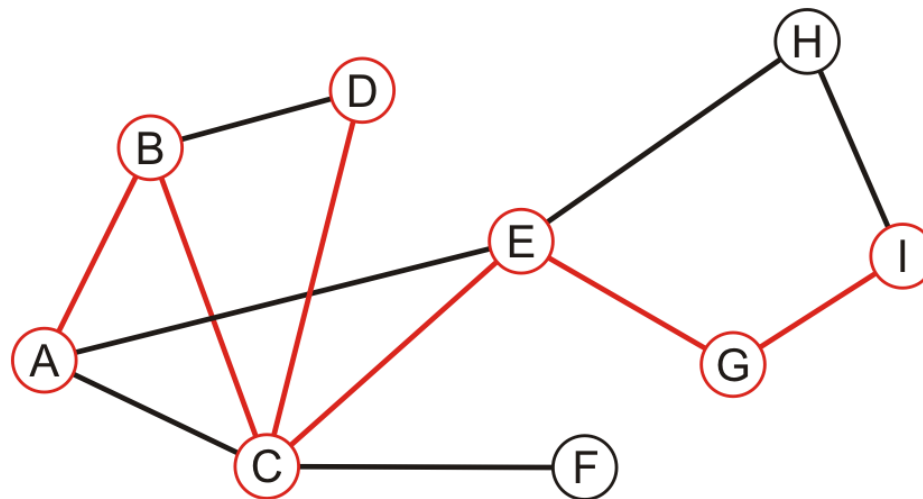


# Example

Performing a recursive depth-first traversal:

- F has an unvisited neighbor

A, B, C, D, E, G, I



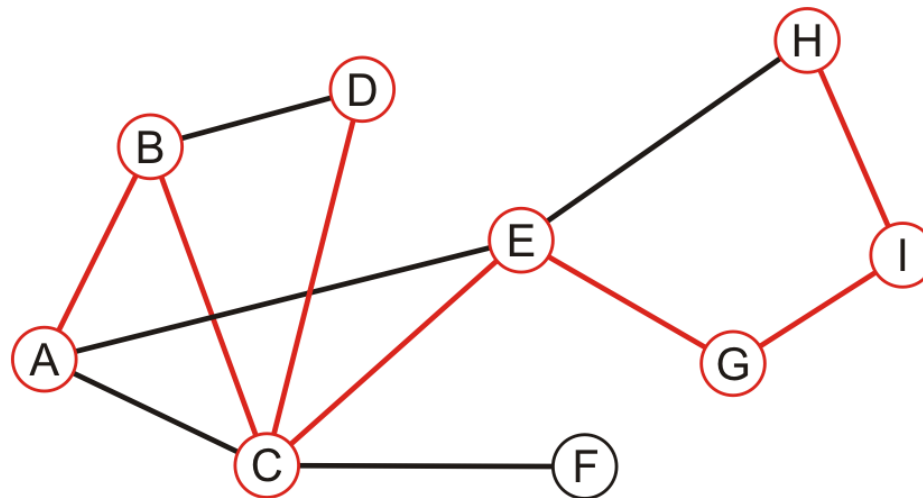


# Example

Performing a recursive depth-first traversal:

- H has an unvisited neighbor

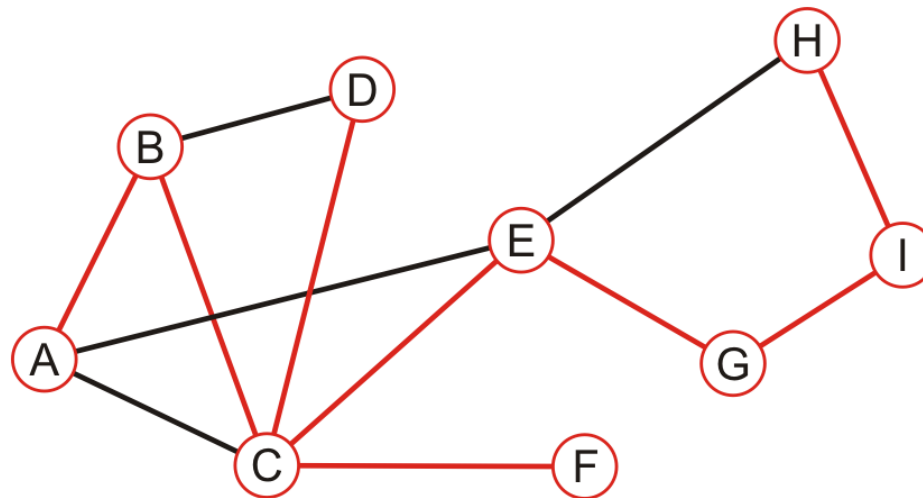
A, B, C, D, E, G, I, H



# Example

Performing a recursive depth-first traversal:

- We recurse back to C which has an unvisited neighbour  
A, B, C, D, E, G, I, H, F




# Example

Performing a recursive depth-first traversal:

- We recurse finding that no other nodes have unvisited neighbours

A, B, C, D, E, G, I, H, F

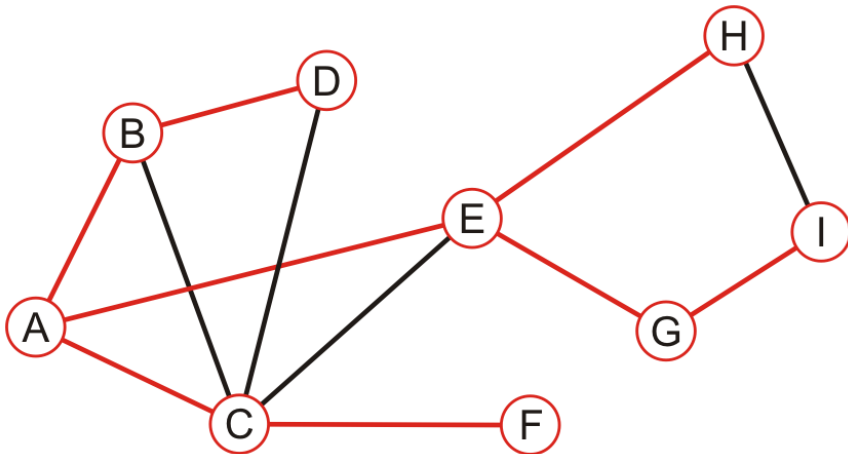


# Comparison

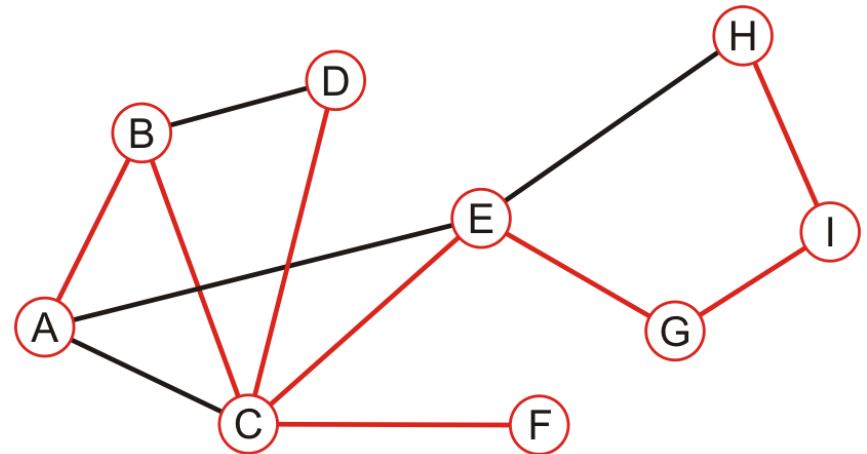
The order in which vertices can differ greatly

- An iterative depth-first traversal may also be different again

A, B, C, E, D, F, G, H, I



A, B, C, D, E, G, I, H, F



# Outline

- Graph traversal
  - Breadth-first
  - Depth-first
- Applications
  - Connectedness
  - Unweighted path length
  - Identifying bipartite graphs

# Connected

First, let us determine whether one vertex is connected to another

- $v_j$  is connected to  $v_k$  if there is a path from the first to the second

Strategy:

- Perform a breadth-first traversal starting at  $v_j$
- If the vertex  $v_k$  is ever found during the traversal, return true
- Otherwise, return false

# Connected

Consider implementing a breadth-first traversal on an undirected graph:

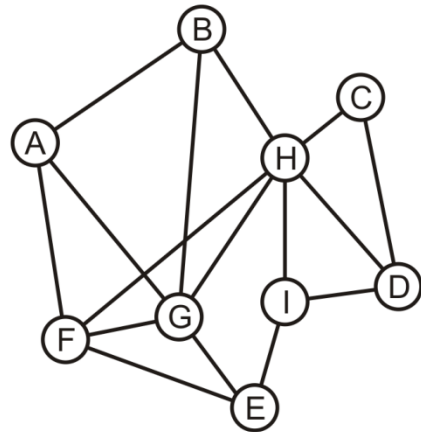
- Choose any vertex, mark it as visited and push it onto queue
- While the queue is not empty:
  - Pop to top vertex  $v$  from the queue
  - For each vertex adjacent to  $v$  that has not been visited:
    - Mark it visited, and
    - Push it onto the queue

This continues until the queue is empty

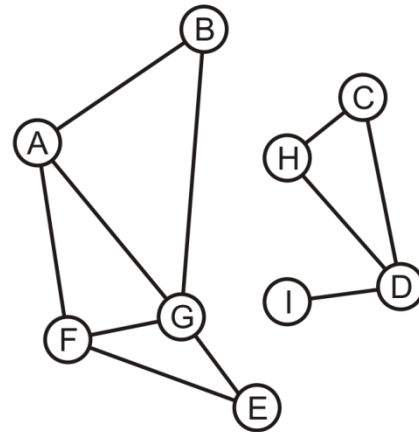
- Note: if there are no unvisited vertices, the graph is connected,

# Determining Connections

Is A connected to D?



--	--	--	--	--	--




--	--	--	--	--	--




# Determining Connections

On the right, the queue is empty and D is not visited

- We determine A is not connected to D



C	I	D			
---	---	---	--	--	--



--	--	--	--	--	--

# Connected Components

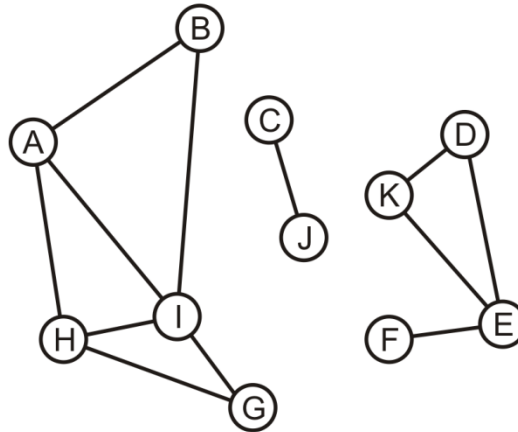
Suppose we want to partition the vertices into connected sub-graphs

- While there are unvisited vertices in the tree:
  - Select an unvisited vertex and perform a traversal on that vertex
  - Each vertex that is visited in that traversal is added to the set initially containing the initial unvisited vertex
- Continue until all vertices are visited

We would use a disjoint set data structure for maximum efficiency

# Connected Components


Here we start with a set of singletons



A	B	C	D	E	F	G	H	I	J	K
<b>A</b>	<b>B</b>	<b>C</b>	<b>D</b>	<b>E</b>	<b>F</b>	<b>G</b>	<b>H</b>	<b>I</b>	<b>J</b>	<b>K</b>

# Connected Components


The vertex A is unvisited, so we start with it



A	B	C	D	E	F	G	H	I	J	K
<b>A</b>	<b>B</b>	<b>C</b>	<b>D</b>	<b>E</b>	<b>F</b>	<b>G</b>	<b>H</b>	<b>I</b>	<b>J</b>	<b>K</b>

# Connected Components

Take the union of with its adjacent vertices: {A, B, H, I}




A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	E	F	G	A	A	J	K

# Connected Components

As the traversal continues, we take the union of the set  $\{G\}$  with the set containing H:  $\{A, B, G, H, I\}$


- The traversal is finished



A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	E	F	A	A	A	J	K

# Connected Components

Start another traversal with C: this defines a new set { C }

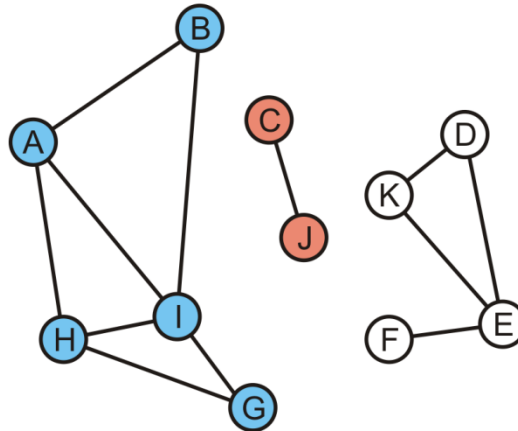


A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	E	F	A	A	A	J	K

# Connected Components

We take the union of  $\{C\}$  and its adjacent vertex J:  $\{C, J\}$

- This traversal is finished

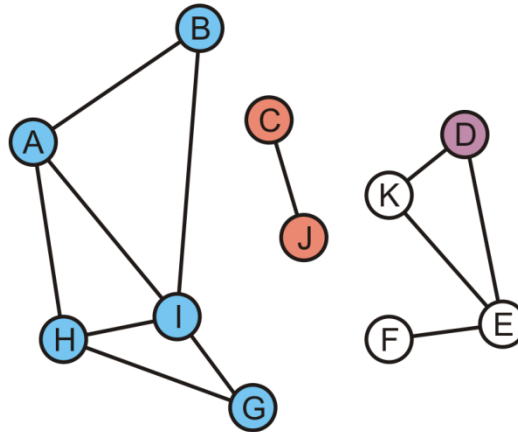


A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	E	F	A	A	A	C	K



# Connected Components

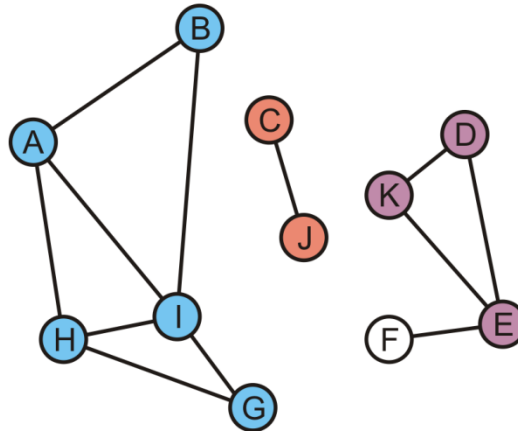
We start again with the set {D}



A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	E	F	A	A	A	C	K

# Connected Components

K and E are adjacent to D, so take the unions creating  $\{D, E, K\}$

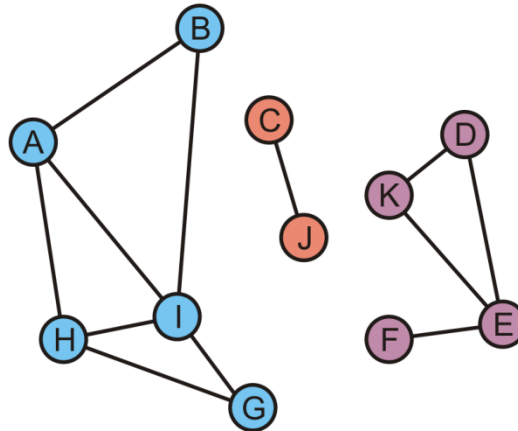


A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	D	F	A	A	A	C	D

# Connected Components

Finally, during this last traversal we find that F is adjacent to E

- Take the union of  $\{F\}$  with the set containing E:  $\{D, E, F, K\}$

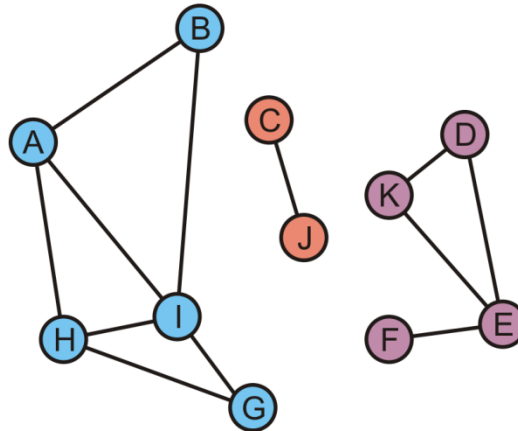


A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	D	D	A	A	A	C	D

# Connected Components

All vertices are visited, so we are done

- There are three connected sub-graphs {A, B, G, H, I}, {C, J}, {D, E, F, K}



A	B	C	D	E	F	G	H	I	J	K
A	A	C	D	D	D	A	A	A	C	D

# Tracking Unvisited Vertices

How do you implement a set of unvisited vertices so as to:

- Find an unvisited vertex in  $\Theta(1)$  time?
- Remove a vertex that has been visited from this list in  $\Theta(1)$  time?

Bad solution

- We can simply flag vertices as visited, but this would require  $O(|V|)$  time to find an unvisited vertex

Good solutions

- A hash table of unvisited vertices
- Or, an array of unvisited vertices, and we store for each vertex its position in the array

# Tracking Unvisited Vertices

Create two arrays:

- One array, `unvisited`, will contain the unvisited vertices
- The other, `loc_in_unvisited`, will contain the location of vertex  $v_i$  in the first array

0	1	2	3	4	5	6	7	8	9	10
A	B	C	D	E	F	G	H	I	J	K

A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	9	10

- Or, instead of a second array, we may add a member variable in the vertex class

# Tracking Unvisited Vertices

Suppose we visit D

- D is in entry 3
- How shall we delete D in the first array?

0	1	2	3	4	5	6	7	8	9	10
A	B	C	D	E	F	G	H	I	J	K


A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	9	10

# Tracking Unvisited Vertices

Suppose we visit D

- D is in entry 3
- Copy the last unvisited vertex into this location and update the location array for this value

0	1	2	3	4	5	6	7	8	9	10
A	B	C	K	E	F	G	H	I	J	



A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	9	3



# Tracking Unvisited Vertices

Suppose we visit G

- G is in entry 6

0	1	2	3	4	5	6	7	8	9	10
A	B	C	K	E	F	G	H	I	J	


A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	9	3

# Tracking Unvisited Vertices

Suppose we visit G

- G is in entry 6
- Copy the last unvisited vertex into this location and update the location array for this value

0	1	2	3	4	5	6	7	8	9	10
A	B	C	K	E	F	J	H	I		



A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	6	3

# Tracking Unvisited Vertices

Suppose we now visit K

- K is in entry 3

0	1	2	3	4	5	6	7	8	9	10
A	B	C	K	E	F	J	H	I		


A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	8	6	3

# Tracking Unvisited Vertices

Suppose we now visit K

- K is in entry 3
- Copy the last unvisited vertex into this location and update the location array for this value

0	1	2	3	4	5	6	7	8	9	10
A	B	C	I	E	F	J	H			



A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	3	6	3

# Tracking Unvisited Vertices

If we want to find an unvisited vertex, we simply return the last entry of the first array and return it

0	1	2	3	4	5	6	7	8	9	10
A	B	C	I	E	F	J	H			

A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	3	6	3

# Tracking Unvisited Vertices

In this case, an unvisited vertex is H

- Removing it is trivial: just decrement the count of unvisited vertices

0	1	2	3	4	5	6	7	8	9	10
A	B	C	I	E	F	J				

A	B	C	D	E	F	G	H	I	J	K
0	1	2	3	4	5	6	7	3	6	3

# Outline

- Graph traversal
  - Breadth-first
  - Depth-first
- Applications
  - Connectedness
  - Unweighted path length
  - Identifying bipartite graphs

# Determining Distances

Problem: in an unweighted graph, find the distances from one vertex  $v$  to all the other vertices

- Distance: the length of the shortest path between two vertices

Method:

- Use a breadth-first traversal
- Vertices are added in *layers*
- The starting vertex  $v$  is defined to be in the zeroth layer,  $L_0$
- While the  $k^{\text{th}}$  layer is not empty, all unvisited vertices adjacent to vertices in  $L_k$  are added to the  $(k + 1)^{\text{st}}$  layer


The distance from  $v$  to vertices in  $L_k$  is  $k$

Any unvisited vertices are said to have an infinite distance from  $v$




# Determining Distances

Consider this graph: find the distance from A to each other vertex



# Determining Distances


A forms the zeroeth layer,  $L_0$



# Determining Distances

The unvisited vertices B, F and G are adjacent to A


- These form the first layer,  $L_1$



# Determining Distances

We now begin popping  $L_1$  vertices: pop B


- H is adjacent to B
- It is tagged  $L_2$



# Determining Distances

Popping F pushes E onto the queue


- It is also tagged  $L_2$



# Determining Distances


We pop G which has no other unvisited neighbours

- G is the last  $L_1$  vertex; thus H and E form the second layer,  $L_2$



# Determining Distances


Popping H in  $L_2$  adds C and I to the third layer  $L_3$



# Determining Distances

E has no more adjacent unvisited vertices

- Thus C and I form the third layer,  $L_3$






# Determining Distances

The unvisited vertex D is adjacent to vertices in  $L_3$

- This vertex forms the fourth layer,  $L_4$



- Distance 1: B, F, G
- Distance 2: H, E
- Distance 3: C, I
- Distance 4: D

# Determining Distances

Theorem:

- If, in a breadth-first traversal of a graph, two vertices  $v$  and  $w$  appear in layers  $L_i$  and  $L_j$ , respectively and  $\{v, w\}$  is an edge in the graph, then  $i$  and  $j$  differ by at most one

Proof:

If  $i = j$ , we are done

If  $i \neq j$ , without loss of generality, assume  $i < j$

Because  $v \in L_i$ ,  $w$  does not appear in any previous layer, and  $\{v, w\}$  is an edge in the graph, it follows that  $w \in L_{i+1}$

Thus,  $j = i + 1$

Therefore,  $i$  and  $j$  differ by at most one

# Outline

- Graph traversal
  - Breadth-first
  - Depth-first
- Applications
  - Connectedness
  - Unweighted path length
  - Identifying bipartite graphs

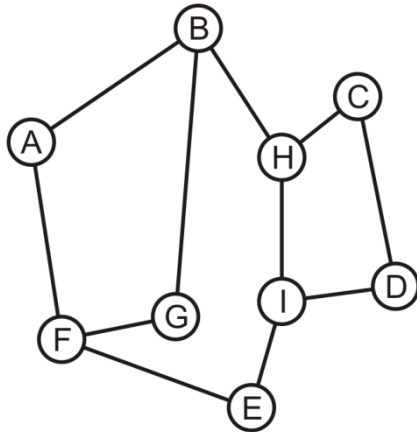
# Definition

## Definition

- A *bipartite graph* is a graph where the vertices  $V$  can be divided into two disjoint sets  $V_1$  and  $V_2$  such that **every** edge has one vertex in  $V_1$  and the other in  $V_2$



# Bipartite Graphs

Consider this graph: is it bipartite?




# Bipartite Graphs

Yes: With a little work, it is possible to determine that we can decompose the vertices into two disjoint sets




# Bipartite Graphs

Is this graph bipartite?



# Bipartite Graphs

In this case, it is not a bipartite graph



How can we determine if a graph is bipartite?



# Bipartite Graphs


Use a breadth-first traversal for a connected graph:

- Choose a vertex, mark it belonging to  $V_1$  and push it onto a queue
- While the queue is not empty, pop the front vertex  $v$  and
  - Any adjacent vertices that are already marked must belong to the set not containing  $v$ , otherwise, the graph is not bipartite (we are done);
  - Any unmarked adjacent vertices are marked as belonging to the other set and they are pushed onto the queue
- If the queue is empty, the graph is bipartite

# Bipartite Graphs


With the first graph, we can start with any vertex

- We will use colours to distinguish the two sets



# Bipartite Graphs


Push A onto the queue and colour it red



# Bipartite Graphs

Pop A and its two neighbours are not marked:


- Mark them as blue and push them onto the queue



# Bipartite Graphs

Pop B—it is blue:


- Its one marked neighbour, A, is red
- Its other neighbours G and H are not marked: mark them red and push them onto the queue



# Bipartite Graphs

Pop F—it is blue:


- Its two marked neighbours, A and G, are red
- Its neighbour E is not marked: mark it red and pus it onto the queue



# Bipartite Graphs

Pop G—it is red:


- Its two marked neighbours, B and F, are blue



# Bipartite Graphs

Pop H—it is red:

- Its marked neighbours, B, is blue
- It has two unmarked neighbours, C and I; mark them blue and push them onto the queue






# Bipartite Graphs

Pop E—it is red:


- Its marked neighbours, F and I, are blue



# Bipartite Graphs

Pop C—it is blue:


- Its marked neighbour, H, is red
- Mark D as red and push it onto the queue



# Bipartite Graphs

Pop I—it is blue:


- Its marked neighbours, H, D and E, are all red



# Bipartite Graphs


Pop D—it is red:

- Its marked neighbours, C and I, are both blue




# Bipartite Graphs

The queue is empty, the graph is bipartite




# Bipartite Graphs

Consider the other graph which was claimed to be not bipartite



# Bipartite Graphs


Push A onto the queue and colour it red



# Bipartite Graphs

Pop A off the queue:

- Its neighbours are unmarked: colour them blue and push them onto the queue






# Bipartite Graphs

Pop B off the queue:


- Its one neighbour, A, is red
- The other neighbour, H, is unmarked: colour it red and push it onto the queue



# Bipartite Graphs

Pop F off the queue:


- Its one neighbour, A, is red
- The other neighbours, E and G, are unmarked: colour them red and push it onto the queue



# Bipartite Graphs

Pop H off the queue—it is red:

- Its one neighbour, G, is already red
- The graph is not bipartite



# Bipartite Graphs

## Definition

Cycles that contains either an even number or an odd number of vertices are said to be *even cycles* and *odd cycles*, respectively

## Theorem

A graph is bipartite if and only if it does not contain any odd cycles

# Outline

- Graph traversal
  - Breadth-first: use a queue
  - Depth-first: use recursion or stack
- Applications
  - Connectedness
  - Unweighted path length
  - Identifying bipartite graphs