



CS120: Computer Networks

Lecture 10. Routing 1

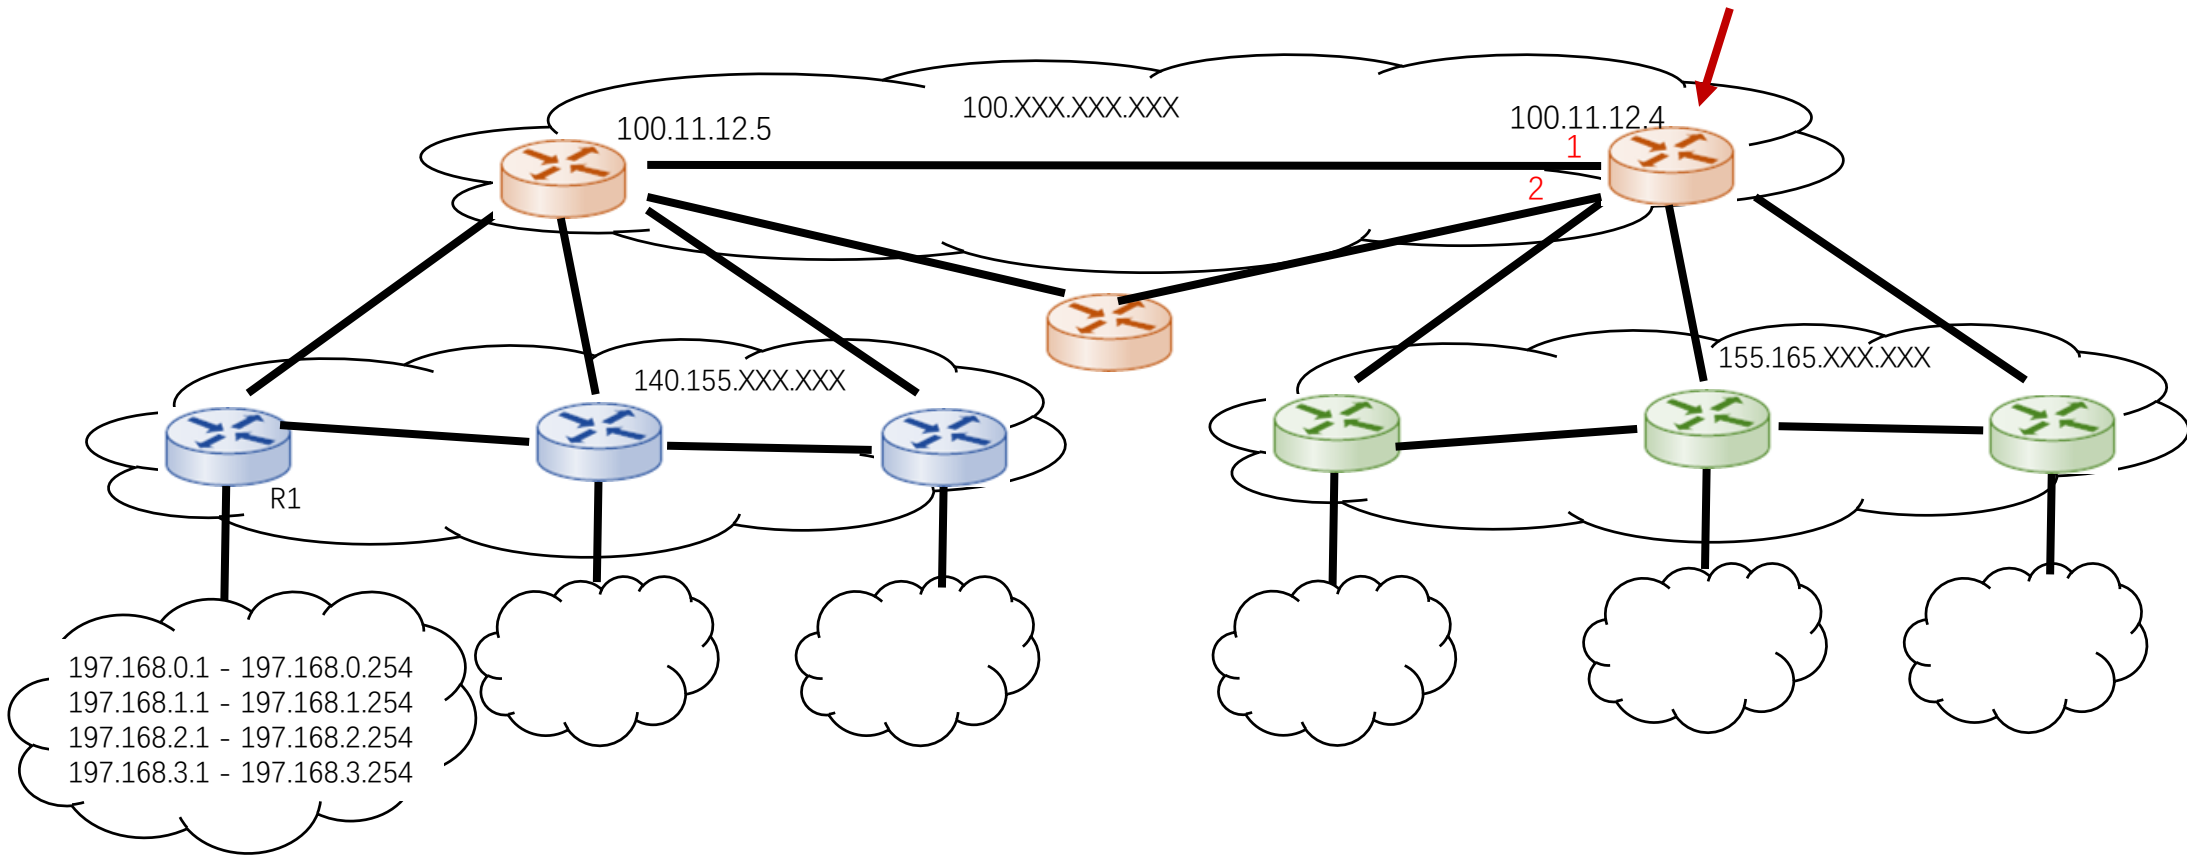
Zhice Yang

Routing Table

SubnetNum	NextHop
197.168.0.0/22	100.11.12.5

Forwarding Table

destaddress	Interface	MAC
100.11.12.5	1	AB.CD.EF.12.34.56




Forwarding Table vs. Routing Table

- Forwarding table
 - Determines local forwarding
 - Optimized for looking up an address when forwarding a packet
 - Normally in hardware
 - Contains the mapping from network numbers to outgoing interfaces and some MAC addresses
- Routing table
 - Built by the routing algorithm as a precursor to build the forwarding table
 - Optimized for calculating changes in network topology
 - Normally in software
 - Contains mapping from network numbers to next hops

Network as a Graph

- The basic problem of routing is to find the **lowest-cost** path between any two nodes
 - Static approach has several shortcomings
 - Can't handle node or link failures
 - Can't handle addition of new nodes or links
 - Edge costs cannot change
 - Centralized solution does not scale
 - Distributed and dynamic protocol



Routing Protocols

- Routing Information Protocol (RIP)
 - Algorithm: Distance Vector
- Open Shortest Path First (OSPF)
 - Algorithm: Link State
- Border Gateway Protocol (BGP)



Intradomain Routing Protocol

Interdomain Routing Protocol

Distance Vector Algorithm

- Bellman-Ford equation

let

$d_x(y)$ = cost of lowest-cost path from x to y

then

$$d_x(y) = \min_v \{c(x, v) + d_v(y)\}$$


min taken over all neighbors v of x

cost to neighbor v

lowest-cost from neighbor v to destination y

Example

- $d_B(A) = 2$
- $d_D(A) = 2$
- $d_C(A) = \min(d_B(A) + 4, d_D(A) + 9) = 6$



Distance Vector Algorithm

- x maintains its distance vector estimate $D_x(y) = \{D_x(y) : y \in N\}$
- x knows:
 - cost to each neighbor v: $c(x, v)$
 - neighbors' distance vectors estimate: $D_v(y) = \{D_v(y) : y \in N\}$
- Algorithm idea:
 - From time-to-time, each node sends its own distance vector estimate to neighbors
 - When x receives new distance vector estimate from neighbor, it updates its own distance vector estimate using Bellman-Ford equation
 - Under minor, natural conditions, the estimate $D_x(y)$ will converge to the actual lowest cost $d_x(y)$

Distance Vector Algorithm

y	$D_A(y)$
A	0
B	inf
C	inf
D	inf
E	inf
F	inf
G	inf

y	$D_B(y)$
A	inf
B	0
C	inf
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	inf
B	inf
C	0
D	inf
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	inf
D	0
E	inf
F	inf
G	inf

y	$D_E(y)$
A	inf
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	inf
B	inf
C	inf
D	inf
E	inf
F	0
G	inf

y	$D_G(y)$
A	inf
B	inf
C	inf
D	inf
E	inf
F	inf
G	0



Distance Vector Algorithm

y	$D_A(y)$
A	0
B	1
C	1
D	inf
E	1
F	1
G	inf

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1

y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm

- Every T seconds each router sends its table to its neighbor
- Each router then updates its table based on the new information

Distance Vector Algorithm

y	$D_A(y)$
A	0
B	1
C	1
D	inf
E	1
F	1
G	inf

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1

y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm

$$\swarrow$$

y	$D_A(y)$
A	0
B	1
C	1
D	inf
E	1
F	1
G	inf

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf

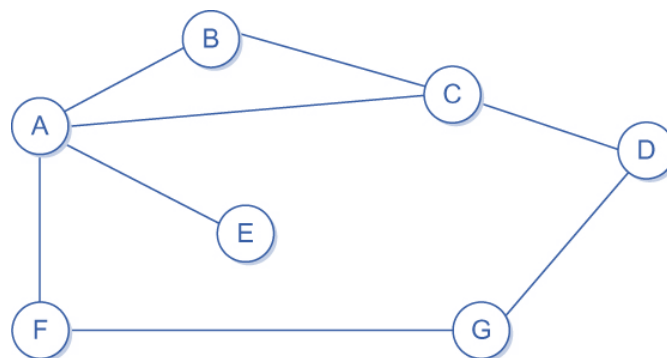
y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1

y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm

y	$D_A(y)$
A	0
B	1
C	1
D	2
E	1
F	1
G	inf

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1

y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm

y	$D_A(y)$
A	0
B	1
C	1
D	2
E	1
F	1
G	inf

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1

y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm

y	$D_A(y)$
A	0
B	1
C	1
D	2
E	1
F	1
G	2

y	$D_B(y)$
A	1
B	0
C	1
D	inf
E	inf
F	inf
G	inf


y	$D_C(y)$
A	1
B	1
C	0
D	1
E	inf
F	inf
G	inf

y	$D_D(y)$
A	inf
B	inf
C	1
D	0
E	inf
F	inf
G	1


y	$D_E(y)$
A	1
B	inf
C	inf
D	inf
E	0
F	inf
G	inf

y	$D_F(y)$
A	1
B	inf
C	inf
D	inf
E	inf
F	0
G	1

y	$D_G(y)$
A	inf
B	inf
C	inf
D	1
E	inf
F	1
G	0



Distance Vector Algorithm




Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0

y	$D_A(y)$	via
A	0	A
B	1	B
C	1	C
D	2	C
E	1	E
F	1	F
G	2	F

Distance Vector Algorithm

- Good news travels fast


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Distance Vector Algorithm

- Good news travels fast


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	1
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	1	1	3	1	0



Distance Vector Algorithm

- Good News Travels Fast


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	2
C	1	1	0	1	2	2	1
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	2	1	1	3	1	0



Distance Vector Algorithm

- Good News Travels Fast


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	2
C	1	1	0	1	2	2	1
D	2	2	1	0	3	2	1
E	1	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	2	1	1	3	1	0



Distance Vector Algorithm

- Bad News Travels Slow


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	inf	2	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Distance Vector Algorithm

- Bad News Travels Slow


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	inf	inf	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Distance Vector Algorithm

- Bad News Travels Slow


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	inf	3	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Distance Vector Algorithm

- Bad News Travels Slow


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	4	3	2	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Distance Vector Algorithm

- Bad News Travels Slow


Information Stored at Node	Distance to Reach Node						
	A	B	C	D	E	F	G
A	0	1	1	2	1	1	2
B	1	0	1	2	2	2	3
C	1	1	0	1	2	2	2
D	2	2	1	0	3	2	1
E	4	3	5	3	0	2	3
F	1	2	2	2	2	0	1
G	2	3	2	1	3	1	0



Count-to-infinity Problem

Routing Information Protocol (RIP)

- Included in BSD-UNIX distribution in 1982
- Use distance vector algorithm
 - Distance metric: # hops (max = 15 hops), each link has cost 1
 - Distance Vectors exchanged with neighbors every 30 sec in response message
 - Each message: list of up to 25 destination subnets



Routing Table A

SubnetNum	Distance	NextHop
1	0	Net1
2	0	Net2
3	1	C
4	3	C
5	2	C
6	2	C

Routing Protocols

- Routing Information Protocol (RIP)
 - Algorithm: Distance Vector
- Open Shortest Path First (OSPF)
 - Algorithm: Link State
- Border Gateway Protocol (BGP)



Intradomain Routing Protocol


Interdomain Routing Protocol

Link-State Routing Algorithm

- Assumptions
 - Net topology, link costs known to all nodes
 - Accomplished via “Reliable Flooding”
 - Send to all nodes (not just neighbors) information about directly connected links (not the entire routing table).
 - Link state packet (LSP)
 - All nodes have same info
- Routing Method: Computes shortest paths from one node ('source') to all other nodes
 - Based on Dijkstra's Algorithm

Reliable Flooding

- Designs
 - Keep LSP up to date
 - Generate new LSP periodically
 - on the order of hours
 - Generate new LSP when link states change
 - Abandon old link state information
 - Differentiates new LSP according to seq number
 - Decreases TTL before flooding
 - Ages stored LSP
 - Limiting the flooding overhead
 - Forward LSP to all nodes but one that sent it



Dijkstra's Algorithm

Initialization:

$M = \{s\}$

for all nodes v

if v adjacent to s

then $D_s(v) = c(u, v)$

else $D_s(v) = \text{inf}$

Loop

find w not in M such that $D_s(w)$ is a minimum

add w to M


update $D_s(v)$ for all v adjacent to w and not in M :

$D_s(v) = \min(D_s(v), D_s(w) + c(w, v))$

until all nodes in M


- M : set of node processed
- S : node of the local router
- v : node of other routers
- $D_s(v)$ distance from s to v
- $c(u, v)$ link weight between node u and v

Dijkstra's Algorithm




M	$D_D(A)$	$D_D(B)$	$D_D(C)$
{D}	Inf, from D	11, from D	2, from D

Dijkstra's Algorithm




M	$D_D(A)$	$D_D(B)$	$D_D(C)$
{D}	Inf, from D	11, from D	2, from D
{D, C}	12, from C	5, from C	2, from D

Dijkstra's Algorithm



M	$D_D(A)$	$D_D(B)$	$D_D(C)$
{D}	Inf, from D	11, from D	2, from D
{D, C}	12, from C	5, from C	2, from D
{D, C, B}	10, from B	5, from C	2, from D

Dijkstra's Algorithm




M	$D_D(A)$	$D_D(B)$	$D_D(C)$
{D}	Inf, from D	11, from D	2, from D
{D, C}	12, from C	5, from C	2, from D
{D, C, B}	10, from B	5, from C	2, from D

Dijkstra's Algorithm (Another notation)

- $\langle \text{Destination, Cost, Nexthop} \rangle$

Step	Confirmed	Tentative
1	(D,0,-)	
2	(D,0,-)	(B,11,B) (C,2,C)
3	(D,0,-) (C,2,C)	(B,11,B)
4	(D,0,-) (C,2,C)	(B,5,C) (A,12,C)
5	(D,0,-) (C,2,C) (B,5,C)	(A,12,C)
6	(D,0,-) (C,2,C) (B,5,C)	(A,10,C)
7	(D,0,-) (C,2,C) (B,5,C) (A,10,C)	



Open Shortest Path First (OSPF)

- “Open”: nonproprietary standard created under Engineering Task Force (IETF).
- Security: all OSPF messages authenticated (to prevent malicious intrusion)
- Hierarchical routing: OSPF in large domains
- Load balancing: multiple same-cost paths allowed (only one path in RIP)

Reference

- Textbook 3.3