



CS120: Computer Networks

Lecture 12. Other Topics in IP


Zhice Yang

Outline

- IPv6
- NAT
- MPLS
- Router Implementation

IPv6 Address

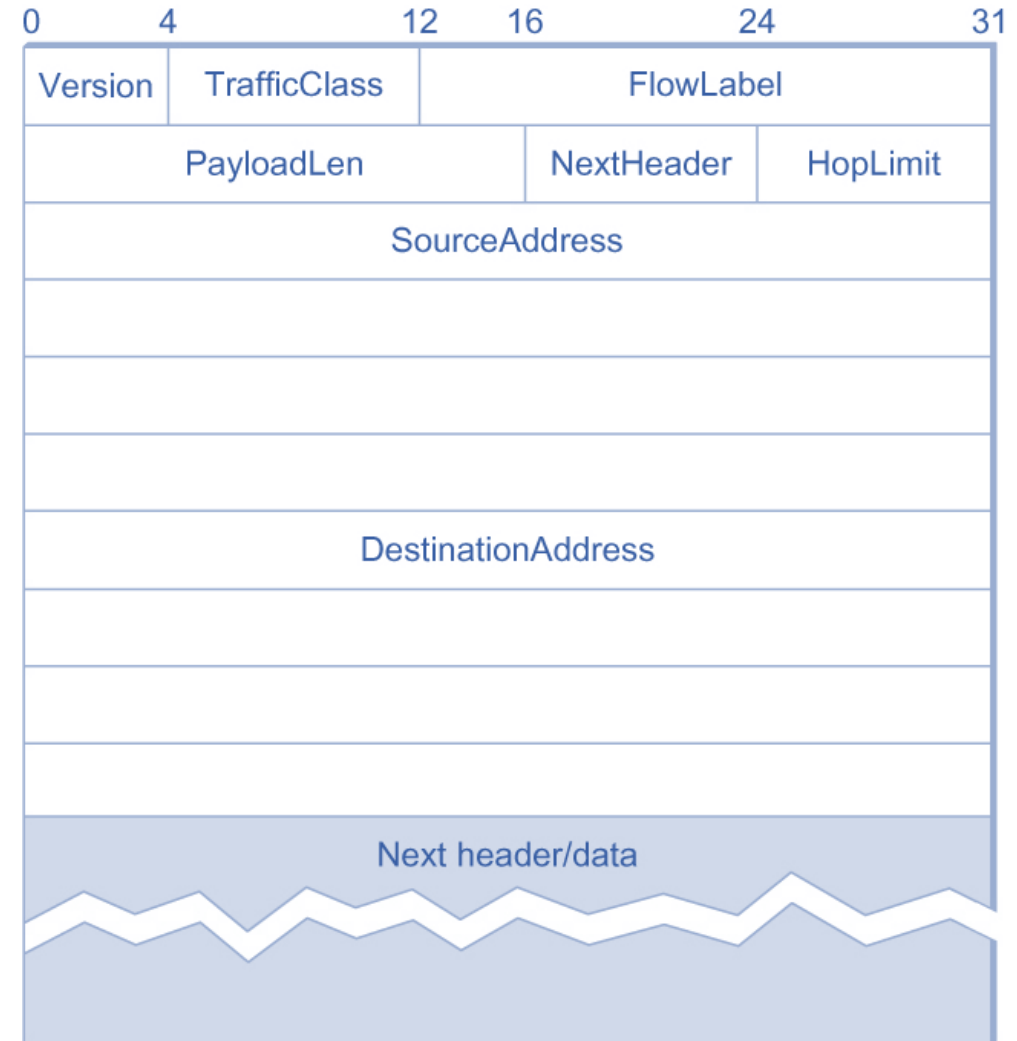
- 16 bytes
 - 1500 addresses per square foot (Earth's surface)
- Classless addressing/routing (similar to CIDR)
 - Notation: x:x:x:x:x:x:x:x (x = 16-bit hex number)
 - Contiguous 0s are compressed: 47CD::A456:0124
 - IPv4-mapped IPv6 address: ::FFFF:123.45.67.8
- Address assignment: more hierarchy



64bit

IP Version 6 (IPv6)

- Motivation
 - 32 bits IPv4 Address is not enough
 - Other Features
 - Stateless auto configuration
 - Source Routing
- IPv6 “base” Header
 - 40 bytes “base” header
 - 16 bytes addresses
- Extension headers
 - Fragmentation
 - Source routing
 - Authentication and security
 - etc.




Transition from IPv4 to IPv6

- IETF began looking at the problem of IPv4 address space in 1991
- Not all routers can be upgraded simultaneously

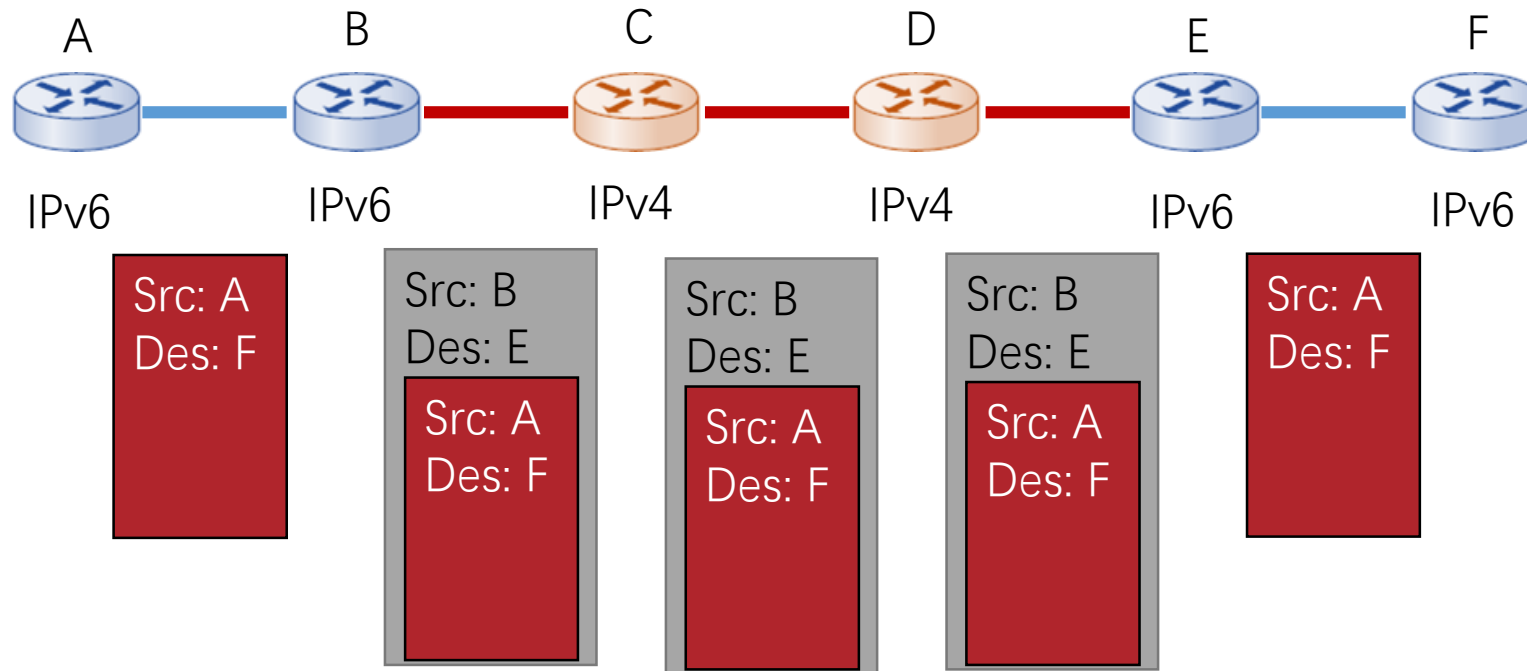
How will network operate with mixed IPv4 and IPv6 routers?

Transition from IPv4 to IPv6

- Tunneling: IPv6 datagram carried as payload in IPv4 datagram among IPv4 routers



Transition from IPv4 to IPv6



Why haven't IPv6 replaced IPv4 ?

NAT: Network Address Translation


Wireless LAN adapter Wi-Fi:

```

Connection-specific DNS Suffix  . :
Link-local IPv6 Address . . . . . : fe80::d1b5:35be:9832:af6c%9
IPv4 Address. . . . . : 192.168.31.143
Subnet Mask . . . . . : 255.255.255.0
Default Gateway . . . . . : 192.168.31.1
  
```

Address range	Subnet mask	Provides	Addresses per LAN
10.0.0.0 - 10.255.255.255.255	255.0.0.0	1 class A LAN	16,777,216
172.16.0.0 - 172.31.255.255	255.255.0.0	16 class B LANs	65,536
192.168.0.0 - 192.168.255.255	25.255.255.0	256 class C LANs	256


NAT: Network Address Translation



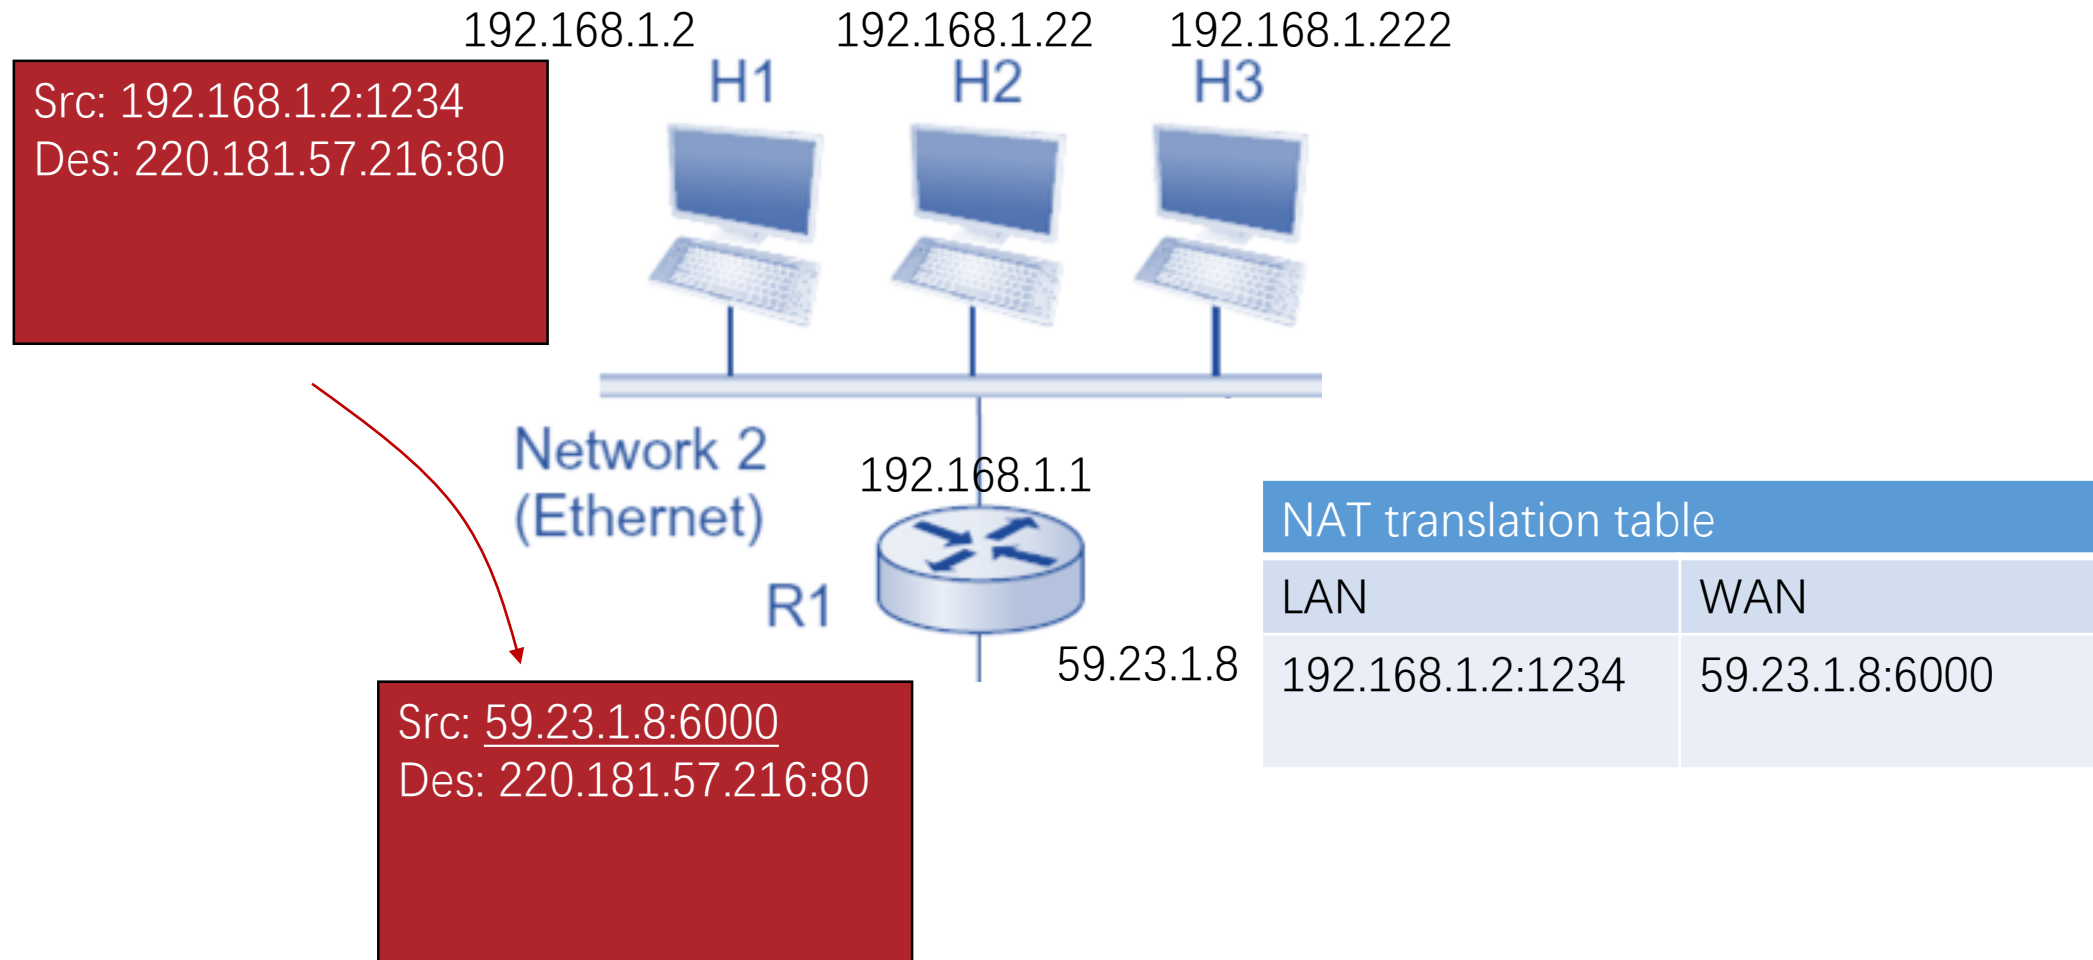
datagrams with source or destination in this network have 192.168.1.0/24 address for source, destination (as usual)

All datagrams leaving local network have same single source NAT IP address: 59.23.1.8, different source port numbers

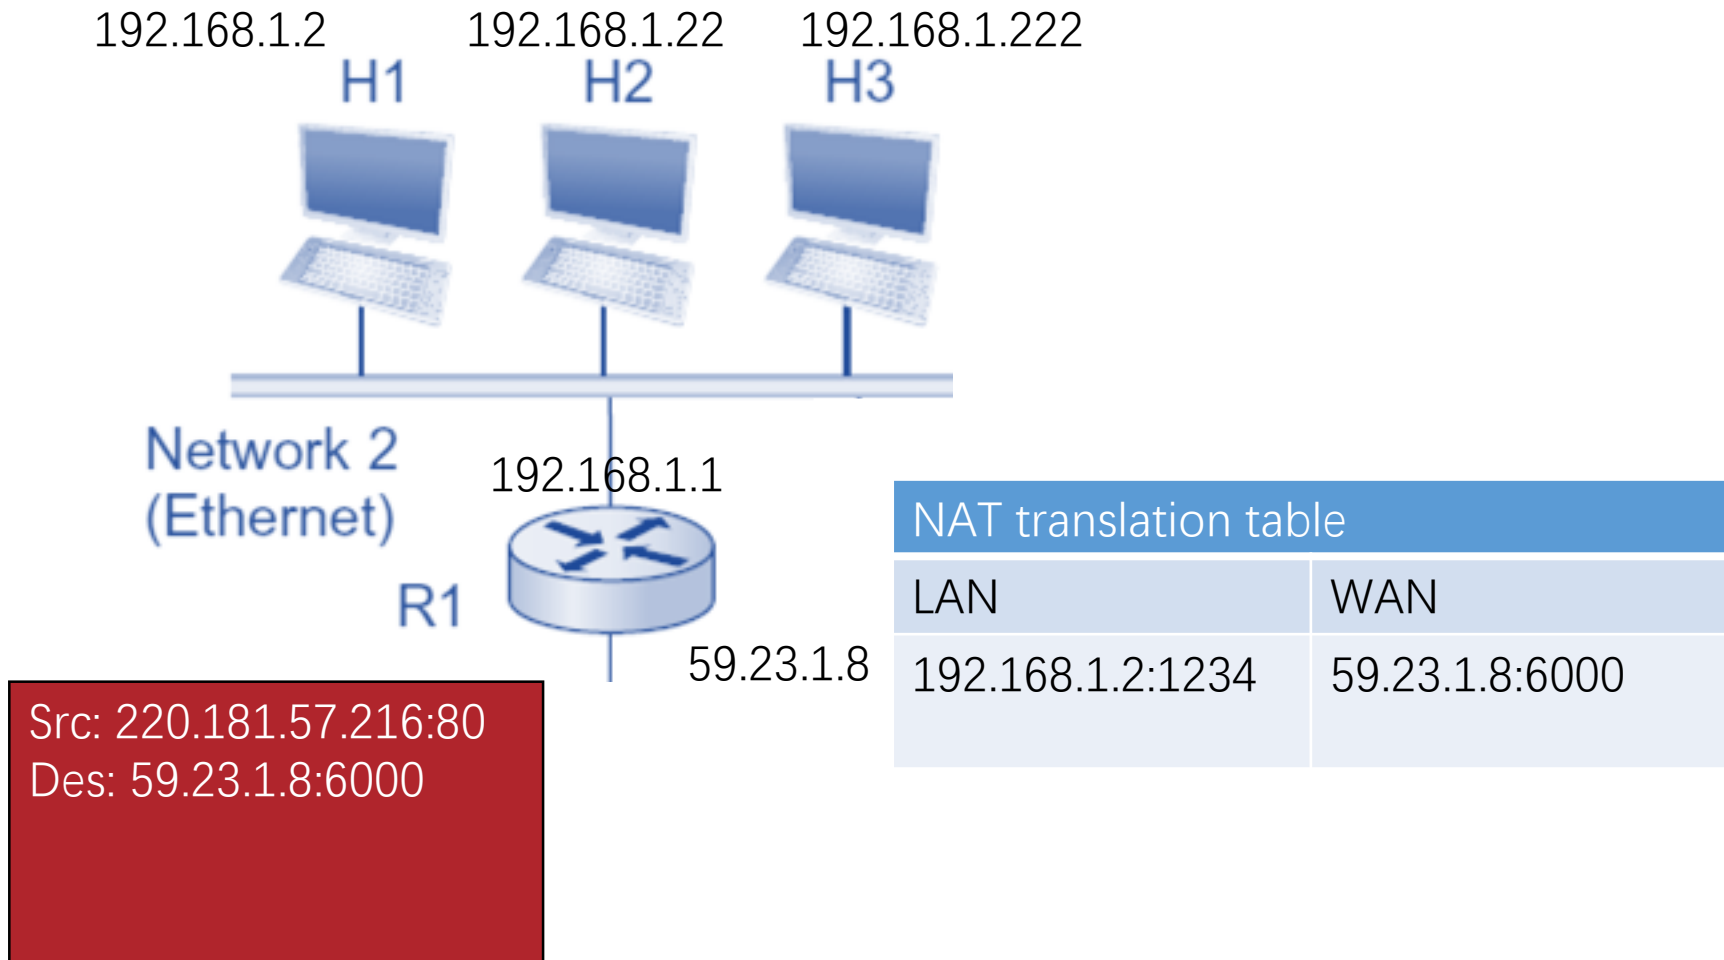
NAT: Network Address Translation



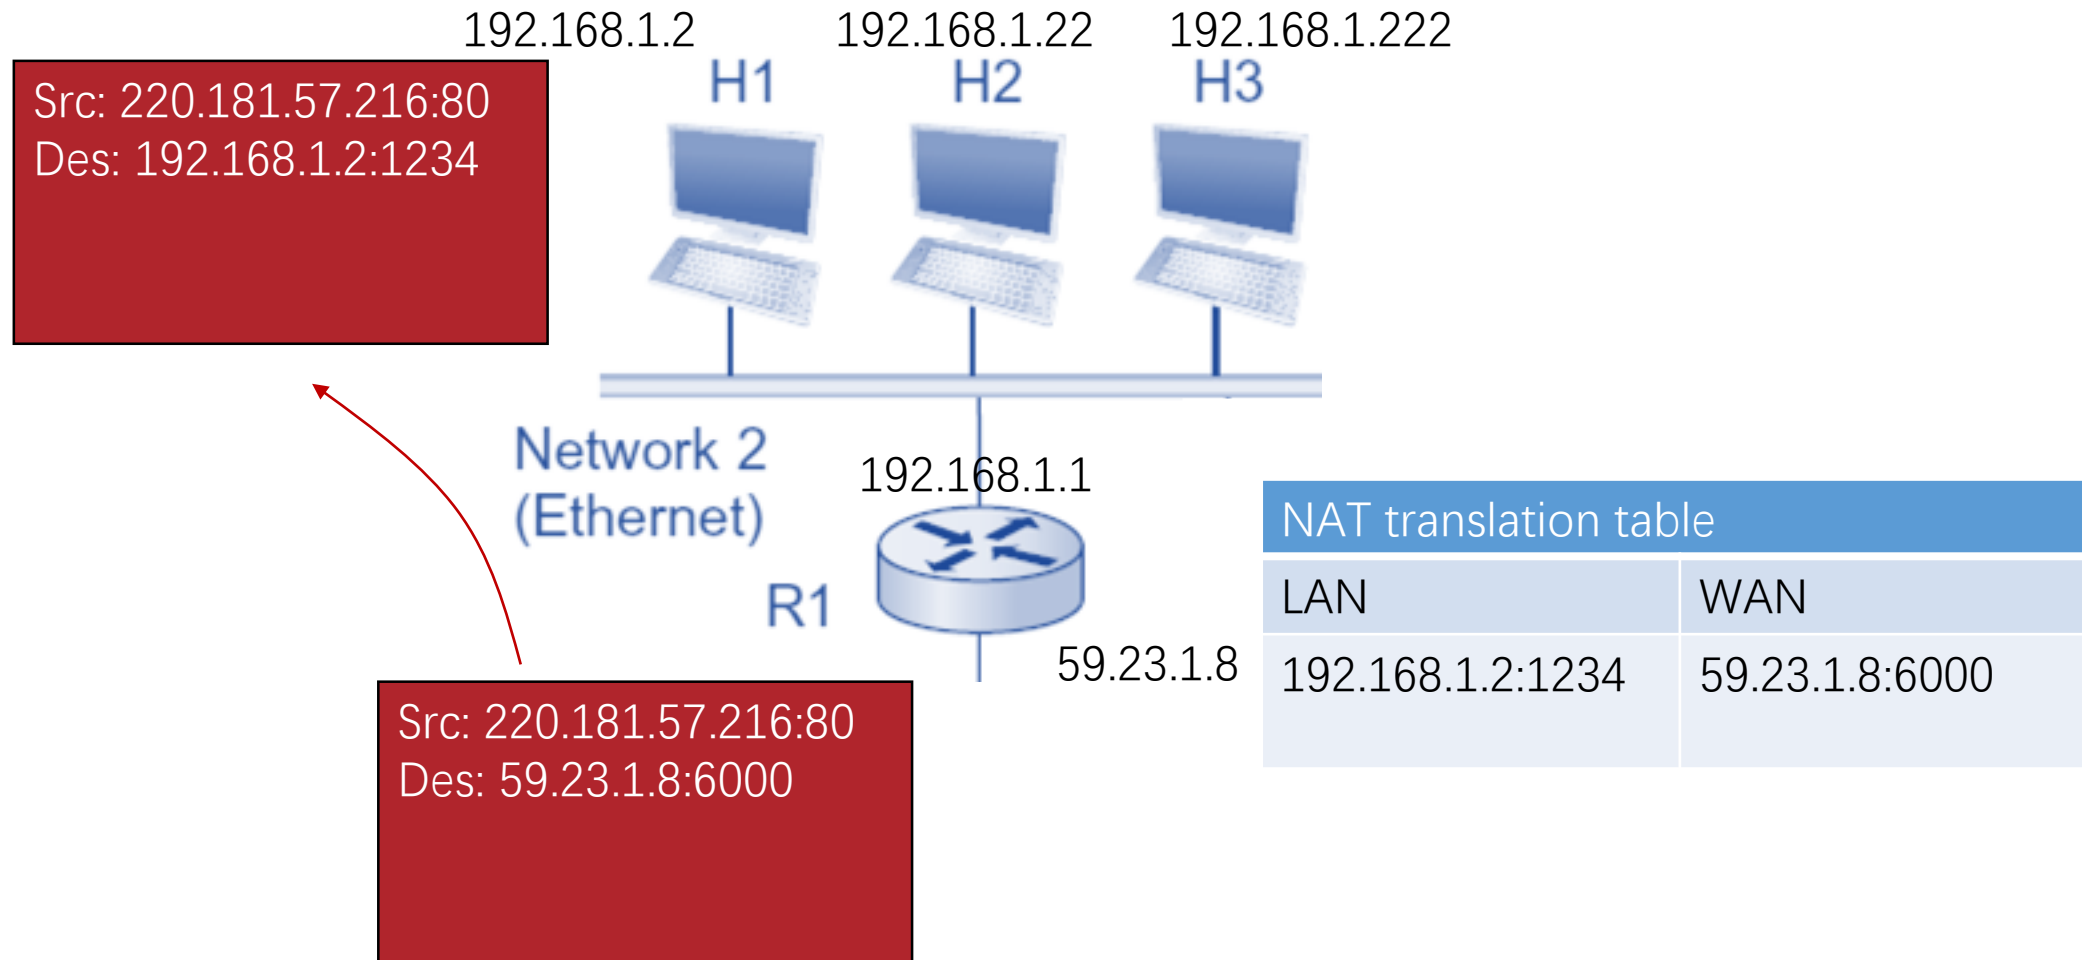
NAT: Network Address Translation



NAT: Network Address Translation




NAT: Network Address Translation



NAT: Network Address Translation

- Multi Layer NAT



NAT translation table R1

LAN	WAN
192.168.1.2:1234	10.1.1.8:6321

NAT translation table R2

LAN	WAN
10.1.1.8:6321	59.23.1.8:6000

NAT: Network Address Translation


- Demo

```
Tracing route to baidu.com [111.13.101.208]
over a maximum of 30 hops:

  1      5 ms      2 ms      <1 ms  XiaoQiang [192.168.31.1]
  2      2 ms      2 ms      3 ms   10.19.14.1
  3     11 ms      1 ms      2 ms   10.13.7.25
  4      *         *         *     Request timed out.
```



NAT: Network Address Translation

- 16-bit port-number field (see later lecture)
 - 60,000 simultaneous connections with a single IPv4 address
 - Hosts uses LAN addresses
 - 192.168.0.0/16
 - 10.0.0.0/8
 - 172.16.0.0/12
 - etc.
- Problems
 - NAT is “impure”
 - Routers should not touch higher layers
 - Efficiency
 - Traversal Connections



NAT Traversal Problem


- How to initiate connections to Host 1 from external network ?
 - e.g. IP Phone



NAT Traversal Problem


- Solution 1: Static Configure
 - Configure NAT to forward incoming connection requests at given port to the host
 - Need to know the port

NAT translation table	
LAN	WAN
192.168.1.2	59.23.1.8: <u>6000</u>



NAT Traversal Problem

- Solution 2: Relay
 - NATed host establishes connection to relay
 - External host connects to relay
 - Relay bridges packets between to connections



Outline

- IPv6
- NAT
- MPLS
- Router Implementation

Longest Prefix Matching


- When looking for forwarding table entry for given destination address, use longest address prefix that matches destination address.

SubnetNum	NextHop	
197.168.0.0/22	R4	11000101.10101000.000000**.*****
197.168.3.0/24	R7	11000101.10101000.00000011.*****
197.168.4.0/22	R9	11000101.10101000.000001**.*****

How to Accelerate Prefix Matching in Backbone Routers ?

11000101.10101000.00000111.11010111 → R9

Virtual Circuit



Virtual Circuit Table


Switch1			
Incoming Interface	Incoming VCI	Outgoing Interface	Incoming VCI
2	5	1	11

Switch2			
Incoming Interface	Incoming VCI	Outgoing Interface	Incoming VCI
3	11	0	7

Switch3			
Incoming Interface	Incoming VCI	Outgoing Interface	Incoming VCI
0	7	3	4

Asynchronous Transfer Mode (ATM)

- Evolved from telephony
 - Connection-oriented packet-switched network
 - Guaranteed Service: strict timing, reliability, etc.
- Packets are called cells



ATM Switch

- Packet Length is Fixed
 - Easier to switch in hardware
- VCI Length is Fixed
 - VCI looking up complexity is $O(1)$

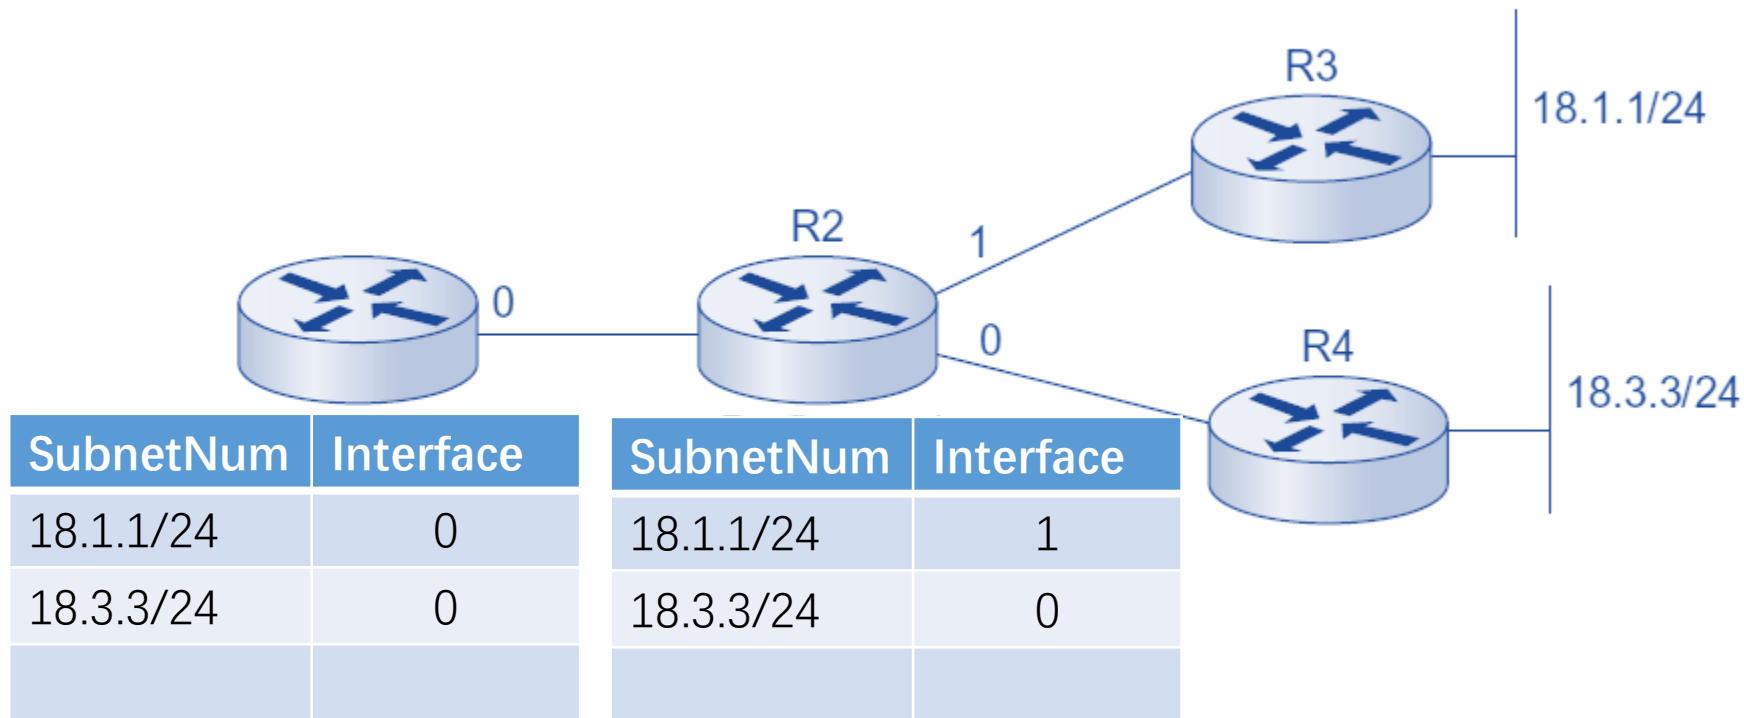


Multiprotocol Label Switching (MPLS)

- Original Motivation
 - Enable IP forwarding on devices that do not have IP forwarding ability (i.e., ATM Switches)
 - Capability (Support IP in ATM)
 - Speed
- Benefit: Destination-Based Forwarding

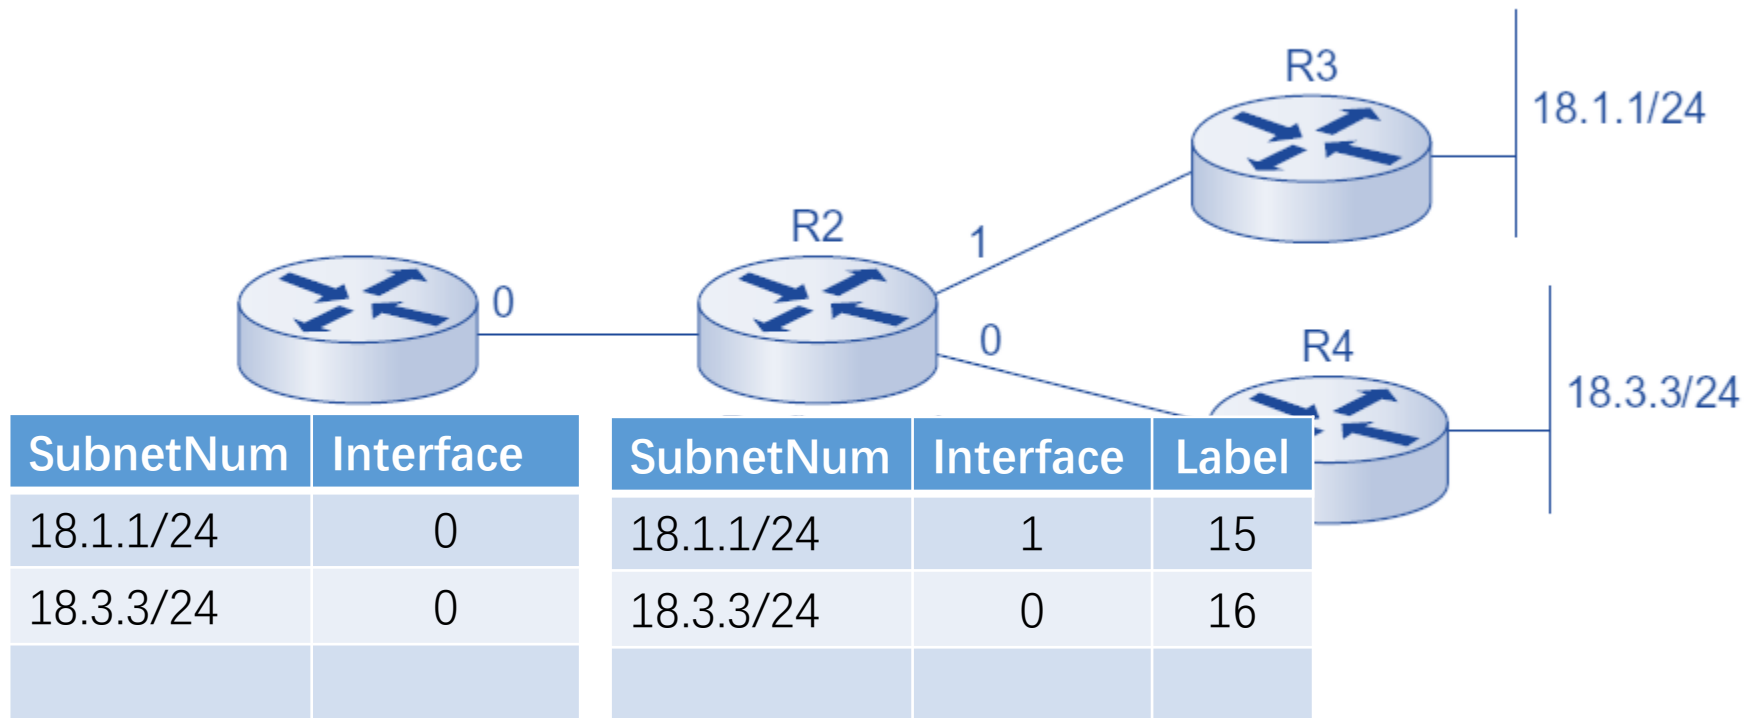
MPLS: Destination-Based Forwarding

- IP Forwarding with Normal Routers



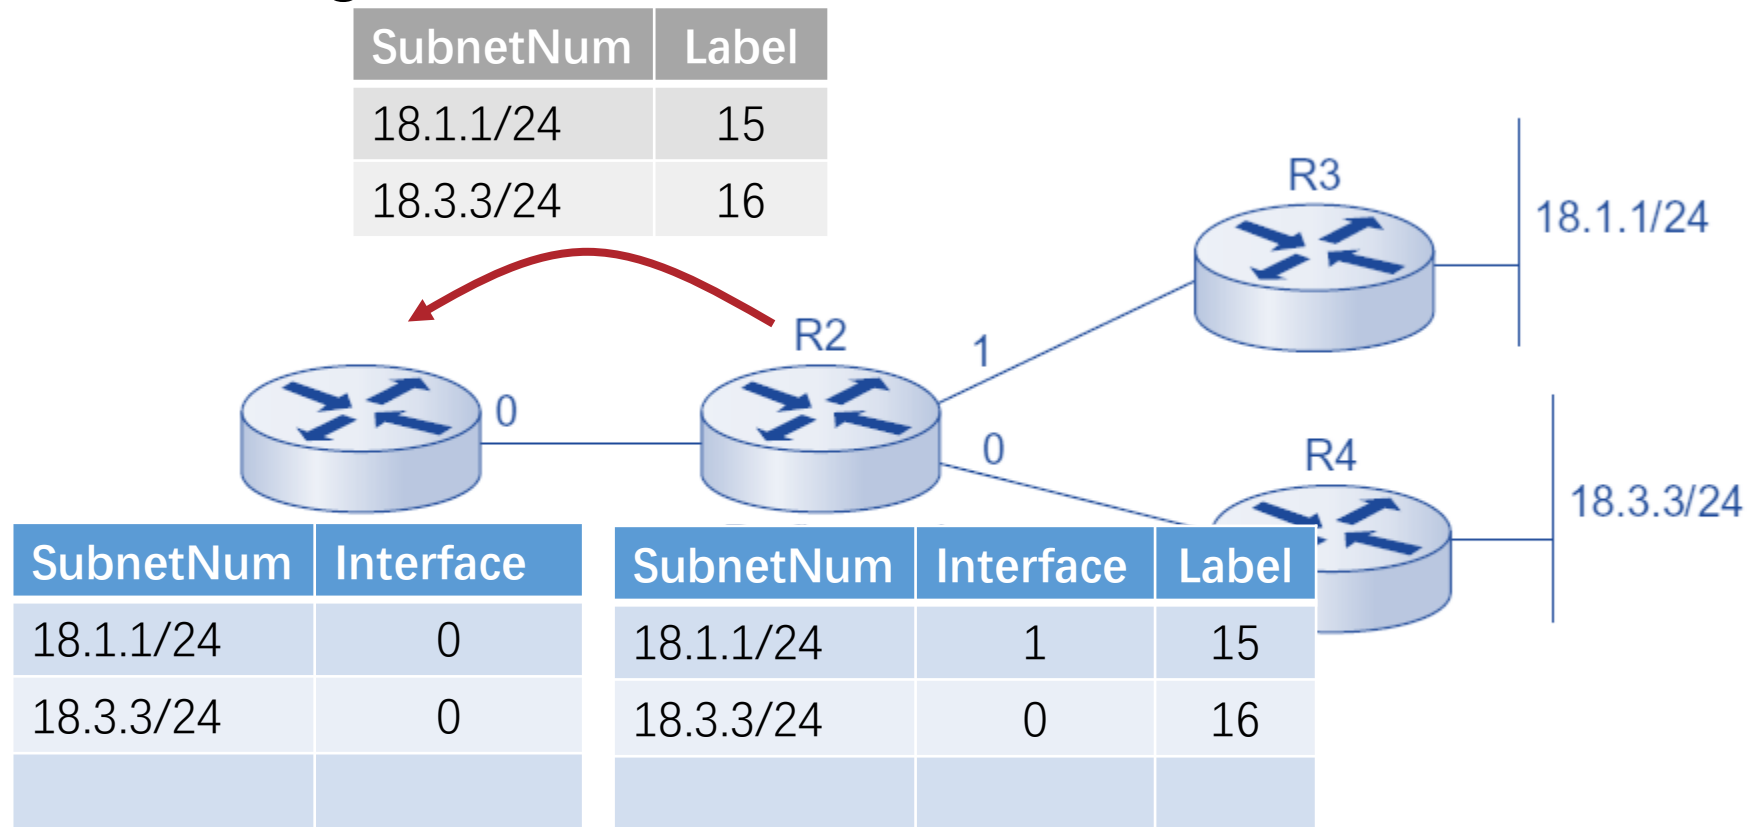
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels



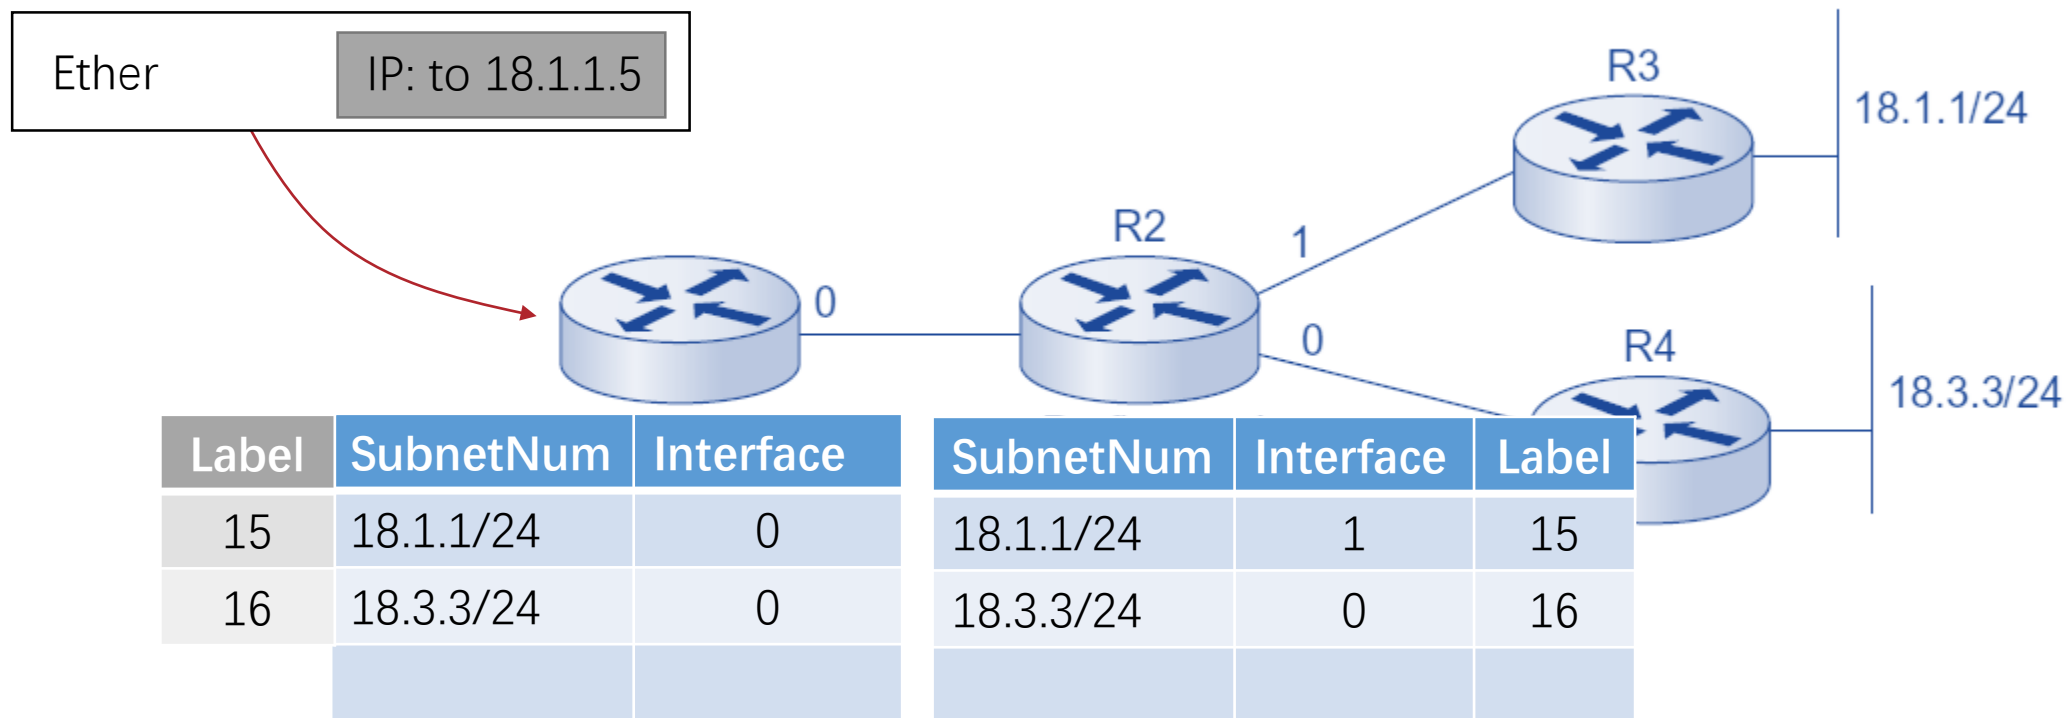
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels



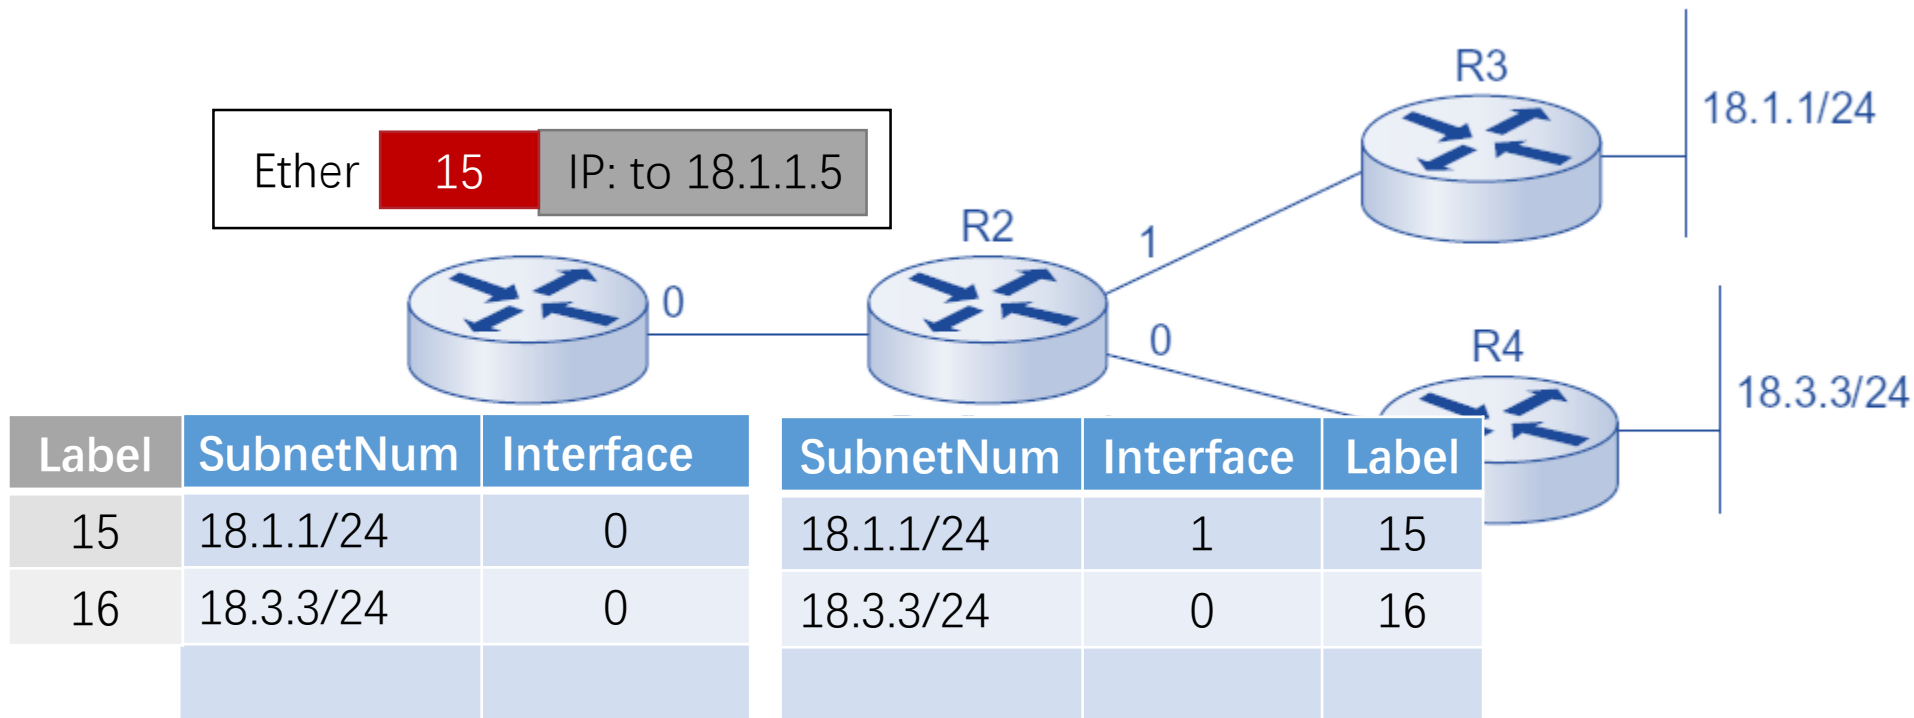
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels



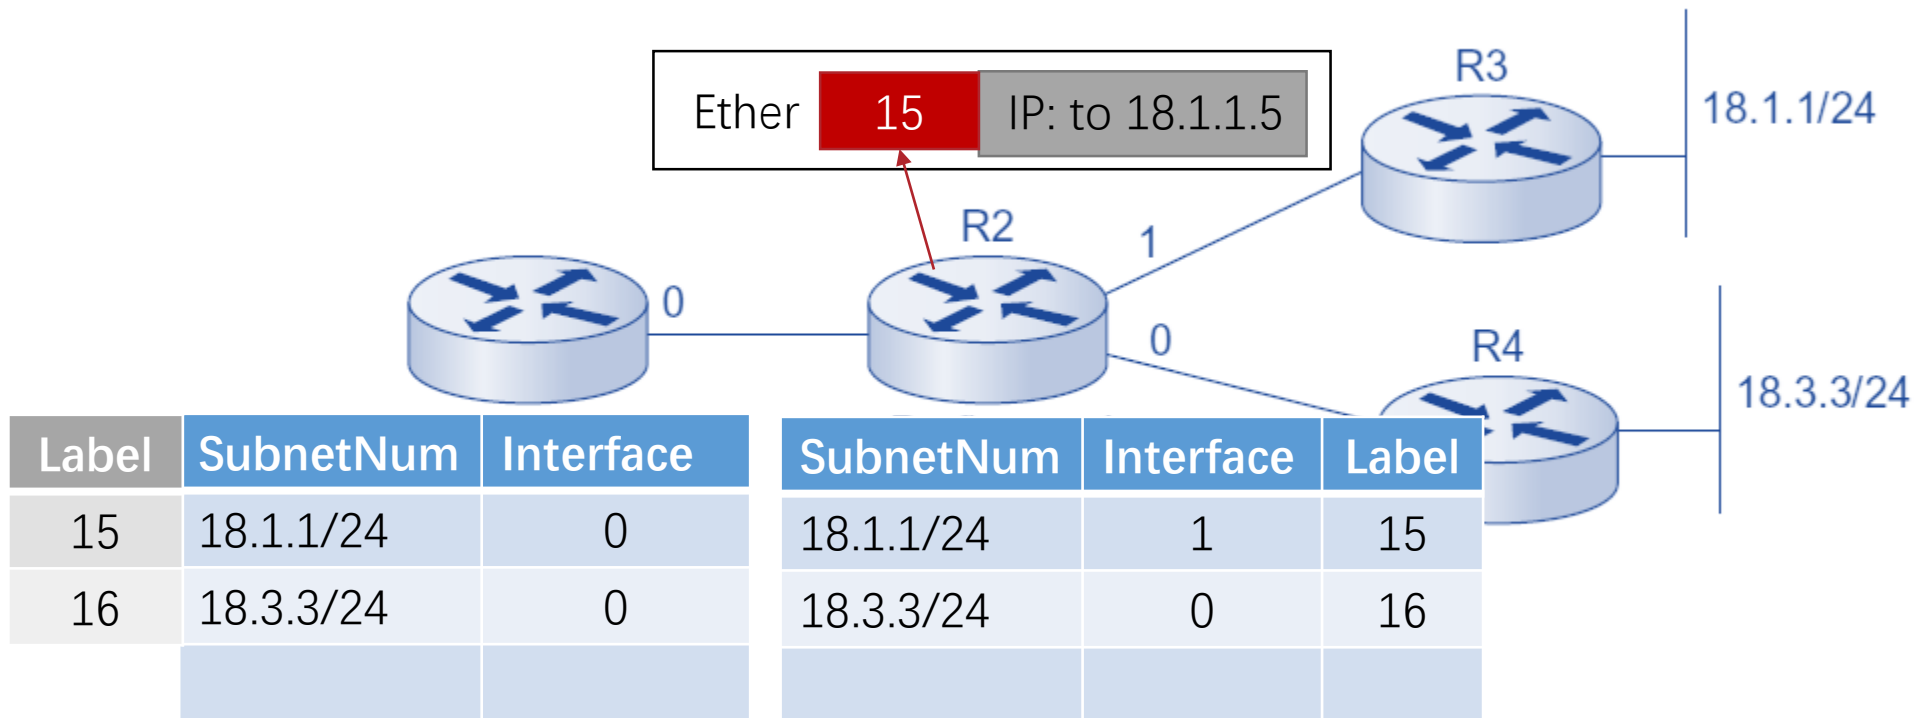
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels



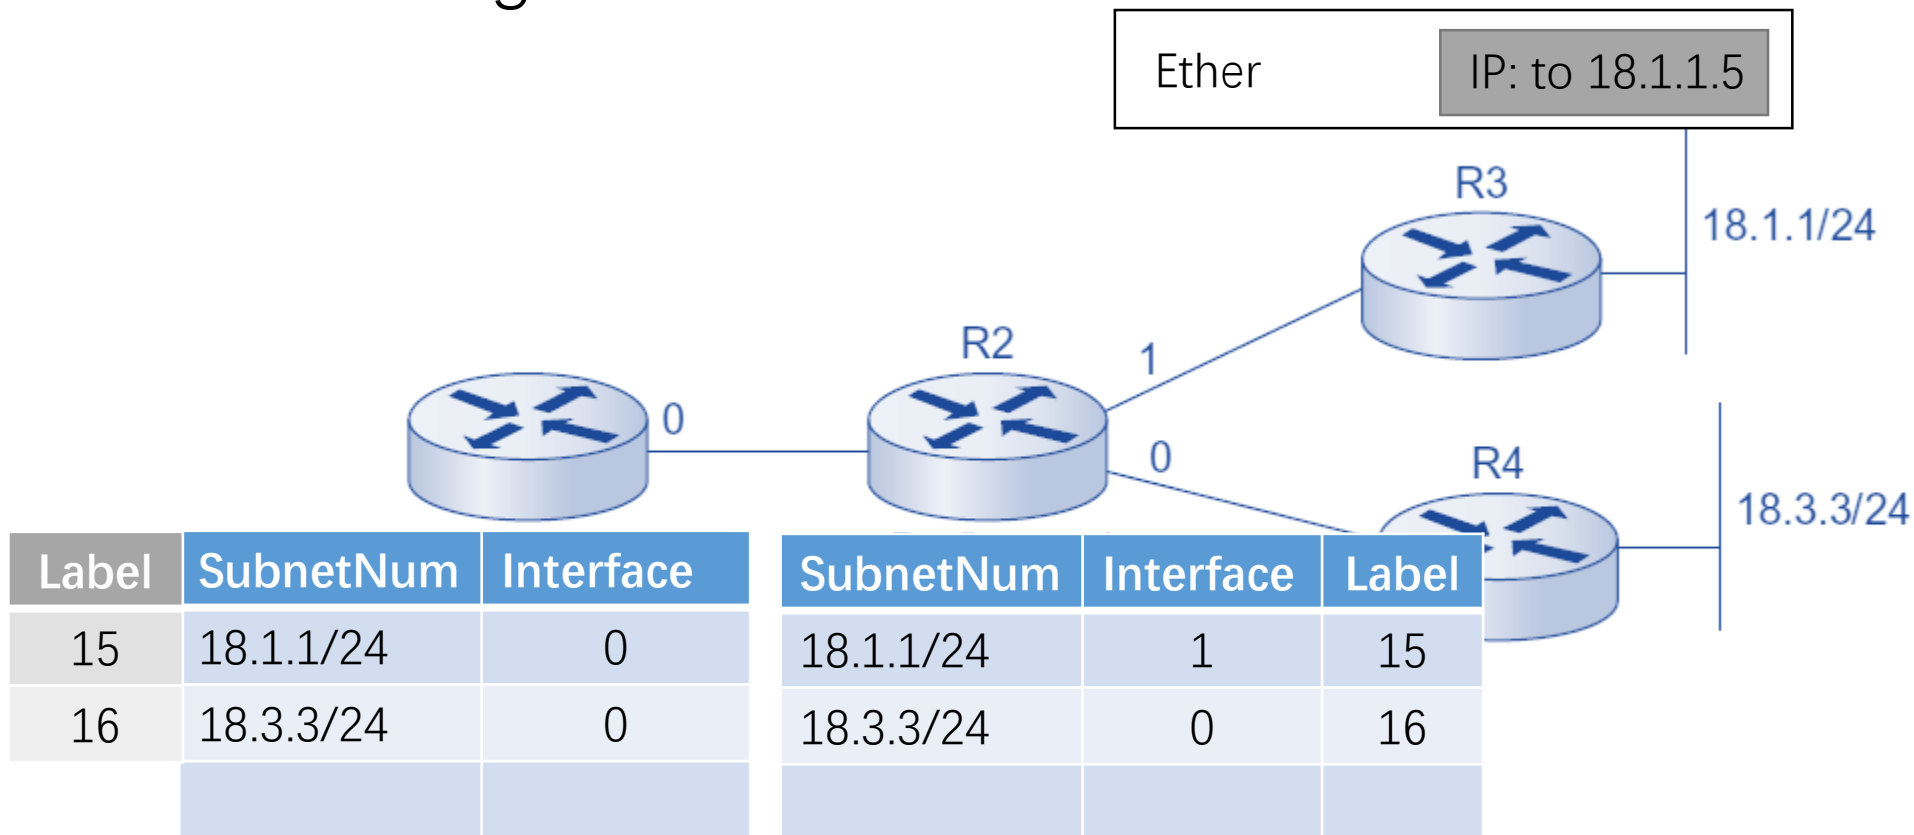
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels



MPLS: Destination-Based Forwarding


- IP Forwarding with Labels



MPLS: Destination-Based Forwarding

- IP Forwarding with Labels
 - Use labels along the path

SubnetNum	Label
18.1.1/24	24

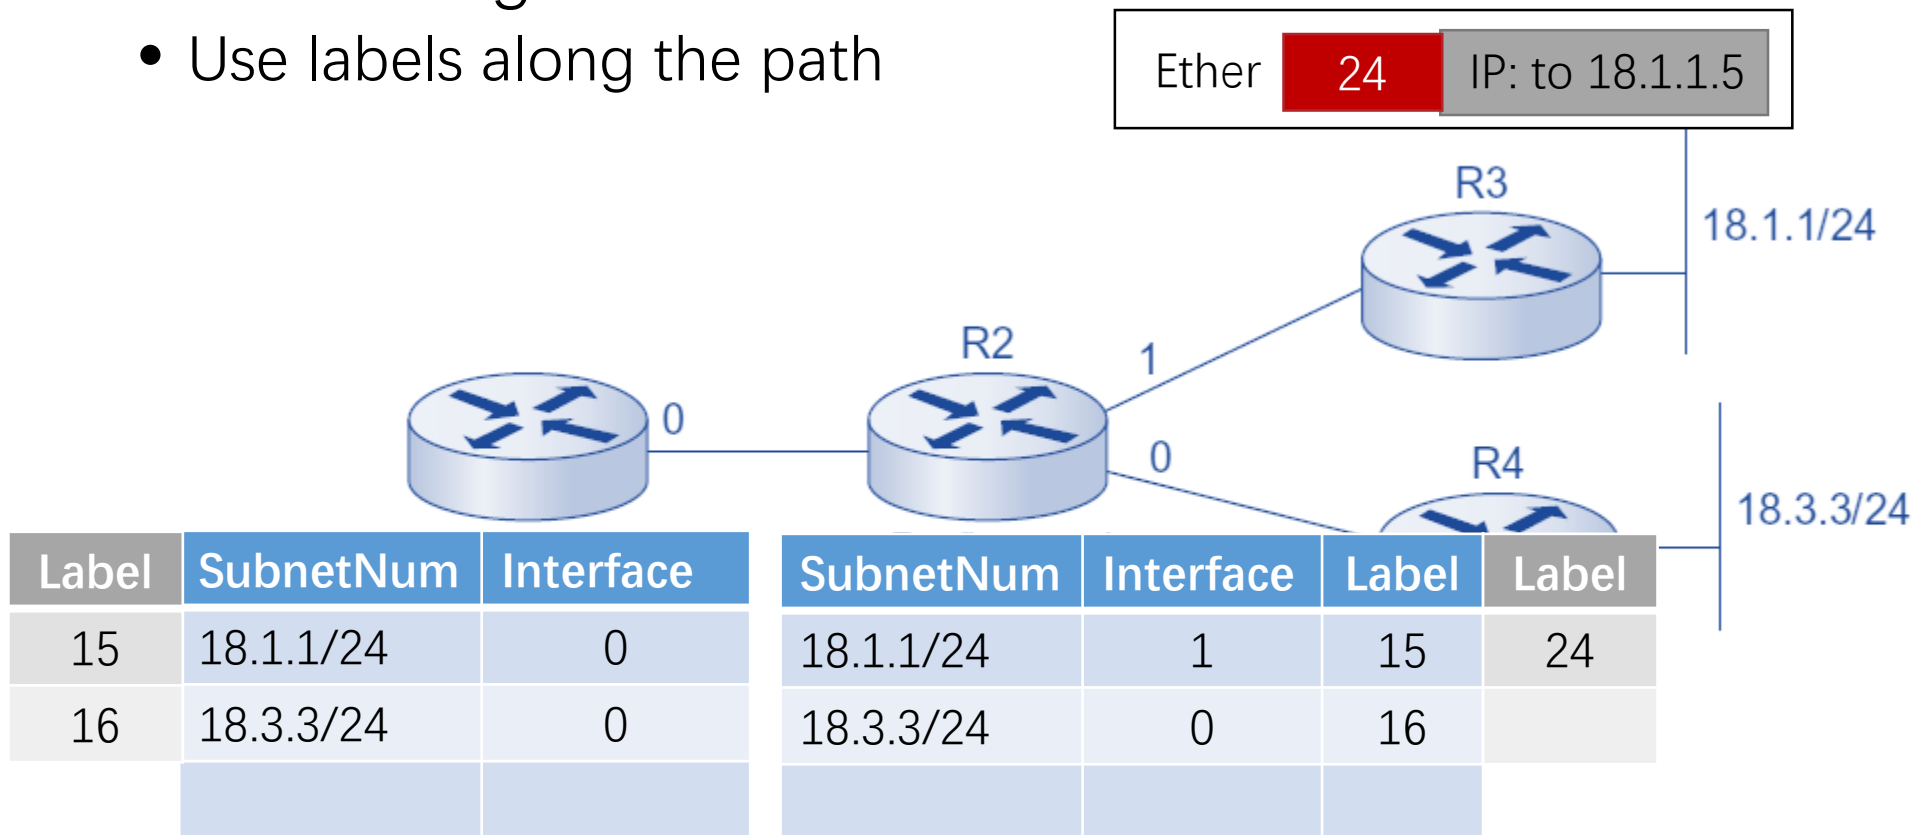


Label	SubnetNum	Interface
15	18.1.1/24	0
16	18.3.3/24	0

SubnetNum	Interface	Label	Label
18.1.1/24	1	15	24
18.3.3/24	0	16	

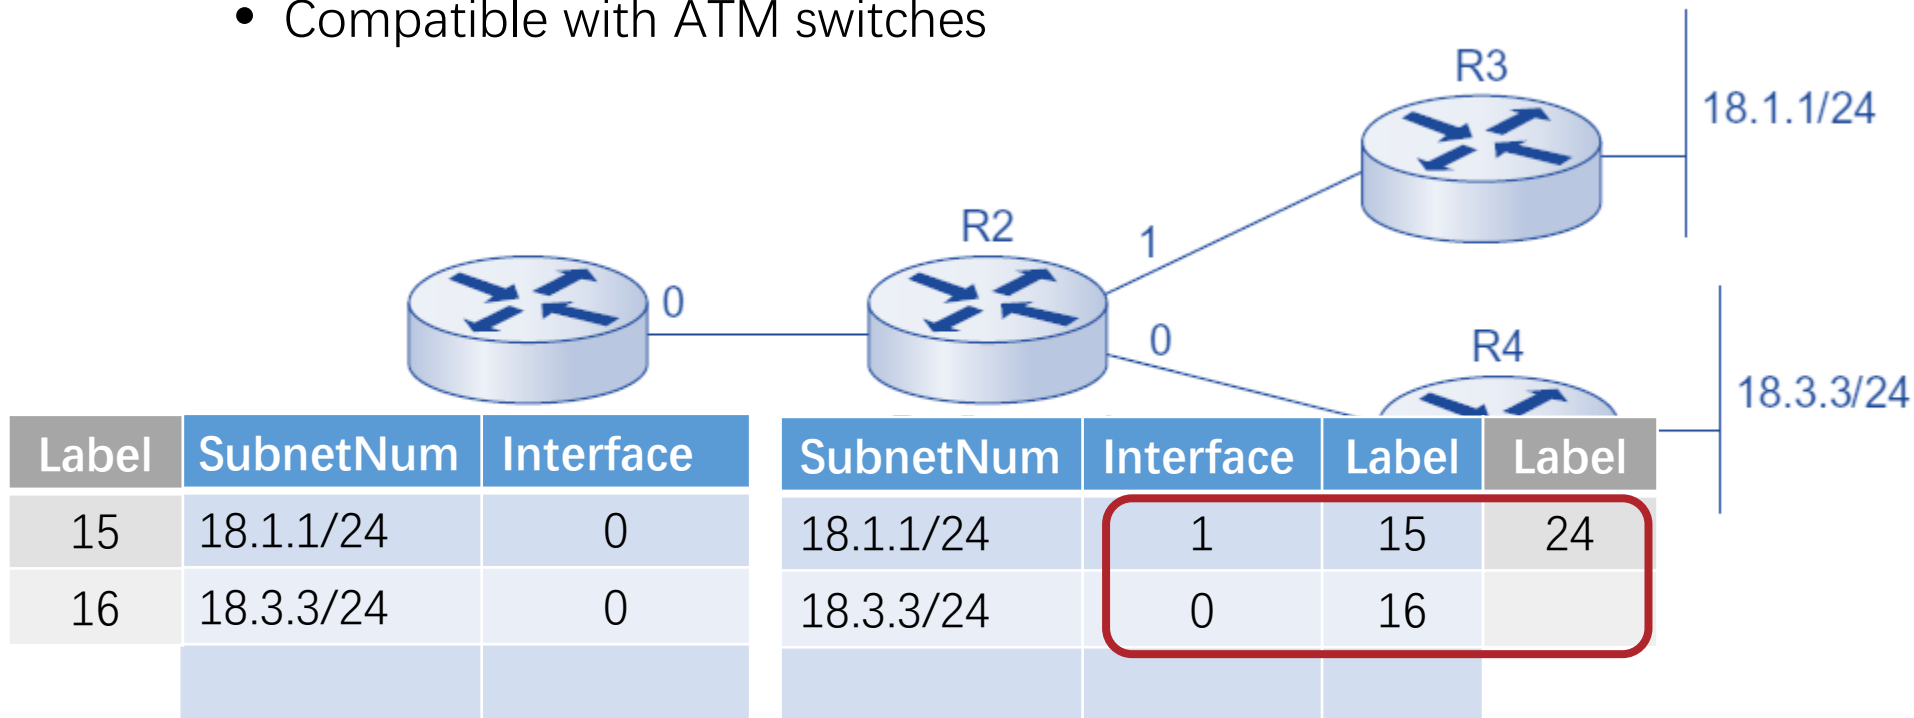
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels
 - Use labels along the path



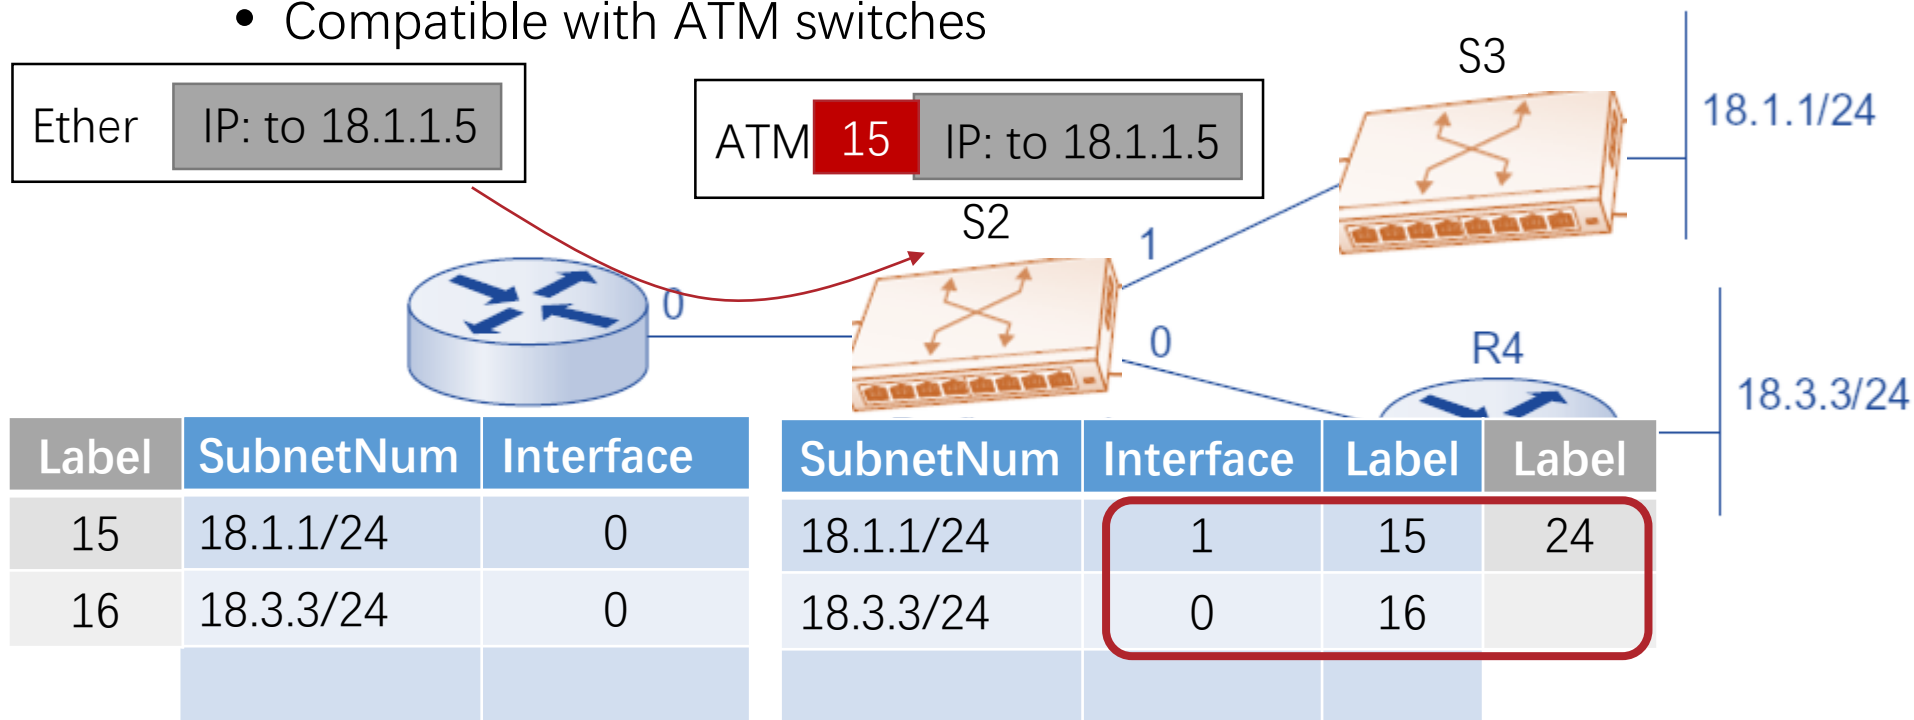
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels
 - The forwarding hardware is based on label
 - Compatible with ATM switches



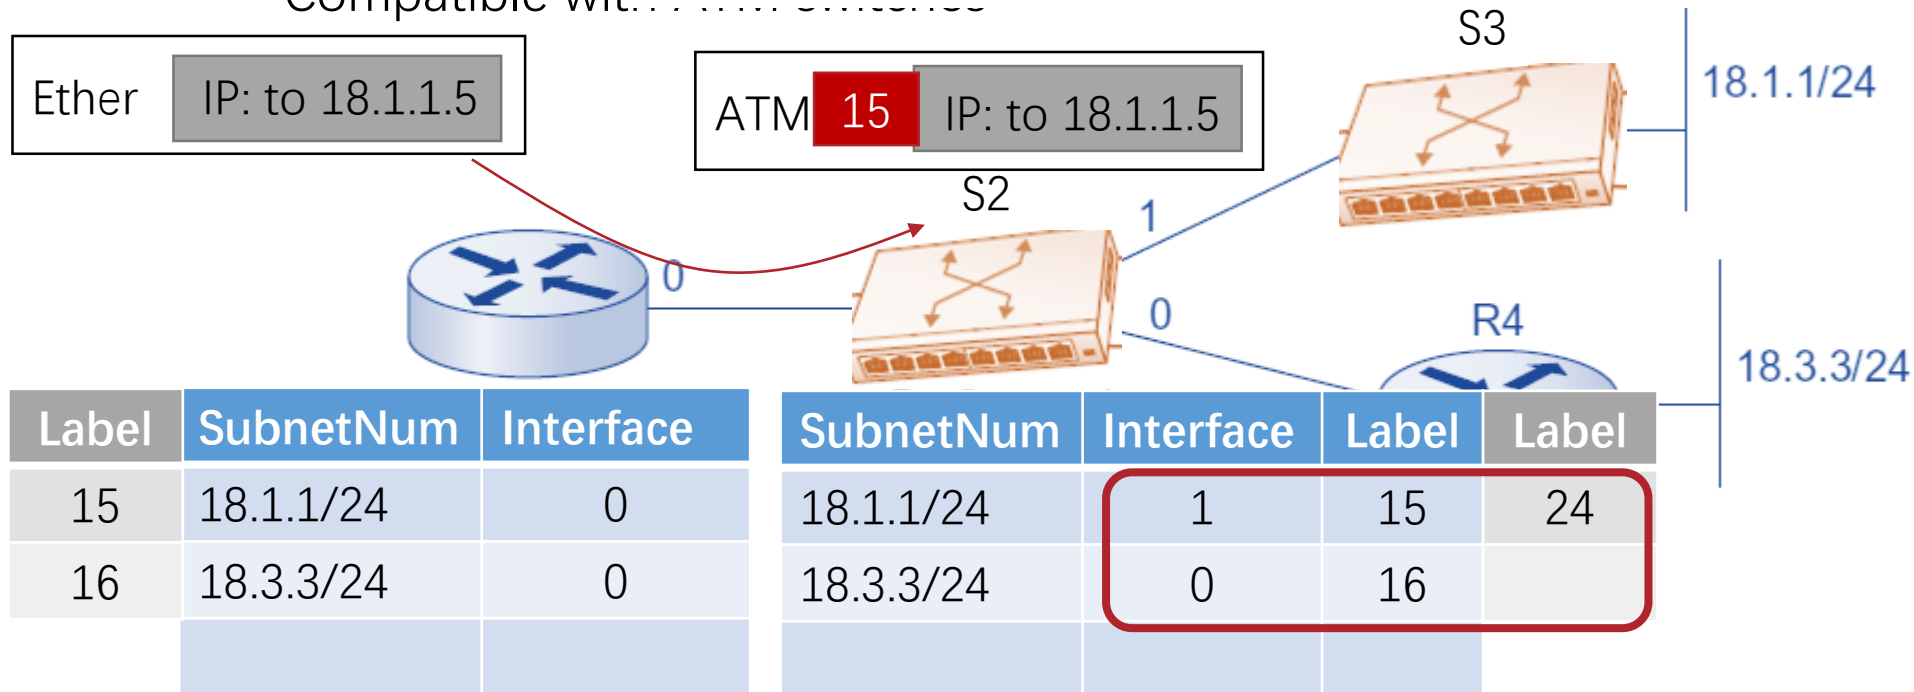
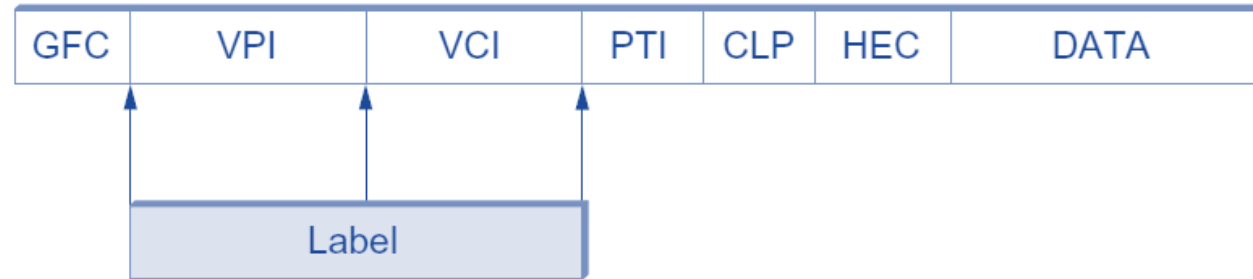
MPLS: Destination-Based Forwarding

- IP Forwarding with Labels
 - The forwarding hardware is based on label
 - Compatible with ATM switches




MPLS: Destination-Based Forwarding

- IP Forwarding with MPLS
 - The forwarding table is based on destination IP
 - Compatible with existing IP networks



MPLS: Destination-Based Forwarding

- Transform ATM switches into IP Routers
 - Only software changes are needed
 - IP routers connected by ATM network are able to see better network topology




MPLS: Explicit Routing

- MPLS Forwarding Decisions
 - Can follow IP routing algorithm
 - Same path as IP forwarding
 - Can differ from IP routing algorithm
- Application
 - Traffic engineering
 - Use destination and source addresses to route flows to same destination differently
 - Fast reroute
 - pre-computed backup paths

MPLS: Explicit Routing

- IP Routing: path to destination determined by destination address alone
 - Contains no information about the traffic of certain path




MPLS: Explicit Routing

- MPLS routing: path to destination can be based on source and destination address

Dest	Interface	L_IN	L_OUT
A	0	/	25

Dest	Interface	L_IN	L_OUT
A	0	/	12

Dest	Interface	L_IN	L_OUT
A	0	25	3
A	1	12	4




MPLS: Explicit Routing

- MPLS routing: path to destination can be based on source and destination address

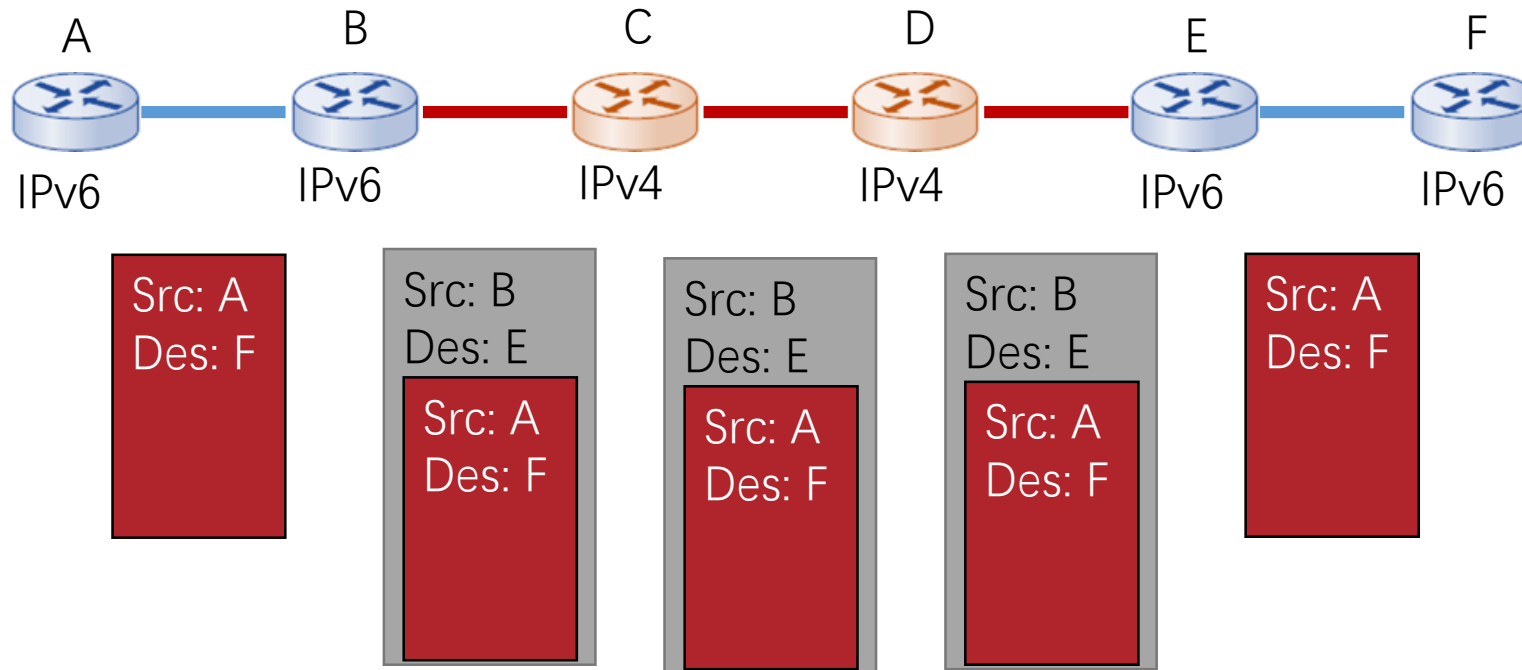
Dest	Interface	L_IN	L_OUT
A	0	/	25
A	0	/	26

Dest	Interface	L_IN	L_OUT
A	0	25	3
A	1	26	4




MPLS: Tunneling

- IPv4 Tunneling



MPLS: Tunneling

- MPLS Tunneling



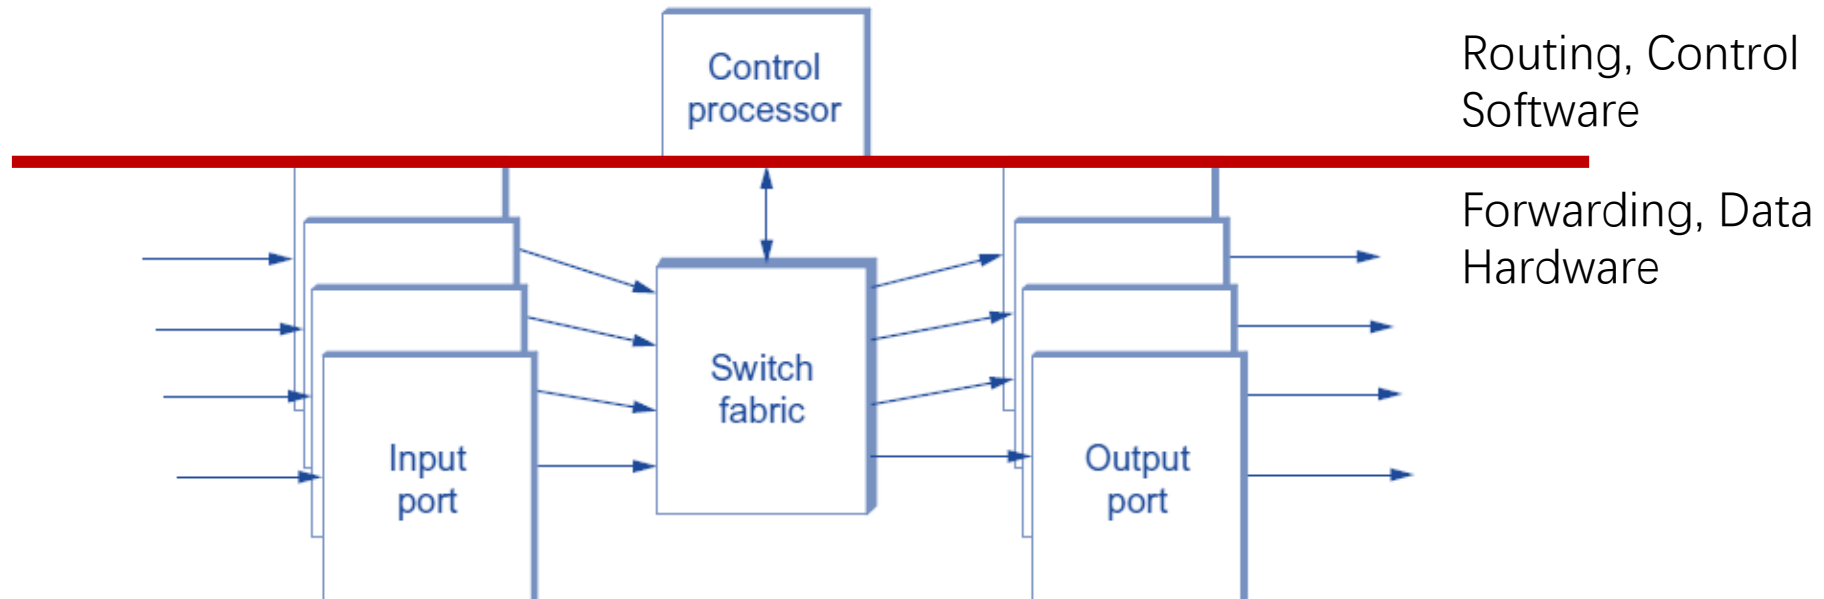
Centralize or Decentralize ?

- That is the question

Router Architecture



- Two Key Functions:
 - Routing algorithms (e.g, RIP, OSPF, etc.)
 - Forwarding packets from input to output ports
- Performance: packet per second




Control Processor

- Function
 - Control
 - Ports and switch fabrics
 - Calculation
 - Routing Algorithm (Router)
 - Push forwarding tables into ports
 - Switches does not run routing algorithms


Input Port

- Deliver Packet to Fabric
 - Handle Different Protocols
- Output Port Looking Up
 - Given packet destination lookup output port
 - Ethernet
 - Using table: destination MAC \Leftrightarrow port
 - ATM
 - Using table: VCI \Leftrightarrow port
 - Goal: complete input port processing at “line speed”
 - 2.48Gbps link, 64byte packet
 - 4.83×10^6 pps
- Buffer




Input Port

- Head-of-line Blocking




Input Port

- Deliver Packet to Network
- Buffer

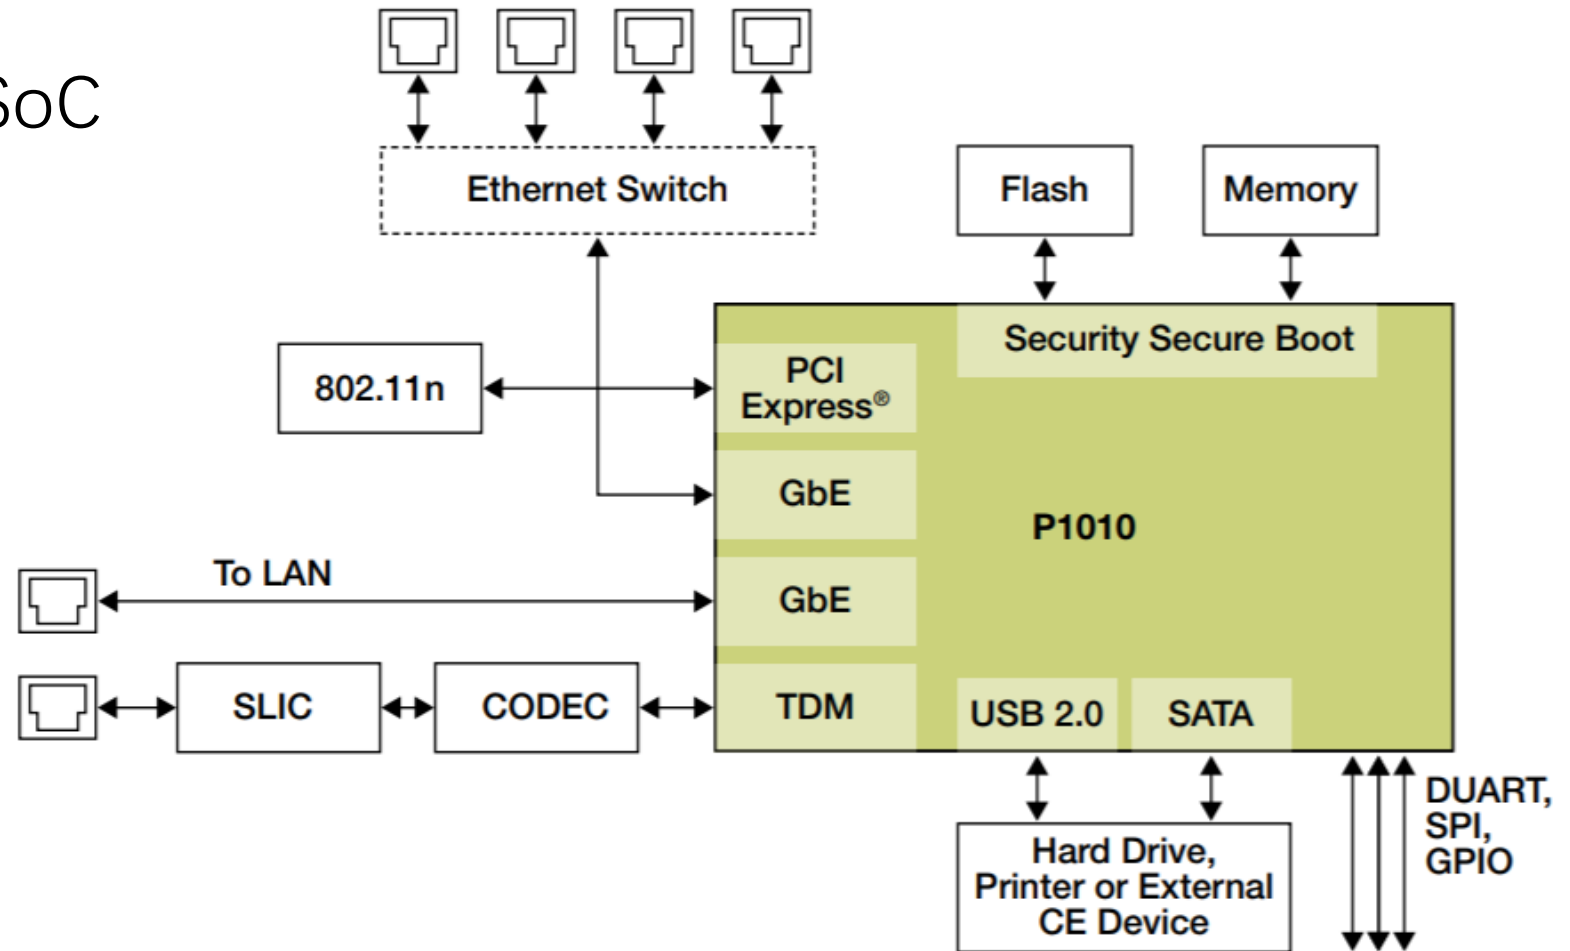


Inside Routers



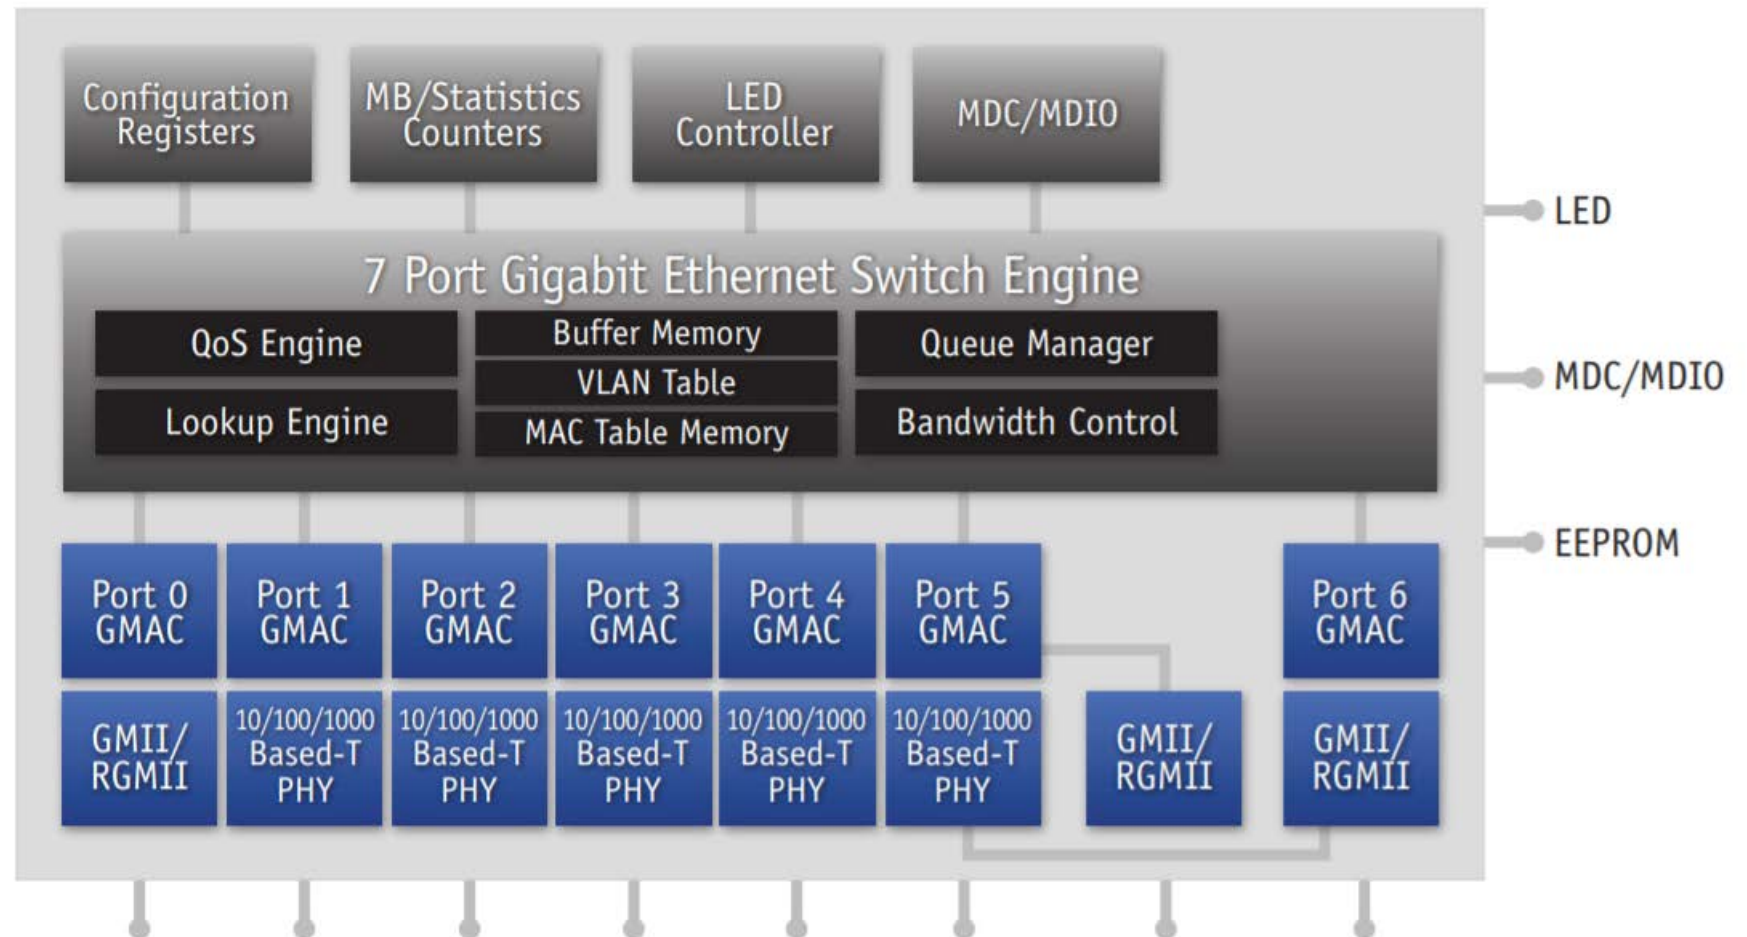
Inside Routers

- Freescale P1010 SoC



Inside Routers

- AR8327 Switch




Switching Fabrics

- Function
 - Transfer packet from input buffer to appropriate output buffer
- Switching Throughput
 - Rate at which packets can be transfer from inputs to outputs
 - N inputs: switching throughput N times line rate desirable
- Four Types
 - Shared Bus
 - Shared Memory
 - Crossbar
 - Self-routing


Shared Bus/Memory

- Datagram from input port to output port via a shared bus
 - 2 bus crossings per datagram
 - Bus and memory bandwidth determines switch throughput



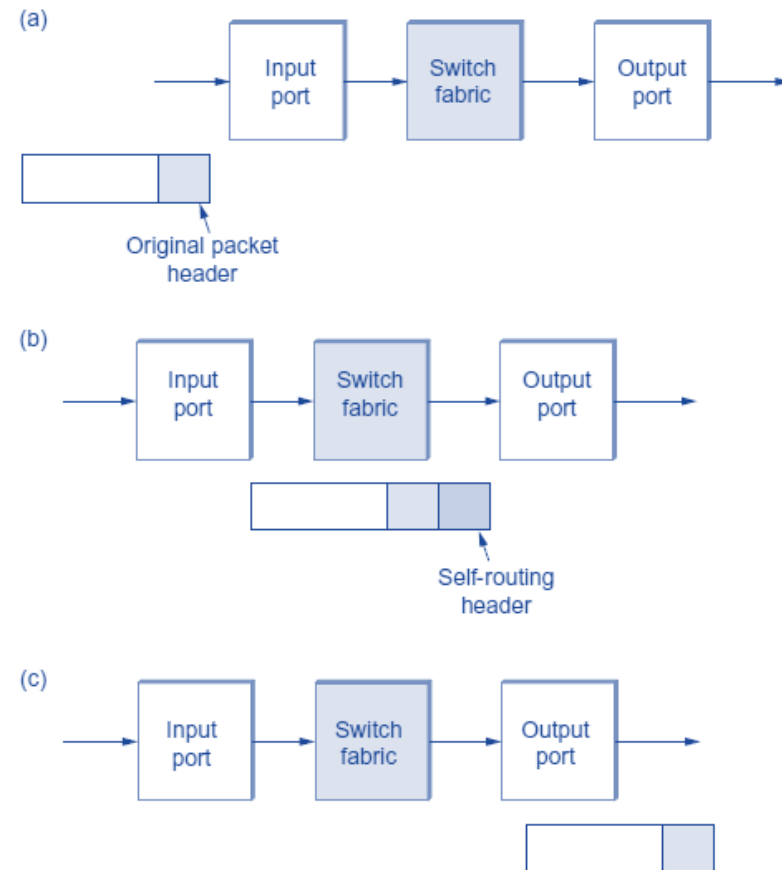
Crossbar

- A crossbar switch is a matrix of pathways that can be configured to connect any input port to any output port




Self-routing

- Routing Header




Self-routing

- Switching Element
 - 0=> up
 - 1=> down




Self-routing

- Switching Element
 - 0=> up
 - 1=> down




Self-routing

- Switching Element
 - 0=> up
 - 1=> down




Self-routing

- Switching Element
 - 0 => up
 - 1 => down




Self-routing

- Switching Element
 - Collision: Two packets with same output ports




Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header




Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header




Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header




Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header




Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header



Self-routing

- Banyan Network
 - Collision Free
 - Input Packets are sorted according to routing header




U.S. Patent

Mar. 20, 1990

Sheet 1 of 28

4,910,730



Open Sourced Wireless Router

- Linksys WRT54G



OpenWrt
Wireless Freedom

Reference

- Textbook 4.1
- Textbook 4.3
- Textbook 3.4