



CS120: Computer Networks

Lecture 16. TCP 2

Zhice Yang

Transmission Control Protocol (TCP)

- RFC: 793,1122,1323, 2018, 2581
- Goal: Reliable, In-order Delivery
 - Connection oriented
 - Flow control
 - Congestion control
- Core Algorithm: Sliding Window


Delay \times Bandwidth

- Quantify the utilization of the link

Table 1.1 Sample Delay \times Bandwidth Products				
Link type	Bandwidth (typical)	One-way distance (typical)	Round-trip delay	RTT \times Bandwidth
Dial-up	56 kbps	10 km	87 μ s	5 bits
Wireless LAN	54 Mbps	50 m	0.33 μ s	18 bits
Satellite	45 Mbps	35,000 km	230 ms	10 Mb
Cross-country fiber	10 Gbps	4,000 km	40 ms	400 Mb

Delay \times Bandwidth

- Efficiency Problem
 - 1.5Mbps bandwidth
 - 45ms RTT
 - 1KB per packet
 - Effective Rate = $1024 \times 8 / (1024 \times 8 / 1.5\text{M} + 45\text{ms}) = 162\text{kbps}$
- Solution
 - Pipeline
 - Full pipe
 - $1.5\text{Mbps} \times 45\text{ms} / 1\text{kB} = 8 \text{ packets}$




Sliding Window


- Sender Buffer
 - Retransmit
- Receiver Buffer
 - Frame order

Sliding Window

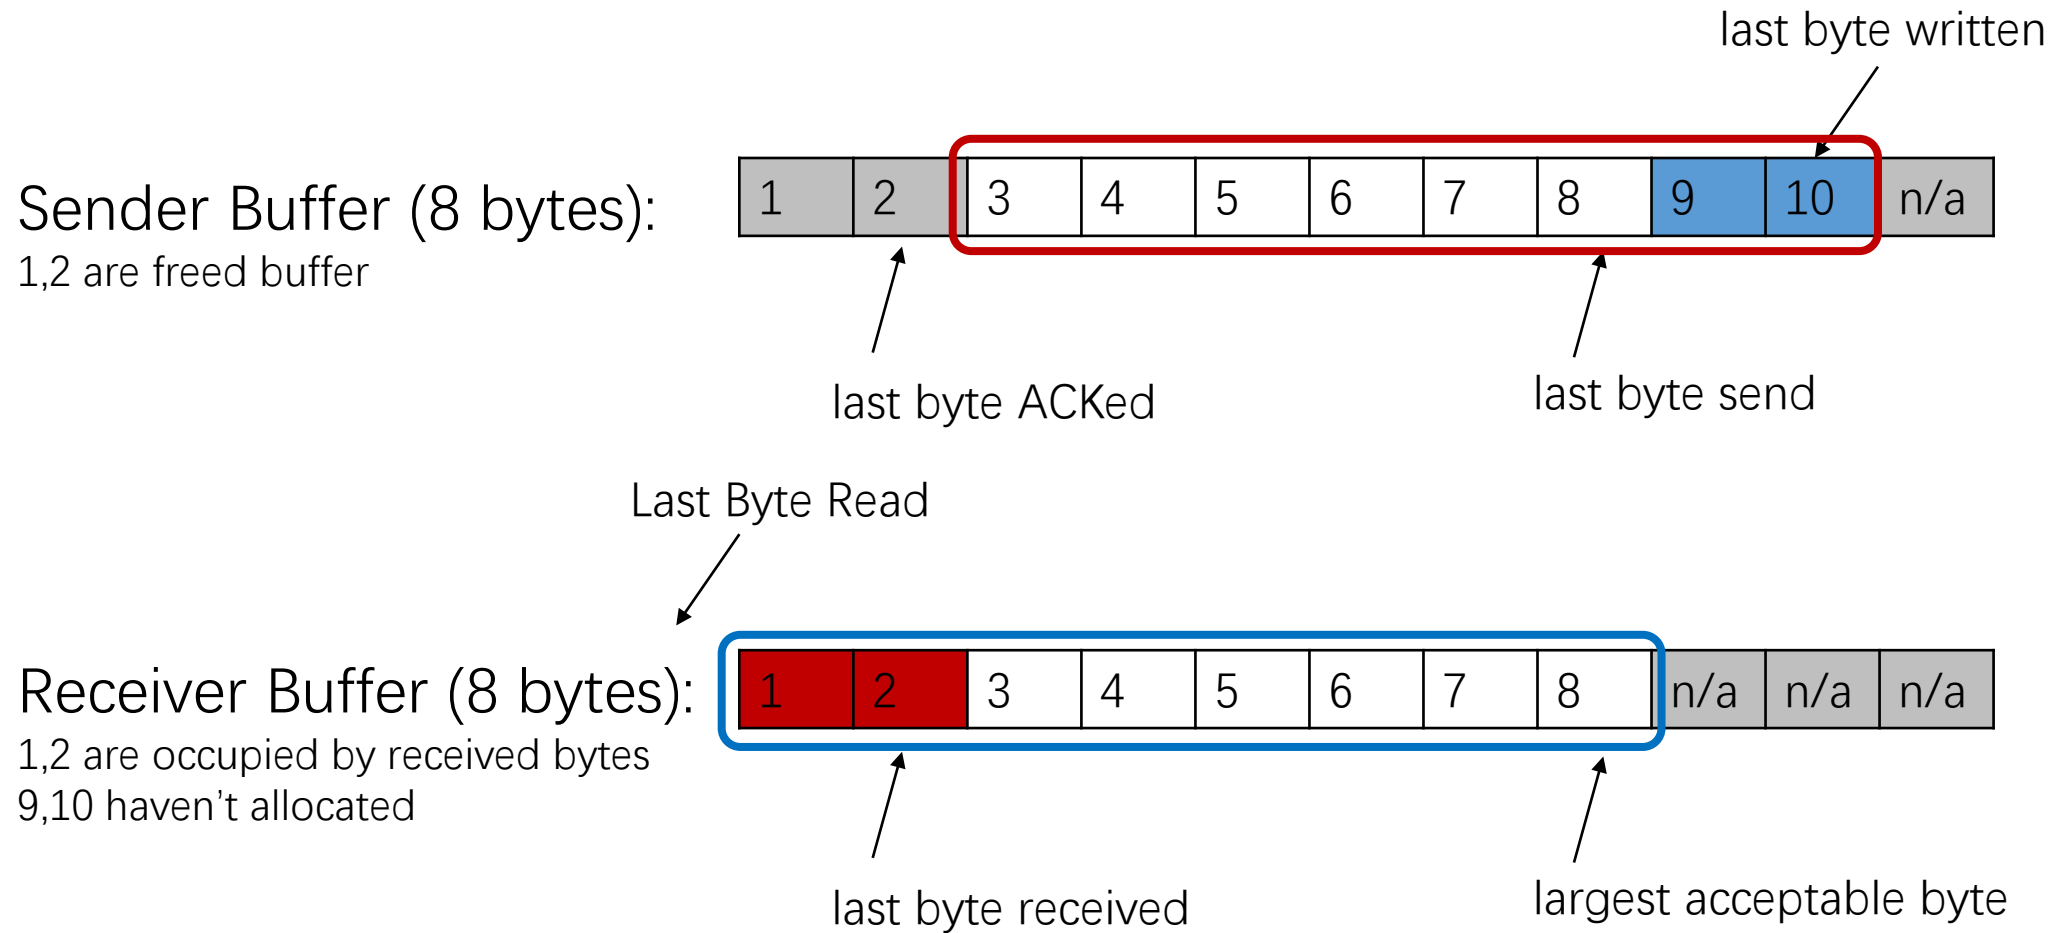
Sender Buffer:



Receiver Buffer:




Sliding Window in TCP




Sliding Window in TCP

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




Sliding Window in TCP

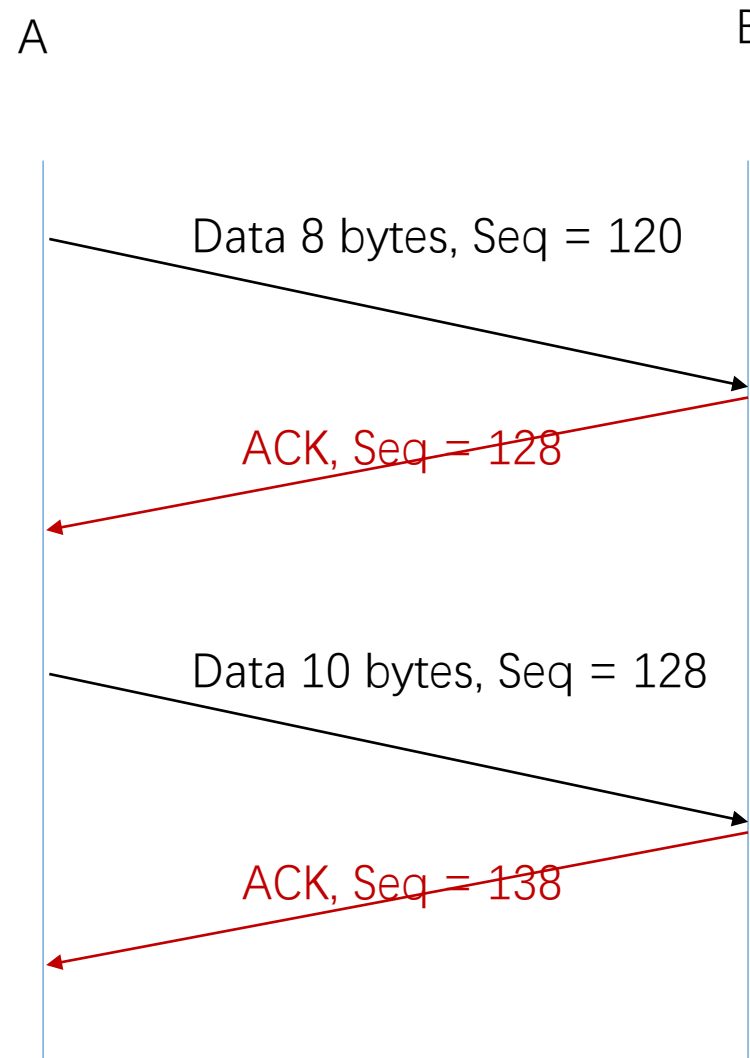
Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (8 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver
9 is allocated for receiving



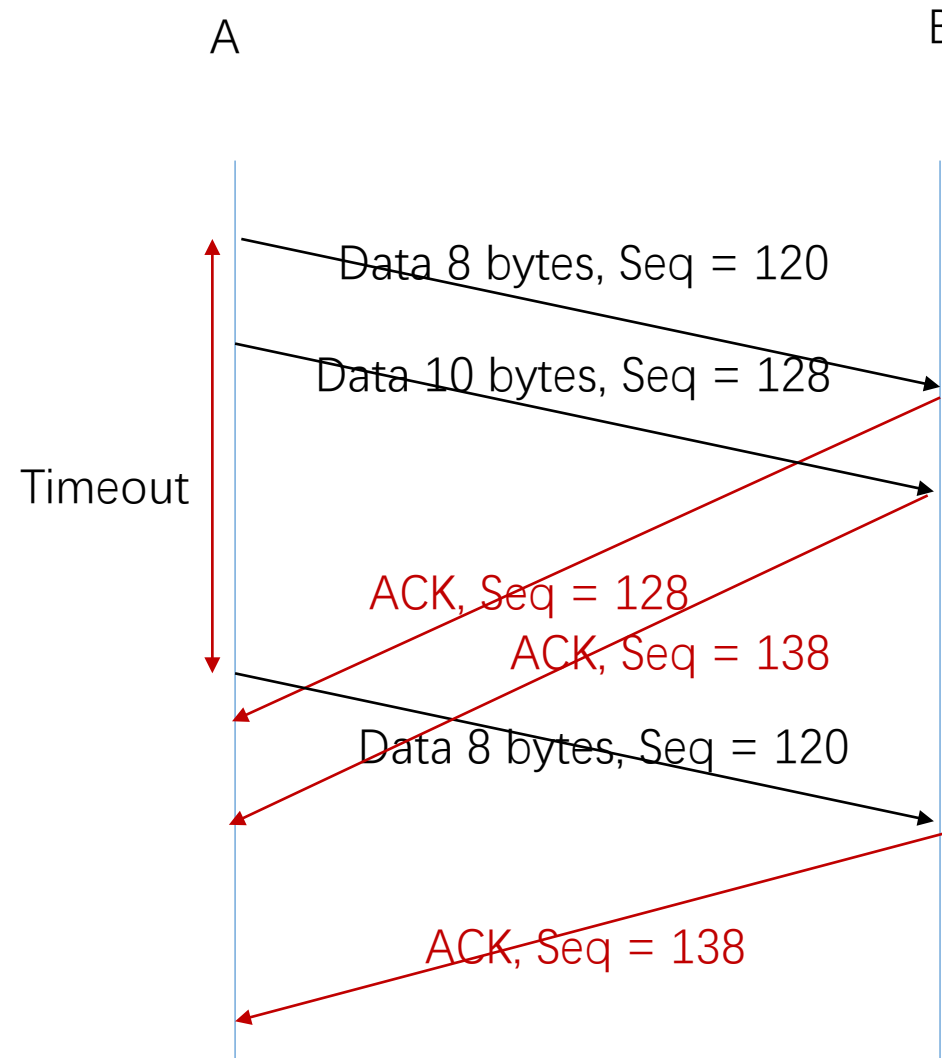
Sliding Window in TCP: Examples



Sliding Window in TCP: Examples

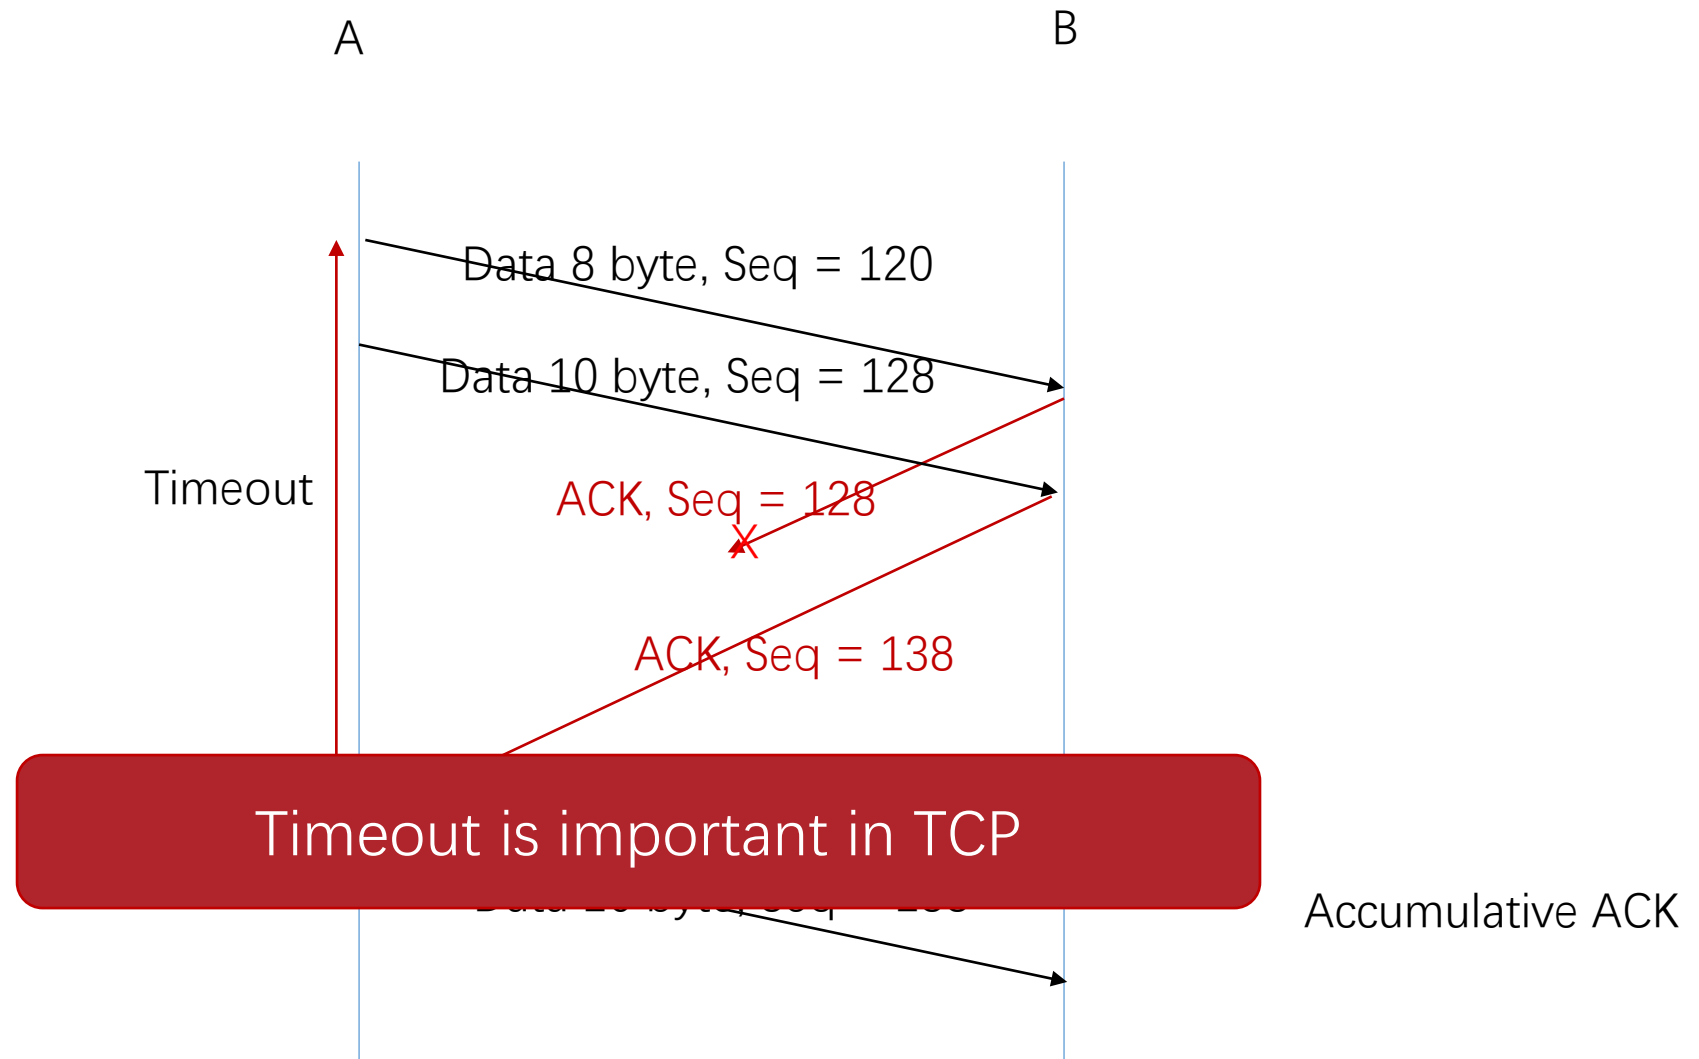


Sliding Window in TCP: Examples



Accumulative ACK


Sliding Window in TCP: Examples



Sliding Window in TCP: Adaptive Timeout

- Original Algorithm
 - Measure SampleRTT for each segment/ ACK pair
 - Compute weighted average of RTT
 - $\text{EstimatedRTT} = a * \text{EstimatedRTT} + (1 - a) * \text{SampleRTT}$
 - a between 0.8 and 0.9
 - Set timeout based on EstimatedRTT
 - $\text{TimeOut} = 2 * \text{EstimatedRTT}$

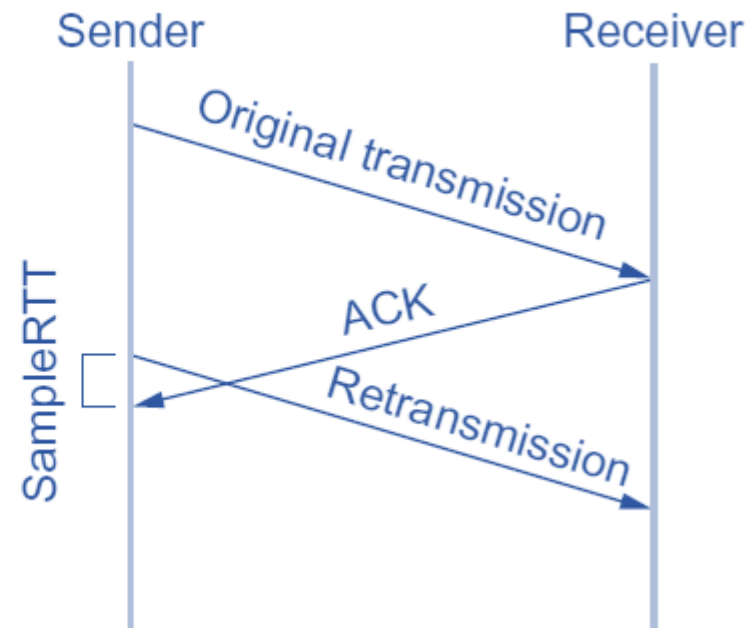





Sliding Window in TCP: Adaptive Timeout

- Original Algorithm

- Measure
- Correct
- Set



Sliding Window in TCP: Adaptive Timeout

- Karn/Partridge Algorithm
 - Do not sample RTT when retransmitting
 - Double timeout after each retransmission
 - (But do not increase sliding window)
 - Retransmission \Rightarrow network is congested, should slow down
 - Congestion Control (will introduce in next few lectures)
- The estimation for RTT is not accurate
 - RTT is a critical signal for identifying network congestion

Sliding Window in TCP: Adaptive Timeout

- Jacobson/Karels Algorithm

- Take RTT variance into account

Difference = SampleRTT - EstimatedRTT

EstimatedRTT = EstimatedRTT + (δ * Difference)

Deviation = Deviation + δ * (|Difference| - Deviation)

- δ is typically set to 1/8

TimeOut = μ * EstimatedRTT + φ * Deviation

- μ is typically set to 1 and φ is set to 4

Sliding Window in TCP: Adaptive Timeout


- Jacobson/Karels Algorithm Implementation

$\text{Difference} = \text{SampleRTT} - \text{EstimatedRTT}$

$\text{EstimatedRTT} = \text{EstimatedRTT} + (\delta * \text{Difference})$

$\text{Deviation} = \text{Deviation} + \delta * (|\text{Difference}| - \text{Deviation})$

$\text{TimeOut} = \mu * \text{EstimatedRTT} + \varphi * \text{Deviation}$



```
SampleRTT -= (EstimatedRTT >> 3);
```

```
EstimatedRTT += SampleRTT;
```

```
if (SampleRTT < 0)
```

```
    SampleRTT = -SampleRTT;
```

```
SampleRTT -= (Deviation >> 3);
```

```
Deviation += SampleRTT;
```

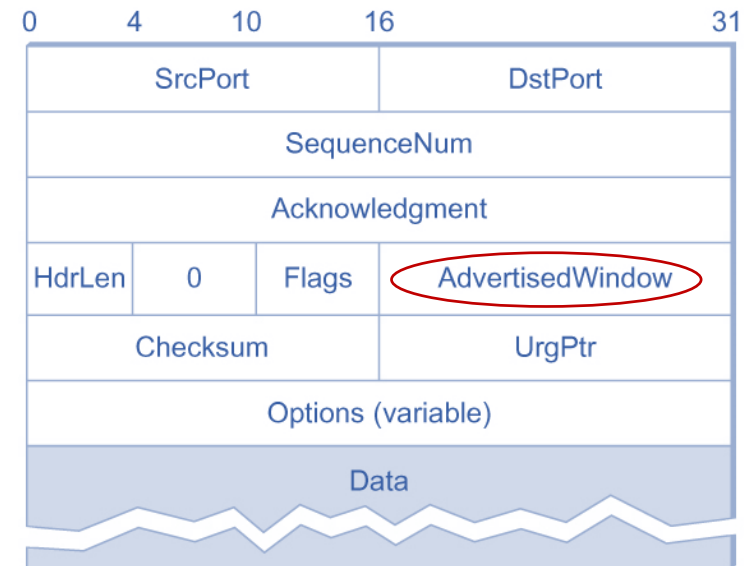
```
TimeOut = (EstimatedRTT >> 3) + (Deviation >> 1);
```

Transmission Control Protocol (TCP)


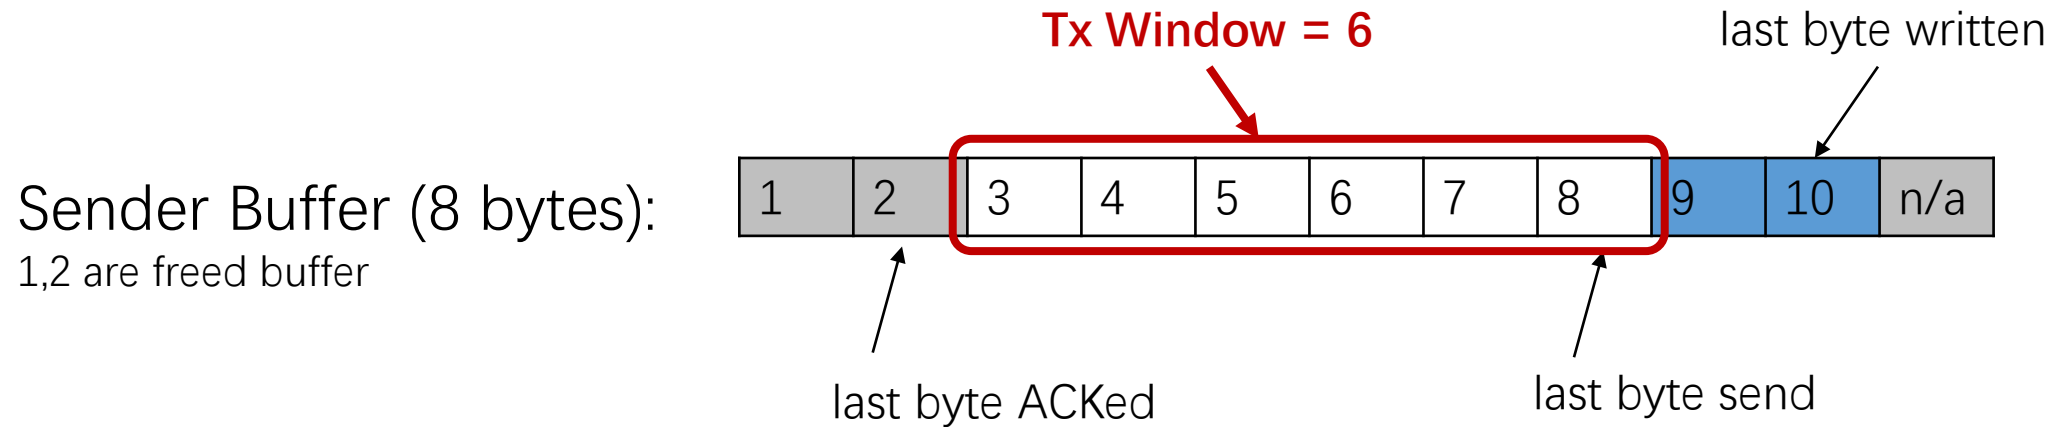
- RFC: 793,1122,1323, 2018, 2581
- Goal: Reliable, In-order Delivery
 - Connection oriented
 - Flow control
 - Congestion control
- Core Algorithm: Sliding Window

TCP Flow Control

- Goal:
 - Sender does not overrun receiver's buffer
- Approach:
 - Rx window is not fixed, adjusted according to receiving buffer
 - The receiver advertises an adjustable window size (AdvertisedWindow field in TCP header)
 - Sender is limited to having no more than AdvertisedWindow bytes of unACKed data at any time




TCP Flow Control




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver

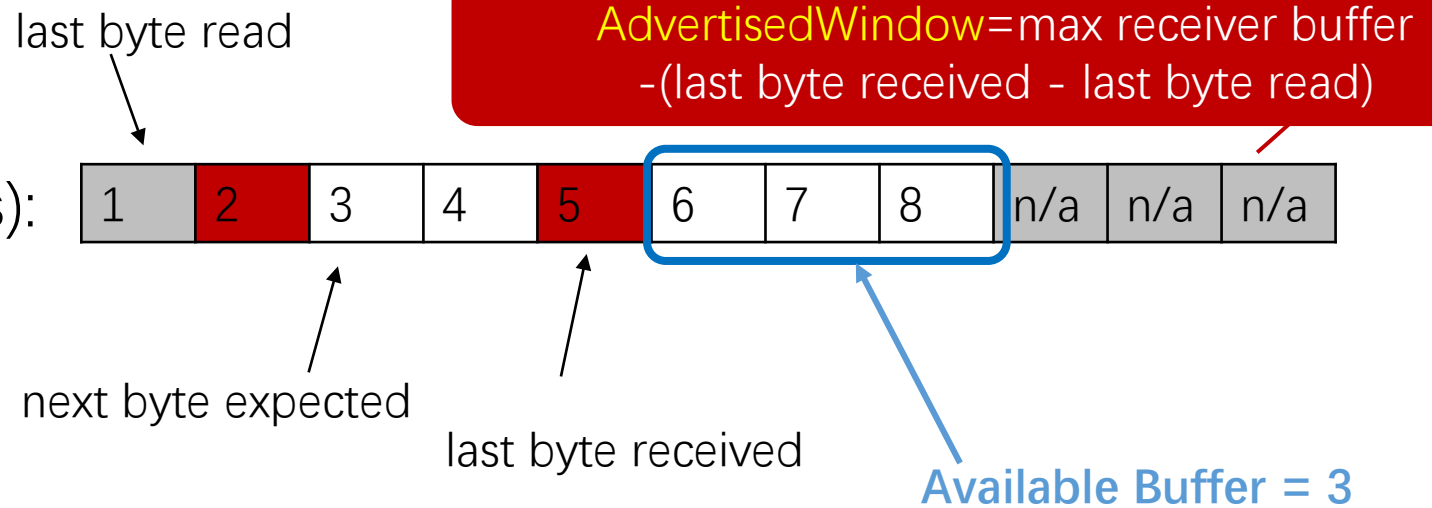


TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

last byte sent – last byte ACKed \leq
AdvertisedWindow

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

last byte sent – last byte ACKed \leq
AdvertisedWindow

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver

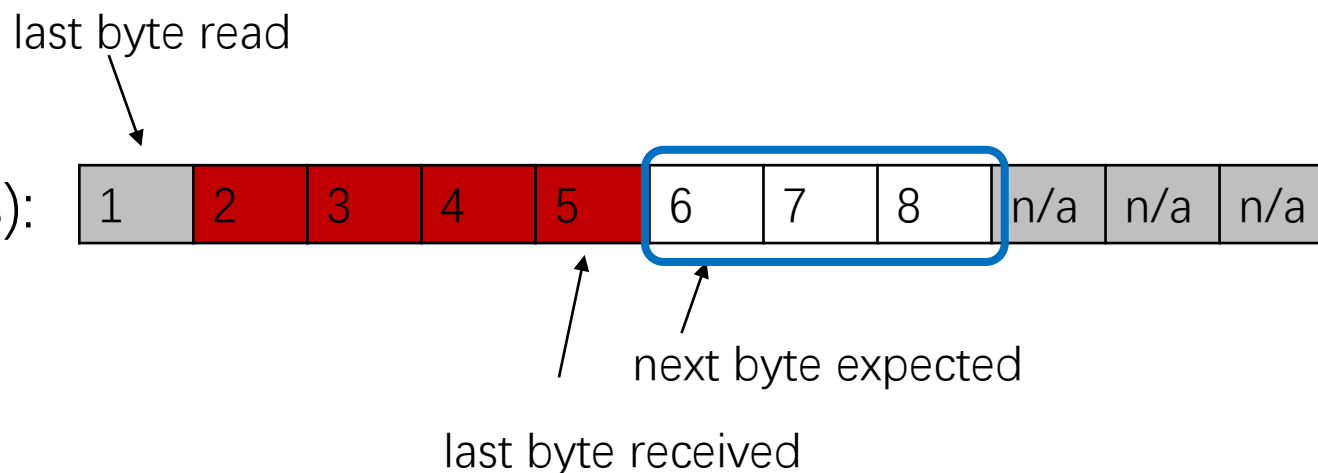


TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver

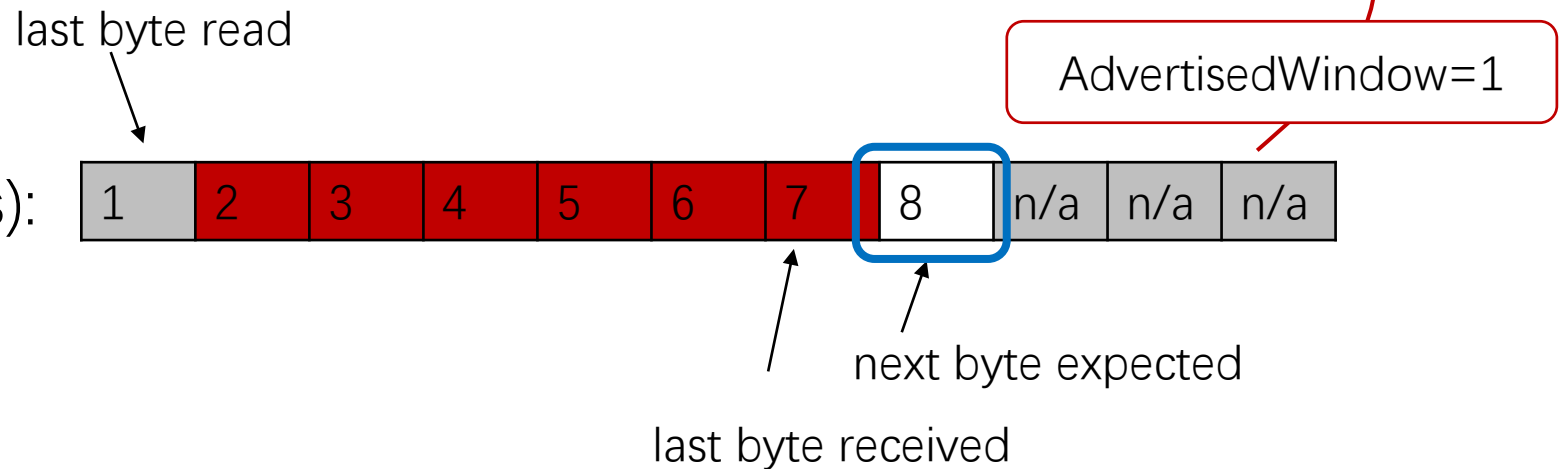


TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver




TCP Flow Control

Sender Buffer (8 bytes):
1,2 are freed buffer




Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver



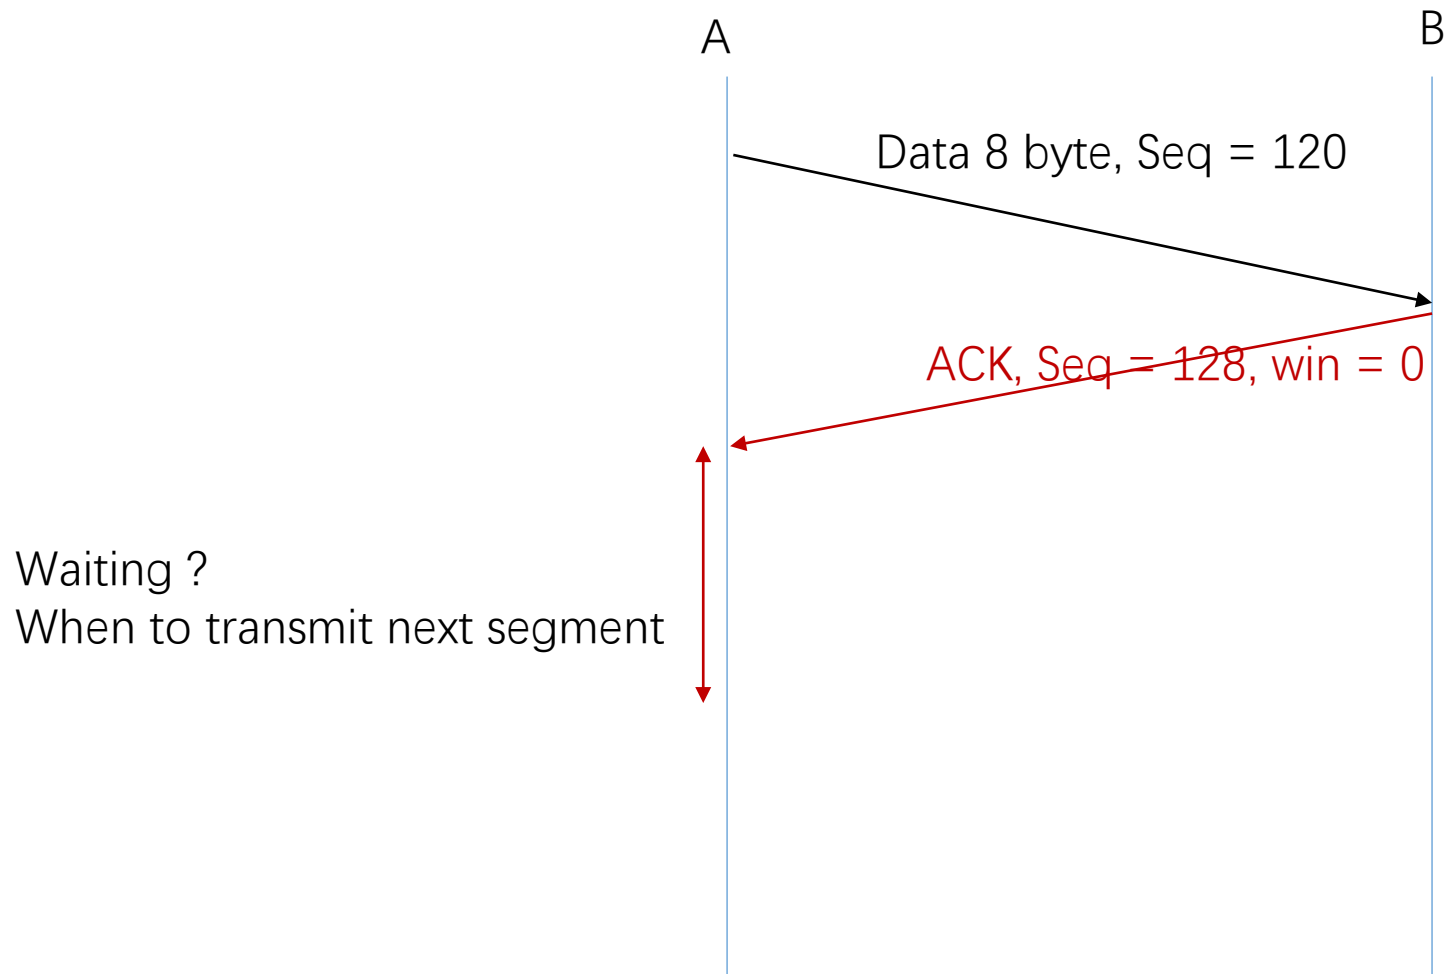
TCP Flow Control

- Receiver throttles sender by advertising a window size no larger than the amount it can buffer:
 - $\text{AdvertisedWindow} = \text{max receiver buffer} - (\text{last byte received} - \text{last byte read})$
 - Indicate the amount of free space available in the receive buffer
- Sender must adhere to AdvertisedWindow from the receiver such that:
 - $\text{last byte sent} - \text{last byte ACKed} \leq \text{AdvertisedWindow}$

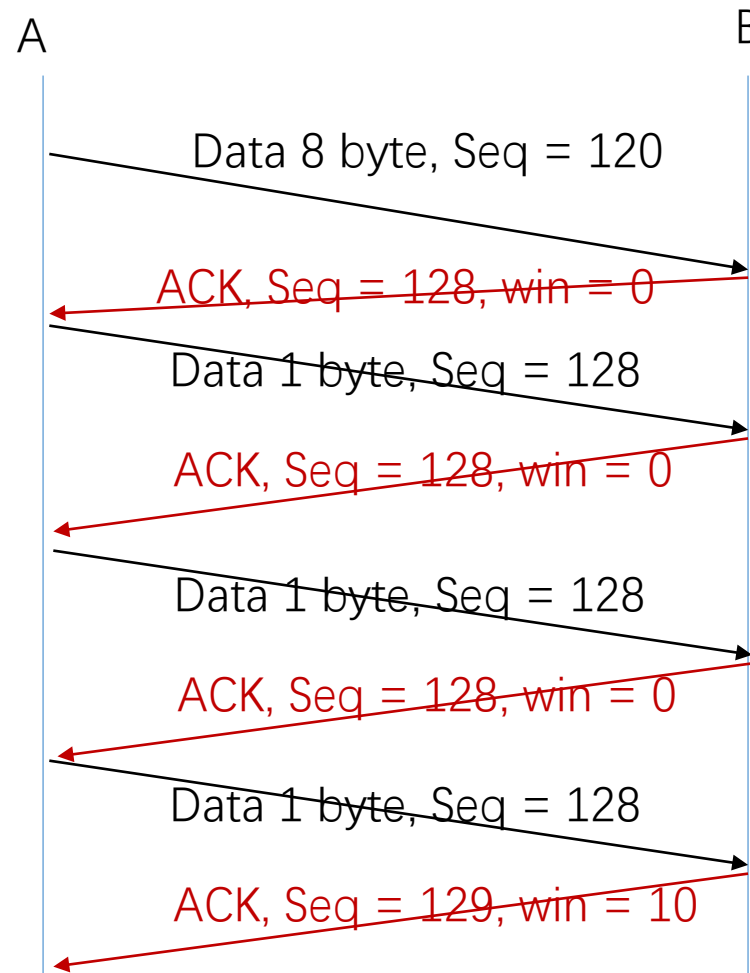
TCP Flow Control: 0 window case



TCP Flow Control: 0 window case



TCP Flow Control: 0 window case




TCP Flow Control: Silly Window Syndrome

Case1: window Size < maximum segment size (6)

Sender Buffer (8 bytes):

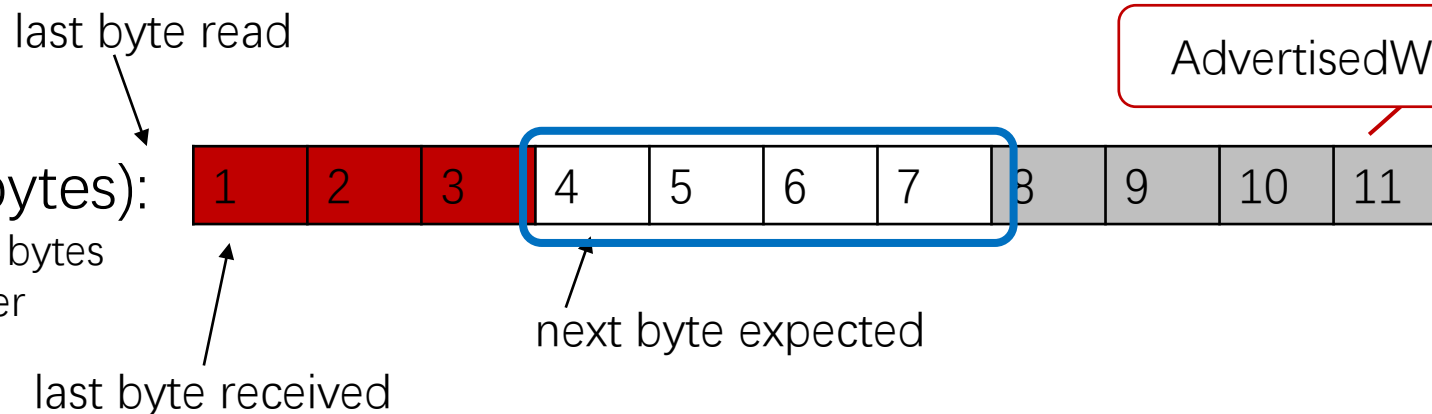
1,2 are freed buffer



AdvertisedWindow=4

Receiver Buffer (7 bytes):


2,5 are occupied by received bytes
1 is read and freed by receiver



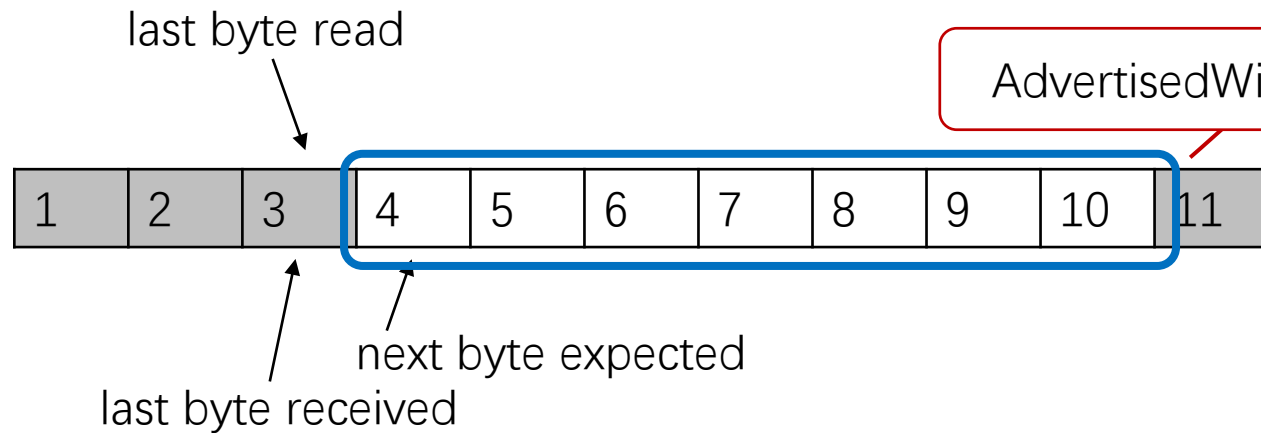
TCP Flow Control: Silly Window Syndrome

Case2: available data < maximum segment size (6)

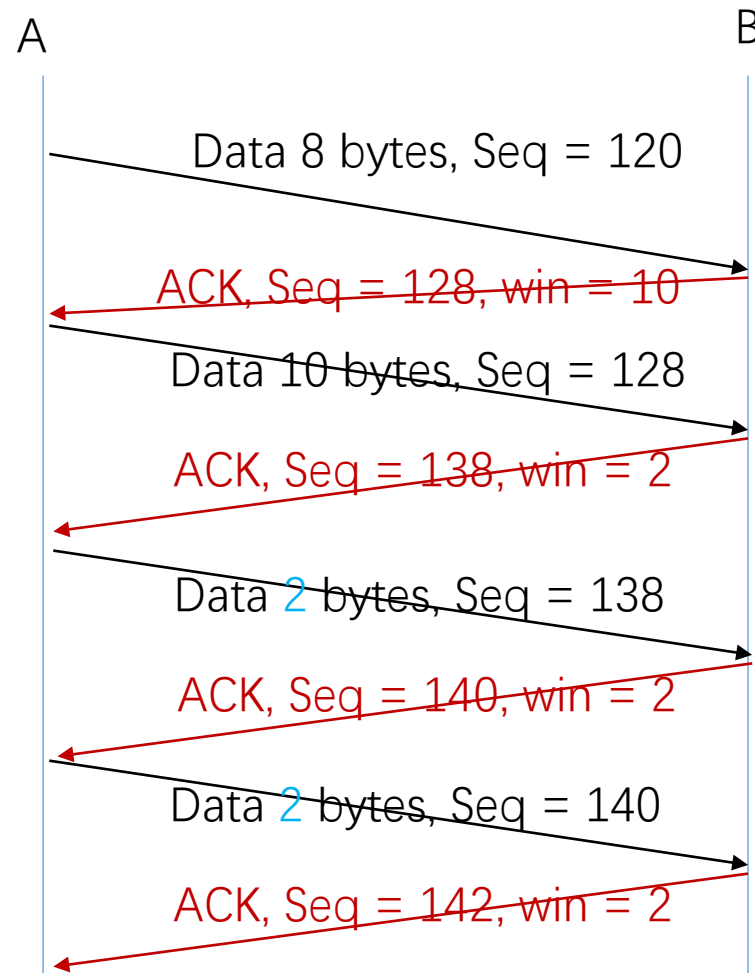
Sender Buffer (8 bytes):
1,2 are freed buffer



Receiver Buffer (7 bytes):
2,5 are occupied by received bytes
1 is read and freed by receiver



TCP Flow Control: Silly Window Syndrome



TCP: How to Trigger Transmission

- Best of Effort
 - Silly Window Syndrome
- Wait Until Full
 - Less Timely
- Hybrid
 - Wait Until Full + Timer

TCP: Nagle's Algorithm

When the application produces data to send

- if both the available data and the window \geq MSS

 - send a full segment

- else

 - if there is unACKed data in flight

 - buffer the new data until an ACK arrives

 - else

 - send all the new data now

Transmission Control Protocol (TCP)

- RFC: 793,1122,1323, 2018, 2581
 - Goal: Reliable, In-order Delivery
 - Connection oriented
 - Flow control
 - Congestion control
 - Core Algorithm: Sliding Window
- Some Patches
- 

TCP Record Boundaries

- Problem: TCP is a byte stream, how to specify a message in the stream
 - UDP: one message per frame

Tx:

10

2

2

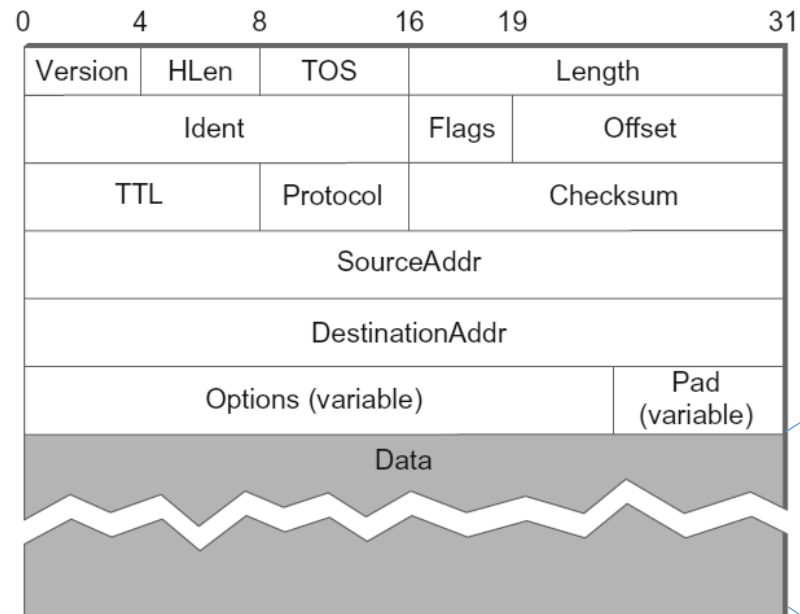
 Rx:

7

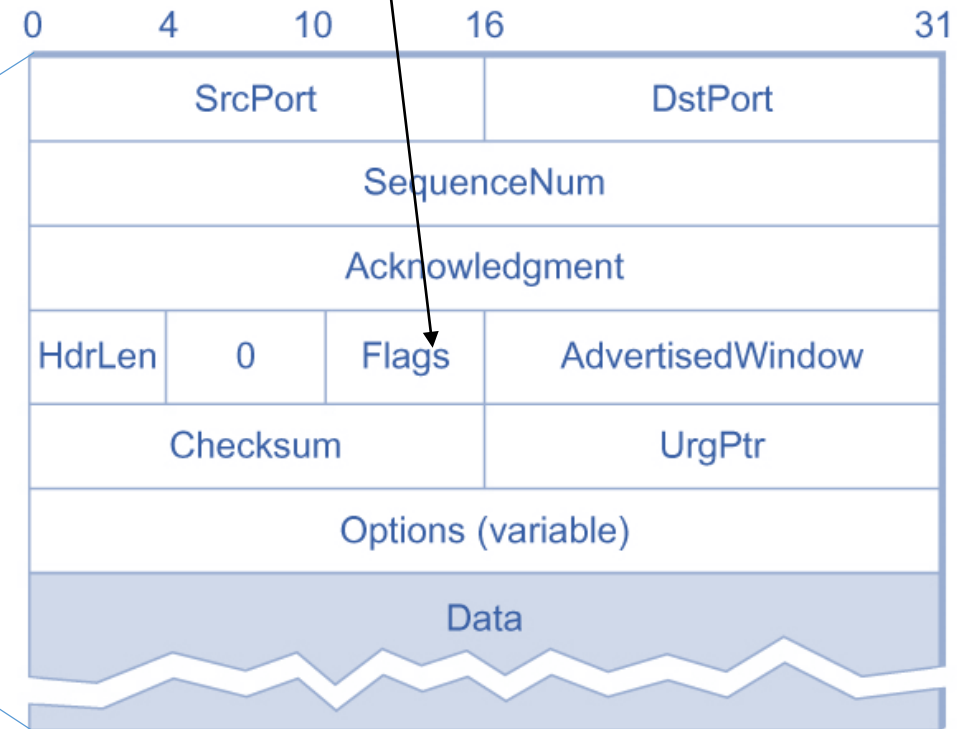
7

- Solution
 - Sentinel-based Checking
 - URG flag
 - PUSH flag
- } TCP will inform the application that a URG/PUSH segment is received

TCP: Header

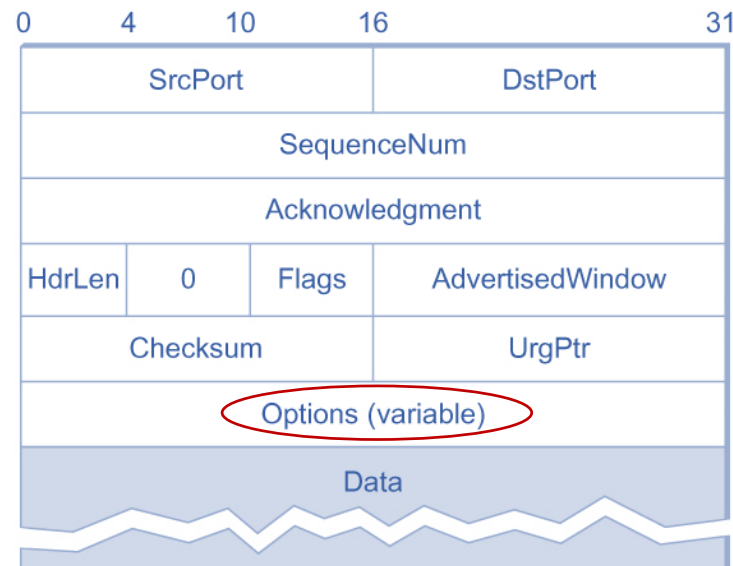


URG|ACK|PUSH|RESET|SYN|FIN



TCP Extensions

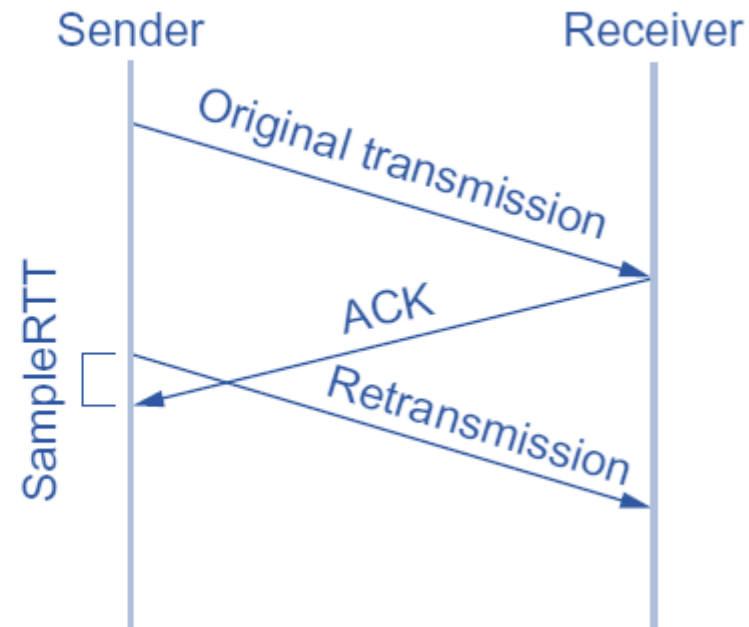

- TCP keeps evolving to incorporate high performance network



Sliding Window in TCP: Adaptive Timeout


- Original Algorithm

- Measure
- Compute
- Set



TCP Extensions

- Better Estimation of RTT
 - Introduce 32-bit timestamp



TCP Extensions

- Sequence Number Wraparound: sequence number can be easily used up in high speed network (too many bytes are out per unit time)
 - Use Seq + Timestamp to identify wraparound

Table 5.1 Time Until 32-Bit Sequence Number Space Wraps Around

Bandwidth	Time until Wraparound
T1 (1.5 Mbps)	6.4 hours
Ethernet (10 Mbps)	57 minutes
T3 (45 Mbps)	13 minutes
Fast Ethernet (100 Mbps)	6 minutes
OC-3 (155 Mbps)	4 minutes
OC-12 (622 Mbps)	55 seconds
OC-48 (2.5 Gbps)	14 seconds

TCP Extensions

- Keeping the Pip Full: Advertised Window is not large enough for the Rx buffer
 - Advertised Window (16bits) + Scaling Factor (0 to 14)
 - Negotiated in SYN packet

Table 5.2 Required Window Size for 100-ms RTT

Bandwidth	Delay × Bandwidth Product
T1 (1.5 Mbps)	18 KB
Ethernet (10 Mbps)	122 KB
T3 (45 Mbps)	549 KB
Fast Ethernet (100 Mbps)	1.2 MB
OC-3 (155 Mbps)	1.8 MB
OC-12 (622 Mbps)	7.4 MB
OC-48 (2.5 Gbps)	29.6 MB

TCP Extensions

- Selective ACK

Demo

- TCP Handshake
 - filter: tcp.stream eq X
- TCP Sliding Window
- TCP Extensions

TCP: Design Summary

- Stream-oriented v.s. Request/reply



- Byte-stream v.s. Message
 - Arbitrary long message v.s. Limited message
- In Order v.s. Out Order



- Explicit Setup v.s. Implicit Setup
- Window-based v.s. Rate-based

Reference

- Textbook 5.2
- TCP Extensions <https://www.ietf.org/rfc/rfc1323.txt>