



# CS120: Computer Networks

## **Lecture 18. Congestion Control 2**

Zhice Yang

# Congestion Control

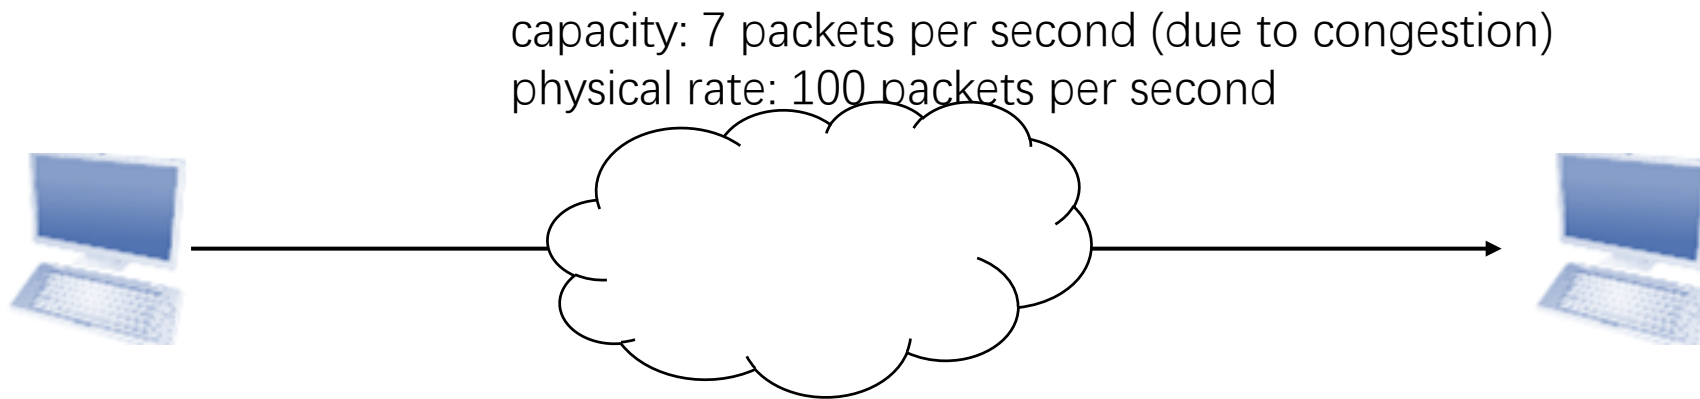
- Queuing
- Connection Control Methods
  - Congestion Control
  - Congestion Avoidance
- QoS

# TCP Congestion Control


- Introduced by Van Jacobson through his Ph.D. dissertation work in late 1980s
  - 8 years after TCP became operational
- Basic ideas
  - Each host determines network capacity for itself
    - Leverage feedback
    - Assumption: FIFO or FQ queue in routers
- Challenges
  - Determining the available capacity
  - Adjusting to changes in capacity

# Congestion Control with Sliding Window

- Size of sliding window is determined by network capacity:  
 $\text{delay} \times \text{bandwidth}$ 
  - Problem
    - bandwidth is unknown
    - How to pace the sender ?




# Congestion Control in Road



How to determine the road capacity ?

# Congestion Control in Road




Determine Road Capacity according to the Reverse Traffic

# Congestion Control in Road



Determine Road Capacity according to the Reverse Traffic

# Congestion Control in Road

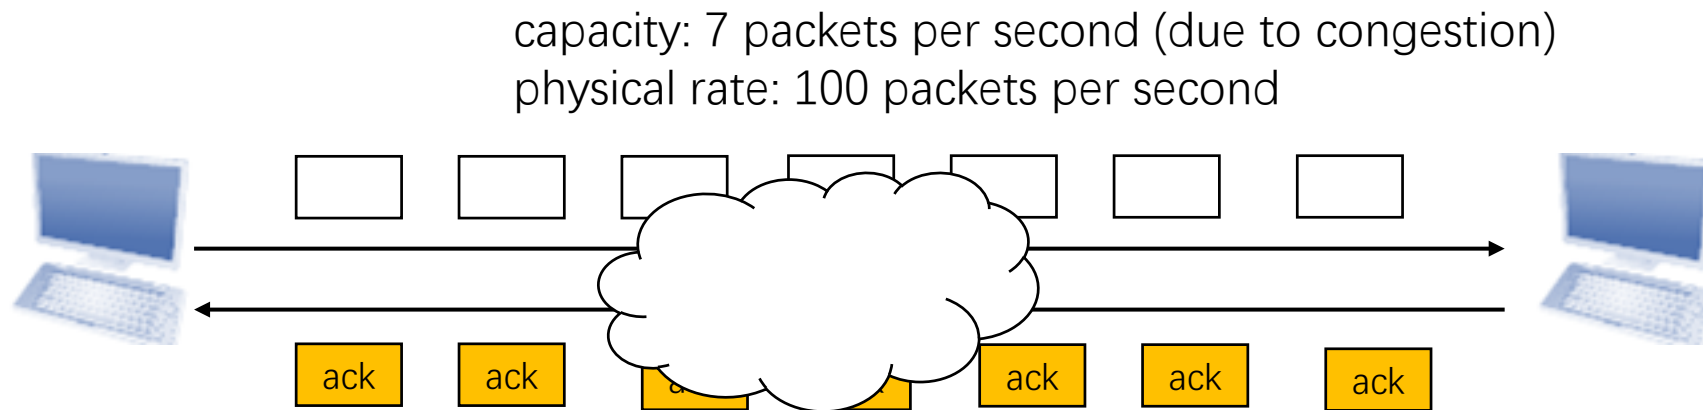


Determine Road Capacity according to the Reverse Traffic



# Congestion Control with Sliding Window

- Size of sliding window is determined by network capacity:  
 $\text{delay} \times \text{bandwidth}$ 
  - Problem
    - bandwidth is unknown
    - How to pace the sender ?



Use ACK to estimate congestion and pace the sending

# TCP Congestion Control


- Objective: Estimate and adapt to (varying) network capacity
- Approach: Adjust Sliding Window
  - $\text{MaxWindow} = \text{MIN}(\text{CongestionWindow}, \text{AdvertisedWindow})$
  - Decrease **CongestionWindow** upon detecting congestion
  - Increase **CongestionWindow** upon lack of congestion
- Basic Components
  - Additive Increase/Multiplicative Decrease (AIMD)
  - Slow Start
  - Fast Retransmission
  - Fast Recovery
- Other Variations

# Additive Increase/Multiplicative Decrease (AIMD)

- Intuition: over-sized window is much worse than an under-sized window
  - Over-sized window: packets dropped and retransmitted
  - Under-sized window: somewhat lower throughput
- Additive Increase
  - If successfully received acks of the **last window** of data
    - $\text{CongestionWindow} = \text{CongestionWindow} + 1$
- Multiplicative Decrease
  - If packet loss
    - $\text{CongestionWindow} = \text{CongestionWindow} / 2$

# AIMD

- TCP sawtooth pattern



# Sliding Window in TCP: Adaptive Timeout


- Jacobson/Karels Algorithm Implementation

Difference = SampleRTT - EstimatedRTT

EstimatedRTT = EstimatedRTT + ( $\delta$ \*Difference)

Deviation = Deviation +  $\delta$ \*(|Difference| - Deviation)

TimeOut =  $\mu$  \* EstimatedRTT +  $\varphi$  \* Deviation



```
SampleRTT -= (EstimatedRTT >> 3);
```

```
EstimatedRTT += SampleRTT;
```

```
if (SampleRTT < 0)
```

```
    SampleRTT = -SampleRTT;
```


```
SampleRTT -= (Deviation >> 3);
```

```
Deviation += SampleRTT;
```

```
TimeOut = (EstimatedRTT >> 3) + (Deviation >> 1);
```


# AIMD

- Additive Increase
  - Increment =  $\text{MSS} / \text{CongestionWindow}$
  - $\text{CongestionWindow} += \text{Increment}$




# AIMD

- Multiplicative Decrease
  - $\text{CongestionWindow} = \text{CongestionWindow} / 2$



# AIMD


- TCP sawtooth pattern






# AIMD

- TCP sawtooth pattern




# AIMD

- TCP sawtooth pattern




No ACKs to guide sending;  
Better start from cwnd = 1



# AIMD

- TCP sawtooth pattern




# Slow Start

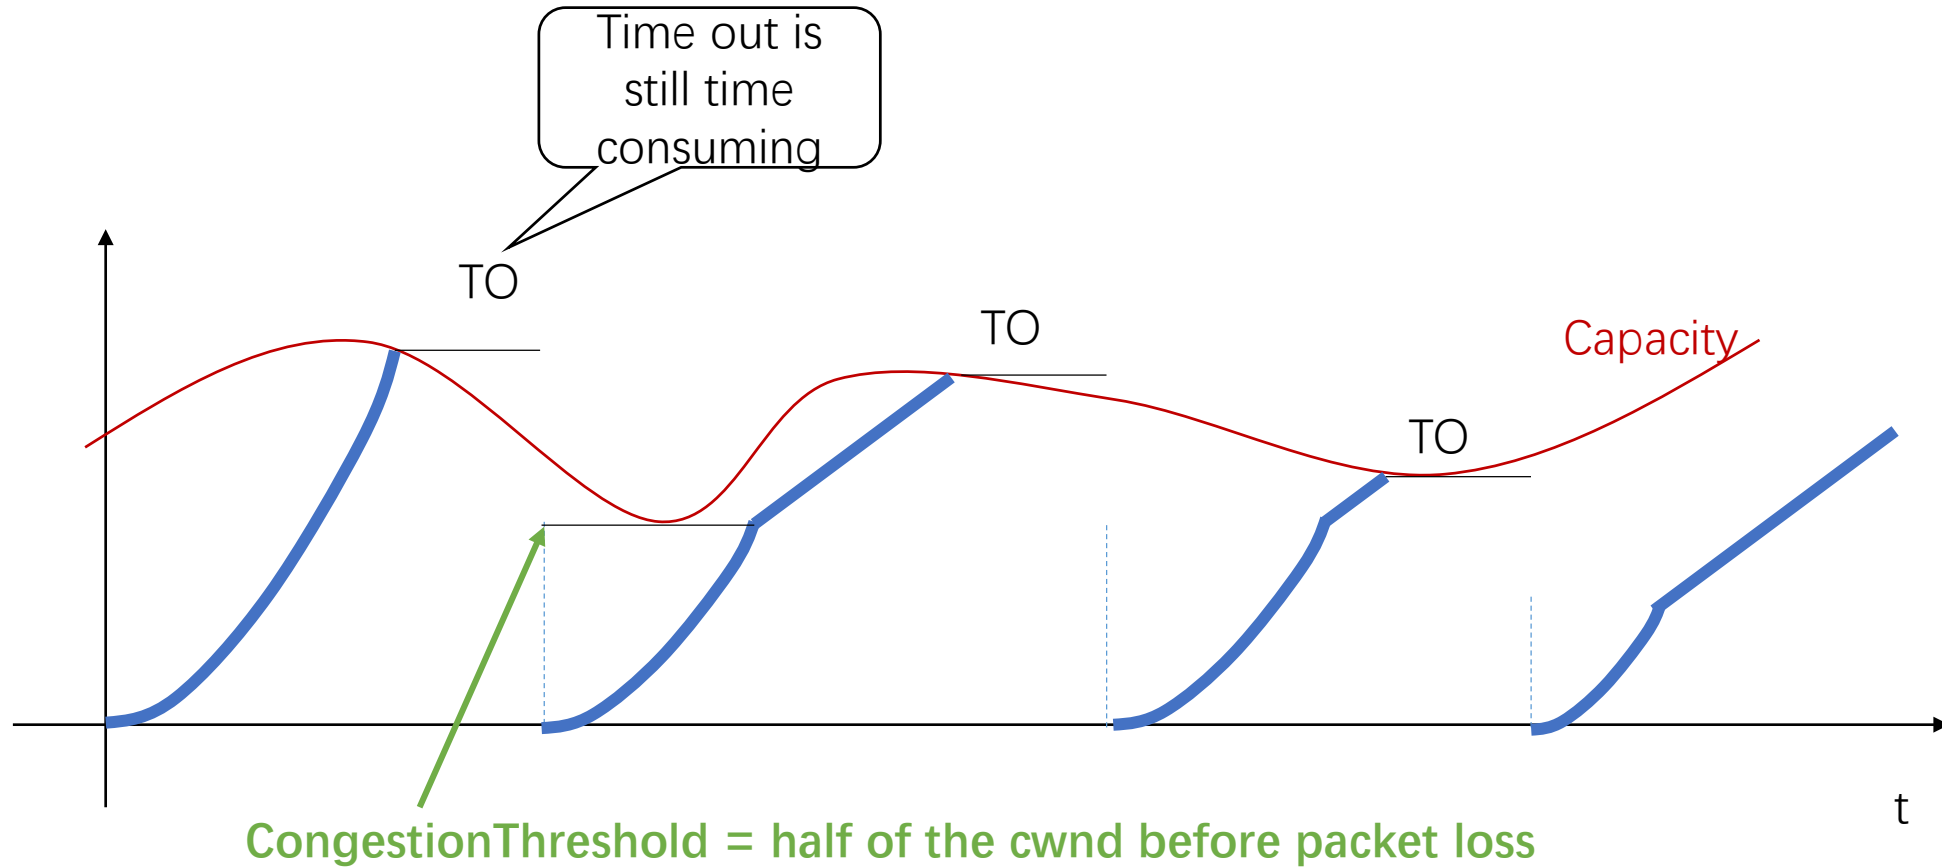
- Intuition: speed up additive Increase when TCP start
- Why “Slow Start”
  - “Slow Start” is not slow compared with additive Increase
  - “Slow Start” is slow compared with sending a whole window’s worth of data (Original TCP)
- **Double** CongestionWindow per round-trip time
  - If successfully received **one ack**
    - $\text{CongestionWindow} = \text{CongestionWindow} + 1$
  - Until  $\text{CongestionWindow} == \text{CongestionThreshold}$
  - Then do Additive Increase

# Slow Start

- If successfully received one ack
  - $\text{CongestionWindow} = \text{CongestionWindow} + 1$




# Slow Start



# Fast Retransmission


- Intuition: use **duplicate ACK** to indicate packet loss
- Approach:
  - Receiver replies every TCP segment with acknum = next byte expected
  - Transmitter resends a segment after 3 duplicate acks
    - 3 duplicate acks => possible packet loss
- Throughput Gain: 20%



# Fast Retransmission

Timeout still exists

- Too many packet loss
- Window may be too small to generate enough duplicate acks




TCP Tahoe



# Fast Retransmission

Timeout still exists


- Too many packet loss
- Window may be too small to generate enough duplicate acks



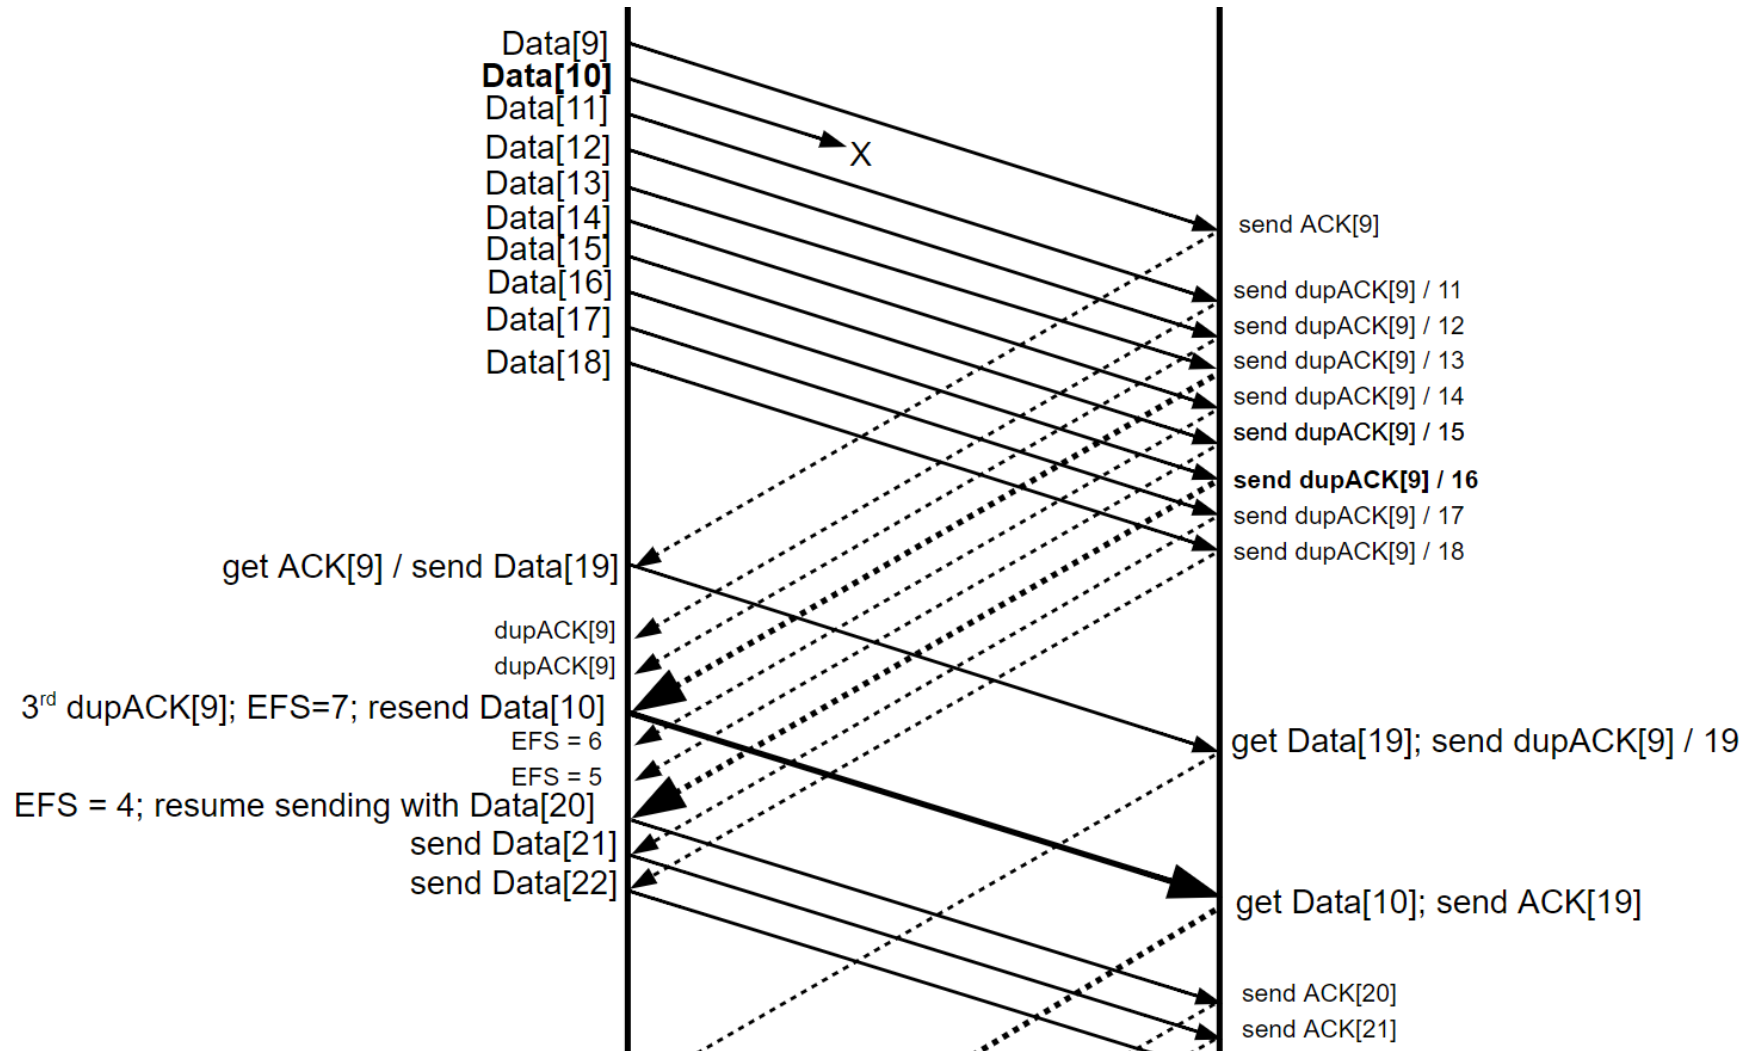
TCP Tahoe

# Fast Recovery


- Intuition:
  - Flying acks can be used as clock
    - No need to start from window size 1



# Fast Recovery



# Fast Recovery



TCP Reno

# TCP Congestion in Wireless

- Challenges
  - Timeout doesn't mean congestions (with very high probability)
    - Reason: wireless channel is not reliable
- Possible Solutions
  - Error Correction
    - Additional traffic overhead
  - MAC layer retransmission (WiFi)
    - Large End-to-end RTT variance

# Demo

- [http://leeo1116.github.io/TCP\\_congestion\\_control/TCP.html](http://leeo1116.github.io/TCP_congestion_control/TCP.html)

# Reference

- Textbook 6.3
- <http://intronetworks.cs.luc.edu/current/html/reno.html>