



# CS120: Computer Networks

## **Lecture 19. Other Topics in Transportation Layer**


Zhice Yang

# Outline

- TCP Fairness
- QUIC
- QoS

# Evaluation Criteria

- Defining fairness is hard
  - In terms of a host, a TCP link, or an application ?
- TCP fairness goal: if  $n$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/n$



- Fairness Index



$$f(x_1 \dots x_n) = \frac{(\sum x_i)^2}{n * \sum x_i^2}$$

# Fairness in TCP

- Consider the steady state, TCP uses a (linear) scheme to adjust its window cwnd
  - $\text{cwnd}' = b * \text{cwnd} + a$
- Possible Designs
  - Additive increase, additive decrease
  - Additive increase, multiplicative decrease (AIMD)
  - Multiplicative increase, additive decrease
  - Multiplicative increase, multiplicative decrease


# Fairness in TCP

- Consider a case with two TCP connections



# Fairness in TCP


- Consider a case with two TCP connections



Increase the throughput to  
utilize more bandwidth


# Fairness in TCP

- Consider a case with two TCP connections



# Fairness in TCP


- Consider a case with two TCP connections






# Fairness in TCP

- Consider a case with two TCP connections




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of additive increase additive decrease
    - Stable but not fair




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of multiplicative increase multiplicative decrease
    - Stable but not fair




# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of multiplicative increase additive decrease
    - Not stable



# Fairness in TCP

- Consider a case with two TCP connections
  - Behavior of AIMD
    - Stable and faire



# Fairness and RTT

- TCP connection with smaller RTT occupies more bandwidth
  - When congestion happens, they recover more quickly
    - TCP adjust cwnd in RTT basis

# Fairness and Parallel TCP Connections

- Application can open multiple parallel connections between two hosts
  - web browsers do this , e.g., link of rate  $R$  with 9 existing connections:
    - new app asks for 1 TCP, gets rate  $R/10$
    - new app asks for 11 TCPs, gets  $R/2$

# Fairness and UDP

- Multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- Instead use UDP:
  - send audio/video at constant rate, tolerate packet loss
- There is no “Internet police” policing use of congestion control



# Outline

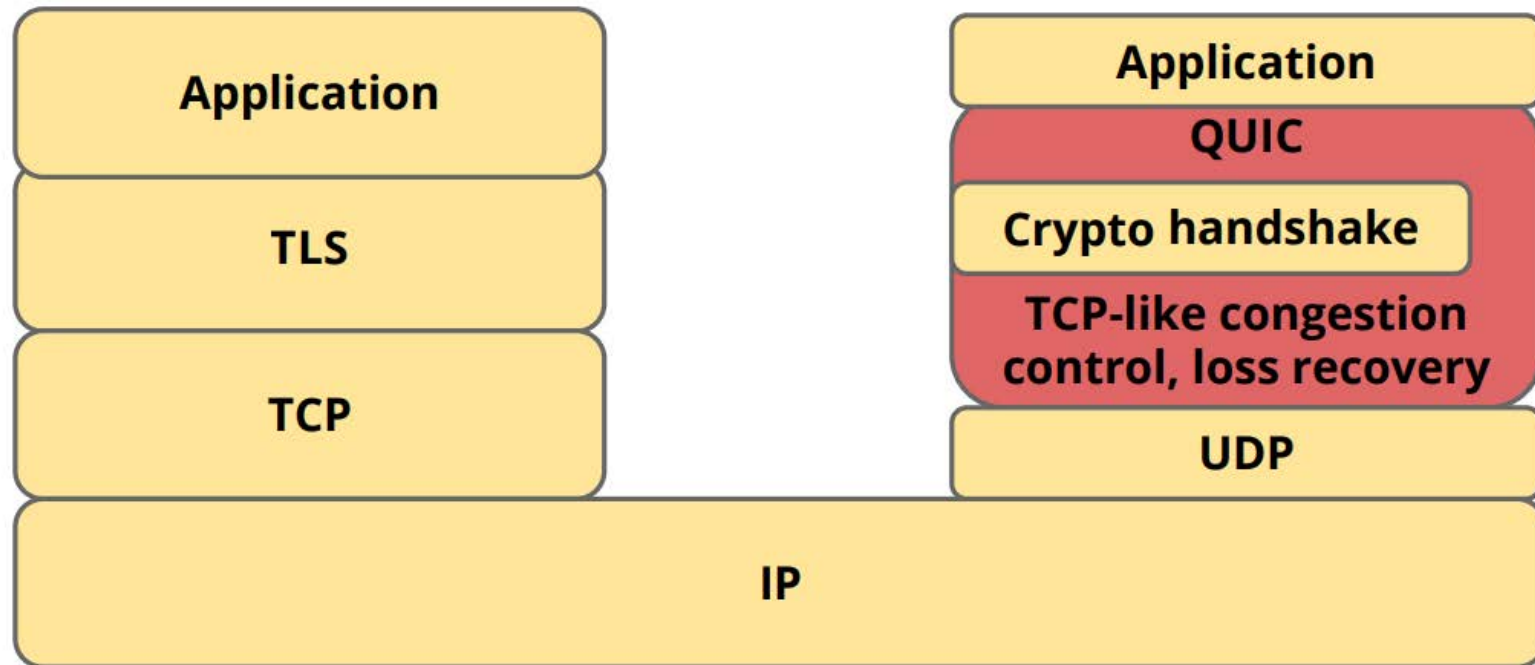
- TCP Fairness
  - QUIC
- QoS

# QUIC

- QUIC: Quick UDP Internet Connections
- Application-layer protocol, on top of UDP
  - Deployed by Google starting at 2014
    - Deployed on many Google servers, apps (Chrome, mobile YouTube app)
  - QUIC working group formed in Oct 2016
- Initial goal: increase performance of HTTP

# QUIC

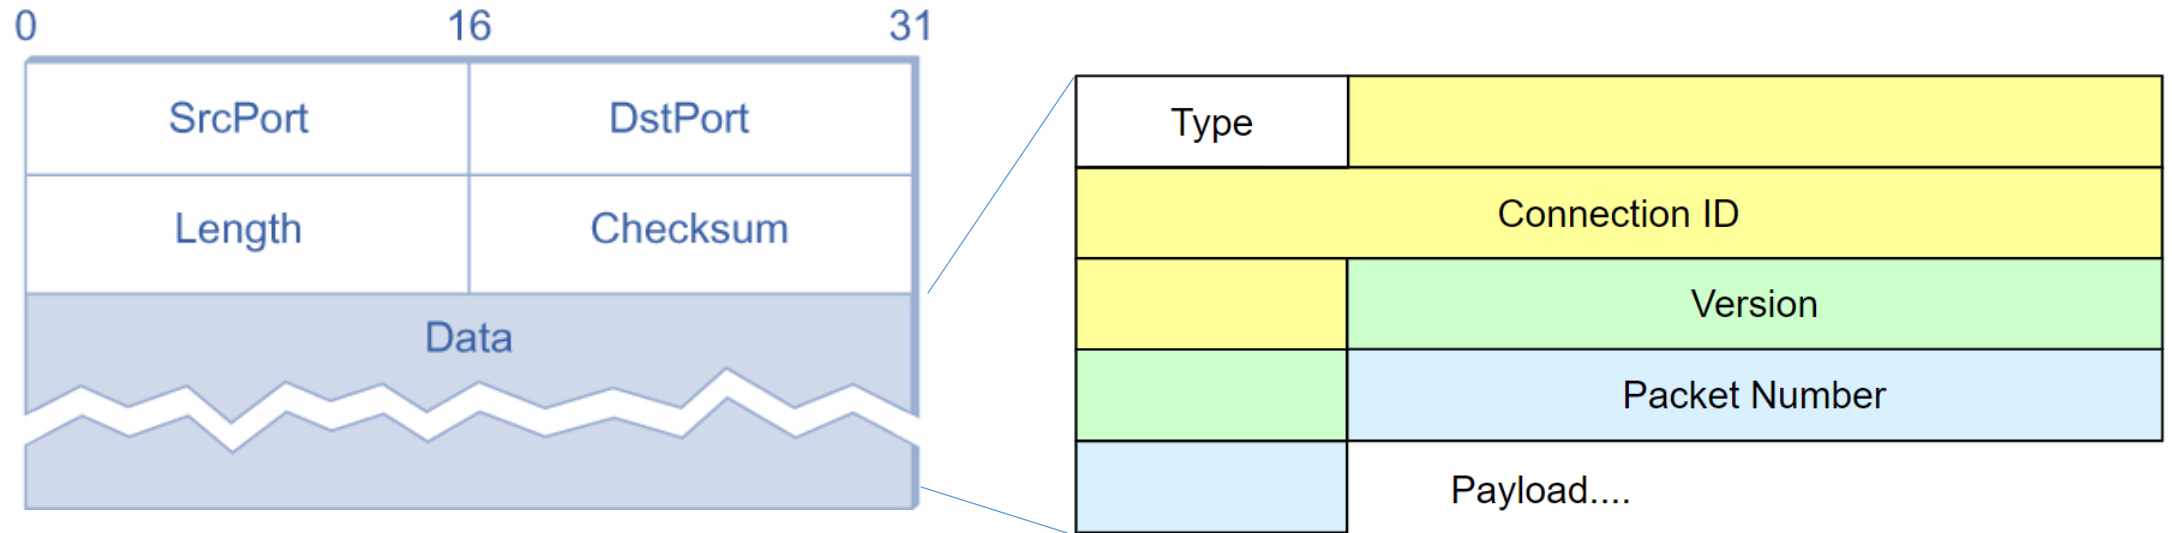
- Protocol Stack




# QUIC

- Key features
  - Always encrypted
  - 0-RTT connection establishment
  - Connection migration
  - Congestion control
  - Parallel Streams


# QUIC - Header




# QUIC Connection Establishment



TCP  
(2RTT)




TCP+TLS 1.2  
(new 4RTT  
resumed 3RTT )



QUIC  
(new 2RTT  
Resumed 1RTT)


# QUIC Connection Establishment

- 1-RTT (First-ever connection)
  - No cached information available
  - First CHLO is inchoate (empty)
    - Simply includes version and server name
  - Server responds with REJ
    - Includes server config, certs, etc.
    - Allows client to make forward progress
  - Second CHLO is complete
    - Followed by initially encrypted request data
  - Server responds with SHLO
    - Followed immediately by forward-secure encrypted response data



# QUIC Connection Establishment

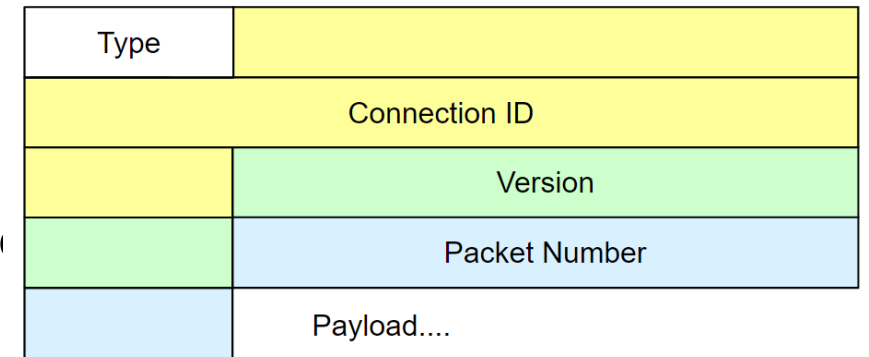
- 0-RTT (Subsequent connection)
  - Motivation: client can cache information about the *origin* it connected to
  - First CHLO is complete
    - Based on information from previous connection
    - Followed by initially encrypted data.
  - Server responds with SHLO
    - Followed immediately by forward-secure encrypted data





# QUIC Connection Migration

- NAT Rebinding
  - NATs remaps port
    - Frequency (~ mins)
    - Why ? to release unused ports
      - According to TCP connection state (if they are closed)
    - UDP does not have connection state, QUIC state is encrypted
- Mobility
  - Switching between different IP
    - Wi-Fi and cellular network
- Connection Migration
  - Keep QUIC connections alive even if port and IP are changing
  - Detect connection path changes via Connection ID and IP/port
    - Connection is identified by connection ID rather than <IP, port>
    - 64-bit connection ID
    - randomly chosen by client



# QUIC Congestion Control

- Incorporates TCP best practices
  - TCP Cubic, Fast Retransmission, Selective ACK, etc.
- Better signaling than TCP
  - Each packet carries a monotonically increasing packet number
    - Better RTT measurement
    - Even ACK
  - Retransmitted packets also consume new sequence numbers
    - no retransmission ambiguity
- More verbose ACK
  - support 256 NAK ranges (vs. TCP's 3SACK ranges)

# QUIC - Parallel Streams

- Handle HOL blocking

# Outline

- TCP Fairness
- QUIC
- QoS