



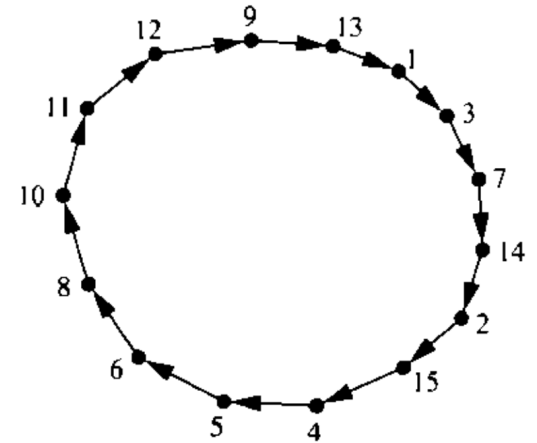
# PRAM 2

## Graph algorithms

CS121 Parallel Computing  
Fall 2021

# Coloring a cycle

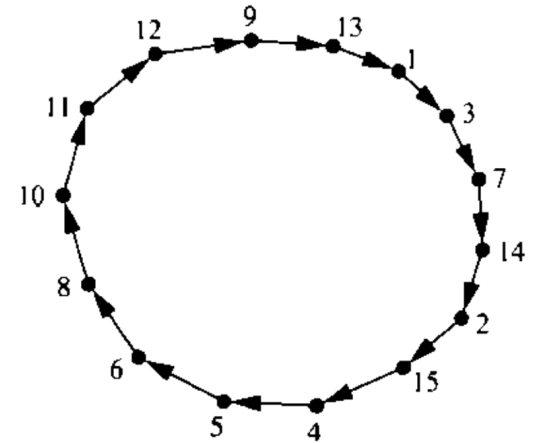
- Given a graph  $G=(V,E)$ , a  $k$ -coloring of  $G$  is a mapping  $c: V \rightarrow \{0,1, \dots, k-1\}$  s.t.  $c(i) \neq c(j)$  whenever  $(i,j) \in E$ .
- We give a super fast algorithm for 3-coloring a (directed) cycle of  $n$  nodes.
  - If  $n$  is odd, any coloring uses at least 3 colors.
- Coloring the cycle is a form of symmetry breaking.
- For any node  $v$ , let  $S(v)$  be the node after  $v$ .
- The main subroutine is the following.
  - Initially, color every node by its node ID.
  - Consider the binary representation of  $c(v)$  for a node  $v$ .
  - Let  $k$  be the least significant digit in which  $c(v)$  and  $c(S(v))$  differ.
  - Set  $c'(v)=2k+c(v)_k$ , where  $c(v)_k$  is the  $k$ 'th digit of  $c(v)$ .



$v$	$c$	$k$	$c'$
1	0001	1	2
3	0011	2	4
7	0111	0	1
14	1110	2	5
2	0010	0	0
15	1111	0	1
4	0100	0	0
5	0101	0	1
6	0110	1	3
8	1000	1	2
10	1010	0	0
11	1011	0	1
12	1100	0	0
9	1001	2	4
13	1101	2	5

# Coloring a cycle

- **Claim** If  $c$  is a valid coloring, then so is  $c'$ .
- **Proof** Since  $c$  is a valid coloring, then  $c(v) \neq c(S(v))$ , so  $k$  exists.
  - Suppose  $c'(v) = c'(u)$ , for some  $v$  and  $u = S(v)$ .
  - Then  $c'(v) = 2k + c(v)_k$  and  $c'(u) = 2l + c(u)_l$  for some  $k$  and  $l$ .
  - Since  $c'(v) = c'(u)$ , then  $k = l$ , because  $c(v)_k, c(u)_l < 2$ .
  - But then  $c(v)_k = c(u)_k$ , contradicting the definition of  $k$ .



$v$	$c$	$k$	$c'$
1	0001	1	2
3	0011	2	4
7	0111	0	1
14	1110	2	5
2	0010	0	0
15	1111	0	1
4	0100	0	0
5	0101	0	1
6	0110	1	3
8	1000	1	2
10	1010	0	0
11	1011	0	1
12	1100	0	0
9	1001	2	4
13	1101	2	5



# Coloring a cycle

- To analyze the time complexity, suppose in some round the max number of bits to represent any color is  $t$ .
- Then the max number of bits to represent any color in the next round is  $\lceil \log t \rceil + 1$ , because any color in the next round is  $\leq 2t + 1$ .
  - So the number of bits used to represent a color decreases from  $t$  to  $\lceil \log t \rceil + 1$  in each round.
- Let  $\log^{(i)} x = \log(\log^{(i-1)} x)$ , i.e. we apply the log function  $i$  times to  $x$ .
- Let  $\log^* x = \min\{i \mid \log^{(i)} x \leq 1\}$  be the number of times we have to take log's until a value becomes  $\leq 1$ .
  - $\log^* x$  is incredibly small. In fact,  $\log^* x \leq 6$  for all  $x \leq 2^{65536}!$




# Coloring a cycle

- Since the number of bits to represent a color in the first round is  $\log n$ , then in  $O(\log^* n)$  rounds, we can represent any color using  $O(1)$  bits.
  - In fact, we can apply the subroutine until we use 6 colors in a round.
  - With 6 colors, need 3 bits to represent a color. So in the next round, colors are between 0 and  $2^3-1=7$ , and we again use up to 6 colors.
- To decrease the number of colors from 6 to 3, we run 3 more rounds.
  - In round  $i$ , take any node colored using color  $i+2$  and color it using the min possible color in  $\{0,1,2\}$ , i.e. the min color not used by its neighbors.
- In total, we 3-color the ring in  $O(\log^* n)$  rounds, using  $O(n \log^* n)$  work.
- The algorithm can be modified to produce a 3-coloring in  $O(\log n)$  time and  $O(n)$  work.

# Independent set on line

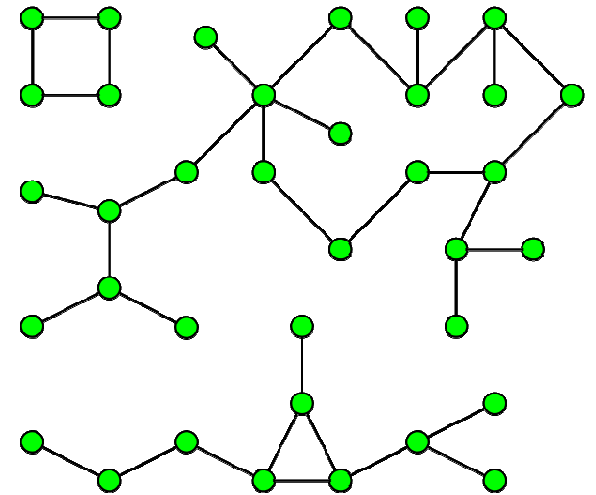
- **Thm** Given a  $k$  coloring of a line graph with  $n$  nodes, we can compute an independent set of size  $\Omega(n/k)$  in  $O(1)$  time.
- **Proof** Every node has a color from 1 to  $k$ .
  - Take the nodes whose colors are local minima as the indep. set  $S$ .
    - No two nodes in  $S$  are neighbors.
    - $S$  can be computed in  $O(1)$  time.
  - Consider two consecutive nodes  $u, v \in S$ .
  - Since  $u, v$  are local minima and consecutive, the colors between  $u, v$  first increase, then decrease.
  - Thus, there are at most  $2k - 3$  nodes between  $u$  and  $v$ .
  - So there are  $\leq 2k - 3$  nodes between any two consecutive nodes in  $S$ , and so  $|S| \geq \frac{n}{2k-3}$ .
- Thus, compute independent set of size  $\Omega(n)$  on line with  $n$  nodes in  $O(\log n)$  time by computing a 3 coloring of the line in  $O(\log n)$  time, then computing an independent set of size  $\Omega(n/3) = \Omega(n)$  in  $O(1)$  time.



- The colors of the nodes are shown in parentheses.
- Nodes 4, 3, 2, 5 are local minima.

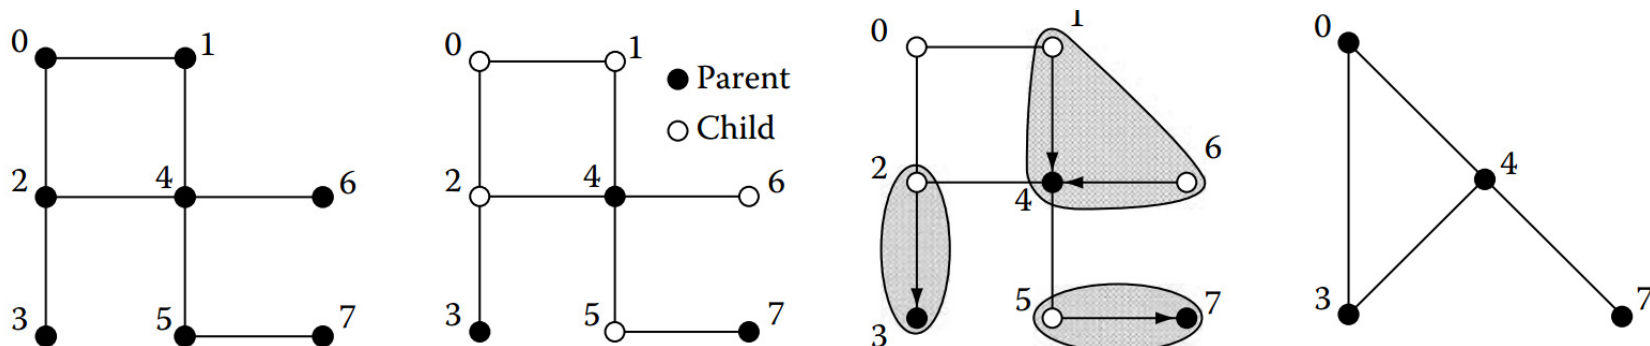
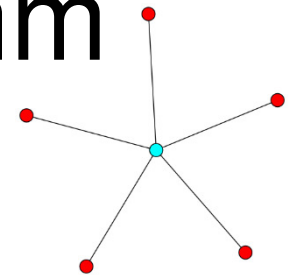
# Connected components

- Given an undirected graph, partition it into maximal sets of nodes that are connected to each other.
- Can be solved sequentially in  $O(m+n)$  time using BFS / DFS.
  - $m$  is number of edges,  $n$  is number of vertices.
- However, no efficient BFS / DFS PRAM algorithms known.
- Instead, use graph contractions.
  - In each phase, merge (contract) a set of connected nodes into a supernode.
  - Form a contracted graph on the supernodes, then apply algorithm recursively.
  - Eventually each connected component is contracted to one node.
  - Many different algorithms, depending on which nodes they contract.



# Randomized parallel algorithm

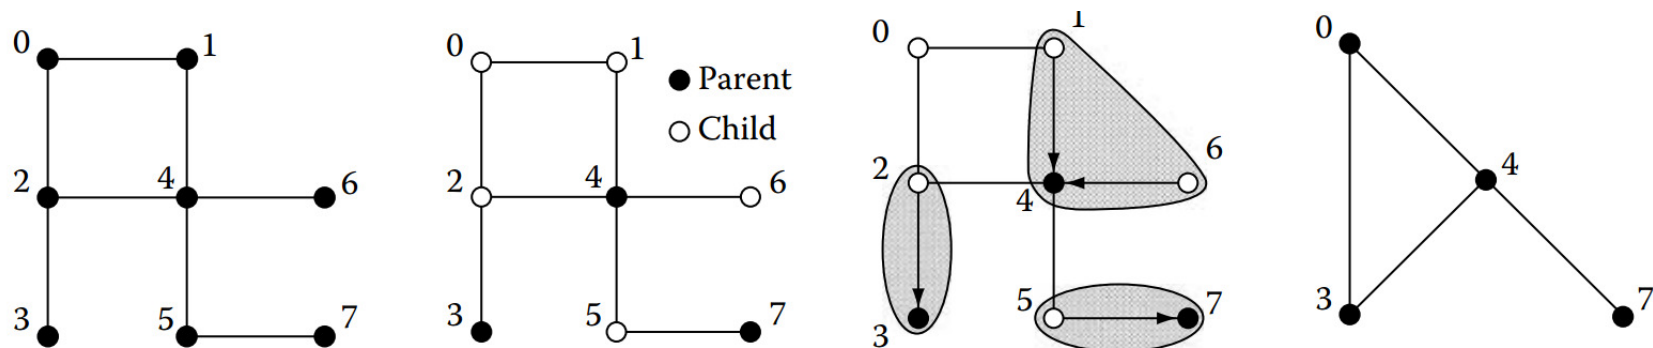
- Break graph into star graphs and contract each star.
- In each phase, do the following steps in parallel.
  - Every node flips a coin and chooses to be a parent or child node.
  - Each child node points to a parent node it's connected to.
    - Now have a set of stars, with the parent nodes as the centers.
    - If child not connected to any parent node, it forms its own star.
  - Contract each star to its center, then apply algorithm recursively.
    - Label all nodes in star by the label of the parent.
    - Keep the edges between differently labeled nodes.
  - After recursion returns, each child with a parent again takes parent's label, which might have changed.






# Implementation

- Use  $n+m$  processors, one processor for each vertex and edge.
  - Call these V and E procs, resp.
  - Each V proc has a label which can change over time.
  - Each E proc responsible for edge  $(u,v)$ , where  $u, v$  are V procs.
  - V and E procs may become inactive over time.
- In each phase, each active V proc flips a coin to decide if it's a child or parent proc.
- Each active E proc  $(u,v)$  checks if  $u$  is a child proc and  $v$  is a parent (or vice versa).
  - If so, it sets  $u$ 's label to  $v$  (or  $v$ 's label to  $u$ ).
  - Another E proc  $(u,v')$  could set  $u$ 's label to  $v'$ . In this case, either the  $v$  or  $v'$  write succeeds.



# Implementation

- Each parent V proc, or child V proc whose label didn't change (i.e. it had no parent), stays active.
  - Other V procs become inactive.
- Each active E proc  $(u,v)$  where  $u, v$  have different labels stays active.
  - Other E procs become inactive.
  - From now on, E will be responsible for V processes  $(u',v')$ , where  $u'$  and  $v'$  are the labels of  $u$  and  $v$ , resp.
- The active V and E procs run the algorithm recursively.
- After recursion returns, inactive E procs  $(u,v)$  (where  $u$  is the child) write  $v$ 's label to  $u$ .
- At end, all V procs in a connected component have same label.



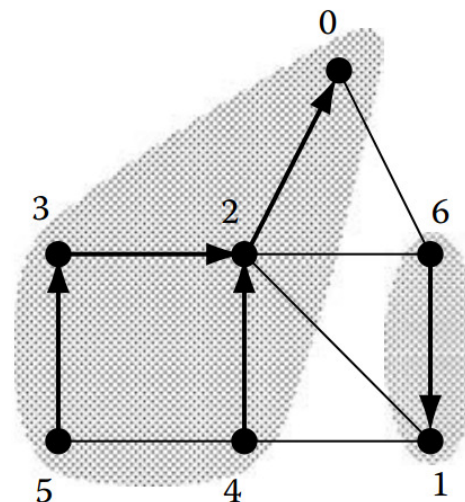
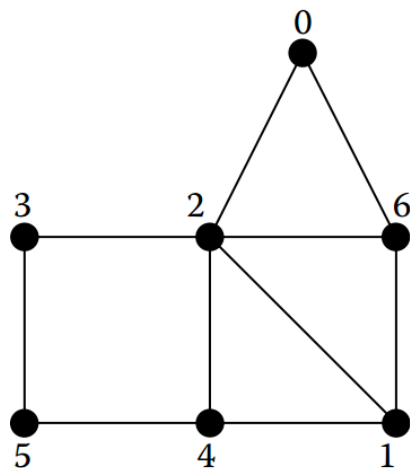


# Complexity

- Key fact is that number of active  $V$  procs decreases by  $1/4$  fraction in expectation each phase.
  - A  $V$  proc becomes inactive if it's a child node and one its neighbors is a parent node.
  - The former probability is  $1/2$ , and the latter is  $\geq 1/2$ .
  - Thus each  $V$  proc becomes inactive with probability  $\geq 1/4$ .
- With high probability, after  $O(\log n)$  phases, there's only one  $V$  proc and recursion ends.
- Each phase takes  $O(1)$  time, and does  $O(m+n)$  work.
- Total time is  $O(\log n)$ , total work is  $O((m+n) \log n)$ .
  - This algorithm isn't work efficient.
  - There exist work efficient randomized CC algorithms.

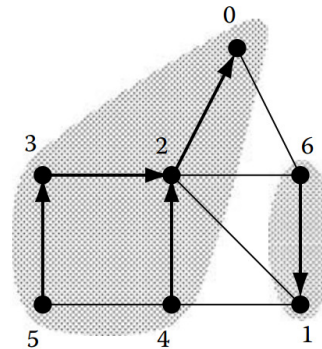
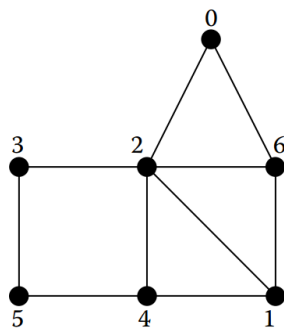
# Deterministic parallel algorithm

- Again work in phases, with following parallel steps.
  - Each node points to a neighbor with lower ID.
  - This breaks graph into a directed forest.
  - Contract each forest to the lowest ID node using pointer jumping.
  - Recurse on contracted graph.






# Implementation

- As before, use  $n+m$  procs, for  $V$  and  $E$ .
  - But now,  $V$  procs always active.  $E$  procs may become inactive.
- Each  $E$  proc  $(u,v)$  checks if  $u < v$ , and if so sets  $v$ 's label to  $u$ .
  - Again, conflicts resolved arbitrarily.
- $V$  procs then apply pointer jumping on the labels, taking the label of the proc it points to.
- Each active  $E$  proc  $(u,v)$  where  $u, v$  have different labels stays active. Other  $E$  procs become inactive.
- The  $V$  procs and active  $E$  procs run algorithm recursively.
  - Note that in recursive call, all  $V$  procs apply pointer jumping.



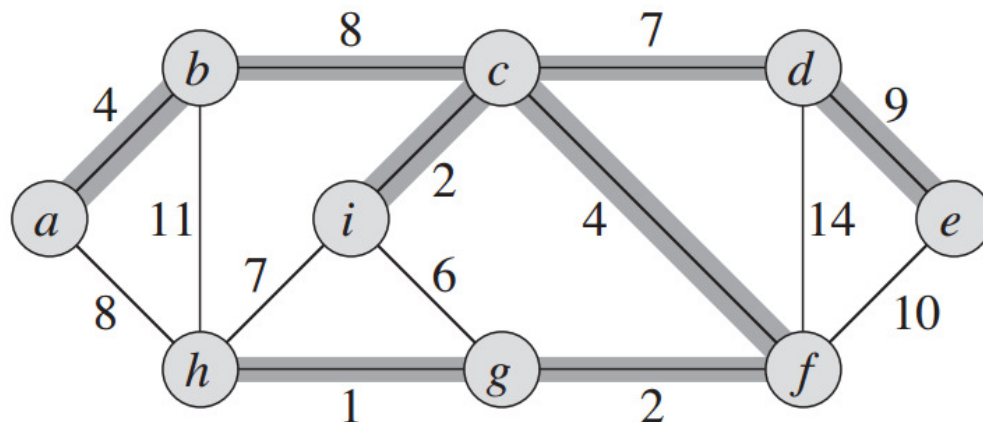
# Complexity

- The basic algorithm may take  $O(n)$  time on the graph below.
- But notice that if we made nodes point to higher neighbors, the graph would be solved in  $O(1)$  time.
- In each phase, if we consider either having nodes point to smaller neighbors, or pointing to higher neighbors.
  - If  $v$  doesn't point to any nbr in "high $\rightarrow$ low" round, then it's smaller than all nbrs. So in "low $\rightarrow$ high" round, it points to some nbr.
  - In one of these cases,  $\geq n/2$  nodes point to another node, and are contracted.
- Thus the algorithm finishes in  $O(\log n)$  phases.
- Each phase does pointer jumping, using  $O(\log n)$  time and  $O(n)$  work.
- Total time is  $O(\log^2 n)$ , and work is  $O((m+n) \log n)$ .



# Minimum spanning tree

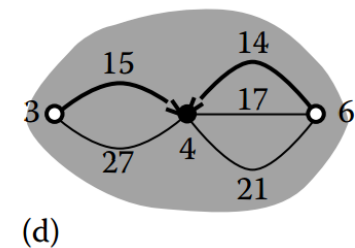
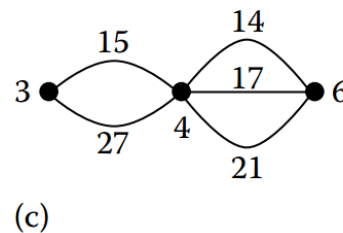
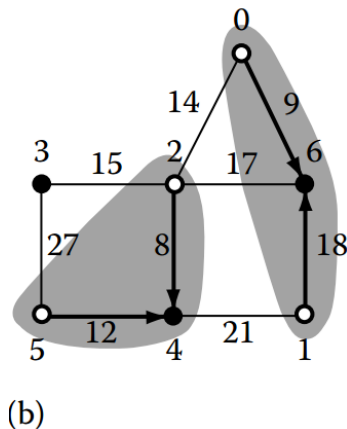

- Given an undirected graph with edge weights, an MST is a connected subgraph containing all the vertices, which has minimum total weight.
- Can be solved in  $O(m + n \log n)$  time sequentially by a greedy algorithm.
- Key property is that for any set of vertices  $W$ , the minimum cost edge from  $W$  to  $V \setminus W$  is in the MST.
  - So for any vertex  $v$ , min cost edge containing  $v$  is in MST.
- Will describe a parallel MST algorithm based on the randomized parallel algorithm for connected component.



Source: Introduction to Algorithms, Cormen et al.

# Randomized parallel algorithm

- Each node randomly chooses to be a parent or child node.
- Each child node  $u$  finds min weight incident edge  $(u,v)$ , and points to  $v$  if  $v$  is parent.
  - This forms a set of stars with parents as centers.
  - If  $v$  isn't a parent,  $u$  forms its own star.
- Contract each star to the parent, and run algorithm recursively.
- How do we find min weight incident edge in  $O(1)$  time?
  - One possibility is to use priority CRCW.
  - Presort the edges by nondecreasing weight. Each edge processor's priority is its value.
  - So when each  $E$  proc  $(u,v)$  writes to  $u$ , min weight edge wins.







# Complexity

- At least  $1/4$  of vertex processors become inactive each phase in expectation.
  - Given a node  $u$  and min weight edge  $(u,v)$ , there's  $1/4$  probability  $u$  is child and  $v$  is parent.
- Thus, there are  $O(\log n)$  phases with high probability.
  - Finding min weight incident edge takes  $O(1)$  time after presorting edge weights.
- Total time is  $O(\log n)$ , total work is  $O((m+n) \log n)$ .