

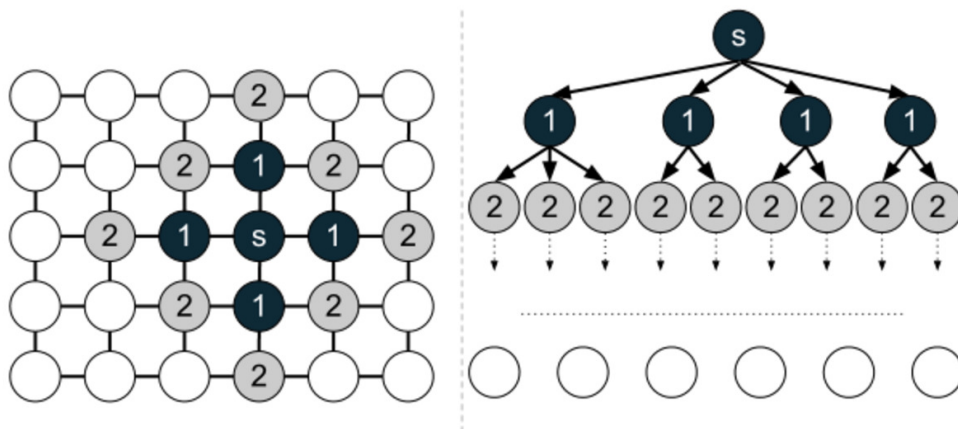


# CUDA 6 Breadth-First Search

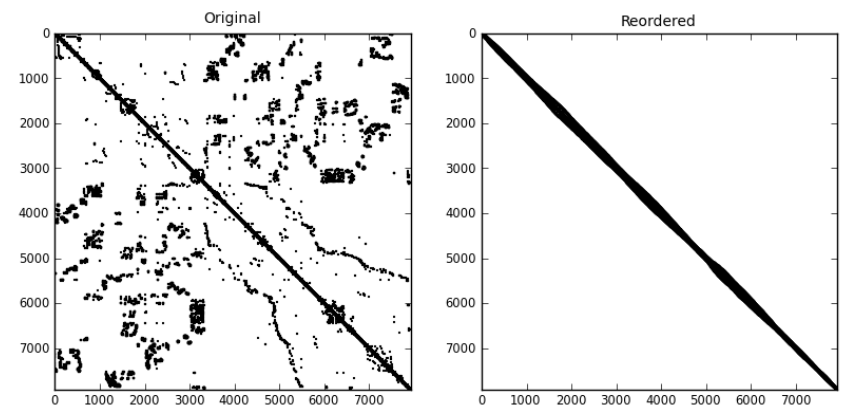
CS121 Parallel Computing  
Spring 2017

# Breadth-first search

- Given a graph, explore it layer by layer.
  - Go wide, then go deep.
- Large number of applications.
  - Connected components, path finding, Ford-Fulkerson max flow algorithm, Cuthill-McKee ordering, bipartiteness testing, search engine crawlers, garbage collection, etc.
- Used in benchmarks such as Graph500 and Parboil to test parallel computer's memory performance.

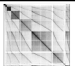





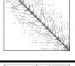
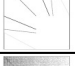
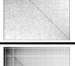
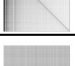
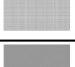
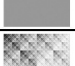



Source: <http://www.stoimen.com/blog/2012/10/08/computer-algorithms-shortest-path-in-a-graph/>



Source: [http://dpo.github.io/pyorder/\\_images/commanche\\_dual\\_rcmk.png](http://dpo.github.io/pyorder/_images/commanche_dual_rcmk.png)

# Real world graphs

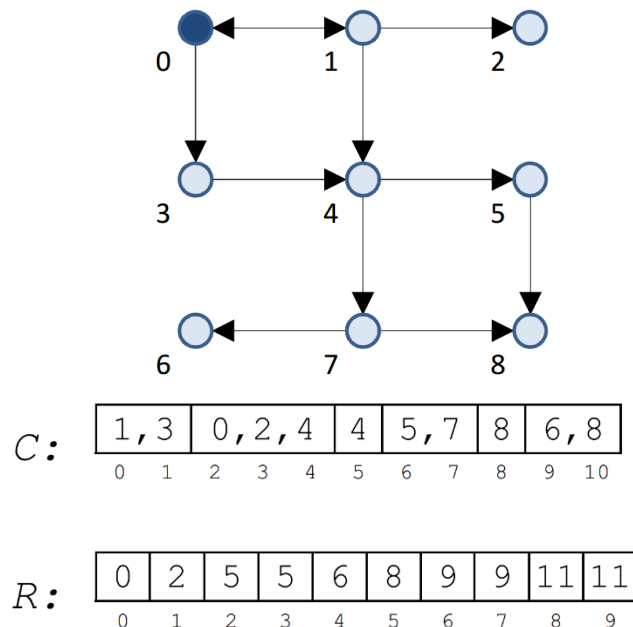
Name	Sparsity Plot	Description	$n$ ( $10^6$ )	$m$ ( $10^6$ )	$\bar{d}$	Avg. Search Depth
europe.osm		European road network	50.9	108.1	2.1	19314
grid5pt.5000		5-point Poisson stencil (2D grid lattice)	25.0	125.0	5.0	7500
hugebubbles-00020		Adaptive numerical simulation mesh	21.2	63.6	3.0	6151
grid7pt.300		7-point Poisson stencil (3D grid lattice)	27.0	188.5	7.0	679
nlpkkt160		3D PDE-constrained optimization	8.3	221.2	26.5	142
audikw1		Automotive finite element analysis	0.9	76.7	81.3	62
case15		Electrophoresis transition probabilities	5.2	94.0	18.2	37
kkt_power		Nonlinear optimization (KKT)	2.1	13.0	6.3	37
coPapersCiteseer		Citation network	0.4	32.1	73.9	26
wikipedia-20070206		Links between Wikipedia pages	3.6	45.0	12.6	20
kron_g500-logn20		Graph500 RMAT ( $A=0.57$ , $B=0.19$ , $C=0.19$ )	1.0	100.7	96.0	6
random.2Mv.128Me		$G(n, M)$ uniform random	2.0	128.0	64.0	6
rmat.2Mv.128Me		RMAT ( $A=0.45$ , $B=0.15$ , $C=0.15$ )	2.0	128.0	64.0	6

- Hundreds of millions of nodes and edges.
  - Some graphs have billions or trillions of edges. But these don't fit into the memory of a single GPU.
- Low average degree (sparse), but high variation in degree.
  - Some nodes have a few neighbors, some nodes 100K's.
- “Small world” graphs have low diameter ( $\sim 10$ ).
- Grids and maps have high diameter ( $\sim 1$ -10K).

Source: Scalable GPU Graph Traversal, Merrill, Garland, Grimshaw

# Sequential algorithm

- Assume graph is sparse, and stored in compressed sparse row format.
- Maintain a queue of unvisited nodes.
  - Dequeue a node, add its unvisited neighbors to the queue.
- Running time  $O(|V|+|E|)$ .



Traversal from source vertex $v_0$		
BFS Iteration	Vertex frontier	Edge frontier
1	{0}	{1,3}
2	{1,3}	{0,2,4,4}
3	{2,4}	{5,7}
4	{5,7}	{6,8,8}
5	{6,8}	{}

```

10  if ( $dist[j] == \infty$ )
11       $dist[j] := dist[i] + 1;$ 
12       $Q.\text{Enqueue}(j)$ 
    
```

# Parallelizing BFS

- First BFS algorithms for GPUs focused on data parallelism.
- Initially set source distance to 0.
- Run for  $|V|$  rounds. In round  $i$ , distance  $i$  nodes are marked. Mark their unvisited neighbors as distance  $i+1$  nodes.
- Works well in small diameter graphs, e.g. social networks.
- Very inefficient for large diameter graphs, e.g. maps, since only a few nodes marked per round.
- $O(|V|^2)$  running time.

```
parallel for (i in V) :  
    dist[i] := ∞  
dist[s] := 0  
iteration := 0  
do :  
    done := true  
    parallel for (i in V) :  
        if (dist[i] == iteration)  
            done := false  
            for (offset in R[i] .. R[i+1]-1)  
                j := C[offset]  
                dist[j] = iteration + 1  
            iteration++  
    while (!done)
```


# Parallelizing BFS

- Linear work parallel BFS algorithms follow the sequential algorithm.
- Two main bottlenecks
  - Maintaining explicit queue of unvisited nodes requires expensive locking operations.
  - If nodes have very different degrees (e.g. power law graphs), there's high load imbalance.

```
parallel for (i in V) :  
    dist[i] := ∞  
dist[s] := 0  
iteration := 0  
inQ := {}  
inQ.LockedEnqueue(s)  
while (inQ != {}) :  
    outQ := {}  
    parallel for (i in inQ) :  
        for (offset in R[i] .. R[i+1]-1)  
            j := C[offset]  
            if (dist[j] == ∞)  
                dist[j] = iteration + 1  
                outQ.LockedEnqueue(j)  
    iteration++  
    inQ := outQ
```

# Gathering neighbors

- We use two queues, one for nodes in current layer of BFS, other for nodes in next layer.
  - After every phase of BFS we swap the queues, to reuse memory.
  - To synchronize the layers, use a separate kernel for each layer.
- For each node in first queue, we first add all its neighbors into the second queue (gather).
  - Some of the neighbors don't belong in the next BFS layer because they've already been visited.
  - Also, we may add duplicates into the second queue.
  - We'll address both problems later.
- To add neighbors of a node into the queue, we use prefix sum instead of locking.
  - If node has  $n_i$  neighbors, we reserve  $n_i$  queue spots for them by adding  $n_i$  into the prefix sum.



# Gathering neighbors

```
1  GatherScan(cta_offset, Q_vfront, C) {
2      shared comm[CTA_THREADS];
3      {r, r_end} = Q_vfront[cta_offset + thread_id];
4      // reserve gather offsets
5      {rsv_rank, total} = CtaPrefixSum(r_end - r);
6      // process fine-grained batches of adjlists
7      cta_progress = 0;
8      while ((remain = total - cta_progress) > 0) {
9          // share batch of gather offsets
10         while((rsv_rank < cta_progress + CTA_THREADS)
11             && (r < r_end))
12         {
13             comm[rsv_rank - cta_progress] = r;
14             rsv_rank++;
15             r++;
16         }
17         CtaBarrier();
18         // gather batch of adjlist(s)
19         if (thread_id < Min(remain, CTA_THREADS) {
20             volatile neighbor = C[comm[thread_id]];
21         }
22         cta_progress += CTA_THREADS;
23         CtaBarrier();
24     }
25 }
```

- Here, CTA means a thread block, an CtaBarrier is a \_\_syncthreads().
- This code uses a block of threads to gather the neighbors of the nodes in the queue.
- After the prefix sum, a thread knows an index into the queue where it can add its neighbors without interference from other threads (while loop from line 10 to 16).
- However, if the thread has many neighbors there might be load imbalance.
  - I.e., the while loop from lines 10-16 might be different for different threads.





# Load balanced gathering

- To load balance, we assign different numbers of threads to gather the neighbors of some node in parallel.
- If the node has moderate number of neighbors, assign a warp of threads to gather its neighbors.
  - Each thread in the warp might initially want to gather neighbors of a different node. All threads in warp write to a common location. The last to write “wins”, and other threads help gather its node’s neighbors.
- If node has large number of neighbors, use entire thread block for gather.
- For remaining nodes, use prefix sum based method.
  - There’s load imbalance, but only for low degree nodes.
- “Moderate” and “large” need to be tuned.
- Eliminates most, but not all load imbalance.

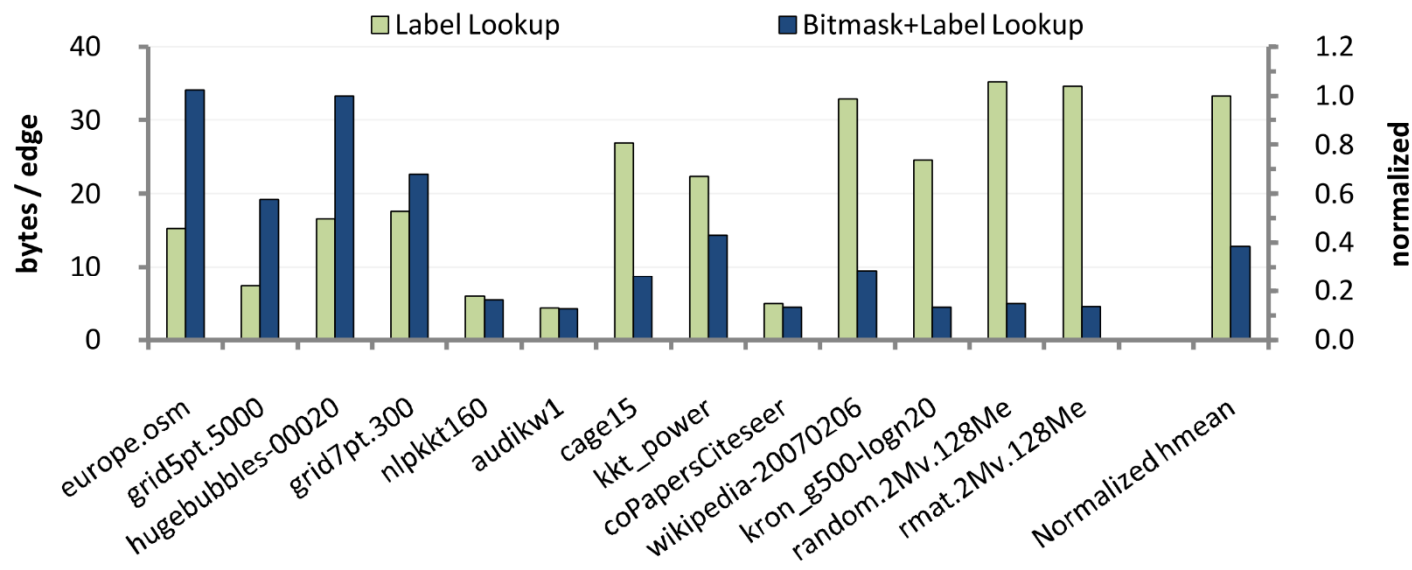


# Visited status lookup

- A node should only be added into the frontier queue if it hasn't been visited.
  - Before adding a node, lookup its visited status.
- To reduce memory traffic, use an integer bitmask to store status of 32 nodes.
- But then two threads might “clobber” each other by setting (different) bits in same int.
  - Can avoid using atomics, but they're very slow.
  - Instead, use normal read writes, but treat bitmask conservatively.
    - For each node, maintain both a shared bitmask bit, and a private integer label.
    - Usually only access bitmask, saving memory traffic. Occasionally access the label.
    - If bit for a node is set, it's definitely visited.
    - If bit is unset, then not sure about node's visited status, so do another lookup on node's label.

# Visisted status lookup





- Bitmasks are cached in texture caches.
- This is effective for low diameter graphs.
  - Even though only 48 KB texture cache per SM.
- Works less well for high diameter graphs, because each layer is processed in separate kernel, and cache flushed after each kernel launch.
  - Graphs on left side have high diameter. The ones on the left are low diameter.



# Duplicates in frontier

- May add same node into frontier multiple times, due to concurrent discovery.
- Problem especially severe in GPU because of SIMD and high parallelism.

<i>BFS Iteration</i>	<i>Actual Vertex-frontier</i>	<i>Actual Edge-frontier</i>
1	0	1,3
2	1,3	2,4,4,6
3	2,4,4,6	5,5,7,5,7,7
4	5,5,7,5,7,7	8,8,8,8,8,8,8




Iteration 1

Iteration 2


Iteration 3

Iteration 4


# Duplicates in frontier




(a) wikipedia-20070206




(b) europe.osm



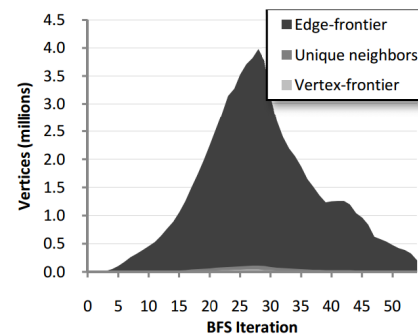
(c) grid7pt.300



(d) nlpkkt60



(e) rmat.2Mv.128Me



(f) audikw1

- Edge frontier: Number of nodes added to queue, allowing duplicates.
- Unique neighbors: Number of nodes added to queue, removing duplicates, but allowing visited nodes.
- Vertex frontier: Unique neighbors which haven't been visited.
- Allowing duplicates can lead to large amount of redundant work.

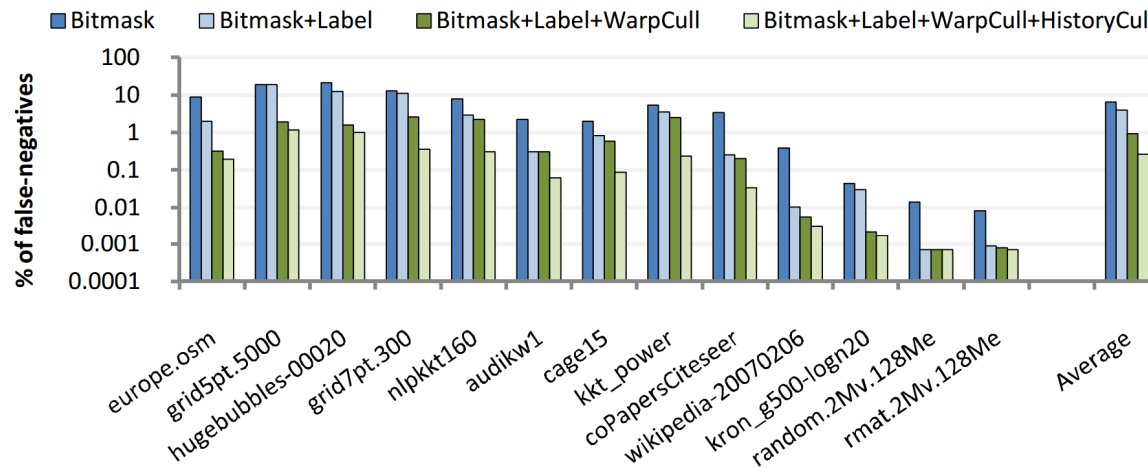
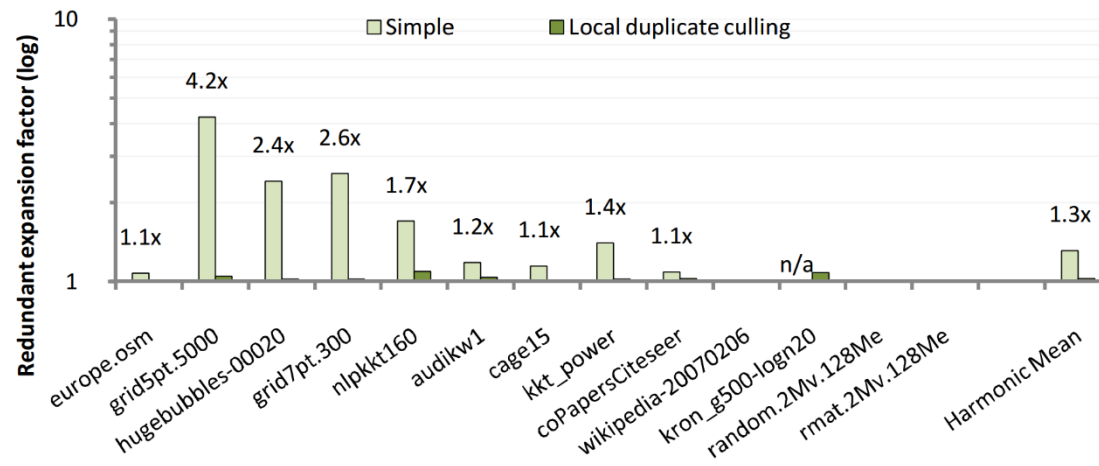


# Duplicate culling

- Try to remove duplicates using hash table.
  - Won't remove all duplicates, but quite effective.
- Warp culling
  - Each warp allocates a hash table in shared memory.
  - When inserting a node, hash it into hash table.
    - If table entry empty, store the node in entry, and add node to queue.
    - If table entry filled, then if entry equals the node, don't add node to queue. Otherwise, add it.
- History culling
  - Same idea, but use the SM's L1 cache.

# Duplicate culling

- Despite small hash table, culling surprisingly effective.



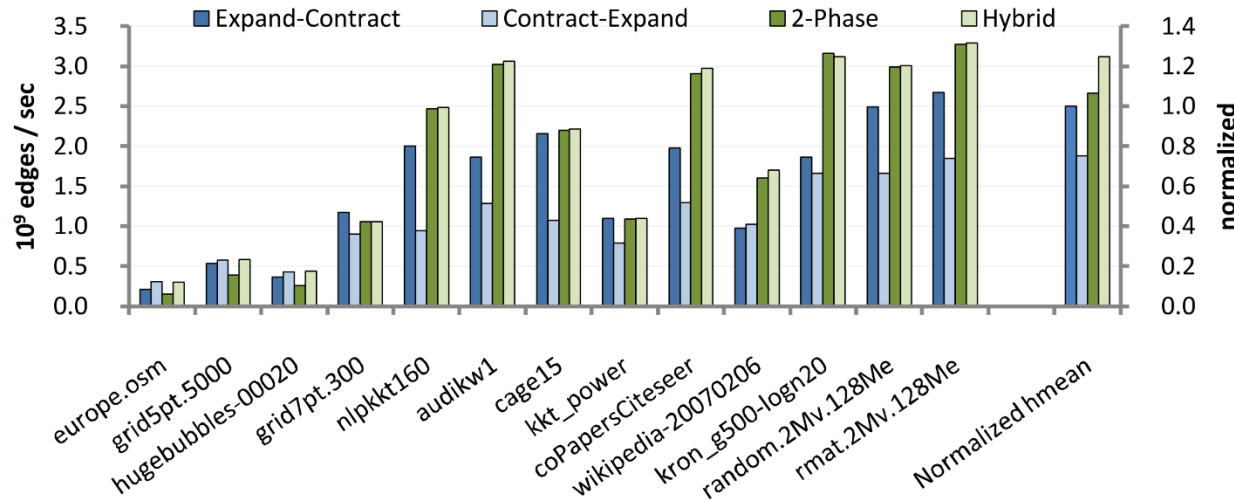


# Putting it together

- Each kernel expands one layer of the BFS.
  - Input is queue containing last BFS layer (possibly with duplicate nodes).
- Threads assigned nodes from queue.
- A thread first uses warp and history culling to determine if its vertex is a duplicate.
- If not, thread gathers node's neighbors.
  - Based on neighbor list size, use a block, warp, or prefix sum gather.
    - Each thread wants to gather neighbors of a different node, and tries to “enlist” a block or warp of threads to help it.
    - Each thread writes into a variable shared by warp or block, then reads it.
    - One thread from the warp / block “wins”. All other threads help it.
- Before adding a gathered node to (current layer's) queue, check if it's already visited.
- If not, the thread contributes 1 to a blockwide prefix sum.
- Synchronize the block and do a blockwide prefix sum to get number of enqueued nodes for block.
- First thread in block atomically adds sum to global queue index, then shares old global index with block.
- Using old global offset and prefix sum offset, each thread adds its gathered neighbor into queue.



# Performance



- Previous algorithm called “contract-expand”, because it first takes current layer’s edge frontier, contracts it (removes duplicates), then expands into next layer’s edge frontier (containing duplicates).
- “Expand-contract”, algorithm expands current vertex frontier, then contracts it (removes duplicates) to next layer’s vertex frontier.
- 2-phase expands then contracts in two kernels.
- Hybrid combines contract-expand with 2-phase.
- Variants differ in amount of memory traffic, latency and parallelism.

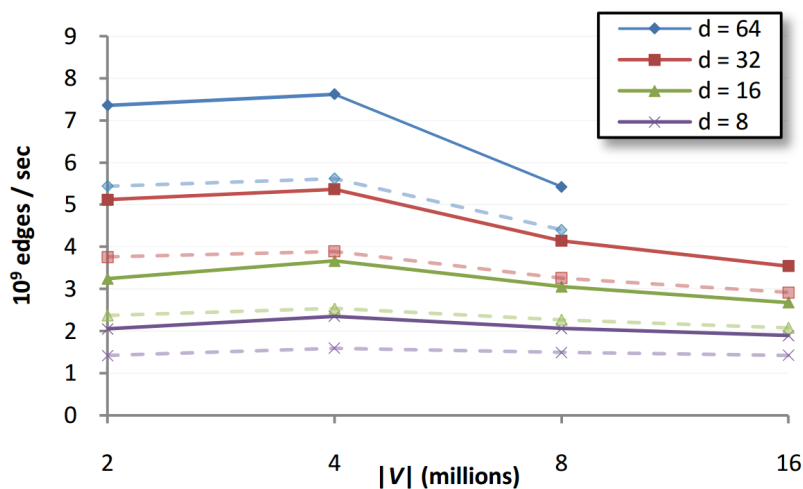

Graph Dataset	CPU	CPU	NVIDIA Tesla C2050 ( <i>hybrid</i> )			
	Sequential <sup>†</sup>	Parallel	Label Distance		Label Predecessor	
	10 <sup>9</sup> TE/s	10 <sup>9</sup> TE/s	10 <sup>9</sup> TE/s	Speedup	10 <sup>9</sup> TE/s	Speedup
europe.osm	0.029		0.31	11x	0.31	11x
grid5pt.5000	0.081		0.60	7.3x	0.57	7.0x
hugebubbles-00020	0.029		0.43	15x	0.42	15x
grid7pt.300	0.038	0.12 <sup>††</sup>	1.1	28x	0.97	26x
nlpkkt160	0.26	0.47 <sup>††</sup>	2.5	9.6x	2.1	8.3x
audikw1	0.65		3.0	4.6x	2.5	4.0x
cage15	0.13	0.23 <sup>††</sup>	2.2	18x	1.9	15x
kkt_power	0.047	0.11 <sup>††</sup>	1.1	23x	1.0	21x
coPapersCiteseer	0.50		3.0	5.9x	2.5	5.0x
wikipedia-20070206	0.065	0.19 <sup>††</sup>	1.6	25x	1.4	22x
kron_g500-logn20	0.24		3.1	13x	2.5	11x
random.2Mv.128Me	0.10	0.50 <sup>†††</sup>	3.0	29x	2.4	23x
rmat.2Mv.128Me	0.15	0.70 <sup>†††</sup>	3.3	22x	2.6	18x



# Multi-GPU BFS

- Multiple GPUs can use a single logical address space.
  - Communicate through PCI-e 2.0 (6.6 GB/s).
- Given  $p$  GPUs, assign  $n/p$  vertices and corresponding edges per GPU.
  - Vertices assigned in round robin order for load balancing.
  - Poor locality if  $p$  large.
- Each GPU expands / contracts its own vertex queue, as in the single GPU algorithm (\*).
- Then sort the new frontier into  $p$  bins, corresponding to vertices from different GPUs.
- Barrier across all GPUs.
- Run  $p-1$  kernels, where in  $i$ 'th kernel, the  $i$ 'th GPU collects bin  $i$  from each other GPU.
- Then go back to step (\*) to form the next layer. Continue until all nodes visited.

# Performance



- Only achieved speedup on graphs with small diameters and large average degrees.
  - Smaller diameter requires less synchronization.
  - Larger degree makes duplicate culling more effective.
  - Max speedups 1.5X, 2.1X and 2.5X on 2, 3, 4 GPUs.
  - Sometimes parallel algorithm performed much worse.