



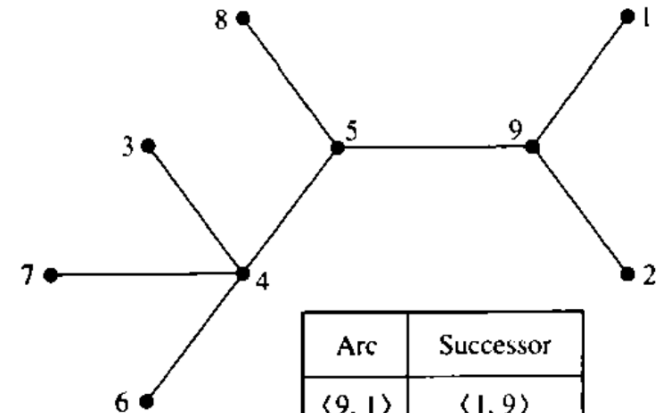
PRAM 2

Graph algorithms

CS121 Parallel Computing
Spring 2017

Euler tours

- An Euler tour of a graph is a cycle that goes through every edge of the graph.
 - It may go through a vertex multiple times.
- A connected, directed graph has an Euler tour if and only if the indegree and outdegree of every vertex are equal.
- Suppose we take an undirected graph, and for edge (u,v) , create two directed edges (u,v) and (v,u) .
 - Then every vertex has equal indegree and outdegree, and so has an Euler tour.
- Consider a tree where each edge has been doubled.
 - To find an Euler tour of the tree, first order the edges adjacent to each node arbitrarily.
 - Say the neighbors of a node v are ordered u_0, \dots, u_{d-1} . Then set the successor of edge (u_i, v) on the tour to $(v, u_{(i+1) \bmod d})$.
- The Euler tour of a tree can be computed in $O(1)$ parallel time.



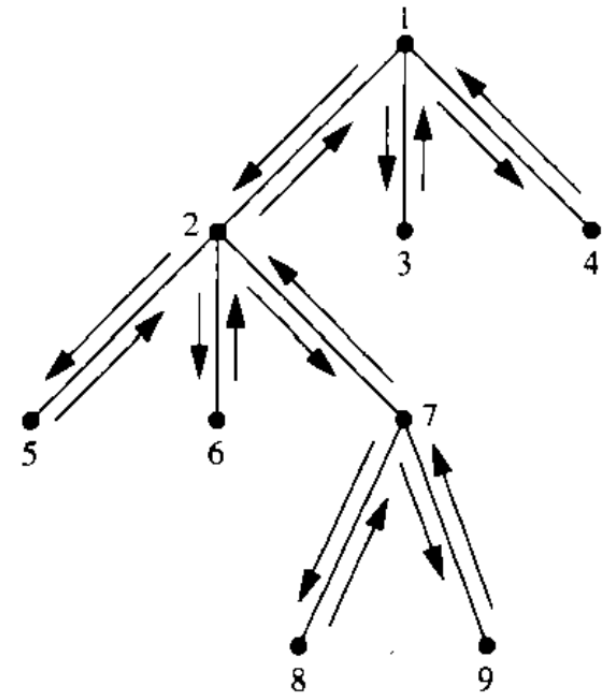
Arc	Successor
$\langle 9, 1 \rangle$	$\langle 1, 9 \rangle$
$\langle 9, 2 \rangle$	$\langle 2, 9 \rangle$
$\langle 4, 3 \rangle$	$\langle 3, 4 \rangle$
$\langle 5, 4 \rangle$	$\langle 4, 3 \rangle$
$\langle 3, 4 \rangle$	$\langle 4, 7 \rangle$
$\langle 7, 4 \rangle$	$\langle 4, 6 \rangle$
$\langle 6, 4 \rangle$	$\langle 4, 5 \rangle$
$\langle 8, 5 \rangle$	$\langle 5, 4 \rangle$
$\langle 4, 5 \rangle$	$\langle 5, 9 \rangle$
$\langle 9, 5 \rangle$	$\langle 5, 8 \rangle$
$\langle 4, 6 \rangle$	$\langle 6, 4 \rangle$
$\langle 4, 7 \rangle$	$\langle 7, 4 \rangle$
$\langle 5, 8 \rangle$	$\langle 8, 5 \rangle$
$\langle 5, 9 \rangle$	$\langle 9, 2 \rangle$
$\langle 2, 9 \rangle$	$\langle 9, 1 \rangle$
$\langle 1, 9 \rangle$	$\langle 9, 5 \rangle$

$\langle 9, 1 \rangle \rightarrow \langle 1, 9 \rangle \rightarrow \langle 9, 5 \rangle \rightarrow \langle 5, 8 \rangle \rightarrow \langle 8, 5 \rangle \rightarrow \langle 5, 4 \rangle \rightarrow$
 $\langle 4, 3 \rangle \rightarrow \langle 3, 4 \rangle \rightarrow \langle 4, 7 \rangle \rightarrow \langle 7, 4 \rangle \rightarrow \langle 4, 6 \rangle \rightarrow \langle 6, 4 \rangle \rightarrow$
 $\langle 4, 5 \rangle \rightarrow \langle 5, 9 \rangle \rightarrow \langle 9, 2 \rangle \rightarrow \langle 2, 9 \rangle \rightarrow \langle 9, 1 \rangle$

Source: Introduction to Parallel Algorithms, Jaja

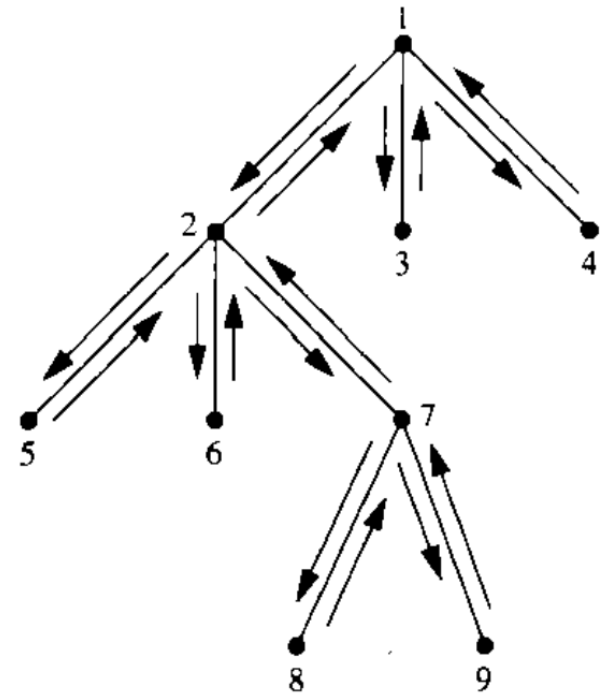
Parallel tree operations

- Many operations on trees can be done in parallel using Euler tours and prefix sum.
- These operations in turn are used in other parallel graph algorithms.
- We first root a tree in parallel.
 - I.e. set an arbitrary node r as the tree's root. Then each node v needs to compute $p(v)$, its parent in the rooted tree.
 - To do this, assign a weight of 1 to each edge in an Euler tour of the tree.
 - Then compute the parallel prefix sum of the edges.
 - For each edge (u,v) , set $u=p(v)$ whenever the prefix sum of (u,v) is less than the prefix sum of (v,u) .
 - Thus, we can root a tree with n nodes in $O(\log n)$ time and $O(n)$ work.



Node depths


- For each node, compute its depth in a rooted tree.
 - For each node v , let $p(v)$ be its parent.
 - Set the weight of edge $(p(v), v)$ to 1, and the weight of edge $(v, p(v))$ to -1.
 - Compute a parallel prefix sum of the Euler tour starting at the root.
 - The depth of node v is the prefix sum of edge $(p(v), v)$.
- For a tree with n nodes, this takes $O(\log n)$ time using $O(n)$ work.



Postorder numbering

- Traverse a rooted tree in postorder, starting from the root r .
 - Start an Euler tour from r . For each node v , we want to visit v 's children in the order of the tour, then visit v itself.
 - Ex


v	1	2	3	4	5	6	7	8	9
$n(v)$	9	6	7	8	1	2	5	3	4
 - For each node v , set the weight of edge $(v, p(v))$ to 1, and the weight of $(p(v), v)$ weight 0.
 - Compute a parallel prefix sum of the Euler tour.
 - For each $v \neq r$, set $n(v)$ to the prefix sum of edge $(v, p(v))$. Set $n(r)=n$.
- In a tree with n nodes, we can compute the postorder numbering in $O(\log n)$ time and $O(n)$ work.



Euler Path	Weight	Prefix Sums
$\langle 1, 2 \rangle$	0	0
$\langle 2, 5 \rangle$	0	0
$\langle 5, 2 \rangle$	1	1
$\langle 2, 6 \rangle$	0	1
$\langle 6, 2 \rangle$	1	2
$\langle 2, 7 \rangle$	0	2
$\langle 7, 8 \rangle$	0	2
$\langle 8, 7 \rangle$	1	3
$\langle 7, 9 \rangle$	0	3
$\langle 9, 7 \rangle$	1	4
$\langle 7, 2 \rangle$	1	5
$\langle 2, 1 \rangle$	1	6
$\langle 1, 3 \rangle$	0	6
$\langle 3, 1 \rangle$	1	7
$\langle 1, 4 \rangle$	0	7
$\langle 4, 1 \rangle$	1	8

Number of descendant


- For each node v in a rooted tree, compute the number of nodes in the subtree rooted at v .
- To do this, we compute prefix sums as in the postorder numbering.
- Then the number of descendants of a node v equals the prefix sum of $(v, p(v))$ minus the prefix sum of $(p(v), v)$.
- In a tree with n nodes, we can compute the number of descendants in $O(\log n)$ time and $O(n)$ work.



Euler Path	Weight	Prefix Sums
$\langle 1, 2 \rangle$	0	0
$\langle 2, 5 \rangle$	0	0
$\langle 5, 2 \rangle$	1	1
$\langle 2, 6 \rangle$	0	1
$\langle 6, 2 \rangle$	1	2
$\langle 2, 7 \rangle$	0	2
$\langle 7, 8 \rangle$	0	2
$\langle 8, 7 \rangle$	1	3
$\langle 7, 9 \rangle$	0	3
$\langle 9, 7 \rangle$	1	4
$\langle 7, 2 \rangle$	1	5
$\langle 2, 1 \rangle$	1	6
$\langle 1, 3 \rangle$	0	6
$\langle 3, 1 \rangle$	1	7
$\langle 1, 4 \rangle$	0	7
$\langle 4, 1 \rangle$	1	8

Coloring a cycle

- Given a graph $G=(V,E)$, a k -coloring of G is a mapping $c: V \rightarrow \{0,1, \dots, k-1\}$ s.t. $c(i) \neq c(j)$ whenever $(i,j) \in E$.
- We give a super fast algorithm for 3-coloring a (directed) cycle of n nodes.
 - If n is odd, any coloring uses at least 3 colors.
 - Coloring the cycle is a form of symmetry breaking.
 - For any node v , let $S(v)$ be the node after v .
- The main subroutine is the following.
 - Initially, color every node by its node ID.
 - Consider the binary representation of $c(v)$ for a node v .
 - Let k be the least significant digit i which $c(v)$ and $c(S(v))$ differ.
 - Set $c'(v)=2k+c(v)_k$, where $c(v)_k$ is the k 'th digit of $c(v)$.
- **Claim** If c is a valid coloring, then so is c' .
- **Proof** Since c is a valid coloring, then $c(v) \neq c(S(v))$, so k exists.
 - Suppose $c'(v) = c'(u)$, for some v and $u = S(v)$.
 - Then $c'(v) = 2k + c(v)_k$ and $c'(u) = 2l + c(u)_l$ for some k and l .
 - Since $c'(v) = c'(u)$, then $k=l$, because $c(v)_k, c(u)_l < 2$.
 - But then $c(v)_k = c(u)_k$, contradicting the definition of k .



v	c	k	c'
1	0001	1	2
3	0011	2	4
7	0111	0	1
14	1110	2	5
2	0010	0	0
15	1111	0	1
4	0100	0	0
5	0101	0	1
6	0110	1	3
8	1000	1	2
10	1010	0	0
11	1011	0	1
12	1100	0	0
9	1001	2	4
13	1101	2	5



Coloring a cycle

- To analyze the time complexity, suppose in some round the max number of bits to represent any color is t .
- Then the max number of bits to represent any color in the next round is $\lceil \log t \rceil + 1$, because any color in the next round is $\leq 2t + 1$.
 - So the number of bits used to represent a color decreases from t to $\lceil \log t \rceil + 1$ in each round.
- Let $\log^{(i)} x = \log(\log^{(i-1)} x)$, i.e. we apply the log function i times to x .
- Let $\log^* x = \min\{i \mid \log^{(i)} x \leq 1\}$ be the number of times we have to take log's until a value becomes ≤ 1 .
 - $\log^* x$ is incredibly small. In fact, $\log^* x \leq 5$ for all $x \leq 2^{65526}!$




Coloring a cycle

- Since the number of bits to represent a color in the first round is $\log n$, then in $O(\log^* n)$ rounds, we can represent any color using $O(1)$ bits.
 - In fact, we can apply the subroutine until we use 6 colors in a round.
 - With 6 colors, need 3 bits to represent a color. So in the next round, colors are between 0 and $2^3-1=7$, and we again use up to 6 colors.
- To decrease the number of colors from 6 to 3, we run 3 more rounds.
 - In round i , take any node colored using color $i+2$ and color it using the min possible color in $\{0,1,2\}$, i.e. the min color not used by its neighbors.
- In total, we 3-color the ring in $O(\log^* n)$ rounds, using $O(n \log^* n)$ work.



Connected components

- Given an undirected graph, partition it into maximal sets of nodes that are connected to each other.
- Can be solved sequentially in $O(m)$ time using BFS / DFS.
 - m is number of edges, n is number of vertices.
- However, no efficient BFS / DFS PRAM algorithms known.
- Instead, use graph contractions.
 - In each phase, merge (contract) a set of connected nodes into a supernode.
 - Form a contracted graph on the supernodes, then apply algorithm recursively.
 - Eventually each connected component is contracted to one node.
 - Many different algorithms, depending on which nodes they contract.




Randomized parallel algorithm

- Break graph into star graphs and contract each star.
- In each phase, do the following steps in parallel.
 - Every node flips a coin and chooses to be a parent or child node.
 - Each child node points to a parent node it's connected to.
 - Now have a set of stars, with the parent nodes as the centers.
 - If child not connected to any parent node, it forms its own star.
 - Contract each star to its center, then apply algorithm recursively.
 - Label all nodes in star by the label of the parent.
 - Keep the edges between differently labeled nodes.
 - After recursion returns, each child with a parent again takes parent's label, which might have changed.




Implementation

- Use $n+m$ processors, one processor for each vertex and edge.
 - Call these E and V procs, resp.
 - Each V proc has a label which can change over time.
 - Each E proc responsible for edge (u,v) , where u, v are V procs.
 - V and E procs may become inactive over time.
- In each phase, each active V proc flips a coin to decide if it's a child or parent proc.
- Each active E proc (u,v) checks if u is a child proc and v is a parent (or vice versa).
 - If so, it sets u 's label to v (or v 's label to u).
 - Another E proc (u,v') could set u 's label to v' . In this case, either the v or v' write succeeds.



Implementation

- Each parent V proc, or child V proc whose label didn't change (i.e. it had no parent), stays active.
 - Other V procs become inactive.
- Each active E proc (u,v) where u, v have different labels stays active.
 - Other E procs become inactive.
 - From now on, E will be responsible for V processes (u',v') , where u' and v' are the labels of u and v , resp.
- The active V and E procs run the algorithm recursively.
- After recursion returns, inactive E procs (u,v) (where u is the child) write v 's label to u .
- At end, all V procs in a connected component have same label.







Complexity

- Key fact is that number of active V procs decreases by $1/4$ fraction in expectation each phase.
 - A V proc becomes inactive if it's a child node and one its neighbors is a parent node.
 - The former probability is $1/2$, and the latter is $\geq 1/2$.
 - Thus each V proc becomes inactive with probability $\geq 1/4$.
- With high probability, after $O(\log n)$ phases, there's only one V proc and recursion ends.
- Each phase takes $O(1)$ time, and does $O(m+n)$ work.
- Total time is $O(\log n)$, total work is $O((m+n) \log n)$.
 - This algorithm isn't work efficient.
 - There exist work efficient randomized CC algorithms.



Deterministic parallel algorithm

- Again work in phases, with following parallel steps.
 - Each node points to a neighbor with lower ID.
 - This breaks graph into a directed forest.
 - Contract each forest to the lowest ID node using pointer jumping.
 - Recurse on contracted graph.






Implementation

- As before, use $n+m$ procs, for V and E .
 - But now, V procs always active. E procs may become inactive.
- Each E proc (u,v) checks if $u < v$, and if so sets v 's label to u .
 - Again, conflicts resolved arbitrarily.
- V procs then apply pointer jumping on the labels, taking the label of the proc it points to.
- Each active E proc (u,v) where u, v have different labels stays active. Other E procs become inactive.
 - From now on, E will be responsible for V processes (u',v') , where u' and v' are the labels of u and v , resp.
 - There may be multiple E procs responsible for same (u',v') .
- The V procs and active E procs run algorithm recursively.
 - Note that in recursive call, all V procs apply pointer jumping.




Complexity

- The basic algorithm may take $O(n)$ time on the graph below.
- But notice that if we made nodes point to higher neighbors, the graph would be solved in $O(1)$ time.
- In each phase, if we consider either having nodes point to smaller neighbors, or pointing to higher neighbors, in one case $\geq n/2$ nodes point to another node.
- Thus the algorithm finishes in $O(\log n)$ phases.
- Each phase does pointer jumping, which takes $O(\log n)$ time and $O(n \log n)$ work.
- Total time is $O(\log^2 n)$, and work is $O((m+n \log n) \log n)$.



Minimum spanning tree





- Given an undirected graph with edge weights, an MST is a connected subgraph containing all the vertices, which has minimum total weight.
- Can be solved in $O(m + n \log n)$ time sequentially by a greedy algorithm.
- Key property is that for any set of vertices W , the minimum cost edge from W to $V \setminus W$ is in the MST.
 - So for any vertex v , min cost edge containing v is in MST.
- Will describe a parallel MST algorithm based on the randomized parallel algorithm for connected component.



Source: Introduction to Algorithms, Cormen et al.

Randomized parallel algorithm

- Each node randomly chooses to be a parent or child node.
- Each child node u finds min weight incident edge (u,v) , and points to v if v is parent.
 - This forms a set of stars with parents as centers.
 - If v isn't a parent, u forms its own star.
- Contract each star to the parent, and run algorithm recursively.
- Finding min weight incident edge quickly is a little tricky.
 - One possibility is to use priority CRCW.
 - Presort the edges by nondecreasing weight. Then when each E proc (u,v) writes to u , min weight edge wins.





Complexity

- At least $1/4$ of vertex processors become inactive each phase in expectation.
 - Given a node u and min weight edge (u,v) , there's $1/4$ probability u is child and v is parent.
- Thus, there are $O(\log n)$ phases with high probability.
 - Finding min weight incident edge takes $O(1)$ time after presorting edge weights.
- Total time is $O(\log n)$, total work is $O((m+n) \log n)$.