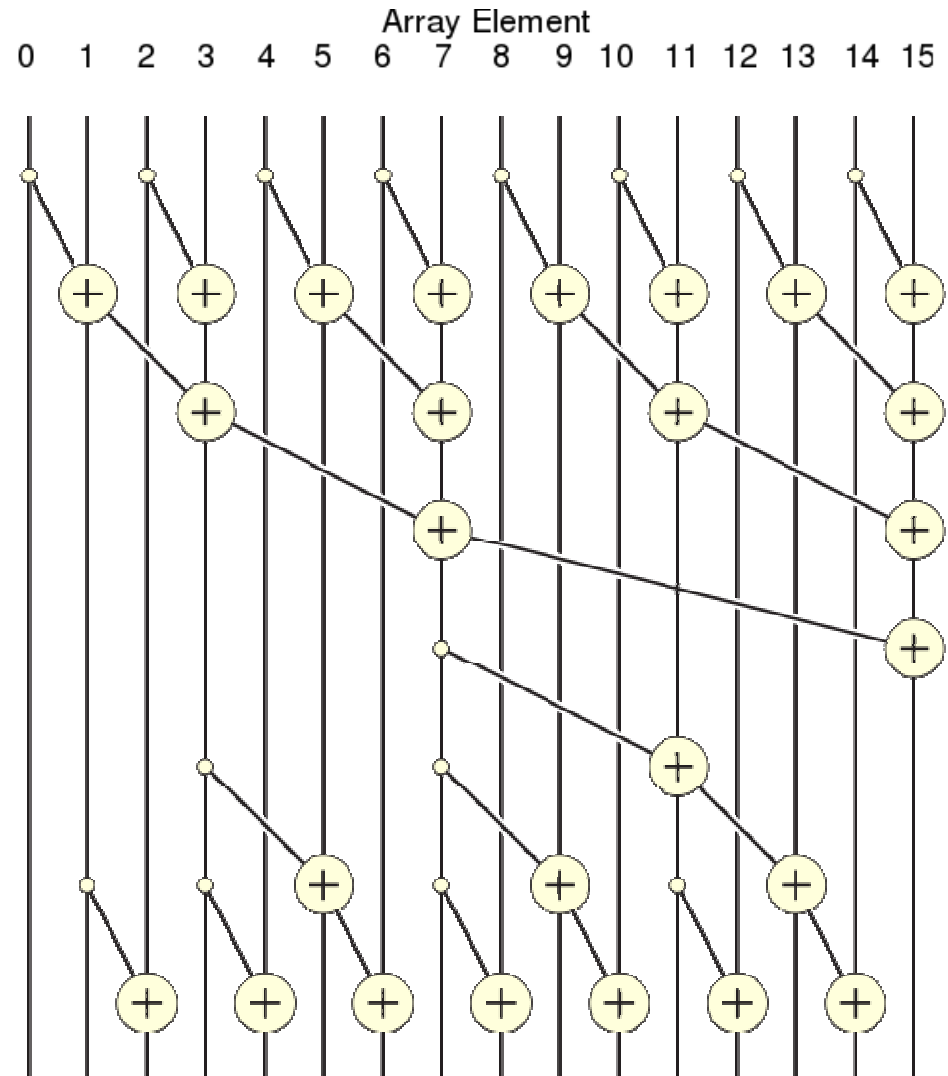


# Efficient parallel prefix sum

```
int stride = 1;
while (stride <= blockDim.x) {
    int i = 2*stride*(threadIdx.x+1)-1;
    if (i < 2*blockDim.x)
        sum[i] += sum[i-stride];
    stride *= 2;
    __syncthreads();
}

int stride = blockDim.x/2;
while (stride > 0) {
    int i = 2*stride*(threadIdx.x+1)-1;
    if (i+stride < 2*dimBlock.x)
        sum[i+stride] += sum[i];
    stride /= 2;
    __syncthreads();
}
```

- A thread block computes prefix sum of array sum in shared memory.
  - Size of sum is  $2 \times (\text{block size})$ .
  - In example, block size = 8.
- In down sweep, threads 0 to (block size) / stride – 1 work in iteration stride.
- In up sweep, threads 0 to (block size) / (2\*stride) – 1 work in iteration stride.

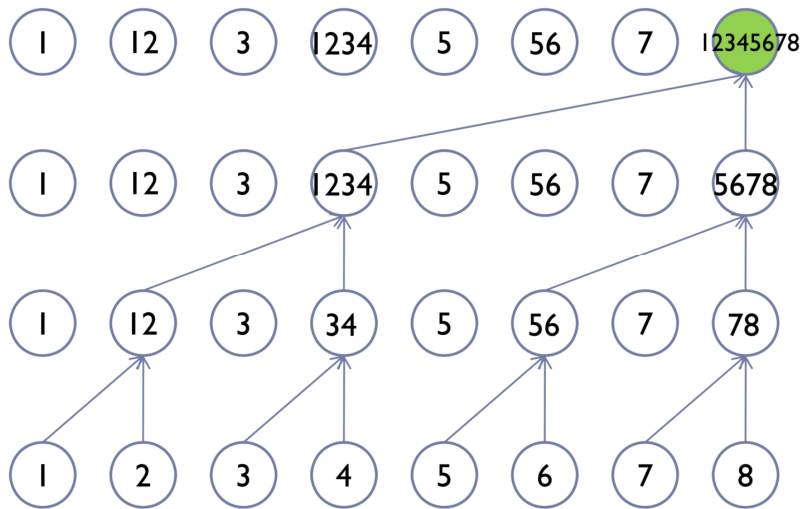





# Exclusive scans

- Just like a normal scan, except each input value shouldn't include itself in its output.
  - Ex  $[1,2,3,4] \Rightarrow [0,1,3,6]$ .
- Up-sweep is the same as in inclusive scan.
- But during down-sweep, first zero out the final output value.
- Then follow a half butterfly pattern downwards.
  - Each right child sums its parents' values.
  - Each left child takes its parent's value.

# Exclusive scans



Up-sweep



Down-sweep

Up-sweep (reduce):

```

1: for  $d = 0$  to  $\log_2 n - 1$  do
2:   for all  $k = 0$  to  $n - 1$  by  $2^{d+1}$  in parallel do
3:      $x[k + 2^{d+1} - 1] \leftarrow x[k + 2^d - 1] + x[k + 2^{d+1} - 1]$ 
  
```

Down-sweep:


```

1:  $x[n - 1] \leftarrow 0$ 
2: for  $d = \log_2 n - 1$  down to  $0$  do
3:   for all  $k = 0$  to  $n - 1$  by  $2^{d+1}$  in parallel do
4:      $t \leftarrow x[k + 2^d - 1]$ 
5:      $x[k + 2^{d+1} - 1] \leftarrow x[k + 2^{d+1} - 1]$ 
6:      $x[k + 2^{d+1} - 1] \leftarrow t + x[k + 2^{d+1} - 1]$ 
  
```

Source: <http://courses.me.berkeley.edu/ME290R/S2009/lectures/lec15.PDF>

# Arbitrary input size

- The inclusive scan algorithm only works for array size  $\leq 2 \times (\text{block size})$ .
- For bigger inputs, break it into segments of size  $2 \times (\text{block size})$ .
- Compute prefix sum on each segment using block algorithm.
- Copy sum of whole segment (stored in `sum[blockDim.x-1]`) to `segment_sum` array.
- Do this for all blocks until they all finish.
  - Ensure blocks finished by ending kernel.
- Compute prefix sum of `segment_sum` array in a second kernel.
- In a third kernel, distribute prefix sums to each segment.
  - Segment increases all values by prefix sum received.



# Bank conflicts

- Recall memory address  $x$  stored at  $x \% n$  if shared memory has  $n$  banks.
  - Current GPUs have 32 banks.
- Current algorithm has many bank conflicts, causing serialized accesses.

bank 0	0	4	8	12	16
bank 1	1	5	9	13	17
bank 2	2	6	10	14	18
bank 3	3	7	11	15	19

16 banks, stride = 1. 2 way bank conflicts

<b>tid</b>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>i</b>	1	3	5	7	9	11	13	15	17	19	21	23	25	27	29	31
<b>bank</b>	1	3	5	7	9	11	13	15	1	3	5	7	9	11	13	15

16 banks, stride = 2. 4 way bank conflicts

<b>tid</b>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>i</b>	3	7	11	15	19	23	27	31	35	39	43	47	51	55	59	63
<b>bank</b>	3	7	11	15	3	7	11	15	3	7	11	15	3	7	11	15

```
...  
int i = 2*stride*  
    (threadIdx.x+1)-1;  
if (i < 2*blockDim.x)  
    sum[i] += sum[i-  
        stride];  
...
```

# Removing bank conflicts

- Remove bank conflicts by padding the sum array.
- Store  $i$ 'th item at address  $i + \text{floor}(i / (\# \text{ banks}))$  instead of address  $i$ .
  - Do this for reads and writes.
  - Waste some space ( $\sim 3\%$  with 32 banks), but get faster performance.
- Ex 4 banks.

array	0	1	2	3	4	5	6	7	8	9	10	11		
padded array	0	1	2	3	P	4	5	6	7	P	8	9	10	11

- Padding is a general strategy for removing bank conflicts, though exact scheme depends on problem.

# Removing bank conflicts

16 banks, stride = 2. 4 way bank conflicts

<b>tid</b>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>i</b>	3	7	11	15	19	23	27	31	35	39	43	47	51	55	59	63
<b>bank</b>	3	7	11	15	3	7	11	15	3	7	3	15	3	7	11	15

16 banks, stride = 2,  $i' = i + \text{floor}(i / \# \text{ banks})$ . No bank conflicts

<b>tid</b>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>i'</b>	3	7	11	15	20	24	28	32	37	41	45	49	54	58	62	66
<b>bank</b>	3	7	11	15	4	8	12	0	5	9	13	1	6	10	14	2

# Segmented scan

## Work-efficient segmented scan

### Up-sweep:

```
for d=0 to (log2n - 1) do
  forall k=0 to n-1 by 2d+1 do
    if flag[k + 2d+1 - 1] == 0:
      data[k + 2d+1 - 1] ← data[k + 2d - 1] + data[k + 2d+1 - 1]
      flag[k + 2d+1 - 1] ← flag[k + 2d - 1] || flag[k + 2d+1 - 1]
```

### Down-sweep:

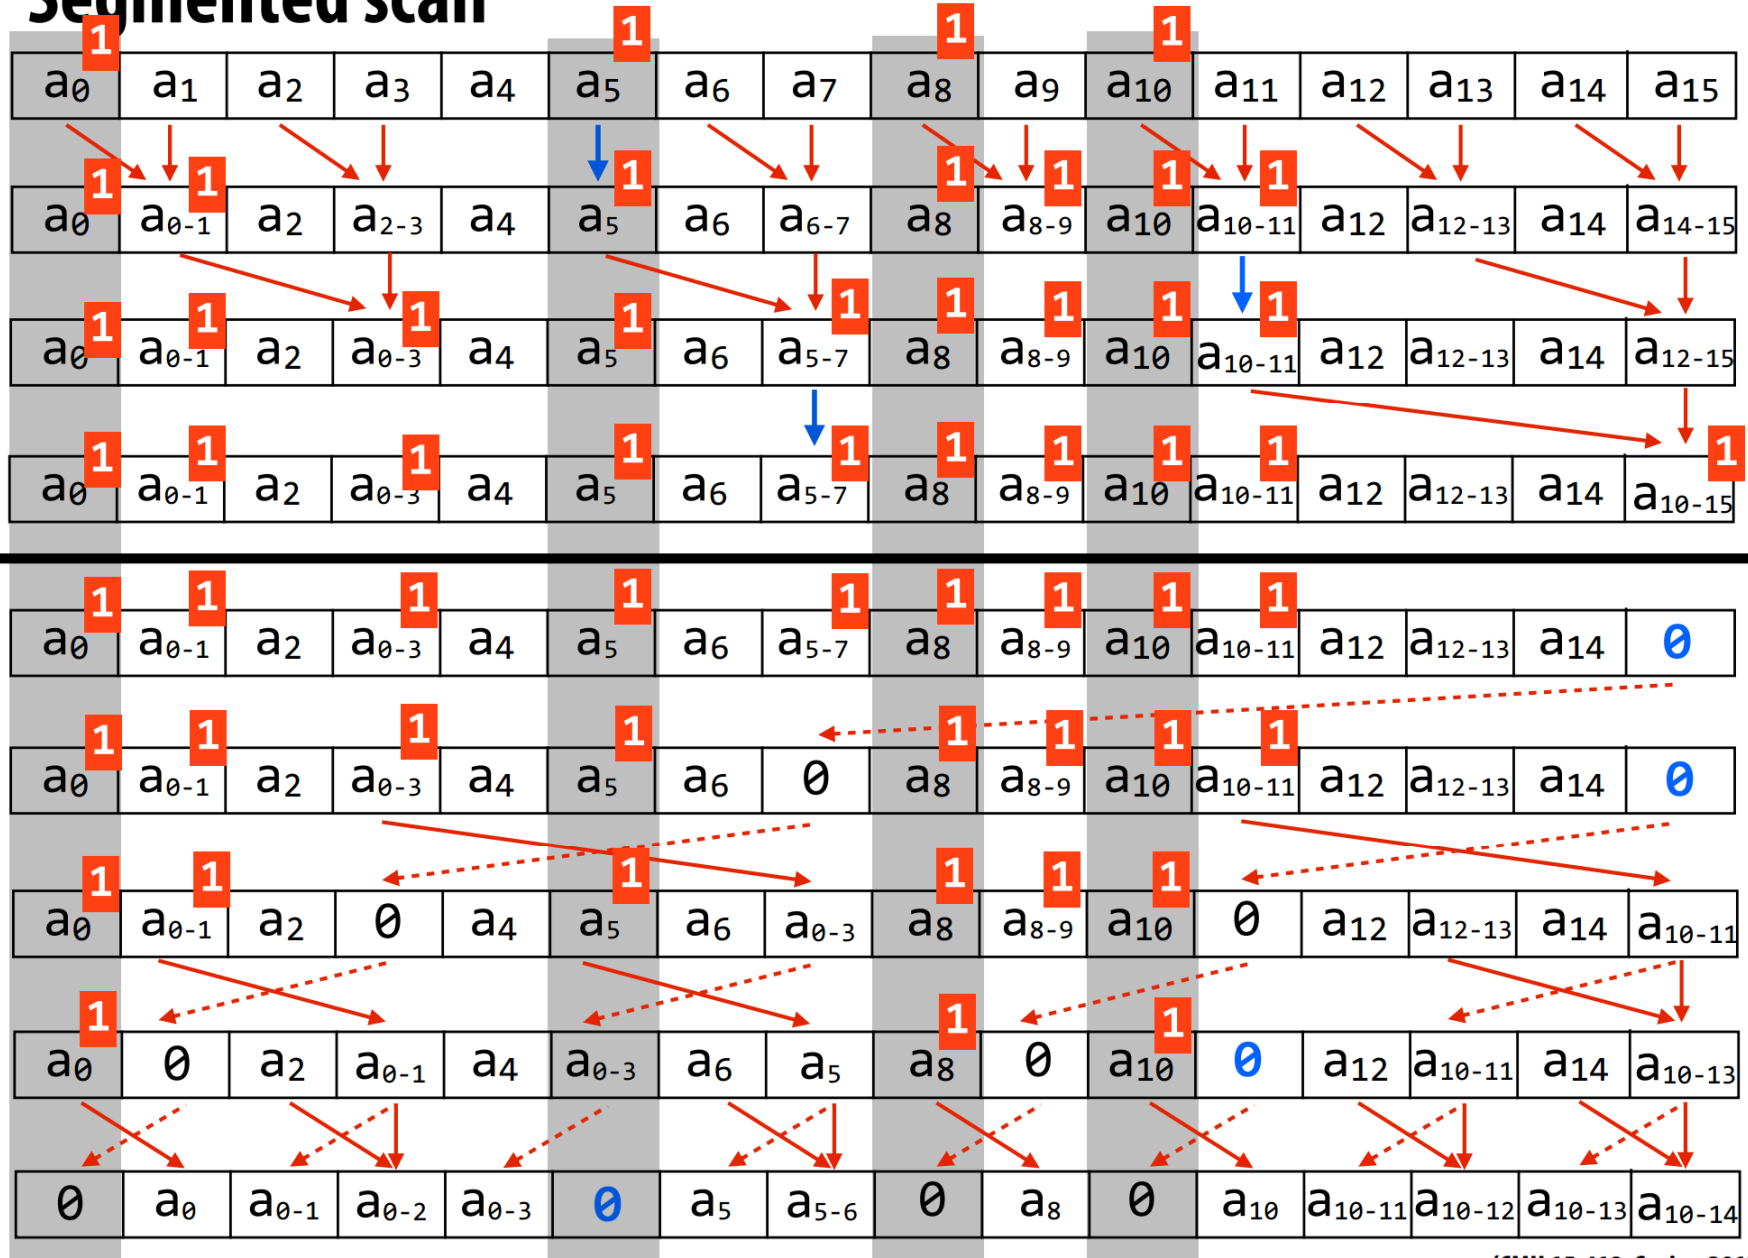
```
data[n-1] ← 0
for d=(log2n - 1) down to 0 do
  forall k=0 to n-1 by 2d+1 do
    tmp ← data[k + 2d - 1]
    data[k + 2d - 1] ← data[k + 2d+1 - 1]
    if flagoriginal[k + 2d] == 1:           // maintain copy of original flags
      data[k + 2d+1 - 1] ← 0
    else if flag[k + 2d - 1] == 1:
      data[k + 2d+1 - 1] ← tmp
    else:
      data[k + 2d+1 - 1] ← tmp + data[k + 2d+1 - 1]
      flag[k + 2d - 1] ← 0
```

- ❑ Sometimes need to run (exclusive) prefix sum on several segments at once.
- ❑ Ex [1 2 3 4] [6 5] [1 3 5] ⇒ [0 1 3 6] [0 6] [0 1 4]
- ❑ If do m scans, each of size n individually, then  $O(n \log m)$  time.
- ❑ We do all the scans in  $O(\log(mn))$  time.
- ❑ Up sweep distributes partial sums.
  - ❑ Use flags to delimit segments.
  - ❑ No value “crosses” a segment.
- ❑ Down sweep collects values from one segment and sums them.

Source: <http://www.cs.cmu.edu/afs/cs/academic/class/15418-s12/www/>



# Segmented scan





# Application: compaction

- Create array containing elements of input array satisfying a condition.
- **Ex** Move all odd numbers in  $A$  to front of *output*.
  - Create filter array that's 1 if element satisfies condition.
  - Prefix sum the filter array.
  - For each element, if it satisfies condition, move it to index given by prefix sum.

$A =$  [1 3 2 4 8 6 5 4 9 7 3]

*filter* = [1 1 0 0 0 0 1 0 1 1 1]

*sums* = [1 2 2 2 2 2 3 3 4 5 6]

*output* = [1 3 5 9 7 3]




# Application: string comparison


- Compare two strings alphabetically.
- Ex  $\text{parallax} < \text{parallel}$ .
- Let strings be  $S, T$ . Let  $S[i], T[i]$  denote  $i$ 'th letter of  $S, T$ .
- ❖ In parallel,  $i$ 'th processor compares  $S[i]$  to  $T[i]$ .
  - ❖ If  $S[i] > T[i]$ , set  $A[i] = 1$ .
  - ❖ If  $S[i] = T[i]$ , set  $A[i] = 0$ .
  - ❖ If  $S[i] < T[i]$ , set  $A[i] = -1$ .
  - ❖ If  $S[i]$  or  $T[i]$  doesn't exist, set  $A[i] = 0$ .
- ❖ Compact  $A$  to remove all 0's.
- ❖ If  $\text{output}[1] = 1$ , then  $S > T$ .
- ❖ If  $\text{output}[1] = -1$ , then  $T > S$ .
- ❖ If output is empty, then  $S = T$ .
- Ex  $S = \text{parallax}, T = \text{parallel}, A = [0, 0, 0, 0, 0, 0, -1, 1], \text{output} = [-1, 1]$ , so  $T > S$ .

# Application: line of sight

```
procedure line-of-sight(altitude)
  in parallel for each index  $i$ 
     $\text{angle}[i] \leftarrow \arctan(\text{scale} \times (\text{altitude}[i] - \text{altitude}[0]) / i)$ 
  max-previous-angle  $\leftarrow$  max-prescan(angle)
  in parallel for each index  $i$ 
    if ( $\text{angle}[i] > \text{max-previous-angle}[i]$ )
       $\text{result}[i] \leftarrow \text{"visible"}$ 
    else
       $\text{result}[i] \leftarrow \text{not "visible"}$ 
```



Altitude Map



Ray Vectors

- Given a contour map, an observation point X and a direction, want to know which points are visible.
- First, draw a line from X in the observing direction and record the altitudes along the line in an altitude vector.
- Then for each point calculate its angle, based on its altitude and distance from X.
- Then do a max-scan over the angle vectors.
- A point is visible iff its angle is larger than all the preceding angles.