



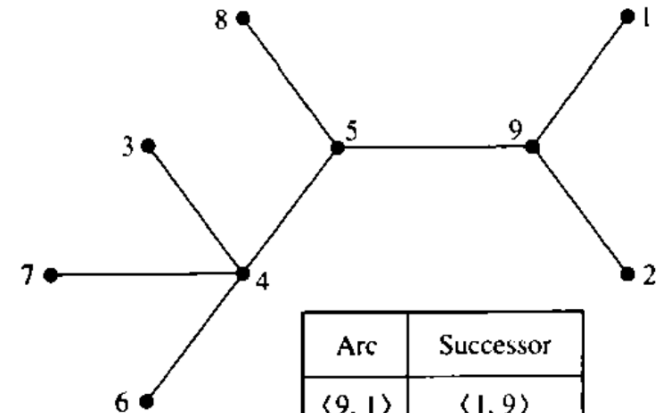
PRAM 3

Tree algorithms

CS121 Parallel Computing
Spring 2020

Euler tours

- An Euler tour of a graph is a cycle that goes through every edge of the graph.
 - It may go through a vertex multiple times.
- A connected, directed graph has an Euler tour if and only if the indegree and outdegree of every vertex are equal.
- Suppose we take an undirected graph, and for edge (u,v) , create two directed edges (u,v) and (v,u) .
 - Then every vertex has equal indegree and outdegree, and so has an Euler tour.
- Consider a tree where each edge has been doubled.
 - To find an Euler tour of the tree, first order the edges adjacent to each node arbitrarily.
 - Say the neighbors of a node v are ordered u_0, \dots, u_{d-1} . Then set the successor of edge (u_i, v) on the tour to $(v, u_{(i+1) \bmod d})$.
- The Euler tour of a tree can be computed in $O(1)$ parallel time.



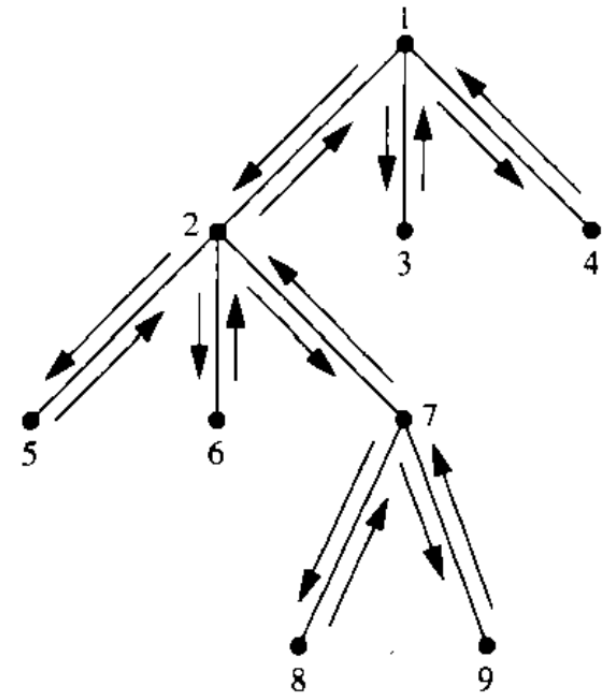
Arc	Successor
$\langle 9, 1 \rangle$	$\langle 1, 9 \rangle$
$\langle 9, 2 \rangle$	$\langle 2, 9 \rangle$
$\langle 4, 3 \rangle$	$\langle 3, 4 \rangle$
$\langle 5, 4 \rangle$	$\langle 4, 3 \rangle$
$\langle 3, 4 \rangle$	$\langle 4, 7 \rangle$
$\langle 7, 4 \rangle$	$\langle 4, 6 \rangle$
$\langle 6, 4 \rangle$	$\langle 4, 5 \rangle$
$\langle 8, 5 \rangle$	$\langle 5, 4 \rangle$
$\langle 4, 5 \rangle$	$\langle 5, 9 \rangle$
$\langle 9, 5 \rangle$	$\langle 5, 8 \rangle$
$\langle 4, 6 \rangle$	$\langle 6, 4 \rangle$
$\langle 4, 7 \rangle$	$\langle 7, 4 \rangle$
$\langle 5, 8 \rangle$	$\langle 8, 5 \rangle$
$\langle 5, 9 \rangle$	$\langle 9, 2 \rangle$
$\langle 2, 9 \rangle$	$\langle 9, 1 \rangle$
$\langle 1, 9 \rangle$	$\langle 9, 5 \rangle$

$\langle 9, 1 \rangle \rightarrow \langle 1, 9 \rangle \rightarrow \langle 9, 5 \rangle \rightarrow \langle 5, 8 \rangle \rightarrow \langle 8, 5 \rangle \rightarrow \langle 5, 4 \rangle \rightarrow$
 $\langle 4, 3 \rangle \rightarrow \langle 3, 4 \rangle \rightarrow \langle 4, 7 \rangle \rightarrow \langle 7, 4 \rangle \rightarrow \langle 4, 6 \rangle \rightarrow \langle 6, 4 \rangle \rightarrow$
 $\langle 4, 5 \rangle \rightarrow \langle 5, 9 \rangle \rightarrow \langle 9, 2 \rangle \rightarrow \langle 2, 9 \rangle \rightarrow \langle 9, 1 \rangle$

Source: Introduction to Parallel Algorithms, Jaja


Parallel tree operations

- Many operations on trees can be done in parallel using Euler tours and prefix sum.
- These operations in turn are used in other parallel graph algorithms.
- We first root a tree in parallel.
 - I.e. set an arbitrary node r as the tree's root. Then each node v needs to compute $p(v)$, its parent in the rooted tree.
 - To do this, assign a weight of 1 to each edge in an Euler tour of the tree.
 - Then compute the parallel prefix sum of the edges.
 - For each edge (u,v) , set $u=p(v)$ whenever the prefix sum of (u,v) is less than the prefix sum of (v,u) .
 - Thus, we can root a tree with n nodes in $O(\log n)$ time and $O(n)$ work.



Node depths

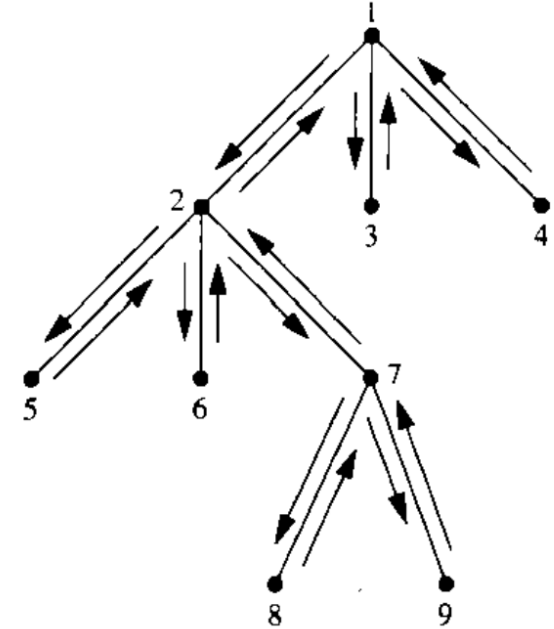
- For each node, compute its depth in a rooted tree.
 - For each node v , let $p(v)$ be its parent.
 - Set the weight of edge $(p(v), v)$ to 1, and the weight of edge $(v, p(v))$ to -1.
 - Compute a parallel prefix sum of the Euler tour starting at the root.
 - The depth of node v is the prefix sum of edge $(p(v), v)$.
- For a tree with n nodes, this takes $O(\log n)$ time using $O(n)$ work.



Postorder numbering

- Traverse a rooted tree in postorder, starting from the root r .
 - Start an Euler tour from r . For each node v , we want to visit v 's children in the order of the tour, then visit v itself.
 - Ex


v	1	2	3	4	5	6	7	8	9
$n(v)$	9	6	7	8	1	2	5	3	4
 - For each node v , set the weight of edge $(v, p(v))$ to 1, and the weight of $(p(v), v)$ weight 0.
 - Compute a parallel prefix sum of the Euler tour.
 - For each $v \neq r$, set $n(v)$ to the prefix sum of edge $(v, p(v))$. Set $n(r)=n$.
- In a tree with n nodes, we can compute the postorder numbering in $O(\log n)$ time and $O(n)$ work.



Euler Path	Weight	Prefix Sums
$\langle 1, 2 \rangle$	0	0
$\langle 2, 5 \rangle$	0	0
$\langle 5, 2 \rangle$	1	1
$\langle 2, 6 \rangle$	0	1
$\langle 6, 2 \rangle$	1	2
$\langle 2, 7 \rangle$	0	2
$\langle 7, 8 \rangle$	0	2
$\langle 8, 7 \rangle$	1	3
$\langle 7, 9 \rangle$	0	3
$\langle 9, 7 \rangle$	1	4
$\langle 7, 2 \rangle$	1	5
$\langle 2, 1 \rangle$	1	6
$\langle 1, 3 \rangle$	0	6
$\langle 3, 1 \rangle$	1	7
$\langle 1, 4 \rangle$	0	7
$\langle 4, 1 \rangle$	1	8

Number of descendant

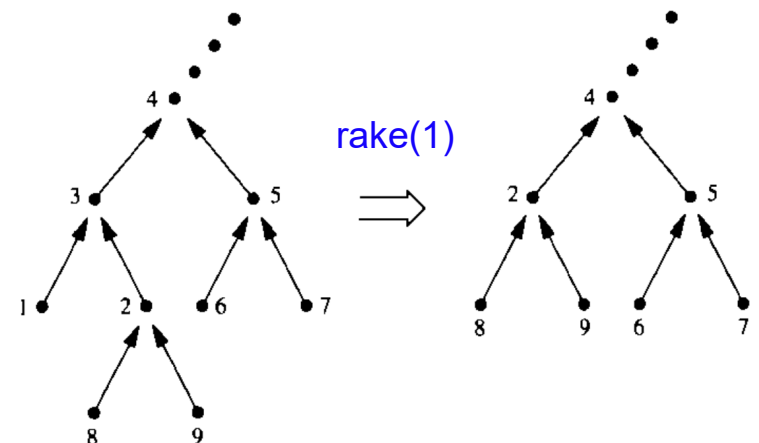
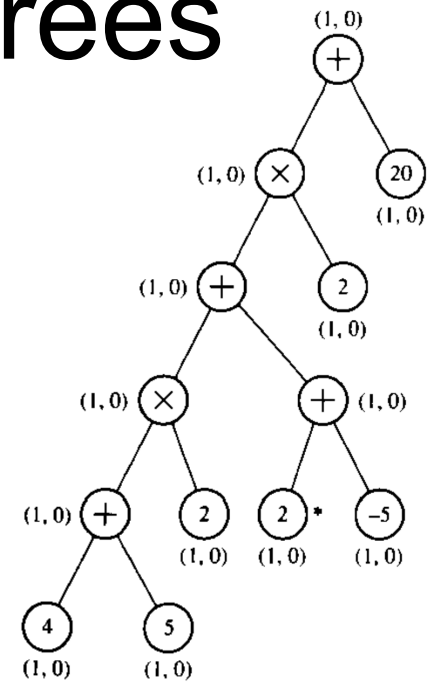
- For each node v in a rooted tree, compute the number of nodes in the subtree rooted at v .
- To do this, we compute prefix sums as in the postorder numbering.
- Then the number of descendants of a node v equals the prefix sum of $(v, p(v))$ minus the prefix sum of $(p(v), v)$.
- In a tree with n nodes, we can compute the number of descendants in $O(\log n)$ time and $O(n)$ work.



Euler Path	Weight	Prefix Sums
$\langle 1, 2 \rangle$	0	0
$\langle 2, 5 \rangle$	0	0
$\langle 5, 2 \rangle$	1	1
$\langle 2, 6 \rangle$	0	1
$\langle 6, 2 \rangle$	1	2
$\langle 2, 7 \rangle$	0	2
$\langle 7, 8 \rangle$	0	2
$\langle 8, 7 \rangle$	1	3
$\langle 7, 9 \rangle$	0	3
$\langle 9, 7 \rangle$	1	4
$\langle 7, 2 \rangle$	1	5
$\langle 2, 1 \rangle$	1	6
$\langle 1, 3 \rangle$	0	6
$\langle 3, 1 \rangle$	1	7
$\langle 1, 4 \rangle$	0	7
$\langle 4, 1 \rangle$	1	8

Evaluating expression trees

- An expression tree is a binary tree with values at the leaves and operators (+ or \times) in the interior nodes.
- We want to quickly evaluate an expression tree in parallel.
- The main tool is the rake operation.
 - Given a node u with sibling v , parent p and grandparent p' , $\text{rake}(u)$ removes u and p , and connects v with p' .
- We repeatedly rake an expression tree in parallel to contract it to 3 nodes with the same value as the original tree.





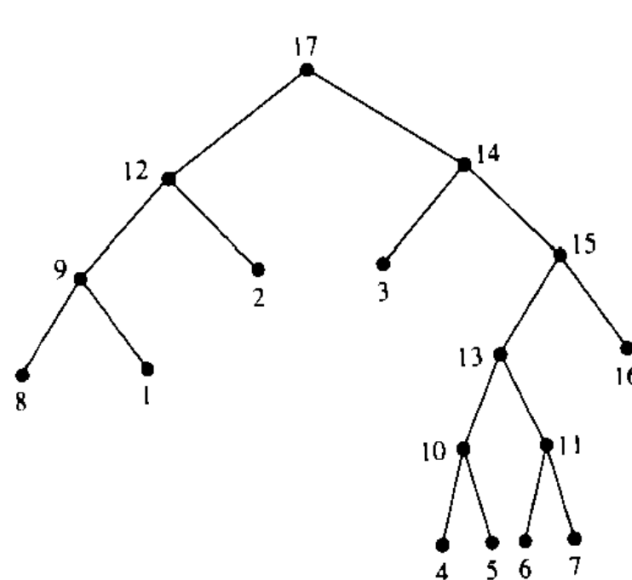
Raking in parallel

- When raking multiple nodes in parallel, we have to avoid concurrent changes to the same node.
- Given an expression tree, label the leaves from left to right (except for the first and last leaf), and call the set A .
 - Let A_{odd} and A_{even} be the subset of A with odd and even labels, resp.
- Repeat for $\lceil \log(n + 1) \rceil$ rounds, where $n = |A|$.
 - Concurrently rake all the leaves in A_{odd} that are left children.
 - Concurrently rake the rest of the leaves in A_{odd} .
 - Set $A = A_{even}$.

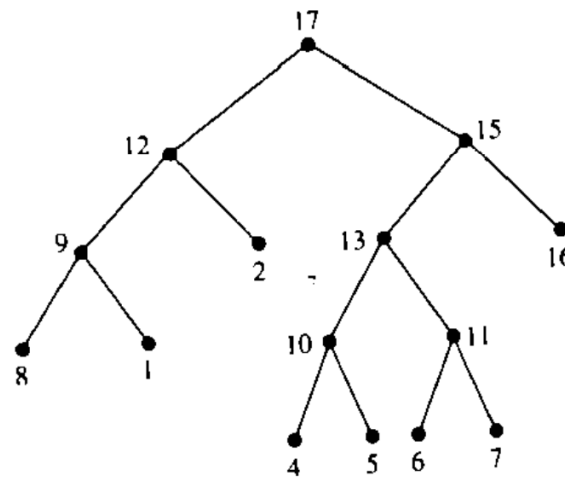
Example

- ❑ First rake node 3.
- ❑ Then rake nodes 1, 5, 7.
- ❑ Then rake leaves 2, 6.
- ❑ Then rake leaf 4.

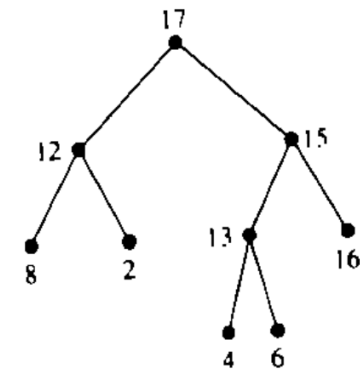
- ❑ Nodes raked concurrently don't have the same parent. So the raking process works correctly.
- ❑ Each round reduces the number of leaves by a factor of 2.
- ❑ After $\lceil \log(n + 1) \rceil$ rounds, there will be two leaves.




(a)



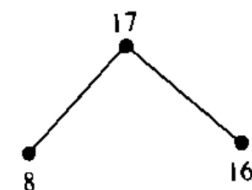
(b)



(c)




(d)



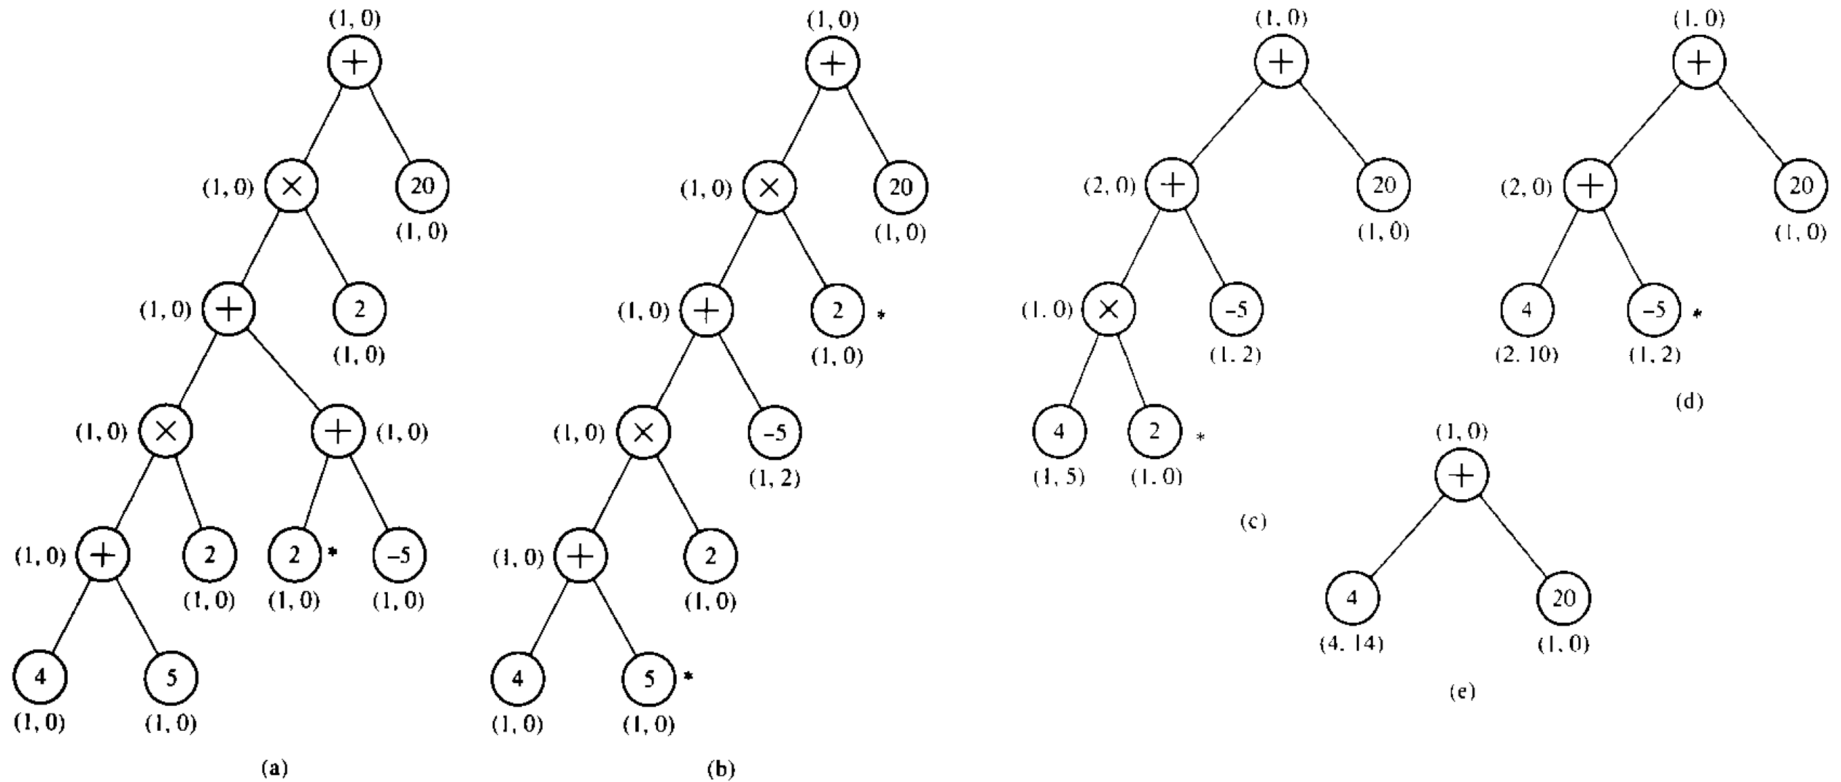
(e)

Parallel expression evaluation

- When raking an expression tree, combine raked node's value with its sibling's value.
- When raking many nodes in parallel, some nodes may not have values yet.
 - Imagine a chain graph with leaves hanging off the side. Nodes in the middle of the chain do not have values yet when they take part in a rake.
- Given a node v , we assign a label (a_v, b_v) to v .
 - If v is a leaf, v also has a numerical label x , and the numerical value of v is $a_v x + b_v$.
- When nodes are raked, we also update their labels.
 - **Ex** In example below, suppose w has a still undermined value X . Then after raking v , w 's value is $a_u(a_v c_v + b_v)(a_w X + b_w) + b_u = [a_u(a_v c_v + b_v)a_w]X + [a_u(a_v c_v + b_v)b_w + b_u]$.
 - Thus, the new label of w is $(a_u(a_v c_v + b_v)a_w, a_u(a_v c_v + b_v)b_w + b_u)$.
 - Leaves always have a numerical label, so their value can be immediately evaluated.
- Since raking in parallel reduces the tree to 3 nodes in $O(\log n)$ time, the tree's value can be evaluated in $O(\log n)$ time.



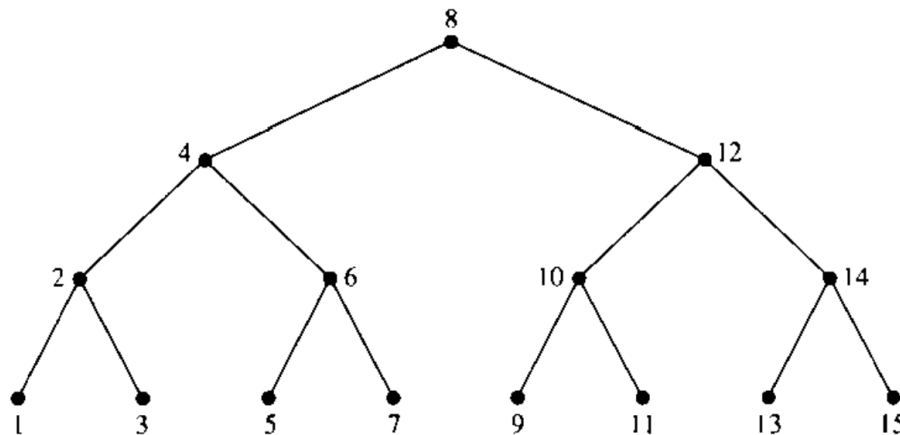
Example



Nodes marked by * are raked in parallel in each round.

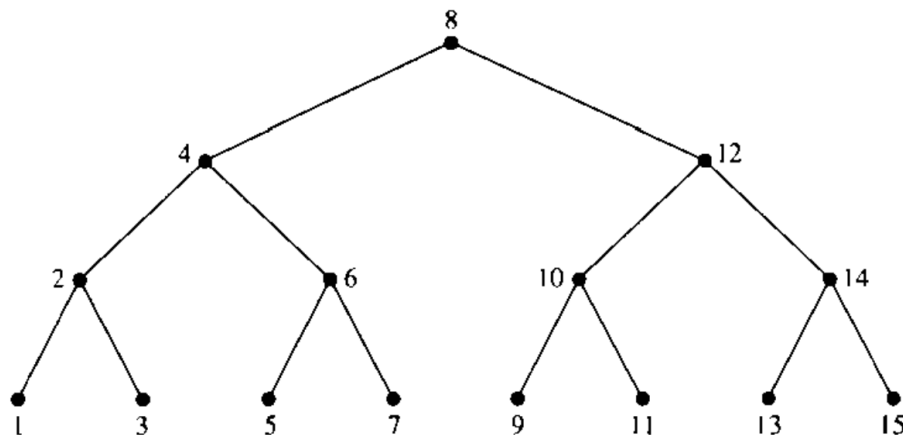
Lowest common ancestors

- Given two nodes u, v in a tree, $LCA(u, v)$ is the lowest node in the tree with u and v as descendants.
 - Ex $LCA(1, 5) = 4$, $LCA(3, 10) = 8$.
- Given a tree, we want to preprocess it so that LCA queries can be answered in $O(1)$ time.



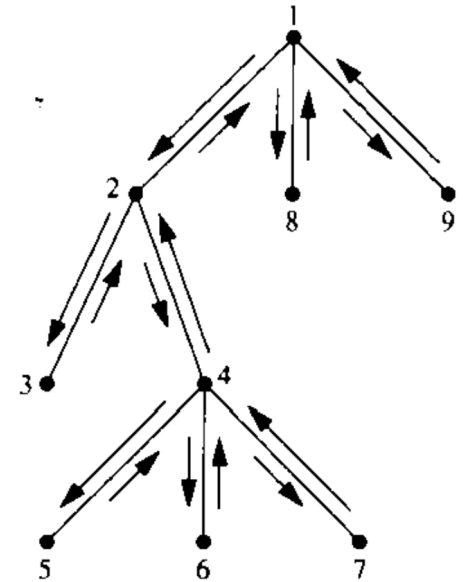
Simple cases

- For a path graph, the LCA of two nodes is whichever node is closer to the root, which can be computed in $O(1)$ time.
- For a complete binary tree, suppose all nodes are labeled by their inorder number.
 - Given u and v , write their labels in binary. Suppose $u = (z_1 z_2 \dots)_2$.
 - Let i be the most significant bit on which u and v differ.
 - Then $LCA(u, v) = (z_1 z_2 \dots z_{i-1} 10 \dots 0)_2$.
 - Ex $9 = (1001)_2$ and $13 = (1101)_2$. The MSB on which they differ is 2. So $LCA(9, 13) = (1100)_2 = 12$.
 - Thus, the LCA can be computed in $O(1)$ time.



LCA and Euler tours

- Consider a general tree with n nodes.
- First compute an Euler tour of the tree, in $O(1)$ parallel time.
 - Label nodes by their order of appearance in the tour.
- Next compute the depths of all nodes, in $O(\log n)$ parallel time.
- The Euler tour can be described by a size $2n-1$ array A listing the order of nodes visited.
- Create a corresponding array B giving the levels (i.e. depths) of the nodes in A .

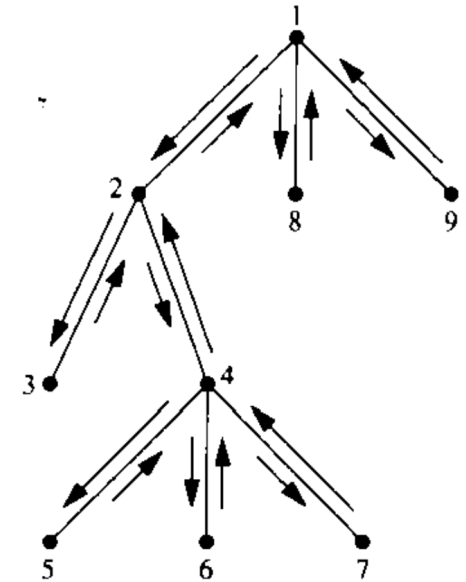


$A = (1, 2, 3, 2, 4, 5, 4, 6, 4, 7, 4, 2, 1, 8, 1, 9, 1)$

$B = (0, 1, 2, 1, 2, 3, 2, 3, 2, 3, 2, 1, 0, 1, 0, 1, 0)$

LCA and Euler tours

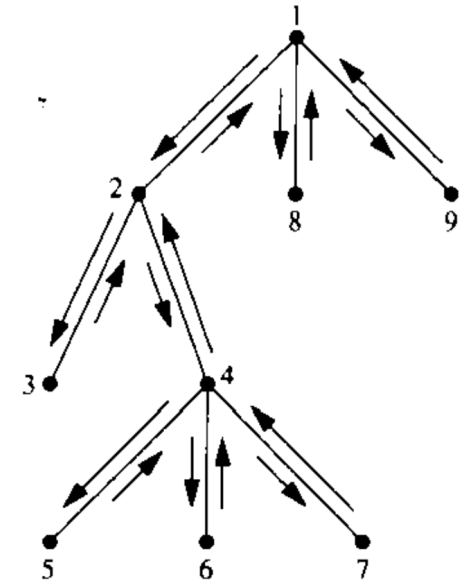
- Given a node v , define the following
 - $level(v)$ is the level of v .
 - $l(v), r(v)$ are the indices of the first, resp. last occurrence of v in A .
 - Ex $l(4) = 4, r(4) = 10$.
- **Lem** Given a node v , we have the following.
 - $l(v) = i$ iff $level(A[i - 1]) = level(v) - 1$.
 - $r(v) = i$ iff $level(A[i + 1]) = level(v) - 1$.
- **Ex** $l(4) = 4$, and $level(A[3]) = level(2) = 1 = level(4) - 1$.
- **Ex** $r(4) = 10$, and $level(A[11]) = level(2) = 1 = level(4) - 1$.
- Assume A and B are given. Then we can compute $l(v)$ and $r(v)$ for all nodes v in parallel $O(1)$ time.
 - Use one processor for each node in A and check the conditions in the lemma.
 - If either condition holds for a node v , record $l(v)$ or $r(v)$.



$A = (1, 2, 3, 2, 4, 5, 4, 6, 4, 7, 4, 2, 1, 8, 1, 9, 1)$
 $B = (0, 1, 2, 1, 2, 3, 2, 3, 2, 3, 2, 1, 0, 1, 0, 1, 0)$

LCA and Euler tours

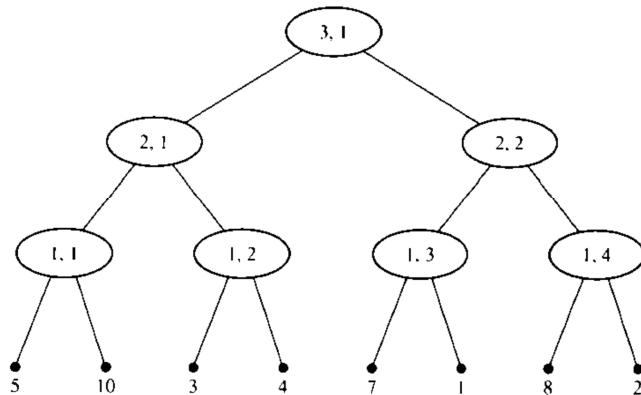
- **Lem** Given nodes u and v in the tree, the following properties hold
 - u is an ancestor of v if and only if $l(u) < l(v) < r(u)$.
 - If $r(u) < l(v)$, then $\text{LCA}(u,v)$ is the node with the minimum level in the interval $[r(u), l(v)]$ of A .
- **Ex** 2 is an ancestor of 4, and $l(2) = 1 < l(4) = 4 < r(2) = 11$.
- **Ex** $r(4) = 10 < l(8) = 13$, and $\text{LCA}(4,8)$ is the node with min level in $A[10:13]$, namely node 1 with level 0.
- Using the lemma, we can find LCAs in $O(1)$ time, if we can find the node with min level in second property of lemma in $O(1)$ time.
- In particular, given an array X and any interval $[i, j]$, we want to find the min value in $X[i:j]$ in $O(1)$ time.



$A = (1, 2, 3, 2, 4, 5, 4, 6, 4, 7, 4, 2, 1, 8, 1, 9, 1)$
 $B = (0, 1, 2, 1, 2, 3, 2, 3, 2, 3, 2, 1, 0, 1, 0, 1, 0)$

Range minima problem

- Given an array X of size n , preprocess X so that given any $1 \leq i, j \leq n$, we can find $Rmin[i, j] = \min(X_i, \dots, X_j)$ in $O(1)$ time.
 - The preprocessing will take $O(\log n)$ parallel time and $O(n \log n)$ work.
- Assume for simplicity n is a power of 2.
- Form a complete binary tree with the values of X as the leaves.
 - Label each node in the tree by its height and index within a layer.
- For each node (h, j) , create two arrays $P(h, j)$ and $S(h, j)$.
 - Let the values of the leaf nodes in the subtree rooted at (h, j) be X_p, X_{p+1}, \dots, X_q .
 - $P(h, j)$ is the array of prefix minima of $(X_p, X_{p+1}, \dots, X_q)$, i.e. $P(h, j)[k] = \min(X_p, \dots, X_{p+k-1})$.
 - $S(h, j)$ is the array of suffix minima of $(X_q, X_{q-1}, \dots, X_p)$, i.e. $S(h, j)[k] = \min(X_q, X_{q-1}, \dots, X_{q-k+1})$.



Ex Array $X = [5, 10, 3, 4, 7, 1, 8, 2]$.

$$\begin{aligned}
 P(2, 1) &= (5, 5, 3, 3), & S(2, 1) &= (4, 3, 3, 3) \\
 P(2, 2) &= (7, 1, 1, 1), & S(2, 2) &= (2, 2, 1, 1) \\
 P(3, 1) &= (5, 5, 3, 3, 3, 1, 1, 1) \\
 S(3, 1) &= (2, 2, 1, 1, 1, 1, 1, 1)
 \end{aligned}$$

Computing range minima

- Given the P and S arrays, we can compute range minima in $O(1)$ time.
- Given an interval $[i,j]$, let $v = LCA(i, j)$.
 - Since X is represented by a complete binary tree, LCA can be computed in $O(1)$ time.
 - Let u and w be the left, resp. right child of v.
 - Let X_i be the p 'th node counting from the right in the subtree rooted at u.
 - Let X_j be the q 'th node counting from the left in the subtree rooted at w.
 - Then $Rmin[i, j] = \min(S(u)[p], P(w)[q])$.
- **Ex** To find $Rmin[2,5]$, we have $LCA(2,5)$ is the root (3,1).
 - Then $u = (2,1)$, $w = (2,2)$. Also, $S(u) = (4,3,3,3)$ and $P(w) = (7,1,1,1)$.
 - Leaf 2 is the second child from the right in u's subtree, and leaf 5 is the second child from the left in w's subtree.
 - So $Rmin[2,5] = \min(3,4,7,1) = \min(S(u)[2], P(w)[2]) = \min(3,1) = 1$.

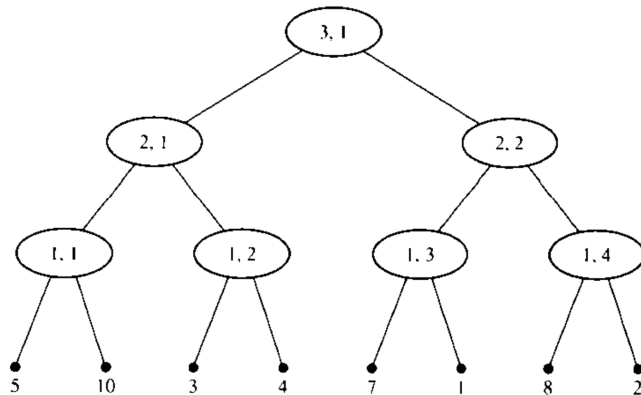
- Given the P and S arrays, we can compute range minima in $O(1)$ time.
- Given an interval $[i,j]$, let $v = LCA(i, j)$.
 - Since X is represented by a complete binary tree, LCA can be computed in $O(1)$ time.
 - Let u and w be the left, resp. right child of v.
 - Let X_i be the p 'th node counting from the right in the subtree rooted at u.
 - Let X_j be the q 'th node counting from the left in the subtree rooted at w.
 - Then $Rmin[i, j] = \min(S(u)[p], P(w)[q])$.
- **Ex** To find $Rmin[2,5]$, we have $LCA(2,5)$ is the root (3,1).
 - Then $u = (2,1)$, $w = (2,2)$. Also, $S(u) = (4,3,3,3)$ and $P(w) = (7,1,1,1)$.
 - Leaf 2 is the second child from the right in u's subtree, and leaf 5 is the second child from the left in w's subtree.
 - So $Rmin[2,5] = \min(3,4,7,1) = \min(S(u)[2], P(w)[2]) = \min(3,1) = 1$.

Ex Array $X = [5,10,3,4,7,1,8,2]$.

$$\begin{aligned}
 P(2,1) &= (5,5,3,3), & S(2,1) &= (4,3,3,3) \\
 P(2,2) &= (7,1,1,1), & S(2,2) &= (2,2,1,1) \\
 P(3,1) &= (5,5,3,3,3,1,1,1) \\
 S(3,1) &= (2,2,1,1,1,1,1,1)
 \end{aligned}$$

Computing P and S arrays

- We compute the P and S arrays from the bottom up.
- For a leaf with value X_i , the P and S arrays for it are just (X_i) .
- Given a node v , let u and w be its left and right children, and suppose we've computed the P and S arrays for u and w .
 - Let x_u be the last value in $P(u)$, and x_w be the last value in $S(w)$.
 - Let $P'(w)$ be the elementwise min of $P(w)$ and x_u .
 - Let $S'(u)$ be the elementwise min of $S(u)$ and x_w .
- Then $P(v) = P(u) \circ P'(w)$, i.e. $P(u)$ followed by $P'(w)$. Also, $S(v) = S(w) \circ S'(u)$.

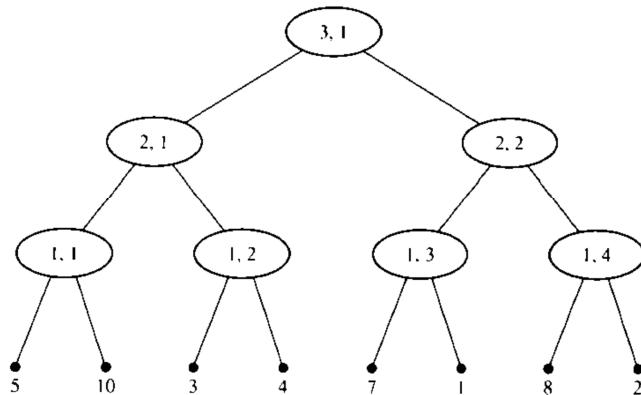


Ex Array $X = [5, 10, 3, 4, 7, 1, 8, 2]$.

$$\begin{aligned}
 P(2,1) &= (5, 5, 3, 3), & S(2,1) &= (4, 3, 3, 3) \\
 P(2,2) &= (7, 1, 1, 1), & S(2,2) &= (2, 2, 1, 1) \\
 P(3,1) &= (5, 5, 3, 3, 3, 1, 1, 1) \\
 S(3,1) &= (2, 2, 1, 1, 1, 1, 1, 1)
 \end{aligned}$$

Computing P and S arrays

- **Ex** Suppose we computed the P and S arrays for nodes $u = (2,1)$ and $w = (2,2)$, and want to compute it for node $v = (3,1)$.
 - $x_u = 3$ and $x_w = 1$, so $P'(w) = (3,1,1,1)$ and $S'(u) = (1,1,1,1)$.
 - Then $P(3,1)$ and $S(3,1)$ are as shown below.
- **Thm** Given a tree with n leaves, all the P and S arrays in the tree can be computed in $O(\log n)$ parallel time and $O(n \log n)$ work.
- **Proof** Given two P (resp. S) arrays of size k , computing the parent P (resp. S) array takes $O(1)$ parallel time and $O(k)$ work.
 - On each layer of the tree, the total sizes of all P and S arrays is $O(n)$.
 - So can compute each layer in $O(1)$ parallel time and $O(n)$ work.
 - There are $O(\log n)$ layers. So the theorem follows.
- The work can be reduced to $O(n)$ using accelerated cascading.



Ex Array $X = [5,10,3,4,7,1,8,2]$.

$$\begin{aligned}
 P(2,1) &= (5,5,3,3), & S(2,1) &= (4,3,3,3) \\
 P(2,2) &= (7,1,1,1), & S(2,2) &= (2,2,1,1) \\
 P(3,1) &= (5,5,3,3,3,1,1,1) \\
 S(3,1) &= (2,2,1,1,1,1,1,1)
 \end{aligned}$$