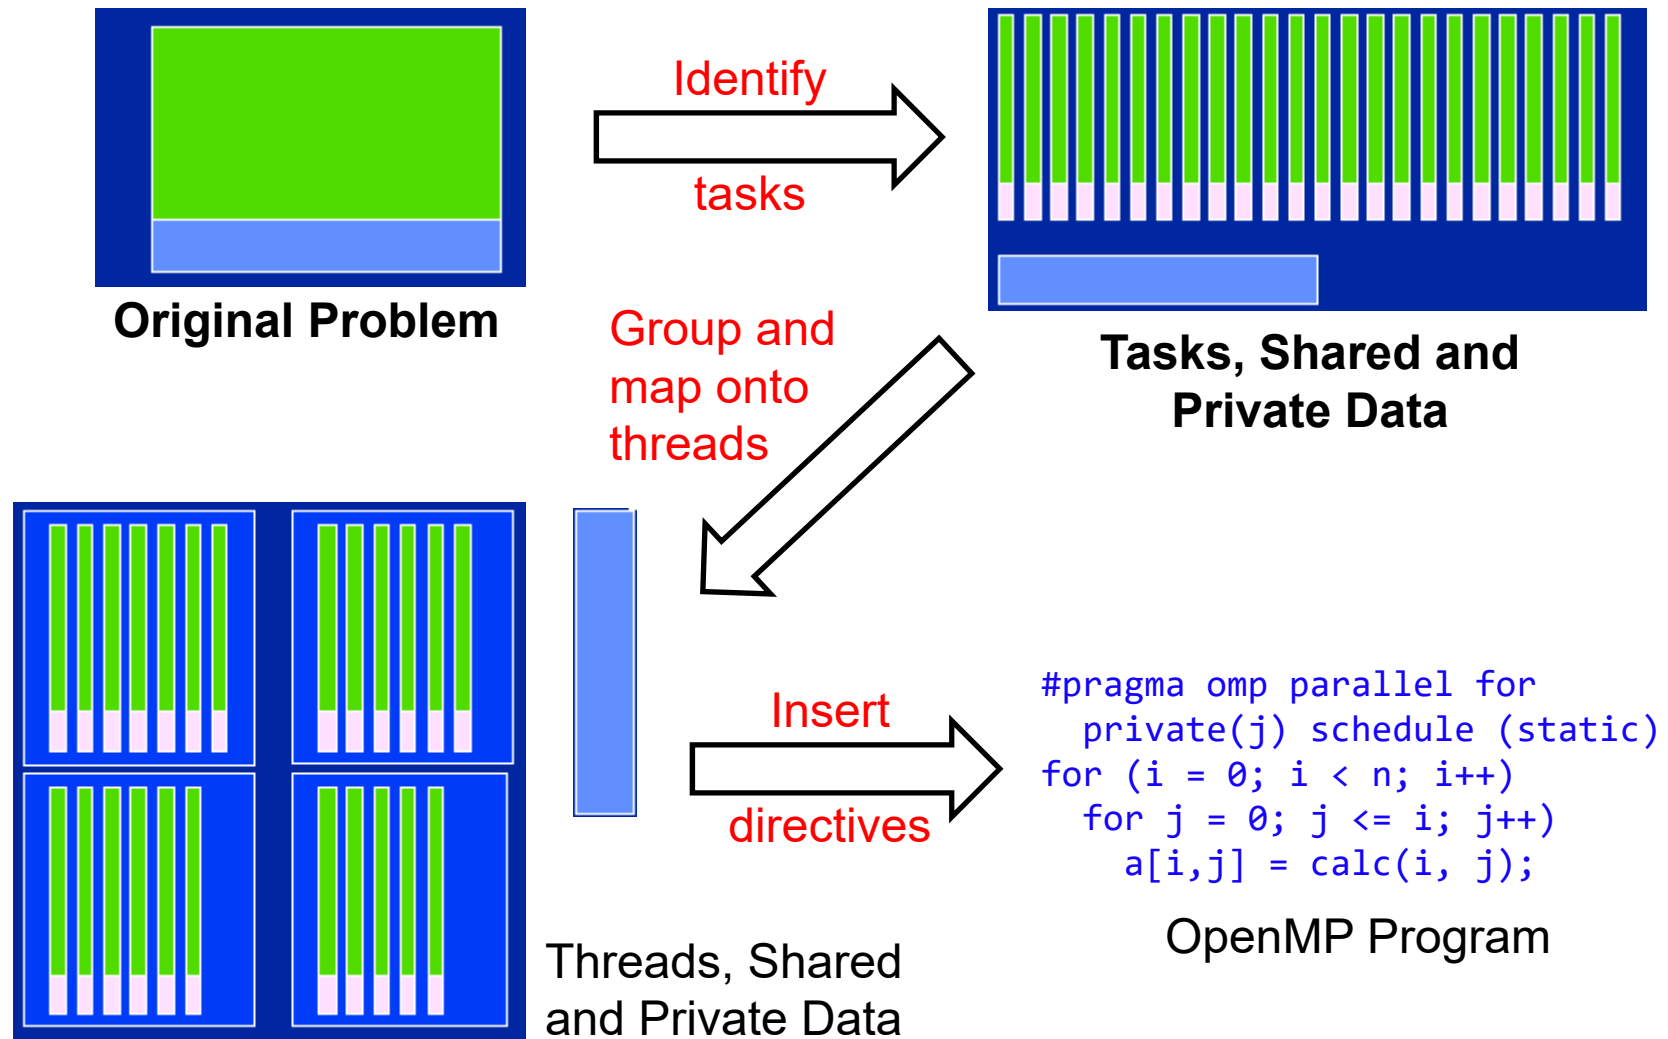




Loop Parallelism

CS121 Parallel Computing
Spring 2021

Shared memory algorithm design





Design considerations

- Break program into tasks, consisting of statements that must be executed in order.
 - Use data dependence analysis.
- Map independent tasks to different processors.
 - Mapping needs to consider load balancing, e.g. static vs dynamic, block vs cyclic work assignment.
- Variable specification
 - Shared vs. private vs. reduction
 - Shared variables cause cache coherence traffic and much lower performance.
 - Private and reduction variables don't need synchronization (except possibly at end of a loop).
- Dimension mapping, e.g. row-wise vs column-wise.
 - Matching mapping to access pattern improves cache locality.

Data dependence analysis

- Let S1 and S2 be two statements in a sequential execution of a program.
- S1 and S2 are independent if running them in different orders produces the same result. Otherwise they're dependent.
- Suppose S1 occurs before S2. They can have the following dependencies.
 - $S1 \rightarrow^T S2$ denotes true dependence (RAW), i.e. S1 writes to a location that is read by S2.
 - $S1 \rightarrow^A S2$ denotes anti dependence (WAR), i.e. S1 reads a location written by S2.
 - $S1 \rightarrow^O S2$ denotes output dependence (WAW), i.e. S1 writes to the same location written by S2.

```
S1: x = 2;  
S2: z = 3;  
S3: y = x;  
S4: y = x + z;  
S5: z = 6;
```

□ Dependences

- $S1 \rightarrow^T S3$
- $S1 \rightarrow^T S4$
- $S2 \rightarrow^T S4$
- $S2 \rightarrow^O S5$
- $S3 \rightarrow^O S4$
- $S4 \rightarrow^A S5$



Parallelizing loops

- When parallelizing a program, must ensure dependent statements run in the same order in the sequential and parallel programs.
 - Guarantees the parallel and sequential programs behave the same way.
 - A parallel program may run correctly without satisfying this condition, but it's not guaranteed.
 - The ordering requirement is transitive. I.e. if $S1 \rightarrow^* S2$ and $S2 \rightarrow^* S3$, then $S1$ must run before $S3$ in any parallelization.
- Dependent statements cannot on run on different processors, since we can't enforce the order of execution (interleaving) on different processors.
- Independent statements can run on different processors if they haven't been ordered by transitivity.



Parallelizing loops

- Goal is to identify all independent statements, to maximize parallelism.
- Focus on parallelizing loops, since these are common in shared memory programs and are the main performance hotspots.
- Notation
 - Let S denote a statement in the source program.
 - Given a nested loop with iteration variables i, j, \dots , let $S[i, j, \dots]$ denote a statement in loop iteration $[i, j, \dots]$.

Loop dependence analysis

- Loop-carried dependence
 - Dependence exists across different iterations of loop.
- Loop-independent dependence
 - Dependence exists within the same iteration of loop.

```
for (i=1; i<=n; i++) {  
    S1: a[i] = a[i-1] + 1;  
    S2: b[i] = a[i];  
}
```

$S1[i] \rightarrow_T S1[i+1]$
- loop-carried dependence
 $S1[i] \rightarrow_T S2[i]$
- loop-independent dependence

```
for (i=1; i<=n; i++)  
    for (j=1; j<=n; j++)  
        S3: a[i][j] = a[i][j-1] + 1;
```

$S3[i,j] \rightarrow_T S3[i,j+1]$
- loop-carried dependence in j loop
- no loop-carried dependence in i loop

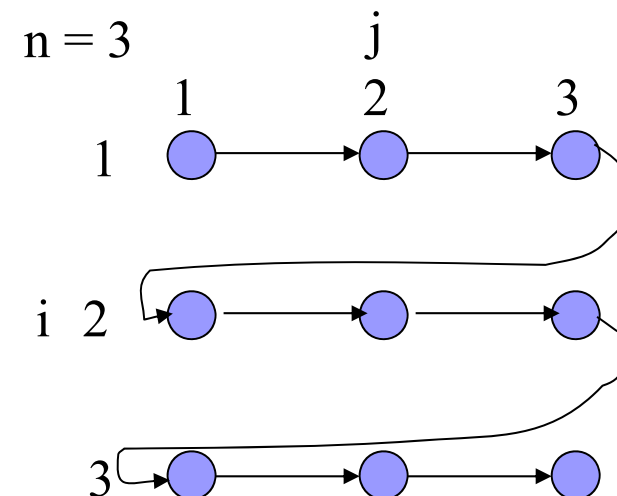
```
for (i=1; i<=n; i++)  
    for (j=1; j<=n; j++)  
        S4: a[i][j] = a[i-1][j] + 1;
```

$S4[i,j] \rightarrow_T S4[i+1,j]$
- loop-carried dependence in i loop
- no loop-carried dependence in j loop

Iteration-space traversal graph

- Iteration-space traversal graph (ITG) is a line graph showing the order of traversal in the iteration space.
- Node in ITG is a point in the iteration space, i.e. a particular iteration.
- Directed edge in ITG gives the next iteration that will be executed after the current iteration.

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++)  
    S3: a[i][j] = a[i][j-1] + 1;
```



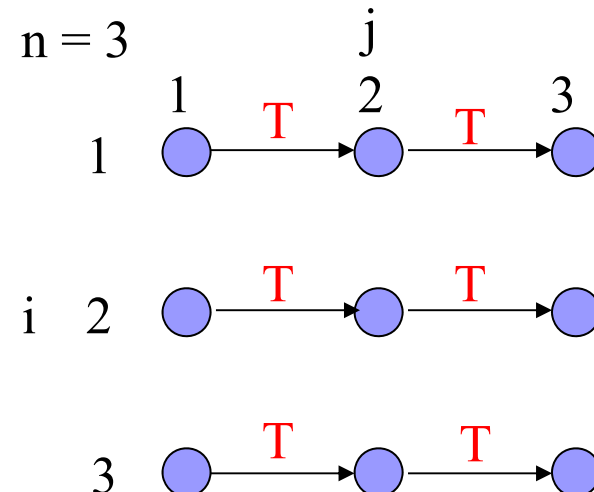
Loop-carried dependence graph

- Given the ITG, can determine the dependence between different loops.
- Loop-carried Dependence Graph (LDG) shows the loop-carried true/anti/output dependence relationships.
- Node in LDG is a point in the iteration space.
- Directed edge in LDG is the dependence.
- LDG helps identify parts of the loop that can be done in parallel.

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++)  
    S3: a[i][j] = a[i][j-1] + 1;
```

$S3[i,j] \rightarrow T S3[i,j+1]$

- loop-carried dependence in j loop
- no loop-carried dependence in i loop



Example 1

```

for (i=1; i<=n; i++)
  for (j=1; j<=n; j++) {
    S1: a[i][j] = b[i][j] + c[i][j];
    S2: b[i][j] = a[i][j-1] * d[i][j];
  }

```

True dependences

$S1[i,j] \rightarrow T S2[i,j+1]$

- loop-carried
dependence

Anti dependences

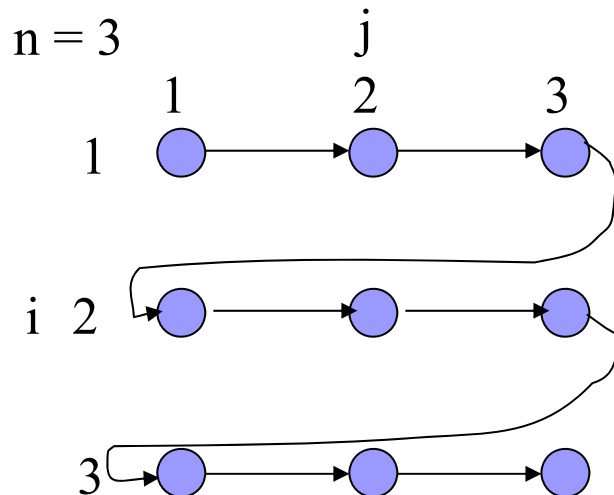
$S1[i,j] \rightarrow A S2[i,j]$

- loop-independent
dependence

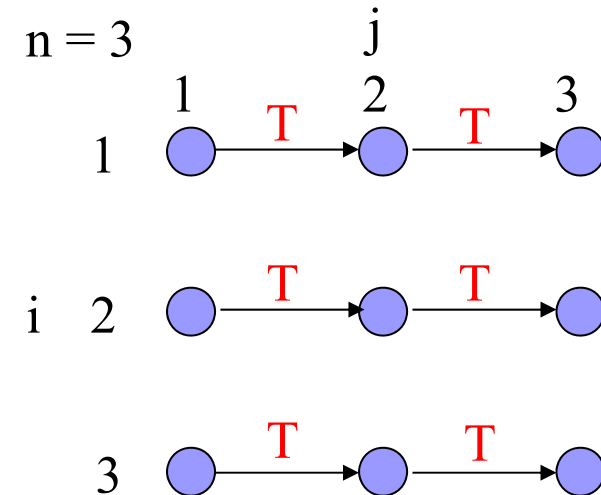
Output dependences

None

■ ITG



■ LDG





Example 1

- Task Identification

- n parallel tasks - one task per iteration of i loop.

- Grouping / mapping

- Static block as different iterations have same work.

- OpenMP

```
#pragma omp parallel for private(j) schedule(static)
for (i=1; i<=n; i++)
    for (j=1; j<=n; j++) {
        S1: a[i][j] = b[i][j] + c[i][j];
        S2: b[i][j] = a[i][j-1] * d[i][j];
    }
```

Example 2

```
for (i=1; i<=n; i++)  
  S1: a[i] = a[i-2];
```

True dependences

$S1[i] \rightarrow T S1[i+2]$

- loop-carried dependence

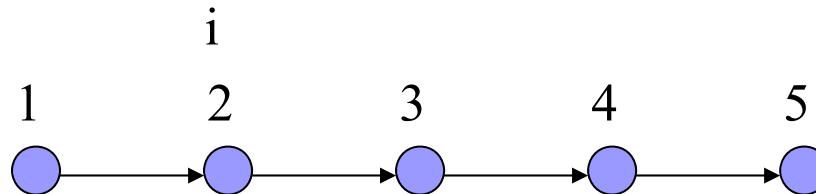
Anti dependences

None

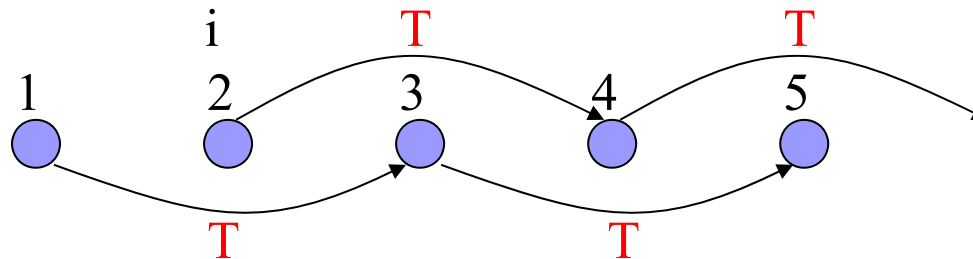
Output dependences

None

■ ITG



■ LDG





Example 2

- Task Identification

- There are opportunities when some dependences are missing.
- Can divide the for loop into two parallel tasks, one with even iterations and another with odd iterations.

- Grouping / mapping

- One task per thread

- OpenMP

```
#pragma omp parallel sections
private(i)
{
    #pragma omp section
    for (i=1; i<=n; i+=2)
        a[i] = a[i-2];
    #pragma omp section
    for (i=2; i<=n; i+=2)
        a[i] = a[i-2];
}
```

Example 3

```
for (i=1; i<=n; i++)  
  S1: a[i] = a[i-1] + b[i]*c[i] + d[i];
```

True dependences

$S1[i] \rightarrow T S1[i+1]$

- loop-carried
dependence

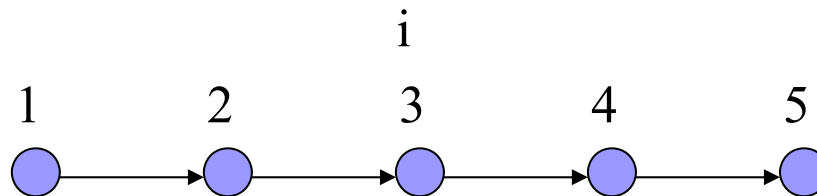
Anti dependences

None

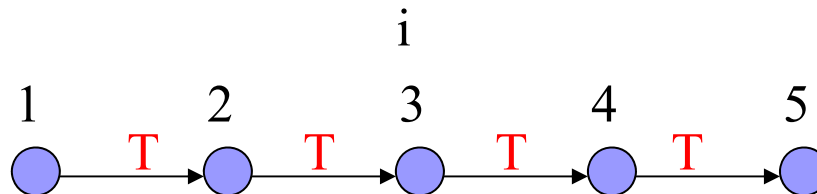
Output dependences

None

■ ITG



■ LDG



Must respect loop-carried dependence

Example 3

■ Task Identification

- Any opportunities for parallelism? Loop-carried dependence $S[i] \rightarrow T S[i+1]$ must be respected.
- But no loop-carried dependence in $b[i]*c[i]+d[i]$ part.
- Loop fission
 - Distribute into two separate loops .

■ Code after loop fission

```
for (i=1; i<=n; i++) {  
    S1: temp[i] = b[i]*c[i] + d[i];  
}  
  
for (i=1; i<=n; i++) {  
    S2: a[i] = a[i-1] + temp[i];  
}
```

No dependences in first loop,
so can be parallelized.

In second loop

True dependences:

$S2[i] \rightarrow T S2[i+1]$

- loop-carried dependence

Anti dependences None

Output dependences None



Example 3

□ OpenMP

```
#pragma omp parallel for schedule(static)
for (i=1; i<=n; i++) {
    temp[i] = b[i]*c[i] + d[i];
}
for (i=1; i<=n; i++) {
    a[i] = a[i-1] + temp[i];
}
```

- Note array temp[] introduces storage overhead.

□ Better OpenMP solution

```
#pragma omp parallel for ordered private(t) schedule(static,1)
for (i=1; i<=n; i++) {
    t = b[i]*c[i] + d[i]; /* one copy of t per thread */
    #pragma omp ordered
    a[i] = a[i-1] + t; }
```

- ordered statement enforces a[i] assignment ordering.
- With k threads, uses k extra storage.
- Typically $k \ll n$ so we save space.

Example 4

```
for (i=1; i<=n; i++) {  
  S1: a[i] = b[i+1]*a[i];  
  S2: b[i] = b[i]*coef;  
  S3: c[i] = 0.5*(c[i] + a[i]);  
  S4: d[i] = d[i-1] + d[i];  
}
```

True dependencies

$S1[i] \rightarrow T S3[i]$

-loop-independent dependence

$S4[i] \rightarrow T S4[i+1]$

-loop-carried dependence

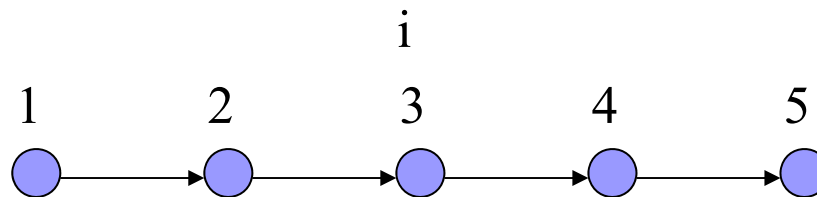
Anti dependencies

$S1[i] \rightarrow A S2[i+1]$

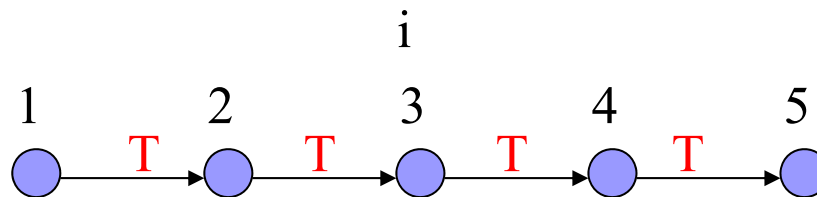
-loop-carried dependence

Output dependencies None

■ ITG



■ LDG



Must respect loop-carried dependence



Example 4

■ Task Identification

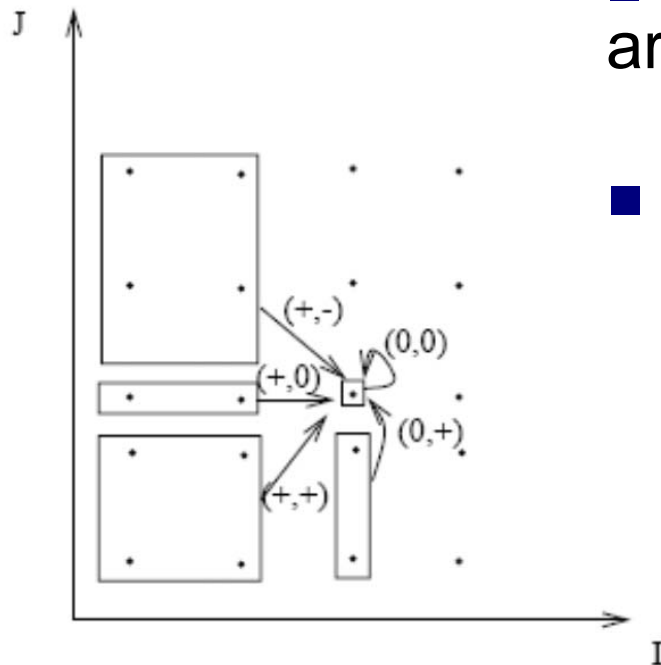
- S4 has no dependences with other statements, so can distribute into two separate loops (loop fission).
- This gives two parallel tasks, a loop containing S1, S2 and S3 and a loop containing S4.

■ OpenMP

```
#pragma omp parallel sections
private(i)
{
    #pragma omp section
    for (i=0; i<=n; i++) {
        S1: a[i] = b[i+1]*a[i];
        S2: b[i] = b[i]*coef;
        S3: c[i] = 0.5*(c[i] + a[i]);
    }
    #pragma omp section
    for (i=0; i<n; i++) {
        S4: d[i] = d[i-1] + d[i];
    }
}
```

Distance and direction vectors

- Let T_1 and T_2 be iterations s.t. $T_1 \rightarrow^* T_2$.
- Distance vector from T_1 to T_2 is $T_2 - T_1$.
- Direction vector from T_1 to T_2 is $\text{sign}(T_2 - T_1)$.
- Consider a nested loop over (i,j) (j is the inner loop).



- The following direction vectors are possible

- ☐ $(+, +), (+, 0), (+, -), (0, +), (0, 0)$.

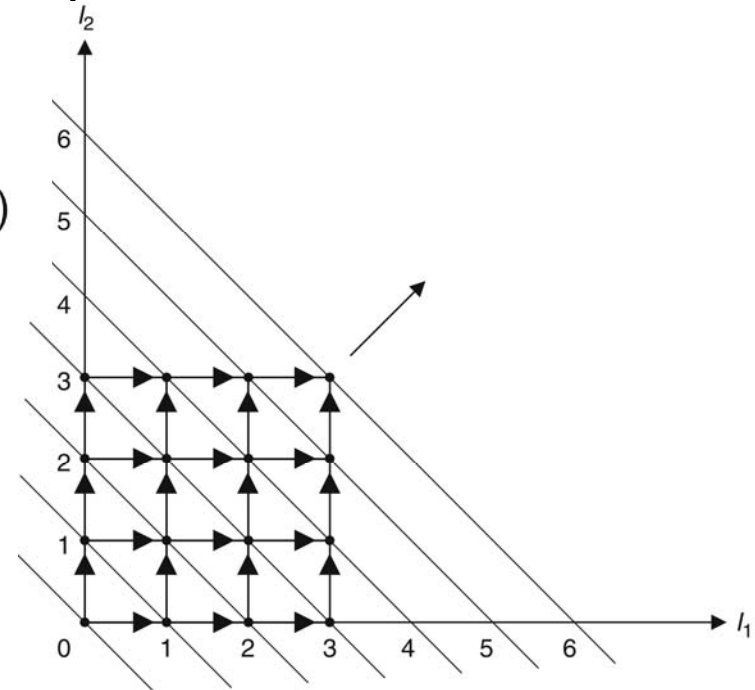
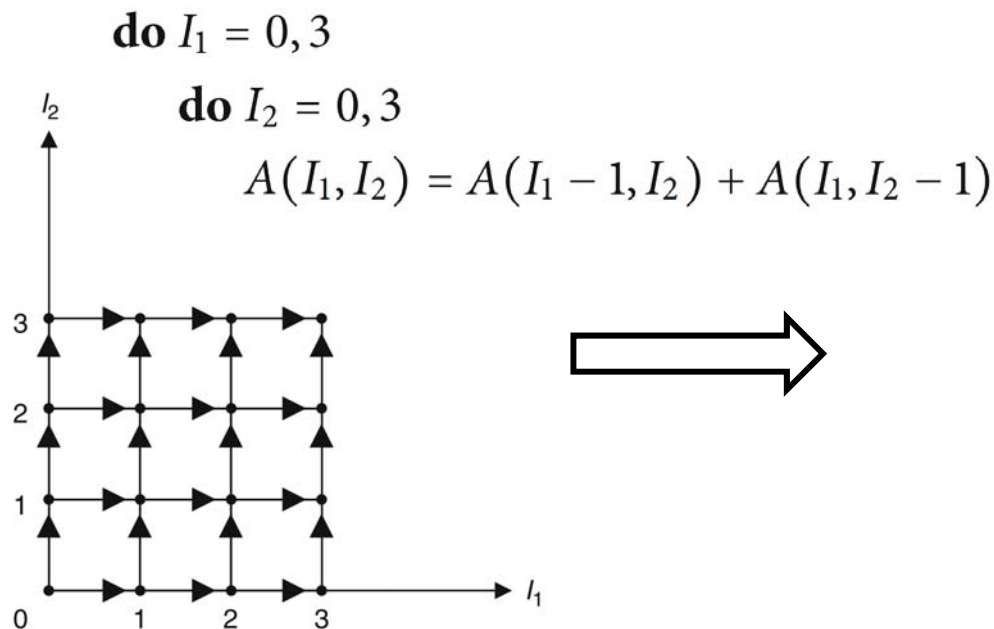
- The following directions are not possible.

- ☐ $(0, -), (-, +), (-, 0), (-, -)$.

- ☐ **Ex** Direction vector $(-, +)$ would mean, e.g. iteration (i, j) depends on iteration $(i+1, j-1)$. But (i, j) occurs before $(i+1, j-1)$.

Loop skewing

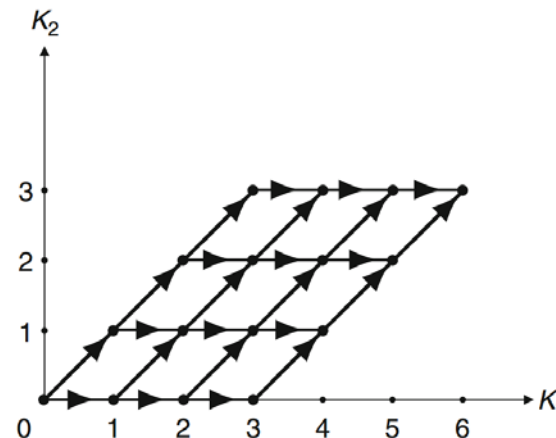
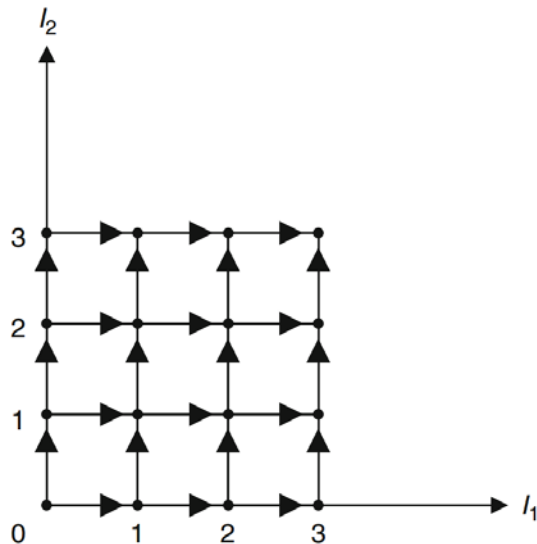
- If a loop has dependencies that prevent it being parallelized, we can try to transform (skew) the loop indices to create parallelizable loops.
- **Ex** Create new loop indices, the outer loop across the diagonal lines, the inner loop along the diagonal lines.
 - Inner loop iterations can be done in parallel.



Source: Unimodular Transformations, Utpal Banerjee

Unimodular transformations

- Can define new indices using a linear transformation.
 - We use unimodular linear transformations, where the matrix has determinant -1 or 1.
 - Unimodularity ensures all iterations in original loop are performed in the transformed loop.
- Ex Let $U = \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix}$, and $(K_1, K_2) = (I_1, I_2)U = (I_1 + I_2, I_2)$.
 - Let K_1, K_2 be the outer and inner loops, resp. Then for each K_1 iteration, can run all the K_2 iterations in parallel.



Extracting parallelism

- Transformed loop must be legal. Also, we want it to be parallelizable.
- **Thm** Let U be a unimodular transformation. If vU is legal for all distance vectors v , then U is a legal transformation.
- **Ex** For $U = \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix}$, $(0,1)U = (1,1)$, $(1,0)U = (1,0)$.
 - New direction vectors are $(+,+)$ and $(+,0)$, so the loop with the new loop indices is legal.
- **Thm** Suppose all the direction vectors for a loop are $+$ in the i 'th coordinate, for some i . Then all loops deeper than level i can be run in parallel.
- **Ex** For the above U , direction vectors are all $+$ in K_1 coordinate, so the K_2 loop can be parallelized.

$$A[I_1, I_2] = A[I_1-1, I_2] + A[I_1, I_2-1]$$



$$A[K_1, K_2] = A[K_1-1, K_2-1] + A[K_1-1, K_2]$$



Data accesses after skewing

- Since $(K_1, K_2) = (I_1, I_2)U$, then $(I_1, I_2) = (K_1, K_2)U^{-1}$.
- So in iteration (K_1, K_2) of the transformed loop, we access data from iteration $(K_1, K_2)U^{-1}$ of the original loop.
- **Ex** For $U = \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix}$, $U^{-1} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix}$, and $(K_1, K_2)U^{-1} = (K_1 - K_2, K_2) = (I_1, I_2)$.
So in iteration (K_1, K_2) , we do

$$\begin{aligned} A[I_1, I_2] &= A[K_1 - K_2, K_2] = \\ &A[K_1 - K_2 - 1, K_2] + A[K_1 - K_2, K_2 - 1] \end{aligned}$$



- $$+ A(K_1 - K_2, K_2 - 1)$$



Algorithmic analysis

- Sometimes there is no way to restructure a loop to increase parallelism.
- We can try to restructure the algorithm to eliminate dependences and improve parallelism.
- Need to understand the purpose of the algorithm and how it is used.
- For example, some algorithms are nondeterministic or calculate an approximation (e.g., Jacobi iteration).
 - In this case, restructuring the algorithm or ignoring some dependences may still give a valid result.

Fibonacci numbers

- Fibonacci numbers

- $F_1 = F_2 = 1.$

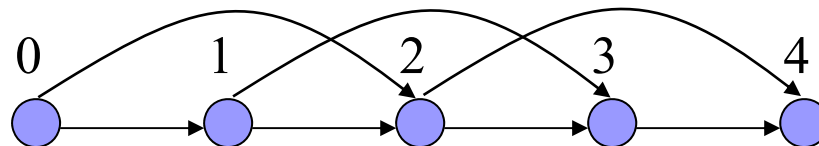
- $F_n = F_{n-1} + F_{n-2},$ for $n > 2.$

- $1, 1, 2, 3, 5, 8, 13, 21, 34, \dots$

- Computing F_n sequentially takes $O(n)$ time.

- Can we compute F_n in parallel?

- Looking at the LDG, it seems there's no parallelism available.

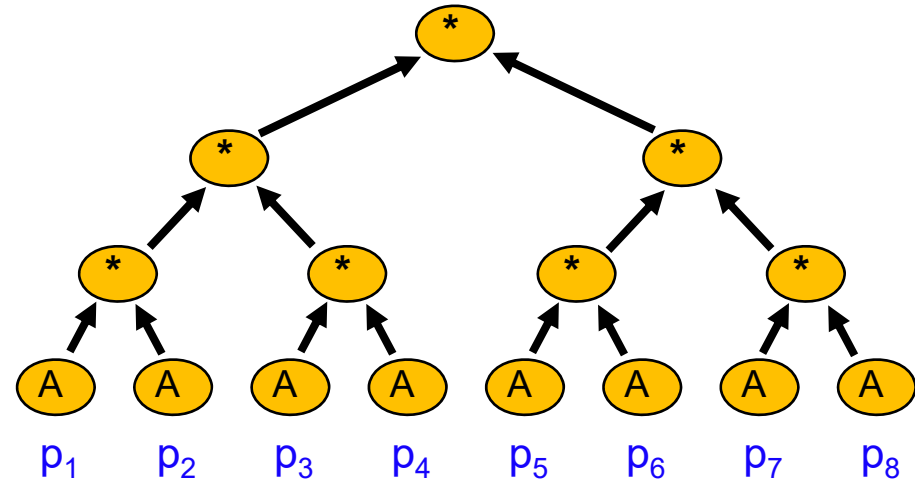


Fibonacci numbers in parallel

- A simple identity.

$$\begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} F_{n-1} \\ F_{n-2} \end{bmatrix} = \begin{bmatrix} F_n \\ F_{n-1} \end{bmatrix}$$

- Let $A = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}$



- By repeatedly applying the identity, we get

$$A^n \cdot \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} F_{n+2} \\ F_{n+1} \end{bmatrix}$$

- So if we can quickly compute A^n in parallel, we can compute F_{n+2} .
 - Can compute F_n in $O(\log n)$ time with n processors.