



# Parallel Iterative Matrix Algorithms

CS121 Parallel Computing  
Spring 2021



# Iterative matrix algorithms

- To solve a dense system of linear equations  $Ax = b$ , we saw direct methods such as Gaussian Elimination.
- When the  $A$  is very large (millions of variables), GE is too slow.
- For structured matrices, e.g. banded matrices, special solution methods can be developed.
- For general sparse matrices, we use iterative algorithms that compute approximate solutions which eventually converge to the true solution.



# Iterative matrix algorithms

- Given  $A \in \mathbb{R}^{n \times n}$ , write  $A = M - N$ , where  $M$  and  $N$  are matrices such that  $M^{-1}$  is easy to compute (e.g. a diagonal matrix).
  - Let  $x^*$  be the solution to  $Ax = b$ . Then  $Mx^* = Nx^* + b$ .
- Let  $C = M^{-1}N$ ,  $d = M^{-1}b$ .
  - So  $x^* = Cx^* + d$ .
- Starting from an initial  $x$ , repeatedly compute  $Cx + d$ .
  - Denote  $k$ 'th iterate of  $x$  as  $x^{(k)}$ .
  - Then  $x^{(k+1)} = Cx^{(k)} + d$ .



# Convergence criteria

- We want the iterations to converge, starting from any initial vector  $x^{(0)} \in \mathbb{R}^n$ .
  - I.e. we want  $\lim_{k \rightarrow \infty} x^{(k)} = x^*$ , so that  $x^* = Cx^* + d$ .
- Since  $x^{(k+1)} = Cx^{(k)} + d$ , then subtracting, we get  $x^{(k+1)} - x^* = C(x^{(k)} - x^*)$ .
  - Also,  $x^{(k)} - x^* = C(x^{(k-1)} - x^*)$ , etc.
  - So in general  $x^{(k)} - x^* = C^k(x^{(0)} - x^*)$ .
- Let  $\rho(C)$  be the magnitude of the largest eigenvalue of  $C$ .
- **Thm** The following are equivalent
  - The iterative algorithm converges for any initial  $x^{(0)}$ .
  - $\lim_{k \rightarrow \infty} C^k = 0$ .
  - $\rho(C) < 1$ .



# Jacobi method

- Write  $A = D - L - R$ , where  $D$  is the diagonal elements of  $A$ ,  $-L$  is the lower triangular part of  $A$  without  $D$ , and  $-R$  is the upper triangular part without  $D$ .
- Let  $M = D, N = L + R$ . Note that  $M$  is easy to invert.
- Then  $C = D^{-1}(L + R)$ , and  $d = D^{-1}b$ .
  - So  $c_{ij} = -a_{ij}/a_{ii}$  if  $j \neq i$ , and  $c_{ii} = 0$  for all  $i$ .
  - Also,  $d_i = b_i/a_{ii}$  for all  $i$ .
- Convergence is guaranteed if the matrix is diagonally dominant, i.e.  $|a_{ii}| > \sum_{j=1, j \neq i}^n |a_{ij}|$  for all  $i$ .
- Recall  $x^{(k+1)} = Cx^{(k)} + d$ . So the  $i$ 'th component of  $x^{(k+1)}$  is
$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1, j \neq i}^n a_{ij} x_j^{(k)} \right).$$
- $x^{(k+1)}$  depends only on  $x^{(k)}$ , and different components of  $x^{(k+1)}$  do not have any dependencies.
  - Thus, all components of  $x^{(k+1)}$  can be computed in parallel.

# Jacobi method example

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1, j \neq i}^n a_{ij} x_j^{(k)} \right)$$

## Jacobi Method

Consider 4x4 case

$$\begin{bmatrix} 10 & -1 & 2 & 0 \\ -1 & 11 & -1 & 3 \\ 2 & -1 & 10 & -1 \\ 0 & 3 & -1 & 8 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} 6 \\ 25 \\ -11 \\ 15 \end{bmatrix}$$

**Example**

$$\begin{aligned} 10x_1 - x_2 + 2x_3 &= 6 \\ -x_1 + 11x_2 - x_3 + 3x_4 &= 25 \\ 2x_1 - x_2 + 10x_3 - x_4 &= -11 \\ 3x_2 - x_3 - 8x_4 &= 15 \end{aligned}$$

$$\begin{aligned} x_1 &= (x_2 - 2x_3 + 6)/10 \\ x_2 &= (x_1 + x_3 - 3x_4 + 25)/11 \\ x_3 &= (-2x_1 + x_2 + x_4 - 11)/10 \\ x_4 &= (-3x_2 + x_3 + 15)/(-8) \end{aligned}$$

$$\text{given } x^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix} \Rightarrow \begin{aligned} x_1^{(1)} &= (x_2^{(0)} - 2x_3^{(0)} + 6)/10 \\ x_2^{(1)} &= (x_1^{(0)} + x_3^{(0)} - 3x_4^{(0)} + 25)/11 \\ x_3^{(1)} &= (-2x_1^{(0)} + x_2^{(0)} + x_4^{(0)} - 11)/10 \\ x_4^{(1)} &= (-3x_2^{(0)} + x_3^{(0)} + 15)/(-8) \end{aligned} \Rightarrow x^{(1)} = \begin{bmatrix} 0.6000 \\ 2.2727 \\ -1.1000 \\ 1.8750 \end{bmatrix}$$

# Parallel Jacobi method

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1, j \neq i}^n a_{ij} x_j^{(k)} \right)$$

- Since all components of  $x^{(k+1)}$  are independent, we can use up to n processors.
- In distributed memory, matrix A and vector b are stored in row-wise block format across the processors.
- $x^{(k+1)}$  and  $x^{(k)}$  are computed in x\_new and x\_old, resp.
- Each processor needs all the values of x.
- After all processors compute their part of x\_new, the whole vector is distributed to all processors using MPI\_Allgather.

Source: Parallel Programming for Multicore and Cluster Systems, Rauber and Runger

```
int Parallel_jacobi(int n, int p, int max_it, float tol)
{
    int i_local, i_global, j, i;
    int n_local, it_num;
    float x_temp1[GLOB_MAX], x_temp2[GLOB_MAX], local_x[GLOB_MAX];
    float *x_old, *x_new, *temp;

    n_local = n/p; /* local blocksize */
    MPI_Allgather(local_b, n_local, MPI_FLOAT, x_temp1, n_local,
                  MPI_FLOAT, MPI_COMM_WORLD);
    x_new = x_temp1;
    x_old = x_temp2;
    it_num = 0;
    do {
        it_num++;
        temp = x_new; x_new = x_old; x_old = temp;
        for (i_local = 0; i_local < n_local; i_local++) {
            i_global = i_local + me * n_local;
            local_x[i_local] = local_b[i_local];
            for (j = 0; j < i_global; j++)
                local_x[i_local] = local_x[i_local] -
                                   local_A[i_local][j] * x_old[j];
            for (j = i_global+1; j < n; j++)
                local_x[i_local] = local_x[i_local] -
                                   local_A[i_local][j] * x_old[j];
            local_x[i_local] = local_x[i_local] / local_A[i_local][i_global];
        }
        MPI_Allgather(local_x, n_local, MPI_FLOAT, x_new, n_local,
                      MPI_FLOAT, MPI_COMM_WORLD);
    } while ((it_num < max_it) && (distance(x_old, x_new, n) >= tol));
    output(x_new, global_x);
    if (distance(x_old, x_new, n) < tol) return 1;
    else return 0;
}
```



# Gauss-Seidel method

- The Gauss-Seidel method usually converges faster than the Jacobi method.
- Again write  $A = D - L - R$ , but set  $M = D - L$  and  $N = R$ .
  - Thus, we have  $C = (D - L)^{-1}R$ .
  - Since  $D - L$  is lower triangular, it can be inverted by forward substitution.
- Convergence is again guaranteed if the matrix is diagonally dominant, i.e.  $|a_{ii}| > \sum_{j=1, j \neq i}^n |a_{ij}|$  for all  $i$ .
- As before,  $x^{(k+1)} = Cx^{(k)} + d$ . So in components form we have
$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{(k+1)} - \sum_{j=i+1}^n a_{ij} x_j^{(k)} \right).$$
- Unlike Jacobi method,  $x_i^{(k+1)}$  depends on all  $x_j^{(k+1)}$  for  $j < i$  and  $a_{ij} \neq 0$ .
  - Thus, unless many values of  $a_{ij} = 0$ , different  $x_i^{(k+1)}$  cannot be computed in parallel.
- While Gauss-Seidel converges faster than Jacobi, it has less parallelism, and may not run faster.





# SOR method

- Successive over-relaxation modifies the Gauss-Seidel method to obtain faster convergence.
- $x$  is updated as a linear combination of Gauss-Seidel update and its previous value.

- In components form,

$$x_i^{(k+1)} = \frac{\omega}{a_{ii}} \left( b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{(k+1)} - \sum_{j=i+1}^n a_{ij} x_j^{(k)} \right) + (1 - \omega) x_i^{(k)}.$$

- Convergence depends on properties of  $A$  and  $\omega$ .  
E.g. if  $A$  is symmetric and positive definite and  $\omega \in (0,2)$  then SOR converges.

# Parallel Gauss-Seidel method




- Since in general  $x_i^{(k+1)}$  depends on  $x_j^{(k+1)}$  for all  $j < i$ , we compute the  $x_i^{(k+1)}$  sequentially, for  $i=0,1,2,\dots$
- Each  $x_i^{(k+1)}$  is a dot product of  $(x_1^{(k+1)}, \dots, x_{i-1}^{(k+1)}, 0, x_{i+1}^{(k)}, x_n^{(k)})$  with the  $i$ 'th row of A.
  - This dot product can be split into multiple parts and computed in parallel.
- Use block column-wise decomposition of A and x across the processors.
- Each processor computes part of  $x_i^{(k+1)}$ .
  - The parts are then summed and distributed to all the processors using MPI\_Allreduce.
- Repeat the do loop until x converges.
- Each processor only does  $n/p$  computations for each reduce communication step, so speedup is limited unless  $n \gg p$ .

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{(k+1)} - \sum_{j=i+1}^n a_{ij} x_j^{(k)} \right)$$

```
n_local = n/p;
do {
    delta_x = 0.0;
    for (i = 0; i < n; i++) {
        s_k = 0.0;
        for (j = 0; j < n_local; j++)
            if (j + me * n_local != i)
                s_k = s_k + local_A[i][j] * x[j];
        root = i/n_local;
        i_local = i % n_local;
        MPI_Reduce(&s_k, &x[i_local], 1, MPI_FLOAT, MPI_SUM, root,
                  MPI_COMM_WORLD);
        if (me == root) {
            x_new = (b[i_local] - x[i_local]) / local_A[i][i_local];
            delta_x = max(delta_x, abs(x[i_local] - x_new));
            x[i_local] = x_new;
        }
    }
    MPI_Allreduce(&delta_x, &global_delta, 1, MPI_FLOAT,
                  MPI_MAX, MPI_COMM_WORLD);
} while(global_delta > tol);
```


# Poisson's equation

- Poisson's equation is a partial differential equation (PDE) to describe the potential field caused by a mass or electrostatic density distribution.
  - We'll look at Poisson's equation in 2D space.
- Given a function  $f(x, y)$ , we want to find a function  $\phi(x, y)$  with  $-\Delta\phi = f$ .
  - Here  $\Delta\phi = \nabla^2\phi = \frac{\partial^2\phi}{\partial^2x} + \frac{\partial^2\phi}{\partial^2y}$ .
- Poisson's equation can be solved numerically by discretizing 2D space.
  - For simplicity, we divide  $[0,1] \times [0,1]$  evenly into  $N+1$  points along each axis.
  - Let  $h = 1/(N + 1)$ , and let  $u_{ij} = \phi(x_i, y_j)$ , and  $f_{ij} = f(x_i, y_j)$ .
- Then  $\frac{1}{h^2}(4u_{ij} - u_{i+1,j} - u_{i-1,j} - u_{i,j+1} - u_{i,j-1}) = f_{ij}$ , for  $0 \leq i, j \leq N + 1$ .
- For simplicity, fix the value of  $\phi$  on the boundary of the square, and divide out both sides by  $1/h^2$ .
- Look for the value of  $\phi$  in the square's interior.
  - This leads to a set of  $N^2$  linear equations, one for each  $1 \leq i, j \leq N$ .



# Matrix form of Poisson's equation

- Let  $x_k = u_{ij}$  for  $k = i + (j - 1)N$ ,  $1 \leq i, j \leq N$ .
- Each equation from the discretization has the form  $4x_k - x_{k+1} - x_{k-1} - x_{k+N} - x_{k-N} = b_k$ , for some  $1 \leq k \leq N^2$  and  $b_k$ .
- Let  $n = N^2$ , and create an  $n \times n$  matrix  $A$  for the nonzero coefficients of all the equations.
- $A$  has the following nonzero structure.
  - There are three bands of nonzeros, on the diagonal, and above and below the diagonal.
  - There are two additional bands of nonzeros distance  $N = \sqrt{n}$  above and below the diagonal.



# Gauss-Seidel for Poisson's equation



- Recall the Gauss-Seidel iteration is

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{(k+1)} - \sum_{j=i+1}^n a_{ij} x_j^{(k)} \right).$$

- Applied to Poisson's equation, we have

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left( b_i - a_{i,i-\sqrt{n}} \cdot x_{i-\sqrt{n}}^{(k+1)} - a_{i,i-1} \cdot x_{i-1}^{(k+1)} - a_{i,i+1} \cdot x_{i+1}^{(k)} - a_{i,i+\sqrt{n}} \cdot x_{i+\sqrt{n}}^{(k)} \right), \text{ for } i = 1, \dots, n.$$

- The values  $x_{i-\sqrt{n}}^{(k+1)}$  and  $x_{i-1}^{(k+1)}$  need to be computed before  $x_i^{(k+1)}$ .
- Place the x values on the grid in row major order.
- Each x value depends on value directly above and to its left.
  - Ex Point 9 depends on 3 and 8.
- Notice the x values along each diagonal are all independent.
- There are  $2\sqrt{n} - 1$  diagonals.
  - Each diagonal has  $O(\sqrt{n})$  points, giving a large amount of parallelism.
- The first  $\sqrt{n}$  diagonals  $l = 1, \dots, \sqrt{n}$  each contain  $l$  points, with indices  $i = l + j(\sqrt{n} - 1)$ , for  $0 \leq j \leq l$ .
- The last  $\sqrt{n} - 1$  diagonals  $l = 2, \dots, \sqrt{n}$  contain  $\sqrt{n} - l + 1$  points, with indices  $i = l\sqrt{n} + j(\sqrt{n} - 1)$ , for  $0 \leq j \leq \sqrt{n} - l$ .



# Gauss-Seidel for Poisson's equation

- We parallelize Gauss-Seidel for Poisson's equation by iterating through the diagonals sequentially, and computing all the values in each diagonal in parallel.
- Given  $p$  processors, each processor computes every  $p$ 'th value  $x[i]$  on the  $l$ 'th diagonal.
- Notice the sizes of the diagonals first increases, then decreases.
  - The two for loops compute  $x$  during the increasing and decreasing phase, resp.
- The function `collect_elements` sends the  $x$  values from the  $l$ 'th diagonal to neighboring processors to compute the  $l+1$ 'st diagonal.
- Repeat the `do` loop until  $x$  converges.


$$x_i^{(k+1)} = \frac{1}{a_{ii}} (b_i - a_{i,i-\sqrt{n}} \cdot x_{i-\sqrt{n}}^{(k+1)} - a_{i,i-1} \cdot x_{i-1}^{(k+1)} - a_{i,i+1} \cdot x_{i+1}^{(k)} - a_{i,i+\sqrt{n}} \cdot x_{i+\sqrt{n}}^{(k)})$$

```
sqn = sqrt(n);
do {
    for (l = 1; l <= sqn; l++) {
        for (j = me; j < l; j+=p) {
            i = l + j * (sqn-1) - 1; /* start numbering with 0 */
            x[i] = 0;
            if (i-sqn >= 0) x[i] = x[i] - a[i][i-sqn] * x[i-sqn];
            if (i > 0) x[i] = x[i] - a[i][i-1] * x[i-1];
            if (i+1 < n) x[i] = x[i] - a[i][i+1] * x[i+1];
            if (i+sqn < n) x[i] = x[i] - a[i][i+sqn] * x[i+sqn];
            x[i] = (x[i] + b[i]) / a[i][i];
        }
        collect_elements(x,l);
    }
    for (l = 2; l <= sqn; l++) {
        for (j = me - l + 1; j <= sqn - l; j+=p) {
            if (j >= 0) {
                i = l * sqn + j * (sqn-1) - 1;
                x[i] = 0;
                if (i-sqn >= 0) x[i] = x[i] - a[i][i-sqn] * x[i-sqn];
                if (i > 0) x[i] = x[i] - a[i][i-1] * x[i-1];
                if (i+1 < n) x[i] = x[i] - a[i][i+1] * x[i+1];
                if (i+sqn < n) x[i] = x[i] - a[i][i+sqn] * x[i+sqn];
                x[i] = (x[i] + b[i]) / a[i][i];
            }
        }
        collect_elements(x,l);
    }
} while(convergence_test() < tol);
```

# Red-black ordering


- In Gauss-Seidel, each  $x_i^{(k+1)}$  depends on all  $x_j^{(k+1)}$  for  $j < i$  and  $a_{ij} \neq 0$ .
- The method is valid for any ordering of the  $x$  values, so we can choose an ordering that gives best parallelism.
- The diagonals method used the top ordering. We now consider the bottom red-black ordering.
- Assign each mesh point a color, red or black.
  - For each mesh point  $(i,j)$ , if  $i+j$  is even, it is colored red (grey in the picture on right). Otherwise color it black.
- Since a point only depends on the points above it and to its left, none of the red points depend on each other, and similarly for the black points.
  - So the red and black  $x$  values can be computed in parallel.

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24



# Red-black ordering

- Matrix A has a different structure after reordering.
- The red points for  $x^{(k+1)}$  only depend on the black points for  $x^{(k)}$ , and the black points for  $x^{(k+1)}$  only depend on the red points for  $x^{(k+1)}$ .
  - Thus, we can compute all  $n/2$  red points in  $x^{(k+1)}$  in parallel, then compute all the  $n/2$  black points in  $x^{(k+1)}$  in parallel.
- Write matrix  $A = \begin{pmatrix} D_R & F \\ E & D_B \end{pmatrix}$ , where  $D_R$  and  $D_B$  are diagonal matrices corresponding to the red and black points resp, and E and F are banded matrices.
- Also write  $x = (x_R \ x_B)$ , where  $x_R, x_B$  are the set of red and black x values, resp.





# Red-black ordering

- Let  $D = \begin{pmatrix} D_R & 0 \\ 0 & D_B \end{pmatrix}$ ,  $L = \begin{pmatrix} 0 & 0 \\ -E & 0 \end{pmatrix}$  and  $U = \begin{pmatrix} 0 & -F \\ 0 & 0 \end{pmatrix}$ .
- Write the Gauss-Seidel iteration as  $\begin{pmatrix} D_R & 0 \\ E & D_B \end{pmatrix} \begin{pmatrix} x_R^{(k+1)} \\ x_B^{(k+1)} \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} - \begin{pmatrix} 0 & F \\ 0 & 0 \end{pmatrix} \begin{pmatrix} x_R^{(k)} \\ x_B^{(k)} \end{pmatrix}$ .
- So  $D_R \cdot x_R^{(k+1)} = b_1 - F \cdot x_B^{(k)}$  for  $k = 1, 2, \dots$ , and  $D_B \cdot x_B^{(k+1)} = b_2 - E \cdot x_R^{(k+1)}$ , for  $k = 1, 2, \dots$
- Let  $n_R$  and  $n_B$  be the number of red and black points, resp.
- For each point  $i$ , let  $N(i)$  be its neighbors in the grid.
- In component form, we have

$$\left(x_R^{(k+1)}\right)_i = \frac{1}{a_{ii}} \left( b_i - \sum_{j \in N(i)} a_{ij} \cdot \left(x_B^{(k)}\right)_j \right) \text{ for } i = 1, \dots, n_R$$

$$\left(x_B^{(k+1)}\right)_i = \frac{1}{a_{i+n_R, i+n_R}} \left( b_{i+n_R} - \sum_{j \in N(i)} a_{i+n_R, j} \cdot \left(x_R^{(k+1)}\right)_j \right) \text{ for } i = 1, \dots, n_B$$

- Thus, we first compute  $x^{(k+1)}$  for all the red points, then  $x^{(k+1)}$  for all the black points.



# Parallel red-black algorithm

- Use a block row-wise decomposition of  $A$  and  $x$  across the processors.
- Use barrier synchronization between the two loops to compute black values after red ones.
- `collect_elements` sends newly computed values of  $x$  that lie on the boundary between two processors to the other processor.

```
local_nr = nr/p; local_nb = nb/p;
do {
    mestarttr = me * local_nr;
    for (i= mestarttr; i < mestarttr + local_nr; i++) {
        xr[i] = 0;
        for (j ∈ N(i))
            xr[i] = xr[i] - a[i][j] * xb[j];
        xr[i] = (xr[i]+b[i]) / a[i][i] ;
    }
    collect_elements(xr);
    mestarttb = me * local_nb + nr;
    for (i= mestarttb; i < mestarttb + local_nb; i++) {
        xb[i] = 0;
        for (j ∈ N(i))
            xb[i] = xb[i] - a[i+nr][j] * xr[j];
        xb[i] = (xb[i] + b[i+nr]) / a[i+nr][i+nr];
    }
    collect_elements(xb);
} while (convergence_test());
```