

## L9e Virtual Machine Intro

### ① Benefits of Virtualization

- Multiple Secure Environment

A system VM provides a sandbox that isolates one system environment from other environments.

- Failure Isolation

Virtualization helps isolate the effects of a failure to the VM where the failure occurred.

- Better System Utilization

A virtualized system can be (dynamically or statically) reconfigured for changing needs.

- Mixed-OS Environment

A single hardware platform can support multiple operating systems concurrently.

### ② Virtualization

A virtualized system (or subsystem) is a mapping of its interface, and all resources visible through that interface, to the interface and resources of a real system. Virtualization involves the construction of an isomorphism that maps a virtual guest system to a real host system.

### ③ Virtualization Properties

- Isolation: Fault Isolation, Software Isolation, Performance Isolation.
- Encapsulation: All VM state can be captured into a file, Complexity is proportional to virtual HW model and independent of guest.
- Interposition: All guest actions go through the virtualizing software which can inspect, modify, and deny operations.

### ④ Virtual Machine Monitor

A VM is implemented by adding a layer of software to a real machine so as to support the desired VM's architecture. This layer of software is often referred to as virtual machine monitor (VMM). VMMs are often implemented as a co-designed firmware-software layer, referred to as the hypervisor.

### ⑤ Types of Virtual Machines

- Process VM: Capable of supporting an individual process.
- System VM: Provides a complete system environment. Supports an OS with potentially many types of processes.

### ⑥ Partitioning

Using partitioning, multiple applications can simultaneously exploit the available resources of the system. in-space: physical partitioning. in-time: logical partitioning.

### ⑦ Virtualizing Devices

The technique that is used to virtualize an I/O device depends on whether the device is shared and, if so, the ways in which it can be shared. The common categories of devices are: Dedicated devices, Partitioned devices, Shared devices, Spooled devices.

## L10 Memory Management

### ① Address Space


Definition: Set of accessible addresses and the state associated with them.

### ② Important Aspects of Memory Multiplexing

Protection, Controlled overlap, Translation.

### ③ Base and Bound

Translation happens at execution. Issues: Fragmentation problem over time, Missing support for sparse address space, Hard to do inter-process sharing.



Multiple separate segments: each segment is given region of contiguous memory.

Fragmentation: wasted space

- External: free gaps between allocated chunks
- Internal: don't need all memory within allocated chunks

### ④ Page Sharing


- The "kernel region" of every process has the same page table entries.

- Different processes running same binary
- User-level system libraries (execute only)
- Shared-memory segments between different processes

### ⑤ Security Measures

- Address Space Randomization
- Kernel address space isolation
- Context switch
- Needs to be switched: Page table pointer and limit.
- Protection: Translation (per process) and dual-mode.

### ⑥ Multi-Level Translation



Pros: Only need to allocate as many page table entries as we need for application, Easy memory allocation, Easy Sharing.  
Cons: One pointer per page, Page tables need to be contiguous, Two lookups per reference.

## L10e I/O Intro

### ① Linux: | pipe

### ② POSIX I/O

POSIX: Portable Operating System Interface  
Design Patterns/Concepts: Uniformity - everything is a file, Open before use, Byte-oriented, Kernel buffered reads, Kernel buffered writes, Explicit close.

### ③ File System Abstraction

Files live in hierarchical namespace of filenames.

### ④ C High-Level File API – Streams.

Three predefined streams: stdin, stdout, stderr.

### ⑤ C Low Level: Standard Descriptors

int fsync (int filedes) – wait for i/o to finish

void sync (void) – wait for ALL to finish

int pipe(int fileds[2]);

// Writes to fileds[1] read from fileds[0]

int dup2(int old, int new);

int dup(int old);

Operations: open, read, write, create, close

Operations specific to terminals, devices, networking; Duplicating descriptors; Pipes: bi-directional channel; File Locking; Memory-mapping files; Asynchronous I/O.

### ⑥ Streams vs. File Descriptors

Stream are buffered in user memory (sleep won't work). Operations on file descriptors are visible immediately (sleep will work).

Why Buffer in User-space? Avoid system call overhead; Functionality, Simplifies operating system.

### ⑦ Device Drivers

Device-specific code in the kernel that interacts directly with the device hardware.

Typically divided into two pieces: Top half: accessed in call path from system calls; Bottom half: run as interrupt routine.

### ⑧ File abstraction works for inter-processes communication (local or Internet)

## L11 Caching, TLB

### ① Page Table Discussion

Pros: Simple memory allocation, Easy to share

Cons: Too big table.

### ② Page Table Entry (PTE)

Pointer to next-level page table or to actual page + Permission bits: valid, read-only, read-write, write only.

How to use:

- Demand Paging: Keep only active pages in memory, Place others on disk and mark their PTEs invalid.
- Copy on Write: UNIX fork gives copy of parent address space to child, Address spaces disconnected after child created. To do this cheaply, make copy of parent's page tables (point at same memory); mark entries in both

sets of page tables as read-only; page fault on write creates two copies.

- Zero Fill On Demand: New data pages must carry no information (say be zeroed). Mark PTEs as invalid; page fault on use gets zeroed page.

### ③ How are segments used

One set of global segments (GDT) for everyone, different set of local segments (LDT) for every process.

### ④ Inverted Page Table (P29)

Forward Page Tables: Size of page table is at least as large as amount of virtual memory allocated to processes.

Inverted Page Table: Size directly related to amount of physical memory. Cons: Complexity of managing hash chains: Often in hardware! Poor cache locality of page table.

### ⑤ Comparison

- Simple Segmentation. Pros: Fast context switching; Segment mapping maintained by CPU. Cons: External fragmentation;

- Paging (single-level page). Pros: No external fragmentation, fast easy allocation. Cons: Large table size ~ virtual memory, Internal fragmentation.


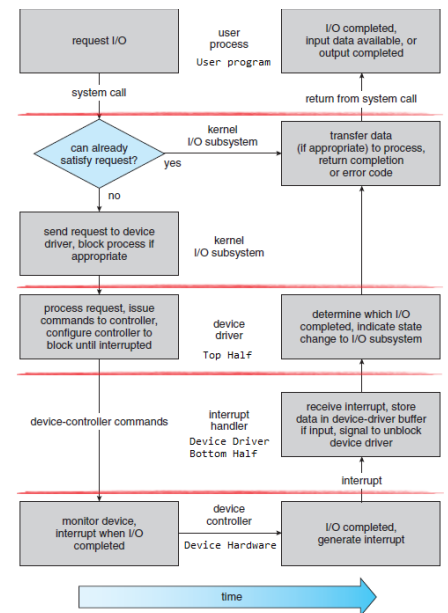
- Paged segmentation & Two-level pages. Pros: Table size ~ # of pages in virtual memory, fast easy allocation. Cons: Multiple memory references per page access.

- Inverted Table. Pros: Table size ~ # of pages in physical memory. Cons: Hash function more complex, No cache locality of page table.

### ⑥ Memory Management Unit (MMU)

The processor requests READ Virtual-Address to memory system, through the MMU to the cache (to the memory).

On every reference (I-fetch, Load, Store) MMU read (multiple levels of) page table entries to get physical frame or FAULT.



SIZE	PREFIX	SYMBOL	SIZE	PREFIX	SYMBOL
10 <sup>3</sup>	Kilo-	K	2 <sup>10</sup>	Kibi-	Ki
10 <sup>6</sup>	Mega-	M	2 <sup>20</sup>	Mebi-	Mi
10 <sup>9</sup>	Giga-	G	2 <sup>30</sup>	Gibi-	Gi
10 <sup>12</sup>	Tera-	T	2 <sup>40</sup>	Tebi-	Ti
10 <sup>15</sup>	Peta-	P	2 <sup>50</sup>	Pebi-	Pi
10 <sup>18</sup>	Exa-	E	2 <sup>60</sup>	Exbi-	Ei
10 <sup>21</sup>	Zetta-	Z	2 <sup>70</sup>	Zebi-	Zi
10 <sup>24</sup>	Yotta-	Y	2 <sup>80</sup>	Yobi-	Yi
10 <sup>27</sup>	milli-	m	10 <sup>-9</sup>	femto-	f
10 <sup>30</sup>	micro-	μ	10 <sup>-18</sup>	atto-	a
10 <sup>33</sup>	nano-	n	10 <sup>-27</sup>	zepto-	z
10 <sup>36</sup>	pico-	p	10 <sup>-36</sup>	yocto-	y

## ⑦ Cache (P61)

A repository for copies that can be accessed more quickly than the original.  
Average Memory Access Time (AMAT)  
 $\text{AMAT} = (\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$   
Temporal Locality, Spatial Locality.  
Misses: Compulsory, Capacity, Conflict, Coherence.  
Organizations: Direct Mapped: single block per set; Set associative: more than one block per set; Fully associative: all entries equivalent.  
Cache line: Tag + Index + Block (offset)  
Write through: The information is written to both the line in the cache and to the block in the lower-level memory; Write back: The information is written only to the block in the Cache.  
⑧ Translation Look-Aside Buffer (TLB)  
Placed before cache.  
Thrashing: continuous conflicts between accesses.  
Conflict misses expensive: Fully Associative.  
When Context Switch: Invalidate TLB, or Include ProcessID in TLB.  
TLB Consistency: If translation tables change (evict and swap), must invalidate TLB entry.  
⑨ Page Fault  
Not an interrupt, synchronous to instruction execution (Trap).

## L12 Demand Paging, General I/O

① Uses of Demand Paging  
Extend the stack, Extend the heap, Process Fork, Exec, MMAP to explicitly share region.  
② Cache size needs to fit working set size.  
③ Effective Access Time  
 $\text{EAT} = \text{Hit Rate} \times \text{Hit Time} + \text{Miss Rate} \times \text{Miss Time}$   
 $\text{Time} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$   
④ Misses in Page Cache  
Compulsory Misses, Capacity Misses, Conflict Misses, Policy Misses (bad eviction policy).  
⑤ Thrashing  
A process is busy swapping pages in and out with little or no actual progress (low CPU utilization).  
⑥ I/O Subsystem  
Devices have many different speeds.  
Goal: Provide Uniform Interfaces, Despite Wide Range of Different Devices.  
Devices: Block Devices, Character Devices, Network Devices.  
Timing: Blocking Interface: "Wait"; Non-blocking Interface: "Don't Wait"; Asynchronous Interface: "Tell Me Later".  
⑦ Transferring Data to/from Controller  
- Programmed I/O: Each byte transferred via processor in/out or load/store. Pro: Simple hardware, easy to program, Con: Consumes processor cycles proportional to data size.  
- Direct Memory Access: Give controller access to memory bus. Ask it to transfer data blocks to/from memory directly.  
⑧ I/O Device Notifying the OS  
- I/O Interrupt: Device generates an interrupt whenever it needs service. Pro: handles unpredictable events well. Con: interrupts relatively high overhead.  
- Polling: OS periodically checks a device-specific status register. Pro: low overhead. Con: may waste many cycles on polling if infrequent or unpredictable I/O operations.  
- Actual devices combine both polling and interrupts

## L13 File Systems

① Logical Directory Organization  
Goals: Efficiency - locating a file quickly; Naming - convenient to users; Grouping  
② Hard links, soft links (P19)  
Links (hard links) make directory a DAG, not just a tree. Softlinks (aliases) are another name for an entry.  
③ File: permanent storage  
Contains: Data (Blocks on disk), Metadata (Attributes - Owner, size, last opened, ... Access rights (R, W, X, Owner, Group, Other, ...)).

## ④ File Allocation Table (FAT) (P70)

In FAT, directory is a file containing <file\_name: file\_number> mappings. File attributes are kept in directory. Each directory is a linked list of entries. Root directory is at block 2 (no 0 or 1). Security Holes: FAT has no access rights; FAT has no header in the file blocks (all processes have access of FAT table); Just gives an index into the FAT to read data (Could start in middle of file or access deleted data).  
⑤ Inode Structure (P73)  
File defined by header, called "inode".  
- Problem 1: When create a file, don't know how big it will become. Solution: Fast File System (FFS) Allocation and placement policies for BSD 4.2  
- Problem 2: Missing blocks due to rotational delay. Solution 1: Skip sector positioning ("interleaving"); Solution 2: Read ahead: read next block right after first, even if application hasn't asked for it yet.  
⑥ Directories  
- link / unlink (rm): Link existing file to a directory (Not in FAT); Forms a DAG.  
When can file be deleted?  
- Maintain ref-count of links to the file  
- Delete after the last reference is gone  
⑦ Links  
- Hard link  
Sets another directory entry to contain the file number for the file; Creates another name (path) for the file; Each is "first class".  
- Soft link or Symbolic Link or Shortcut  
Directory entry contains the path and name of the file; Map one name to another name.  
⑧ New Technology File System (NTFS) (P82)  
Variable extents not fixed blocks, tiny files data is in header.  
Master File Table (MFT)  
⑨ File System Caching  
Buffer Cache: Memory used to cache kernel resources, including disk blocks and name translations.  
Cache Size: adjust boundary dynamically so that the disk access rates for paging and file access are balanced.  
Read Ahead Prefetching: fetch sequential blocks early.  
Delayed Writes: Writes to files not immediately sent out to disk. Flushed to disk periodically.  
- Advantages: 1. Disk scheduler can efficiently order lots of requests; 2. Disk allocation algorithm can be run with correct size value for a file; 3. Some files need never get written to disk! (e.g. temporary scratch files written /tmp often don't exist for 30 sec)  
- Disadvantages: 1. What if system crashes before file has been written out? 2. Worse yet, what if system crashes before a directory file has been written out? (lose pointer to inode!)  
⑩ Important "ilities"  
Availability: the probability that the system can accept and process requests.  
Durability: the ability of a system to recover data despite faults. Doesn't necessarily imply availability.  
Reliability: the ability of a system or component to perform its required functions under stated conditions for a specified period of time.

## L14 Advanced File Systems

① Transaction  
- An atomic sequence of actions (reads/writes) on a storage system (or database) that takes it from one consistent state to another.  
- Typical Structure: Begin a transaction (get transaction id); Do a bunch of updates. If any fail along the way, or any conflicts with other transactions, roll-back; Commit the transaction.  
- ACID properties: Atomicity: all actions in the transaction happen, or none happen;  
Consistency: transactions maintain data integrity (e.g. Balance cannot be negative);  
Isolation: execution of one transaction is isolated from that of all others; no problems from concurrency;  
Durability: if a transaction

commits, its effects persist despite crashes.

② Difference between "Log Structured" and "Journaled": In a Log Structured filesystem, data stays in log form; In a Journaled filesystem, Log used for recovery.  
Crash During Logging: No commit, discard log entries.  
Recovery After Commit: Redo it as usual.  
Expensive Journaling: Record modifications to file system data structures, but apply updates to a file's contents directly.  
③ Flash Translation Layer (FTL)  
Random reads are as fast as sequential reads; Random writes are bad for flash storage.  
Node Address Table (NAT): Independent of FTL.  
Updates to data sorted by predicted write frequency to optimize FLASH management.  
Checkpoint (CP): Keeps the file system status.  
Segment Information Table (SIT): garbage collection. Per segment information.  
Segment Summary Area (SSA): Summary representing the owner information of all blocks in the Main area.  
④ Distributed File Systems (DFS)  
Virtual File System (VFS): Virtual abstraction similar to local file system.  
Four primary object types: superblock object; inode object; dentry object; file object. (no specific directory object, treated as files.)  
⑤ Problems and solutions  
Simple Distributed File System - Problem: Performance. Caching to reduce network load  
- Problem: Failure, Cache consistency (What if client removes a file but server crashes before acknowledgement?).  
Stateless Protocol: A protocol in which all information required to service a request is included with the request.  
Idempotent Operations: repeating an operation multiple times is same as executing it just once. Client: timeout expires without reply, just run the operation again.  
Remote Procedure Call Protocol (RPC Protocol).  
Call procedure on remote machine or in remote domain.  
⑥ Network File System (NFS)  
Defines an RPC protocol for clients to interact with a file server, Keeps most operations idempotent, Don't buffer writes on server side cache.  
Three Layers: UNIX file-system interface, VFS layer, NFS service layer.  
Write-through caching. Weak consistency: Client polls server periodically to check for changes. Multiple write: arbitrary result.  
⑦ Key-value storage systems (K-V Store)  
Distributed Hash Tables (DHT): partition set of key-values across many machines.  
Challenges: Scalability, Fault Tolerance, Consistency, Heterogeneity (inconsistency in latency and bandwidth).  
⑧ Iterative vs. Recursive  
> Recursive Directory Architecture: Have a node maintain the mapping between keys and the machines (nodes) that store the values associated with the keys.  
+Faster, as directory server is typically close to storage nodes; +Easier for consistency: directory can enforce an order for all puts and gets; -Directory is a performance bottleneck.  
> Iterative Directory Architecture: Return node to requester and let requester contact node.  
+More scalable, clients do more work;  
-Harder to enforce consistency

## HW2

① Advantages of inode to FAT: Fast random access to files. Support for hard links.  
② Write-behind policy:  
Advantage 1: The disk scheduling algorithm (i.e. SCAN) has more dirty blocks to work with at any one time and can thus do a better job of scheduling the disk arm.  
Advantage 2: Temporary files may be written and deleted before data is written to disk.  
Disadvantage: File data may be lost if the computer crashes before data is written to disk.