

CS150 Discussion IV

Jiachun Jin
jinjch@shanghaitech.edu.cn

1. B^+ Tree

A quick review
Cost model for search
Examine the codes
find
insert

2. Buffer Management

Exercises

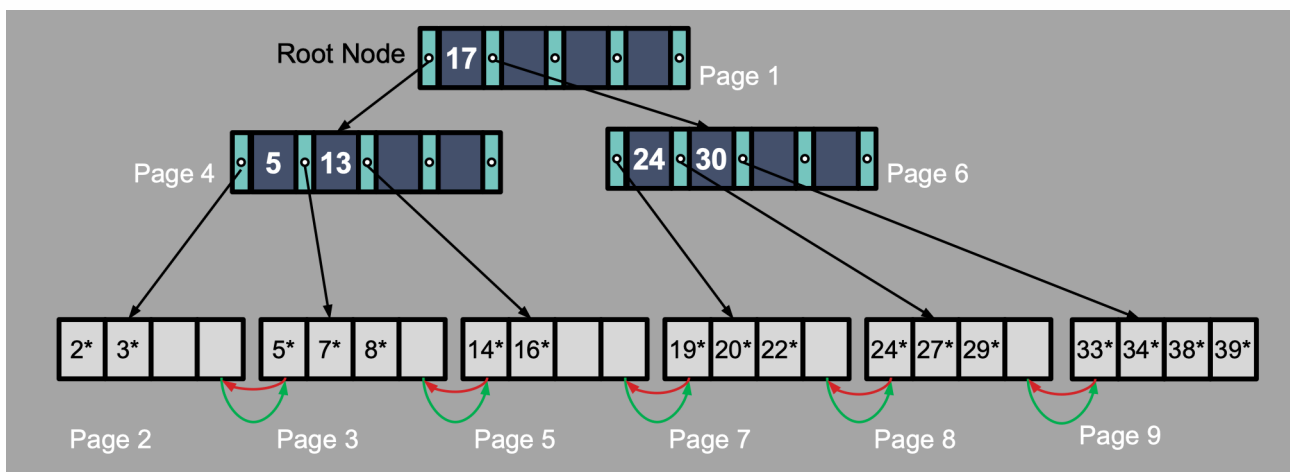
1. B^+ Tree

A quick review

- balanced tree
- internal nodes direct the search, the leaf nodes contain the data entries
- doubly linked list
- a parameter d := the order of the tree

$d \leq \# \text{ entries} \leq 2d$ for interior nodes

$1 \leq \# \text{ entries} \leq 2d$ for root node



Cost model for search

d := order of the B^+ Tree

f := fanout, $f \in [d + 1, 2d + 1]$, here assume it's constant for simplicity

N := total number of pages we'd like to index

F := fill factor (~usually $2/3$)

B := #available buffer pages

- our B^+ Tree needs to have room to index N/F pages

- what is the height h of our B^+ Tree?

$$h = \lceil \log_f \frac{N}{F} \rceil$$

- L_B is the number the number of levels such that the sum of all the levels' nodes fit in the buffer

$$B \geq 1 + f + \dots + f^{L_B-1} = \sum_{l=0}^{L_B-1} f^l$$

- IO cost: $\lceil \log_f \frac{N}{F} \rceil - L_B + 1$, where $B \geq 1 + f + \dots + f^{L_B-1} = \sum_{l=0}^{L_B-1} f^l$


$+\lceil \log_f \frac{N}{F} \rceil$: We read in one page per level of the tree

$-L_B$: However, levels that we can fit in buffer are free!

$+1$: Finally we read in the actual record

Examine the codes

find



```
function find(value V)
/* Returns leaf node C and index i such that C.Pi points to first record
* with search key value V */
Set C = root node
while (C is not a leaf node) begin
    Let i = smallest number such that V ≤ C.Ki
    if there is no such number i then begin
        Let Pm = last non-null pointer in the node
        Set C = C.Pm
    end
else if (V = C.Ki)
    then Set C = C.Pi+1
else C = C.Pi /* V < C.Ki */
end
/* C is a leaf node */
Let i be the least value such that Ki = V
if there is such a value i
    then return (C, i)
else return null /* No record with key value V exists */
```

insert

```
procedure insert(value K, pointer P)
if (tree is empty) create an empty leaf node L, which is also the root
else Find the leaf node L that should contain key value K
if (L has less than n-1 key values)
    then insert.in.leaf(L, K, P)
else begin /* L has n-1 key values already, split it */
    Create node L'
    Copy L.P1 ... L.Kn-1 to a block of memory T that can
        hold n (pointer, key-value) pairs
    insert.in.leaf(T, K, P)
    Set L'.Pn = L.Pn; Set L'.Pn-1 = L'
    Erase L.P1 through L.Kn-1 from L
    Copy T.P1 through T.K⌊n/2⌋} from T into L starting at L.P1
    Copy T.P⌊n/2⌋+1} through T.Kn from T into L' starting at L'.P1
    Let K' be the smallest key-value in L'
    insert.in.parent(L, K', L')
end

procedure insert.in.leaf(node L, value K, pointer P)
if (K < L.K1)
    then insert P, K into L just before L.P1
else begin
    Let Ki be the highest value in L that is less than K
    Insert P, K into L just after T.Ki
end

procedure insert.in.parent(node N, value K', node N')
if (N is the root of the tree)
    then begin
        Create a new node R containing N, K', N' /* N and N' are pointers */
        Make R the root of the tree
        return
    end
Let P = parent(N)
if (P has less than n pointers)
    then insert(K', N') in P just after N
else begin /* Split P */
    Copy P to a block of memory T that can hold P and (K', N')
    Insert (K', N') into T just after N
    Erase all entries from P; Create node P'
    Copy T.P1 ... T.P⌊n/2⌋} into P
    Let K'' = T.K⌊n/2⌋+1}
    Copy T.P⌊n/2⌋+1} ... T.Pn into P'
    insert.in.parent(P, K'', P')
end
```

2. Buffer Management

Exercises

1. What does it mean to say that a page is pinned in the buffer pool? Who is responsible for pinning pages? Who is responsible for unpinning pages?
2. Name an important capability of a DBMS buffer manager that is not supported by a typical operating system's buffer manager.
3. What happens if a page is requested when all pages in the buffer pool are dirty?