

Discussion 7 Query Optimization

Weijie Lyu




Reminder

1. Congratulations on finished Midterm!
2. Homework 2 due on Nov. 13, 23:59

Query Optimization - Background

- We can represent relational algebra expressions as trees
- Order of operators affects I/Os and resource usage, but not necessarily output



Query Optimization - Alternate Plans




- Given a plan, some things we can do are:
 - **Push selections/projections down the tree**
 - The earlier we reduce the size of our input data, the fewer I/Os are incurred as we traverse up the tree
 - Only affects I/O cost if materialized, or if operator only makes one pass (so not right relation of BNLJ)

Query Optimization - Alternate Plans

- Given a plan, some things we can do are:
 - **Push selections/projections down the tree**
 - **Materialize intermediate relations** (write to a temp file)
 - Results in additional write I/Os, but is better in the long run
 - **Use indices** (e.g. INLJ)

Query Optimization - Materializing




- Table A: takes 50 I/Os to perform a scan
- Table B: takes 100 I/Os to perform a scan
- $\text{Sel}(\text{B.age} < 25) = 0.5$, $[\text{B}] = 100$

Without materializing, we're performing $\sigma_{\text{age} < 25}$ on the fly each time in PNLJ, and scanning the entire table B for each page of A.

Cost = Scan A (50) + PNLJ (50*100)

→ 5,050 I/Os in total

Query Optimization - Materializing



- Table A: takes 50 I/Os to perform a scan
- Table B: takes 100 I/Os to perform a scan
- $\text{Sel}(B.\text{age} < 25) = 0.5$, $[B] = 100$

By materializing the intermediate relation, we're applying $\sigma_{\text{age} < 25}$ before PNLJ, and performing the join on the *result* of the selection.

Cost = *Scan A* (50) + *Scan B* (100)
+ *Materialize* (100 * 0.5) + *PNLJ* (50 * 50)
→ 2,700 I/Os in total

Query Optimization

- A query optimizer takes in a query plan (e.g. one directly translated from a SQL query), and outputs a better (hopefully optimal) query plan
 - Works on and optimizes over a **plan space** (set of all plans considered)
 - Performs **cost estimation** on query plans
 - Uses a **search algorithm** to search through plan space to find plan with lowest cost estimate
 - May not be optimal (bad estimate, or small plan space)


Query Optimization - Selinger

- We'll be looking at the System R optimizer (aka Selinger optimizer)
 - **Plan space:** only **left-deep** trees, avoid cartesian products unless they're the only option.
 - **Left-deep trees** represent a plan where all new tables are joined one at a time from the right.
 - **Cost estimation:** actual Selinger optimizer incorporates both CPU and I/O cost; we'll only use I/O cost for this class
 - **Search algorithm:** dynamic programming

Query Optimization - Selinger


- Why only left-deep trees?
 - Join new tables one at a time from the right
 - Create an ordering in which to add tables to the query being executed
 - Too many possible trees for joins
 - Using only left-deep trees: $N!$ different ways to order relations
 - Including all permutations tree layouts: A very large number of ways to parenthesize given an ordering (superexponential in N)

Query Optimization - Selinger




Left-deep

$((A \bowtie B) \bowtie C) \bowtie D$



Right-deep


$A \bowtie (B \bowtie (C \bowtie D))$



Bushy


$(A \bowtie B) \bowtie (C \bowtie D)$

Only consider left-deep plans!




Left-deep

$((A \bowtie B) \bowtie C) \bowtie D$



Right-deep

$A \bowtie (B \bowtie (C \bowtie D))$



Bushy

$(A \bowtie B) \bowtie (C \bowtie D)$

Query Optimization - Selinger

- Search algorithm for Selinger: use dynamic programming
 - Runtime drops from $n!$ to around $n \cdot 2^n$
- To be considered, must be:
 - Left deep
 - No cartesian products (i.e. if we join R and S on <cond1> and we join S and T on <cond2>, we don't consider joining R and T if there's no condition)

Query Optimization - Selinger

- For n relations joined, perform n passes
 - on the i -th path, output only the best plan for joining any i of the n relations
 - Also keep around plans that have higher cost but have an **interesting order**

Query Optimization - Interesting Orders

- **Interesting orders** are orderings on intermediate relations that *may* help reduce the cost of later joins
 - ORDER BY attributes
 - GROUP BY attributes
 - *downstream* join attributes

Query Optimization - Selinger

- **Pass 1:** find minimum cost access method for each (relation, interesting order) pair
 - Index scan, full table scans

A toy example:

SELECT *

FROM A, B, C

Pass 1:

- Full scan on A: 2 I/Os
- Index scan on A.b: 1 I/Os
- Full scan on B: 2 I/Os
- Full scan on C: 4 I/Os
- Index scan on C.c: 2 I/Os
- Index scan on C.d: 3 I/Os

Query Optimization - Selinger

- **Pass i** (Repeat until all relations are joined):
take in list of optimal plans for (i - 1 relations, interesting order) from Pass i-1, and compute minimum cost plan for (i relations, interesting orders)
(every size i subset of the n relations)

Pass 2:

- Index scan on A.b: 1 I/Os
- Full scan on B: 2 I/Os
- Index scan on C.c: 2 I/Os

- A BNLJ B: 5 I/Os

- B INLJ A: 6 I/Os

- C PNLJ A: 6 I/Os

- B BNLJ C: 5 I/Os

- C INLJ B: 6 I/Os

Query Optimization - Selinger

- **Pass i** (Repeat until all relations are joined):
take in list of optimal plans for (i - 1 relations, interesting order) from Pass i-1, and compute minimum cost plan for (i relations, interesting orders)
(every size i subset of the n relations)

Pass 3:

- A BNLJ B: 5 I/Os
- C PNLJ A: 6 I/Os
- B BNLJ C: 5 I/Os

- (AB) BNLJ C: 14 I/Os
- (CA) INLJ B: 13 I/Os
- (CA) BNLJ B: 12 I/Os
- (BC) PNLJ A: 13 I/Os

Selectivity Estimation

- To estimate the cost of a query, we add up the estimated costs of each operator in the query
 - Need to know the size of the **intermediate relations** (generated from one operator and passed into another) in order to do this!
 - Need to know the **selectivity** of predicates - what % of tuples are selected by a predicate
- These are all estimates: if we don't know, we make up a value for it (selectivity = 1/10)

Selectivity Estimation - Equalities

Predicate	Selectivity	Assumption
$c = v$	$1 / (\text{number of distinct values of } c \text{ in index})$	We know $ c $.
$c = v$	$1 / 10$	We don't know $ c $.
$c1 = c2$	$1 / \text{MAX}(\text{number of distinct values of } c1, \text{ number of distinct values of } c2)$	We know $ c1 $ and $ c2 $.
$c1 = c2$	$1 / (\text{number of distinct values of } c_i)$	We know $ c_i $ but not $ \text{other column} $.
$c1 = c2$	$1 / 10$	We don't know $ c1 $ or $ c2 $.

- **lcolumnl** = the number of distinct values for the column
- If you have an index on the column, you can assume you know **lcolumnl**, **max(c)**, and **min(c)**
- When applying selectivity to # of records, take the floor of the result. (e.g. $256.3 \rightarrow 256$ records)

Selectivity Estimation - Inequalities on Integers

Predicate	Selectivity	Assumption
$c < v$ $c > v$	$(v - \min(c)) / (\max(c) - \min(c) + 1)$ $(\max(c) - v) / (\max(c) - \min(c) + 1)$	We know $\max(c)$ and $\min(c)$. c is an integer.
$c < v$ $c > v$	$1 / 10$	We don't know $\max(c)$ and $\min(c)$. c is an integer.
$c \leq v$ $c \geq v$	$(v - \min(c)) / (\max(c) - \min(c) + 1) + (1 / c)$ $(\max(c) - v) / (\max(c) - \min(c) + 1) + (1 / c)$	We know $\max(c)$ and $\min(c)$. c is an integer.
$c \leq v$ $c \geq v$	$1 / 10$	We don't know $\max(c)$ and $\min(c)$. c is an integer.

* We add 1 to the denominator in order for our [low, high] range to be inclusive.

E.g. range [2, 4] = 2, 3, 4 $\rightarrow (4 - 2) + 1 = 3$

Selectivity Estimation - Inequalities on Floats

Predicate	Selectivity	Assumption
$c \geq v$	$(\max(c) - v) / (\max(c) - \min(c))$	We know $\max(c)$ and $\min(c)$. c is a float.
$c \geq v$	$1 / 10$	We don't know $\max(c)$ and $\min(c)$. c is a float.
$c \leq v$	$(v - \min(c)) / (\max(c) - \min(c))$	We know $\max(c)$ and $\min(c)$. c is a float.
$c \leq v$	$1 / 10$	We don't know $\max(c)$ and $\min(c)$. c is a float.

* We don't add 1 to the denominator. floats are continuous, integers are discrete)

E.g. range $[2.0, 4.0] = 2.0, 2.1, \dots, 3.9, 4.0 \rightarrow 4.0 - 2.0 = 2.0$

Selectivity Estimation - Connectives

Predicate	Selectivity	Assumption
p1 AND p2	$S(p1) * S(p2)$	Independent predicates
p1 OR p2	$S(p1) + S(p2) - S(p1) * S(p2)$	Independent predicates
NOT p	$1 - S(p)$	

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
```

1000 tuples

(no predicates, select all)

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE a = 42;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in

c = v	1 / (number of distinct values of c in index)	We know c .
-------	---	--------------

• NO INDEX ON C

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R  
WHERE a = 42;
```

50 unique values in **a**

$1/50 * (1000 \text{ tuples}) = 20 \text{ tuples}$

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE c = 42;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in

c = v	1 / 10	We don't know c .
-------	--------	--------------------

- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R  
WHERE c = 42;
```

no information about **c**

$1/10 * (1000 \text{ tuples}) = 100 \text{ tuples}$

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE a <= 25;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in

c <= v	$(v - \min(c)) / (\max(c) - \min(c) + 1) + (1 / c)$	We know max(c) and min(c). c is an integer.
c >= v	$(\max(c) - v) / (\max(c) - \min(c) + 1) + (1 / c)$	

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE a <= 25;
```

Sel(a <= 25)

$$= (25 - 1) / (50 - 1 + 1) + 1/50 = 1/2$$

$$1/2 * (1000 \text{ tuples}) = 500 \text{ tuples}$$

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE b <= 25;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float

c <= v	$(v - \min(c)) / (\max(c) - \min(c))$	We know max(c) and min(c). c is a float.
--------	---------------------------------------	---

- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE b <= 25;
```

Sel(b <= 25)

= (25 - 1)/(100 - 1) = 24/99 = 0.2424...

floor(0.2424... * (1000 tuples)) = 242 tuples

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
  WHERE a <= 25
        AND b <= 25;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]

p1 AND p2	$S(p1) * S(p2)$	Independent predicates
------------------	-----------------	------------------------

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE a <= 25
      AND b <= 25;
```

$$\text{Sel}(a \leq 25) * \text{Sel}(b \leq 25) \\ = \frac{1}{2} * \frac{24}{99} = 0.1212\dots$$

$$\text{floor}(0.1212\dots * (1000 \text{ tuples})) = 121 \text{ tuples}$$

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE a = c;
```

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in

c1 = c2	1 / (number of distinct values of ci)	We know ci but not other column .
---------	---------------------------------------	--------------------------------------

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R  
WHERE a = c;
```

no information about **c**

$1/50 * (1000 \text{ tuples}) = 20 \text{ tuples}$

- R(a, b, c) has 1000 tuples
- index on a with 50 unique integer values, uniformly distributed in the range [1, 50]
- index on b with 100 unique float values, uniformly distributed in the range [1, 100]
- no index on c

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE R.a = S.a;
```

- R(a, b, c) has 1000 tuples
- index on R.a with 50 unique integer values, uniformly distributed in the range [1, 50]
- S(a) has 500 tuples
- index on S.a with 25 unique

c1 = c2	1 / MAX(number of distinct values of c1, number of distinct values of c2)	We know c1 and c2 .
---------	---	------------------------

Selectivity Estimation - Worksheet

How many tuples are selected by the following query?

```
SELECT * FROM R
WHERE R.a = S.a;
```

$\text{Sel}(\mathbf{R.a = S.a})$

$= 1/\text{MAX}(50, 25) = 1/50$

$1/50 * (1000 \text{ tuples} * 500 \text{ tuples}) = 10,000 \text{ tuples}$

- R(a, b, c) has 1000 tuples
- index on R.a with 50 unique integer values, uniformly distributed in the range [1, 50]
- S(a) has 500 tuples
- index on S.a with 25 unique integer values, uniformly distributed in the range [1, 25]