

# CS150 Discussion 8

## Transactions & Concurrency

Jiachun Jin  
jinjch@shanghaitech.edu.cn

# Motivation

- ▶ Many users work with the application at the same time



- ▶ Multiple clients running SQL at the same time, over one database
  - ▶ concurrency provides efficiency
    - ▶ improper algorithms introduces weird bugs
  - ▶ fault tolerant
  - ▶ reliable

# Outline

- ▶ Transaction
- ▶ Concurrency

# Transaction

- ▶ A **transaction** is a sequence of one or more operations (reads or writes) which reflects a single real-world transition
  - ▶ transfer money between accounts
- ▶ In a program, multiple statements (rd & wt) can be grouped together as a transaction:

```
START TRANSACTION
  UPDATE Bank SET amount = amount - 100
  WHERE name = 'Bob'
  UPDATE Bank SET amount = amount + 100
  WHERE name = 'Joe'
COMMIT
```

# Property

- ▶ **A tomicity**: All actions in the **transaction** happen, or none happen
- ▶ **C onsistency**: If the DB starts out consistent, it ends up consistent at the end of the **transaction**
  - ▶ tables must always satisfy user-specified integrity constraints
- ▶ **I solation**: Execution of each **transaction** is isolated from that of others
  - ▶ concurrency relies on this property
- ▶ **D urability**: If a **transaction** commits, its effects persist

# Outline

- ▶ Transaction
- ▶ Concurrency

# Interleave transactions

- ▶ Individual **transaction** might be slow– don't want to block other users during!
- ▶ We must pick an interleaving or schedule such that isolation and consistency are maintained

# An example

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Interleaved schedule A:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Same result!

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$112

Different result than serial  $T_1, T_2$ !

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_2, T_1$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$153	\$112

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$112

Different result than serial  $T_2, T_1$  ALSO!

Scheduling examples

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

This schedule is different than *any serial order*! We say that it is not serializable

what does serializable mean?



# Scheduling Definitions

- ▶ A **serial** schedule is one that does not interleave the actions of different transactions
- ▶ A and B are **equivalent** schedules if, for any database state, the effect on DB of executing A is identical to the effect of executing B
- ▶ A **serializable** schedule is a schedule  $S$  that is equivalent to **some** serial execution of the transactions  $S_{serial}$ 
  - ▶ revisit the above example

# An example

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Interleaved schedule A:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Same result!

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_1, T_2$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$106

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$112

Different result than serial  $T_1, T_2$ !



: serializable



: not serializable

Scheduling examples

Starting Balance

A	B
\$50	\$200

Serial schedule  $T_2, T_1$ :

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$153	\$112

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

$T_2$

$A *= 1.06$  $B *= 1.06$

A	B
\$159	\$112

Different result than serial  $T_2, T_1$  ALSO!

Scheduling examples

Interleaved schedule B:

$T_1$

$A += 100$  $B -= 100$

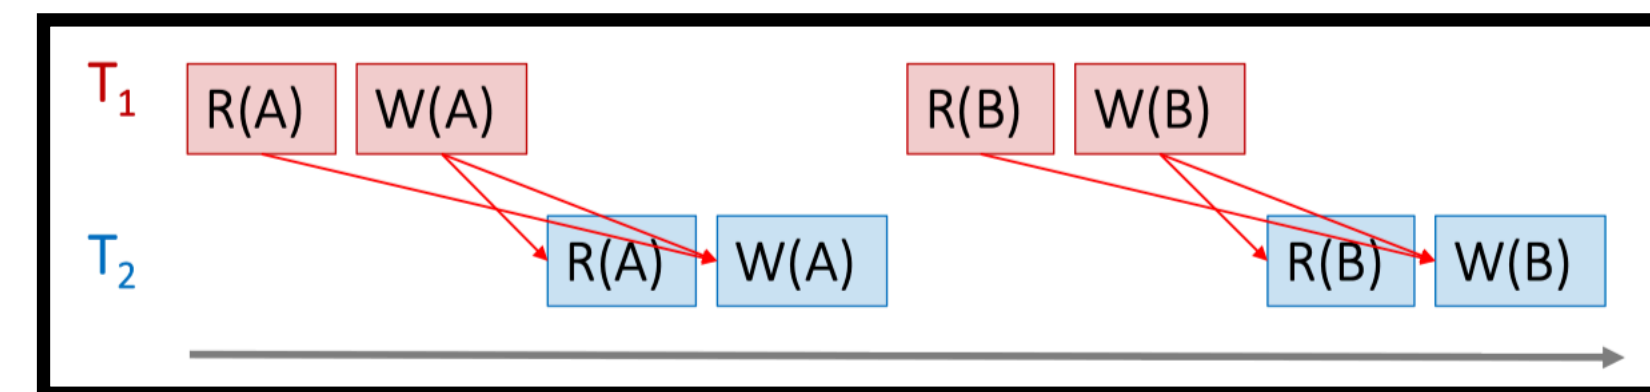
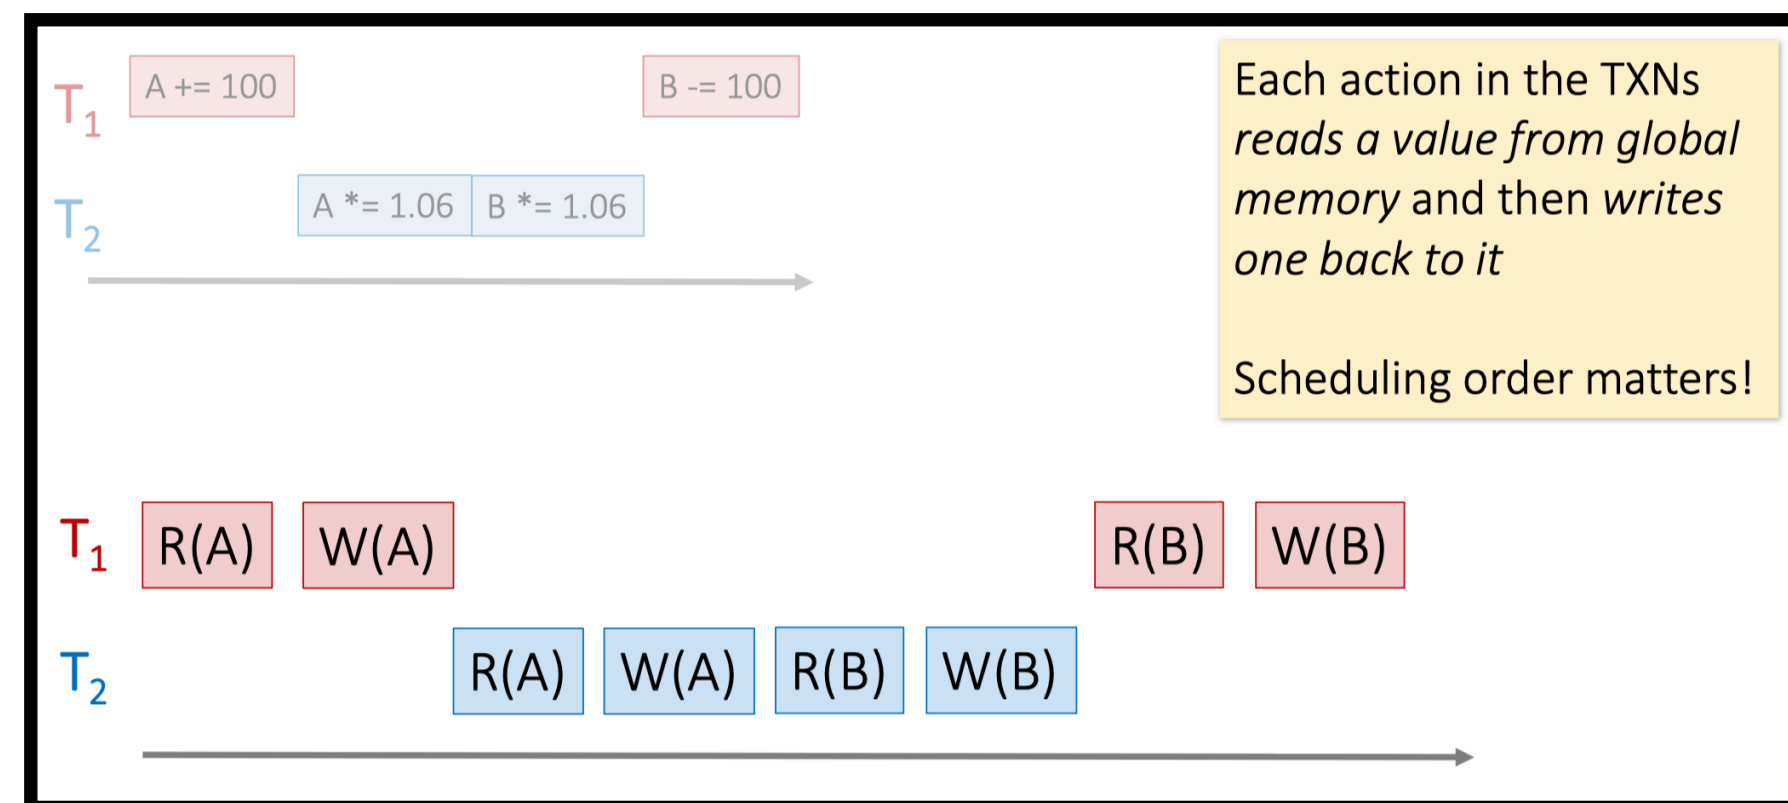
$T_2$

$A *= 1.06$  $B *= 1.06$

This schedule is different than *any serial order*! We say that it is not serializable

# Conflicts

- ▶ Goal: discerning “good” vs. “bad” schedules
  - ▶ serializable schedule will maintain isolation & consistency (to some extend “good”)
  - ▶ a stricter, but very useful variant: **conflict serializability**
- ▶ The DBMS’s view of the schedule



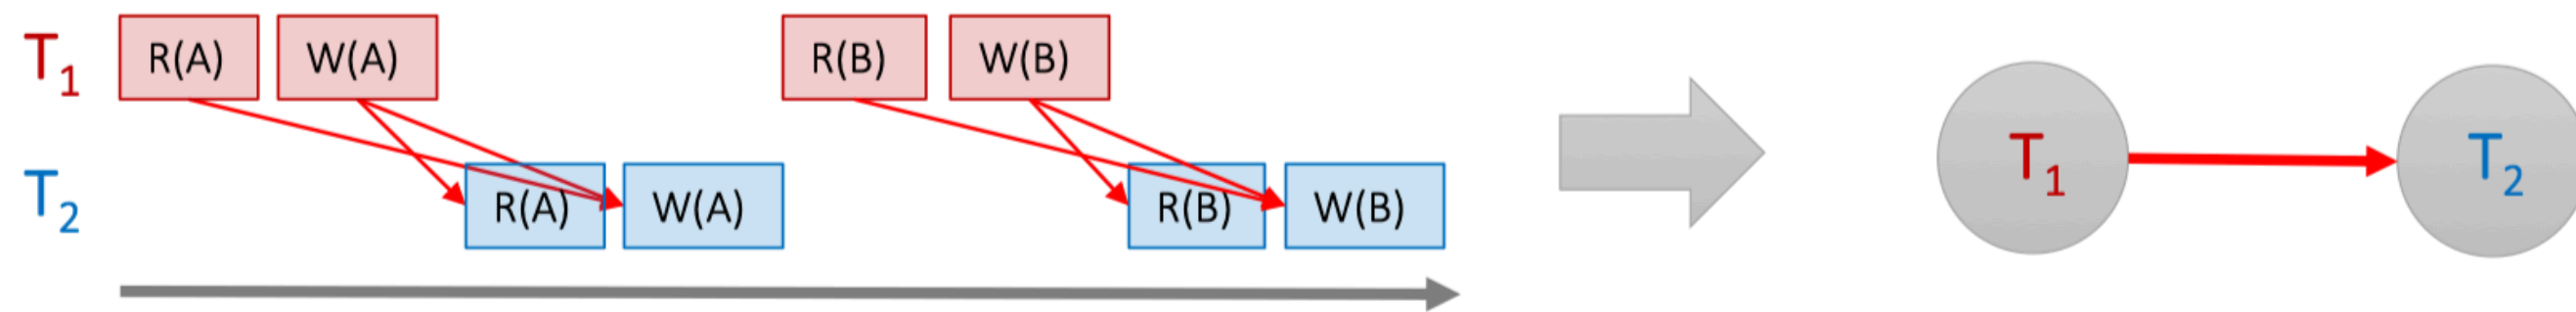
- ▶ Two actions **conflict** if they are part of **different transactions**, involve the **same variable**, and at least one of them is a **write**

# Conflict serializable

- ▶ The order of non-conflicting operations has no effect on the final state of the database
- ▶ We need more definitions
  - ▶ two schedules are **conflict equivalent** if:
    - ▶ they involve the same actions of the same transactions
    - ▶ **every pair** of **conflicting actions** of two transactions are ordered in the same way
  - ▶ schedule  $S$  is **conflict serializable** if  $S$  is **conflict equivalent** to some serial schedule
    - ▶ we like **conflict serializable** schedules


# Conflict graph

- Looking at conflicts at the transaction level
  - there is an edge from  $T_i \rightarrow T_j$  if an action in  $T_i$  precede and conflict with an action in  $T_j$



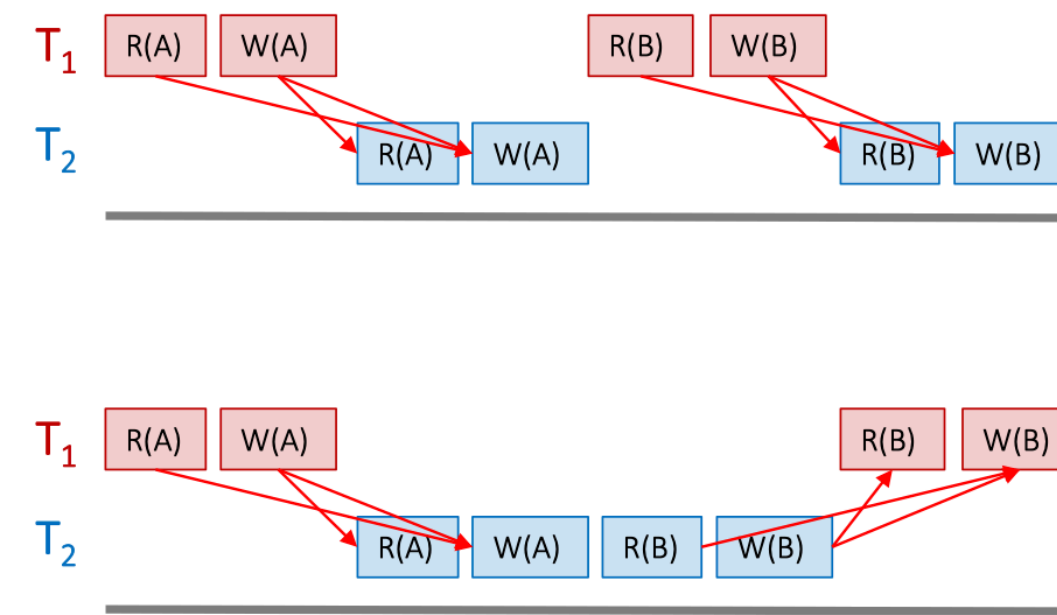


- Theorem: schedule is conflict serializable  $\Leftrightarrow$  its conflict graph is acyclic

Serial Schedule:



Interleaved Schedules:



# A mindmap

