

Discussion 9

Lock & Recovery

Transactions-ACID

Transactions

- We want transactions to obey **ACID**
 - **atomicity**: all operations in a transaction happen, or none of them
 - **consistency**: database consistency (unique constraints, etc) is maintained
 - **isolation**: should look like we only run 1 transaction at a time (even if we run multiple concurrently)
 - **durability**: once a transaction commits, it persists
- **commit**: indicates successful transaction (save changes)
- **abort**: indicates unsuccessful transaction (revert changes)

Types of Serializability


All Schedules

Serializable

View Serializable

Conflict Serializable

Serial



lock

TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
 - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
 - If a TXN holds, no other TXN can get an X lock on that object
- TXNs cannot get new locks after releasing any locks.

Simple Locking

- The lock compatibility matrix looks like:

	NL (no lock held)	S	X
NL (no lock held)	✓	✓	✓
S	✓	✓	
X	✓		

Deadlocks

- What if T_1 is waiting for T_2 to release a lock, but T_2 is also waiting for T_1 to release a lock?
 - This is called a **deadlock** - when a bunch of transactions are waiting on each other in a cycle
- We can either avoid them in the first place (**deadlock avoidance**) or catch them (**deadlock detection**) and abort a transaction in the deadlock

Deadlock Detection

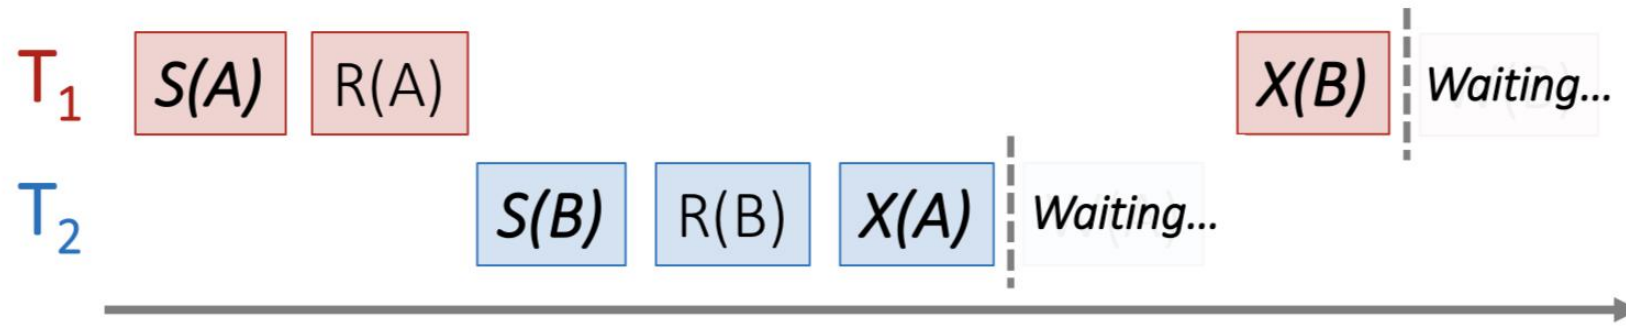
- We draw out a “waits-for” graph

- One node for each transaction
- If T_i holds a lock that conflicts with the lock that T_j wants (or T_j “waits for” T_i), we add an edge from T_j to T_i




- A cycle indicates a deadlock (between the transactions in the cycle)
 - we can abort one to end the deadlock
- Alternative approach (used in some real databases): just kill transactions if they aren't doing anything for a while

Deadlock Detection: Example



Finally, T_1 requests an exclusive lock on B to write to it- **now T_1 is waiting on T_2 ... DEADLOCK!**


Waits-for graph:



Cycle =
DEADLOCK

2-Phase Locking (2PL)

- One way to enforce conflict serializability
- In 2-phase locking,
 - a transaction may not acquire a lock after it has released any lock
 - two “phases”
 - from start to until a lock is released, the transaction is only acquiring locks
 - then until the end of the transaction, it is only releasing locks



T1	T2
Lock_X(A)	
Read(A)	
	Lock_S(A)
A: = A-50	
Write(A)	
Lock_X(B)	
Unlock(A)	
	Read(A)
	Lock_S(B)
Read(B)	
B := B +50	
Write(B)	
Unlock(B)	
	Unlock(A)
	Read(B)
	Unlock(B)
	PRINT(A+B)

Cascading Aborts


- Consider the following (+ for acquire, - for release):

T_1	+X(A)	-X(A)						abort!
T_2			+S(A)					
T_3			+S(A)	+X(B)	-X(B)			
T_4						+S(B)		
T_5							+S(B)	

- T_1 aborting means all of the transactions have to be rolled back
 - Even though T_4 and T_5 didn't read A at all
 - This is a **cascading abort**

Strict 2-Phase Locking (Strict 2PL)

- The problem is that 2PL lets another transaction read new values before the transaction commits (since locks can be released long before commit)
- **Strict 2PL** avoids cascading aborts
 - Same as 2PL, except only allow releasing locks at end of transaction



T1	T2
Lock_X(A)	
Read(A)	
	Lock_S(A)
A: = A-50	
Write(A)	
Lock_X(B)	
Read(B)	
B := B +50	
Write(B)	
Unlock(A)	
Unlock(B)	
	Read(A)
	Lock_S(B)
	Read(B)
	PRINT(A+B)
	Unlock(A)
	Unlock(B)

Strict 2PL


Theorem: Strict 2PL allows only schedules whose dependency graph is acyclic

Proof Intuition: In strict 2PL, if there is an edge $T_i \rightarrow T_j$ (i.e. T_i and T_j conflict) then T_j needs to wait until T_i is finished – so *cannot* have an edge $T_j \rightarrow T_i$

Therefore, Strict 2PL only allows conflict serializable \Rightarrow serializable schedules


logging

Two Important Logging Decisions



- **Decision 1: STEAL or NO-STEAL**
 - Impacts ATOMICITY and UNDO
 - Steal: allow the buffer pool (or another txn) to “steal” a pinned page of an uncommitted txn by flushing to disk
 - No-steal: disallow
 - If we allow “Steal”, then need to deal with uncommitted txn edits appearing on disk
 - To ensure Atomicity we need to support UNDO of uncommitted txns
 - OTOH “No-steal” has poor performance (**pinned pages limit buffer replacement**)
 - But no UNDO required. Atomicity for free.

Two Important Logging Decisions



- **Decision 2: FORCE or NO-FORCE**

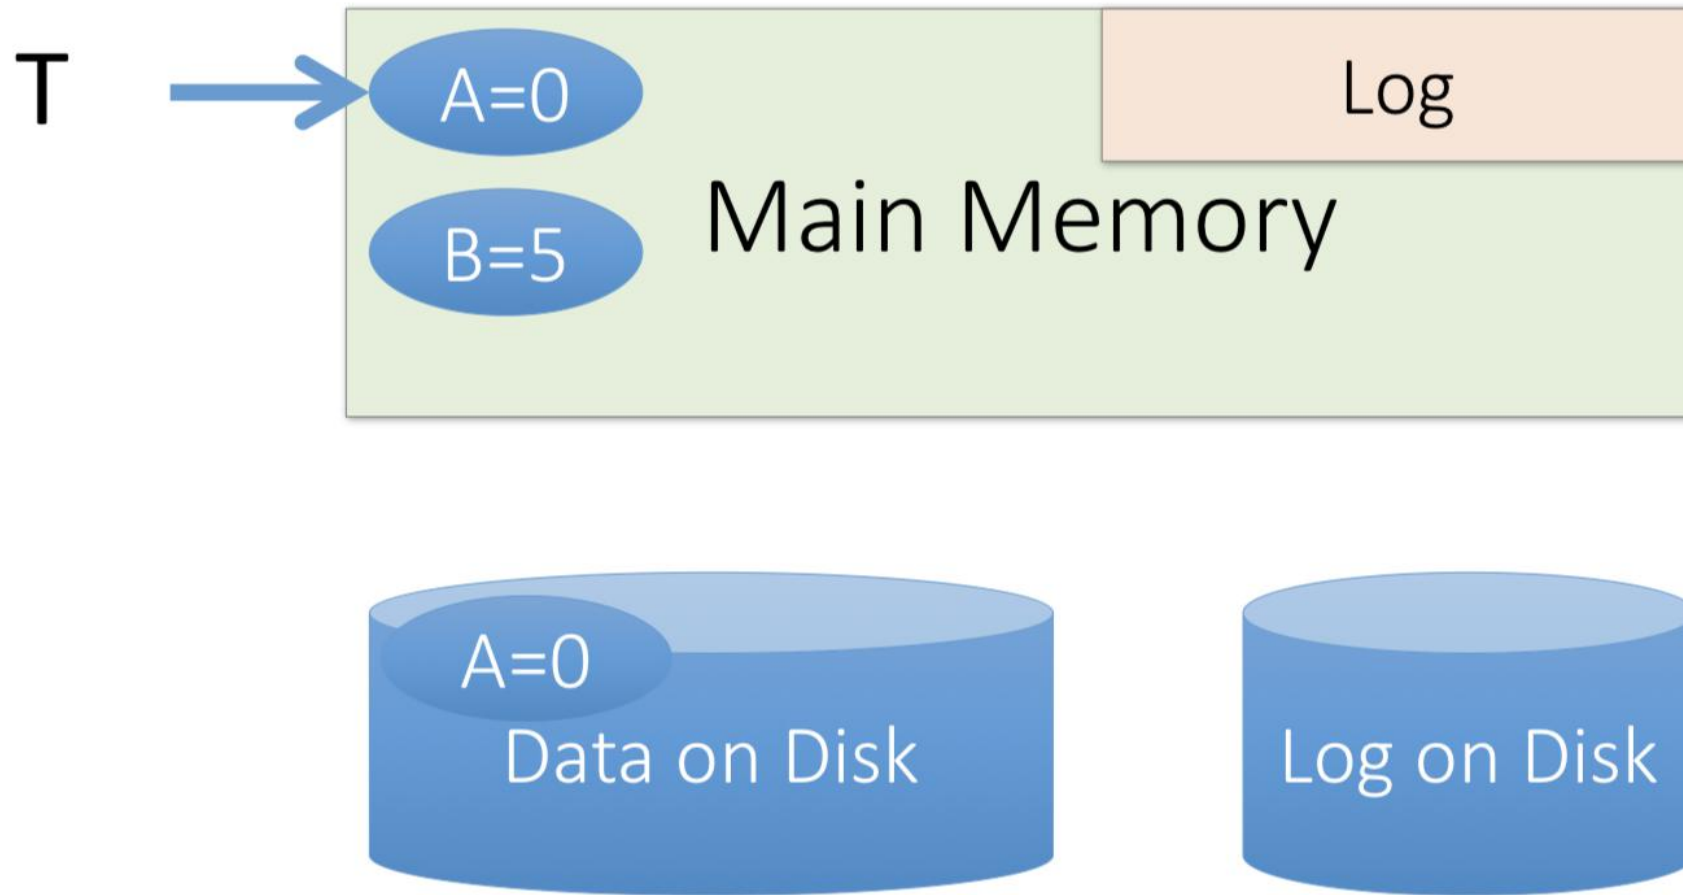
- Impacts DURABILITY and REDO
- Force: ensure that all updates of a transaction is “forced” to disk prior to commit
- No-force: no need to ensure
- If we allow “No-force”, then need to deal with committed txns not being durable
 - To ensure Durability we need to support REDO of committed txns
- OTOH, “Force” has poor performance (lots of random I/O to commit)
 - But no REDO required, Durability for free.

- The **log** consists of **an ordered list of actions**
 - Log record contains:
<XID, location, old data, new data>

This is sufficient to UNDO any transaction!

A picture of logging


T: R(A), W(A)



Incorrect Commit Protocol #1

T: R(A), W(A)

A: 0 → 1



Let's try committing *before* we've written either data or log to disk...

OK, Commit!

If we crash now, is T durable?



Lost T's update!

Incorrect Commit Protocol #2

T: R(A), W(A)

A: 0 → 1

T



Let's try committing
after we've written
data but *before* we've
written log to disk...


OK, Commit!

If we crash now, is T
durable? Yes! Except...

***How do we know
whether T was
committed??***

Write-ahead Logging (WAL) Commit Protocol

T: R(A), W(A)



This time, let's try committing after we've written log to disk but before we've written data to disk... this is WAL!

OK, Commit!

If we crash now, is T durable?

USE THE LOG!

WAL

- The **Write -Ahead Logging Protocol** :
 1. Must **force** the **log record** for an update **before** the corresponding **data page** gets to the DB disk.
 2. Must **force all log records** for a Xact **before commit**.
 - I.e. transaction is not committed until all of its log records including its “commit” record are on the stable log.