
Homework 1 Writing

School of Information Science and Technology

XXX

April 2, 2019

1 SINGLE-RELATION QUERIES

1. Consider the following relation:

Graph(n1 , n2)

A tuple (n1, n2) in Graph stores a directed edge from a node n1 to a node n2 in the corresponding graph. Your goal is to, for *every* node in the graph, count the number of outgoing edges of that node. Note that for nodes without any outgoing edges, their edge count would be zero; you need to output this as well. You can assume that:

- a) there are no duplicates or null values in the table; and
- b) every node in the graph is involved in at least one edge.

Solution:

```
SELECT n1, COUNT(n1)
FROM Graph
GROUP BY Graph.n1;
UNION
SELECT n2, 0
FROM Graph as g1
WHERE g1.n2 NOT IN (SELECT g2.n1
FROM Graph as g2
);
```

2. Consider the following relation:

Trained(student, master, year)

A tuple (S, M, Y) in Trained specifies that a SQL Master M trained student S who graduated in year Y. Your goal is to find *the count of* SQL Masters who trained a student who

graduated in the same year that 'Alice' or 'Bob' graduated.

Solution:

```
SELECT COUNT(DISTINCT t1.master) as CountofMaster
FROM Trained as t1
WHERE t1.year IN ( SELECT t2.year
FROM Trained as t2
WHERE t2.student = "Alice"
OR t2.student = "Bob"
);
```

3. Consider the following relation:

DBMS(operator, system, performance)

A tuple (O, S, P) in DBMS specifies an operator O in system S and has the performance value P. Your goal is to find those systems whose operators achieves a higher performance value on average than the average performance value in a system named 'PostgreSQL'.

Solution:

```
SELECT d3.system
FROM ( SELECT d1.system as system, AVG(d1.performance) as performan
FROM DBMS as d1
GROUP BY d1.system
) as d3
WHERE d3.performance > ( SELECT AVG(d2.performance)
FROM DBMS as d2
WHERE d2.system = "PostgreSQL"
);
```

2 MULTI-RELATION QUERIES

Consider the following relations representing student information at UIUC:

Mentorship(mentee_sid, mentor_sid)

Study(sid, credits)

Enrollment(did, sid)

Student(sid, street, city)

- A tuple (M1, M2) in Mentorship specifies that M2 is a mentor of another student M1.
- A tuple (S, C) in Study specifies that the student S has taken C credits.
- A tuple in Enrollment (D, S) specifies that student S is enrolled in department D.
- A (ST, S, C) in Student specifies that student ST lives on street S in city C.

1. Find all students who live in the same city and on the same street as their mentor.

Solution:

```

SELECT m1.mentee_sid
FROM Mentorship as m1, Student as s1, Student as s2
WHERE m1.mentee_sid = s1.sid
AND m1.mentor_sid = s2.sid
AND s1.street = s2.street
AND s1.city = s2.city;

```

2. Find all students(i.e., distinct sid) who have taken more credits than the average credits of all of the students of their department.

Solution:

```

SELECT s1.sid
FROM Study as s1, Enrollment as e1
WHERE e1.sid = s1.sid
AND s1.credits > ( SELECT AVG(s2.credits)
FROM Study as s2, Enrollment as e2
WHERE s2.sid = e2.sid
AND e1.did = e2.did
);

```

3 DATABASE MANIPULATION AND VIEWS

1. In the Study relation, insert a new student, whose id is 66666 and has 0 credits.

Solution:

```

INSERT INTO Study
VALUES('66666', '0');

```

2. In the Study relation, delete students who have graduated (i.e., the ones who have more than 200 credits).

Solution:

```

DELETE FROM Study
WHERE Study.credits > 200;

```

3. In the Study relation, add 2 credits for students who are mentors.

Solution:

```

UPDATE Study
SET Study.credits = Study.credits + 2
WHERE Study.sid IN ( SELECT m1.mentor_sid
FROM Mentorship as m1
);

```

4. Incoming students are those who have been accepted (i.e., exist in the **Student** relation) but have not registered in any department (i.e., do not exist in the **Enrollment** relation). Create a View that contains **sid** of all incoming students.

Solution:

```
CREATE VIEW Incoming AS
SELECT s1.sid
FROM Student as s1
WHERE s1.sid NOT IN ( SELECT e1.sid
FROM Enrollment as e1
);
```