

Tree Indexes




R & G - Chapter 10



Reminder on Heap Files

- Two access APIs:
 - fetch by recordId (pageId, slotId)
 - scan (starting from some page)




Wouldn't it be nice...

- ...if we could look things up by value?
- Toward a Declarative access API
- But ... efficiency?

“If you don't find it in the index, look very carefully through the entire catalog.”

—Sears, Roebuck, and Co., Consumers' Guide, 1897





We've seen this before

- Data structures ... in RAM:
 - Search trees (Binary, AVL, Red-Black, ...)
 - Hash tables
- Needed: disk-based data structures
 - “paginated”: made up of disk pages!



Index

An **index** is data structure that enables fast **lookup** and **modification** of **data entries** by **search key**

- **Lookup:** may support many different operations
 - **Equality**, 1-d range, 2-d region, ...
- **Search Key:** any subset of columns in the relation
 - Do not need to be unique
 - —e.g. (firstname) or (firstname, lastname)



Index Part 2

An **index** is data structure that enables fast **lookup** and **modification** of **data entries** by **search key**

- **Data Entries:** items stored in the index
 - Assume for today: a pair (**k**, recordId) ...
 - Pointers to records in Heap Files!
 - Easy to generalize later
- **Modification:** want to support fast insert and delete

Many Types of indexes exist: B+-Tree, Hash, R-Tree, GiST, ...




Simple Idea?

Input
Heap
File



- **Step 1:** Sort heap file & leave some space
 - Pages physically stored in logical order (sequential access)
 - Do we need “next” pointers to link pages?
 - No. Pages are physically sorted in logical order




- **Step 2:** Build the index data structure over this...
 - Why not just use binary search in this heap file?
 - Fan-out of 2 → deep tree → lots of I/Os
 - Examine entire records just to read key during search



Build a high fan-out search tree


- Start simple: *Sorted (key, page id) file*
 - No record data
 - Binary search in the key file. Better!
 - **Forgot:** Need to break across pages!





Build a high fan-out search tree


- Start simple: *Sorted (key, page id) file*
 - No record data
 - Binary search in the key file. Better!
 - **Forgot:** Need to break across pages!





Build a high fan-out search tree Part 2


- Start simple: *Sorted (key, page id) file*
 - No record data
 - Binary search in the key file. Better!
 - **Complexity?**





Build a high fan-out search tree Part 3


- Start simple: *Sorted (key, page id) file*
 - No record data
 - Binary search in the key file. Better!
 - **Complexity**: Still binary search, just a constant factor smaller input





Build a high fan-out search tree Part 4


- Recursively “index” key file
- **Key Invariant:**
 - Node $[..., (K_L, P_L), (K_R, P_R), ...]$ \rightarrow All tuples in range $K_L \leq K < K_R$ are in tree P_L





Search a high fan-out search tree


- Searching for **5**?
 - Binary Search each node (page) starting at root
 - Follow pointers to next level of search tree
- Complexity? $O(\log_F(\#Pages))$





Left Key Optimization?


- Optimization
 - Do we need the left most key?





Build a high fan-out search tree

- Disk Layout? All in a single file, Data Pages first.






Status Check

- Some design goals:
 - Fast sequential access?
 - High Fan-out
 - Support insertions

Indexed File

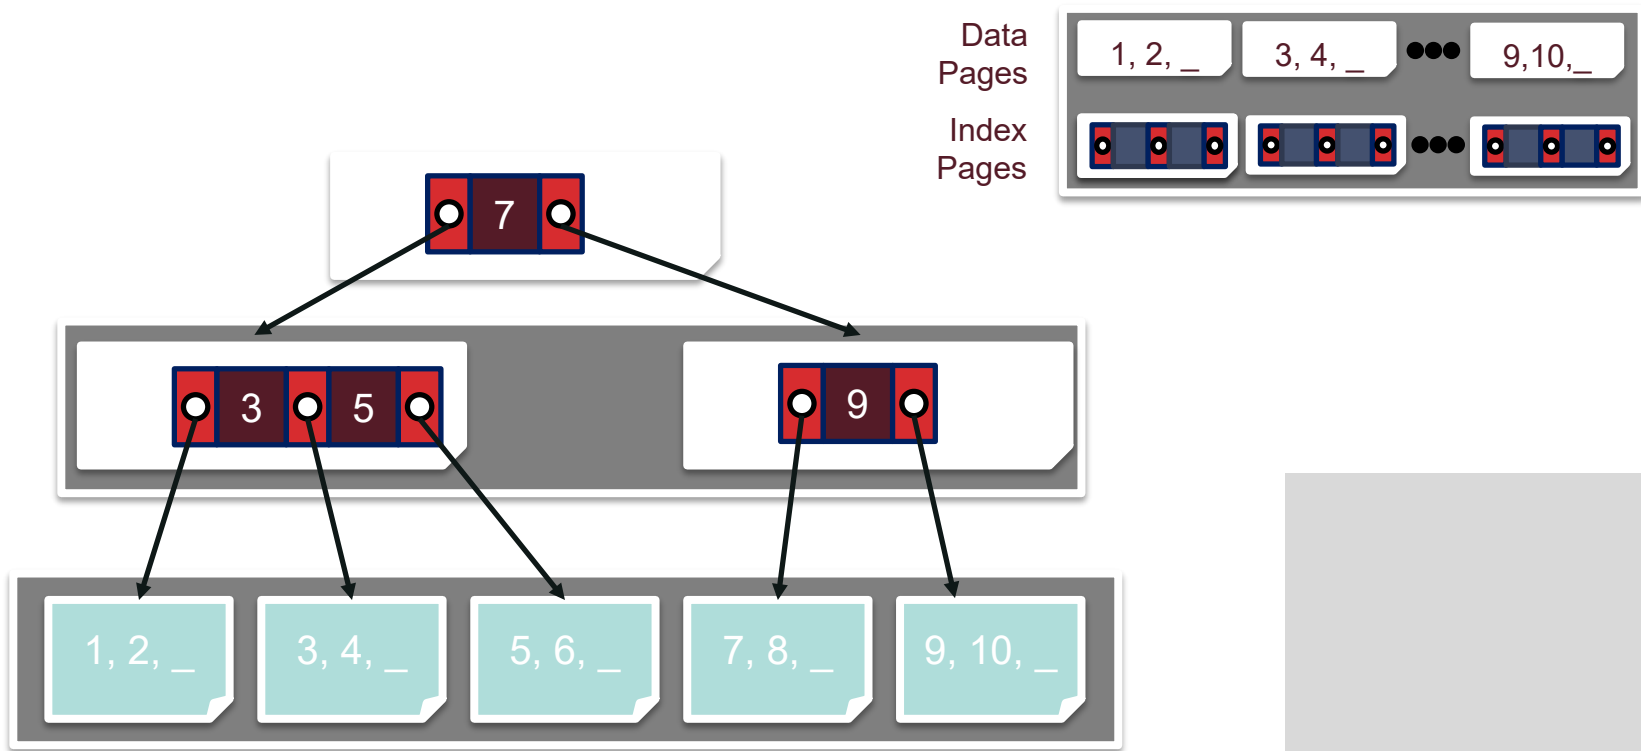
ISAM Indexed Sequential Access Method

(Early IBM Indexing Technology)





Insert 11, Before







Insert 11, After

- Find location
- Place in data page
 - Re-sort page ...



Data
Pages

Index
Pages





- Find location
- Place in data page
- Add overflow page if necessary ...





Recap: ISAM

- Data entries in sorted heap file
- High fan-out static tree index
- Fast search + good locality
 - Assuming nothing changes
- Insert into overflow pages



A Note of Caution

- ISAM is an old-fashioned idea
 - Introduced by IBM in 1960s
 - B+ trees are usually better, as we'll see
 - Though not always (← we'll come back to this)
- But, it's a good place to start
 - Simpler than B+ tree, many of the same ideas
- Upshot
 - Don't brag about ISAM on your resume
 - Do understand ISAM, and tradeoffs with B+ trees

B+-TREE




Enter the B+ Tree

- Similar to ISAM
 - Same interior node structure
 - $\langle \text{Key}, \text{Page Ptr} \rangle$ pairs with same key invariant
 - Same search routine as before
- **Dynamic Tree Index**
 - Always Balanced
 - Support efficient insertion & deletion
 - Grows at root not leaves!
- “+”? B-tree that stores data entries in leaves only



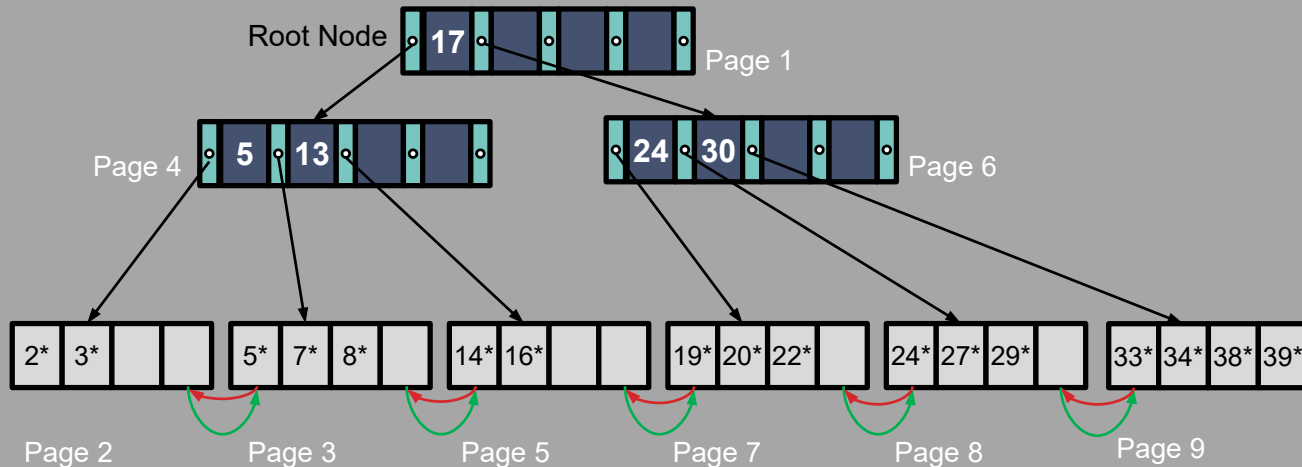
Example of a B+ Tree



- Occupancy Invariant
 - Each interior node is at least partially full:
 - $d \leq \#entries \leq 2d$
 - d : order of the tree (max fan-out = $2d + 1$)
- Data pages at bottom need not be stored in logical order
 - Next and prev pointers



Sanity Check



What is the value of d?

2

What about the root?


The root is special

Why not in sequential order?

Data pages allocated dynamically




B+ Trees and Scale



- How big is a height 1 B+ tree
 - $d = 2 \rightarrow$ Fan-out?
 - Fan-out = $2d + 1 = 5$
 - **Height 1:** $5 \times 4 = 20$ Records



B+ Trees and Scale Part 2



- How big is a height **3** B+ tree
 - $d = 2 \rightarrow$ Fan-out?
 - Fan-out = $2d + 1 = 5$
 - **Height 3:** $5^3 \times 4 = 500$ Records

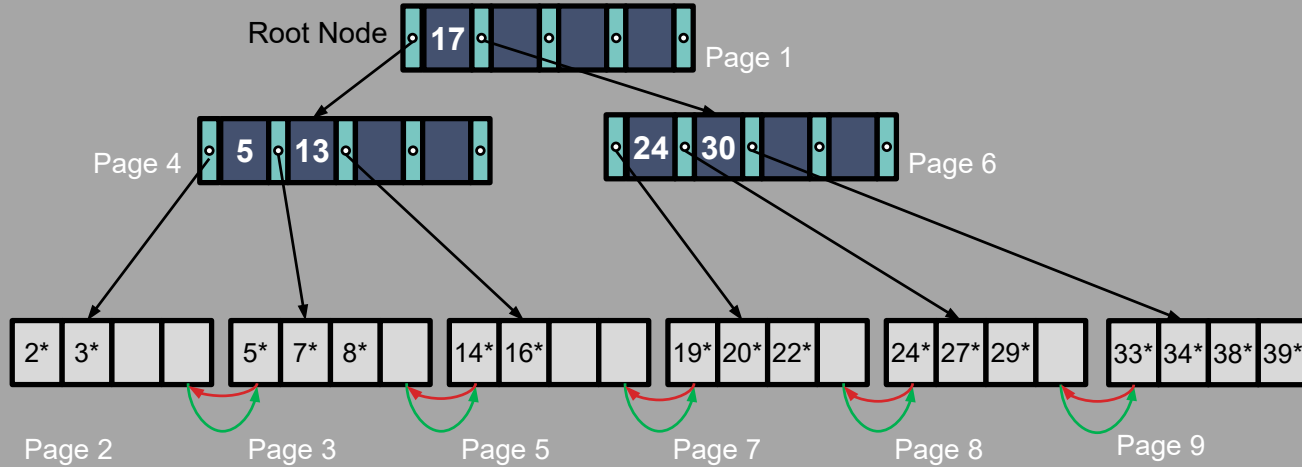


B+ Trees in Practice

- Typical order: 1600. Typical fill-factor: 67%.
 - average fan-out = 2144
 - (assuming 128 Kbytes pages at 40Bytes per record)
- At typical capacities
 - Height 1: $2144^2 = \mathbf{4,596,736 \text{ records}}$
 - Height 2: $2144^3 = \mathbf{9,855,401,984 \text{ records}}$




Searching the B+ Tree



- Same as ISAM
- Find key = 27
 - Find split on each node (Binary Search)
 - Follow pointer to next node



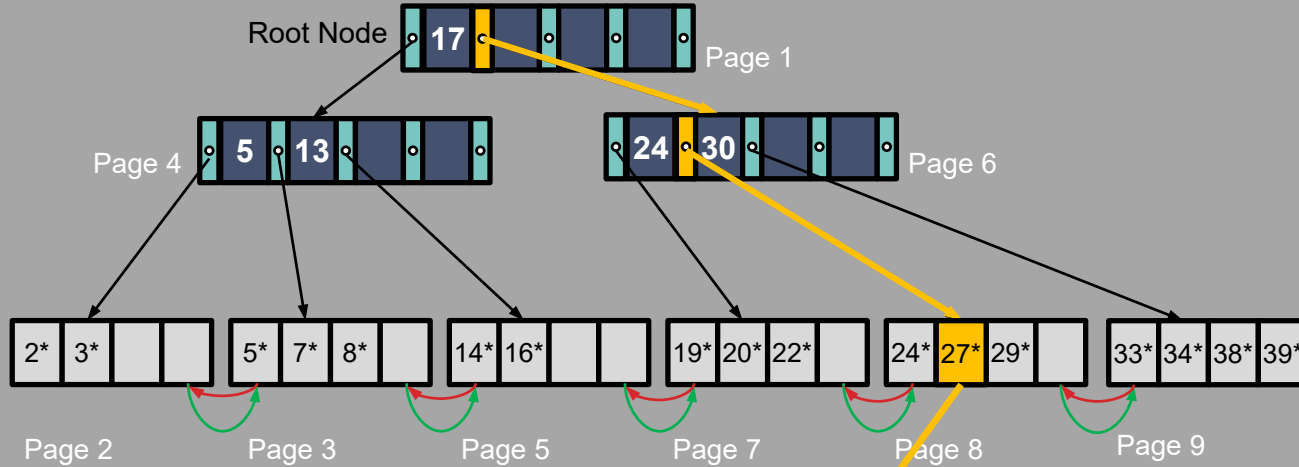
Searching the B+ Tree: Find 27



- Same as ISAM
- Find key = 27
 - Find split on each node (Binary Search)
 - Follow pointer to next node



Searching the B+ Tree: Fetch Data



PageId,
Record Id

Page 1

(20, Tim) (7, Dan)

Page 2

(5, Kay) (3, Jim)

Page 3


(27, Joe) (34, Kit)

Page 4

(1, Kim) (42, Hal)




Inserting 25* into a B+ Tree Part 1



- Find the correct leaf




Inserting 25* into a B+ Tree Part 2



- Find the correct leaf
- If there is room in the leaf just add the entry


Inserting 25* into a B+ Tree Part 3



- Find the correct leaf
- If there is room in the leaf just add the entry
 - Sort the leaf page by key




Inserting 8^* into a B+ Tree: Find Leaf



- Find the correct leaf




Inserting 8* into a B+ Tree: Insert



- Find the correct leaf
 - Split leaf if there is not enough room




Inserting 8* into a B+ Tree: Split Leaf



- Find the correct leaf
 - Split leaf if there is not enough room
 - Redistribute entries evenly




Inserting 8* into a B+ Tree: Split Leaf, cont



- Find the correct leaf
 - Split leaf if there is not enough room
 - Redistribute entries evenly
 - Fix next/prev pointers




Inserting 8* into a B+ Tree: Fix Pointers



- Find the correct leaf
 - Split leaf if there is not enough room
 - Redistribute entries evenly
 - Fix next/prev pointers




Inserting 8^* into a B+ Tree: Mid-Flight



- Something is still wrong!




Inserting 8^* into a B+ Tree: Copy Middle Key



- **Copy up from leaf** the middle key
- No room in parent? Recursively split index nodes




Inserting 8^* into a B+ Tree: Split Parent, Part 1



- **Copy up from leaf** the middle key
- No room in parent? Recursively split index nodes
 - Redistribute the rightmost d keys



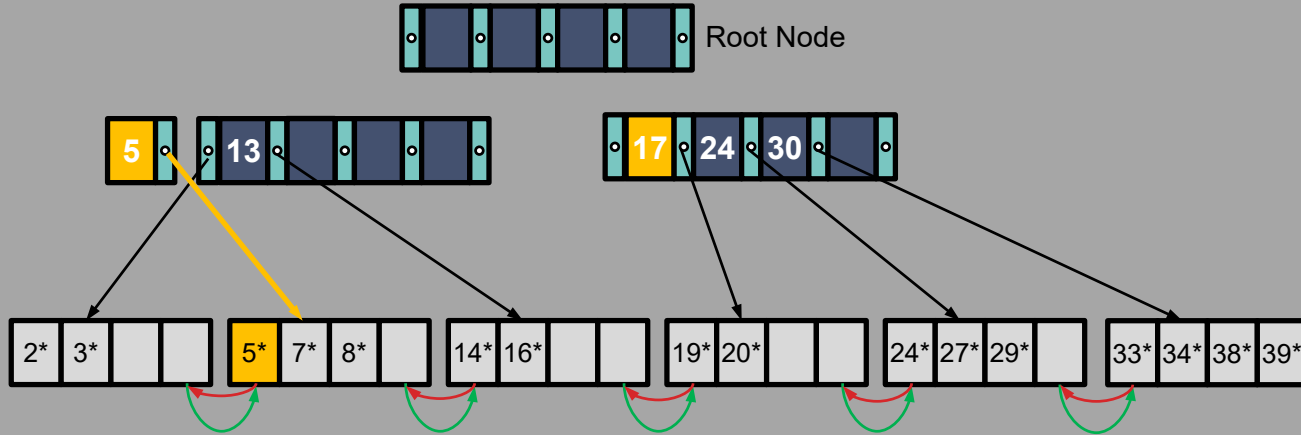
Inserting 8^* into a B+ Tree: Split Parent, Part 2



- **Copy up from leaf** the middle key
- No room in parent? Recursively split index nodes
 - Redistribute the rightmost d keys



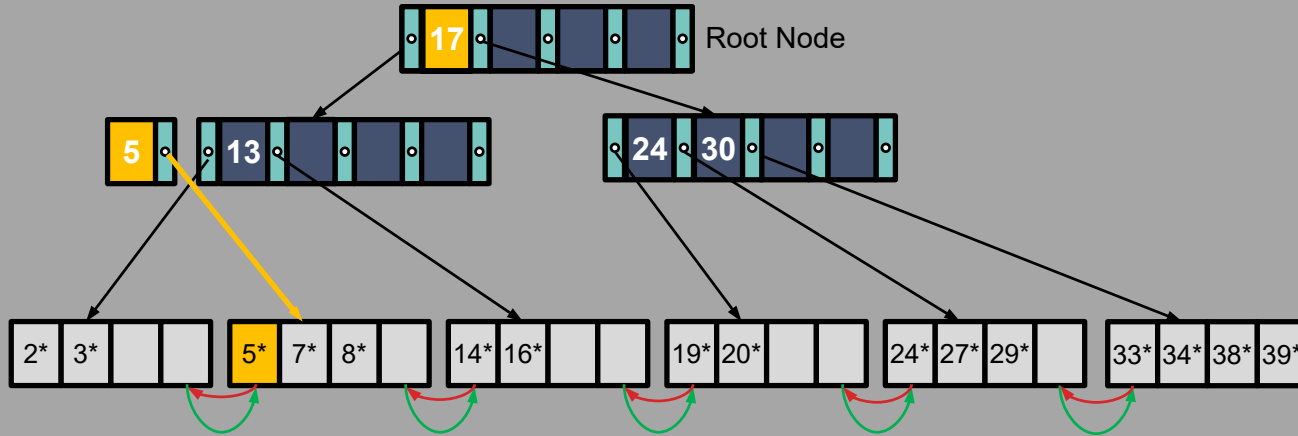
Inserting 8^* into a B+ Tree: Root Grows Up



- **Push up from interior node** the middle key
 - Now the last key on left
- No room in parent? Recursively split index nodes
 - Redistribute the rightmost d keys



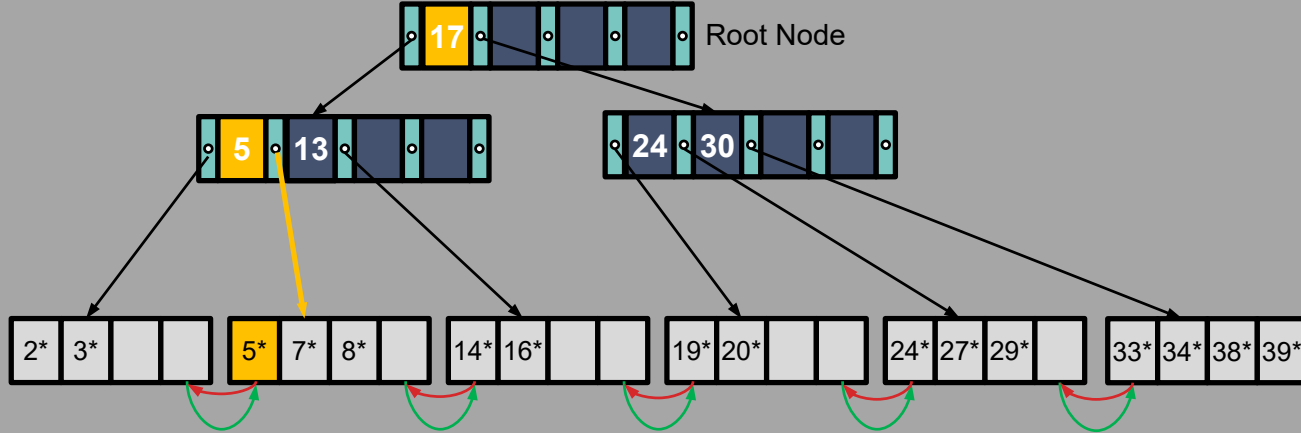
Inserting 8* into a B+ Tree: Root Grows Up, Pt 2



- Recursively split index nodes
 - Redistribute right d keys
 - **Push** up middle key



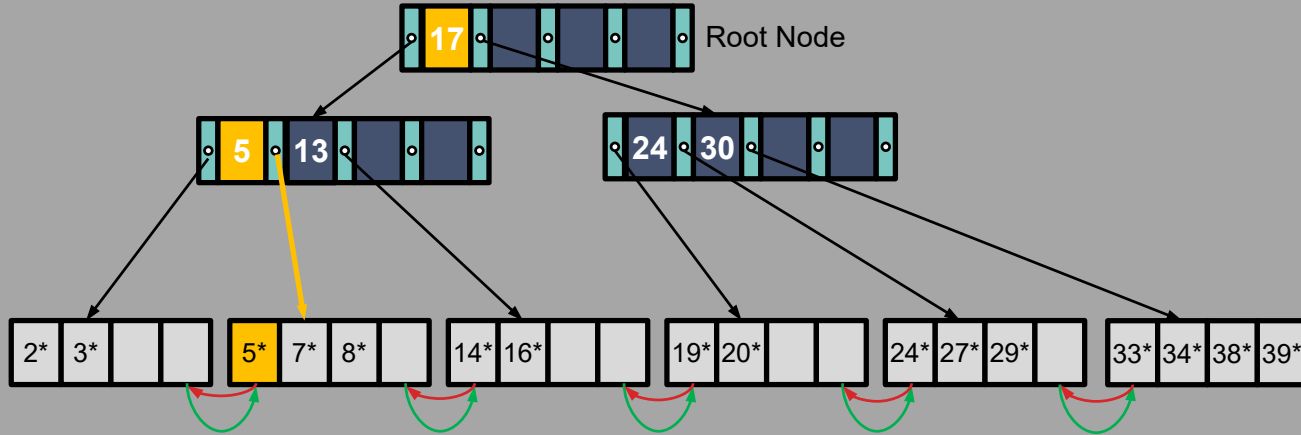
Inserting 8^* into a B+ Tree: Root Grows Up, Pt 3



- Recursively split index nodes
 - Redistribute right d keys
 - **Push** up middle key




Copy up vs Push up!



- Notice:
 - The **leaf** entry (5) was **copied** up
 - The **index** entry (17) was **pushed** up



Inserting 8^* into a B+ Tree: Final



- Check invariants
- **Key Invariant:**
 - $\text{Node}[\dots, (K_L, P_L), \dots] \rightarrow K_L \leq K$ for all K in P_L Sub-tree
- **Occupancy Invariant:**
 - $d \leq \# \text{ entries} \leq 2d$