



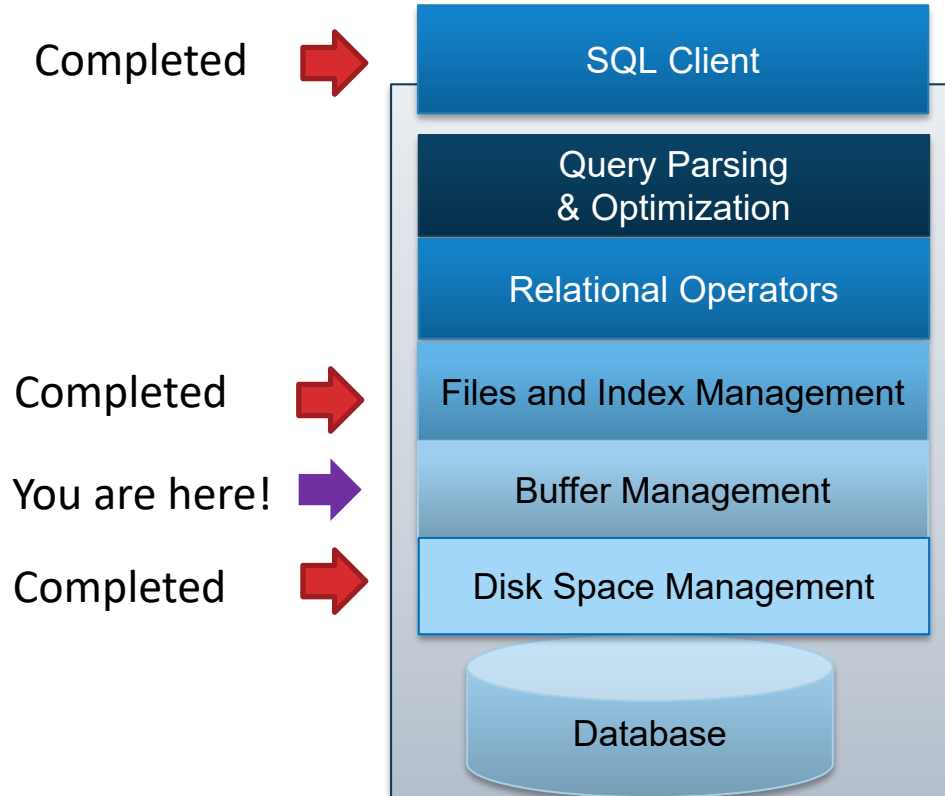
# Buffer Management

R & G - Chapter 9.4






# Architecture of a DBMS: What we've learned



# Lower Architecture of a DBMS





# Buffer Management Levels of Abstraction


Files and Index Management

RAM

Buffer Management

Disk

Disk Space Management






# Buffer Management, cont


RAM

Buffer Manager

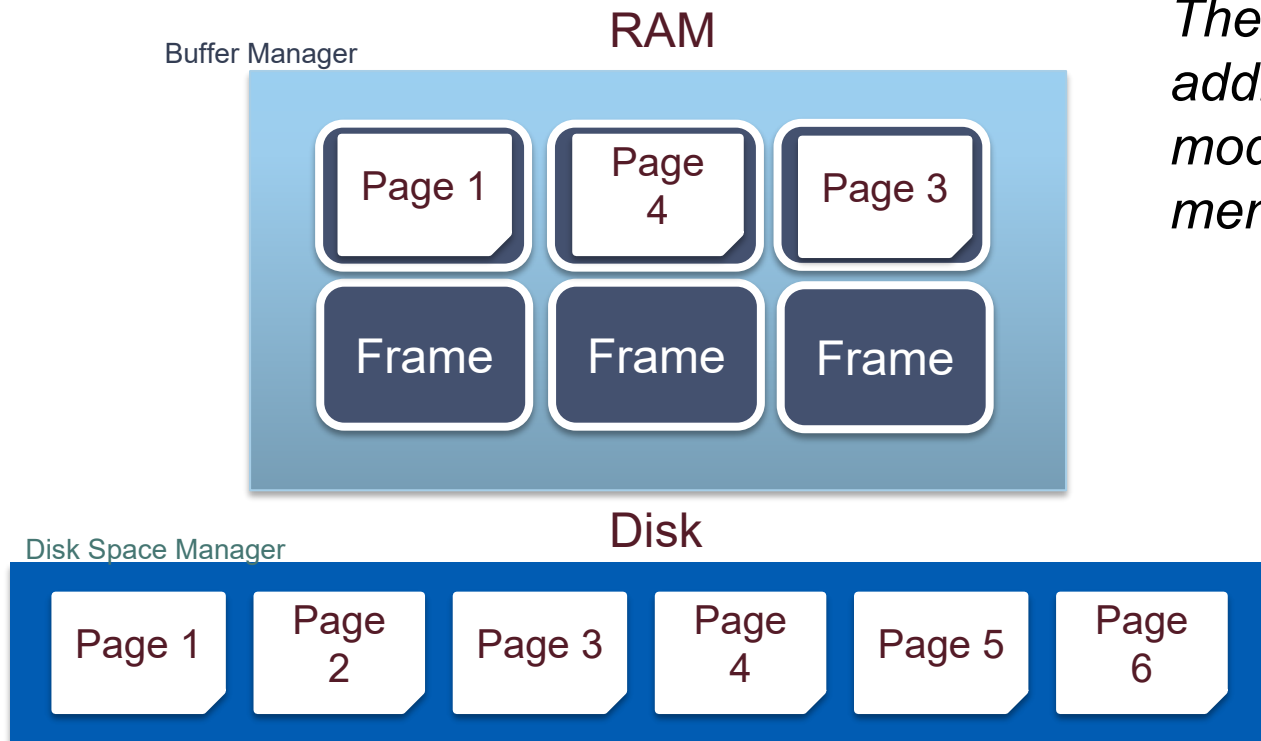


Disk

Disk Space Manager

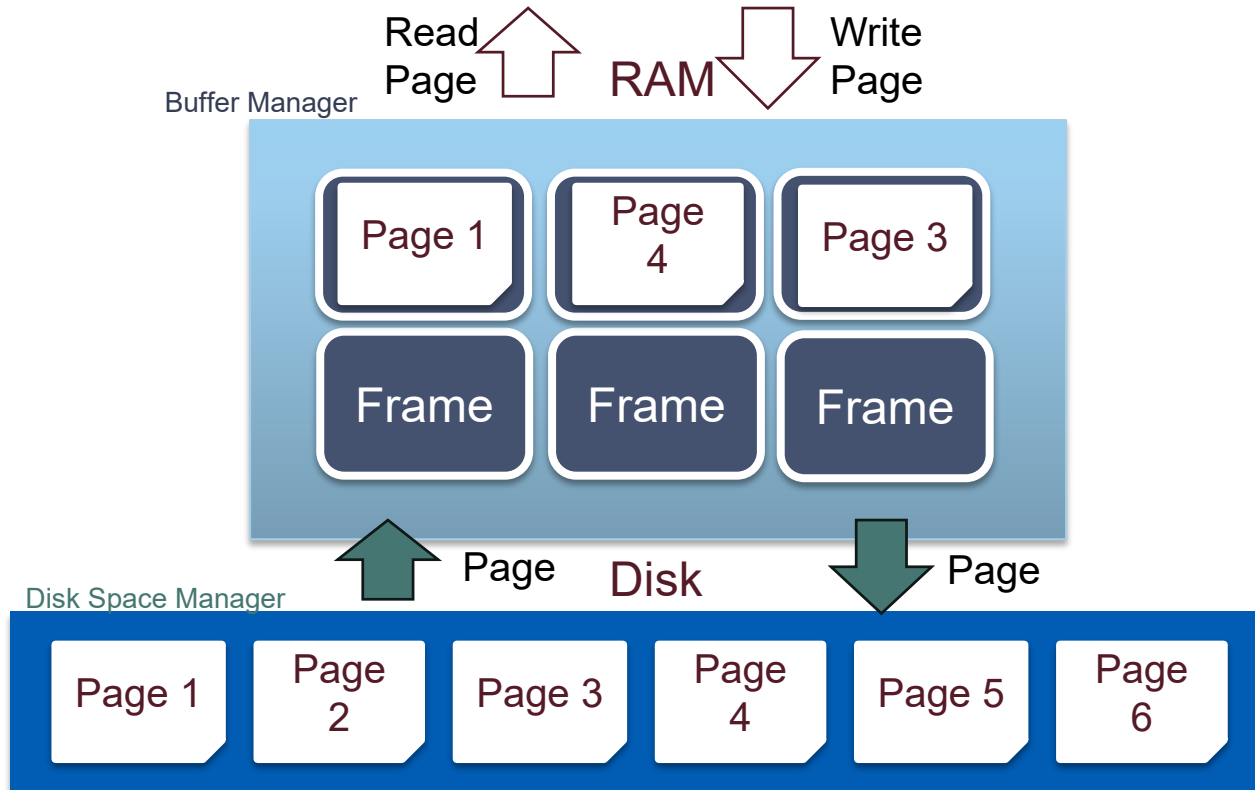


# Buffer Management Read

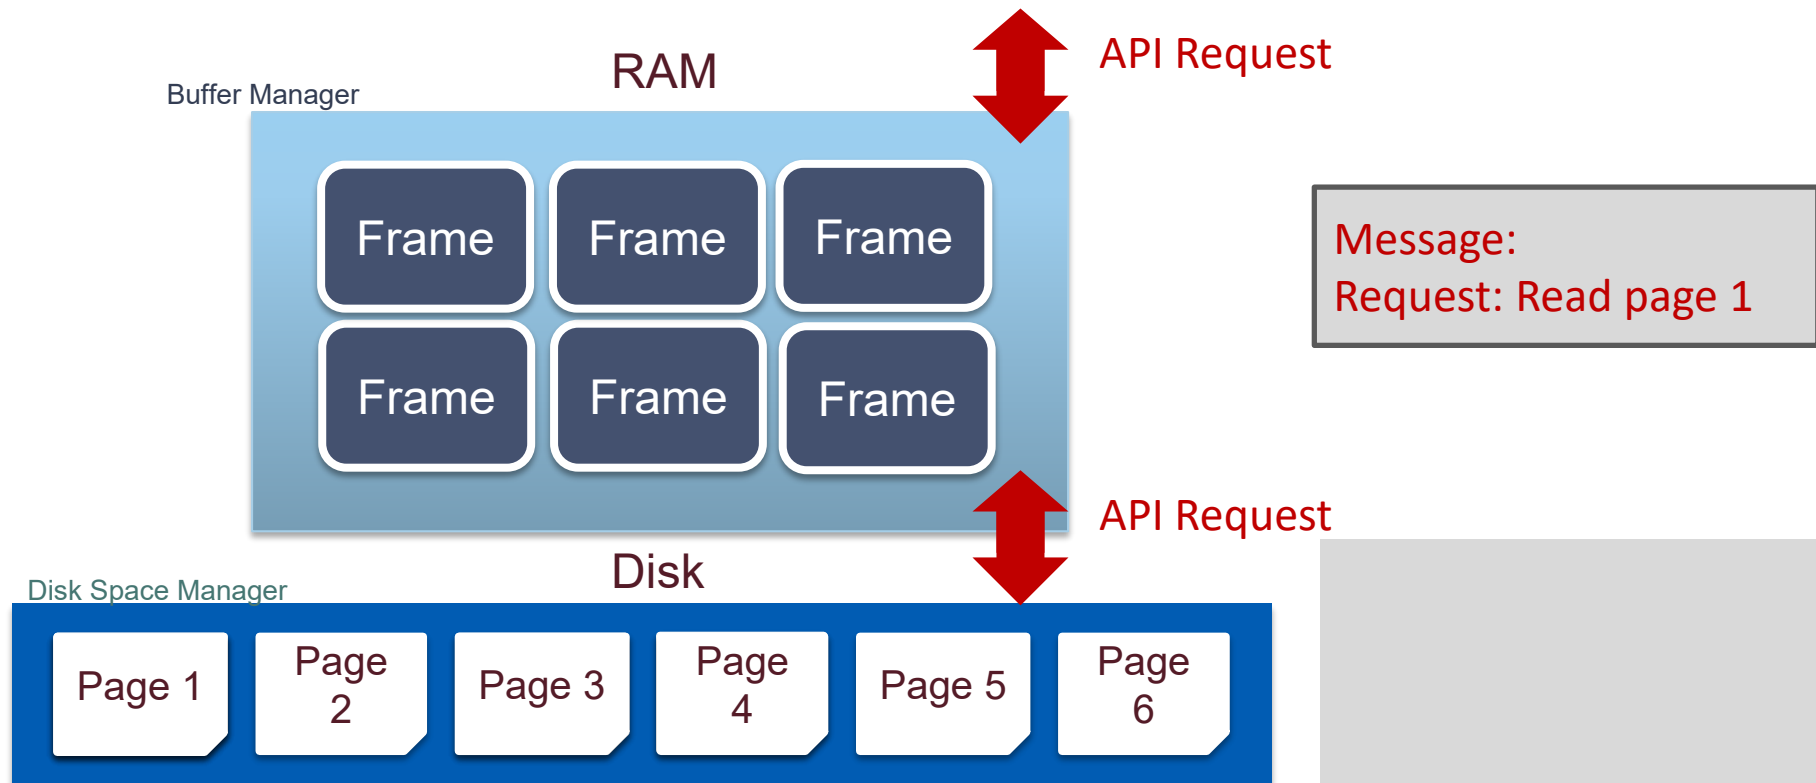


*The illusion of addressing and modifying disk pages in memory.*

# APIs

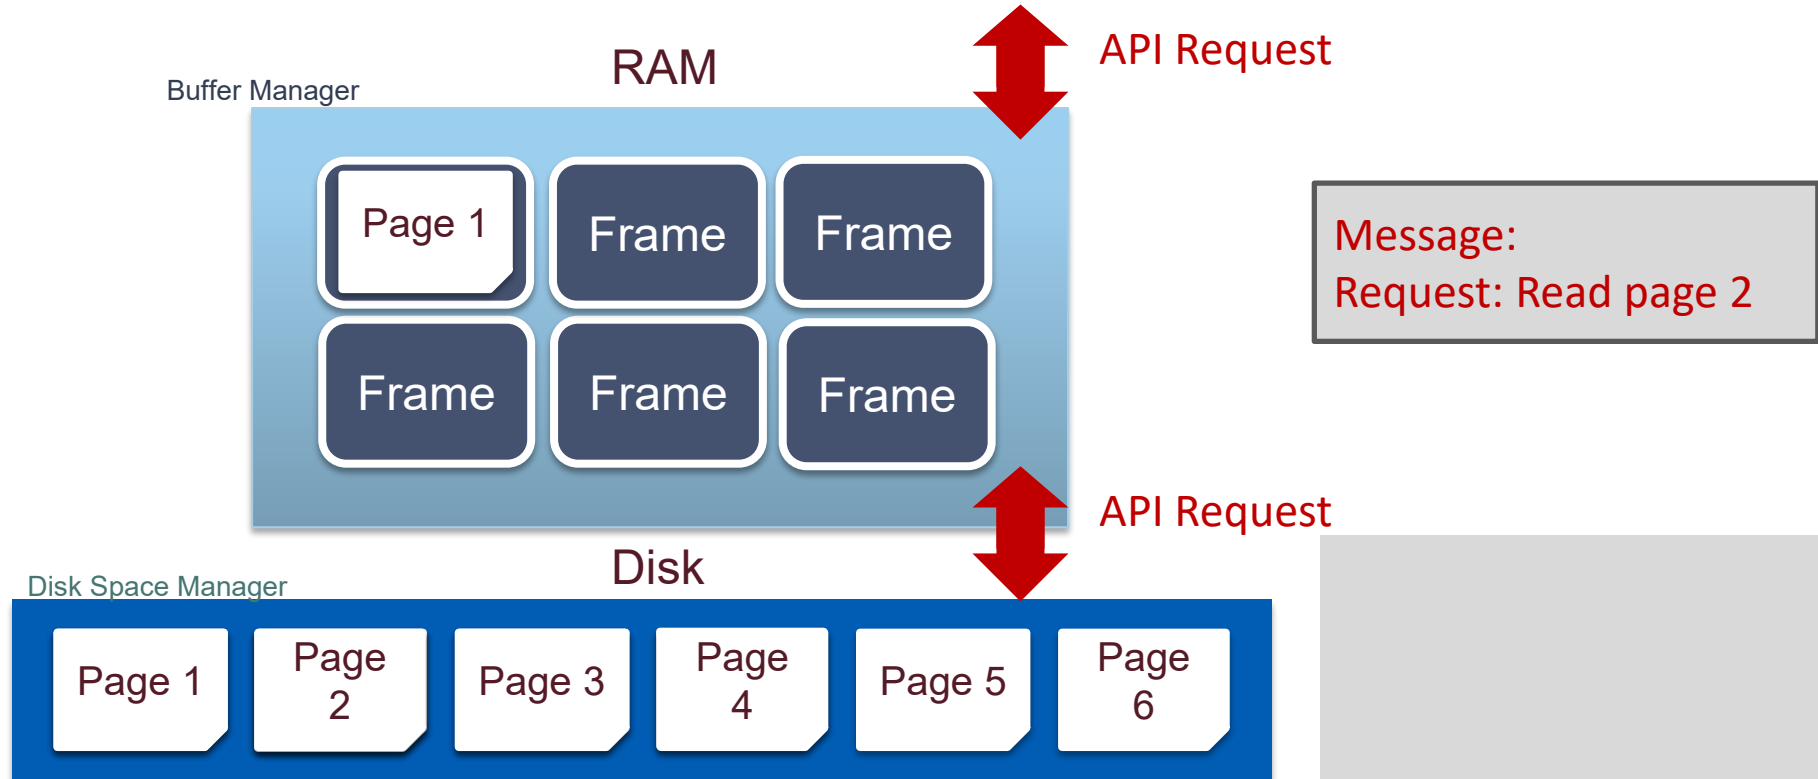


# Mapping Pages Into Memory, Pt 1

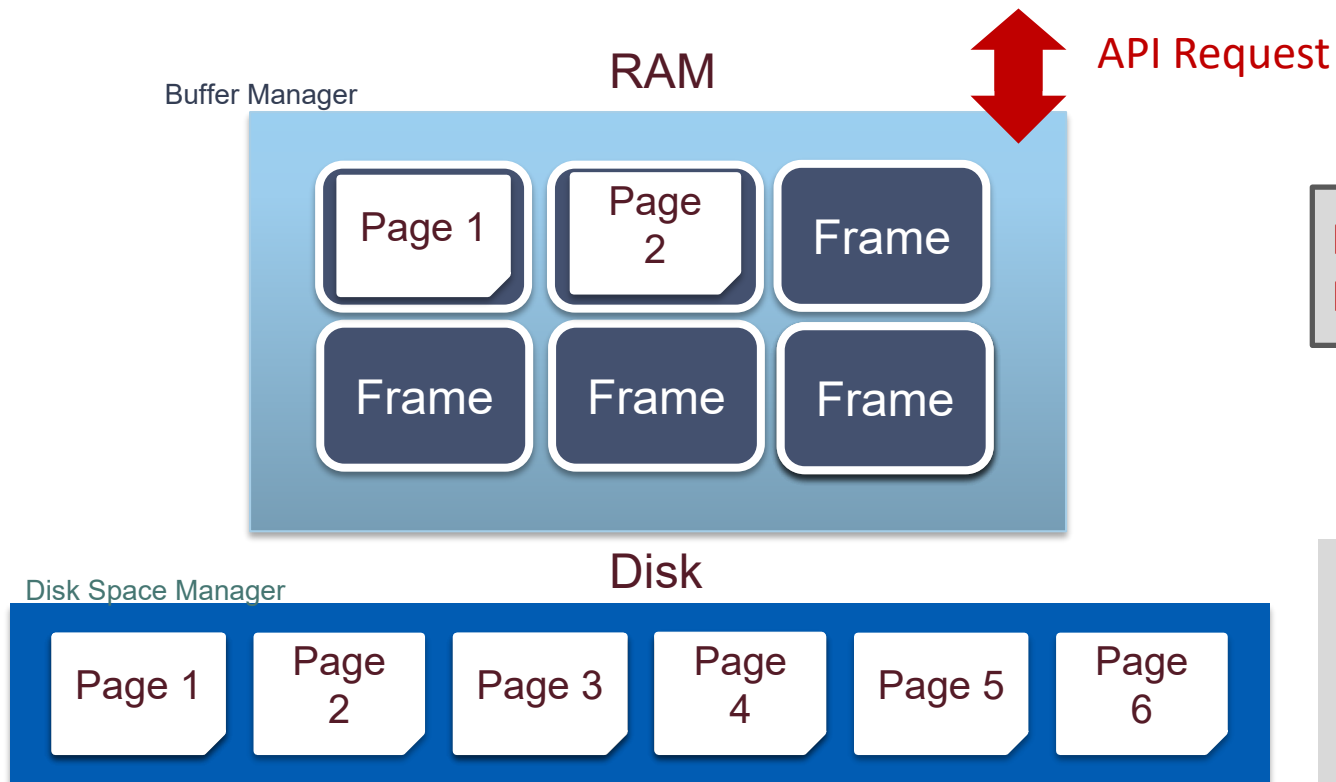




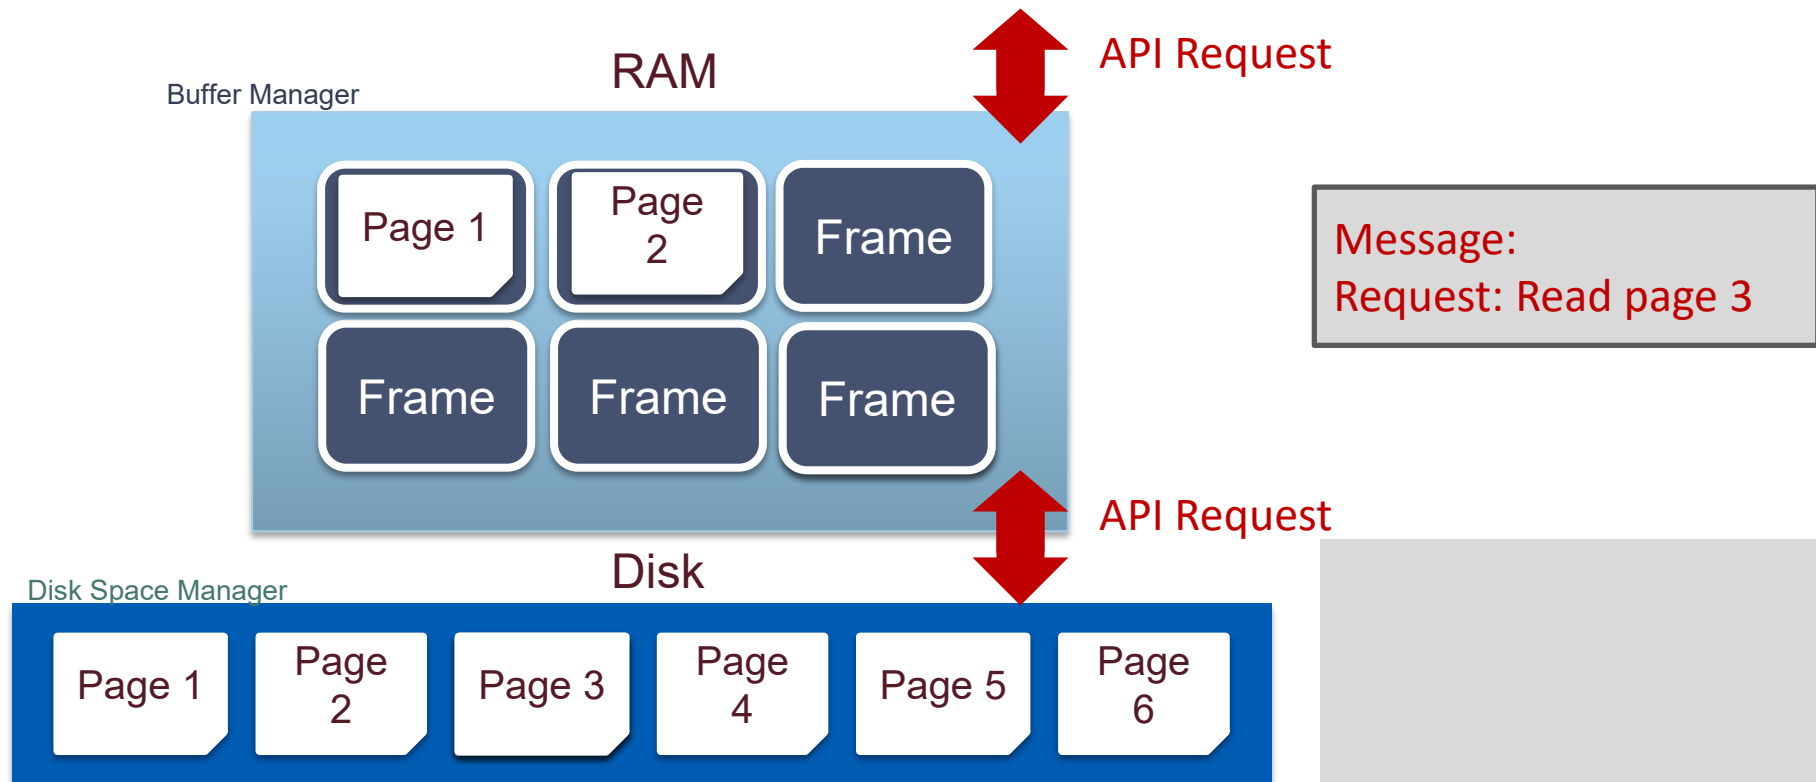
# Mapping Pages Into Memory, Pt 2



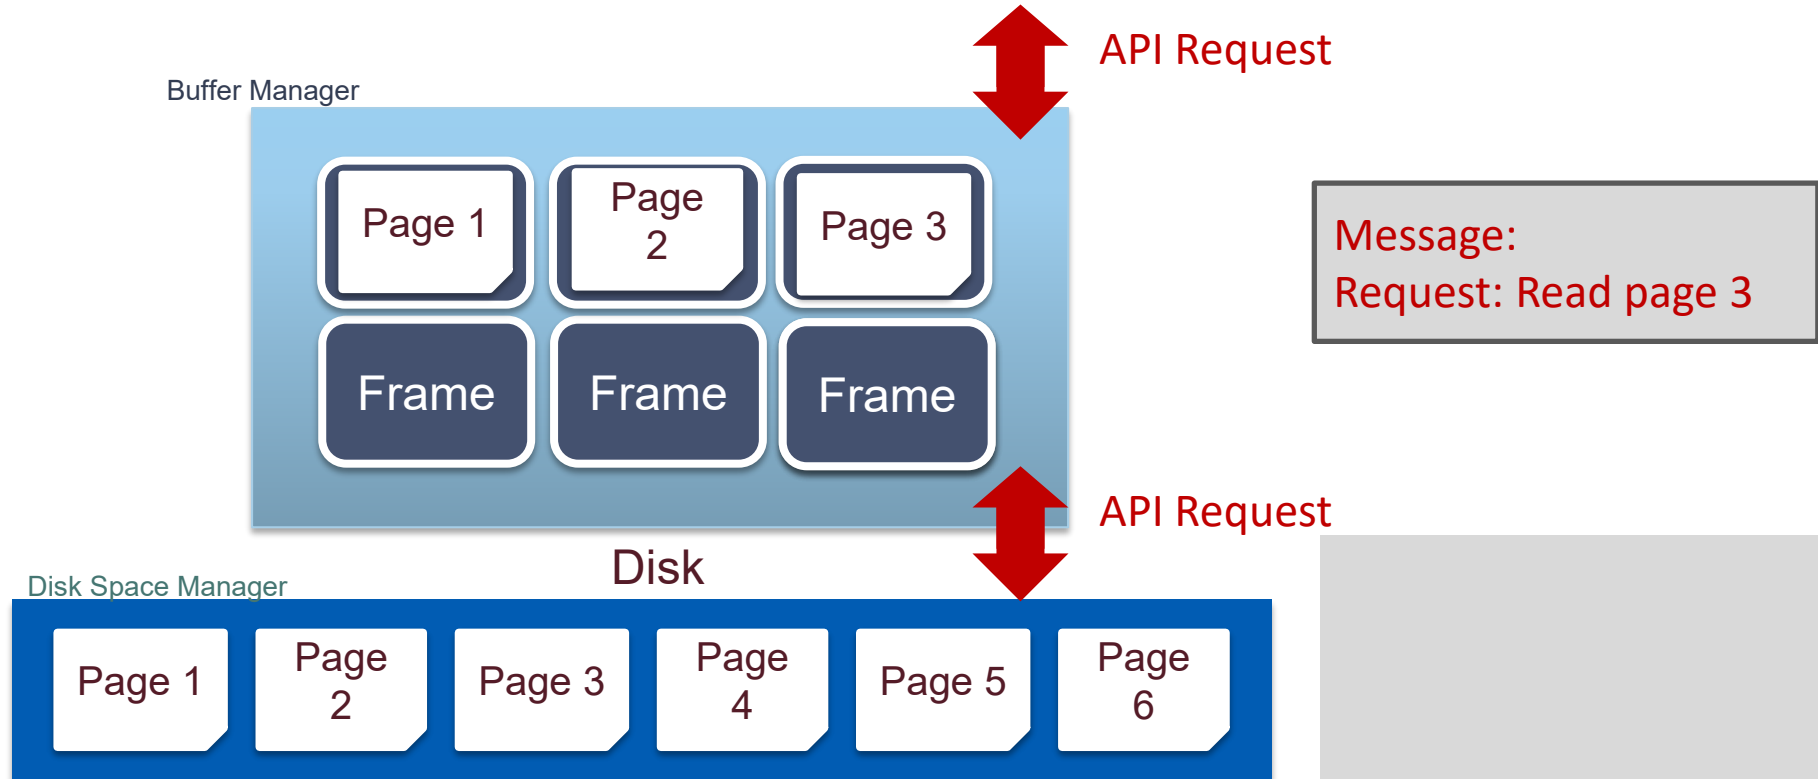
# Mapping Pages Into Memory, Pt 3



# Mapping Pages Into Memory, Pt 4



# Mapping Pages Into Memory







# Questions We Need to Answer

1. Handling dirty pages
2. Page Replacement



# Q1: Dirty Pages?



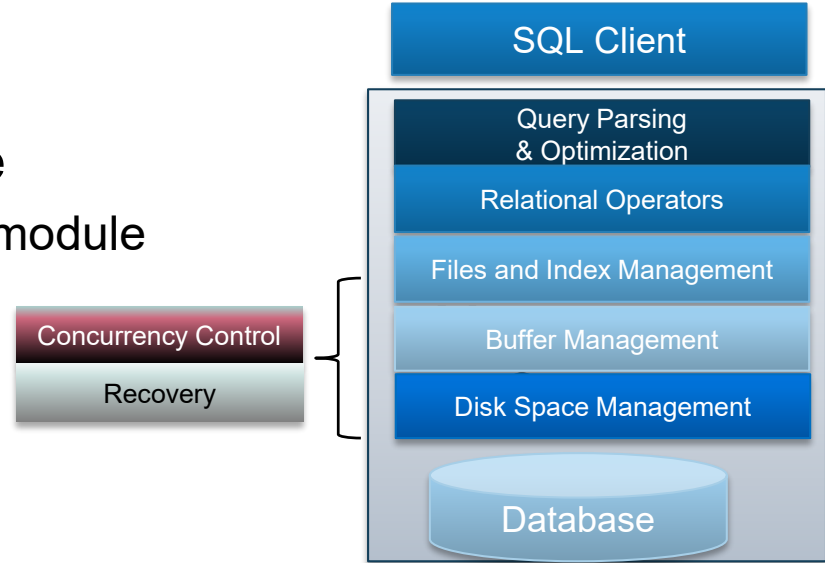


# Handling Dirty Pages

- Handling dirty pages
  - How will the buffer manager find out?
    - Dirty bit on page
  - What to do with a dirty page?
    - Write back via disk manager

# Advanced Questions


- Concurrent operations on a page
  - Solved by Concurrency Control module
- System Crash before write-back
  - Solved by Recovery module







# BufMgr State





# BufMgr State: Explicit

Buffer pool: Large range of memory, malloc'ed at DBMS server boot time (MBs-GBs)



Frameld	Pageld	Dirty?	Pin Count
1			
2			
3			
4			
5			
6			



# BufMgr State: Explicit Pt 2

Buffer pool: Large range of memory, malloc'ed at DBMS server boot time (MBs-GBs)



Buffer Manager metadata: Smallish array in memory, malloc'ed at DBMS server boot time

Frameld	Pageld	Dirty?	Pin Count
1	1	N	0
2	2	Y	1
3	3	N	0
4	6	N	2
5	4	N	0
6	5	N	0

Keep an in-memory index (hash table) on Pageld



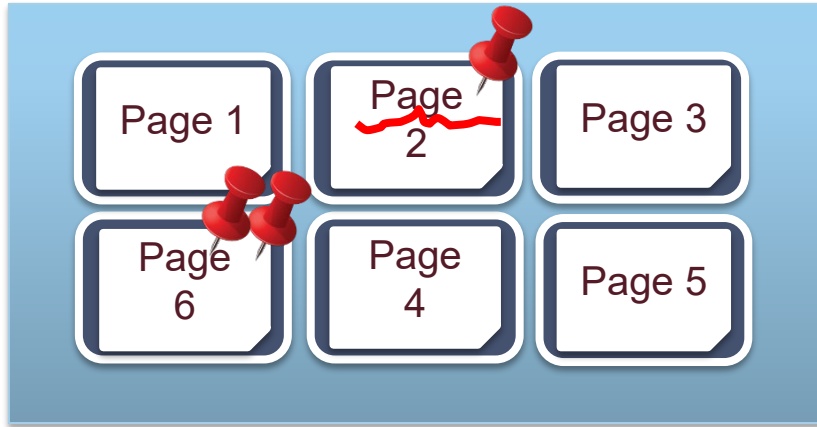
# BufMgr State: Illustrated



Frameld	Pageld	Dirty?	Pin Count
1	1	N	0
2	2	Y	1
3	3	N	0
4	6	N	2
5	4	N	0
6	5	N	0



# BufMgr State: Illustrated 2





# Page Replacement Terminology Review

- How will the buffer mgr know if a page is “in use”?
  - **Page pin count**
- If buffer manager is full, what page should be replaced?
  - **Page replacement policy**




# When a Page is Requested ...

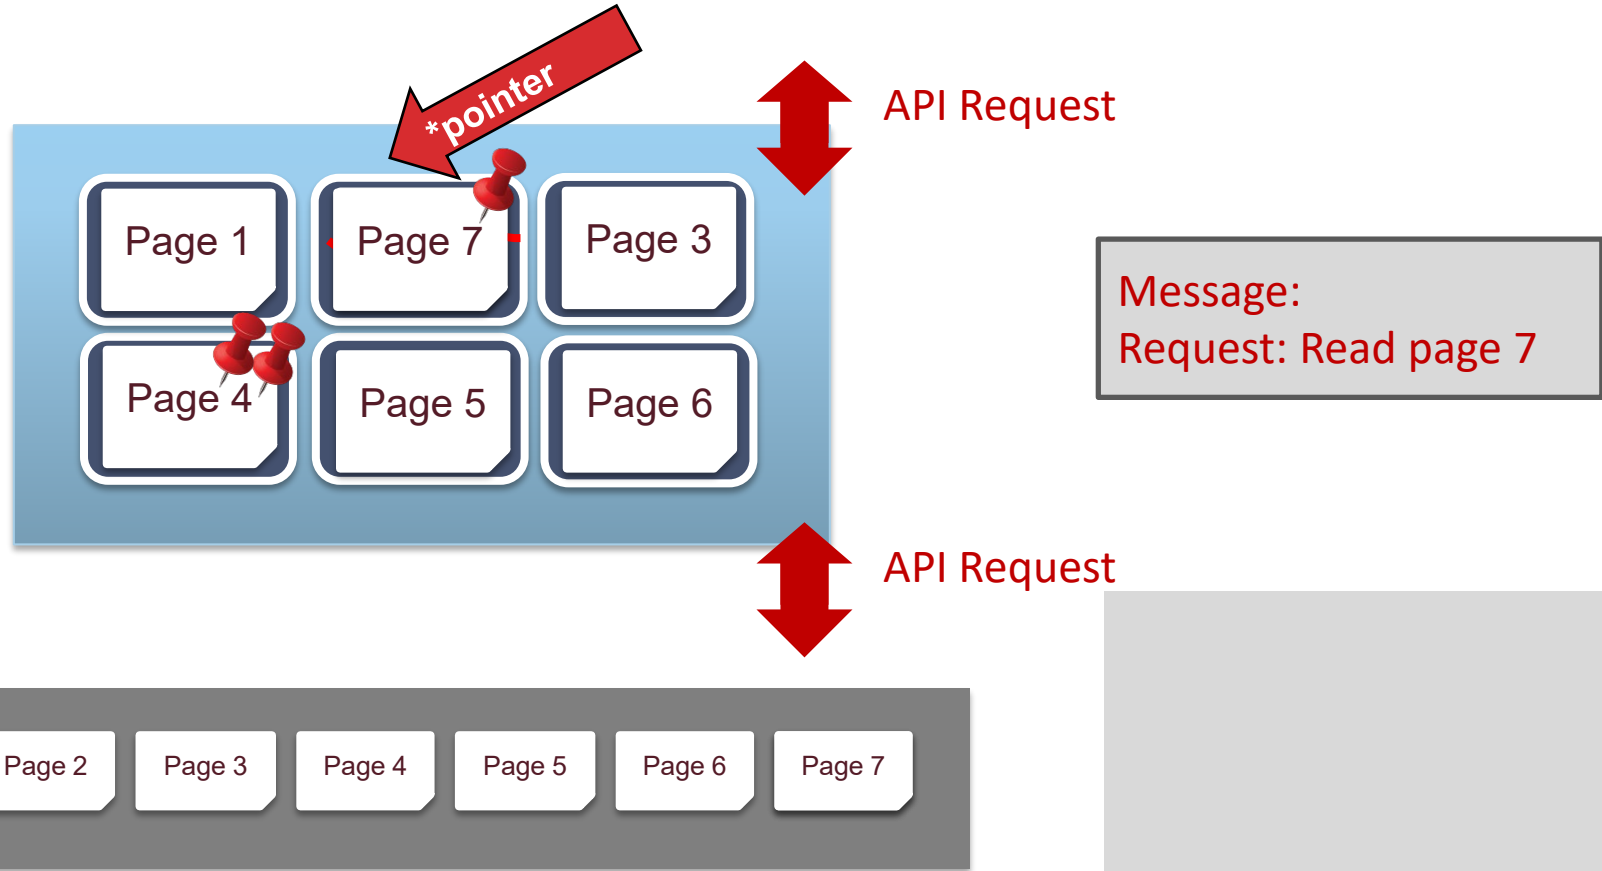
1. If requested page is not in pool:
  - a. Choose an **un-pinned** (`pin_count = 0`) frame for replacement.
  - b. If frame “dirty”, write current page to disk, mark “clean”
  - c. Read requested page into frame
2. Pin the page and return its address

If requests can be predicted (e.g., sequential scans) pages can be pre-fetched

- several pages at a time!



## Q2: Page Replacement







# After Requestor Finishes

1. Requestor of page must:
  - set dirty bit if page was modified
  - unpin the page (preferably soon!)
    - Why does requestor unpin?
    - What happens if they don't do it soon?
2. Page in pool may be requested many times
  - a pin count is used.
  - To pin a page: `pin_count++`
  - A page is a candidate for replacement iff
    - `pin_count == 0` ("unpinned")
3. CC & recovery may do additional I/Os upon replacement
  - Write Ahead Log protocol; more later!



# Answers to Our Previous Questions

## 1. Handling dirty pages

- How will the buffer manager find out?
  - Dirty bit on page
- What to do with a dirty page?
  - Write back via disk manager

## 2. Page Replacement

- How will the buffer mgr know if a page is “in use”?
  - Page pin count
- **If buffer manager is full, which page should be replaced?**
  - **Page replacement policy**



# Page Replacement Policy Intro

- Page is chosen for replacement by a **replacement policy**:
  - Least-recently-used (LRU), Clock
  - Most-recently-used (MRU)
- Policy can have big impact on #I/Os
  - Depends on the **access pattern**.



# LRU Replacement Policy

- Least Recently Used (LRU)
  - Pinned Frame: not available to replace
  - Track time each frame last unpinned (end of use)
  - Replace the frame which was least recently used

FrameId	PageId	Dirty?	Pin Count	Last Used
1	1	N	0	43
2	2	Y	1	21
3	3	N	0	22
4	6	N	2	11
5	4	N	0	24
6	5	N	0	15




# LRU Replacement Policy, Pt 2

- Very common policy: intuitive and simple
  - Good for repeated accesses to popular pages (temporal locality)
  - Can be costly. Why?
    - Need to “find min” on the last used attribute (priority heap data structure)
- Approximate LRU: CLOCK policy

FrameId	PageId	Dirty?	Pin Count	Last Used
1	1	N	0	43
2	2	Y	1	21
3	3	N	0	22
4	6	N	2	11
5	4	N	0	24
6	5	N	0	15

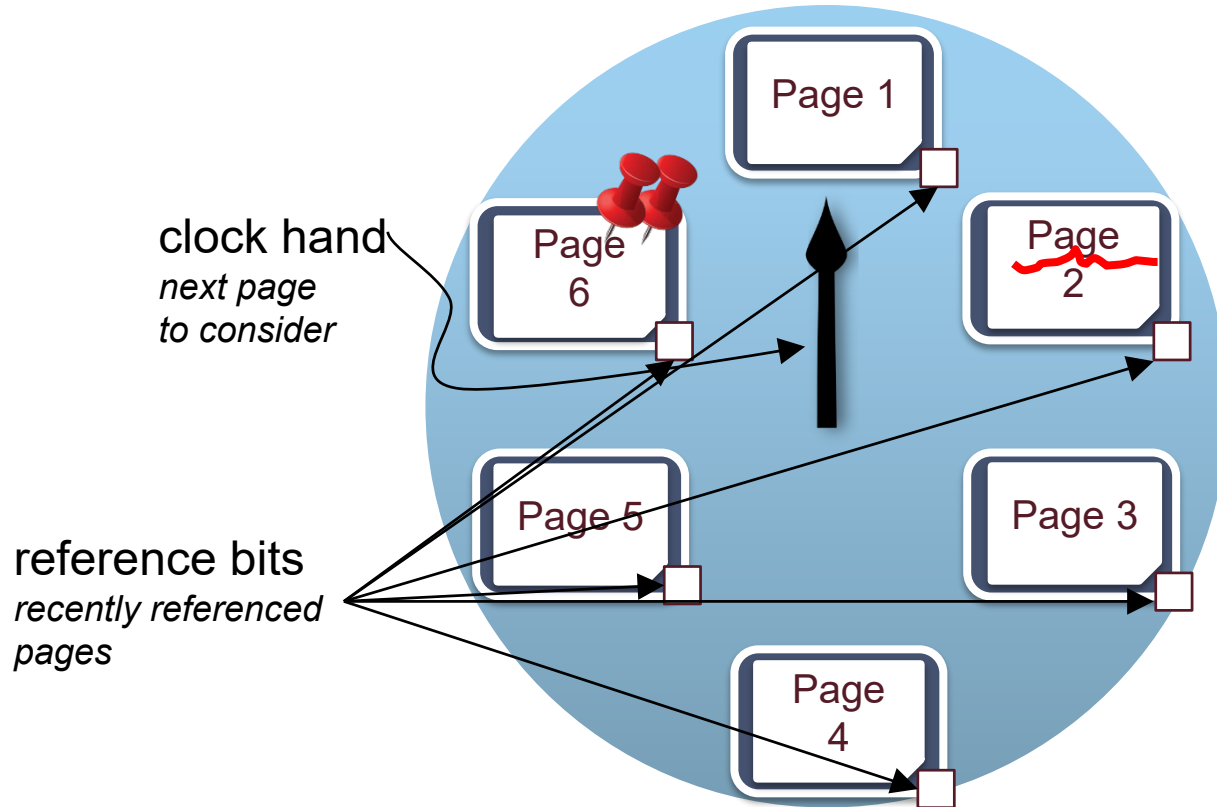


# BufMgr State: Illustrated





# Clock Policy State: Illustrated





# Clock Policy State: Explicit

FrameId	PageId	Dirty?	Pin Count	Ref Bit
1	1	N	1	1
2	2	N	1	1
3	3	N	0	1
4	4	N	0	0
5	5	N	0	0
6	6	N	0	1

Clock Hand

1

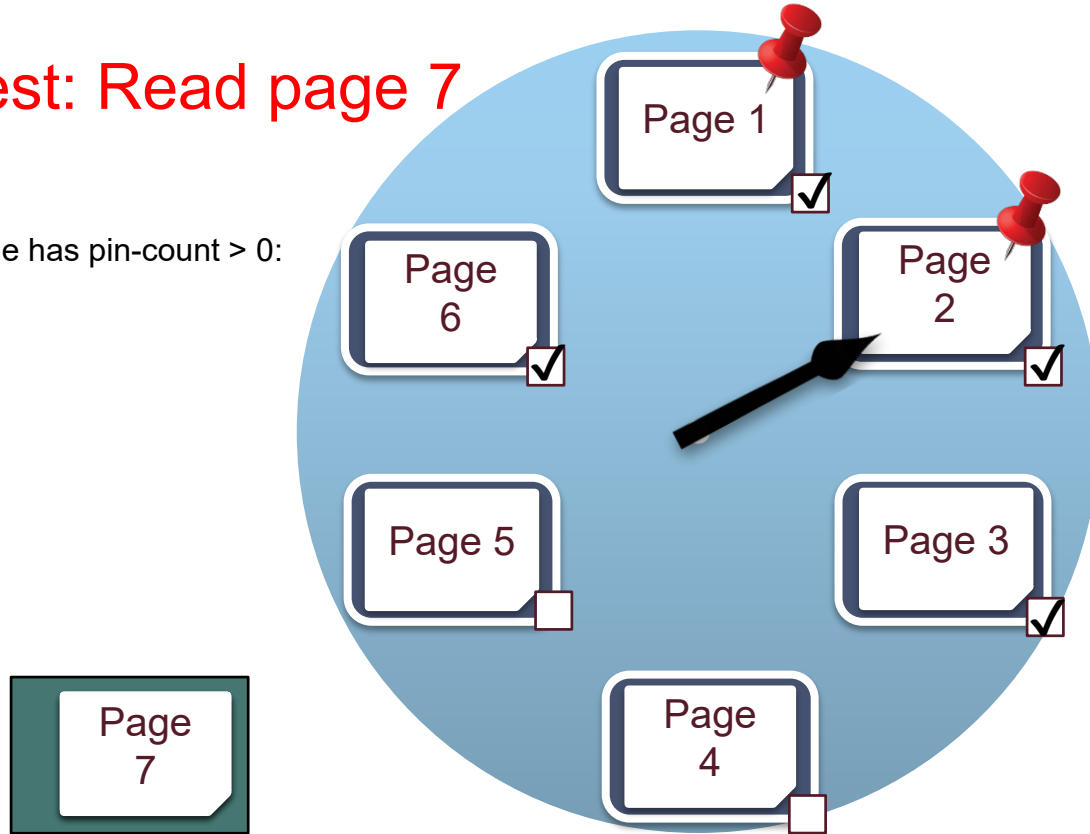




# Clock Policy State: Illustrated Part 1

Request: Read page 7

Current frame has pin-count > 0:  
**Skip**






# Clock Policy State: Illustrated, Part 2

Request: Read page 7

Current frame not pinned,  
Ref bit set:

**Clear ref bit**  
**Skip**







# Clock Policy State: Illustrated, Pt 3

Request: Read page 7

Current frame not pinned  
Ref bit unset:

**Replace**



# Clock Policy State: Illustrated, Pt 4

Request: Read page 7


Current frame not pinned  
Ref bit unset:

**Replace**

**Set pinned**

**Set ref bit**

**Advance clock**



# Clock Policy State: Illustrated, Pt 5

Request: Read page 7

Current frame not pinned  
Ref bit unset:


**Replace**

**Set pinned**

**Set ref bit**

**Advance clock**

**Return pointer**





# Clock Policy Pseudocode

```
1  page *clock_request_page(int &clk_hand, int pg_num) {  
2      retval = NULL;  
3      while (retval == NULL) {  
4          current = frame_table[clk_hand];  
5          // the happy case: replace current page  
6          if (current.pin_count == 0 && current.refbit == 0) {  
7              if (current.dirty == 1)  
8                  write_page(fi.page, frames[clk_hand]);  
9              read_page(pg_num, frames[clk_hand]);  
10             retval = frames[clk_hand];  
11             current.dirty = 0;  
12             current.pin_count = 1;  
13             current.refbit = 1; // referenced!  
14         }  
15         // second chance: unset reference bit  
16         else if (current.pin_count == 0 && current.refbit == 1) {  
17             current.refbit = 0;  
18         }  
19         // else pin_count > 1, so skip  
20  
21         clk_hand += (clk_hand + 1) % MAX_FRAME; // advance clock hand  
22     }  
23     return retval;  
24 }
```



# Clock Policy Pseudocode, Pt 2

```
1  page *clock_request_page(int &clk_hand, int pg_num) {
2      retval = NULL;
3      while (retval == NULL) {
4          current = frame_table[clk_hand];
5          // the happy case: replace current page
6          if (current.pin_count == 0 && current.refbit == 0) {
7              if (current.dirty == 1)
8                  write_page(fi.page, frames[clk_hand]);
9              read_page(pg_num, frames[clk_hand]);
10             retval = frames[clk_hand];
11             current.dirty = 0;
12             current.pin_count = 1;
13             current.refbit = 1; // referenced!
14         }
15         // second chance: unset reference bit
16         else if (current.pin_count == 0 && current.refbit == 1) {
17             current.refbit = 0;
18         }
19         // else pin_count > 1, so skip
20
21         clk_hand += (clk_hand + 1) % MAX_FRAME; // advance clock hand
22     }
23     return retval;
24 }
```

# Clock Policy Pseudocode, Pt 3

```
1  page *clock_request_page(int &clk_hand, int pg_num) {
2      retval = NULL;
3      while (retval == NULL) {
4          current = frame_table[clk_hand];
5          // the happy case: replace current page
6          if (current.pin_count == 0 && current.refbit == 0) {
7              if (current.dirty == 1)
8                  write_page(fi.page, frames[clk_hand]);
9              read_page(pg_num, frames[clk_hand]);
10             retval = frames[clk_hand];
11             current.dirty = 0;
12             current.pin_count = 1;
13             current.refbit = 1; // referenced!
14         }
15         // second chance: unset reference bit
16         else if (current.pin_count == 0 && current.refbit == 1) {
17             current.refbit = 0;
18         }
19         // else pin_count > 1, so skip
20
21         clk_hand += (clk_hand + 1) % MAX_FRAME; // advance clock hand
22     }
23     return retval;
24 }
```






# Is LRU/Clock Always Best?

- Very common policy: intuitive and simple
- Works well for repeated accesses to popular pages
  - Temporal locality
- LRU can be costly → Clock policy is cheap
  - Quite similar
  - If you like, try to find cases where they differ.
- When might they perform poorly
  - What about repeated scans of big files?



# Repeated Scan (LRU)

- Cache Hits: 0
- Attempts: 0




Disk Space Manager






# Repeated Scan (LRU): Read Page 1

- Cache Hits: 0
- Attempts: 1




Disk Space Manager






# Repeated Scan (LRU): Read Page 2

- Cache Hits: 0
- Attempts: 2



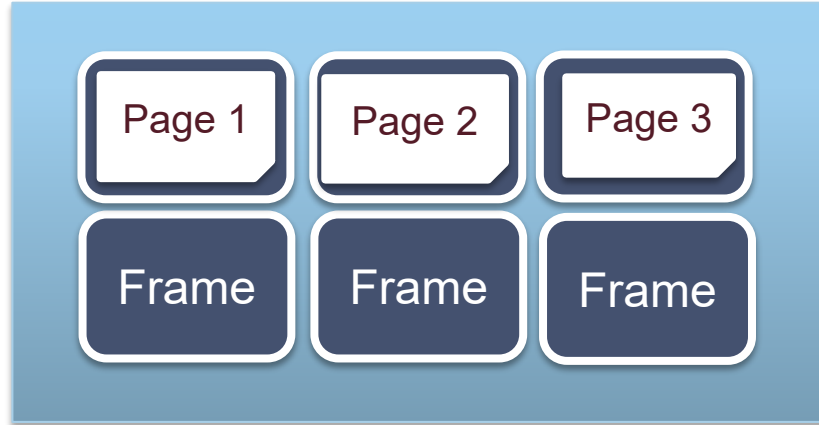
Disk Space Manager






# Repeated Scan (LRU): Read Page 3

- Cache Hits: 0
- Attempts 3:




Disk Space Manager






# Repeated Scan (LRU): Read Page 4

- Cache Hits 0:
- Attempts: 4




Disk Space Manager






# Repeated Scan (LRU): Read Page 5

- Cache Hits: 0
- Attempts: 5




Disk Space Manager





# Repeated Scan (LRU): Read Page 6

- Cache Hits: 0
- Attempts 6



Disk Space Manager









# Repeated Scan (LRU): Read Page 7

- Cache Hits: 0
- Attempts: 7




Disk Space Manager






# Repeated Scan (LRU): Reset to beginning

- Cache Hits: 0
- Attempts: 7




Disk Space Manager





# Repeated Scan (LRU): Read Page 1 (again)

- Cache Hits: 0
- Attempts: 8




Disk Space Manager






# Repeated Scan (LRU): Read Page 2 (again)

- Cache Hits: 0
- Attempts: 9




Disk Space Manager





# Repeated Scan (LRU): Read Page 3 (again)

- Cache Hits: 0
- Attempts: 10




Disk Space Manager





# Repeated Scan (LRU): Page 4 (again)

- Cache Hits: 0
- Attempts: 11




Disk Space Manager






# Repeated Scan (LRU): Read Page 5, cont

- Cache Hits: 0
- Attempts: 12



Disk Space Manager





# Sequential Scan + LRU


- Sequential flooding
- 0% hit rate in cache!
- Repeated sequential scan very common in database workloads
  - We will see it in nested-loops join
- What could be better?





# Repeated Scan (MRU)

- Cache Hits: 0
- Attempts: 6




Disk Space Manager






# Repeated Scan (MRU): Read Page 7

- Cache Hits: 0
- Attempts: 7




Disk Space Manager





# Repeated Scan (MRU): Reset

- Cache Hits: 0
- Attempts: 7




Disk Space Manager





# Repeated Scan (MRU): Read Page 1 (again)

- Cache Hits: 1
- Attempts: 8




Disk Space Manager





# Repeated Scan (MRU): Read Page 2 (again)

- Cache Hits: 2
- Attempts: 9




Disk Space Manager





# Repeated Scan (MRU): Read Page 3 (again)

- Cache Hits: 3
- Attempts: 10




Disk Space Manager





# Repeated Scan (MRU): Read Page 4 (again)

- Cache Hits: 4
- Attempts: 11




Disk Space Manager





# Repeated Scan (MRU): Read Page 5 (again)

- Cache Hits: 5
- Attempts: 12



Disk Space Manager








# Repeated Scan (MRU): Read Page 6 (again)

- Cache Hits: 5
- Attempts: 13





Disk Space Manager





# Repeated Scan (MRU): Read Page 7 (again)

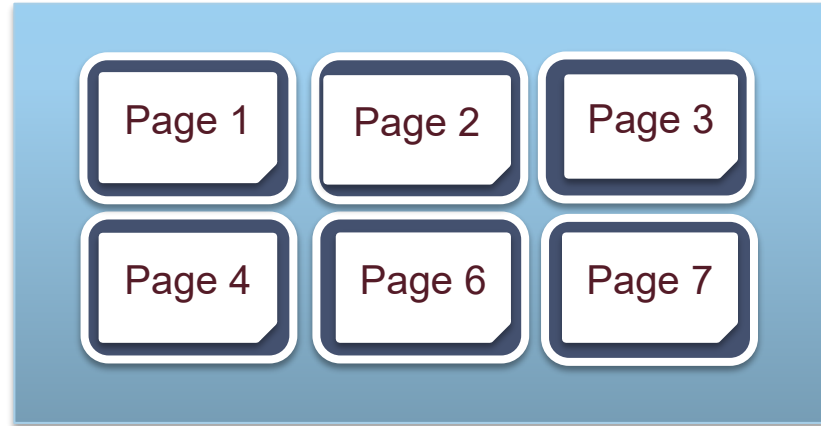
- Cache Hits: 6
- Attempts: 14





# Repeated Scan (MRU): Reset (again)



- Cache Hits: 6
- Attempts: 14





# Repeated Scan (MRU): Read Page 1 (again x2)


- Cache Hits: 7
- Attempts: 15





# Repeated Scan (MRU): Read Page 2 (again x2)


- Cache Hits: 8
- Attempts: 16





# Repeated Scan (MRU): Read Page 3 (again x2)


- Cache Hits: 9
- Attempts: 17





# Repeated Scan (MRU): Read Page 4 (again x2)


- Cache Hits: 10
- Attempts: 18





# Repeated Scan (MRU): Read Page 5 (again x2)

- Cache Hits: 10
- Attempts: 19







# General Case: SeqScan + MRU

B buffers

$N > B$  pages in file

1. First N attempts: 0 hits
2. Next N attempts: B-1 hits  
Pages 1 through B-1
3. Next N attempts: B-1 hits  
Pages N through B-2
4. Next N attempts: B-1 hits  
Pages N-1 through B-3
- ...

In limit:  $\sim(B-1)/N$  hit rate



# Improvement for sequential scan: prefetch

- Prefetch: Ask disk space manager for a run of sequential pages
  - E.g. On request for Page 1, ask for Pages 2-5
- Why does this help?
  - Amortize random I/O overhead
  - Allow computation while I/O continues in background
    - Disk and CPU are “parallel devices”



# We seem to need a hybrid!

- LRU wins for random access (hot vs. cold)
  - When might we see that behavior?
- MRU wins for repeated sequential
  - E.g. for certain joins



# Two General Approaches

- Use DBMS information to hint to BufMgr
  - For big queries: we can predict I/O patterns from the handful of query processing algorithms we'll learn shortly
  - For simple lookups: LRU often does well
- Find fancier stochastic policies
  - E.g. 2Q, LRU-2, ARC.
  - See [Page Replacement Algorithm](#) on Wikipedia but beware the OS-centric history
- Hybrids are not uncommon in modern DBMSs
  - E.g. special-case for indexes, use LRU-2 otherwise
  - FWIW, PostgreSQL currently uses CLOCK
  - Imagine workloads for a big cloud DBMS like AWS Aurora!



# DBMS vs OS Buffer Cache

- Doesn't the filesystem (OS) manage buffers and pages too?
- Issues:
  - Portability: different FS, different behavior
  - OS limitations: DBMS requires ability to force pages to disk
    - Required for recovery, as we'll see
  - OS limitations: DBMS can predict its own page reference patterns
    - E.g. consider scanning the leaves of a B+-tree
    - Affects both page replacement and prefetching



# Summing Up

- Buffer Manager provides a level of indirection
  - Maps disk page Ids to RAM addresses
- Ensures that each requested page is “pinned” in RAM
  - To be (briefly) manipulated in-memory
  - And then unpinned by the caller!
- Attempts to minimize “cache misses”
  - By replacing pages unlikely to be referenced
  - By prefetching pages likely to be referenced



# Make Sure You Know

- Pin Counts and Dirty Bits:
  - When do they get set/unset?
  - By what layer of the system?
- LRU, MRU and Clock
  - Be able to run each by hand
  - For Clock:
    - What pages are eligible for replacement
    - When is reference bit set/unset
    - What is the point of the reference bit?
- Sequential flooding
  - And how it behaves for LRU (Clock), MRU