

# Iterators, Relational Operators and Joins

R&G Chapters 12 & 14





# Recall from Last Lecture

## SQL Query

```
SELECT S.name
FROM Reserves R, Sailors S
WHERE R.sid = S.sid
AND R.bid = 100
AND S.rating > 5
```


Query Parser  
& Optimizer

## Relational Algebra

$$\pi_{S.name}(\sigma_{bid=100 \wedge rating > 5}(\text{Reserves} \bowtie_{R.sid=S.sid} \text{Sailors}))$$


Equivalent to...

## (Logical) Query Plan:



But actually will  
produce...

## Optimized (Physical) Query Plan:



## Operator Code

B+-Tree  
Indexed Scan  
Iterator

On-the-fly  
Project Iterator

On-the-fly  
Select Iterator

Indexed Nested  
Loop Join Iterator


Heap Scan  
Iterator



# Relational Operators and Query Plans

$$\pi_{\text{sname}}(\pi_{\text{sid}}(\pi_{\text{bid}}(\sigma_{\text{color}=\text{'red'}}(\text{Boats})) \bowtie \text{Res}) \bowtie \text{Sailors})$$

- Query plan
  - Edges encode “flow” of tuples
  - Vertices = Relational Alg Operators
  - Source vertices = table access operators ...
- Also called dataflow graph

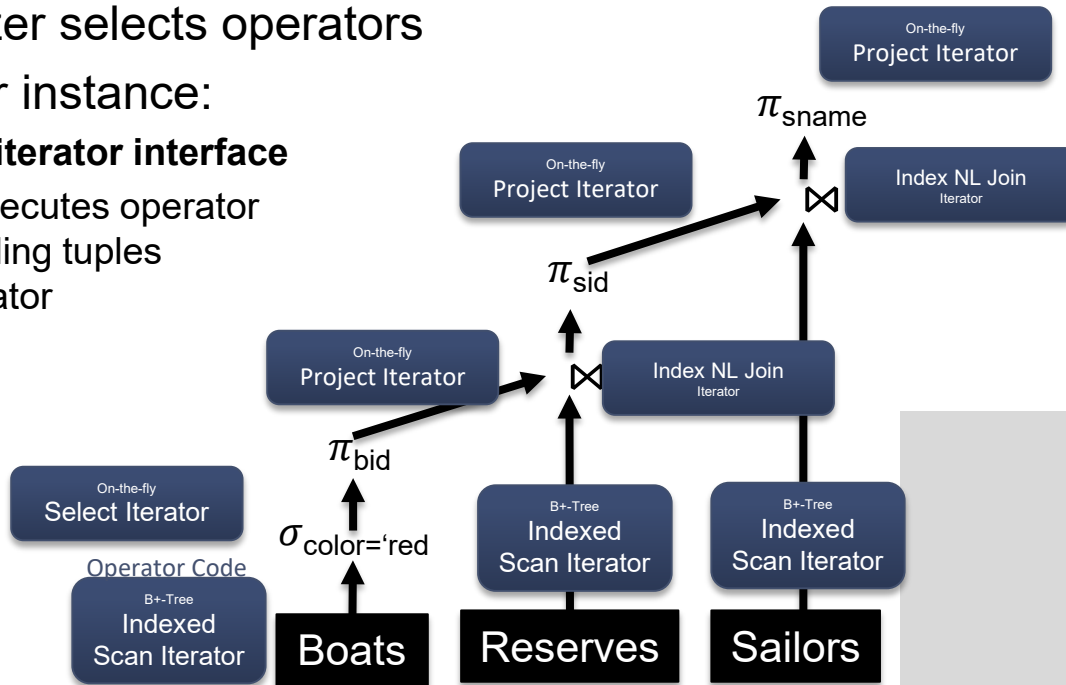




# Query Executor Instantiates Operators

$$\pi_{\text{sname}}(\pi_{\text{sid}}(\pi_{\text{bid}}(\sigma_{\text{color}='red'}(\text{Boats})) \bowtie \text{Res}) \bowtie \text{Sailors})$$

- Query optimizer selects operators
- Each operator instance:
  - Implements **iterator interface**
  - Efficiently executes operator logic forwarding tuples to next operator





# Iterator Interface

The relational operators implemented as subclasses of the class `Iterator`:

```
abstract class iterator {  
    void setup(List<Iterator> inputs);  
    void init(args);  
    tuple next();  
    void close();  
}
```

Notes:

- **Pull-based** computation model
  - e.g., Console calls **init** and **next** which propagates down graph
  - *init/next* can result in either *streaming* (“on-the-fly”) or *blocking* (“batch”) algorithm:
    - streaming: small, constant amount of work per call
    - blocking: does not produce output until it consumes its entire input!
- **Encapsulation**: any iterator can be input to any other!
- **State**: iterators may maintain substantial “internal” state
  - e.g., hash tables, running counts, large sorted files ...



# Example: Select (on-the-fly)

- `init(predicate):`  
    `child.init()`  
    `pred = predicate;`  
    `current = NULL;`
- `next():`  
    `while (current != EOF && !pred(current))`  
        `current = child.next();`  
    `}`  
    `return current;`
- `close():`  
    `child.close()`



# Example: Heap Scan

- **init(relation):**  
heap = open heap file for this relation;  
cur\_page = heap.first\_page(); // first page  
cur\_slot = cur\_page.first\_slot(); // first slot on that page
- **next():**  
if (cur\_page == NULL) return EOF; // End Of Fun  
current = [cur\_page, cur\_slot]; // we will return this recordId  
// advance the slot  
cur\_slot = cur\_slot.next();  
if (cur\_slot == NULL) {  
// advance to next page, first slot  
cur\_page = cur\_page.next();  
if (cur\_page != NULL)  
cur\_slot = cur\_page.first\_slot();  
}  
return current;
- **close():**  
heap.close()



# Example: Sort (2-pass)

- **init(keys):**                   // all of pass 0 in init, a blocking call  
    child.init()  
    repeatedly call child.next() and generate the sorted runs on disk, until child gives EOF  
    // set up for pass 1, assumes enough buffers to merge  
    open each sorted run file and load into input buffer for pass 1
- **next():**                   // pass 1 (assumes enough buffers to merge)  
    output = min tuple across all buffers  
    if min tuple was last one in its buffer, fetch next page from that run into buffer  
    return output (or EOF -- “End of Fun” -- if no tuples remain)
- **close():**  
    deallocate the runs files  
    child.close()






# Example: Group By on Sorted input

- `init(group_keys, aggs):`  
    `child.init()`  
    `cur_group = NULL;`
- `next():`  
    `result = NULL`  
    do {  
        `tup = child.next();`  
        if (`group(tup) != cur_group`) { // New group!  
            if (`cur_group != NULL`) // Form a result for current group  
                `result = [cur_group, final() of all aggs]`  
            `cur_group = group(tup);`  
            call `init()` on all the aggs  
        }  
        call `merge(tup)` on all the aggs  
    } while (!`result`);  
    return `result`;
- `close():`  
    `child.close()`


agg_type	state	init	merge(x)	final
COUNT	count	0	count ++	count
SUM	sum	0	sum += x	sum
AVG	[count, sum]	[0, 0]	[count++, sum+=x]	sum / count
MIN	min	+infinity	min > x ? x : min	min





# A Full Query Plan

- A Query Plan is Single-threaded!
- Trace the calls:
  - Call `init()` on the root `GroupBy`
    - How does `init()` recurse down the chain and return?
  - call `next()` on root
    - How does `next()` recurse down the chain and return a tuple?
- Note how the blocking operator (sort) interacts with the other, streaming operators
- Note how we don't store operator output on disk; tuples stream through the plan's call stack
  - Some operators like `Sort` use disk internally





# Join Operators

R&G 14.4





# Schema for Examples



- Cost Notation
  - $[R]$  : the number of pages to store  $R$
  - $p_R$  : number of records per page of  $R$
  - $|R|$  : the cardinality (number of records) of  $R$ 
    - $|R| = p_R * [R]$
- Reserves (sid: int, bid: int, day: date, rname: string)
  - $[R]=1000$ ,  $p_R=100$ ,  $|R| = 100,000$
- Sailors (sid: int, sname: string, rating: int, age: real)
  - $[S]=500$ ,  $p_S=80$ ,  $|S| = 40,000$



# Simple Nested Loops Join




foreach **record**  $r$  in  $R$  do

  foreach **record**  $s$  in  $S$  do

    if  $\theta(r_i, s_j)$  then add  $\langle r_i, s_j \rangle$  to result buffer

*Note: for simplicity we do not present iterator implementations for the join algorithms.*



$[R]=1000, p_R=100, |R| = 100,000$

$[S]=500, p_S=80, |S| = 40,000$

Cost:

$[R] + |R| [S]$

$= 50,001,000$




# Changing the Join Order



foreach **record**  $s$  in  $S$  do

  foreach **record**  $r$  in  $R$  do

    if  $\theta(r_i, s_j)$  then add  $\langle r_i, s_j \rangle$  to result buffer



$[R]=1000, p_R=100, |R| = 100,000$

$[S]=500, p_S=80, |S| = 40,000$

Cost:


$$\begin{aligned} & [S] + |S| [R] \\ &= 40,000,500 \\ &\text{vs. } 50,001,000 \end{aligned}$$



# Page Nested Loop Join



```
for each rpage in R:
  for each spage in S:
    for each rtuple in rpage:
      for each stuple in spage:
        if join_condition(rtuple, stuple):
          add <rtuple, stuple> to result buffer
```



$$\begin{aligned}\text{Cost} &= [R] + ([R] * [S]) \\ &= 501,000\end{aligned}$$



“Chunk”

# ~~“Block”~~ Nested Loop Join




for each rchunk of B-2 pages of R:

for each spage of S:

for all matching tuples in spage and rchunk:

add <rtuple, stuple> to result buffer



$$\begin{aligned}
 \text{Cost} &= [R] + \lceil [R]/(B-2) \rceil * [S] \\
 &= 1000 + \lceil 1000/(B-2) \rceil * 500 \\
 &= 6,000 \text{ for } B=102 \text{ (~100x better than Page NL!)}
 \end{aligned}$$





# Index Nested Loops Join




foreach **tuple**  $r$  in  $R$  do


    foreach **tuple**  $s$  in  $S$  where  $r_i == s_j$  do

        add  $\langle r_i, s_j \rangle$  to result buffer

$R$



lookup( $r_i$ )





# Index Nested Loops Join Cost



```
foreach tuple r in R do
    foreach tuple s in S where ri == sj do
        add <ri, sj> to result
```

$$\text{Cost} = [R] + |R| * \text{cost to find matching } S \text{ tuples}$$

- If index uses Alt. 1  $\rightarrow$  cost to traverse tree from root to leaf. (e.g., 2-4 IOs)
- For Alt. 2 or 3:
  - Cost to lookup RID(s); typically 2-4 IOs for B+Tree.
  - Cost to retrieve records from RID(s)
    - Clustered index: 1 I/O per *page of matching S tuples*.
    - Unclustered: up to 1 I/O per matching S tuple



# Index Nested Loops Join Cost, Part 2



- **Reserves** (sid: int, bid: int, day: date, *rname*: string)
  - $[R]=1000$ ,  $p_R=100$ ,  $|R| = 100,000$
- **Sailors** (sid: int, *sname*: string, *rating*: int, *age*: real)
  - $[S]=500$ ,  $p_S=80$ ,  $|S| = 40,000$
  - Index on sid



# Index Nested Loops Join Cost, Part 3



- Unclustered Cost(R,S) =  $[R] + |R| * (\text{Search} + \# \text{ matching tuples})$
- Clustered Cost(R,S):  $[R] + |R| * (\text{Search} + \# \text{ of matching pages})$
- Here, sid is the primary key for Sailors, so there is exactly one matching sailor for each tuple in R
- Unclustered B+-Tree height 2 (3 I/Os from root to leaf):
  - $R \bowtie S$ :  $1000 + (100,000) * (3 + 1) = 401,000$
- Clustered B+-tree height 2 (3 I/Os from root to leaf)
  - $R \bowtie S$ :  $1000 + (100,000) * (3 + 1) = 401,000$



# Sort-Merge Join



- Requires equality predicate:
  - Equi-Joins & Natural Joins
- Two Stages:
  - Sort tuples in R and S by join key
    - All tuples with same key in consecutive order
    - Input might already be sorted ... why?
  - Join Pass: Merge-scan the sorted partitions and emit tuples that match



# Sort-Merge Join, Part 1

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```


sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107






# Sort-Merge Join, Part 2

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty



sid	bid
28	103
28	104
31	101
31	102
42	142
58	107





# Sort-Merge Join, Part 3

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy	→	28	104
31	lubber		31	101
31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107









# Sort-Merge Join, Part 4


```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty




sid	bid
28	103
28	104
31	101
31	102
42	142
58	107






# Sort-Merge Join, Part 5


```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty




sid	bid
28	103
28	104
31	101
31	102
42	142
58	107






# Sort-Merge Join, Part 6


```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty



sid	bid
28	103
28	104
31	101
31	102
42	142
58	107





# Sort-Merge Join, Part 7

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107





# Sort-Merge Join, Part 8

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103





# Sort-Merge Join, Part 9

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103





# Sort-Merge Join, Part 10

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107



sid	sname	bid
28	yuppy	103





# Sort-Merge Join, Part 11

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103








# Sort-Merge Join, Part 12

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty



sid	bid
28	103
28	104
31	101
31	102
42	142
58	107



sid	sname	bid
28	yuppy	103



# Sort-Merge Join, Part 13

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 14

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107


sid	sname	bid
28	yuppy	103
28	yuppy	104






# Sort-Merge Join, Part 15

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```



sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty



sid	bid
28	103
28	104
31	101
31	102
42	142
58	107



sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 16

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid	
22	dustin		28	103	
28	yuppy	→	28	104	
31	lubber		31	101	
31	lubber2		31	102	
44	guppy		42	142	
58	rusty		58	107	

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 17

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber	→	31	101
31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 18



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber	→	31	101
31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104



# Sort-Merge Join, Part 19

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber	→	31	101
31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 20

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber		31	101
31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 21



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104



# Sort-Merge Join, Part 22

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 23

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 24

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 25

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104





# Sort-Merge Join, Part 26

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101





# Sort-Merge Join, Part 27

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101







# Sort-Merge Join, Part 28

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101





# Sort-Merge Join, Part 29

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101





# Sort-Merge Join, Part 30

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101





# Sort-Merge Join, Part 31

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102





# Sort-Merge Join, Part 32

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102





# Sort-Merge Join, Part 33

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102





# Sort-Merge Join, Part 34

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102





# Sort-Merge Join, Part 35

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102







# Sort-Merge Join, Part 36



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102



# Sort-Merge Join, Part 37



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102



# Sort-Merge Join, Part 38

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber	→	31	101
31	lubber2	→	31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102





# Sort-Merge Join, Part 39

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname		sid	bid
22	dustin		28	103
28	yuppy		28	104
31	lubber	→	31	101
→ 31	lubber2		31	102
44	guppy		42	142
58	rusty		58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102



# Sort-Merge Join, Part 40

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102



# Sort-Merge Join, Part 41

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
  }
  return result
}
else {
  reset s to mark
  advance r
  mark = NULL
}
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102



# Sort-Merge Join, Part 42

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101





# Sort-Merge Join, Part 43

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101





# Sort-Merge Join, Part 44

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101





# Sort-Merge Join, Part 45

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101





# Sort-Merge Join, Part 46

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
  }
  return result
}
else {
  reset s to mark
  advance r
  mark = NULL
}
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101



# Sort-Merge Join, Part 47

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 48

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 49

```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 50

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 51

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 52



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102

# Sort-Merge Join, Part 53



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 54



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 55



```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 56

```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 57



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 57



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 58



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102





# Sort-Merge Join, Part 59



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 60



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 61



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 62



```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102

# Sort-Merge Join, Part 63



```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
    return result
  }
  else {
    reset s to mark
    advance r
    mark = NULL
  }
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 64



```
do {
  if (!mark) {
    while (r < s) { advance r }
    while (r > s) { advance s }
    // mark start of "block" of S
    mark = s
  }
  if (r == s) {
    result = <r, s>
    advance s
  }
  return result
}
else {
  reset s to mark
  advance r
  mark = NULL
}
}
```

sid	sname
22	dustin
28	yuppy
31	lubber
31	lubber2
44	guppy
58	rusty

sid	bid
28	103
28	104
31	101
31	102
42	142
58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102



# Sort-Merge Join, Part 65



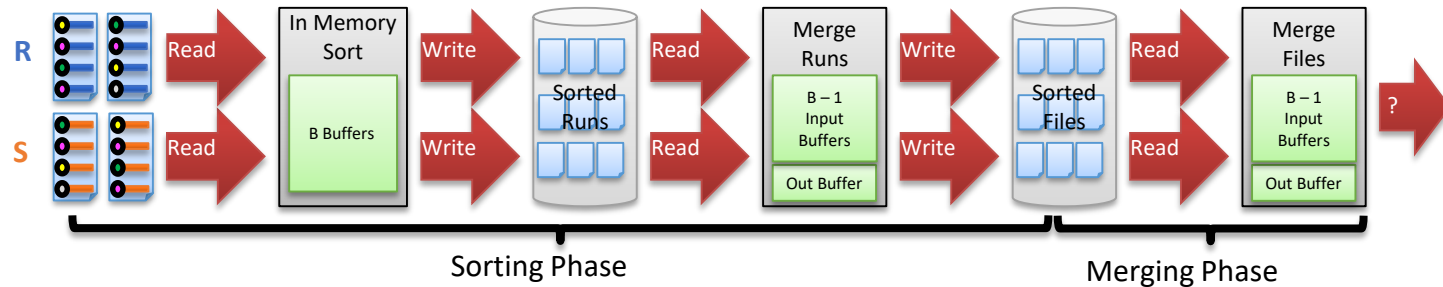
```
do {  
  if (!mark) {  
    while (r < s) { advance r }  
    while (r > s) { advance s }  
    // mark start of "block" of S  
    mark = s  
  }  
  if (r == s) {  
    result = <r, s>  
    advance s  
    return result  
  }  
  else {  
    reset s to mark  
    advance r  
    mark = NULL  
  }  
}
```

sid	sname	sid	bid
22	dustin	28	103
28	yuppy	28	104
31	lubber	31	101
31	lubber2	31	102
44	guppy	42	142
58	rusty	58	107

sid	sname	bid
28	yuppy	103
28	yuppy	104
31	lubber	101
31	lubber	102
31	lubber2	101
31	lubber2	102
58	rusty	107



# Cost of Sort-Merge Join



- Cost: Sort **R** + Sort **S** + ( $[R] + [S]$ )
  - But in worst case, last term could be  $|R| * [S]$  (very unlikely!)
  - Q: what is worst case?
- Question: How big does the buffer have to be to sort both R and S in two passes each?
- Suppose buffer  $B > \sqrt{(\max([R], [S]))}$ 
  - Both R and S can be sorted in 2 passes
  - $4 * 1000 + 4 * 500 + (1000 + 500) = 7500$