




# Relational Query Optimization I: The Plan Space

R&G 15





# Architecture of a DBMS





# Query Optimization is Magic

- The bridge between a *declarative* domain-specific language...
  - “What” you want as an answer
- ... and custom *imperative* computer programs
  - “How” to compute the answer
- In 2018 terms:
  - This is AI-driven Software Synthesis
  - That’s not just marketing!
    - Similar to cutting edge AI work today
    - Optimization + heuristic pruning
    - Research exploring the use of modern AI techniques to improve that pruning (e.g. Deep Reinforcement Learning)

# Invented in 1979 by Pat Selinger et al.

- We'll focus on “System R” (“Selinger”) optimizers
- “Cascades” optimizer is the other common one
  - Later, with notable differences, but similar big picture

## Access Path Selection in a Relational Database Management System

P. Griffiths Selinger  
M. M. Astrahan  
D. D. Chamberlin  
R. A. Lorie  
T. G. Price

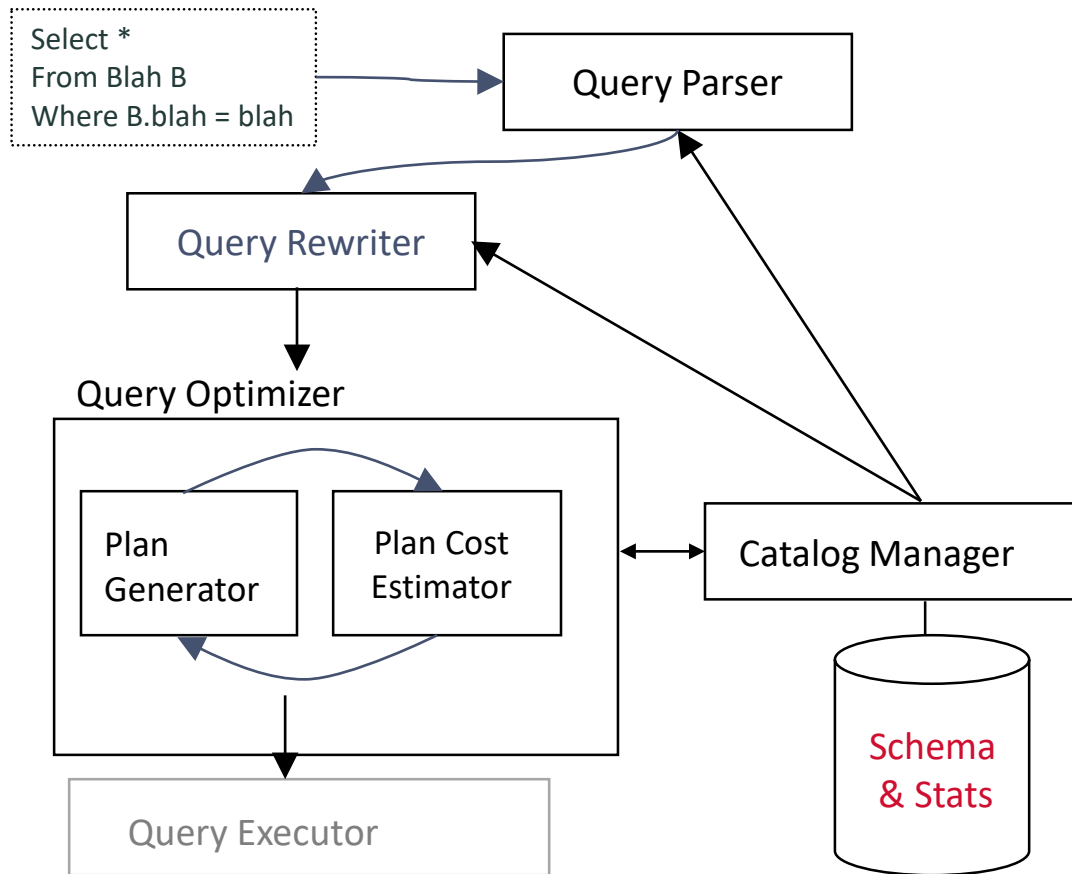
IBM Research Division, San Jose, California 95193

ABSTRACT: In a high level query and data manipulation language such as SQL, requests are stated declaratively, without

retrieval. Nor does a user specify in what order joins are to be performed. The System R optimizer chooses both join order




# Query Parsing & Optimization



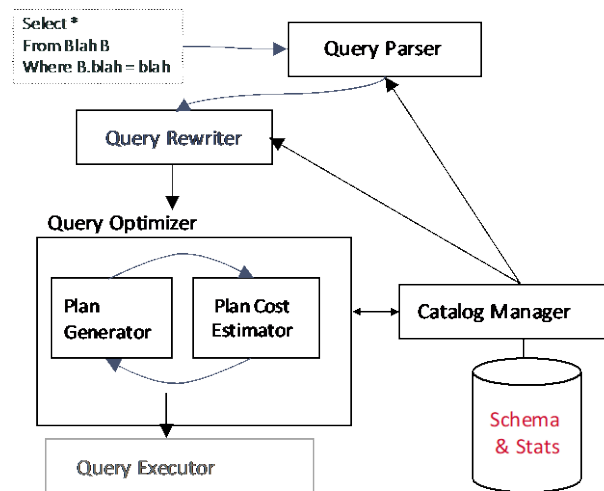
# Query Parsing & Optimization Part 2

- Query parser
  - Checks correctness, authorization
  - Generates a parse tree
  - Straightforward



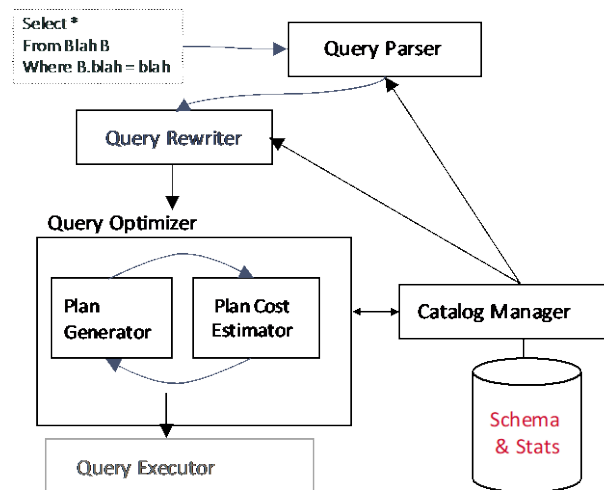
# Query Parsing & Optimization Part 3

- Query rewriter
  - Converts queries to canonical form
    - flatten views
    - subqueries into fewer query blocks
  - Weak spot in many open-source DBMSs



# Query Parsing & Optimization Part 4

- “Cost-based” Query Optimizer
  - Optimizes 1 query block at a time
    - Select, Project, Join
    - GroupBy/Agg
    - Order By (if top-most block)
  - Uses catalog stats to find least-“cost” plan per query block
  - “Soft underbelly” of every DBMS
    - Sometimes not truly “optimal”






# Query Optimization Overview

- Query block can be converted to relational algebra
- Rel. Algebra converts to tree
- Each operator has implementation choices
- Operators can also be applied in different orders!

```
SELECT S.sname
  FROM Reserves R, Sailors S
 WHERE R.sid=S.sid
    AND R.bid=100
    AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)}$   
(Reserves  $\bowtie$  Sailors)





# Query Optimization: The Components

- Three beautifully orthogonal concerns:
  - Plan space:
    - for a given query, what plans are considered?
  - Cost estimation:
    - how is the cost of a plan estimated?
  - Search strategy:
    - how do we “search” in the “plan space”?



# Query Optimization: The Goal

- Optimization goal:
  - Ideally: Find the plan with least actual cost.
  - Reality: Find the plan with least estimated cost.
    - And try to avoid really bad actual plans!

# Today

- We will get a feel for the plan space
- Explore one simple example query



# Relational Algebra Equivalences: Selections

- Selections:
  - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$  (cascade)
  - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$  (commute)



# Relational Algebra Equivalences: Projections

- Selections:
  - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$  (cascade)
  - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$  (commute)
- Projections:
  - $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1, \dots, an-1}(R))\dots)$  (cascade)



# Relational Algebra Equivalences: Cartesian Product

- Selections:
  - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$  (cascade)
  - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$  (commute)
- Projections:
  - $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1, \dots, an-1}(R))\dots)$  (cascade)
- Cartesian Product
  - $R \times (S \times T) \equiv (R \times S) \times T$  (associative)
  - $R \times S \equiv S \times R$  (commutative)



# Are Joins Associative and Commutative?





- After all, just Cartesian Products with Selections
- You can think of them as associative and commutative...
- ...But beware of join turning into cross-product!
  - Consider  $R(a,z)$ ,  $S(a,b)$ ,  $T(b,y)$
  - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \not\equiv S \bowtie_{b=b} (T \bowtie_{a=a} R)$  (*not* legal!!)
    - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \not\equiv S \bowtie_{b=b} (T \times R)$  (*not* the same!!)
  - $(S \bowtie_{b=b} T) \bowtie_{a=a} R \alpha S \bowtie_{b=b \wedge a=a} (T \times R)$  (the same!!)



# Join ordering, again

- Similarly, note that some join orders have cross products, some don't
- Equivalent for the query above:

```
SELECT *  
  FROM R, S, T  
 WHERE R.a = S.a  
       AND S.b = T.b;
```





# Some Common Heuristics: Selections

- Selection cascade and pushdown
  - Apply selections as soon as you have the relevant columns
  - Ex:
    - $\pi_{\text{sname}} (\sigma_{\text{bid}=100 \wedge \text{rating} > 5} (\text{Reserves} \bowtie_{\text{sid}=\text{sid}} \text{Sailors}))$
    - $\pi_{\text{sname}} (\sigma_{\text{bid}=100} (\text{Reserves}) \bowtie_{\text{sid}=\text{sid}} \sigma_{\text{rating} > 5} (\text{Sailors}))$





# Some Common Heuristics: Projections

- Projection cascade and pushdown
  - Keep only the columns you need to evaluate downstream operators
  - Ex:
    - $\pi_{\text{sname}} \sigma_{(\text{bid}=100 \wedge \text{rating} > 5)} (\text{Reserves} \bowtie_{\text{sid}=\text{sid}} \text{Sailors})$
    - $\pi_{\text{sname}} (\pi_{\text{sid}} (\sigma_{\text{bid}=100} (\text{Reserves})) \bowtie_{\text{sid}=\text{sid}} \pi_{\text{sname}, \text{sid}} (\sigma_{\text{rating} > 5} (\text{Sailors})))$



# Some Common Heuristics

- Avoid Cartesian products
  - Given a choice, do theta-joins rather than cross-products
  - Consider  $R(a,b)$ ,  $S(b,c)$ ,  $T(c,d)$
  - Favor  $(R \bowtie S) \bowtie T$  over  $(R \times T) \bowtie S$





# Physical Equivalences

- Base table access,  
with single-table selections and projections
  - Heap scan
  - Index scan (if available on referenced columns)
- Equijoins
  - Block (Chunk) Nested Loop: simple, exploits extra memory
  - Index Nested Loop: often good if 1 rel small and the other indexed properly
  - Sort-Merge Join: good with small memory, equal-size tables
  - Grace Hash Join: even better than sort with 1 small table
    - Or Hybrid if you have it
- Non-Equijoins
  - Block Nested Loop



# Schema for Examples


Sailors (sid: integer, sname: text, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: date, rname: text)

- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
  - Assume there are 100 boats
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
  - Assume there are 10 different ratings
- Assume we have 5 pages to use for joins.


# Motivating Example: Plan 1

- Here's a reasonable query plan:



```
SELECT S.sname
FROM Reserves R, sailors S
WHERE R.sid=S.sid
AND R.bid=100
AND S.rating>5
```


# Motivating Example: Plan 1 Cost



- Let's estimate the cost:
- Scan Sailors (500 IOs)
- For each page of Sailors, Scan Reserves (1000 IOs)
- Total:  $500 + 500 \cdot 1000$ 
  - **500,500 IOs**




# Motivating Example: Plan 1 Cost Analysis

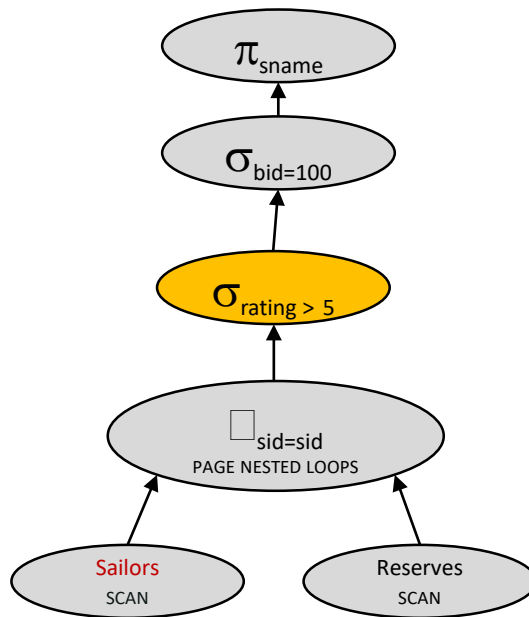


- Cost: **500+500\*1000 I/Os**
- By no means the worst plan!
- Misses several opportunities:
  - selections could be 'pushed' down
  - no use made of indexes
- Goal of optimization:
  - Find faster plans that compute the same answer.


# Selection Pushdown



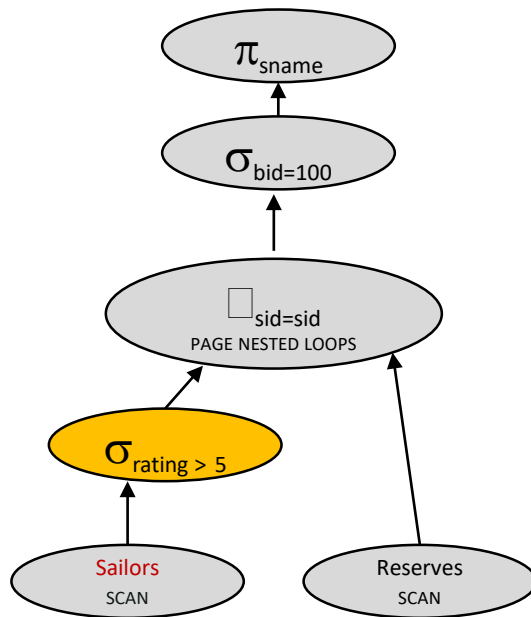
500,500 IOs



# Selection Pushdown, cont



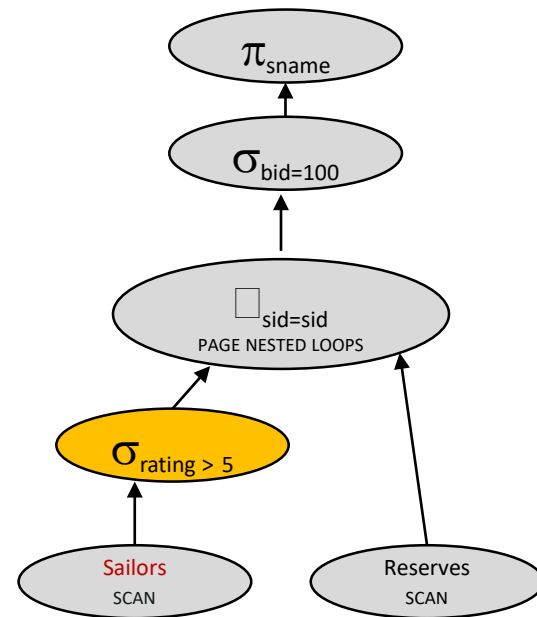
500,500 IOs




Cost?

# Query Plan 2 Cost

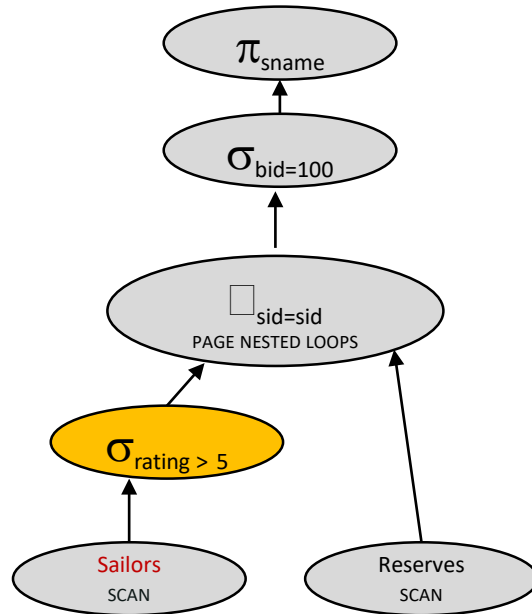
- Let's estimate the cost:
- Scan Sailors (500 IOs)
- For each pageful of high-rated Sailors,  
Scan Reserves (1000 IOs)
- Total:  $500 + ??? * 1000$
- Total:  $500 + 250 * 1000$



# Decision?

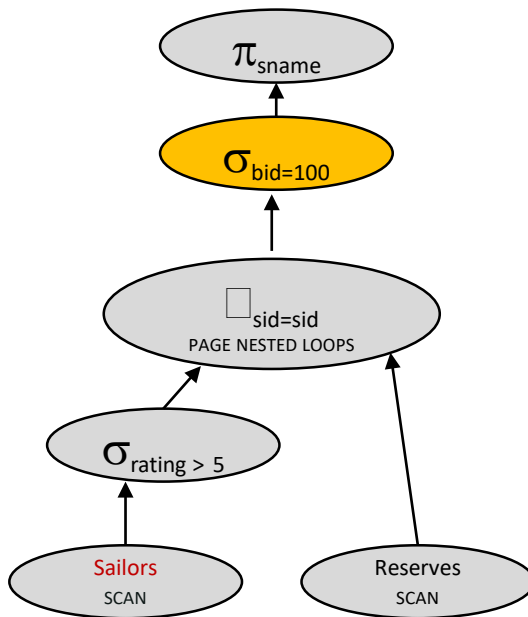
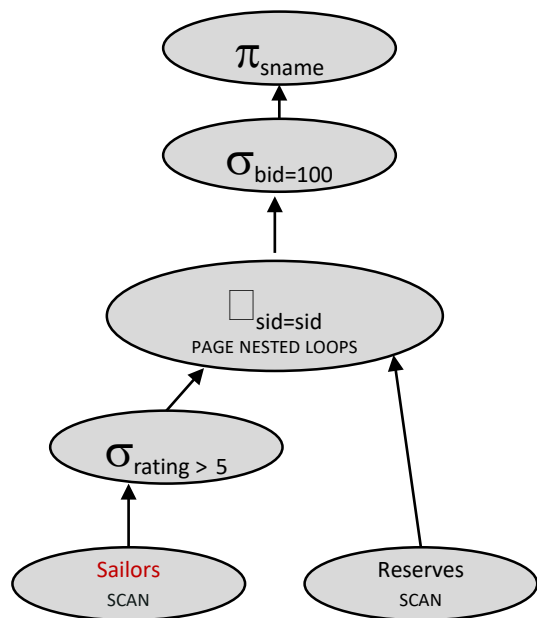


500,500 IOs



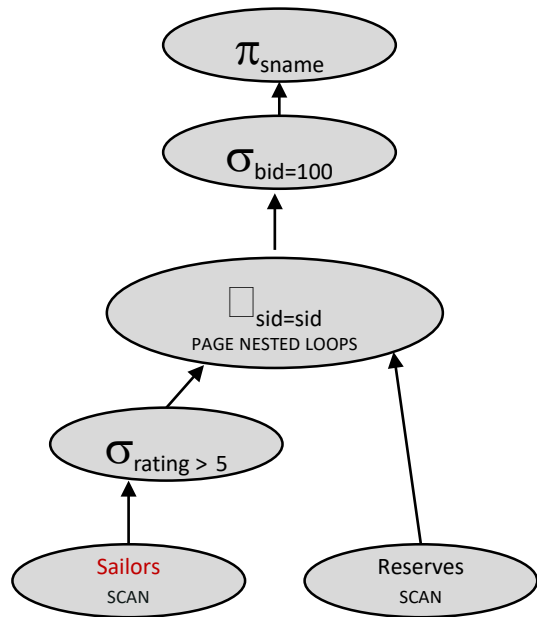
250,500 IOs

# More Selection Pushdown




250,500 IOs

# More Selection Pushdown, cont



250,500 IOs




Cost???

# Query Plan 3 Cost Analysis

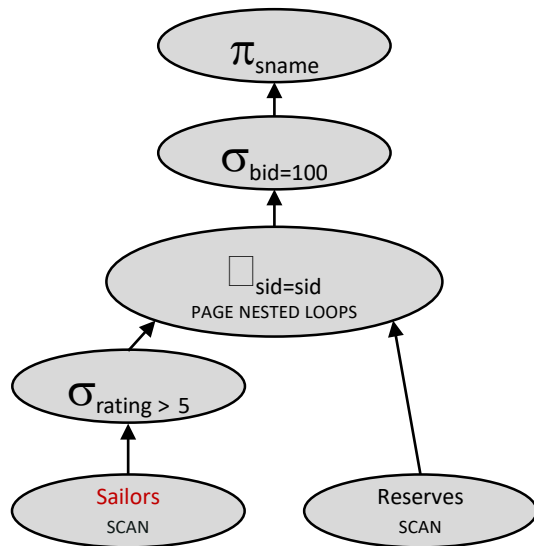
Let's estimate the cost:

- Scan Sailors (500 IOs)
- For each pageful of high-rated Sailors,  
Do what? (??? IOs)
- Total:  $500 + 250 * ???$
- Total:  $500 + 250 * 1000!$






# More Selection Pushdown Analysis



250,500 IOs





250,500 IOs

Pushing a selection into the inner loop of a nested loop join doesn't save I/Os! Essentially equivalent to having the selection above.







# Join Ordering



250,500 IOs


# Join Ordering, cont




250,500 IOs

# Query Plan 4 Cost


- Let's estimate the cost:
- Scan Reserves (1000 IOs)
- For each pageful of Reserves for bid 100,  
Scan Sailors (500 IOs)
- Total: 1000 + ???\*500
- Total: **1000 + 10\*500**



# Decision 3




250,500 IOs




6000 IOs

# Materializing Inner Loops




6000 IOs




Cost???

# Materializing Inner Loops, cont



6000 IOs




Cost???




# Plan 5 Cost Analysis


- Let's estimate the cost:
- Scan Reserves (1000 IOs)
- Scan Sailors (500 IOs)
- Materialize Temp table T1 (??? IOs)
- For each pageful of Reserves for bid 100, Scan T1 (??? IOs)
- Total:  $1000 + 500 + ??? + 10 * ???$
- $1000 + 500 + 250 + (10 * 250)$



# Materializing Inner Loops, cont.




6000 IOs




4250 IOs


# Join Ordering Again



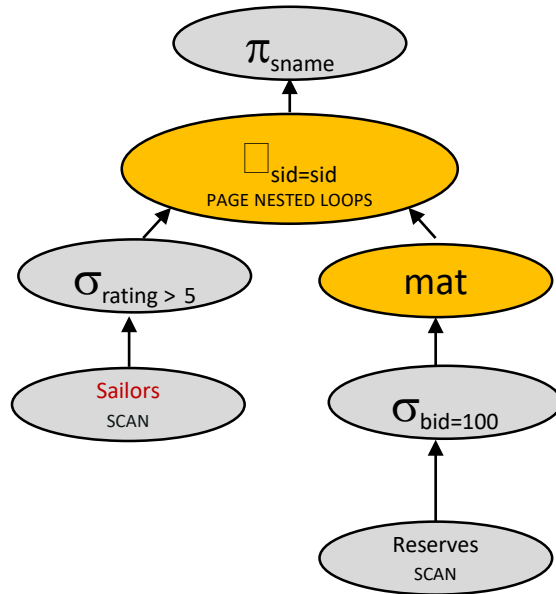
4250 IOs



# Join Ordering Again, Cont



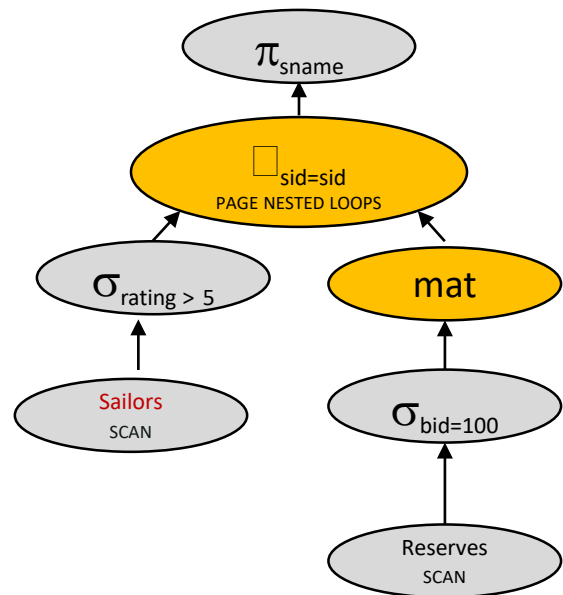
4250 IOs



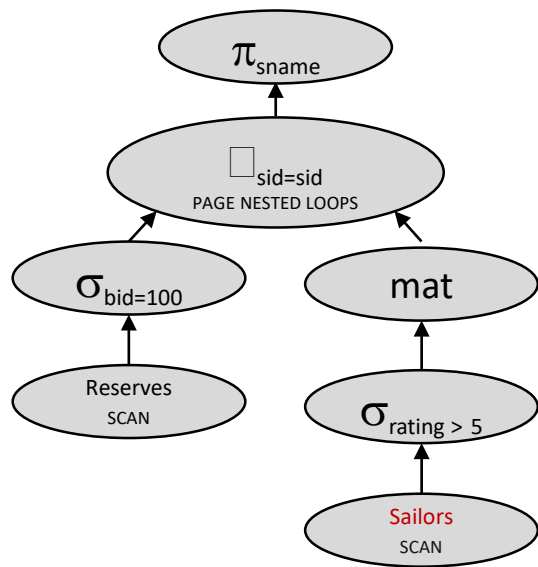
Cost???

# Plan 6 Cost Analysis

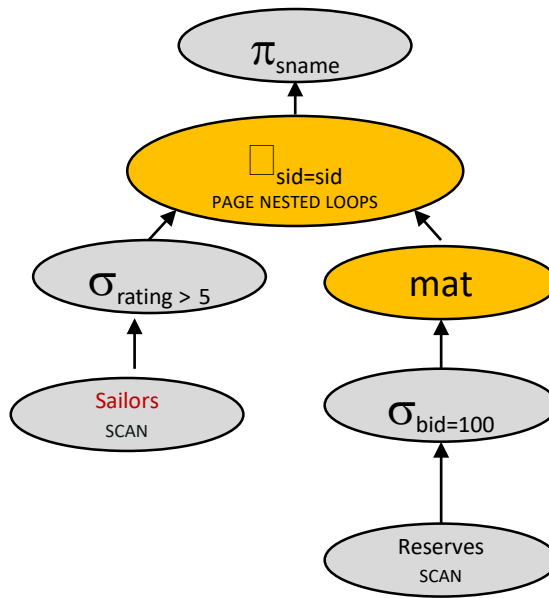
- Let's estimate the cost:
- Scan Sailors (500 IOs)
- Scan Reserves (1000 IOs)
- Materialize Temp table T1 (??? IOs)
- For each pageful of high-rated Sailors, Scan T1 (??? IOs)
- Total:  $500 + 1000 + ??? + 250 * ???$
- $500 + 1000 + 10 + (250 * 10)$



# Decision 4

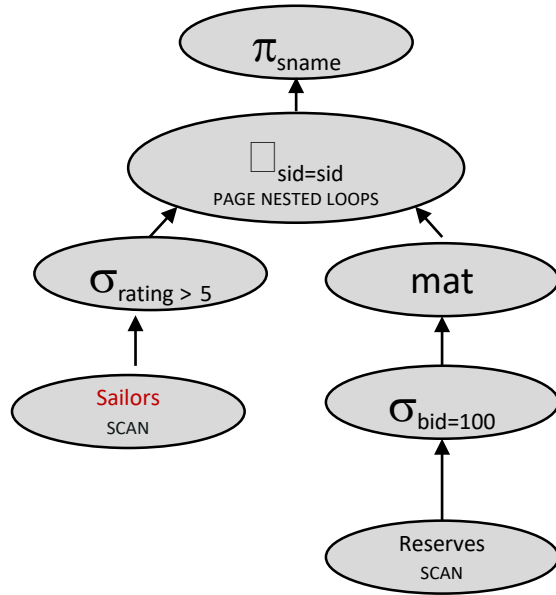


4250 IOs

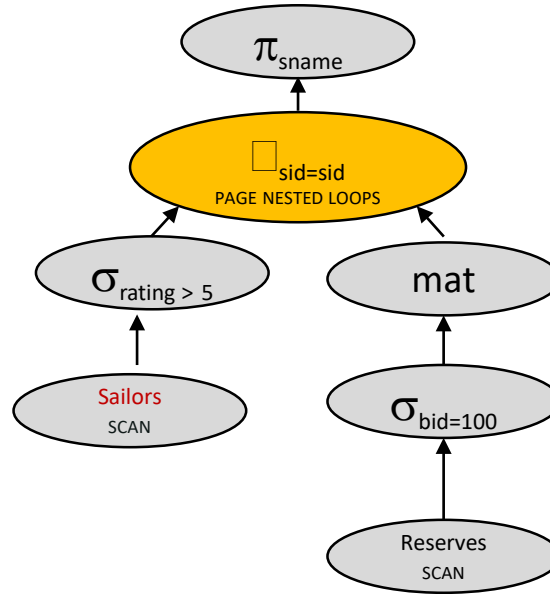


4010 IOs

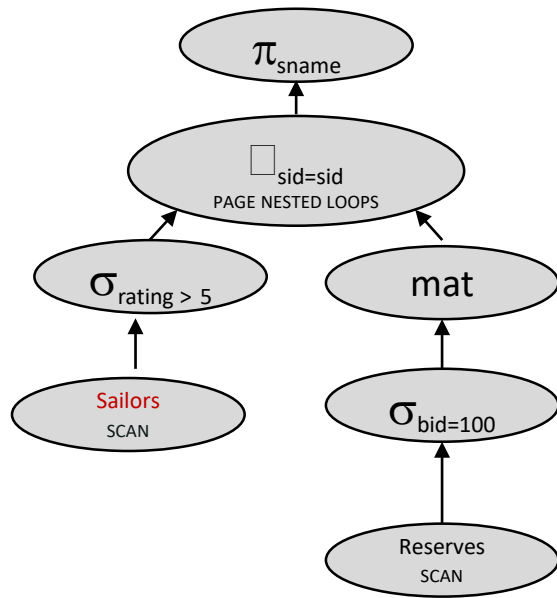
# Join Algorithm




4010 IOs



# Join Algorithm, cont.



4010 IOs




Cost???



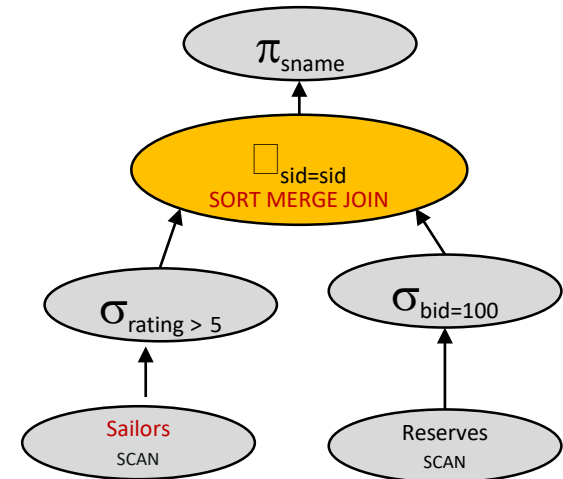
# Query Plan 7 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Reserves (1000)
- Scan Sailors (500)
- Sort high-rated sailors (???)  
Note: pass 0 doesn't do read I/O, just gets input from select.
- Sort reservations for boat 100 (???)  
Note: pass 0 doesn't do read I/O, just gets input from select.
- How many passes for each sort with  $\log_4$ ?
- Merge  $(10+250) = 260$
- Total:




# Query Plan 7 Cost Analysis Part 2


- With 5 buffers, cost of plan:
  - Scan Reserves (1000)
  - Scan Sailors (500)
  - Sort
    - 2 passes for reserves  
pass 0 = 10 to write, pass 1 = 2\*10 to read/write
    - 4 passes for sailors  
pass 0 = 250 to write, pass 1,2,3 = 2\*250 to read/write
  - Merge (10+250) = 260
- 1000 + 500 + sort reserves(10 + 2\*10) + sort sailors  
(250 + 3\*2\*250) + merge (10+250) = **3540**



# Decision 5




4010 IOs



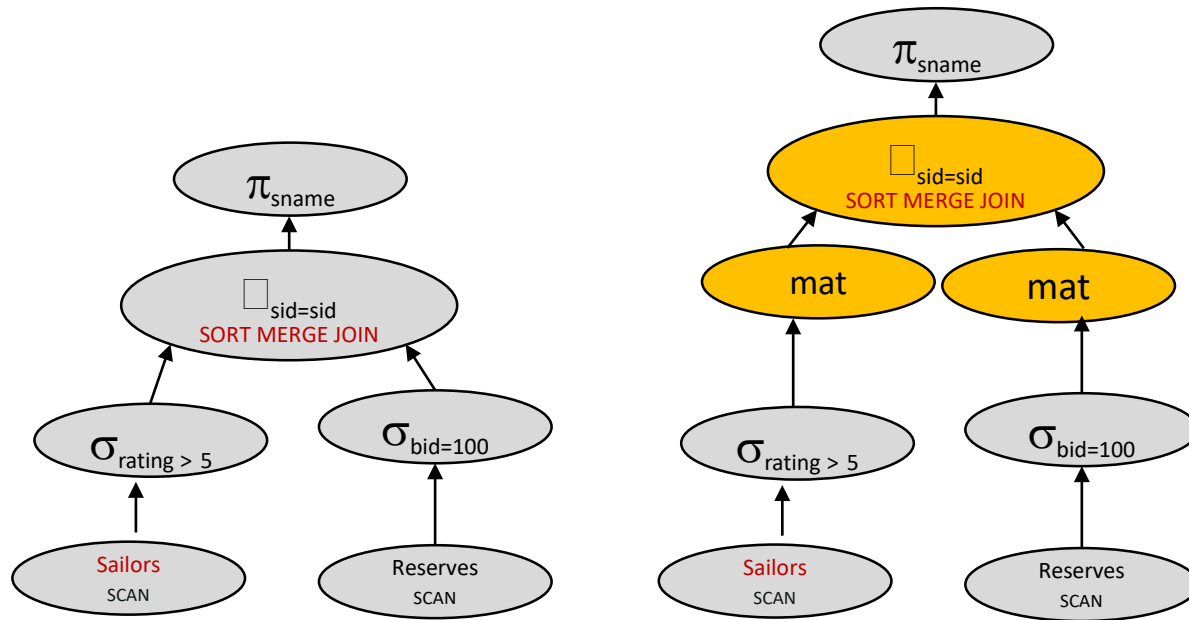
3540 IOs

Textbook considers this:



3540 IOs

# Textbook considers this, cont:

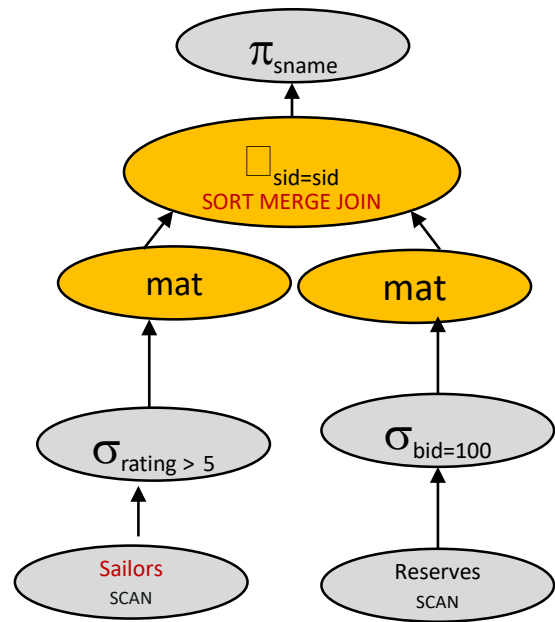


3540 IOs


Cost???

# Plan 8 Cost Analysis

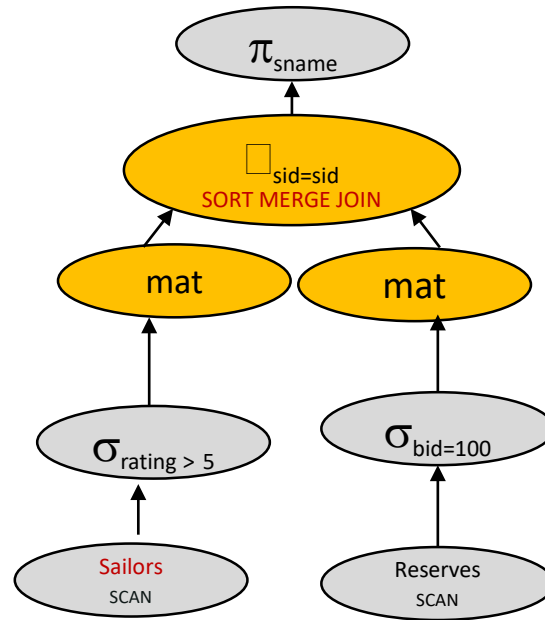
- With 5 buffers, cost of plan:
- Scan Sailors (500), write T1 (250)
- Scan Reserves (1000), write T2 (10)
- Sort T1 (???)
- Sort T2 (???)
- **How many passes for each sort?**
  - 2 passes for reserves (2\*2\*10 to read/write)
  - 4 passes for sailors (4\*2\*250 to read/write)
- Merge (10+250) = 260
- Total:  
1000 + 10 + 500 + 250 + 2\*2\*10 +  
4\*2\*250 + merge (10+250) = **4060**



# Decision 6

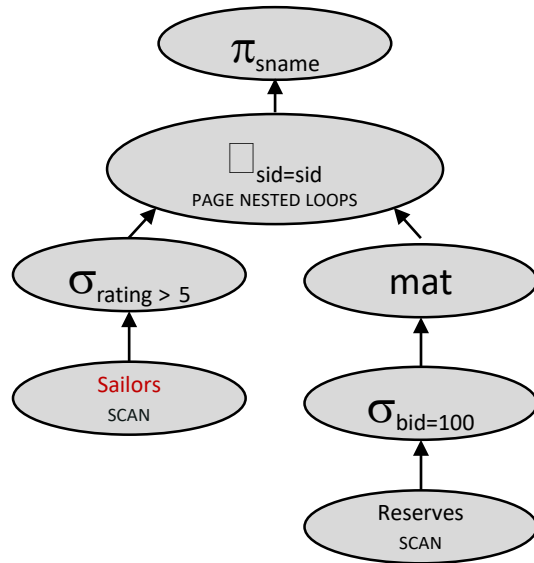


3540 IOs

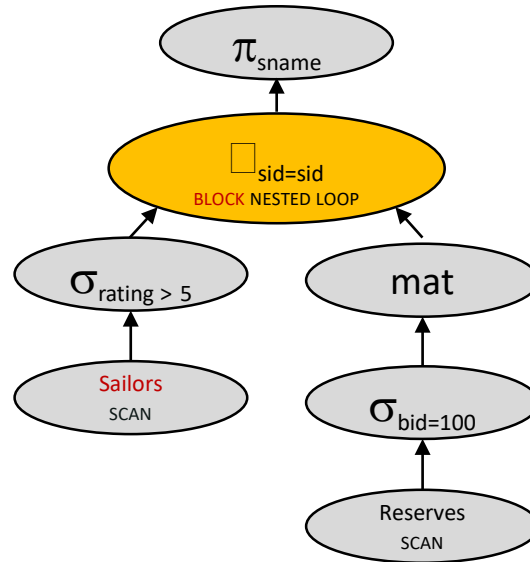


4060 IOs

# Join Algorithm Again, Again



4010 IOs




Cost???




# Query 9 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Sailors (500)
- Scan Reserves (1000)
- Write Temp T1 (10)
- For each blockful of high-rated sailors
  - Loop on T1 (??? \* 10)
- Total:

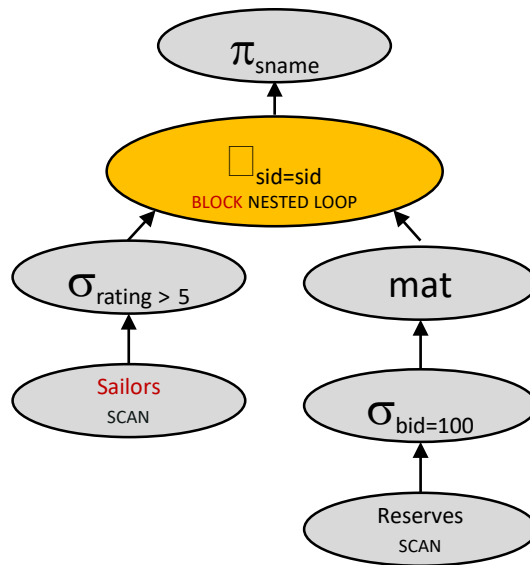
$$500 + 1000 + 10 + (\text{ceil}(250/3) * 10) = 500 + 1000 + 10 + (84 * 10) =$$



# Decision 7



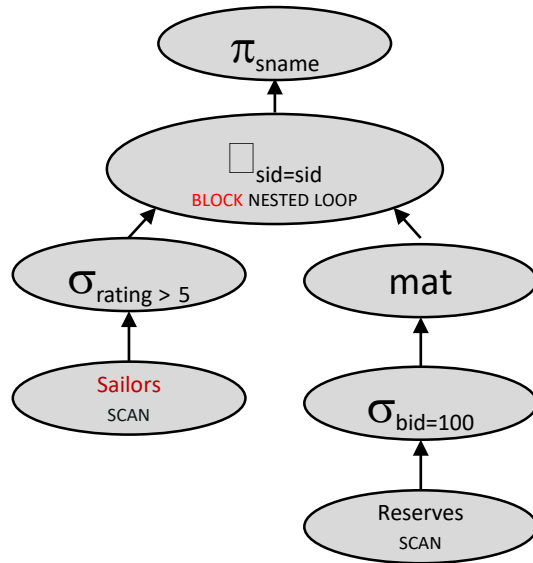
3540 IOs



2350 IOs

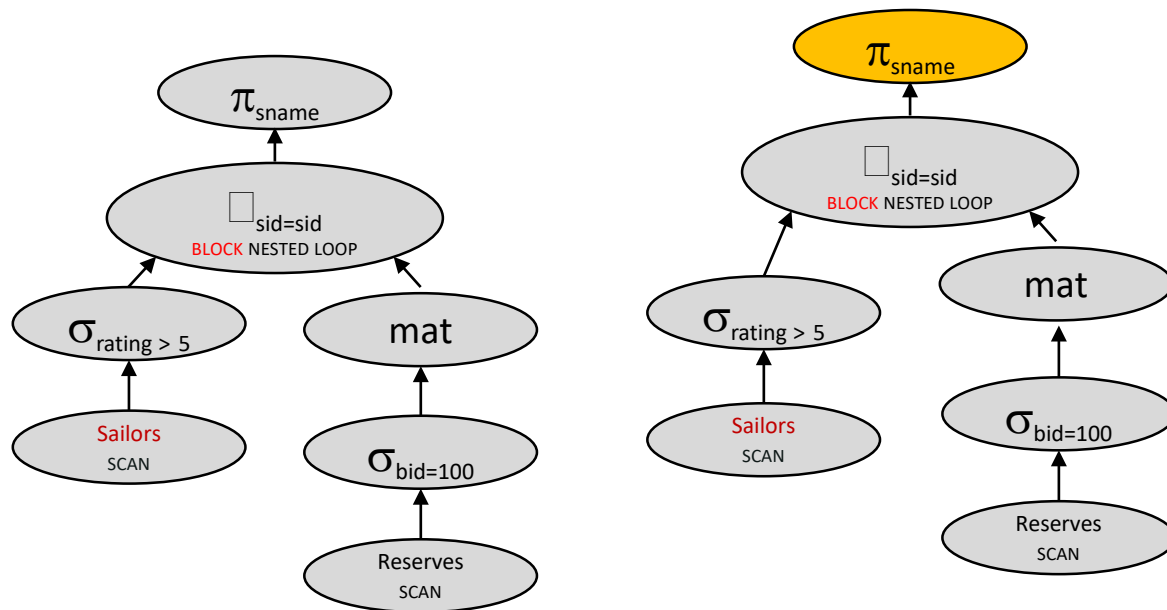


# Projection Cascade & Pushdown



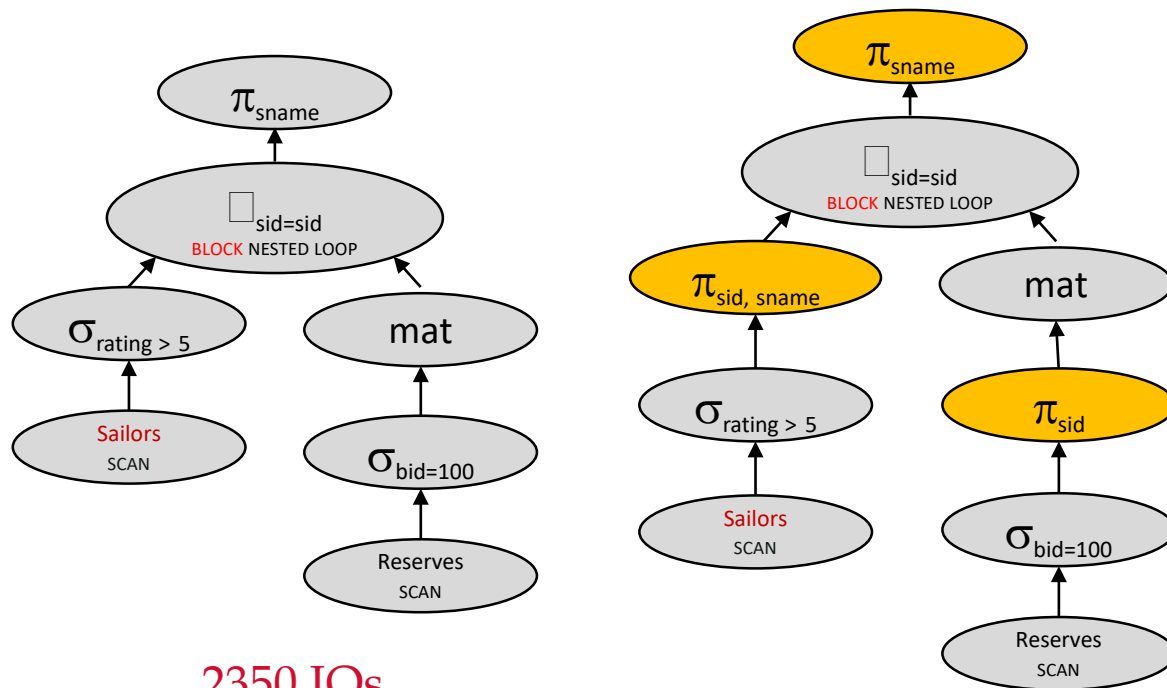
2350 IOs

# Projection Cascade & Pushdown, cont

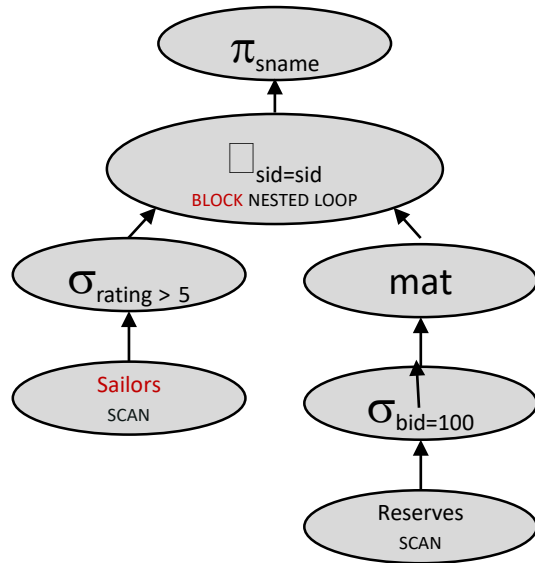


2350 IOs

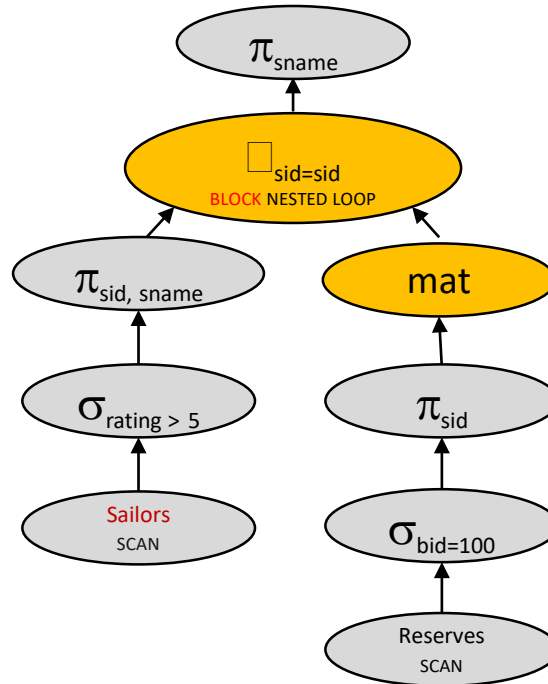
# Projection Cascade & Pushdown, cont



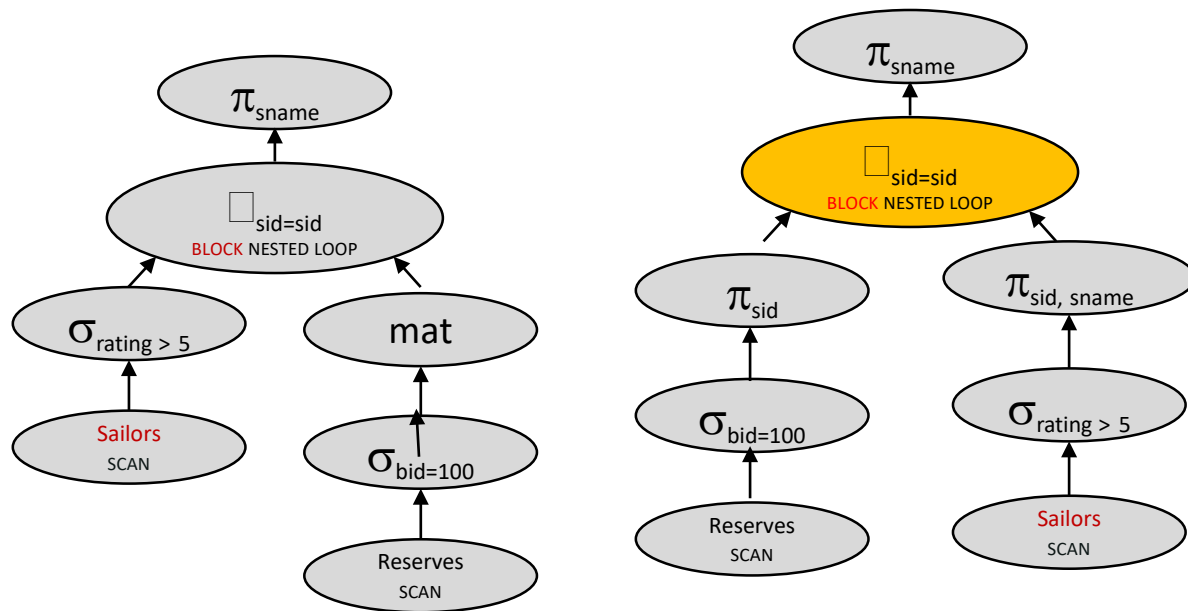
# With Join Reordering, no Mat



2350 IOs




# With Join Reordering, no Mat cont



2350 IOs

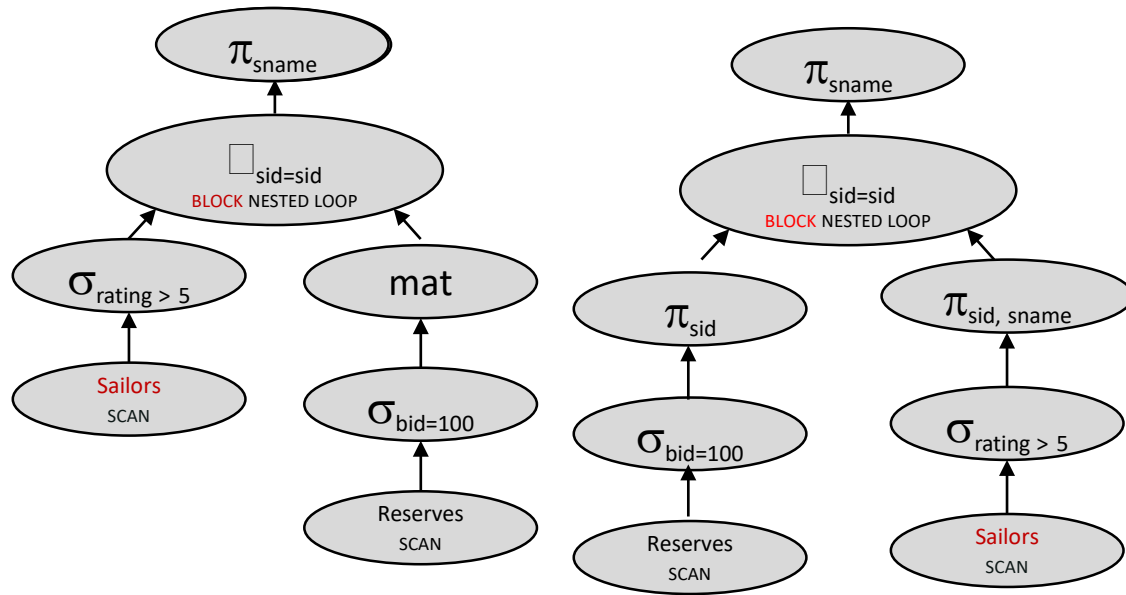
# Plan 11 Cost Analysis

- With 5 buffers, cost of plan:
- Scan Reserves (1000)
- For each blockful of sids that rented boat 100
- (recall Reserve tuple is 40 bytes, assume sid is 4 bytes)
- Loop on Sailors (??? \* 500)
- Total: 1500





# With Join Reordering, no Mat, cont.



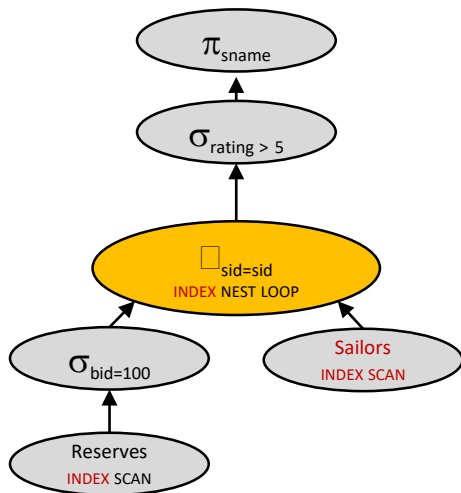


2350 IOs

1500 IOs


# How About Indexes?

- Indexes:
  - Reserves.bid clustered
  - Sailors.sid unclustered
- Assume indexes fit in memory



# Index Cost Analysis

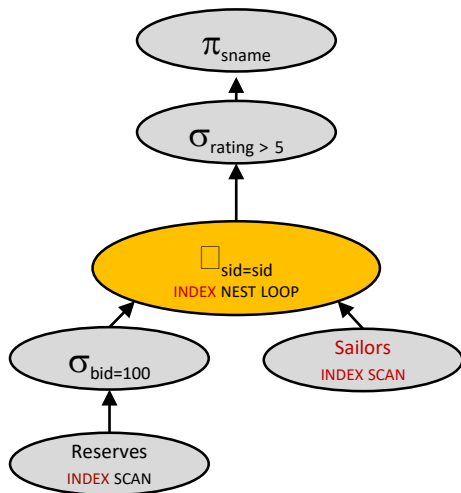
- **No projection pushdown to left** for  $\pi_{sname}$ 
  - Projecting out unnecessary fields from outer of Index NL doesn't make an I/O difference.
- **No selection pushdown to right** for  $\sigma_{rating > 5}$ 
  - Does not affect Sailors.sid index lookup
- With clustered index on bid of Reserves, we access how many pages of Reserves?:
  - $100,000/100 = 1000$  tuples on  $1000/100 = 10$  pages.
- Join column sid is a **key** for Sailors.
  - At most one matching tuple, unclustered index on sid OK



1010 IOs

# Index Cost Analysis Part 2

- With clustered index on bid of Reserves, we access how many pages of Reserves?:
  - $100,000/100 = 1000$  tuples on  $1000/100 = 10$  pages.
- for each Reserves tuple 1000  
get matching Sailors tuple (1 IO)  
(recall: 100 Reserves per page, 1000 pages)
- $10 + 1000*1$
- Cost: Selection of Reserves tuples (10 I/Os); then, for each, must get matching Sailors tuple (1000); total 1010 I/Os.



1010 I/Os



# Summing up

- There are *lots* of plans
  - Even for a relatively simple query
- Engineers often think they can pick good ones
  - E.g. MapReduce API was based on that assumption
  - So was the COBOL API of 1970's!
- Not so clear that's true!
  - Manual query planning can be tedious, technical
  - Machines are better at enumerating options than people
    - Hence AI
  - We will see soon how optimizers make simplifying assumptions