

# Analytics & Machine Learning in Data Systems (Part 2)

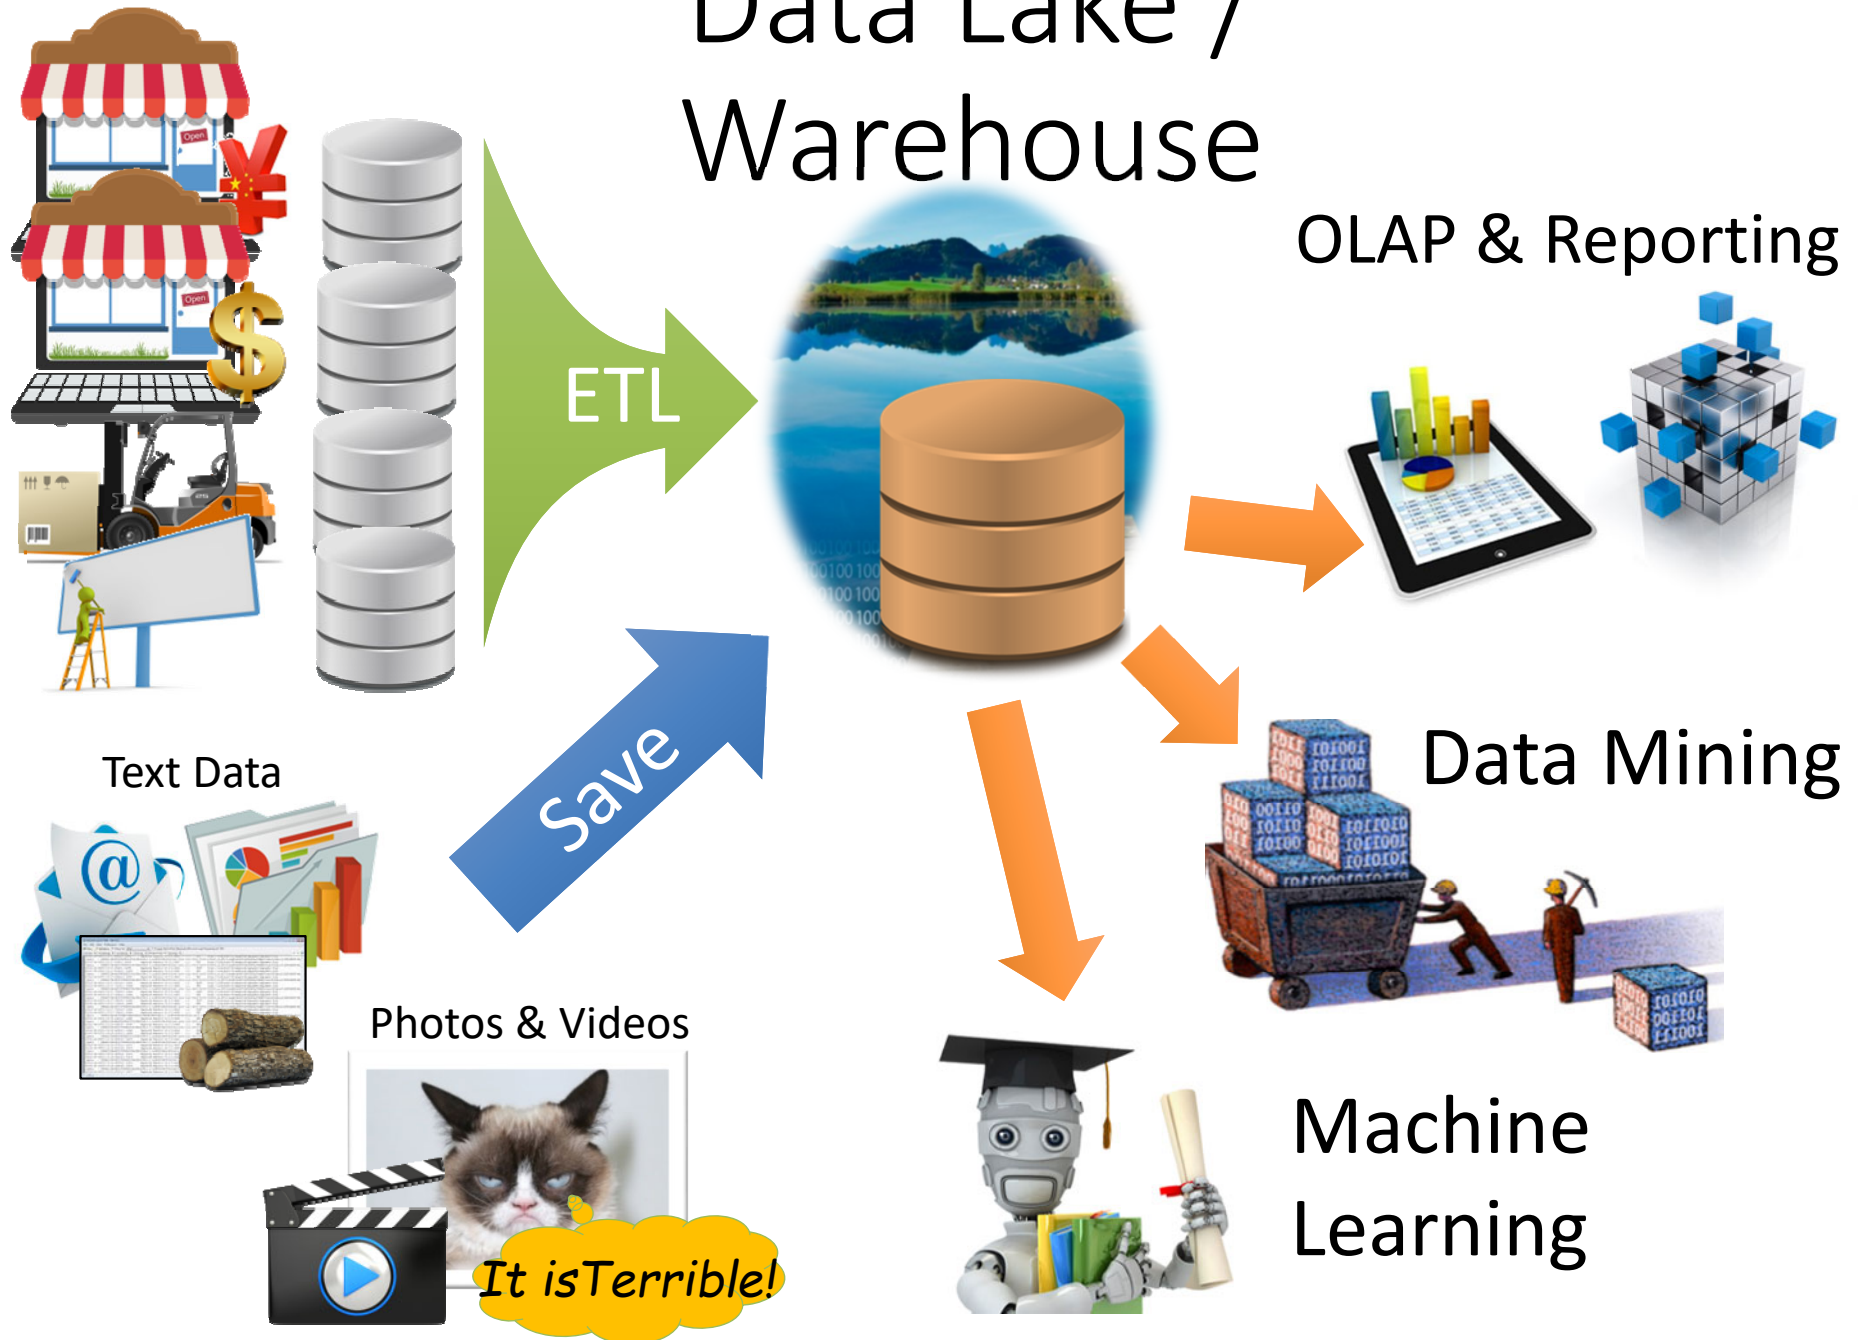
Course Textbook Chapters 26

Newer Material:

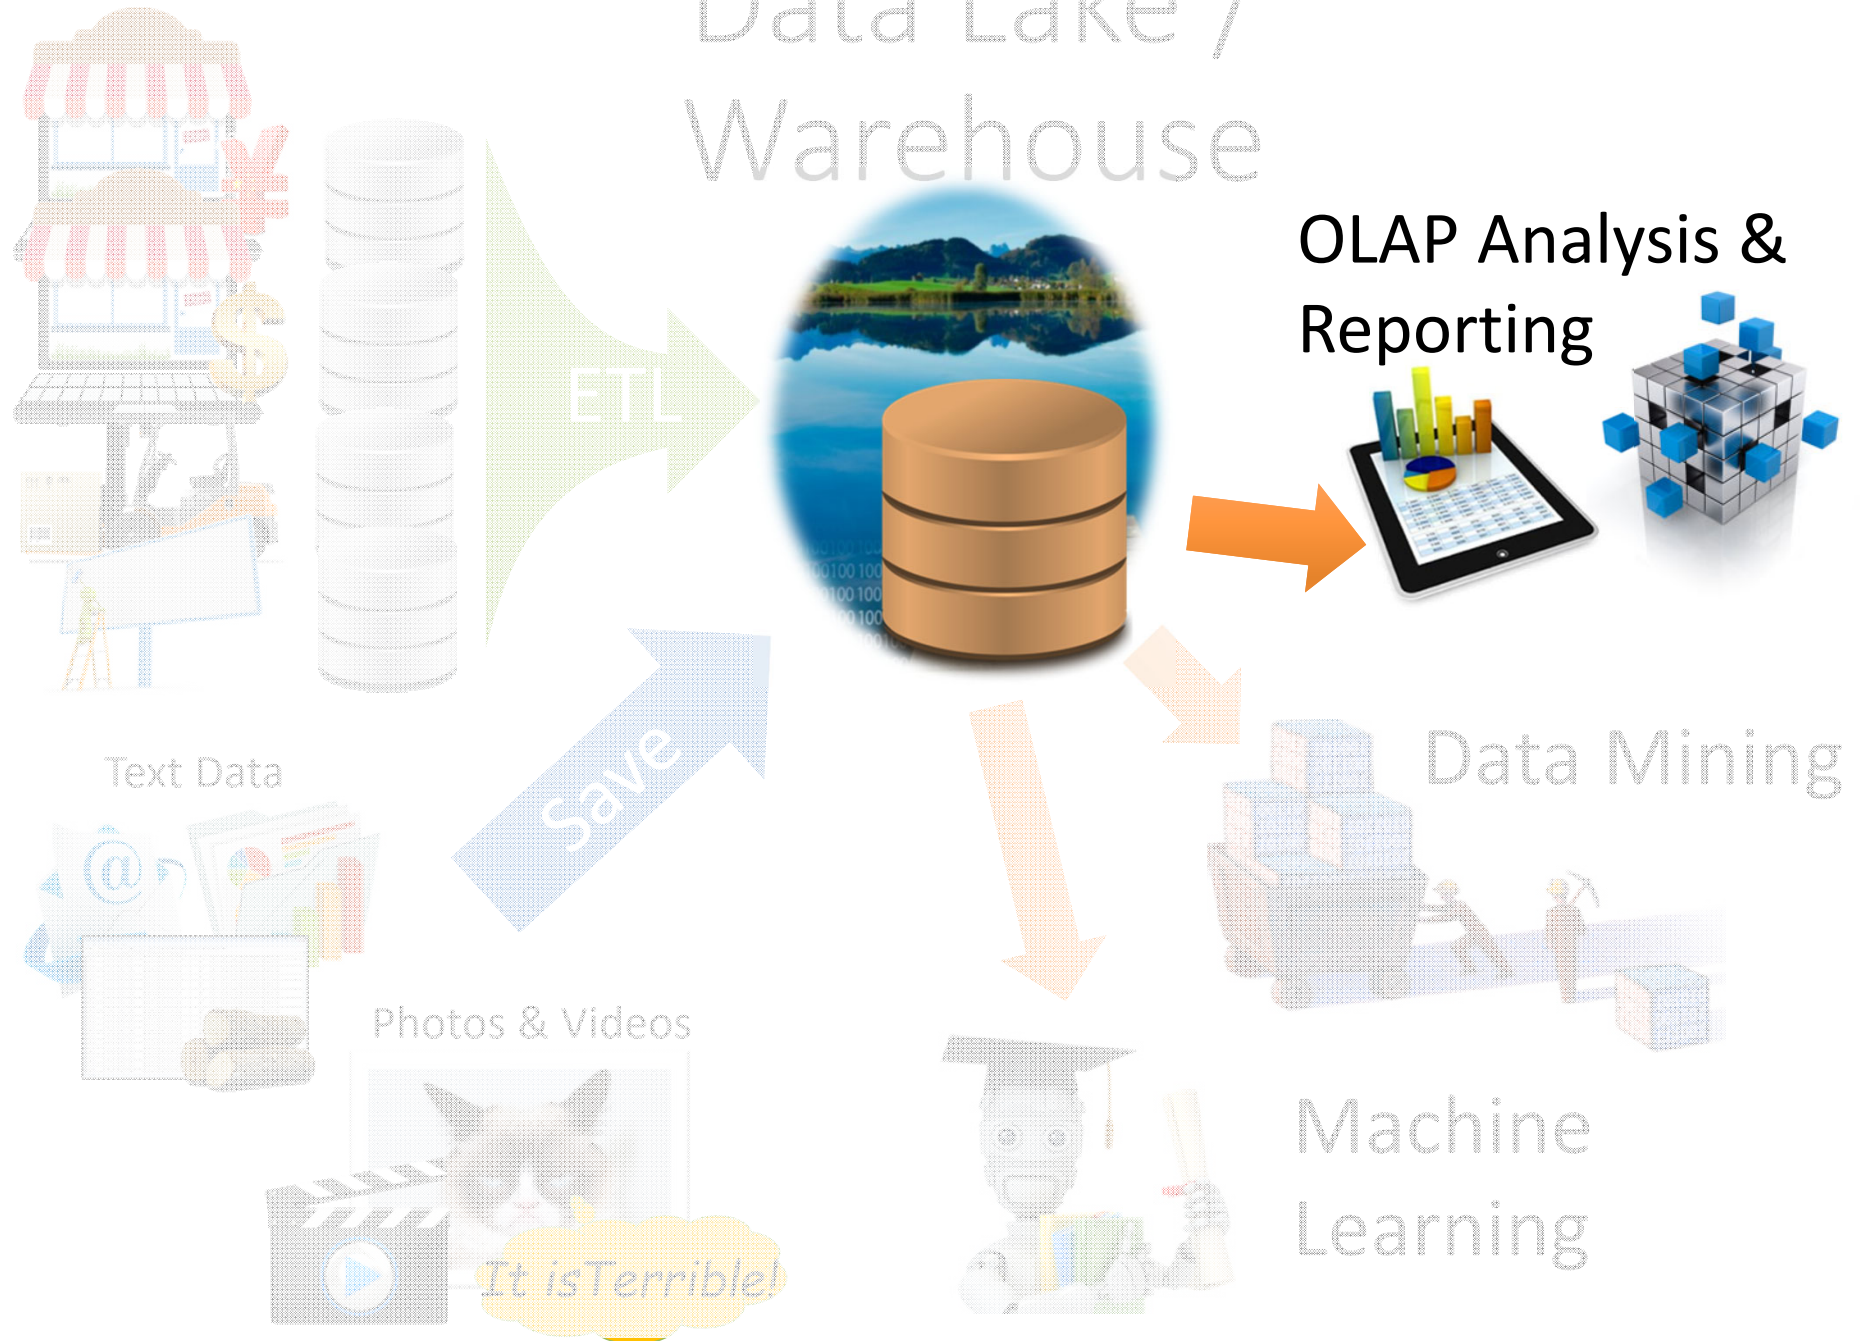
- Data Lake: [https://en.wikipedia.org/wiki/Data\\_lake](https://en.wikipedia.org/wiki/Data_lake)
- K-Means : [https://en.wikipedia.org/wiki/K-means\\_clustering](https://en.wikipedia.org/wiki/K-means_clustering)

Joseph E. Gonzalez  
[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)

# Data Lake / Warehouse



# Data Lake / Warehouse



# Multidimensional Data Model

*Sales* Fact Table

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

**Dimension  
Tables**


Products

pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2

➤ Multidimensional  
“Cube” of data

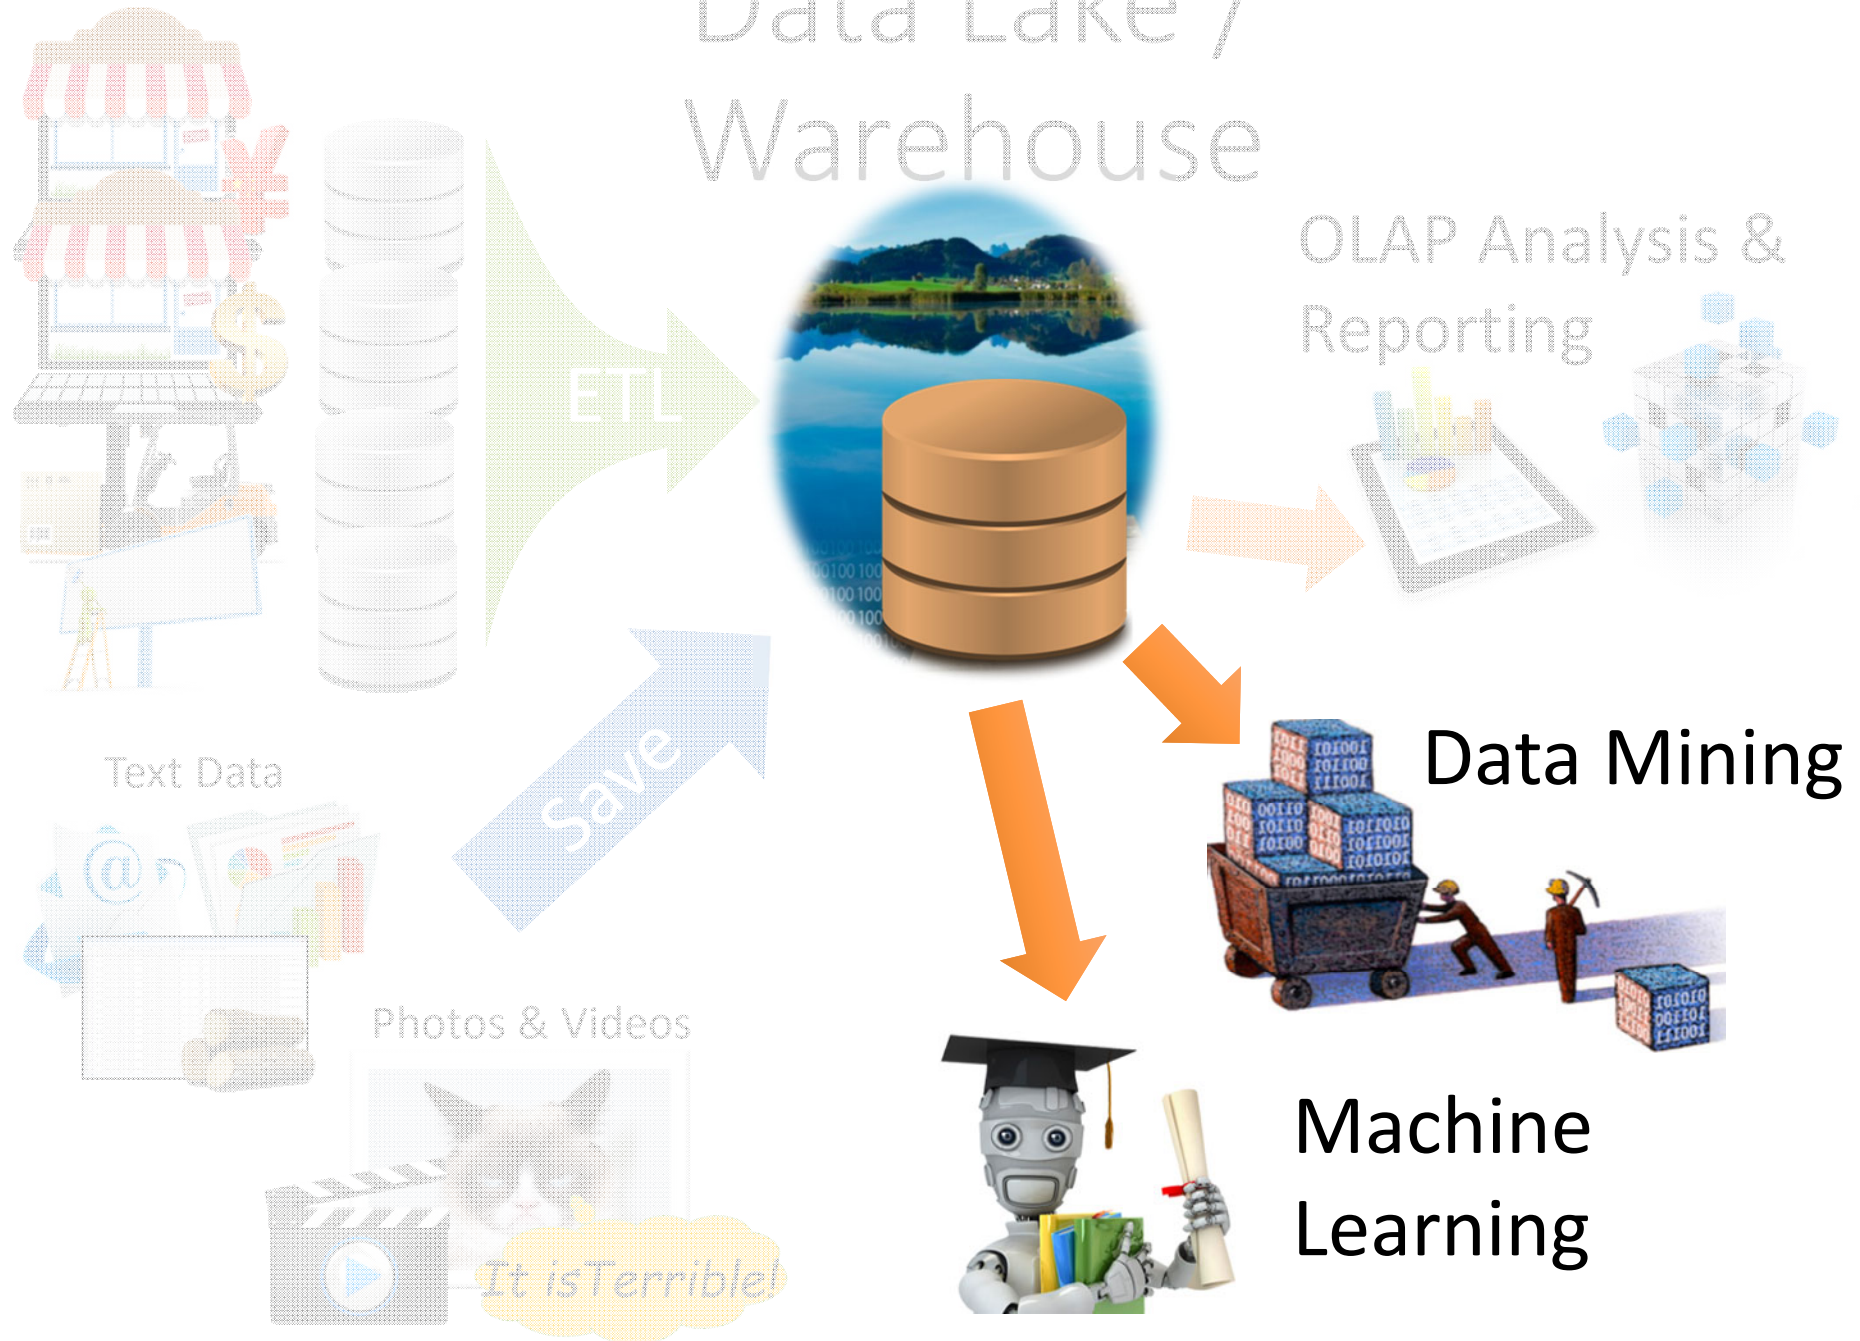
Time

timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

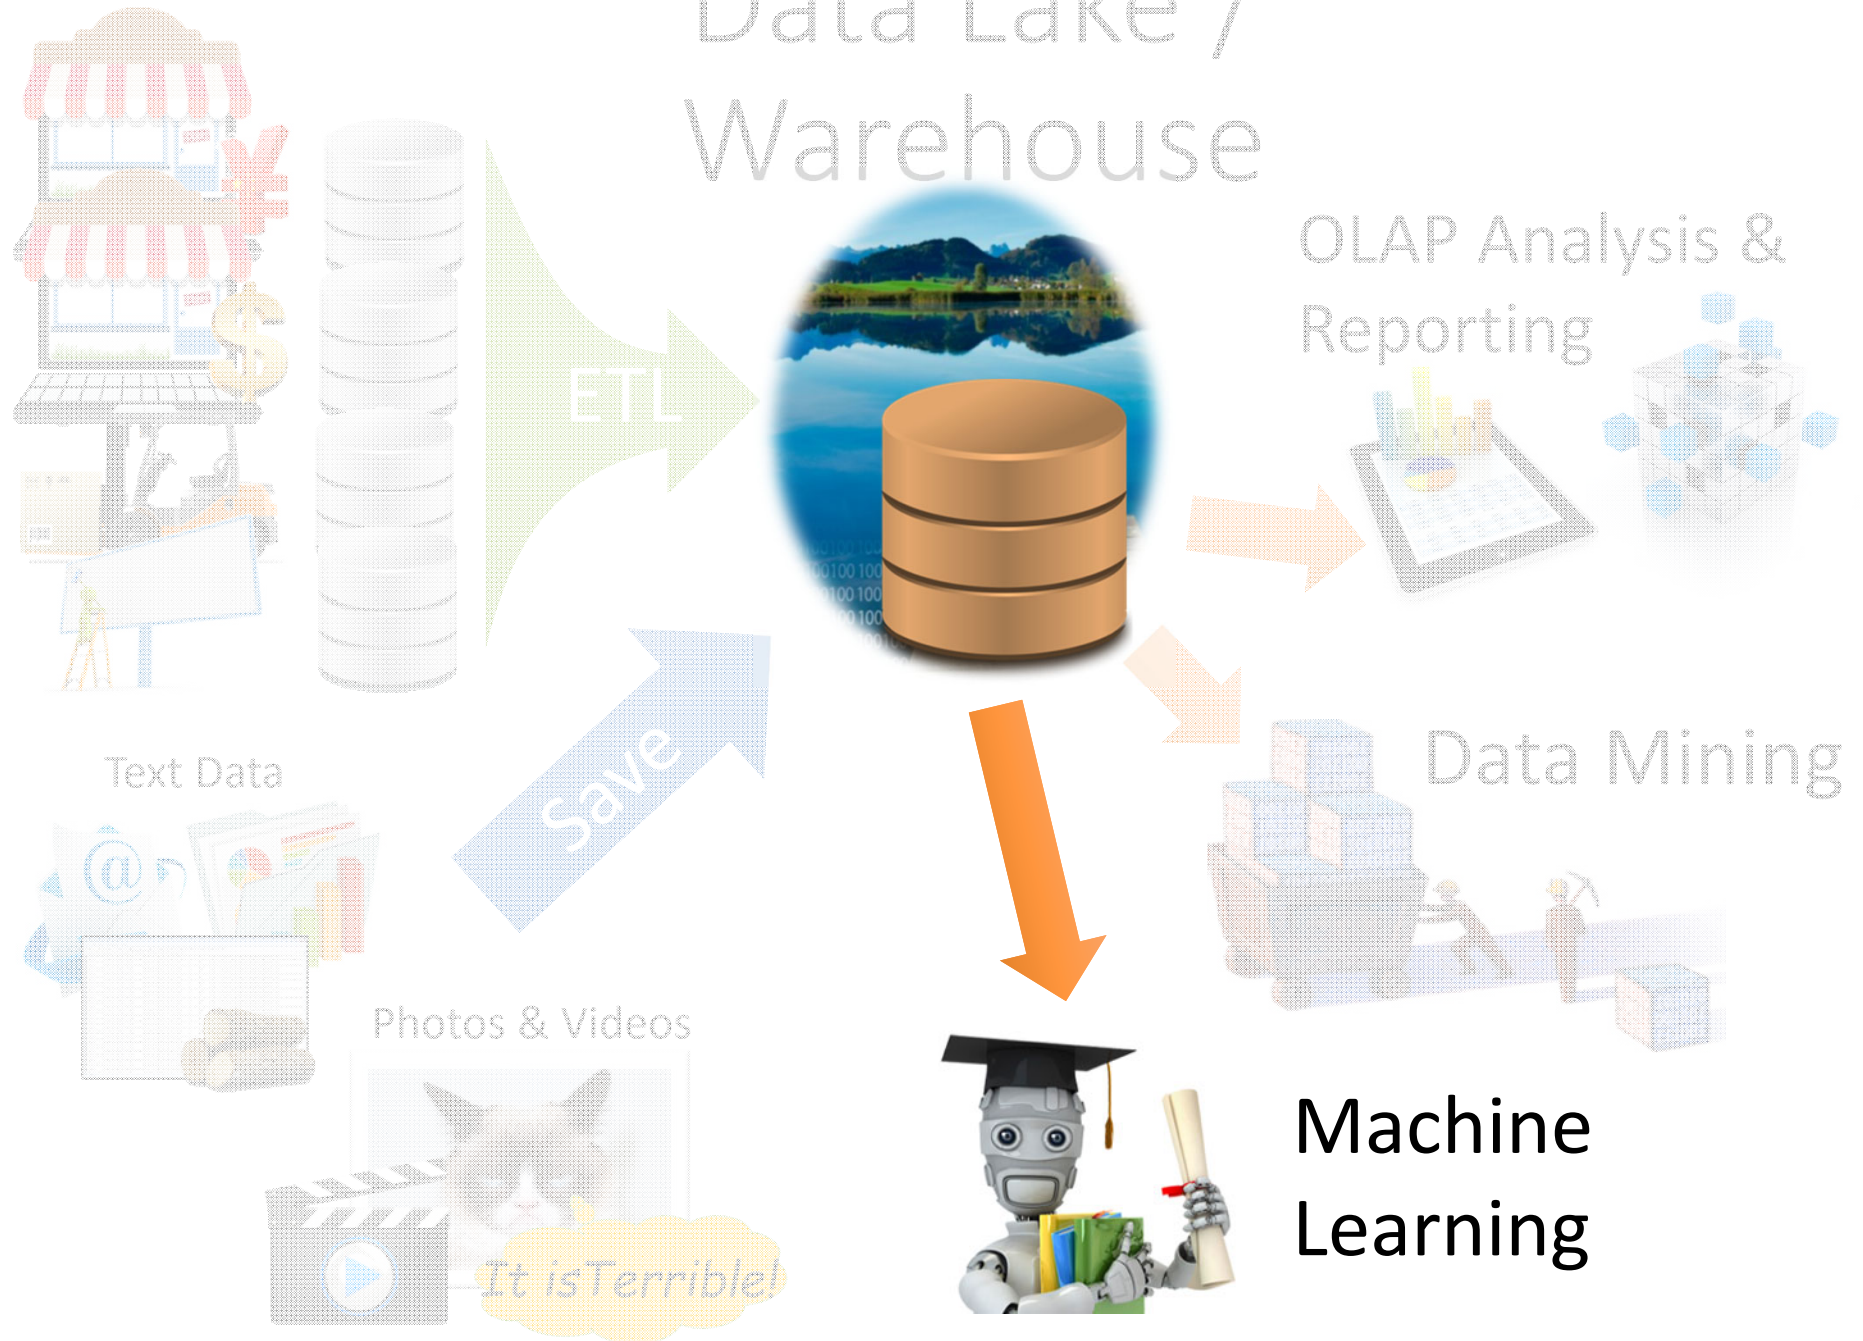




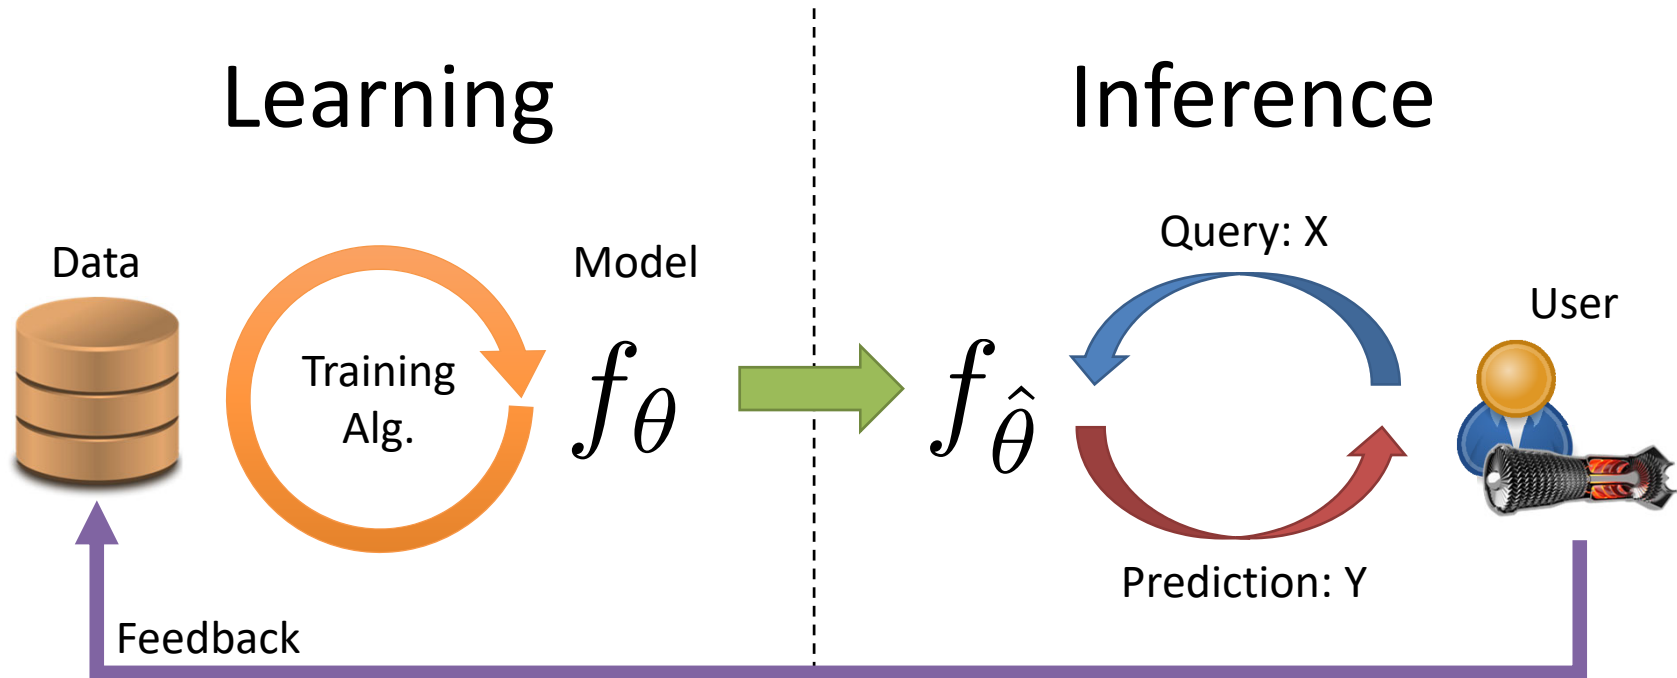
# Data Lake / Warehouse



# Data Lake / Warehouse



# Machine Learning Lifecycle



➤ Typically a time consuming iterative batch process

- Feature engineering
- Validation


➤ Focus is on making fast robust predictions

- Monitoring and tracking feedback
- Materialization + fast model inference






# Taxonomy of Machine Learning





# Clustering Images

- Given a collection of images cluster them into *meaningful groups*.



# Clustering Images

- Given a collection of images cluster them into meaningful groups.

**“Mountains”**



**“Forest”**



**“Beaches”**



# Clustering Images

- Given a collection of images cluster them into meaningful groups.



- **Unsupervised:** The labels of the groups are not given in the training data
- **Exploratory:** overlaps with data mining

# Clustering Images

- Given a collection of images cluster them into meaningful groups.

Simplified Illustration




Image Id	Average Red	Average Green
1	123	200
2	212	103
3	55	35

- How many clusters?
- Where are the clusters?

Features



# Clustering Images

- Given a collection of images cluster them into meaningful groups.

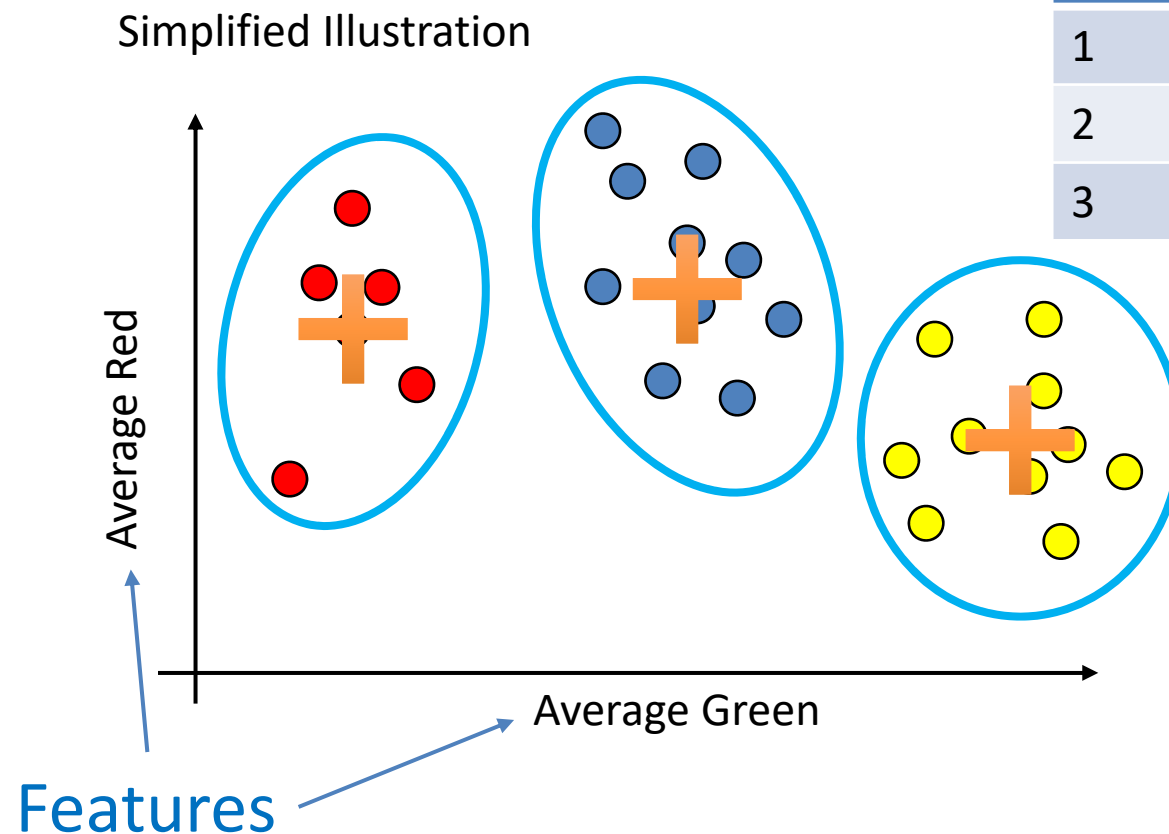



Image Id	Average Red	Average Green
1	123	200
2	212	103
3	55	35

- Where are the clusters?
- How many clusters?

# Clustering Images

- Given a collection of images cluster them into meaningful groups.

Image Id	Average Red	Average Green
1	123	200
2	212	103
3	55	35




What makes a good clustering?

- All points are near the cluster center
- Spread between clusters > spread within clusters


# Clustering Images

- Given a collection of images cluster them into meaningful groups.



# Clustering Images

- Given a collection of images cluster them into meaningful groups.





# Clustering Images

- Given a collection of images cluster them into meaningful groups.




Image Id	Average Red	Average Green
1	123	200
2	212	103
3	55	35

How do we automatically cluster data?

# How do we Compute a Clustering?

Many different clustering models and algorithms:

➤ Feature Based Clustering: *Points in  $R^d$*

- **K-Means:** EM on Symmetric Gaussians ← We will learn this one
- **Mixture Models:** Generalized k-means
- ...

➤ Spectral Methods: *Similarity Function Between Items*


- **Similarity based clustering:** *A and B are co-purchased*
- **Graph clustering:** *Cities based on road network*
- ...

➤ Hierarchical Clustering: *clustering nested items*

- **Latent Dirichlet Allocation:** *Documents based on words*
  - *Developed at Berkeley and widely used!*
- ...


# K-Means Clustering: *Intuition*

- Input K: The number of clusters to find
- Pick an initial set of points as cluster centers



# K-Means Clustering: *Intuition*


- For each data point find the cluster nearest center






# K-Means Clustering: *Intuition*

- For each data point find the cluster nearest center




# K-Means Clustering: *Intuition*

- Compute mean of points in each “cluster”




# K-Means Clustering: *Intuition*

- Adjust cluster centers to be the mean of the cluster



# K-Means Clustering: *Intuition*

- Adjust cluster centers to be the mean of the cluster






# K-Means Clustering: *Intuition*


➤ Improved?

➤ Repeat




# K-Means Clustering: *Intuition*

➤ Assign Points




# K-Means Clustering: *Intuition*

➤ Assign Points




# K-Means Clustering: *Intuition*

- Compute cluster means




# K-Means Clustering: *Intuition*

- Update cluster centers



# K-Means Clustering: *Intuition*

- Update cluster centers






# K-Means Clustering: *Intuition*

➤ Repeat?

- Yes to check that nothing changes → Converged!



# K-Means Algorithm: Details

centers  $\leftarrow$  pick k initial Centers

**while** (centers are changing) {

    // Compute the assignments (E-Step)

    asg  $\leftarrow$  [(x, nearest(centers, x)) for x in data]

What do we mean by “nearest”:

A: Euclidean Distance

$$\arg \min_{c \in \text{centers}} \|c - x\|_2^2 = \sum_{i=1}^d (c_i - x_i)^2$$

# K-Means Algorithm: Details

```
centers ← pick k initial Centers
```

```
while (centers are changing) {
```

```
    // Compute the assignments (E-Step)
```

```
    asg ← [(x, nearest(centers, x)) for x in data]
```

```
    // Compute the new centers (M-Step)
```

```
    for i in range(k):
```

```
        centers[i] =
```

```
            mean([x for (x, c) in asg if c == i])
```

```
}
```

Compute the  
“Expected” Assignment

Find centers that maximize the  
data “likelihood”

# K-Means Algorithm: Details

```
centers ← pick k initial Centers
```

```
while (centers are changing) {  
    // Compute the assignments (E-Step)  
    asg ← [(x, nearest(centers, x)) for x in data]  
  
    // Compute the new centers (M-Step)  
    for i in range(k):  
        centers[i] =  
            mean([x for (x, c) in asg if c == i])  
}
```

Guaranteed to  
converge!

*... to what?*

To a local  
optimum. ☹️

Depends on  
Initial Centers

# K-Means Algorithm: Details

centers  $\leftarrow$  pick k initial Centers

How do we pick initial centers? 

```
while (centers are changing) {  
    // Compute the assignments (E-Step)  
    asg  $\leftarrow$  [(x, nearest(centers, x)) for x in data]  
  
    // Compute the new centers (M-Step)  
    for i in range(k):  
        centers[i] =  
            mean([x for (x, c) in asg if c == i])  
}
```

Guaranteed to  
converge!

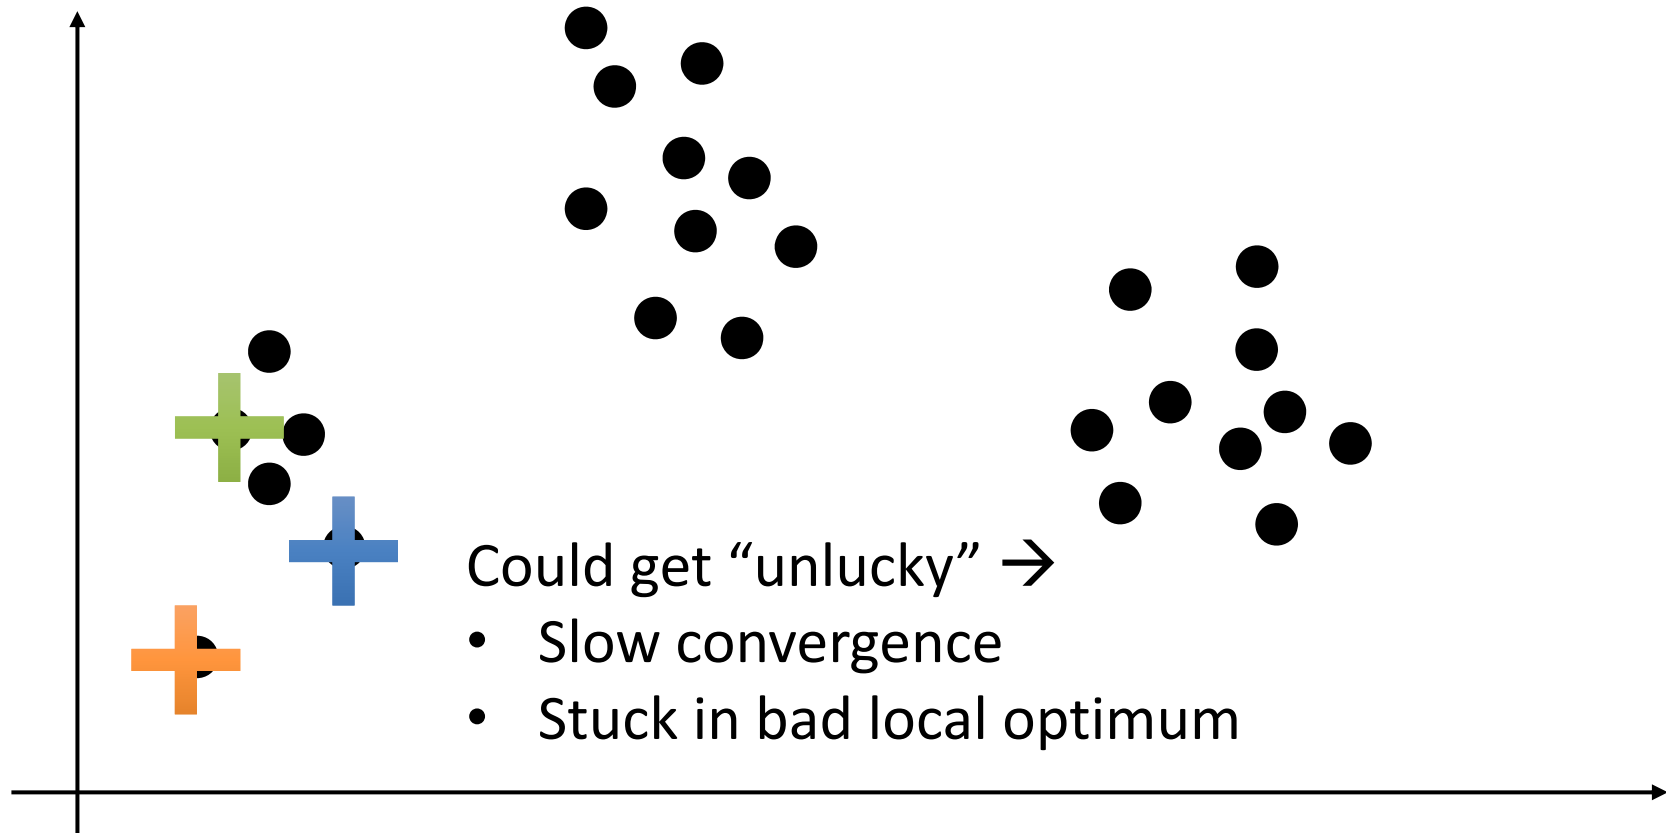
*... to what?*

To a local  
optimum. ☹

Depends on  
Initial Centers

# Picking the Initial Centers

- **Simple Strategy:** select  $k$  points at random
  - What could go wrong?






# Picking the Initial Centers

## ➤ **Better Strategy:** kmeans++

- Randomized approx. algorithm
- Intuition select points that are not near existing centers



# K-Means++ Algorithm

```
centers ← set(randomly select a single point)
```

```
while len(centers) < k:
```

```
    # Compute the distance of each point
```

```
    # to its nearest center  $dSq = d^2$ 
```

```
     $dSq \leftarrow [(x, \text{dist\_to\_nearest}(\text{centers}, x)^2) \text{ for } x \text{ in data}]$ 
```

```
    # Sample a new point with probability
```


```
    # proportional to  $dSq$ 
```

```
     $c \leftarrow \text{sample\_one}(\text{data}, \text{prob} = dSq / \text{sum}(dSq))$ 
```

```
    # Update the clusters
```

```
    centers.add(c)
```

# How do we choose K?




## ➤ Basic Elbow Method (Easy and what you do in HW)


- Try range of K-values and plot average distance to centers

## ➤ Cross-Validation (Better)

- Repeatedly split the data into training and validation datasets
- Cluster the training dataset
- Measure Avg. Dist. To Centers on validation data



# K-Means +



How do we run k-means on the  
data warehouse / data lake?

# Interacting With the Data

Good for smaller datasets

- Faster more natural interaction
- Lots of tools!


Compute  
Locally

$$\Sigma = \bigoplus_{r \in \text{Data}} f_{\theta}(r)$$




Request Data Sample

Sample of Data



Can we send the  
computation to  
the data?

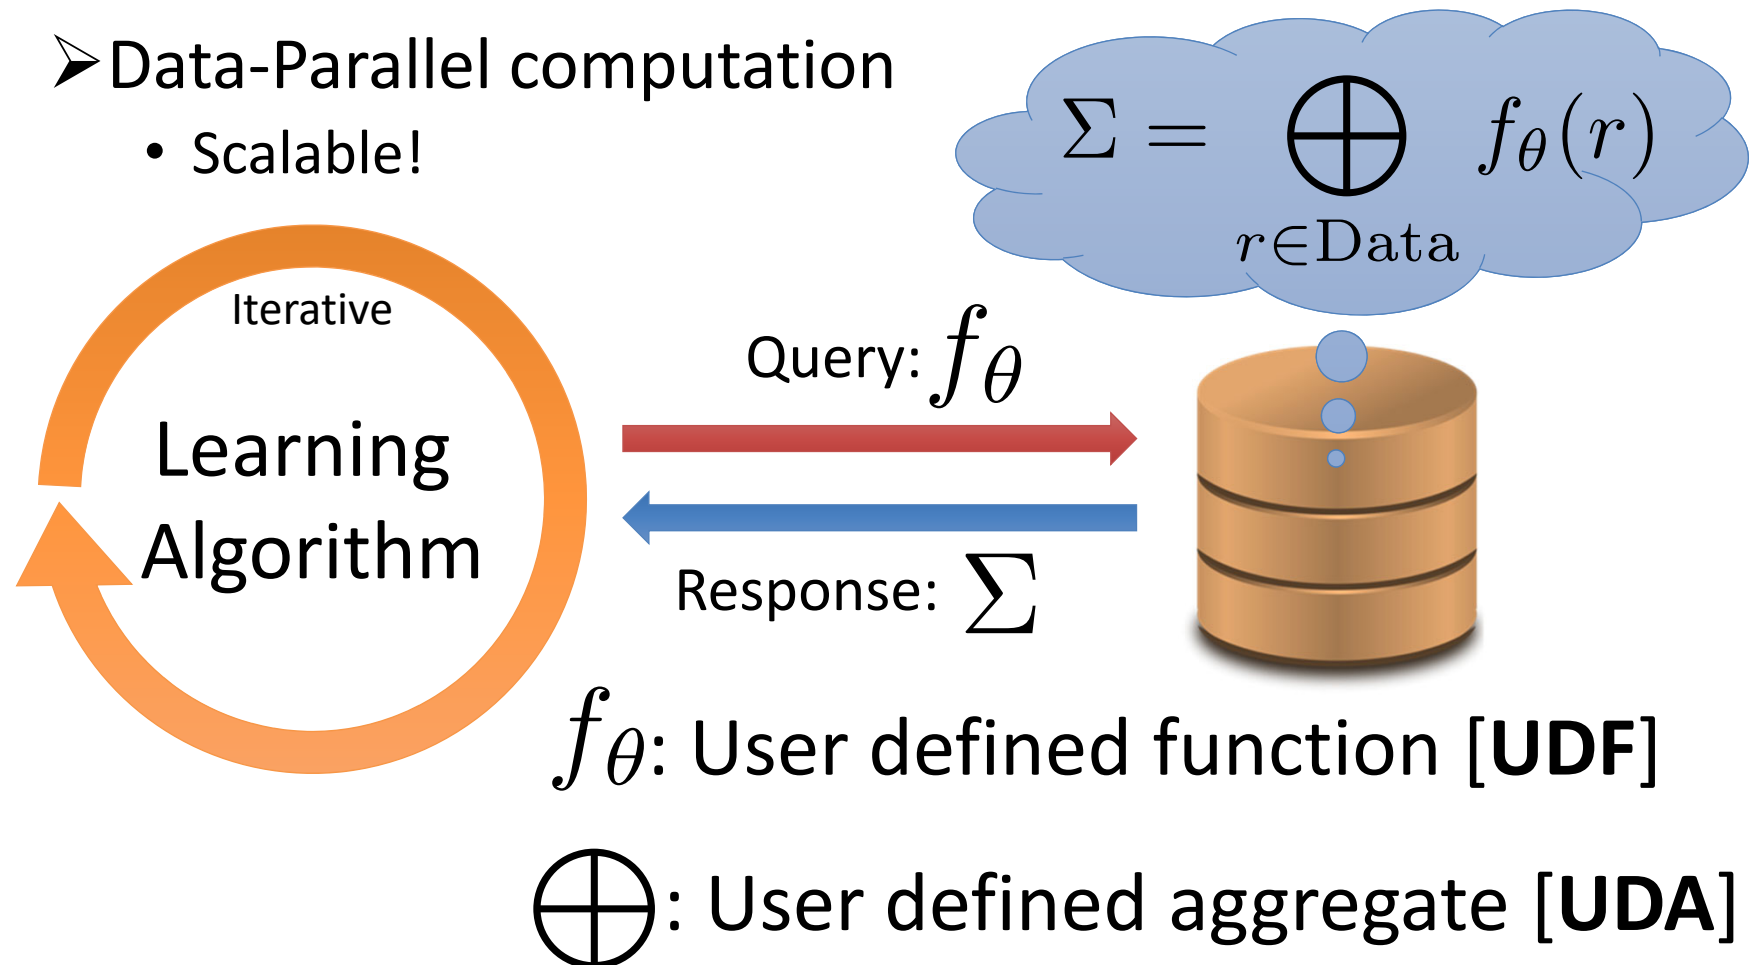
**Yes!**



# Statistical Query Pattern

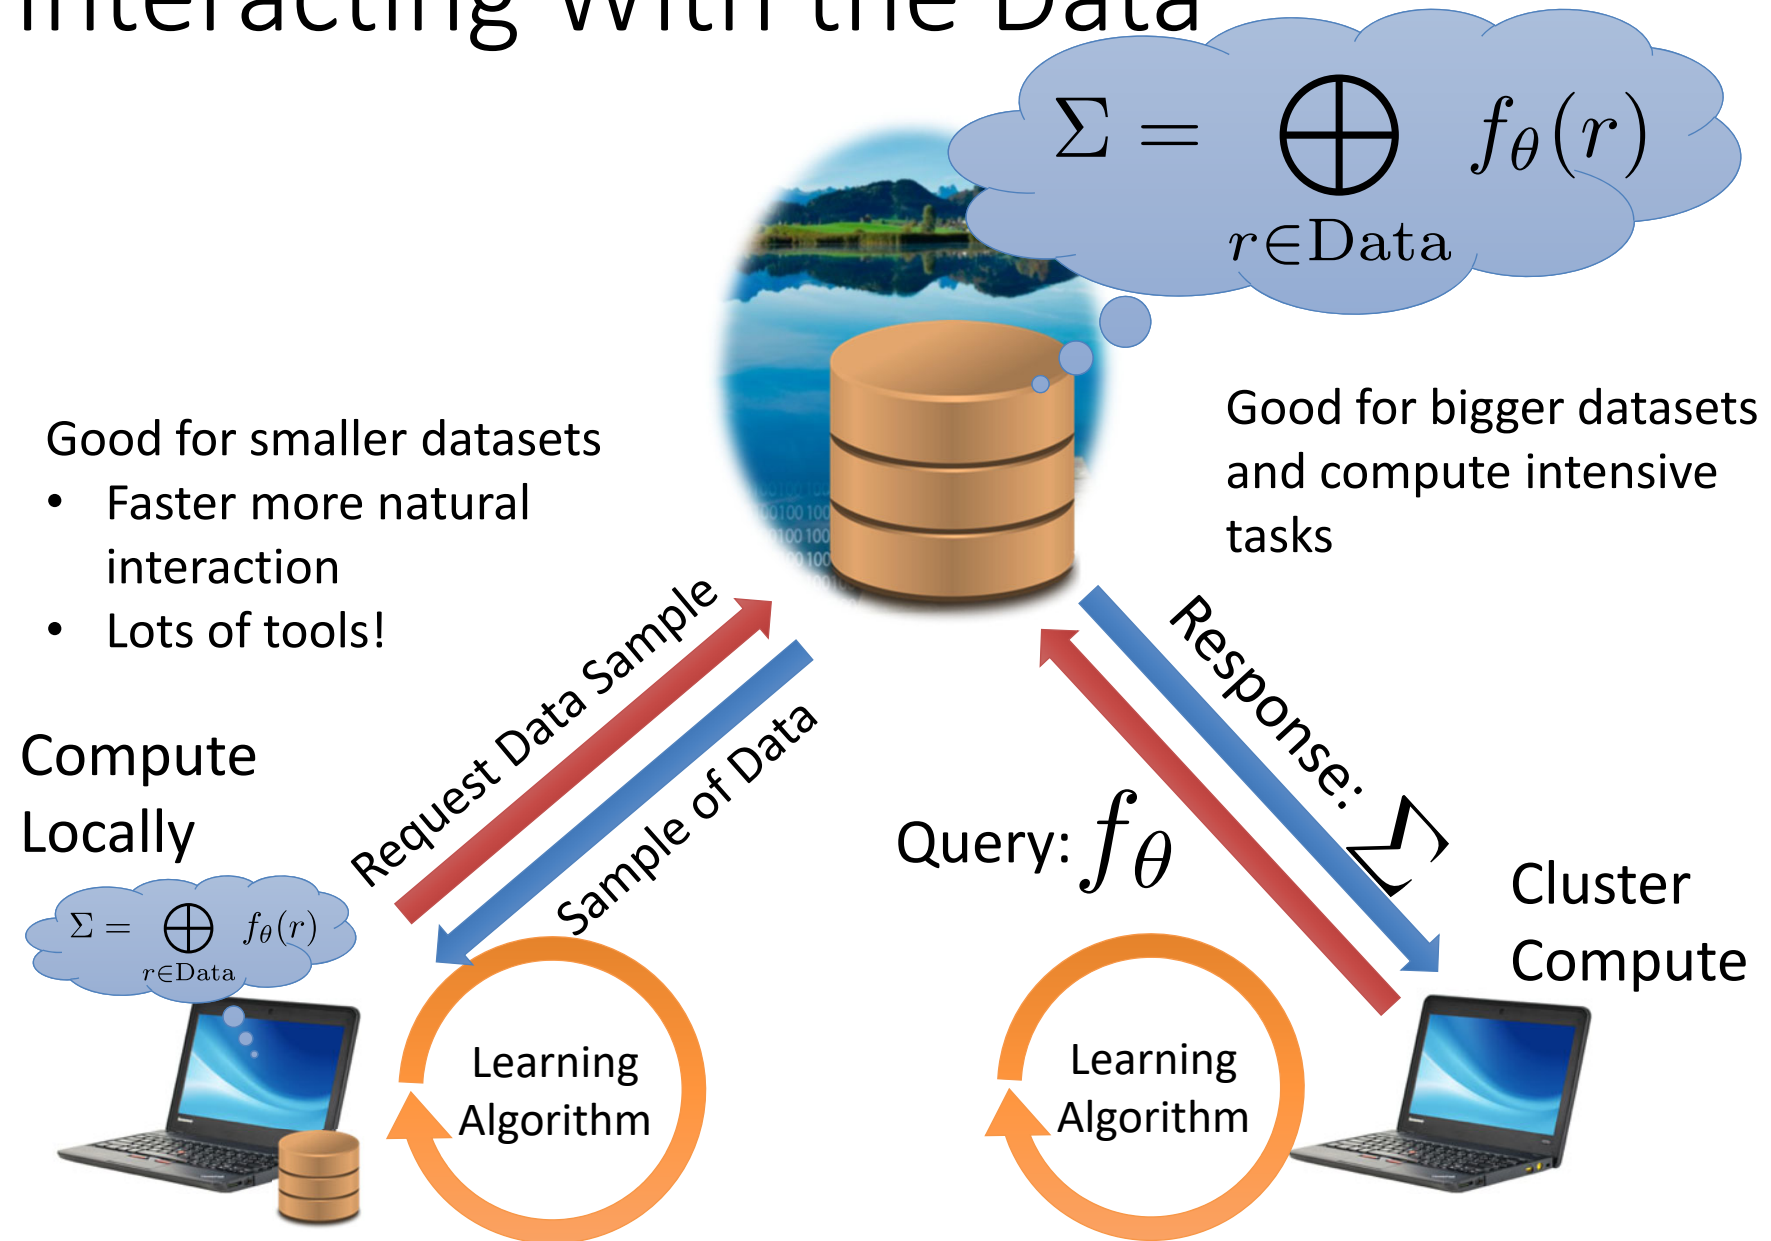
## Common Machine Learning Pattern

- Computing aggregates of user defined functions
- Data-Parallel computation
  - Scalable!





# Interacting With the Data



# Can we express K-Means in the Statistical Query Pattern?

```
centers ← pick k initial Centers
while (centers are changing):
    // Compute the assignments (E-Step)
    asg ← [(x, nearest(centers, x)) for x in data]
    for i in range(k): // Compute the new centers (M-Step)
        centers[i] = mean([x for (x, c) in asg if c == i])
```

Query returns all  
the data ...

Merge with M-Step  
→  
Statistical Query  
Pattern

```
centers ← pick k initial Centers
while (centers are changing):
    for i in range(k):
        new_centers[i] =
            mean([x for x in data if nearest(centers, x) == i])
    centers = new_centers
```

# Can we express K-Means in the Statistical Query Pattern?

```
centers ← pick k initial Centers
while (centers are changing):
    for i in range(k):
        new_centers[i] =
            mean([x for x in data if nearest(centers, x) == i])
    centers = new_centers
```

Group by query:

```
SELECT nearest_UDF(centers, x) AS cid, mean_UDA(x)
FROM data GROUPBY cid
```

➤ UDFs and UDAs are implemented varies across systems.

You can implement this in pure SQL (for two dimensions):

```
CREATE TABLE points (x double precision, y double precision);  
COPY points FROM '~/toy_data.csv' DELIMITER ',' CSV;
```

```
CREATE TABLE centers (  
    id INTEGER,  
    x double precision, y double precision,  
    ver INTEGER);  
INSERT INTO centers VALUES  
    (0, 0.1, 2.3, 0), (1, -0.2, 1.1, 0), (2, 1.4, -2.2, 0), (3, -.2, -3.0, 0);
```

```
CREATE TEMP VIEW maxVer AS  
SELECT max(ver) FROM centers;
```

```
CREATE TEMP VIEW dist AS
```

```
SELECT p.x as x, p.y as y, MIN((p.x - c.x) * (p.x - c.x) + (p.y - c.y) * (p.y - c.y))  
as min_d
```

```
FROM points as p, centers as c
```

```
WHERE (c.ver) in (select * from maxVer)
```

```
GROUP BY p.x, p.y;
```

# Repeatedly invoke the following until convergence

```
INSERT INTO centers
```

```
SELECT c.id, AVG(d.x) as x, AVG(d.y) AS y, max(c.ver) + 1 as ver
```

```
FROM dist as d, centers as c
```

```
WHERE (c.ver) in (select * from maxVer)
```

```
AND d.min_d >= (d.x - c.x) * (d.x - c.x) + (d.y - c.y) * (d.y - c.y)
```

```
GROUP BY c.id;
```

# K-Means in Map-Reduce

## ➤ **MapFunction**(*old\_centers*, *x*)

- Compute the index of the nearest old center
- Return (**key** = *nearest\_centers*, **value** = (*x*, 1))

## ➤ **ReduceFunction** combines values and counts

- For each cluster center (Group By)

## ➤ Data system returns aggregate statistics:

$$s_i = \sum_{x \in \text{Cluster } i} x_i \quad \text{and} \quad n_i = \sum_{x \in \text{Cluster } i} 1$$

## ➤ ML algorithm computes new centers: $\mu_i = s_i / n_i$

# Can we express K-Means++ in the Statistical Query Pattern?

- Yes, however there is a better version: **K-Means||**
  - More complex but much faster
- Or you can parallelize **K-Means++** directly
  - Requires more passes
- Challenging Step?
  - Parallel weighted sampling:

```
sample_one(data, prob = dSq / sum(dSq))
```

- How do you select one point uniformly at random?




# Res-A: weighted reservoir sampling


- **Goal:** Sample  $k$  records from a stream where record  $i$  is included in the sample with probability **proportional to  $w_i$**

How would we normally sample  $k$  records?


Weights:



Draw a random number uniformly between **0** and **3.7**




Sample the corresponding record and remove the weight.



Draw a random number uniformly between **0** and **3.1**

Weights:



Samples

Rec. 2


Rec. 1

Rec. 3

Rec. 4

Rec. 5

We want to do this in **one pass** *without* ever knowing the **sum** of the weights!



# Res-A: weighted reservoir sampling

➤ **Goal:** Sample  $k$  records from a stream where record  $i$  is included in the sample with probability proportional to  $w_i$

➤ **Algorithm:**

- For each record  $i$  draw a uniform random number:

$$u_i \sim \text{Unif}(0, 1)$$

- Select the top- $k$  records ordered by:  $u_i^{1/w_i}$

➤ **Common ML Pattern?**

- **Query Function:**  $[pow(rand(), 1 / record.w), record]$
- **Agg. Function:** *top-k heap*

# Illustrating Res-A Algorithm

Weights:		Uniform Random Number	$u_i^{1/w_i}$
1.0	Rec. 1	0.31	0.31
2.0	Rec. 2	0.20	0.45
1.4	Rec. 3	0.45	0.56
2.5	Rec. 4	0.14	0.46
0.7	Rec. 5	0.55	0.43
3.0	Rec. 6	0.69	0.88

Top 2

The diagram illustrates the Res-A algorithm by comparing a set of weights with a set of uniform random numbers to select the top two recommendations. The weights are represented by red bars of varying lengths, and the recommendations are labeled Rec. 1 through Rec. 6. The uniform random numbers are shown in purple rounded rectangles, and the weighted values  $u_i^{1/w_i}$  are shown in orange rounded rectangles. An arrow points to the top two values, 0.56 and 0.88, which are labeled 'Top 2'.



# Basic Analysis Behind Res-A

➤ Define the random variable:  $X_i = u_i^{1/w_i}$

➤ Then:

$$\mathbf{P}(X_i < \alpha) = \mathbf{P}(u_i^{1/w_i} < \alpha) = \mathbf{P}(u_i < \alpha^{w_i}) = \alpha^{w_i}$$

$$\mathbf{p}(X_i = \alpha) = w_i \alpha^{w_i - 1}$$

Derivative of CDF  $\rightarrow$  PDF

➤ Suppose we want to pick just one element (k=1)

- Probability of selecting  $X_i$  is:

$$\int_0^1 \mathbf{p}(X_i = \alpha) \prod_{j \neq i} \mathbf{P}(X_j < \alpha) d\alpha = \int_0^1 (w_i \alpha^{w_i - 1}) \prod_{j \neq i} \alpha^{w_j} d\alpha$$

$$= \frac{w_i}{\sum_j w_j}$$

We won't test you  
on this derivation

## *People who like Res-A also like...*

### ➤ Algorithm R

- Another reservoir filtering algorithm (recitation?)

### ➤ Bloom Filters

- Efficient set membership with limited memory

### ➤ Count-Min

- Efficient key-counting with limited memory

### ➤ Heavy Hitters Sketch

- Top-k Elements with limited memory


# Algebra details from integration:

$$\begin{aligned}\int_0^1 \mathbf{p}(X_i = \alpha) \prod_{j \neq i} \mathbf{P}(X_j < \alpha) d\alpha &= \int_0^1 (w_i \alpha^{w_i-1}) \prod_{j \neq i} \alpha^{w_j} d\alpha \\&= w_i \int_0^1 (\alpha^{w_i-1}) \alpha^{\sum_{j \neq i} w_j} d\alpha \\&= w_i \int_0^1 \alpha^{-1 + \sum_j w_j} d\alpha \\&= \frac{w_i}{\sum_j w_j} \alpha^{\sum_j w_j} \bigg|_{\alpha=0}^{\alpha=1} \\&= \frac{w_i}{\sum_j w_j}\end{aligned}$$

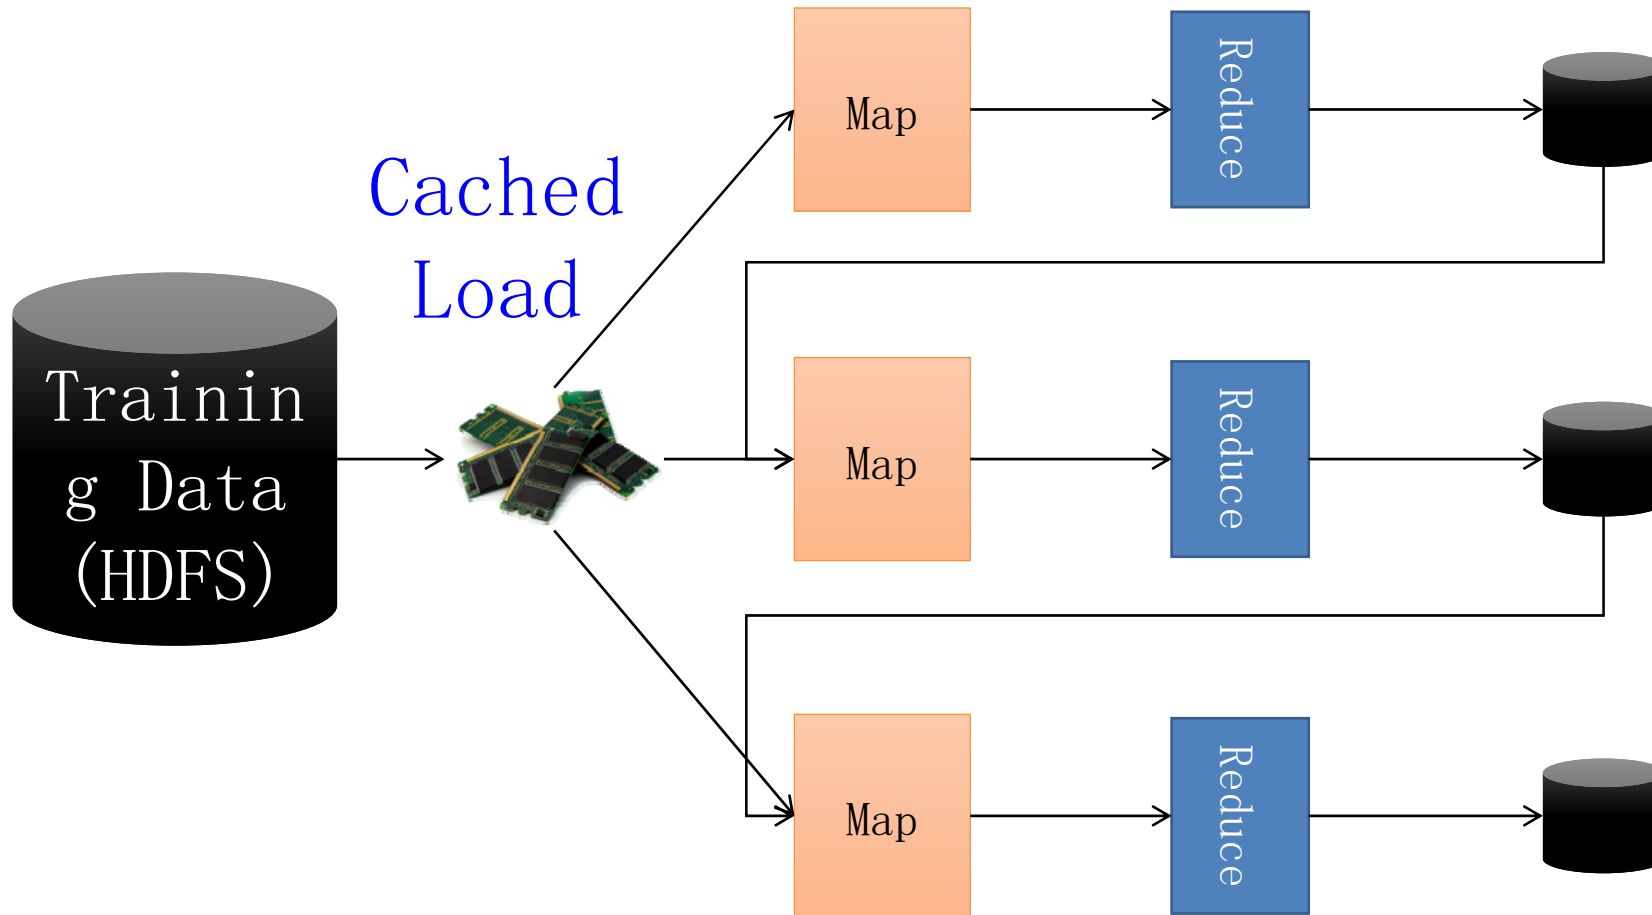
# Implementation Details: Statistical Query Pattern

- **Iterative ML** ➔ Data caching is important
  - Motivation behind Spark project

# Map Reduce Dataflow View

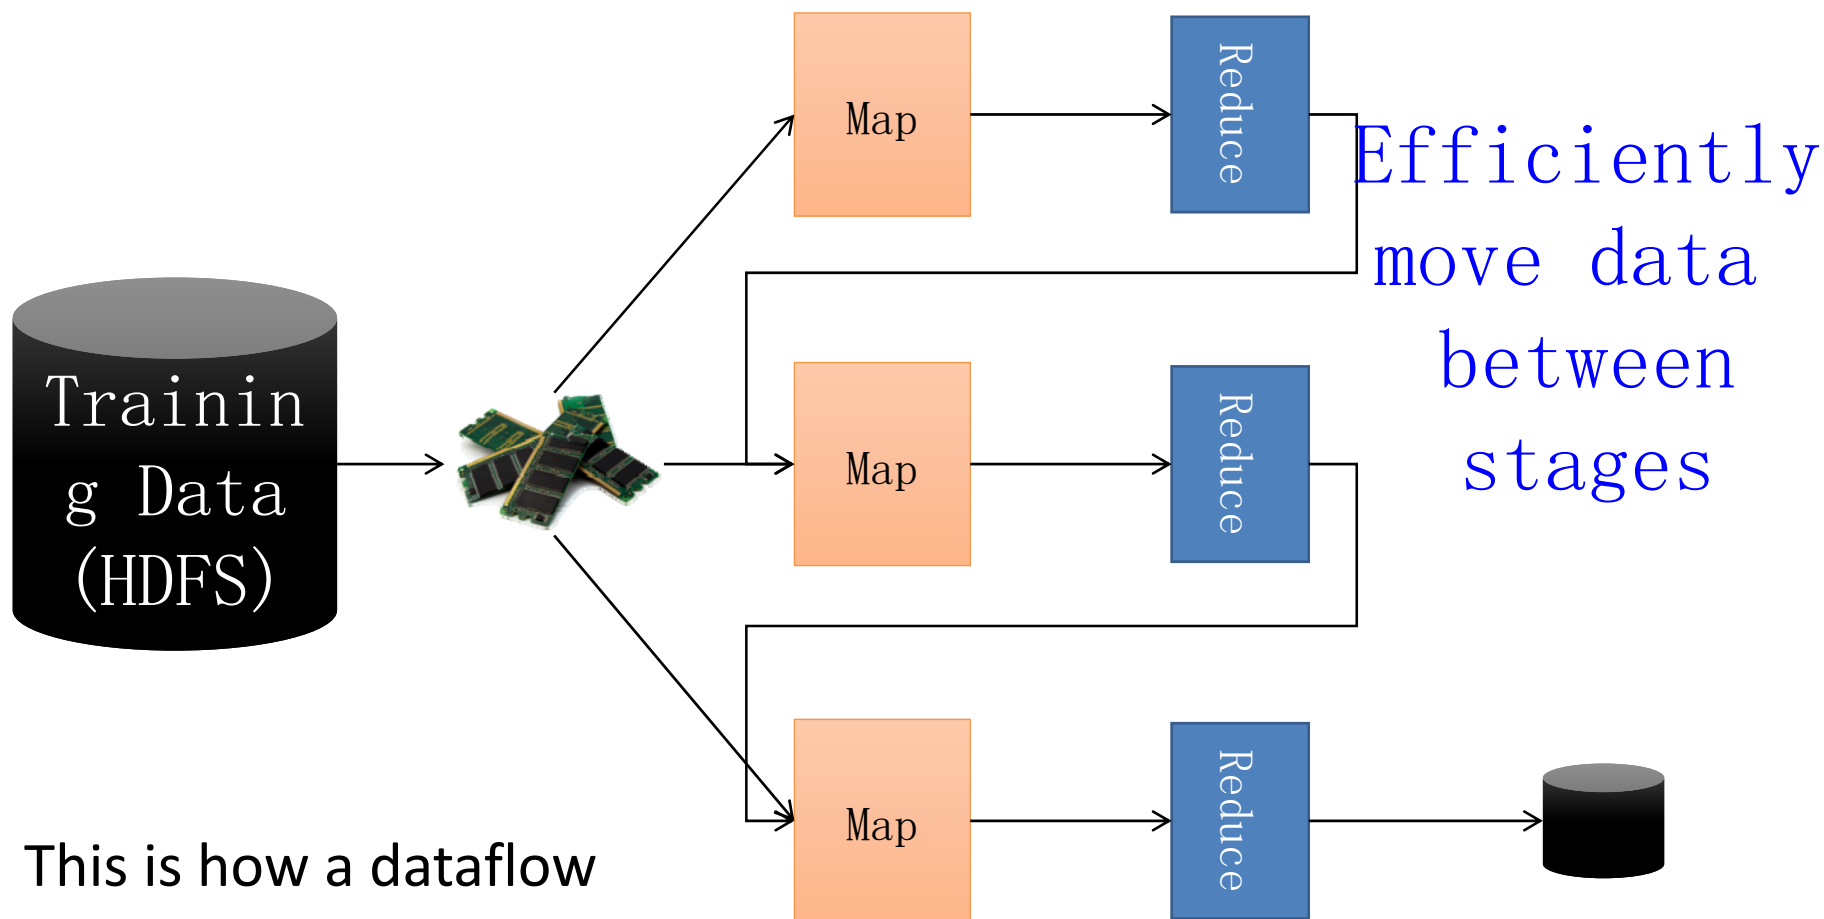


# Spark Opt. Dataflow



10–100× faster than network and disk

# Spark Opt. Dataflow View

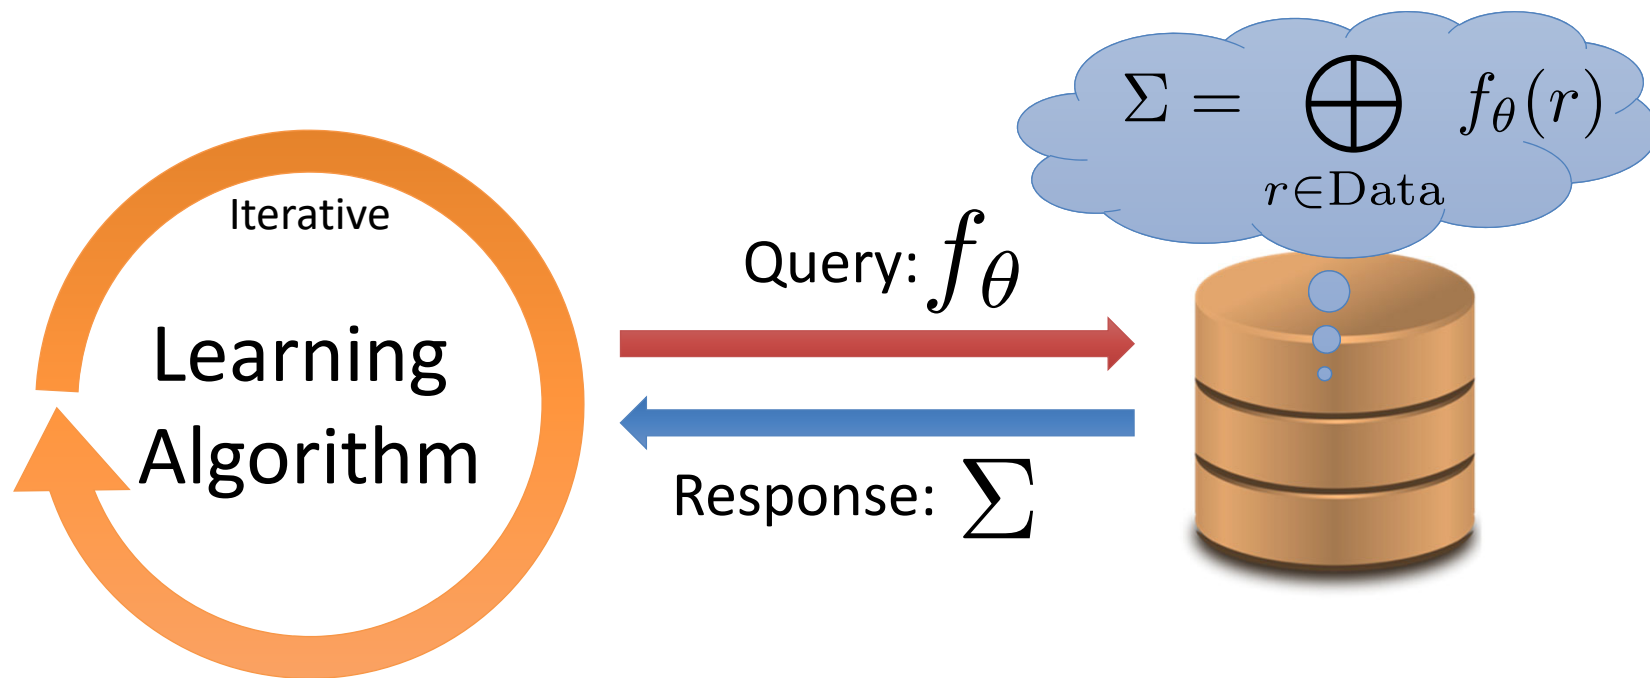


This is how a dataflow system should behave!

- What happened to map-reduce?




# Implementation Details: Common Machine Learning Pattern

- Iterative ML ➔ Data caching is important
  - Motivation behind Spark project
- Need to watch out for large  $\theta$  and  $\Sigma$






# Summary of Clustering



```
SELECT nearest_UDF(centers, x) AS cid, mean_UDA(x)
FROM data GROUPBY cid
```





# Taxonomy of Machine Learning

