

CS150A Homework1 Solutions

I. MULTI-RELATION QUERIES [40 points]

Consider the following relations representing student information at UIUC:

Mentorship(mentee_sid, mentor_sid)

Study(sid, credits)

Enrollment(did, sid)

Student(sid, street, city)

- A tuple (M1, M2) in Mentorship specifies that M2 is a mentor of another student M1.
- A tuple (S, C) in Study specifies that the student S has taken C credits.
- A tuple in Enrollment (D, S) specifies that student S is enrolled in department D.
- A (ST, S, C) in Student specifies that student ST lives on street S in city C.

1. Find all students who live in the same city and on the same street as their mentor.

Solution:

```
SELECT m1.mentee_sid
FROM Mentorship as m1, Student as s1, Student as s2
WHERE m1.mentee_sid = s1 . sid
AND m1.mentor_sid = s2 . sid
AND s1 . street = s2 . street
AND s1 . city = s2 . city;
```

2. Find all students(i.e., distinct sid) who have taken more credits than the average credits of all of the students of their department.

Solution:

```
SELECT s1 . sid
FROM Study as s1, Enrollment as e1
WHERE e1 . sid = s1 . sid
AND s1 . credits > ( SELECT AVG( s2 . credits )
FROM Study as s2, Enrollment as e2
WHERE s2 . sid = e2 . sid
AND e1 . did = e2 . did
);
```

II. Storage: Disk, Files, Buffers [30 points]

1. Write down the letters of true statements *in alphabetical order*. (If none are true, write \emptyset .)
 - A. When querying for a 16 byte record, exactly 16 bytes of data is read from disk. **False, a page's worth of data is always read from disk.**
 - B. Writing to an SSD drive is more costly than reading from an SSD drive. **True, writes to an SSD can involve reorganization to avoid uneven wear and tear.**
 - C. In a heap file, all pages must be filled to capacity except the last page. **False, there is no such requirement. With clustered indices, it is common to keep heap file pages 2/3 full to accommodate inserts.**
 - D. If the file size is smaller than the number of buffer frames, a sequential scan of the file using either MRU or LRU (starting with an empty buffer pool) will have the same hit rate. **True, our eviction policy doesn't matter because the file fits in memory.**
 - E. Assuming integers take 4 bytes and pointers take 4 bytes, a slot directory that is 256 bytes can address 64 records in a page. **False, as shown in lecture, an entry in slot directory is 8 bytes, because a single entry consists of both a pointer and an integer (length).**
 - F. In a page containing fixed-length records with no nullable fields, the size of the bitmap never changes. **True, the size of the records is fixed, so the number we can fit on a page is fixed.**
2. Assume we have 4 empty buffer frames and the following access pattern, in which pages are immediately unpinned.

T A M E T E A M M A T E M E A T L I D

Use the replacement policy listed, and list the four pages in the buffer pool at the end, *in alphabetical order*. Hint: you don't need to draw a big chart for every access — look for patterns.

- 1) MRU. **ADEM**
 - 2) LRU. **DILT**
 - 3) Clock. (Assume the clock hand starts on the first buffer and does not move unless a page needs to be replaced.). **DEIL**
-

III. Indexes and B+ Trees [30 points]

Note: for B+ tree page splits with an odd number of items, assume that the majority of the items is placed on the right-hand page after the split.

1. Alphabetically, write down the letters of statements that apply (or write \emptyset .)
 - A. All internal keys in a B+ tree also appear in its leaf nodes. **False, a leaf node which has been copied up can be deleted.**
 - B. The height of a B+ tree increases whenever any node splits. **False, height increases when root node splits.**
 - C. An ISAM index is similar to a B+ tree, but does not allow for insertion of new values. **False.**
 - D. The column(s) we select for our index key must have a unique value for every row in the table. **False.**

 2. Assume we are trying to construct a B+ Tree of order 2 (max fanout of 5). Each leaf node can hold up to 4 entries. We insert a total of 16 unique keys via bulk loading, with a fill factor of $\frac{3}{4}$.
 - 1) How many leaf nodes will there be? **6**
 - 2) How many internal (non-leaf) nodes will there be? **3**
 - 3) How many internal (non-leaf) nodes do we traverse to do an equality search? **2**
-