


Transactions & Concurrency Control 1

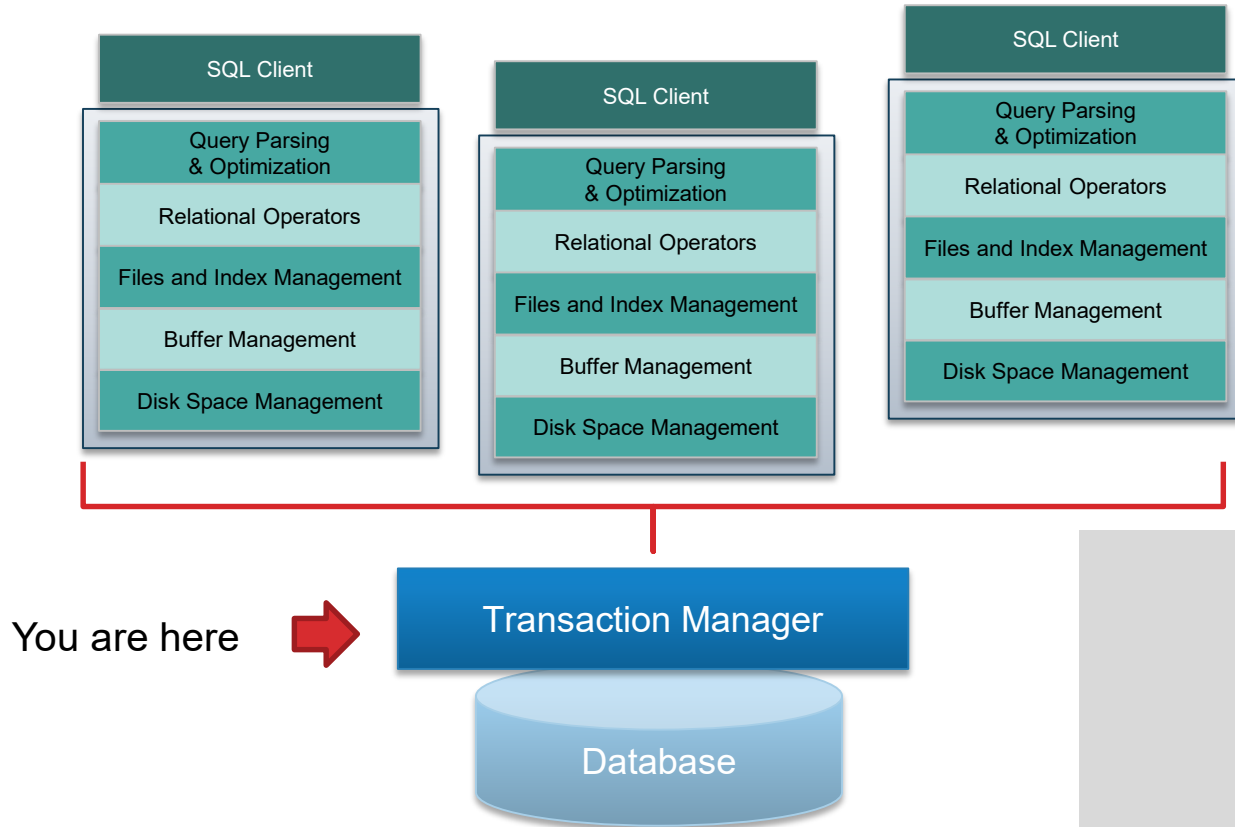
R&G 16/17



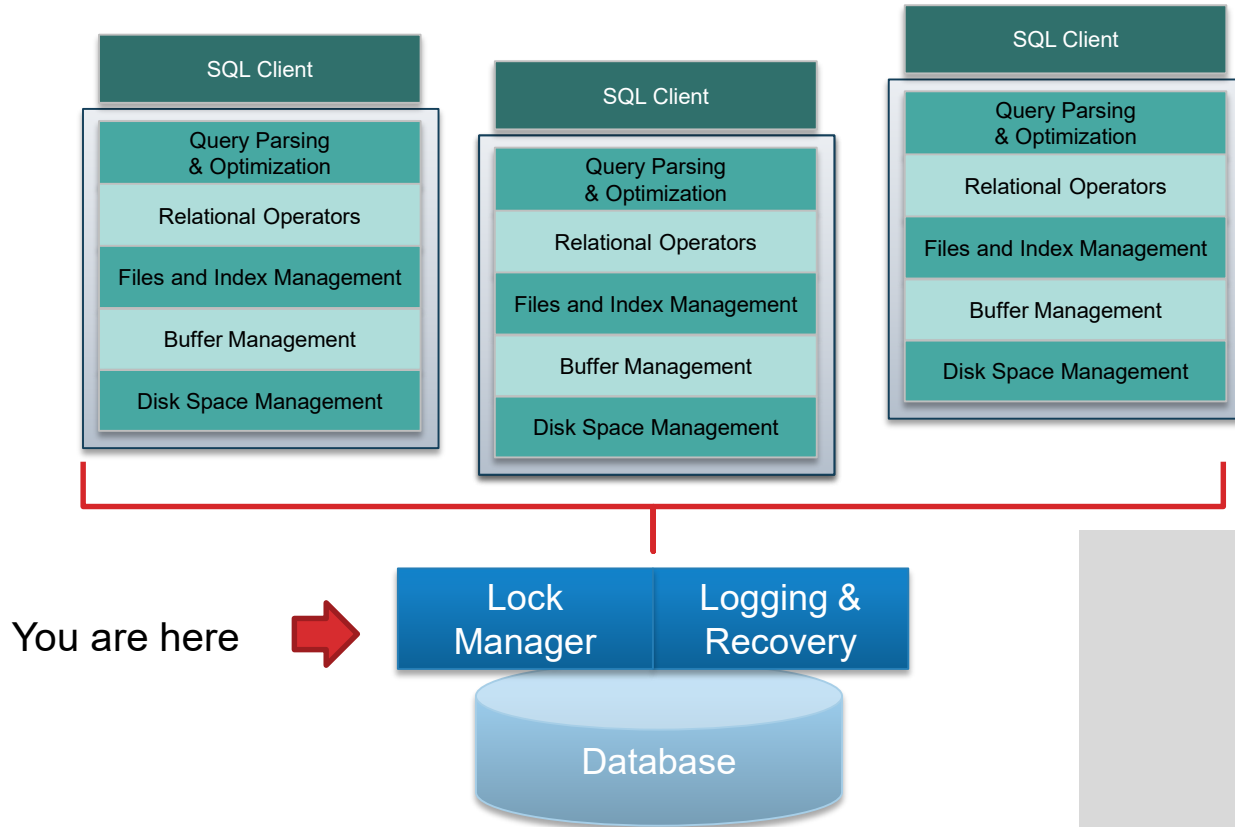
Architecture of a DBMS



Architecture of a DBMS, Part 2



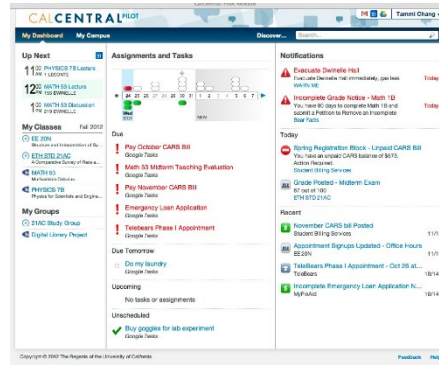


Architecture of a DBMS, Part 3




Applications on DBMS

- Virtually any compute service that maintains state today is an application on top of some kind of DBMS
 - Uber
 - Kayak
 - Amazon.com
 - BankofAmerica
 - Pokemon Go



Applications Want Something from the DBMS

- Queries and updates of course: what you learned so far!
- Real applications are composed of many statements being generated by user behaviors
- Many users work with the application at the same time



Concurrency Control & Recovery

- **Part 1: Concurrency Control**
 - Correct/fast data access in the presence of concurrent work by many users
 - Disorderly processing that provides the illusion of order
- **Part 2: Recovery**
 - Ensure database is fault tolerant
 - Not corrupted by software, system or media failure
 - Storage guarantees for mission-critical data
- **It's all about the programmer!**
 - Systems provide guarantees
 - These guarantees lighten the load of app writers

Concurrent Execution: Why bother?

- Multiple transactions are allowed to run concurrently in the system.
- Advantages are twofold:
 - *Throughput* (transactions per second):
 - Increase processor/disk utilization → more transactions per second (TPS) completed
 - Single core: can use the CPU while another xact is reading to/writing from the disk
 - Multicore: ideally, scale throughput in the number of processors
 - *Latency* (response time per transaction):
 - Multiple transactions can run at the same time
 - So one transaction's latency need not be dependent on another unrelated transaction
 - Or that's the hope
- Both are important!

Motivating Example

```
UPDATE Budget  
SET money = money - 500  
WHERE pid = 1
```

```
UPDATE Budget  
SET money = money + 200  
WHERE pid = 2
```

```
UPDATE Budget  
SET money = money + 300  
WHERE pid = 3
```

```
SELECT sum(money)  
FROM Budget
```

Two Issues:

1. Order matters!
2. Users need a way to say what's OK

Different Types of Problems

User 1

```
INSERT INTO DollarProducts(name, price)
SELECT pname, price
FROM Product
WHERE price <= 0.99

DELETE Product
WHERE price <= 0.99
```

User 2

```
SELECT count(*)
FROM Product
```

```
SELECT count(*)
FROM DollarProducts
```

What could go wrong?

Inconsistent Reads

Different Types of Problems, Part 2

User 1

```
UPDATE Product  
SET Price = Price - 10.99  
WHERE pname = "CoolToy"
```

User 2

```
UPDATE Product  
SET Price = Price*0.6  
WHERE pname = "CoolToy"
```

What could go wrong?

Lost Update

Different Types of Problems, Part 3

User 1

```
UPDATE Account  
SET amount = 1000000  
WHERE number = "my-account"
```

Aborted by
the system

User 2

```
SELECT amount  
FROM Account  
WHERE number = "my-account"
```

What could go wrong? **Dirty Reads**



TRANSACTIONS



Transaction: Concept and Implementation

- Major component of database systems
- Critical for most applications; arguably more so than SQL

An Aside: Database Turing Awards



What is a Transaction?

- A sequence of *multiple actions* to be executed as an *atomic* unit
- Application View (SQL View):
 - Begin transaction
 - Sequence of SQL statements
 - End transaction
- Examples
 - Transfer money between accounts
 - Book a flight, a hotel and a car together on Expedia

Our Transaction Model

- **Transaction (“Xact”):**
 - DBMS’s abstract view of an application program (or activity)
 - A sequence of *reads* and *writes* of database objects
 - Batch of work that must *commit* or *abort* as an atomic unit
- **Xact Manager controls execution of transactions**
- **Program logic is invisible to DBMS!**
 - Arbitrary computation possible on data fetched from the DB
 - The DBMS only sees data read/written from/to the DB
 - (Note: modern systems have started rethinking this assumption, but we’ll stick with it here)

Transaction Example

- Transaction to transfer \$100 from account R to account S

Not seen by the
DBMS transaction
manager!

1. start transaction
2. read(R)
3. $R = R - 100$
4. write(R)
5. read(S)
6. $S = S + 100$
7. write(S)
8. end transaction

ACID: High-Level Properties of Transactions

- **A tomicity:** *All* actions in the Xact happen, or *none* happen.
- **C onsistency:** If the DB *starts* out *consistent*, it *ends* up *consistent* at the end of the Xact
- **I solation:** Execution of *each* Xact is *isolated from* that of *others*
- **D urability:** If a Xact *commits*, its effects *persist*.

Note: This is a mnemonic, not a formalism. We'll do some formalisms shortly.

I Isolation (Concurrency)

- DBMS interleaves actions of many xacts
 - Actions = reads/writes of DB objects
- DBMS ensures 2 xacts do not “interfere”
- Each xact executes as if it ran by itself.
 - Concurrent accesses have no effect on xact’s behavior
 - Net effect must be identical to executing all transactions in some serial order
 - Users & programmers think about transactions in isolation
 - Without considering effects of other concurrent Xacts!

Isolation: An Example

- Think about avoiding problems due to concurrency
 - If another transaction T2 accesses R and S between steps 4 and 5 of T1, it will see a lower value for R+S.

T1

1. start transaction
2. read(R)
3. $R = R - 100$
4. write(R)
5. read(S)
6. $S = S + 100$
7. write(S)
8. end transaction

T2

1. start transaction
2. read(R)
3. print(R+S)
4. end transaction

- Isolation easy to achieve by running one Xact at a time
 - However, recall that serial execution is not desirable

Atomicity and Durability

- **A transaction ends in one of two ways:**
 - **Commit** after completing all its actions
 - “commit” is a contract with the caller of the DB
 - **Abort** (or be aborted by the DBMS) after executing some actions
 - Or **system crash** while the xact is in progress; treat as abort.
- **Two key properties** for a transaction
 - **Atomicity:** Either execute all its actions, or none of them
 - **Durability:** The effects of a committed xact must survive failures.
- DBMS typically ensures the above by **logging** all actions:
 - **Undo** the actions of aborted/failed transactions.
 - **Redo** actions of committed transactions not yet propagated to disk when system crashes

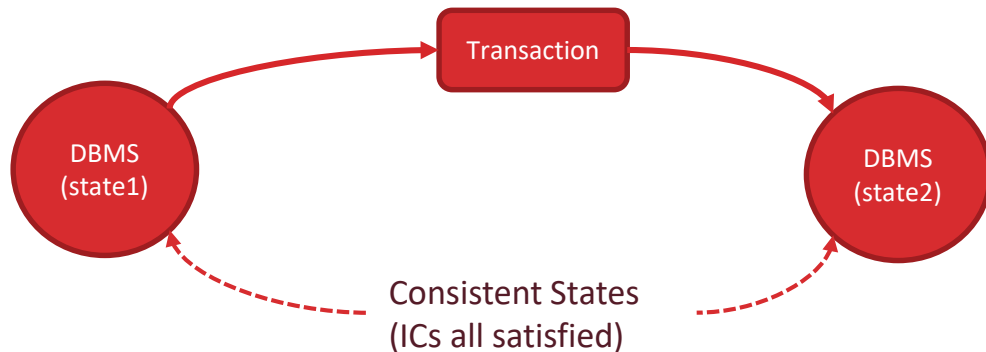
Atomicity and Durability, cont.

- Atomicity
 - If the transaction fails after step 4 and before step 7
 - Money will be “lost” → inconsistent database
 - DBMS should ensure that updates of a partially executed transaction are not reflected
- Durability
 - Once the user hears that the transaction is complete, can rest easy that the \$100M was transferred from R to S.

1. start transaction
2. read(R)
3. $R = R - 100$
4. write(R)
5. read(S)
6. $S = S + 100$
7. write(S)
8. end transaction

Transaction **C**onsistency

- **Transactions preserve DB consistency**
 - Given a consistent DB state, produce another consistent DB state
- DB consistency expressed as a set of **declarative integrity constraints**
 - CREATE TABLE/ASSERTION statements
- **Transactions that violate integrity are aborted**
 - That's all the DBMS can automatically check!



Summary

- We have seen an overview
- ACID Transactions make guarantees that
 - Improve performance (via concurrency)
 - Relieve programmers of correctness concerns
 - Hide concurrency and failure handling!
- Two key issues to consider, and mechanisms
 - Concurrency control (via two-phase locking)
 - Recovery (via write-ahead logging WAL)
- We'll do concurrency control first

CONCURRENCY CONTROL

Concurrency Control: Providing Isolation

- **Naïve approach - serial execution**
 - One transaction runs at a time
 - Safe but slow
- **Execution must be interleaved for better performance**
- With concurrent executions, how does one **define** and **ensure** correctness?

Transaction Schedules

T1	T2
begin	
read(A)	
write(A)	
read(B)	
write(B)	
commit	
	begin
	read(A)
	write(A)
	read(B)
	write(B)
	commit

A **schedule** is a sequence of actions on data from one or more transactions.

Actions: Begin, Read, Write, Commit and Abort.

$R_1(A) W_1(A) R_1(B) W_1(B) R_2(A) W_2(A) R_2(B) W_2(B)$

By convention we only include committed transactions, and omit Begin and Commit.

Serial Equivalence

- We need a “touchstone” concept for correct behavior

- **Definition: Serial schedule**

- Each transaction runs from start to finish without any intervening actions from other transactions

- **Definition: 2 schedules are equivalent** if they:

- involve the same transactions
 - each individual transaction's actions are ordered the same
 - both schedules leave the DB in the same final state

T1	T2
START	
READ(A)	
WRITE(B)	
WRITE(C)	
WRITE(D)	
FINISH	
	START
	READ(B)
	WRITE(C)
	WRITE(D)
	FINISH

T1	T2
START	
WRITE(A)	
WRITE(B)	
WRITE(C)	
FINISH	
	START
	READ(A)
	WRITE(B)
	WRITE(C)
	FINISH

T1	T2
START	
WRITE(A)	
WRITE(B)	
WRITE(C)	
FINISH	
	START
	READ(A)
	WRITE(B)
	WRITE(C)
	FINISH



Serializability

- **Definition:** Schedule S is **serializable** if:
 - S is equivalent to some serial schedule



Schedule 1

T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
begin	
read(A)	
$A = A - 100$	
write(A)	
read(B)	
$B = B + 100$	
write(B)	
commit	
	begin
	read(A)
	$A = A * 1.1$
	write(A)
	read(B)
	$B = B * 1.1$
	write(B)
	commit

- Let T1 transfer \$100 from A to B
- Let T2 add 10% interest to A & B
- Serial schedule in which T1 is followed by T2
 - Final outcome:
 - $A := 1.1 * (A - 100)$
 - $B := 1.1 * (B + 100)$

Schedule 2

T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
	begin
	read(A)
	$A = A * 1.1$
	write(A)
	read(B)
	$B = B * 1.1$
	write(B)
	commit
begin	
read(A)	
$A = A - 100$	
write(A)	
read(B)	
$B = B + 100$	
write(B)	
commit	

- Serial schedule in which T2 is followed by T1
 - Final outcome:
 - $A := (1.1 * A) - 100$
 - $B := (1.1 * B) + 100$
 - Different!
 - But still understandable

Schedule 3

T1: Transfer \$100 from A to B	T2: Add 10% interest to A & B
begin	
read(A)	
A = A - 100	
write(A)	
	begin
	read(A)
	A = A * 1.1
	write(A)
read(B)	
B = B + 100	
write(B)	
commit	
	read(B)
	B = B * 1.1
	write(B)
	commit

- Schedule in which actions of T1 and T2 are interleaved.
- This is not a serial schedule
- But it is equivalent to schedule 1
 - $A := (A-100)*1.1$
 - $B := (B+100)*1.1$
- Hence **serializable!**

Conflicting Operations

- Tricky to check property **“leaves the DB in the same final state”**
- Need an easier equivalence test!
 - Settle for a “conservative” test: always true positives, but some false negatives
 - I.e. sacrifice some concurrency for easier correctness check
- **Use notion of “conflicting” operations (read/write)**
- **Definition: Two operations conflict if they:**
 - Are by different transactions,
 - Are on the same object,
 - At least one of them is a write.
- The order of non-conflicting operations has no effect on the final state of the database!
 - Focus our attention on the order of conflicting operations

Conflict Serializable Schedules

- **Definition: Two schedules are *conflict equivalent* if:**
 - They involve the same actions of the same transactions, and
 - Every pair of conflicting actions is ordered the same way
- **Definition: Schedule S is *conflict serializable* if:**
 - S is conflict equivalent to some serial schedule
 - Implies S is also Serializable

Note: some serializable schedules are NOT conflict serializable

- Conflict serializability gives false negatives as a test for serializability!
- The cost of a conservative test
- A price we pay to achieve efficient enforcement

Conflict Serializability - Intuition

- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- ***Example***

R(A) W(A)


R(B) W(B)

R(A) W(A)

R(B) W(B)


Conflict Serializability – Intuition, Part 2

- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*





Conflict Serializability – Intuition, Part 3

- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*



Conflict Serializability – Intuition, Part 4

- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*



Conflict Serializability – Intuition, Part 5


- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*

R(A) W(A)	R(B) W(B)
R(A) W(A)	R(B) W(B)

R(A) W(A) R(B)	W(B)
R(A)	W(A) R(B) W(B)

Conflict Serializability – Intuition, cont

- **A schedule S is conflict serializable if**
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions
- *Example*



R(A) W(A) R(B) W(B)
R(A)W(A) R(B) W(B)

Conflict Serializability (Continued)


- Here's another example:

$R(A)$ $W(A)$
 $R(A)$ $W(A)$

- Conflict Serializable or not?

NOT!

Conflict Dependency Graph



- **Dependency Graph:**
 - One node per Xact
 - Edge from T_i to T_j if:
 - An operation O_i of T_i conflicts with an operation O_j of T_j and
 - O_i appears earlier in the schedule than O_j
- **Theorem: Schedule is conflict serializable if and only if its dependency graph is acyclic.**

Proof Sketch: Conflicting operations prevent us from “swapping” operations into a serial schedule

Example

- **A schedule that is not conflict serializable**

T1: R(A), W(A)




Dependency graph

Example, pt 2

- **A schedule that is not conflict serializable**

T1:	R(A), W(A),
T2:	R(A)




Example, pt 3

- **A schedule that is not conflict serializable**

T1: R(A), W(A),


T2: R(A), W(A), R(B), W(B)



Example, pt 4


- **A schedule that is not conflict serializable**

T1:	R(A), W(A),	R(B)
T2:	R(A), W(A), R(B), W(B)	



View Serializability

- **Alternative notion of serializability: fewer false negatives**
- **Schedules S1 and S2 are view equivalent if:**
 - *Same initial reads:*
 - If T_i reads initial value of A in S1, then T_i also reads initial value of A in S2
 - *Same dependent reads:*
 - If T_i reads value of A written by T_j in S1, then T_i also reads value of A written by T_j in S2
 - *Same winning writes:*
 - If T_i writes final value of A in S1, then T_i also writes final value of A in S2
- Basically, allows all conflict serializable schedules + “blind writes”



Notes on Serializability Definitions

- **View Serializability allows (a few) more schedules than conflict serializability**
 - But V.S. is difficult to enforce efficiently.
- **Neither definition allows all schedules that are actually serializable.**
 - Because they don't understand the meanings of the operations or the data
- **Conflict Serializability is what gets used, because it can be enforced efficiently**
 - To allow more concurrency, some special cases do get handled separately.
 - (Search the web for "Escrow Transactions" for example)