

Parallel Query Processing

R&G Chapters
22.1-22.4,



A little history

- Relational revolution
 - declarative set-oriented primitives
 - 1970' s
- Parallel relational database systems
 - on commodity hardware
 - 1980' s
- Big Data: MapReduce, Spark, etc.
 - scaling to thousands of machines and beyond
 - 2005-2015


Why Parallelism?

- Scan 100TB
 - At 0.5 GB/sec (see lec 4):
~200,000 sec = ~2.31 days





Why Parallelism? Cont.

- Scan 100TB
 - At 0.5 GB/sec (see lec 4):
~200,000 sec = ~2.31 days
- Run it 100-way parallel:
 - 2,000 sec = 33 minutes
- 1 big problem = many small problems
 - Trick: make them independent




Two Metrics to Shoot For


- Speed-up
 - Increase HW
 - Fix workload
- Scale-up
 - Increase HW
 - Increase workload




Roughly 2 Kinds of Parallelism



: any sequential program,
e.g. a relational operator



Pipeline
scales up to pipeline depth



Partition
scales up to amount of data

We'll get more
refined soon.

Easy for us to say!


- Lots of Data:
 - Batch operations
 - Pre-existing divide-and-conquer algorithms
 - Natural pipelining
- Declarative languages
 - Can adapt the parallelism strategy to the task and the hardware
 - All without changing the program!
 - Codd's Physical Data Independence

DBs: The Parallel Boy that Lived


- 1980s CS challenge: “parallelize” software
 - E.g. via awesome new C compilers
 - E.g. via variants of C designed for parallelism
- In broad terms, a failure
 - Dave Patterson: The “Dead Computer Society”
- Exception: Parallel SQL Databases
 - Why? **Data Independence!**
 - SQL is independent of how many machines you have!
 - The same divide-and-conquer that worked for disks works across machines, as we’ll see
- Big Data is the generalization of these lessons
 - Or in some cases a re-learning of them

Parallel Architectures


Shared Memory



Shared Disk



Shared Nothing
(cluster)



Some Early Systems

- Research
 - XPRS (Berkeley, shared-memory)
 - Gamma (Wisconsin, shared-nothing)
 - Volcano (Colorado, shared-nothing)
 - Bubba (MCC, shared-nothing)
 - Grace (U. Tokyo, shared-nothing)
- Industry
 - Teradata (shared-nothing)
 - Tandem Non-Stop SQL (shared-nothing)

What about the cloud?


- Upshot: not so different from what we'll see here
- Architectural choices and competing systems jockeying for position
- Parallelism in many forms across many users and use cases
- Things still shaking out

Shared Nothing

- We will focus on Shared Nothing here
 - It's the most common
 - DBMS, web search, big data, machine learning, ...
 - Runs on commodity hardware
 - Scales up with data
 - Just keep putting machines on the network!
 - Does not rely on HW to solve problems
 - Good for helping us understand what's going on
 - Control it in SW


Kinds of Query Parallelism

- Inter-query (parallelism across queries)
 - Each query runs on a separate processor
 - Single thread (no parallelism) per query
 - Does require parallel-aware concurrency control
 - A topic for later in the semester
- Note on latin prefixes
 - inter: “between”, “across”.
“Interplanetary travel takes a long time”
 - intra: “within”.
“The political party suffered from intraparty rivalries.”



Intra Query – Inter-operator


- Intra-query (within a single query)
 - **Inter-operator** (between operators)




Pipeline Parallelism

Intra Query – Inter-operator Part 2

- Intra-query
 - Inter-operator




Pipeline Parallelism





Intra Query - Inter-Operator Part 3

- Intra-query
 - Inter-operator




Pipeline Parallelism



Bushy (Tree) Parallelism



Intra Query – Intra-Operator

- Intra-query
 - **Intra-operator** (within a single operator)



Kinds of Query Parallelism, cont.


- Intra-query
 - Intra-operator



Partition Parallelism



Summary: Kinds of Parallelism

- Inter-Query




- Intra-Query

- Inter-Operator



Pipeline Parallelism

- Intra-Operator (partitioned)




Partition Parallelism

INTRA-OPERATOR PARALLELISM

Data Partitioning

- How to partition a table across disks/machines
 - A bit like coarse-grained indexing!




- Shared nothing particularly benefits from "good" partitioning

Parallel Scans


- Scan in parallel, merge (concat) output
- σ_p : skip entire sites that have no tuples satisfying p
 - range or hash partitioning
- Indexes can be built at each partition
- Q: How do indexes differ in the different data partitioning schemes?

Lookup by key

- Data partitioned on function of key?
 - Great! Route lookup only to relevant node
- Otherwise
 - Have to broadcast lookup (to all nodes)




Hash




Round-Robin

What about Insert?

- Data partitioned on function of key?
 - Route insert to relevant node
- Otherwise
 - Route insert to *any* node




Hash




Round-Robin

Insert to Unique Key?

- Data partitioned on function of key?
 - Route to relevant node
 - And reject if already exists




Hash




Round-Robin

Insert to Unique Key cont.

- Otherwise
 - Broadcast lookup
 - Collect responses
 - If not exists, insert anywhere
 - Else reject




Hash




Round-Robin

Remember Hashing?




Parallelize me! Hashing

- Phase 1: shuffle data across machines (h_n)
 - streaming out to network as it is scanned
 - which machine for this record?
 - use (yet another) independent hash function h_n



Parallelize me! Hashing Part 2

- Receivers proceed with phase 1 in a pipeline as data streams in
 - from local disk and network



Nearly same as single-node hashing

*Near-perfect speed-up, scale-up!
Streams through phase 1, during which
time every component works at its top
speed, no waiting.*

Have to wait to start phase 2.


Hash Join?

- Hmmmm....

If you have enough machines...


Naïve parallel hash join

- Phase 1: *shuffle* each table across machines (h_n)
 - Parallel scan streaming out to network
 - Wait for building relation to finish
 - Then stream probing relation through it
- Receivers proceed with naïve hashing in a pipeline *as probe data streams in*
 - from local disk and network
 - Writes are independent, hence parallel
- Note: there is a variation that has *no waiting*: both tables stream
 - Wilschut and Apers' "Symmetric" or "Pipeline" hash join
 - Requires more memory space




Parallel Grace Hash Join Pass 1

- Pass 1 is like hashing above





Parallel Grace Hash Join Pass 1 cont

- Pass 1 is like hashing above
 - But do it 2x: once for each relation being joined



Parallel Grace Hash Join Pass 2

- Pass 2 is local Grace Hash Join per node
 - Complete independence across nodes




Parallel Grace Hash Join

- Pass 1: parallel streaming
 - Stream building and probing tables through shuffle/partition
- Pass 2 is local Grace Hash Join per node
 - Complete independence across nodes in Pass 2
- Near-perfect speed-up, scale-up!
- Every component works at its top speed
 - Only waiting is for Pass 1 to end.
- Note: there is a variant that has no waiting
 - Urhan's Xjoin, a variant of symmetric hash


Parallelize me! Sorting Pass 0

- Pass 0: shuffle data across machines
 - streaming out to network as it is scanned
 - which machine for this record?
Split on value range (e.g. $[-\infty, 10]$, $[11, 100]$, $[101, \infty]$).




Parallelize me! Sorting Pass 1-n

- Receivers proceed with pass 0 as the data streams in
- Passes 1–n done independently as in single-node sorting
- **A Wrinkle: How to ensure ranges are the same #pages?!**
 - i.e. avoid data skew?




Range partitioning

- Goal: equal frequency per machine
- Note: ranges often don't divide x axis evenly
- How to choose?



Range partitioning cont.


- Would be easy if data small
- In general, can sample the input relation prior to shuffling, pick splits based on sample
- Note: Random sampling can be tricky to implement in a query pipeline; simpler if you materialize first.



How to sample a database table?
Advanced topic, we will not discuss in this class.


Parallel Sort-Merge Join

- Pass 0 .. $n-1$ are like parallel sorting above
- Note: this picture is a 2-pass sort ($n=1$); this is pass 0





Parallel Sort-Merge Join Pass 0...n-1

- Pass 0 .. n-1 are like parallel sorting above
 - But do it 2x: once for each relation, with same ranges
 - Note: this picture is a 2-pass sort (n=1); this is pass 0




Pass n (with optimization)

- Pass 0 .. n-1 are like parallel sorting above
 - But do it 2x: once for each relation, with same ranges
- Pass n: merge join partitions locally on each node




Parallel Aggregates

- Hierarchical aggregation
- For each aggregate function, need a global/local decomposition:
 - **sum**(S) = $\Sigma \Sigma (s)$
 - **count** = $\Sigma \text{count} (s)$
 - **avg**(S) = $(\Sigma \Sigma (s)) / \Sigma \text{count} (s)$
 - etc...




Parallel GroupBy

- Naïve Hash Group By
 - Local aggregation: in hash table keyed by group key k_i keep local agg_i
 - E.g. `SELECT SUM(price) group by cart;`




Parallel GroupBy, Cont.

- Naïve Hash Group By
 - Local aggregation: in hash table keyed by group key k_i keep local agg_i
 - For example, k is major, agg is $(avg(gpa), count(*))$
 - Shuffle local aggs by a hash function $h_p(k_i)$
 - Compute global aggs for each key k_i



Parallel Aggregates/GroupBy Challenge!

- Exercise:
 - Figure out parallel 2-pass GraceHash-based scheme to handle # large of groups
 - Figure out parallel Sort-based scheme




Joins: Bigger picture

- Alternatives:
 - Symmetric shuffle
 - What we did so far
 - Asymmetric shuffle
 - Broadcast join


Join: One-sided shuffle

- If R already suitably partitioned,
- just partition S, then run local join at every node and union results.



“Broadcast” Join

- If R is small, send it to all nodes that have a partition of S.
- Do a local join at each node (using any algorithm) and union results.




What are “pipeline breakers”?

- Sort
 - Hence sort-merge join can't start merging until sort is complete
- Hash build
 - Hence Grace hash join can't start probing until hashtable is built
- Is there a join scheme that pipelines?

Symmetric (Pipeline) Hash Join

- Single-phase, streaming
- Each node allocates two hash tables, one for each side
- Upon arrival of a tuple of R:
 - Build into R hashtable by join key
 - Probe into S hashtable for matches and output any that are found
- Upon arrival of a tuple of S:
 - Symmetric to R!




Symmetric (Pipeline) Hash Join cont

- Why does it work?
 - Each output tuple is generated exactly once: when the second part arrives
- Streaming!
 - Can always pull another tuple from R or S, build, and probe for outputs
 - Useful for Stream query engines!

Extensions

- Parallel Symmetric Hash Join
 - Straightforward—part of the original proposal
 - Just add a streaming partitioning phase up front
 - As in naïve hash join
- Out-of-core Symmetric Hash Join
 - Quite a bit trickier. See the [X-Join paper](#).
- Non-blocking sort-merge join
 - See the [Progressive Merge Join](#) paper



Parallel DBMS Summary

- Parallelism natural to query processing:
 - Both pipeline and partition
- Shared-Nothing vs. Shared-Mem vs. Shared Disk
 - Shared-mem easiest SW, costliest HW.
 - Doesn't scale indefinitely
 - Shared-nothing cheap, scales well, harder to implement.
 - Shared disk a middle ground
 - For updates, introduces icky stuff related to concurrency control
- Intra-op, Inter-op, & Inter-query parallelism all possible.

Parallel DBMS Summary, Part 2

- Data layout choices important!
- Most DB operations can be done partition-parallel
 - Sort. Hash.
 - Sort-merge join, hash-join.
- Complex plans.
 - Allow for pipeline-parallelism, but sorts, hashes block the pipeline.
 - Partition parallelism achieved via bushy trees.

Parallel DBMS Summary, Part 3

- Transactions require introducing some new protocols
 - distributed deadlock detection
 - two-phase commit (2PC)
- 2PC not great for availability, latency
 - single failure stalls the whole system
 - transaction commit waits for the slowest worker
- More on this in subsequent lectures