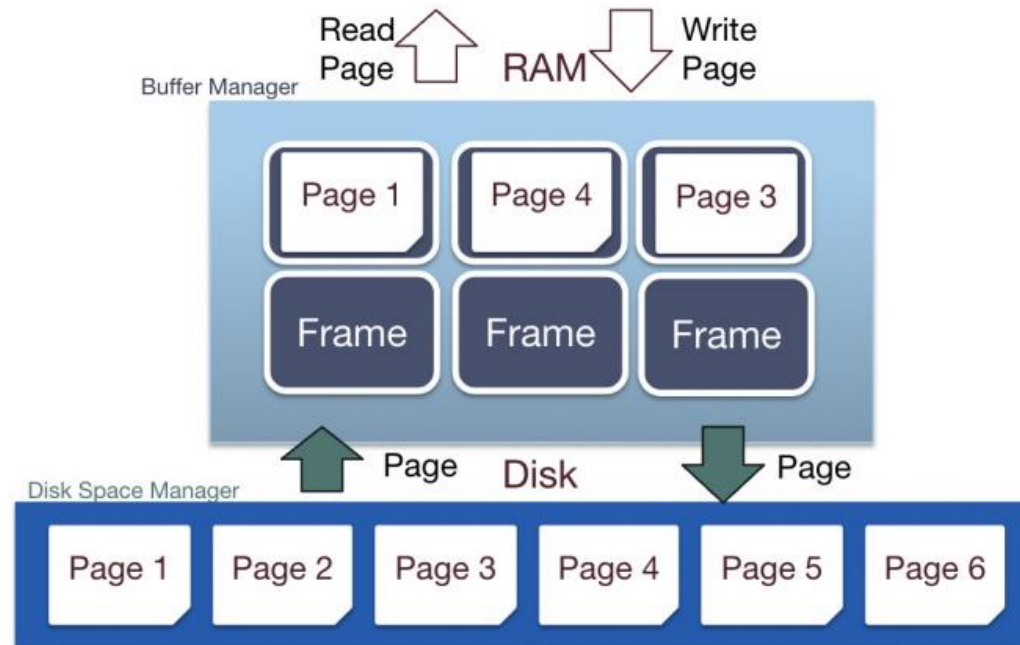


Discussion 4

Buffer Manager & B+ Tree Refinements
& Relational Algebra

Buffer Manager

- manage pages in memory
- process page requests from the file and index manager



Buffer Pool

- **Frame ID** : uniquely associated with a memory address
- **Page ID** : determine which page a frame currently contains
- **Dirty Bit** : verify whether or not a page has been modified
- **Pin Count** : track the number of requestors currently using a page

FrameId	PageId	Dirty?	Pin Count
1	1	N	0
2	2	Y	1
3	3	N	0
4	6	N	2
5	4	N	0
6	5	N	0

A buffer frame can hold the same amount of data as a page can

Handling Page Requests

- Page hit - requested page already exists within memory:
 - the page's pin count is incremented
 - the page's memory address is returned
- Page miss - requested page is not in pool:
 - If there is still space, the next empty frame is found and the page is read into that frame, then pin the page and return its address.
 - Else, Page replacement policy

Page replacement policy

- Choose an un-pinned (`pin_count = 0`) frame for replacement
- If frame “dirty”, write current page to disk, mark “clean”
- Read requested page into frame
- Pin the page and return its address

LRU Replacement Policy

- the Least Recently Used unpinned page which has **pin count = 0**
- a Last Used column is added to the metadata table and measures the latest time at which a page's pin count is decremented -- **costly!**

FrameId	PageId	Dirty?	Pin Count	Last Used
1	1	N	0	43
2	2	Y	1	21
3	3	N	0	22
4	6	N	2	11
5	4	N	0	24
6	5	N	0	15

Clock Policy

- Iterate through frames within the table, skipping pinned pages and wrapping around to frame 0 upon reaching the end, until the first **unpinned frame** with **ref bit = 0** is found.
- if the current frame's ref bit = 1, set the ref bit to 0 and move the clock hand to the next frame.
- if ref bit = 0, evict the existing page (and write it to disk if the dirty bit is set; then set the dirty bit to 0), read in the new page, set the frame's ref bit to 1, and **move the clock hand to the next frame**.

FrameId	PageId	Dirty?	Pin Count	Ref Bit
1	1	N	1	1
2	2	N	1	1
3	3	N	0	1
4	4	N	0	0
5	5	N	0	0
6	6	N	0	1


Clock Hand

1


Sequential Scanning Performance

A B C D A B C D A B C D A B C D

LRU



MRU
Most Recently Used



Search Key and Ordering

In an ordered index (e.g. B+-tree) the keys are ordered lexicographically by the search key columns:

Composite Keys: more than one column

- **Lexicographic order**
- Search a range
✗ Age > 31 & Salary = 400

SSN	Last Name	First Name	Age	Salary
123	Adams	Elmo	31	\$300
443	Grouch	Oscar	32	\$400
244	Oz	Bert	55	\$140
134	Sanders	Ernie	55	\$400
176	Grump	Donald	79	\$300

Three basic alternatives for data entries in any index

- Alternative 1: By Value
 - Record contents are stored in the **index file** --No need to follow pointers
- Alternative 2: By Reference
 - By **Reference**, <**k**, rid of matching data record>
- Alternative 3: By List of references
 - By **Reference**, <**k**, rid of matching data record>
 - For very large rid lists, single data entry **spans multiple blocks**

Alternative 2
Index data entries


Key	Record Id	SSN	Last Name	First Name	Salary
Gonzalez	[3, 1]	123	Gonzalez	Amanda	\$400
Gonzalez	[3, 2]	443	Gonzalez	Joey	\$300
Gonzalez	[3, 3]	244	Gonzalez	Jose	\$140
Hong	[3, 4]	134	Hong	Sue	\$400

Alternative 3
Index data entries


Key	Record Id
Gonzalez	[3, {1, 2, 3}]
Hong	[3,4]

Clustered vs. Unclustered Index

- Heap file records are kept mostly ordered according to **search keys** in index
- To build a clustered index, first sort the heap file
 - Leave some free space on each block for future inserts
- Blocks at end of file may be needed for inserts
 - Order of data records is “close to”, but not identical to, the sort order




B+ Tree Refinement: Variable-Length Keys



- Prefix Key Compression
 - determine minimum splitting prefix and copy up



- Suffix Key Compression
 - Move common prefix to header, leave only (compressed) suffix next to pointer



B+-TREE COSTS

	Heap File	Sorted File	Clustered Index
Scan all records	$O(B)$	$O(B)$	$O(B)$
Equality Search	$O(B)$	$O(\log_2 B)$	$O(\log_F B)$
Range Search	$O(B)$	$O(\log_2 B)$	$O(\log_F B)$
Insert	$O(1)$	$O(B)$	$O(\log_F B)$
Delete	$O(B)$	$O(B)$	$O(\log_F B)$

- B: The number of data blocks
- R: Number of records per block
- D: Average time to read/write disk block
- F: Average internal node fanout
- E: Average # data entries per leaf

Relational Algebra Operators

Unary Operators: on single relation

- Projection (π): Retains only desired columns (vertical)
- Selection (σ): Selects a subset of rows (horizontal)
- Renaming (ρ): Rename attributes and relations.

Binary Operators: on pairs of relations

- **Union** (\cup): Tuples in r1 or in r2.
- **Set-difference** ($-$): Tuples in r1, but not in r2.
- Cross-product (\times): Allows us to combine two relations

Compound Operators: common “macros” for the above

- **Intersection** (\cap): Tuples in r1 and in r2.
- Joins (\bowtie_{θ} , \bowtie): Combine relations that satisfy predicates

Compound Operator: Join

- Joins are compound operators (like intersection):
 - Generally, $R \bowtie_{\theta} S = \sigma_{\theta} (R \times S)$
- Hierarchy of common kinds:
 - **Theta Join** (\bowtie_{θ}): join on logical expression θ
 - **Equi-Join**: theta join with theta being a conjunction of equalities
 - **Natural Join** (\bowtie): equi-join on **all matching column names**