

Discussion 5

Sorting and Hashing

Binbin Chen


chenbb@shanghaitech.edu.cn

Purpose of sorting and hashing

- sorting
 - rendezvous matches (getting together tuples processed at the same time in the same place)
 - like eliminating duplicates, grouping
 - ordering
 - like step in B+ trees
 - infeasibility of virtual memory
 - sort is incompatible with random disk IOs
 - slow with big buffer
- hashing
 - often just rendezvous matches
 - different with traditional hashing with hash tables
 - divide and conquer for out-of-core algorithms

Out-of-Core (RAM) Algorithms


- Two themes
 - Single-pass streaming data from disk through RAM back to disk
 - Read a chunk from INPUT to an Input Buffer
 - Write $f(x)$ to map each item to an Output Buffer
 - When Input Buffer is consumed, read another chunk
 - When Output Buffer fills, write it to OUTPUT
 - Divide (into RAM-sized chunks) and Conquer



Note: the shape in the output buffer is smaller because $f(x)$ maps the item to smaller bytes, so input and output are not synchronize

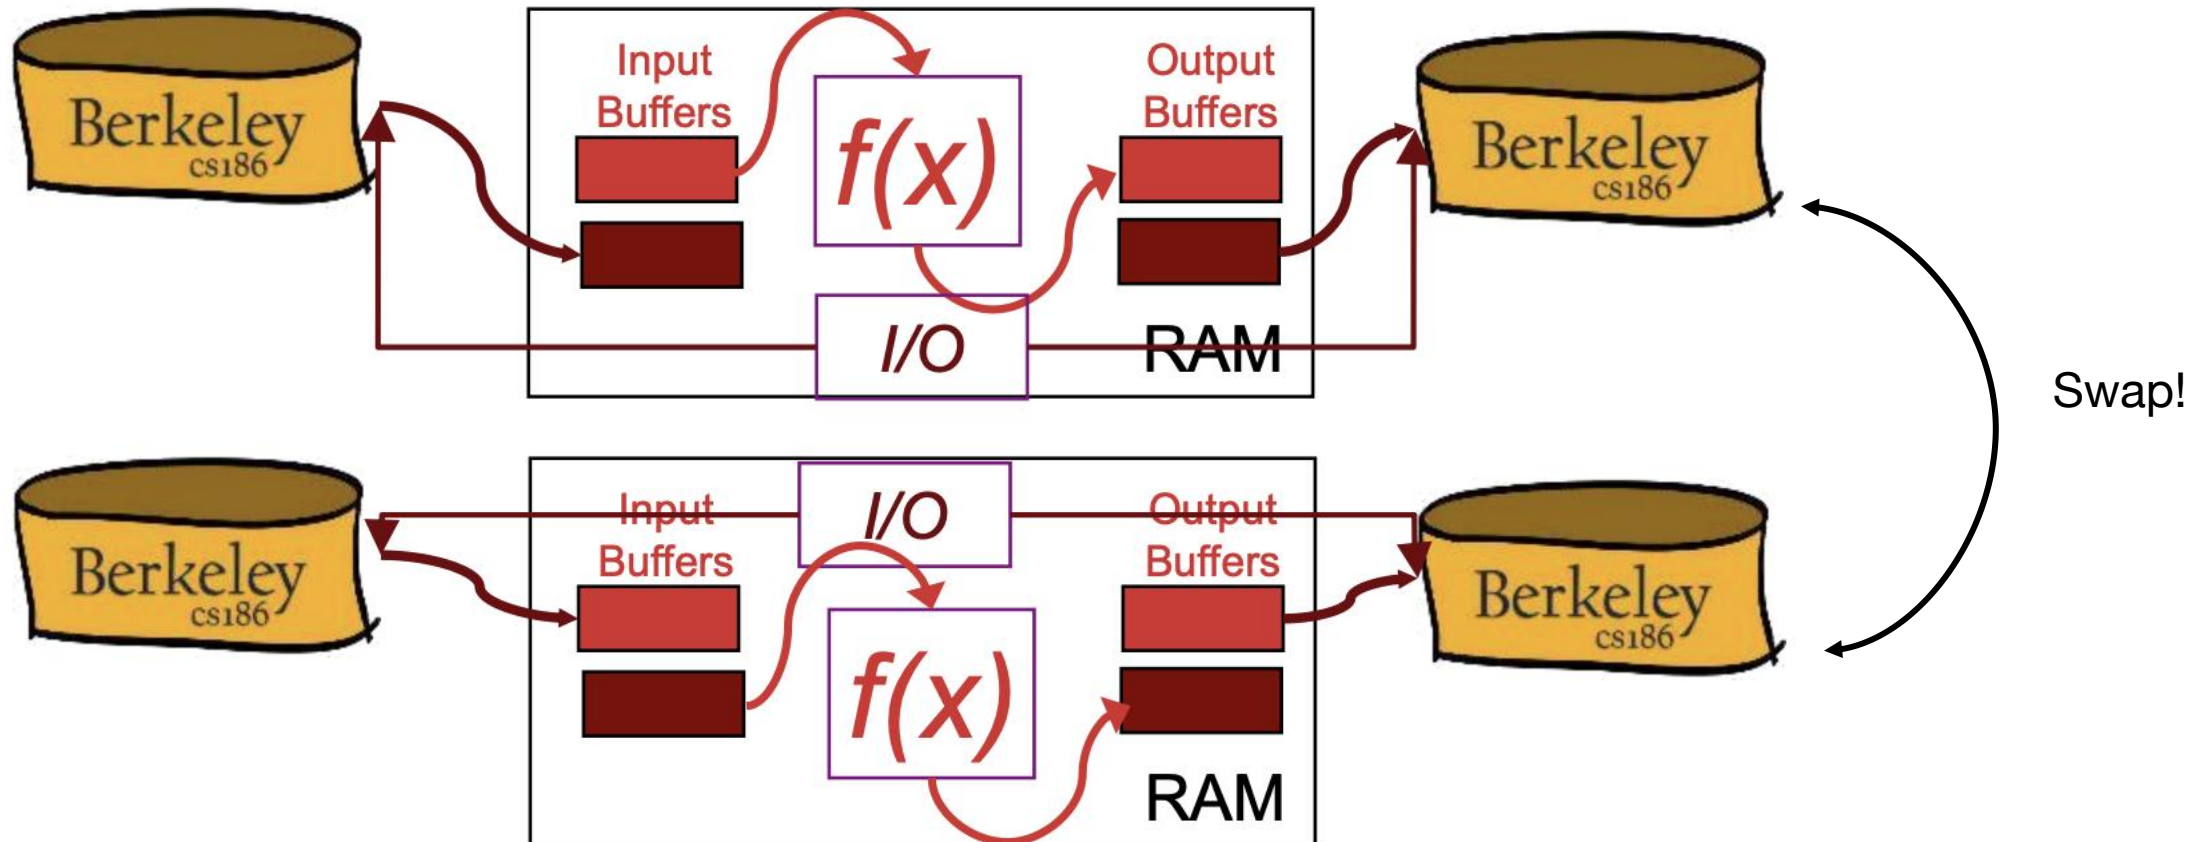
Optimization for Streaming - Double Buffering

- Main thread runs $f(x)$ on one pair I/O bufs
- 2nd I/O thread drains/fills unused I/O bufs in parallel
- so the IOs and $f(x)$ computation are in parallel
- why available: the parallelism of disk



Optimization for Streaming - Double Buffering pt 2

- approach: each thread do their job independently, when they finish their current job, they swap for new buffers.



Sorting: 2-Way


- Pass 0 (conquer a batch):
 - read a page from disk, sort it in RAM, write it to disk.
 - only one buffer page is used in memory
 - a repeated “batch job”
- Pass 1, 2, 3, ..., etc. (merge via streaming):
 - repeatedly pick min item of buffer 1,2 to buffer 3 until full
 - merge pairs of runs into runs twice as long
 - pass1 generates 2-batch-long runs
 - pass2 generates 4-batch-long runs, etc.
 - a streaming algorithm, no double buffering
 - Drain/fill buffers as the data streams through them

#passes:

$$= \lceil \log_2 N \rceil + 1$$

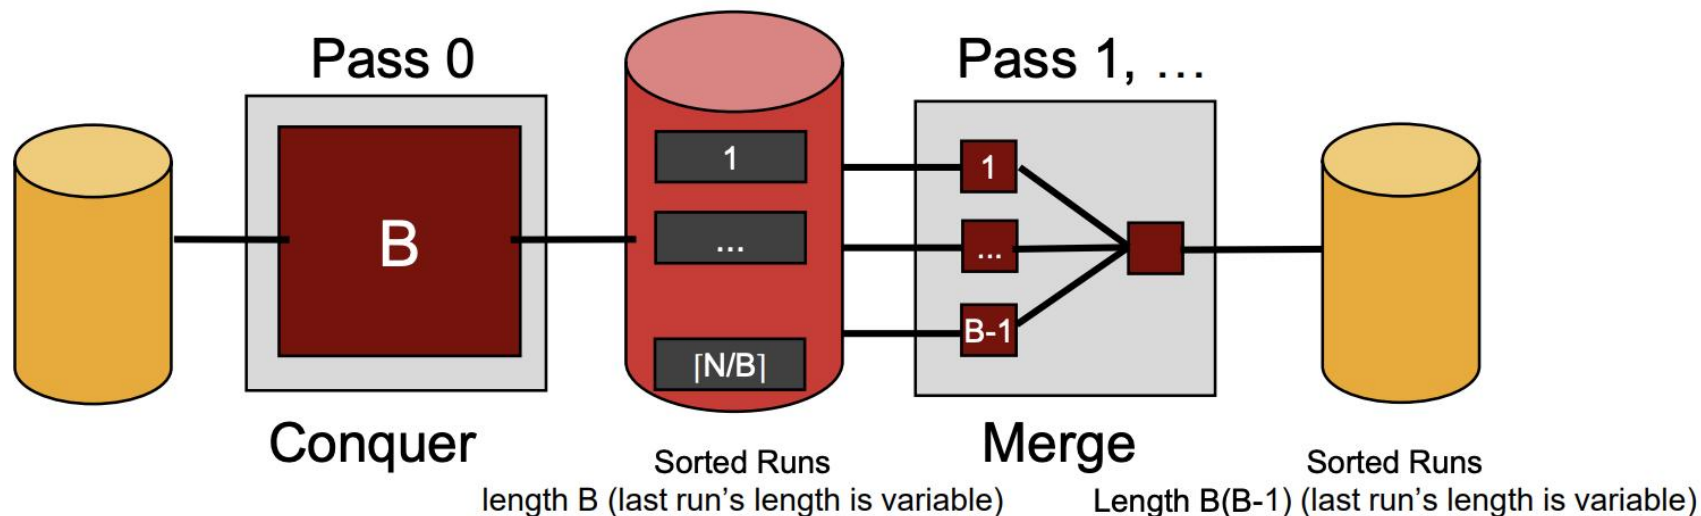
Total IOs:

$$2N(\lceil \log_2 N \rceil + 1)$$



General External Merge Sort

- B buffer pages in total
 - Pass 0: use B buffer pages. Produce $\lceil N/B \rceil$ sorted runs of B pages each.
 - Pass 1, 2, ..., etc.:
 - repeatedly pick min item of buffer 1, 2, ..., B-1 to buffer B until full
 - merge B-1 runs at a time.
 - pass1 generates B(B-1)-batch-long runs, etc.



#passes:

$$1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$$

Total IOs:


$$2 * N * (1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil)$$

Memory Requirement:
B(B-1).


Divide and Conquer in Hashing

- Streaming Partition (divide):
- Use a hash function h_p to stream records to disk partitions
 - All matches rendezvous in the same partition.
 - Each partition a mix of values
 - Streaming alg to create partitions on disk:
 - “Spill” partitions to disk via output buffers
 - Note:
 - same value item may not be contiguous (see later conquer)
 - pages from output after partition can be more than pages from input

Divide and Conquer in Hashing pt 2




Example




Divide and Conquer in Hashing pt 3

- conquer - rehash:
 - Read partitions into RAM hash table one at a time, using different hash function h_r
 - Each bucket contains a small number of distinct values
 - Then read out the RAM hash table buckets and write to disk
 - Ensuring that duplicate values are contiguous (same value same bucket)




Example



Cost of 2-pass external hashing



- cost $\sim 2 \cdot N \cdot (\text{\#passes}) = 4 \cdot N$ IOs
- memory requirement: $B(B-1)$
- Duality: the cost of 2-pass external hashing equals to the cost of 2-pass external sorting
 - Hashing: Divide & Conquer
 - Sorting: Conquer & Merge

Recursive Partitioning



Parallelism of Hashing and Sorting

- 1 scan of entire data set in parallel across all machines
- cost of 2-pass with n machines: $4N/n$ IOs
- range separation for sorting in parallel
 - avoid skew: estimate the distribution



- The question form may appear in exams
- Quiz Overview