

# Discussion 8

Transactions and Concurrency & Quiz 5

# Concurrent Execution

- Advantages:
  - **Throughput** (transactions per second)
  - **Latency** (response time per transaction)
- Problems:
  - **Inconsistent Reads** (Write-Read Conflict)
  - **Lost Update** (Write-Write Conflict)
  - **Dirty Reads** (Write-Read Conflict)
  - **Unrepeatable Reads** (Read-Write Conflict)

# Interleave transactions

- A transaction is a sequence of one or more operations (reads or writes)
- should be executed as a single, logical, atomic unit
- Property
  - **Atomicity**: A transaction ends in two ways: it either **commits** or **aborts**: either all actions in the transaction happen, or none happen.
  - **Consistency**: If the DB starts out consistent, it ends up consistent at the end of the transaction.
  - **Isolation**: Execution of each transaction is isolated from that of others. Each transaction executes as if it ran by itself.
  - **Durability**: If a transaction commits, its effects persist.

# Concurrency Control

Serial schedule:

- run all the operations of one transaction to completion before beginning the operations of next transaction
  - enforce the **isolation** property of transactions
  - not **efficient** to wait for an entire transaction to finish before starting another one

Serializable Schedule: results **equivalent** to a serial schedule

Equivalent:

- involve the same transactions
- operations are ordered the same way within the individual transactions
- each leave the database in the same state

# Example

## Scheduling examples

Serial schedule  $T_1, T_2$ :

$T_1$  A += 100 B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

Starting Balance

A	B
\$50	\$200

A	B
\$159	\$106

Same result!

Interleaved schedule A:

$T_1$  A += 100 B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

A	B
\$159	\$106

## Scheduling examples

Serial schedule  $T_1, T_2$ :

$T_1$  A += 100 B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

Starting Balance

A	B
\$50	\$200

A	B
\$159	\$106

Different result than serial  $T_1, T_2$ !

Interleaved schedule B:

$T_1$  A += 100

B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

A	B
\$159	\$112

## Scheduling examples

Serial schedule  $T_2, T_1$ :

$T_1$  A += 100 B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

Starting Balance

A	B
\$50	\$200

A	B
\$153	\$112

Different result than serial  $T_2, T_1$  ALSO!

Interleaved schedule B:

$T_1$  A += 100

B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

A	B
\$159	\$112

## Scheduling examples

Interleaved schedule B:

$T_1$  A += 100

B -= 100

$T_2$  A \*= 1.06 B \*= 1.06

This schedule is different than *any* serial order! We say that it is not serializable

equivalent to a serial schedule

# Conflict Serializability

## Conflicting operations:

- The operations are from different transactions
- Both operations operate on the same resource
- At least one operation is a write
- The order of non-conflicting operations has no effect on the final state of the database!

## Conflict equivalent schedules:

- involve the same actions of the same transactions
- every pair of conflicting actions is ordered the same way

## Conflict serializable:

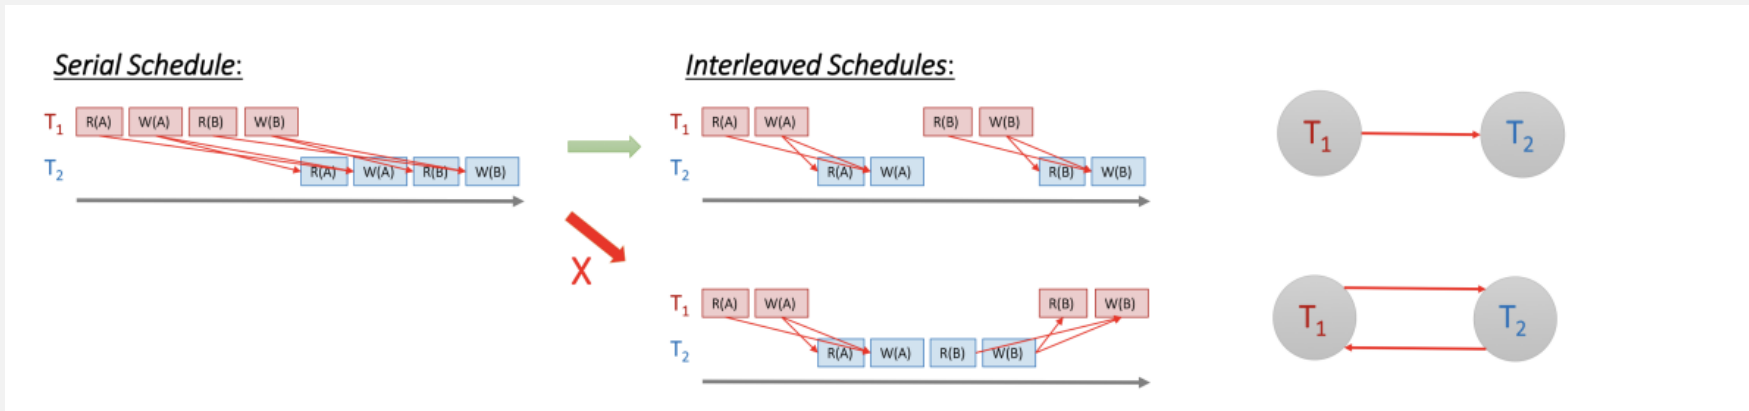
- conflict equivalent to some serial schedule
- also Serializable

# Conflict Dependency Graph

- One node per transaction
- Edge from  $T_i$  to  $T_j$  if:
  - an operation  $O_i$  of  $T_i$  conflicts with an operation  $O_j$  of  $T_j$
  - $O_i$  appears earlier in the schedule than  $O_j$



A schedule is **conflict serializable** if and only if its dependency graph is **acyclic**.





# Two phase locking (2PL)

- Transactions must acquire a S (shared) lock before reading, and an X (exclusive) lock before writing.

Only **one** transaction may hold an **exclusive lock** on a resource, but **many** transactions can hold a **shared lock** on data.

- Transactions cannot acquire new locks after releasing any locks
  - key to enforcing **serializability** through locking!
  - does not prevent **cascading aborts**
- Strict Two Phase Locking
  - locks get released together when the transaction completes





# Lock Management

- manages lock and unlock (or acquire and release) requests
- maintains a hash table, keyed on names of the resources being locked

	Granted Set	Mode	Wait Queue
A	{T1,T2}	S	T3(X)->T4(X)
B	{T6}	X	T5(X)->T7(X)

- **Granted set** (set of granted locks/the transactions holding the locks for each resource)
- **Lock type** (S or X or other)
- **Wait queue** (queue of lock requests that cannot yet be satisfied because they conflict with the locks that have already been granted)

# Lock Management

- When a lock request arrives:
  - Lock Manager checks if any transactions in the Granted Set or in the Wait Queue want a conflicting lock
    - Yes: put the requester into **Wait Queue**
    - No: the requester is granted the lock and put into the **Granted Set**
- Lock upgrade:
  - a transaction with **shared lock** can request to upgrade to exclusive
  - add this upgrade request at the front of the queue

# Deadlock

- a cycle of transactions waiting for locks to be released by each other
- Prevention:
  - resource ordering - Screen < Network Card < Printer
- Avoidance: Assign priorities based on age: (now – start\_time).

If  $T_i$  wants a lock that  $T_j$  holds,

- **Wait-Die**: If  $T_i$  has higher priority,  $T_i$  waits for  $T_j$ ; else  $T_i$  aborts
- **Wound-Wait**: If  $T_i$  has higher priority,  $T_j$  aborts; else  $T_i$  waits
- Detection: detect deadlocks by creating and maintaining a “waits-for” graph

# Deadlock Detection


- “waits-for” graph
- have one node per transaction
- an edge from  $T_i$  to  $T_j$  if :
  - $T_j$  holds a lock on resource  $X$
  - $T_i$  tries to acquire a lock on resource  $X$ , but  $T_j$  must release its lock on resource  $X$  before  $T_i$  can acquire its desired lock.
- periodically check for cycles in the graph
- “shoot” a transaction in the cycle and abort it

Example:

T1: X(A)

T2: S(A) X(B)

T3: S(B) X(A)




# Reminder

- **Conflict Dependency Graph**: Draw an edge from  $T_i$  to  $T_j$  iff
  - $T_j$  and  $T_i$  operate on the same resource, with  $T_i$  operation preceding  $T_j$  op
  - At least one of  $T_i$  and  $T_j$  is a write
  - Used to determine **conflict serializability**
- **Waits-For Graph**: Draw an edge from  $T_i$  to  $T_j$  iff
  - $T_j$  holds a conflicting lock on the resource  $T_i$  wants to operate on, meaning  $T_i$  must wait for  $T_j$
  - Used for **deadlock detection**

# Lock Granularity

- when we place a lock on a node, we implicitly lock all of its children as well
  - Database -> Tables -> Pages -> Records
- Granularity of locking
  - **Fine granularity** (lower in tree): High concurrency, lots of locks (high overhead)
  - **Coarse granularity** (higher in tree): Few locks (low overhead), **lost concurrency**



# Lock Granularity

- 3 new lock modes:
  - IS: Intent to get S lock(s) at finer granularity
  - IX: Intent to get X lock(s) at finer granularity
  - SIX: Like S and IX at the same time
- Compatibility Matrix

Mode	NL	IS	IX	S	SIX	X
NL	Yes	Yes	Yes	Yes	Yes	Yes
IS	Yes	Yes	Yes	Yes	Yes	No
IX	Yes	Yes	Yes	No	No	No
S	Yes	Yes	No	Yes	No	No
SIX	Yes	Yes	No	No	No	No
X	Yes	No	No	No	No	No

# Multiple Granularity Locking Protocol

- Each Transaction starts from the root of the hierarchy.
- To get S or IS lock on a node, must hold IS or IX on parent node.
- To get X or IX on a node, must hold IX or SIX on parent node.
- Must release locks in bottom-up order.
- 2-phase and lock compatibility matrix rules enforced as well.
- Protocol is correct in that it is equivalent to directly setting locks at leaf levels of the hierarchy.



# Quiz 5

**1.Q 1: T/F - If a term has a large reduction factor, the output of the query will have more tuples than if it had a small reduction factor. \***

*Mark only one oval.*

☐ True

☐ False

**True** - a larger reduction factor means higher  $\frac{|\text{output}|}{|\text{input}|}$  ratio

**2.Q 2: T/F - An equidepth histogram gives better resolution on low-frequency entries than a equiwidth histogram. \***

*i.e. it gives more detailed information for these entries.*

*Mark only one oval.*

☐ True

☐ False

**False** - in an equidepth histogram, if you have an entry with low frequency, then it's either going to get clumped together with a bunch of other low frequency entries or with 1 other high frequency entry.

3. **Q 3: When doing a cross join on tables A, B, C, and D, which of the following query plans do we consider? \***

Mark all that apply.

*Check all that apply.*

- ☐ None of the above
- ☐ (A join (B join C)) join D
- ☐ A join ((B join C) join D)
- ☒ ((A join B) join C) join D
- ☐ A join (B join (C join D))
- ☐ (A join B) join (C join D)

left deep joins only

4. **Q4: Which of the following access or join methods will result in an interesting order in a query where we require the output to be sorted? \***

*Check all that apply.*

- ☐ File scan
- ☒ Sort-Merge Join
- ☐ Block-Nested Loops Join
- ☒ Clustered Index Traversal
- ☐ Hash Join

Suppose that we have three tables, R, S, and T. We are running the following query:

```
SELECT *  
FROM R, S, T  
WHERE R.a = S.a  
AND S.b = T.b;
```

Assume that our database has no indices and that none of the relations are sorted in any interesting or useful way. Since we only have one possible single-table access method for each table, we ignore the costs of accessing a single table.

Assume that all provided join costs are for the optimal join algorithm for that join.

These are the two-table join costs:

- 1) S join R = 2,000
- 2) R join S = 6,000
- 3) R join T = 5,000
- 4) T join R = 1,000
- 5) T join S = 3,000
- 6) S join T = 4,000

**5.Q 5: Which of the above two-table join plans will be selected? \***

*Check all that apply.*

1, 5

1)+5) S join R + T join S 5,000

1)+6) S join R + S join T 6,000

2)+5) R join S + T join S 9,000

2)+6) R join S + S join T 10,000

We now add the third table and have the following join costs:

- 1) (R join S) join T = 10,000
- 2) T join (R join S) = 6,000
- 3) R join (S join T) = 12,000
- 4) T join (S join R) = 11,000
- 5) (R join T) join S = 10,000
- 6) S join (R join T) = 7,000
- 7) (T join R) join S = 14,000
- 8) S join (T join R) = 16,000
- 9) (S join T) join R = 13,000
- 10) (S join R) join T = 15,000
- 11) (T join S) join R = 20,000
- 12) R join (T join S) = 9,000

**6.Q 6: Which of these will the optimizer select as your final query plan? \***

*Mark only one oval.*

For example, the best left-deep plan to join tables A, B, C is either:

- (The best plan for joining A, B)  $\bowtie$  C
- (The best plan for joining A, C)  $\bowtie$  B
- (The best plan for joining B, C)  $\bowtie$  A

10

S join R 2,000

R join S 6,000

10) (S join R) join T 15,000

S join T 4,000

T join S 3000

11) (T join S) join R 20,000