

# Discussion 9

# Recovery

Binbin Chen

[chenbb@shanghaitech.edu.cn](mailto:chenbb@shanghaitech.edu.cn)

# Motivation of Recovery

- Atomicity: all actions in a Xact should be either happen or none happen.
- Durability: If a Xact, commits, its effect persist permanently.
- Xact abort & DB crash -> recovery

# No-Steal/Force Scheme for A&D

- NO STEAL: Xact locks the page and pins it in the buffer pool
  - Can't be stolen by replacement policy, no dirty write go to DB
  - enables Atomicity
  - but not scalable and IO inefficient
- FORCE: every update is “forced” onto the DB disk before commit.
  - Provides Durability but somehow violates Atomicity for many IOs will have crash when forcing

# Preferred Scheme: Steal/No-Force

- STEAL (enforcing Atomicity)
  - we allow buffer-pool frames with uncommitted updates to be replaced or flushed to disk.
  - but has problems like Xact aborts and DB crash before end when dirty pages have flushed into DB already
  - so we need UNDO updates that should not happen
- NO FORCE (enforcing Durability)
  - we allow commit without flushing pages to the disk
  - but has problem that System crash before dirty buffer page of a committed transaction is flushed to DB disk.
  - so we need REDO

# Buffer Management Summary

	No Steal	Steal
No Force		Fastest
Force	Slowest	

Performance  
Implications

	No Steal	Steal
No Force	No UNDO REDO	UNDO REDO
Force	No UNDO No REDO	UNDO No REDO

Logging/Recovery  
Implications

Simplified policy:

REDO: deal with commit before flush

UNDO: deal with flush before commit

# Write-Ahead Logging (WAL)

- Log: An ordered list of log records to allow REDO/UNDO, with a write buffer (“tail”) in RAM.
  - Each log record has a unique Log Sequence Number (LSN)
  - we can only write records to tail, tail buffer will periodically flush to log’s end on disk
- The Write-Ahead Logging Protocol:
  - Must force the log record write to log device before the corresponding data page gets to the DB device every time.
  - Must force all log records write to log device for a Xact before commit.

# LSN

- flushedLSN (= largest diskLSN) tracked in RAM
- Each data page in the DB contains a pageLSN, i.e. the LSN of the most recent log record for an update to that page.
- Before page  $i$  is flushed to DB, log must satisfy:  $\text{pageLSN}_i \leq \text{flushedLSN}$ 
  - because pageLSN is now larger than flushLSN, we cannot write page to DB, but with diskLSN increasing, when pageLSN appears in log device, we can flush page to disk
  - this allows steal
- prevLSN is the LSN of the previous log record written by this  
XID

# State in Memory

- Transaction Table
  - One entry per currently active Xact
  - Contains: XID, Status (running, committing, aborting), lastLSN (most recent LSN written by Xact)
- Dirty Page Table
  - One entry per dirty page currently in buffer pool.
  - Contains recLSN (record which first caused the page to be dirty)




# ARIES Big Picture: What's Stored Where



## LogRecords

LSN  
**prevLSN**  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



Data pages  
each with a  
**pageLSN**

Master record



## Xact Table

xid  
**lastLSN**  
status

## Dirty Page Table

pid  
**recLSN**

Log tail  
**flushedLSN**

Buffer pool

# Execution of Xact

- commit: All log records up to Xact's commit record are flushed to disk
- abort: need CLR (compensation log record) with undonextLSN for each undone operation.
  - CLR contains REDO info, CLR's never Undone so exactly undo once
- checkpoint: Store LSN of most recent chkpt record in a safe place (master record)

# Recovery Protocol: 3-phase

- Analysis - Scan log forward from checkpoint.
  - end, commit, update records
  - if commit then remove the Xact from Xact table, the remaining is all active and need abort for Atomicity
- REDO all actions including abort and CLR's. (repeat all history)
  - unless updates are flushed into DB already, like:
    - Affected page is not in the Dirty Page Table, or
    - Affected page is in D.P.T., but has recLSN > LSN, or
    - pageLSN (in DB) >= LSN. (this last case requires I/O)
- UNDO effects of failed Xacts.
  - to deal with Atomicity, we will backward from end to first LSN of oldest Xact alive (running, aborting) after Redo.
  - can optimize via CLR loops