

Discussion 14

No SQL

Two Classes of Relational Database Apps

- OLTP (Online Transaction Processing)
 - Queries are simple lookups: 0 or 1 join
E.g., find customer by ID and their orders
 - Many updates. E.g., insert order, update payment
 - **Consistency** is critical: we need transactions
- OLAP (Online Analytical Processing)
 - aka “Decision Support”
 - Queries have many joins, and group-by's
E.g., sum revenues by store, product, clerk, date
 - **No updates**

Two Classes of Relational Database Apps

Property	OLTP	OLAP
Main read pattern	Small number of records per query, fetched by key	Aggregate over large number of records
Main write pattern	Random-access, low-latency writes from user input	Bulk import or event stream
Primarily used by	End user/customer, via web application	Internal analysis, for decision support
What data represents	Latest state of data (current point in time)	History of events that happend over time
Dataset size	Gigabytes to terabytes	Terabytes to petabytes

Motivation

- Originally motivated by Web 2.0 applications
- Needed: very large scale OLTP workloads
- Give up on consistency, give up OLAP
- NoSQL: reduce functionality
 - Simpler data model
 - Very restricted updates

Scaling the Database

- Two basic approaches:
 - Scale up through **partitioning** – “sharding”
 - **Partition the database** across machines in a cluster
 - Can increase throughput
 - Easy for **writes** but **reads** become expensive!
 - Scale up through **replication**
 - Create **multiple copies** of each database partition. Spread queries across these replicas
 - Can increase throughput and lower latency, improve fault-tolerance
 - Easy for **reads** but **writes** become expensive!
- **Consistency** is much harder to enforce

Relational Model NoSQL

- Relational DB: difficult to replicate/partition. E.g.,
`Supplier(sno,...),Part(pno,...),Supply(sno,pno)`
 - Partition: we may be forced to join across servers
 - Replication: local copy has inconsistent versions
 - **Consistency** is hard in both cases
- NoSQL: simplified data model
 - Given up on functionality
 - Application must now handle joins and consistency

ACID vs BASE

- Relational DB
 - **A**tomicity
 - **C**onsistency
 - **I**solation
 - **D**urability
- NoSQL
 - **B**asic **A**vailability
 - Application must handle partial failures itself
 - **S**oft State
 - DB state can change even without inputs
 - **E**ventually Consistency
 - DB will “eventually” become consistent
- i.e., ACID vs BASE



Data Models

Taxonomy based on data models:

- **Key-value stores**
 - e.g., Amazon Dynamo, Voldemort, Memcached
- **Extensible Record Stores**
 - e.g., HBase, Cassandra, PNUTS
- **Document stores**
 - e.g., SimpleDB, CouchDB, MongoDB

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - `get(key)`, `put(key,value)`
 - Operations on value not supported
- **Distribution / Partitioning** – w/ hash function
 - No replication: key k is stored at server $h(k)$
 - Multi way-way replication: e.g., key k stored at $h1(k),h2(k),h3(k)$

Key-Value Stores Internals

- Partitioning:
 - Use a hash function h
 - Store every (key,value) pair on server $h(\text{key})$
- Replication:
 - Store each key on (say) three servers
 - On update, propagate change to the other servers; *eventual consistency*
 - Issue: when an app reads one replica, it may be stale
- Usually: combine partitioning+replication

Document stores: Motivation

- In Key-Value stores, the Value is often a very complex object
 - Key = '2010/7/1', Value = [all flights that date]
- Better: *value* to be structured data
 - **JSON** or Protobuf or XML
 - Called a “document” but it’s just data

JSON vs Relational

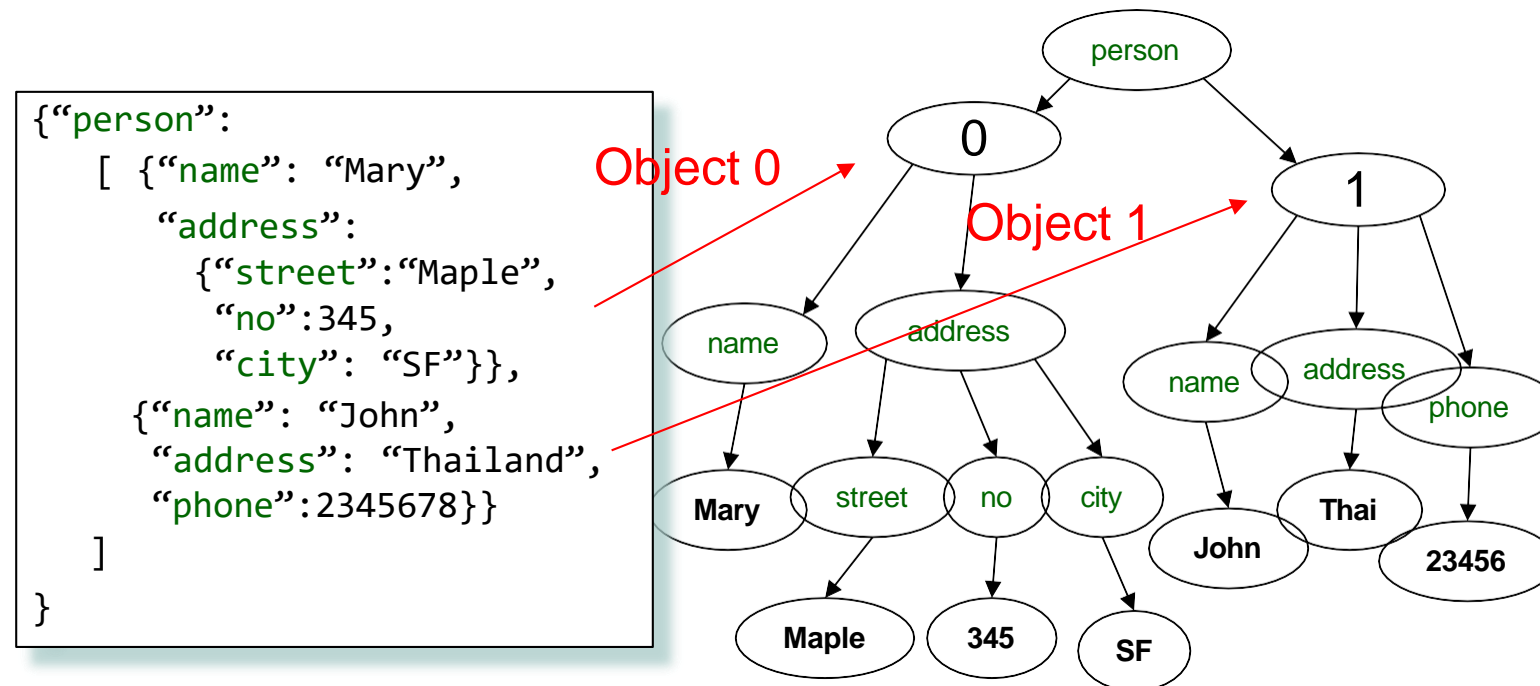
- Relational data model
 - Rigid flat structure (tables)
 - Schema must be fixed in advanced
 - Binary representation: good for performance, bad for exchange
 - Query language based on Relational Algebra
- Semistructured data model / JSON
 - Flexible, nested structure (trees)
 - Does not require predefined schema ("self-describing")
 - Text representation: good for exchange, bad for performance
 - Most common use: Language API; query languages emerging

```
{  "book": [  
    {"id": "01",  
     "language": "Java",  
     "author": "H. Javeson",  
     "year": 2015  
    },  
    {"id": "07",  
     "language": "C++",  
     "edition": "second",  
     "author": "E. Sepp",  
     "price": 22.25  
    }  
  ]  
}
```

JSON Types

- Primitive: number, string, Boolean, null
- Object: collection of name-value pairs:
 - {“name1”: value1, “name2”: value2, ...}
 - “name” is also called a “key”
- Array: *ordered* list of values:
 - [obj1, obj2, obj3, ...]

JSON Semantics: a Tree !



Recall: arrays are *ordered* in JSON!

Relational -> JSON

Single Table:

- key = table name
- value: array of objectives
- objective <-> row in table, with key = table attribute, value = tuple element

One-to-many relationship:

- nested structure
- linked rows in the 2nd table are embedded or “inlined” into the objects representing the row of the 1st table

Relational -> JSON

Many-to-many relationship

First option: represent each relation as a flat JSON array, each element is an object representing a row

- no redundant data, but relies on the application to join the tables to get related data

The second and third options: key off of the one of the primary-keyed tables, inline the contents of the relationship table and the last table within each object

- the elements of the last table will be duplicated for each linked element of the first table
- application does not need to perform any joins to get the full record
- updates to the document store will need to handle updating all duplicated rows

Intro to Semi-structured Data

- JSON is **self-describing**
- Schema elements become part of the data
 - Relational schema: `person(name, phone)`
 - In JSON “`person`”, “`name`”, “`phone`”
are part of the data, and are repeated many times
- \Rightarrow JSON is more flexible
 - Schema can change per tuple

Discussion: Why Semi-Structured Data?

- Semi-structured data works well as *data exchange formats*
 - i.e., exchanging data between different apps
 - Examples: XML, JSON, Protobuf (protocol buffers)
- Increasingly, systems use them as a data model for DBs:
 - SQL Server supports for XML-valued relations
 - CouchBase, MongoDB, Snowflake: JSON
 - Dremel (BigQuery): Protobuf

Storing JSON in RDBMS

- Using JSON as a data type provided by RDBMSs
 - Declare a column that contains either json or jsonb (binary)
 - `CREATE TABLE people (person json)` [or jsonb for binary]
 - In our previous example, we will have one row per person
 - i.e., a row corresponding to that person's attributes
 - Queries now mix relational and semi-structured syntax
 - `SELECT * FROM people`
`WHERE person @> '{"name": "Mary"}';`
- Translate JSON documents into relations

Semistructured Data Model

- Several file formats: JSON, protobuf, XML
- Data model = Tree
- Query language take non first normal form into account as we will see
 - Various “extra” constructs introduced as a result
 - Nesting & Unnesting, strict aggregates, splitting