

Discussion 3

File organization & Index Files
& Quiz3

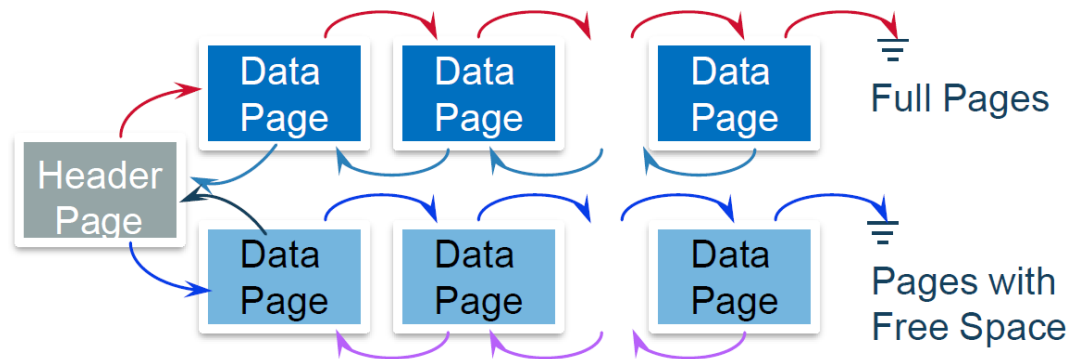
Jiahui Xu

Database File Structures

- Heap Files
 - Implemented as List
 - Better: Use a Page Directory
- Sorted Files
- Index Files
 - ISAM
 - B+ Tree
- Cost of Operations


Heap Files - Linked List

- Each data page contains:
 - records, a free space tracker, and pointers (byte offsets) to the next and previous page
- One header page
 - separates the data pages into full pages and free pages



Heap Files - Linked List


within box = # of free bytes




Insert 20 byte record

Heap Files - Linked List

within box = # of free bytes



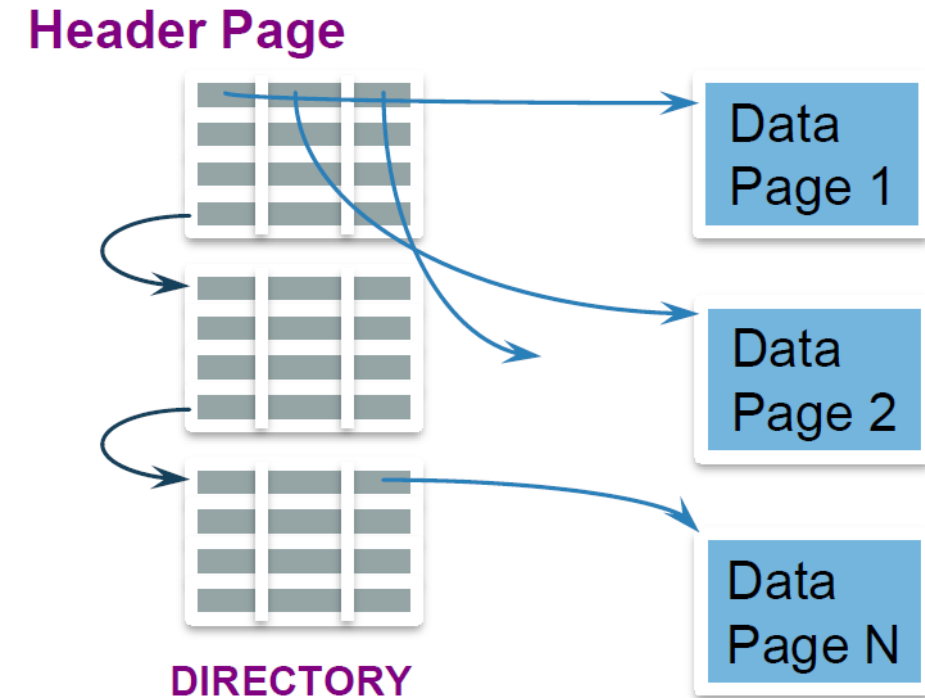
Insert 20 byte record



I/O Cost: 5 (read all pages) + 1 (write last data page) = **6 I/Os**

Heap Files - Page Directory

- using a linked list for header pages
- Each header page contains
 - a pointer (byte offset) to the next header page
 - its entries contain
 - a pointer to a data page
 - the amount of free space left within that data page
- Inserting records is often faster
 - Read the header pages until finding a page with enough free space -> Use pointer to that data page to read the page -> Write data -> Write header



Heap Files - Page Directory


(#, #) within box = (page #, # of free bytes)

(1, 5)	(2, 15)
(3, 10)	(4, 25)

Header Page


Heap Files - Page Directory

(#, #) within box = (page #, # of free bytes)



Header Page

Insert 20 byte record



Header Page

Insert at last page
with enough free
space

I/O Cost: 1 (read header) + 1 (read data) + 1 (write data) + 1 (write header) = **4 I/Os**

Heap Files - Note

- You should **ignore the I/O cost associated with reading/ writing the file's header pages** when the specific file implementation of heap files (i.e., heap file implemented with a linked list or page directory) **is not provided**.
- But if it is provided:
 - You must consider it

Heap Files vs Sorted Files

Heap File



Sorted File



For illustration, records are just integers

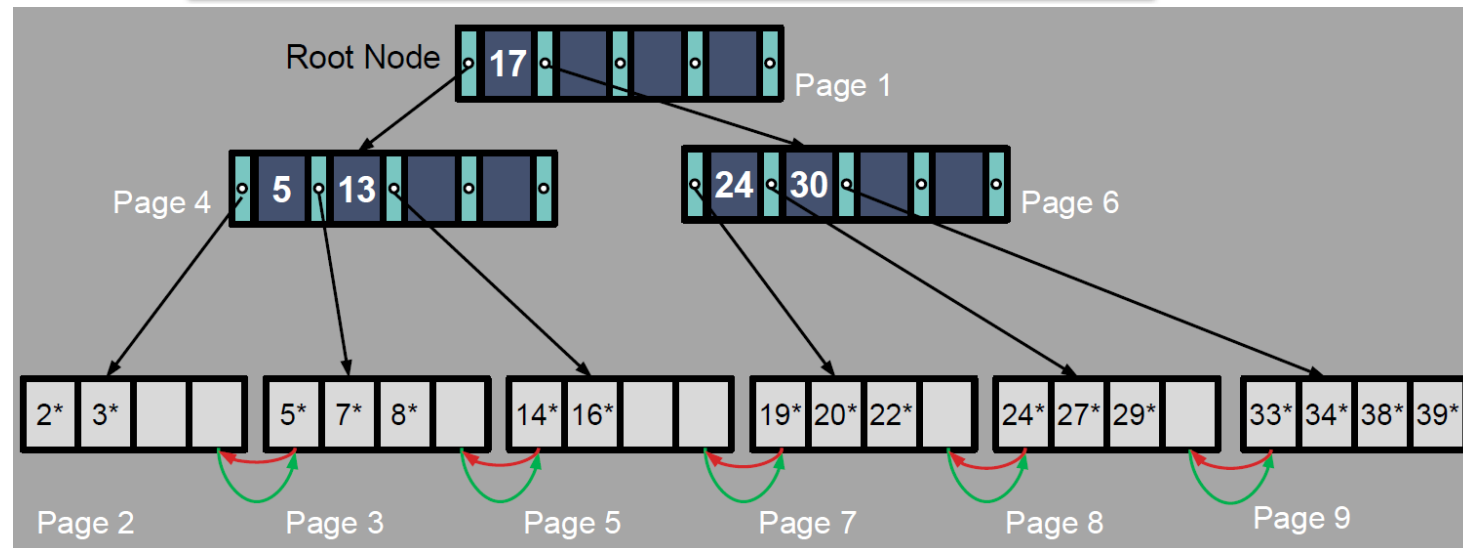

- **B:** The number of data blocks
- **R:** Number of records per block
- **D:** Average time to read/write disk block

	Heap File	Sorted File
Scan all records	$B * D$	$B * D$
Equality Search	$0.5 * B * D$	$(\log_2 B) * D$
Range Search	$B * D$	$\frac{((\log_2 B) + \text{pages}) * D}{D}$
Insert	$2 * D$	$((\log_2 B) + B) * D$
Delete	$(0.5 * B + 1) * D$	$((\log_2 B) + B) * D$


Index Files

ISAM

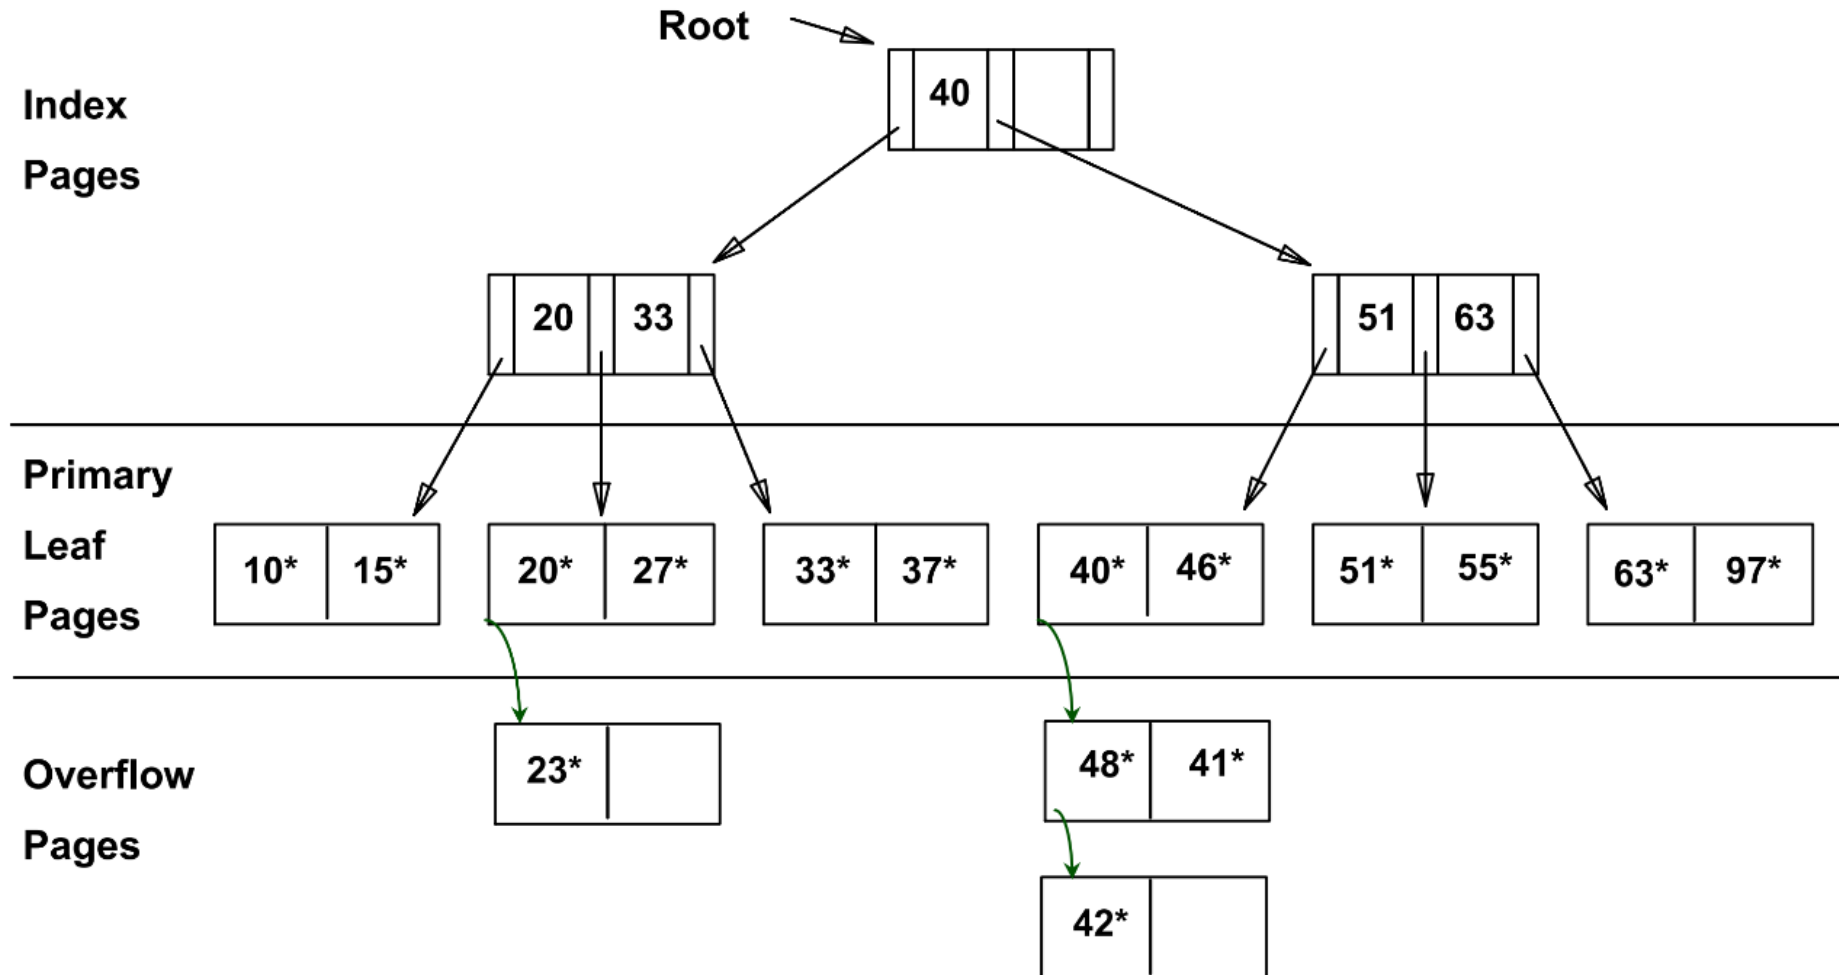
B+ Tree



ISAM - Insert

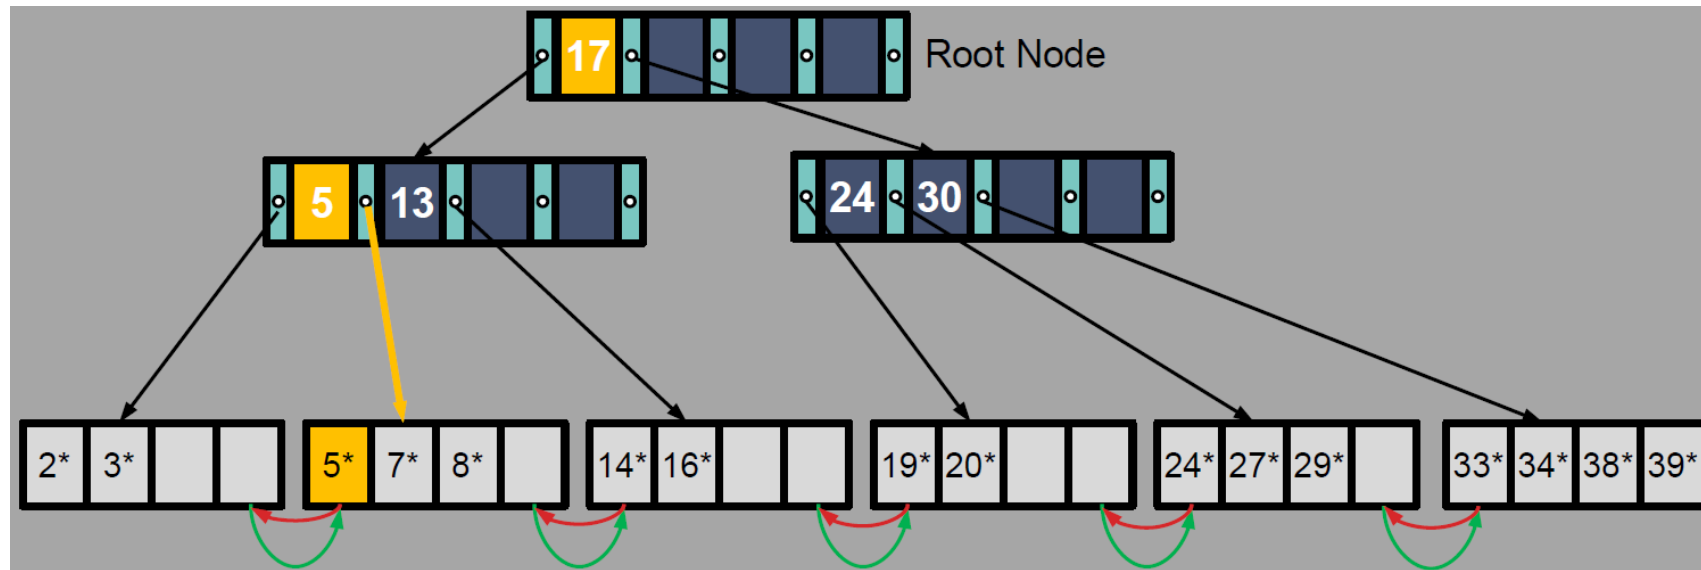


ISAM - Insert



B+ Tree - Insert

- If there is room in the leaf just add the entry
 - **Sort** the leaf page by key
- If no room:
 - Split-> Redistribute entries evenly -> **leaf**: Fix next/ prev pointers -
> **Copy (leaf) / Push (parent)** up from leaf the middle key



Index Files

ISAM

- Data pages at bottom stored in logical order
- Insert into overflow pages
 - Overflow chains can degrade performance unless size of data set and data distribution stay constant
- Only leaf nodes are allowed to be modified, not tree nodes.

B+ Tree

- Data pages at bottom:
 - Use next and prev pointers
 - Don't need to be stored in logical order
- Dynamic Tree Index
 - Always Balanced
 - No overflow pages, more efficient
 - Require Occupancy Invariant (except root node): $d \leq \#entries \leq 2d$
 - High fanout (F) means depth rarely more than 3 or 4
- Can dynamically modify root node, tree node, and leaf nodes.

Quiz 3

Q1: Suppose that all nodes in our B+ tree have an order of 1500. What's the MAXIMUM number of records we can index with a B+ tree of height 2?

Assume our B+ trees are laid out as in lecture.

Quiz 3

Q2: We want to bulk-load a B+ tree, and we increase the fill factor of this bulk load. Which of the following applies, in general?


Check all that apply.

- A. The bulk loading operation is faster
- B. The bulk loading operation is slower
- C. We consume more disk space
- D. A sequence of many consecutive record lookups is faster
- E. A sequence of many consecutive insertions requires fewer disk writes

Quiz 3

Q3: We insert the key 60 into the B+ tree in Figure A. How many I/Os (page reads and writes) does this operation take?

Assume we require zero page reads and one page write to create a new page from scratch. Also assume that we do no key redistribution. Exclude disk I/Os done to data pages. Finally, assume we have 20 pages of memory available for caching pages in memory after reading them.



Quiz 3

