

# Linear Discrimination

Prof. Ziping Zhao

School of Information Science and Technology  
ShanghaiTech University, Shanghai, China

CS182: Introduction to Machine Learning (Fall 2021)  
<http://cs182.sist.shanghaitech.edu.cn>

# Outline

Introduction

Geometric View

Parametric Discrimination Revisited

Logistic Discrimination

# Outline

Introduction

Geometric View

Parametric Discrimination Revisited

Logistic Discrimination

## Likelihood-Based vs. Discriminant-Based Classification

- Classification based on a set of discriminant functions  $g_i(\mathbf{x})$ ,  $i = 1, \dots, K$ :

$$\text{Choose } C_i \text{ if } g_i(\mathbf{x}) = \max_k g_k(\mathbf{x})$$

- Likelihood-based classification:

- Assume a **parametric**, **semiparametric**, or **nonparametric** model for the class-conditional probability densities  $p(\mathbf{x} | C_i)$ .
- Estimate the prior probabilities  $P(C_i)$  and the class likelihoods  $p(\mathbf{x} | C_i)$  from data.
- Apply Bayes' rule to compute the posterior probabilities  $P(C_i | \mathbf{x})$ .
- Perform optimal classification based on  $P(C_i | \mathbf{x})$ , or equivalently based on discriminant functions  $g_i(\mathbf{x})$  such as  $g_i(\mathbf{x}) = \log P(C_i | \mathbf{x})$ .

- Discriminant-based classification:

- Assume a model directly for the discriminant functions, bypassing the estimation of  $p(\mathbf{x} | C_i)$  or  $P(C_i | \mathbf{x})$  from data.
- Perform optimal classification based on the discriminant functions  $g_i(\mathbf{x})$ .

## Likelihood-Based vs. Discriminant-Based Classification (2)

- ▶ **Main difference:** the likelihood-based approach makes an assumption on the form of the **densities** (e.g., whether they are Gaussian, or whether the inputs are correlated, etc.), but the discriminant-based approach makes an assumption on the form of the **discriminants**.

## Discriminant Functions

- ▶ Define a model for the **discriminant functions**:

$$g_i(\mathbf{x} \mid \Phi_i)$$

which are explicitly parameterized with a set of model parameters  $\Phi_i$ .

- ▶ In discriminant-based approach, we make an assumption on the form of the boundaries separating classes.
- ▶ **Learning** is the optimization of  $\Phi_i$  to maximize the quality of the **separation**, that is, the classification accuracy on a given labeled training set.
- ▶ Unlike the likelihood-based approach which performs density estimation separately for each class, the discriminant-based approach typically estimates  $\Phi_i$  for all classes simultaneously to find the **decision boundaries** between classes.
- ▶ Estimating the class boundaries (discriminants) is usually easier than estimating the class densities. E.g., this is true when the discriminant can be approximated by a simple function.

# Linear Discriminant Functions

- ▶ Linear discriminant functions:

$$g_i(\mathbf{x} \mid \mathbf{w}_i, w_{i0}) = \mathbf{w}_i^T \mathbf{x} + w_{i0} = \sum_{j=1}^d w_{ij} x_j + w_{i0}$$

which are linear in  $\mathbf{x}$ .

- ▶ Advantages:
  - **Simplicity**:  $O(d)$  time and space complexity.
  - **Understandability**: final output is a weighted sum of attributes; magnitude and sign of weights have clear physical meaning.
  - **Accuracy**: model is quite accurate if some assumptions are satisfied, e.g., Gaussian densities for classes with shared covariance matrix.
- ▶ We should always use the linear discriminant before trying a more complicated model to make sure that the additional complexity is justified.

## Generalizing the Linear Models

- ▶ When a linear model is not flexible enough, we can use the quadratic discriminant function

$$g_i(\mathbf{x} \mid \mathbf{W}_i, \mathbf{w}_i, w_{i0}) = \mathbf{x}^T \mathbf{W}_i \mathbf{x} + \mathbf{w}_i^T \mathbf{x} + w_{i0}$$

- ▶ An equivalent way is to preprocess the input by adding **higher-order terms** (or called **product terms**).
- ▶ Example: with two inputs  $x_1$  and  $x_2$ , we define new variables

$$z_1 = x_1, \quad z_2 = x_2, \quad z_3 = x_1^2, \quad z_4 = x_2^2, \quad z_5 = x_1 x_2$$

and take  $\mathbf{z} = (z_1, z_2, z_3, z_4, z_5)^T$  as the new input. The **linear function** defined in the new  $\mathbf{z}$ -space corresponds to a **nonlinear function** in the original  $\mathbf{x}$ -space.

- ▶ Compared with defining a nonlinear function (discriminant or regression) in the original input space, defining a linear function in a nonlinearly transformed new space (called a generalized linear model) does not increase the **number of parameters** that need to be estimated significantly.



## Basis Functions

- ▶ More generally, the inputs  $\mathbf{x}$  are (nonlinearly) transformed into **basis functions**  $\phi_{ij}(\mathbf{x})$  which are **linearly combined** to define the discriminant functions:

$$g_i(\mathbf{x}) = \sum_{j=1}^k w_j \phi_{ij}(\mathbf{x})$$

- ▶ Higher-order terms mentioned before are only one set of basis functions.
- ▶ Other examples of basis functions:
  - $\sin(x_1)$
  - $\exp(-(x_1 - m)^2/c)$
  - $\exp(-\|\mathbf{x} - \mathbf{m}\|^2/c)$
  - $\log(x_1)$
  - $\mathbf{1}(x_1 > c)$
  - $\mathbf{1}(ax_1 + bx_2 > c)$

# Outline

Introduction

Geometric View

Parametric Discrimination Revisited

Logistic Discrimination

## Geometric View: Two Classes

- Discriminant function:

$$\begin{aligned}g(\mathbf{x}) &= g_1(\mathbf{x}) - g_2(\mathbf{x}) \\&= (\mathbf{w}_1^T \mathbf{x} + w_{10}) - (\mathbf{w}_2^T \mathbf{x} + w_{20}) \\&= (\mathbf{w}_1 - \mathbf{w}_2)^T \mathbf{x} + (w_{10} - w_{20}) \\&= \mathbf{w}^T \mathbf{x} + w_0\end{aligned}$$


where  $\mathbf{w}$  is the weight vector and  $w_0$  is the threshold.

- Optimal decision rule:

$$\text{Choose } \begin{cases} C_1 & \text{if } g(\mathbf{x}) > 0 \\ C_2 & \text{otherwise} \end{cases}$$

## Hyperplane

- ▶ The discriminant function defines a **hyperplane** ( $g(\mathbf{x}) = 0$ ) that divides the input space into 2 half-spaces:
  - Decision region  $\mathcal{R}_1$  for  $C_1$  ( $g(\mathbf{x}) > 0$ , i.e., positive side of the hyperplane)
  - Decision region  $\mathcal{R}_2$  for  $C_2$  ( $g(\mathbf{x}) < 0$ , i.e., negative side of the hyperplane)



- ▶ When  $\mathbf{x} = \mathbf{0}$  (i.e., the origin),  $g(\mathbf{x}) = w_0$ . If  $w_0 > 0$ , the origin is on the positive side, and if  $w_0 < 0$ , the origin is on the negative side, and if  $w_0 = 0$ , the hyperplane passes through the origin.

## Geometric Interpretation

- ▶ Let  $\mathbf{x}_1$  and  $\mathbf{x}_2$  be two points on the hyperplane, i.e.,  $g(\mathbf{x}_1) = g(\mathbf{x}_2) = 0$ . So

$$\mathbf{w}^T \mathbf{x}_1 + w_0 = \mathbf{w}^T \mathbf{x}_2 + w_0$$

$$\mathbf{w}^T (\mathbf{x}_1 - \mathbf{x}_2) = 0$$

showing that  $\mathbf{w}$  is **normal (orthogonal)** to any vector  $(\mathbf{x}_1 - \mathbf{x}_2)$  lying on the hyperplane.

- ▶ Let us express any point  $\mathbf{x}$  as

$$\mathbf{x} = \mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

where

$\mathbf{x}_p$  : normal projection of  $\mathbf{x}$  onto the hyperplane

$r$  : distance from  $\mathbf{x}$  to the hyperplane ( $r > / < 0$  :  $\mathbf{x}$  is on the positive/negative side)


## Geometric Interpretation (2)

- Calculation of  $r$  (note  $g(\mathbf{x}_p) = 0$ ):

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T \mathbf{x}_p + r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_0 = g(\mathbf{x}_p) + r \|\mathbf{w}\| = r \|\mathbf{w}\|$$

So we have

$$r = \frac{g(\mathbf{x})}{\|\mathbf{w}\|} \quad (\text{sign of } r = \text{sign of } g(\mathbf{x}))$$



## Geometric Interpretation (3)

- ▶ When  $\mathbf{x} = \mathbf{0}$ , the distance from origin to hyperplane is  $\frac{g(\mathbf{0})}{\|\mathbf{w}\|} = \frac{w_0}{\|\mathbf{w}\|}$ .
- ▶ Alternative view: If  $\mathbf{x}$  is a point on the hyperplane, then  $g(\mathbf{x}) = 0$ . So

$$\begin{aligned}\mathbf{w}^T \mathbf{x} + w_0 &= 0 \\ \left( \frac{\mathbf{w}}{\|\mathbf{w}\|} \right)^T \mathbf{x} + \frac{w_0}{\|\mathbf{w}\|} &= 0 \\ \left| \left( \frac{\mathbf{w}}{\|\mathbf{w}\|} \right)^T \mathbf{x} \right| &= \frac{w_0}{\|\mathbf{w}\|}\end{aligned}$$

- ▶ The **orientation** of the hyperplane is determined by  $\mathbf{w}$  and its **distance** from the origin is determined by  $w_0$  and  $\mathbf{w}$ .

## Geometric View: Multiple Classes


- $K$  discriminant functions:

$$g_i(\mathbf{x} \mid \mathbf{w}_i, w_{i0}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$$

- Linearly separable classes:

$$g_i(\mathbf{x} \mid \mathbf{w}_i, w_{i0}) = \begin{cases} > 0 & \text{if } \mathbf{x} \in C_i \\ \leq 0 & \text{otherwise} \end{cases}$$

For each class  $C_i$ , there exists a hyperplane  $H_i$  such that all  $\mathbf{x} \in C_i$  lie on the positive side and all other  $\mathbf{x} \in C_j, j \neq i$  lie on the negative side.





## Linear Classifier


- ▶ During testing, given  $\mathbf{x}$ , ideally, we should have only one  $g_j(\mathbf{x})$ ,  $j = 1, \dots, K$  greater than 0.
- ▶ However, it is possible for **multiple** or **no**  $g_i(\mathbf{x})$  to be  $> 0$ . These may be taken as reject cases, but the usual approach is to assign  $\mathbf{x}$  to the class having the highest discriminant.
- ▶ **Decision rule** for any test case  $\mathbf{x}$ :

$$\text{Choose } C_i \text{ if } g_i(\mathbf{x}) = \max_{j=1}^K g_j(\mathbf{x})$$

- ▶ Geometrically a **linear classifier** partitions the feature space into  $K$  **convex decision regions**  $\mathcal{R}_i$ .

## Pairwise Separation

- ▶ If the classes are not linearly separable, one approach is to divide it into a set of linear problems and linear discriminants can be used to separate the classes.
- ▶ One possibility is to perform **pairwise separation** of classes by considering one pair of distinct classes at a time.
- ▶  $K(K - 1)/2$  linear discriminants are used.
- ▶ It is easier for the classes to be **pairwise linearly separable** than **linearly separable**.



## Pairwise Separation (2)

- **Discriminant function** for classes  $i$  and  $j$  ( $i, j = 1, \dots, K$  and  $j \neq i$ ):

$$g_{ij}(\mathbf{x} \mid \mathbf{w}_{ij}, w_{ij0}) = \mathbf{w}_{ij}^T \mathbf{x} + w_{ij0} = \begin{cases} > 0 & \text{if } \mathbf{x} \in C_i \\ \leq 0 & \text{if } \mathbf{x} \in C_j \\ \text{don't care} & \text{if } \mathbf{x} \in C_k, k \neq i, k \neq j \end{cases}$$

- if  $\mathbf{x}^{(\ell)} \in C_k$  where  $k \neq i, k \neq j$ , then  $\mathbf{x}^{(\ell)}$  is not used during training of  $g_{ij}(\mathbf{x})$ .

- **Decision rule** for any test case  $\mathbf{x}$ :

Choose  $C_i$  if  $\forall j \neq i, g_{ij}(\mathbf{x}) > 0$

- Sometimes we may not be able to find such a class  $C_i$ . If we do not want to reject such cases, a **relaxed** decision rule can be defined based on a new set of discriminant functions:

$$g_i(\mathbf{x}) = \sum_{j \neq i} g_{ij}(\mathbf{x})$$

# Outline

Introduction

Geometric View

Parametric Discrimination Revisited

Logistic Discrimination

## Linear Parametric Discrimination Revisited

- Recall that if the class-conditional densities  $p(\mathbf{x} \mid C_i)$  are **Gaussian** sharing a **common covariance matrix**  $\mathbf{\Sigma}$ , the discriminant functions are **linear**:

$$g_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + w_{i0}$$

where

$$\mathbf{w}_i = \mathbf{\Sigma}^{-1} \boldsymbol{\mu}_i$$

$$w_{i0} = -\frac{1}{2} \boldsymbol{\mu}_i^T \mathbf{\Sigma}^{-1} \boldsymbol{\mu}_i + \log P(C_i)$$

- Given a sample  $\mathcal{X}$ , we find the **ML estimates** for  $\boldsymbol{\mu}_i$  and  $\mathbf{\Sigma}$ , denoted by  $\mathbf{m}_i$  and  $\mathbf{S}$ , and plug them into the discriminant functions.

## Two-Class Example

- ▶ Let

$$P(C_1 | \mathbf{x}) = y \quad P(C_2 | \mathbf{x}) = 1 - y$$

- ▶ Classification rule:

$$\text{Choose } \begin{cases} C_1 & \text{if } y > 0.5 \\ C_2 & \text{otherwise} \end{cases}$$

- ▶ Equivalent tests for classification rule:

$$\frac{y}{1-y} > 1 \quad \text{or} \quad \log \frac{y}{1-y} > 0$$

where  $y/(1-y)$  is called the **odds** (odds ratio) of  $y$  and  $\log[y/(1-y)]$  is called the **log odds** of  $y$  or **logit transformation/function** of  $y$ , written as  $\text{logit}(y)$ .

- ▶ The  $\text{logit}(\cdot)$  is a type of function that maps probability values from  $(0, 1)$  to real numbers in  $(-\infty, +\infty)$ .

## Logit Function

- In the case of two normal classes sharing a common covariance matrix, the **logit function**:

$$\begin{aligned}\text{logit}(P(C_1 | \mathbf{x})) &= \log \frac{P(C_1 | \mathbf{x})}{1 - P(C_1 | \mathbf{x})} = \log \frac{P(C_1 | \mathbf{x})}{P(C_2 | \mathbf{x})} \\ &= \log \frac{p(\mathbf{x} | C_1)}{p(\mathbf{x} | C_2)} + \log \frac{p(C_1)}{p(C_2)} \\ &= \log \frac{(2\pi)^{-\frac{d}{2}} |\mathbf{\Sigma}|^{-\frac{1}{2}} \exp[-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T \mathbf{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_1)]}{(2\pi)^{-\frac{d}{2}} |\mathbf{\Sigma}|^{-\frac{1}{2}} \exp[-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_2)^T \mathbf{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_2)]} + \log \frac{p(C_1)}{p(C_2)} \\ &= \mathbf{w}^T \mathbf{x} + w_0\end{aligned}$$

where

$$\mathbf{w} = \mathbf{\Sigma}^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2), \quad w_0 = -\frac{1}{2}(\boldsymbol{\mu}_1 + \boldsymbol{\mu}_2)^T \mathbf{\Sigma}^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) + \log \frac{p(C_1)}{p(C_2)}$$

## Sigmoid Function

- Sigmoid function or logistic function (inverse function of **logit**):


$$P(C_1 | \mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

which directly computes the posterior class probability  $P(C_1 | \mathbf{x})$ .

- **Training:**
  - Estimate  $\mu_1$ ,  $\mu_2$ , and  $\Sigma$  from data and plug the estimates into the discriminant functions.
- **Testing:**
  - Calculate  $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$  and choose  $C_1$  if  $g(\mathbf{x}) > 0$ , or
  - Calculate  $y = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0)$  and choose  $C_1$  if  $y > 0.5$  (since  $y$  can be interpreted as a posterior probability and  $\text{sigmoid}(0) = 0.5$ ).



## Sigmoid Function (2)



# Outline

Introduction

Geometric View

Parametric Discrimination Revisited

Logistic Discrimination

## Logistic Discrimination

- ▶ In **logistic discrimination** (or **logistic regression**), we do not model the class-conditional densities  $p(x | C_i)$  but rather their ratio.
- ▶ Unlike the parametric classification approach (the likelihood-based approach studied before) which learns the classifier by estimating the parameters of  $p(x | C_i)$ , logistic discrimination (which is a **discriminant-based approach**) estimates the parameters of the discriminant directly.

## Two Classes

- ▶ Let us start with two classes and assume that the log likelihood ratio is linear:

$$\log \frac{p(\mathbf{x} | C_1)}{p(\mathbf{x} | C_2)} = \mathbf{w}^T \mathbf{x} + w_0^0$$

- ▶ Using Bayes' rule, we have

$$\begin{aligned} \text{logit}(P(C_1 | \mathbf{x})) &= \log \frac{P(C_1 | \mathbf{x})}{1 - P(C_1 | \mathbf{x})} = \log \frac{P(C_1 | \mathbf{x})}{P(C_2 | \mathbf{x})} = \log \frac{p(\mathbf{x} | C_1)}{p(\mathbf{x} | C_2)} + \log \frac{p(C_1)}{p(C_2)} \\ &= \mathbf{w}^T \mathbf{x} + w_0 \end{aligned}$$

where  $w_0 = w_0^0 + \log[p(C_1)/P(C_2)]$

- ▶ Then we have

$$y = P(C_1 | \mathbf{x}) = \text{sigmoid}(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{1 + \exp[-(\mathbf{w}^T \mathbf{x} + w_0)]}$$

- equivalent to the case when class-conditional densities are normal
- logistic discrimination is more general, e.g.,  $\mathbf{x}$  may take discrete attributes

## Parameter Learning

- ▶ Training set  $\mathcal{X} = \{(\mathbf{x}^{(\ell)}, r^{(\ell)})\}_{\ell=1}^N$  where

$$r^{(\ell)} = \begin{cases} 1 & \text{if } \mathbf{x}^{(\ell)} \in C_1 \\ 0 & \text{if } \mathbf{x}^{(\ell)} \in C_2 \end{cases}$$

- ▶ Given an input  $\mathbf{x}^{(\ell)}$ , we assume that  $r^{(\ell)}$  is **Bernoulli** with parameter  $y^{(\ell)} = P(C_1 \mid \mathbf{x}^{(\ell)})$ :

$$r^{(\ell)} \mid \mathbf{x}^{(\ell)} \sim \text{Bernoulli}(y^{(\ell)})$$

Here, we see the difference from the likelihood-based methods where we modeled  $p(\mathbf{x} \mid C_i)$ ; in the discriminant-based approach, we model directly  $r^{(\ell)} \mid \mathbf{x}^{(\ell)}$ .

- ▶ **Likelihood**:

$$L(\mathbf{w}, w_0 \mid \mathcal{X}) = \prod_{\ell} (y^{(\ell)})^{r^{(\ell)}} (1 - y^{(\ell)})^{1-r^{(\ell)}}$$

## Parameter Learning (2)

- ▶ Maximizing the likelihood function  $L(\mathbf{w}, w_0 \mid \mathcal{X})$  is equivalent to minimizing an error function (**cross-entropy**)  $E(\mathbf{w}, w_0 \mid \mathcal{X})$  defined as

$$\begin{aligned} E(\mathbf{w}, w_0 \mid \mathcal{X}) &= -\log L(\mathbf{w}, w_0 \mid \mathcal{X}) \\ &= -\sum_{\ell} \left[ r^{(\ell)} \log y^{(\ell)} + (1 - r^{(\ell)}) \log(1 - y^{(\ell)}) \right] \end{aligned}$$

- ▶ No closed-form solution exists. Iterative algorithms such as **gradient descent** or more complicated methods can be used.
- ▶ Please keep in mind once a suitable model and an error function is defined, the **(numerical) optimization** of the model parameters to minimize the error function can be done by using one of many possible techniques.

## Gradient Descent Learning

- ▶ If  $y = \text{sigmoid}(a) = 1/[1 + \exp(-a)]$ , its derivative is

$$\frac{dy}{da} = y(1 - y)$$

- ▶ Update equations for  $w_j$  ( $j = 0, \dots, d$ ):

$$\begin{aligned}\Delta w_j &= -\eta \frac{\partial E}{\partial w_j} = \eta \sum_{\ell} \left( \frac{r^{(\ell)}}{y^{(\ell)}} \frac{\partial y^{(\ell)}}{\partial w_j} - \frac{1 - r^{(\ell)}}{1 - y^{(\ell)}} \frac{\partial y^{(\ell)}}{\partial w_j} \right) \\ &= \eta \sum_{\ell} \left( \frac{r^{(\ell)}}{y^{(\ell)}} - \frac{1 - r^{(\ell)}}{1 - y^{(\ell)}} \right) \frac{\partial y^{(\ell)}}{\partial a^{(\ell)}} \frac{\partial a^{(\ell)}}{\partial w_j}\end{aligned}$$

## Gradient Descent Learning (2)

- Update equations for  $w_j$  ( $j = 1, \dots, d$ ) and  $w_0$ :

Since  $a^{(\ell)} = \mathbf{w}^T \mathbf{x}^{(\ell)} + w_0$ , we have

$$\frac{\partial a^{(\ell)}}{\partial w_j} = x_j^{(\ell)}, \quad j = 1, \dots, d$$

So

$$\begin{aligned} \Delta w_j &= -\eta \frac{\partial E}{\partial w_j} = \eta \sum_{\ell} \left( \frac{r^{(\ell)}}{y^{(\ell)}} - \frac{1 - r^{(\ell)}}{1 - y^{(\ell)}} \right) y^{(\ell)} (1 - y^{(\ell)}) x_j^{(\ell)} \\ &= \eta \sum_{\ell} (r^{(\ell)} - y^{(\ell)}) x_j^{(\ell)}, \quad j = 1, \dots, d \end{aligned}$$

$$\Delta w_0 = -\eta \frac{\partial E}{\partial w_0} = \eta \sum_{\ell} (r^{(\ell)} - y^{(\ell)})$$




## Gradient Descent Algorithm

```
For  $j = 0, \dots, d$   
     $w_j \leftarrow \text{rand}(-0.01, 0.01)$   
Repeat  
    For  $j = 0, \dots, d$   
         $\Delta w_j \leftarrow 0$   
    For  $t = 1, \dots, N$   
         $o \leftarrow 0$   
        For  $j = 0, \dots, d$   
             $o \leftarrow o + w_j x_j^t$   
         $y \leftarrow \text{sigmoid}(o)$   
         $\Delta w_j \leftarrow \Delta w_j + (r^t - y)x_j^t$   
    For  $j = 0, \dots, d$   
         $w_j \leftarrow w_j + \eta \Delta w_j$   
Until convergence
```

For  $w_0$ , we assume that there is an extra input  $x_0$ , which is always  $+1$ :  $x_0^t = +1, \forall t$ .

## A Univariate Two-Class Example



*Both  $wx + w_0$  and  $\text{sigmoid}(wx + w_0)$  are shown as the learning develops.*

- ▶ We see that to get outputs of 0 and 1, the sigmoid hardens, which is achieved by increasing the magnitude of  $w$ , or  $\|\mathbf{w}\|$  in the multivariate case.
- ▶ After training, during testing, given  $\mathbf{x}^{(\ell)}$ , we calculate  $y^{(\ell)} = \text{sigmoid}(\mathbf{w}^T \mathbf{x}^{(\ell)} + w_0)$  and choose  $C_1$  if  $y^{(\ell)} > 0.5$  and choose  $C_2$  otherwise.

## Remarks on Parameter Learning

- ▶ To minimize # of misclassifications, we do not need to continue learning until all  $y^{(\ell)}$  are 0 or 1, but only until  $y^{(\ell)}$  are less than or greater than 0.5 (i.e., on the correct side of the decision boundary).
- ▶ If we do continue training beyond this point, cross-entropy will continue decreasing ( $|w_j|$  will continue increasing to harden the sigmoid), but the number of misclassifications will not decrease.
- ▶ We continue training until the number of misclassifications does not decrease (which will be 0 if the classes are linearly separable).
- ▶ Actually stopping early before we have 0 training error is a form of regularization. Because we start with weights almost 0 and they move away as training continues, stopping early corresponds to a model with more weights close to 0 and effectively fewer parameters.

## Multiple Classes

- ▶ One of the  $K > 2$  classes, e.g.,  $C_K$ , is taken as the **reference class**.
- ▶ Assume that

$$\log \frac{p(\mathbf{x} \mid C_i)}{p(\mathbf{x} \mid C_K)} = \mathbf{w}_i^T \mathbf{x} + w_{i0}^0, \quad i = 1, \dots, K - 1$$

So we have

$$\begin{aligned} \frac{P(C_i \mid \mathbf{x})}{P(C_K \mid \mathbf{x})} &= \frac{p(\mathbf{x} \mid C_i)p(C_i)}{p(\mathbf{x} \mid C_K)p(C_K)} \\ &= \exp(\mathbf{w}_i^T \mathbf{x} + w_{i0}^0) \cdot \exp\left(\log \frac{p(C_i)}{p(C_K)}\right) \\ &= \exp(\mathbf{w}_i^T \mathbf{x} + w_{i0}) \end{aligned} \tag{1}$$

where  $w_{i0} = w_{i0}^0 + \log[p(C_i)/P(C_K)]$

## Generalization of Sigmoid Function

- Summing (1) over  $i = 1, \dots, K - 1$ :

$$\sum_{i=1}^{K-1} \frac{P(C_i | \mathbf{x})}{P(C_K | \mathbf{x})} = \frac{1 - P(C_K | \mathbf{x})}{P(C_K | \mathbf{x})} = \sum_{i=1}^{K-1} \exp(\mathbf{w}_i^T \mathbf{x} + w_{i0})$$

we get

$$P(C_K | \mathbf{x}) = \frac{1}{1 + \sum_{i=1}^{K-1} \exp(\mathbf{w}_i^T \mathbf{x} + w_{i0})} \quad (2)$$

- From (1) and (2), we get

$$P(C_i | \mathbf{x}) = \frac{\exp(\mathbf{w}_i^T \mathbf{x} + w_{i0})}{1 + \sum_{j=1}^{K-1} \exp(\mathbf{w}_j^T \mathbf{x} + w_{j0})}, \quad i = 1, \dots, K - 1$$

## Softmax Function

- ▶ If we want to treat all classes uniformly without having to choose a reference class, we can use the **softmax function** instead for the posterior class probabilities:

$$y_i = \hat{P}(C_i | \mathbf{x}) = \frac{\exp(\mathbf{w}_i^T \mathbf{x} + w_{i0})}{\sum_{j=1}^K \exp(\mathbf{w}_j^T \mathbf{x} + w_{j0})}, \quad i = 1, \dots, K$$

- ▶ If  $\mathbf{w}_i^T \mathbf{x} + w_{i0}$  for one class is sufficiently larger than for the others, its corresponding  $y_i$  will be close to 1 and the others will be close to 0.
- ▶ The softmax function behaves like taking a **maximum**, but it has the advantage of being **differentiable**.

## Parameter Learning

- ▶ Each sample point is a **multinomial trial** with one draw, i.e.

$$\mathbf{r}^{(\ell)} \mid \mathbf{x}^{(\ell)} \sim \text{Mult}_K(1, \mathbf{y}^{(\ell)})$$

where  $y_i^{(\ell)} \equiv P(C_i \mid \mathbf{x}^{(\ell)})$

- ▶ **Likelihood**:

$$L(\{\mathbf{w}_i, w_{i0}\}_i \mid \mathcal{X}) = \prod_{\ell} \prod_i (y_i^{(\ell)})^{r_i^{(\ell)}}$$

- ▶ **Cross-entropy** error function:

$$E(\{\mathbf{w}_i, w_{i0}\}_i \mid \mathcal{X}) = - \sum_{\ell} \sum_i r_i^{(\ell)} \log y_i^{(\ell)}$$

## Gradient Descent Learning

- ▶ If  $y_i = \exp(a_i) / \sum_j \exp(a_j)$ , its derivative is

$$\frac{\partial y_i}{\partial a_j} = y_i(\delta_{ij} - y_j)$$

where  $\delta_{ij}$  is the Kronecker delta, which is 1 if  $i = j$  and 0 if  $i \neq j$ .

- ▶ **Update equations** given  $\sum_i r_i^{(\ell)} = 1$ :

$$\begin{aligned}\Delta \mathbf{w}_j &= \eta \sum_{\ell} \sum_i r_i^{(\ell)} (\delta_{ij} - y_j^{(\ell)}) \mathbf{x}^{(\ell)} = \eta \sum_{\ell} \left[ \sum_i r_i^{(\ell)} \delta_{ij} - y_j^{(\ell)} \sum_i r_i^{(\ell)} \right] \mathbf{x}^{(\ell)} \\ &= \eta \sum_{\ell} (r_j^{(\ell)} - y_j^{(\ell)}) \mathbf{x}^{(\ell)}, \quad j = 1, \dots, K \\ \Delta w_{j0} &= \eta \sum_{\ell} (r^{(\ell)} - y^{(\ell)})\end{aligned}$$

Note that because of the normalization in softmax,  $\mathbf{w}_j$  and  $w_{j0}$  are affected not only by  $\mathbf{x}^{(\ell)} \in C_j$  but also by  $\mathbf{x}^{(\ell)} \in C_i, i \neq j$ .






# Gradient Descent Algorithm

```
For  $i = 1, \dots, K$ , For  $j = 0, \dots, d$ ,  $w_{ij} \leftarrow \text{rand}(-0.01, 0.01)$ 
Repeat
  For  $i = 1, \dots, K$ , For  $j = 0, \dots, d$ ,  $\Delta w_{ij} \leftarrow 0$ 
  For  $t = 1, \dots, N$ 
    For  $i = 1, \dots, K$ 
       $o_i \leftarrow 0$ 
      For  $j = 0, \dots, d$ 
         $o_i \leftarrow o_i + w_{ij} x_j^t$ 
      For  $i = 1, \dots, K$ 
         $y_i \leftarrow \exp(o_i) / \sum_k \exp(o_k)$ 
      For  $i = 1, \dots, K$ 
        For  $j = 0, \dots, d$ 
           $\Delta w_{ij} \leftarrow \Delta w_{ij} + (r_i^t - y_i) x_j^t$ 
    For  $i = 1, \dots, K$ 
      For  $j = 0, \dots, d$ 
         $w_{ij} \leftarrow w_{ij} + \eta \Delta w_{ij}$ 
  Until convergence
```

We take  $x_0^t = 1, \forall t$ .

## A Two-dimensional Three-class Example



## Remarks on Parameter Learning

- ▶ We do not need to continue training to minimize cross-entropy as much as possible; we train only until the correct class has the highest weighted sum, and therefore we can stop training earlier by checking the number of misclassifications.

## Logistic Discriminant

- ▶ When data are normally distributed, the logistic discriminant has a comparable performance to the parametric, normal-based linear discriminant.
- ▶ Logistic discrimination can still be used when the class-conditional densities are nonnormal or when they are not unimodal, as long as classes are linearly separable.

## Generalizing the Linear Model

- ▶ The ratio of class-conditional densities is of course not restricted to be linear.
- ▶ Quadratic discriminant:

$$\log \frac{p(\mathbf{x} | C_i)}{p(\mathbf{x} | C_K)} = \mathbf{x}^T \mathbf{W}_i \mathbf{x} + \mathbf{w}_i^T \mathbf{x} + w_{i0}$$

which corresponds to parametric discrimination with multivariate normal class-conditionals having **different covariance matrices**.

- ▶ Sum of nonlinear basis functions:

$$\log \frac{p(\mathbf{x} | C_i)}{p(\mathbf{x} | C_K)} = \mathbf{w}_i^T \Phi(\mathbf{x}) + w_{i0}$$

where  $\Phi(\cdot)$  are basis functions which transform the original input variables to a new set of variables.

- ▶ Basis functions are related to:
  - Hidden units like sigmoid function in **neural networks** (studied later)
  - Kernels in kernel methods such as **support vector machines (SVM)** (studied later).