

Decision Trees

Prof. Ziping Zhao

School of Information Science and Technology
ShanghaiTech University, Shanghai, China

CS182: Introduction to Machine Learning (Fall 2021)
<http://cs182.sist.shanghaitech.edu.cn>

Outline

Introduction

Univariate Trees

Tree Pruning

Rule Extraction from Trees

Multivariate Trees

Outline

Introduction

Univariate Trees

Tree Pruning

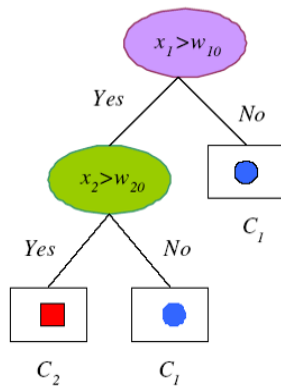
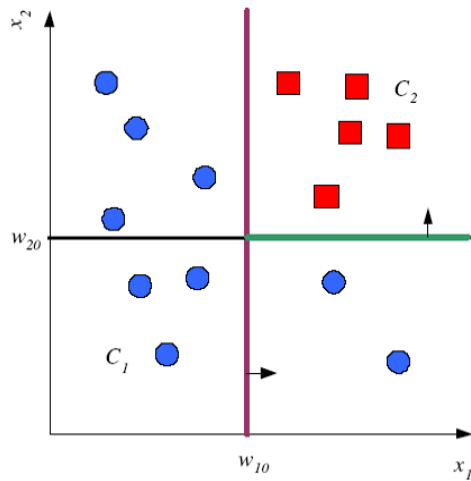
Rule Extraction from Trees

Multivariate Trees

Introduction

- ▶ In nonparametric estimation, the input space is divided into local regions each of which corresponds to a **local model** computed from the training data in that region.
- ▶ A **decision tree** is a **nonparametric hierarchical model** for supervised learning whereby the local region is identified through a sequence of **recursive splits** in a small number of steps – **divide-and-conquer** approach.
- ▶ Decision trees were first made popular in **statistics** and later in machine learning.
- ▶ Two types of nodes in a decision tree:
 - **(Internal) decision node**: a test function with discrete outcomes labeling the branches.
 - **(Terminal) leaf node**: the value associated with it constitutes the output (class label for classification; numeric value for regression).

Data Set and Corresponding Decision Tree



Discriminants

- ▶ The test function $f_m(\mathbf{x})$ at each decision node m defines a **discriminant** in the input space dividing it into smaller regions which are further subdivided as we take a path from the root down.
- ▶ A leaf node defines a **localized region** in the input space where instances falling in the region have the same output.
- ▶ The **region boundaries** are defined by the discriminants that are coded in the internal nodes along the path from the root to the leaf node.
- ▶ Advantages of decision trees:
 - **Fast localization** of the region covering the input as a result of the hierarchical placement of decisions.
 - **High interpretability**: can be converted easily into a set of **IF-THEN** rules that are easily understandable.

Outline

Introduction

Univariate Trees

Tree Pruning

Rule Extraction from Trees

Multivariate Trees

Univariate Trees I

- ▶ In a **univariate tree**, the test in each internal decision node uses only **one** of the input dimensions.
- ▶ **n -way split**: a discrete-valued input dimension x_j with n possible values leads to n branches from the decision node (**binary split** is a special case with $n = 2$).
- ▶ A numeric input should be **discretized** into $n \geq 2$ values using suitably chosen threshold(s). Usually we choose $n = 2$, and the discriminant is

$$f_m(\mathbf{x}) : x_j > w_{m0}$$

- ▶ Successive splits are **orthogonal** to each other. The leaf nodes define **hyperrectangles** in the input space.

Univariate Trees II

- ▶ **Tree induction** or tree learning is the construction of a tree given a training sample.
- ▶ For a given training set, there exists many trees that code it with no error.
- ▶ We are interested in finding the smallest one, where tree size is measured as the number of nodes in the tree and the complexity of the decision nodes.
- ▶ Given a training sample, finding the **smallest** tree to code the data is **NP-complete**, making it necessary to use greedy **local search** algorithms for tree learning.
 - starting at the root with the complete training data, we look for the **best split** at each step.
 - we continue splitting recursively with the corresponding subset until we do not need to split anymore, at which point a leaf node is created.

Classification Trees

- ▶ In the case of a decision tree for classification, namely, a **classification tree**.
- ▶ For node m :
 - N_m training instances
 - N_m^i of N_m instances belong to class C_i , i.e., $\sum_i N_m^i = N_m$.
- ▶ Given that an instance reaches node m , an estimate for the probability of class C_i is

$$\hat{P}(C_i | \mathbf{x}, m) \equiv p_m^i = \frac{N_m^i}{N_m}$$

- ▶ Node m is considered **pure** if there is a class C_i with $p_m^i = 1$ (and hence all other p_m^j with $j \neq i$ are 0). No further splitting is needed and the node becomes a **leaf node** with class label C_i .
- ▶ The goodness of a split in classification trees is quantified by an **impurity measure**.
- ▶ Some impurity measures:
 - Entropy
 - Gini index
 - Misclassification error

Entropy

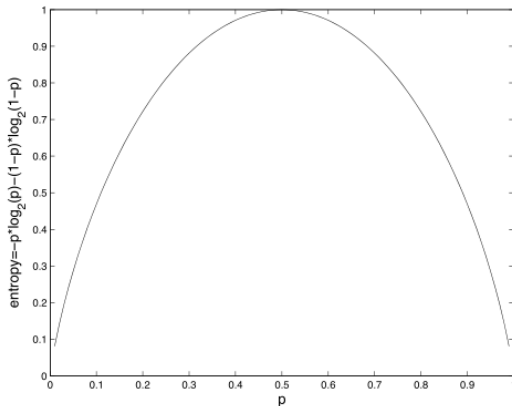
- Entropy at node m :

$$\mathcal{I}_m = - \sum_{i=1}^K p_m^i \log_2 p_m^i$$

where we assume $0 \log 0 = 0$.

- The largest entropy is $\log_2 K$ when all $p_m^i = 1/K$.

- Entropy function for a two-class problem:



Other Impurity Measures

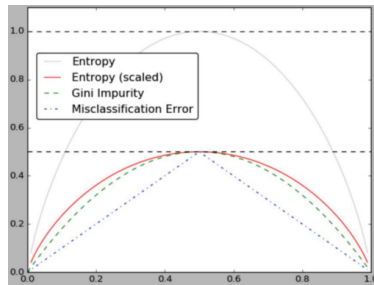
- ▶ For a two-class problem where $p^1 = p$ and $p^2 = 1 - p$, a nonnegative function $\phi(p, 1 - p)$ measures the impurity of a split if it satisfies the following properties:
 - $\phi(1/2, 1/2) \geq \phi(p, 1 - p), \forall p \in [0, 1]$
 - $\phi(0, 1) = \phi(1, 0) = 0$
 - $\phi(p, 1 - p)$ is increasing in p on $[0, 1/2]$ and decreasing in p on $[1/2, 1]$ (very often we need it to be symmetric).
- ▶ Examples other than entropy:
 - Gini index:

$$\phi(p, 1 - p) = 2p(1 - p)$$

which is used in economics as a measure of unequal distribution of wealth.

- Misclassification error:

$$\phi(p, 1 - p) = 1 - \max(p, 1 - p)$$



Best Split

- ▶ At node m , let N_{mj} of the N_m instances take branch j , i.e., $f_m(\mathbf{x}) = j$. Also, N_{mj}^i of the N_{mj} instances belong to class C_i . So,

$$\sum_{j=1}^n N_{mj} = N_m \quad \sum_{i=1}^K N_{mj}^i = N_{mj} \quad \sum_{j=1}^n N_{mj}^i = N_m^i.$$

- ▶ If $f_m(\mathbf{x}) = j$, the estimate for the probability of class C_i is

$$\hat{P}(C_i | \mathbf{x}, m, j) \equiv p_{mj}^i = \frac{N_{mj}^i}{N_{mj}}$$

and the total impurity after the split is

$$\mathcal{I}'_m = - \sum_{j=1}^n \frac{N_{mj}}{N_m} \sum_{i=1}^K p_{mj}^i \log p_{mj}^i$$

- ▶ All attributes are tried to implement the split and the one that gives the **minimum entropy** is chosen for the test function.
- ▶ Tree construction continues recursively and in parallel for all the branches that are not pure, until all are pure.

Classification Tree Construction Algorithm

```
GenerateTree( $\mathcal{X}$ )
  If NodeEntropy( $\mathcal{X}$ ) <  $\theta_I$ 
    Create leaf labelled by majority class in  $\mathcal{X}$ 
    Return
   $i \leftarrow \text{SplitAttribute}(\mathcal{X})$ 
  For each branch of  $\mathbf{x}_i$ 
    Find  $\mathcal{X}_i$  falling in branch
    GenerateTree( $\mathcal{X}_i$ )
SplitAttribute( $\mathcal{X}$ )
  MinEnt  $\leftarrow$  MAX
  For all attributes  $i = 1, \dots, d$ 
    If  $\mathbf{x}_i$  is discrete with  $n$  values
      Split  $\mathcal{X}$  into  $\mathcal{X}_1, \dots, \mathcal{X}_n$  by  $\mathbf{x}_i$ 
       $e \leftarrow \text{SplitEntropy}(\mathcal{X}_1, \dots, \mathcal{X}_n)$ 
      If  $e < \text{MinEnt}$  MinEnt  $\leftarrow$   $e$ ; bestf  $\leftarrow$   $i$ 
    Else /*  $\mathbf{x}_i$  is numeric */
      For all possible splits
        Split  $\mathcal{X}$  into  $\mathcal{X}_1, \mathcal{X}_2$  on  $\mathbf{x}_i$ 
         $e \leftarrow \text{SplitEntropy}(\mathcal{X}_1, \mathcal{X}_2)$ 
        If  $e < \text{MinEnt}$  MinEnt  $\leftarrow$   $e$ ; bestf  $\leftarrow$   $i$ 
  Return bestf
```

Regression Trees

- ▶ A **regression tree** is constructed in a similar manner as a classification tree, except that the impurity measure for classification is replaced by a measure appropriate for regression.
- ▶ For node m , let $\mathcal{X}_m \subset \mathcal{X}$ denote the set of instances reaching m . We define the following indicator function:

$$b_m(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_m \\ 0 & \text{otherwise} \end{cases}$$

- ▶ Estimated value at node m :

$$g_m = \frac{\sum_{\ell} b_m(\mathbf{x}^{(\ell)}) r^{(\ell)}}{\sum_{\ell} b_m(\mathbf{x}^{(\ell)})}$$

- ▶ **Mean squared error** of estimated value measures goodness of split:

$$E_m = \frac{1}{N_m} \sum_{\ell} (r^{(\ell)} - g_m)^2 b_m(\mathbf{x}^{(\ell)})$$

where $N_m = |\mathcal{X}_m| = \sum_{\ell} b_m(\mathbf{x}^{(\ell)})$.

Tree Expansion

- ▶ If $E_m < \theta_r$ for some threshold θ_r , the error is acceptable and hence node m is designated as a **leaf node** with value g_m stored.
- ▶ If E_m is too large, we look for a **split threshold** w_{m0} for **further splitting** so that the sum of the errors in the branches is minimum, and then we continue recursively.
- ▶ At node m , let \mathcal{X}_{mj} be the subset of \mathcal{X}_m taking branch j . We define

$$b_{mj}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_{mj} \\ 0 & \text{otherwise} \end{cases}$$

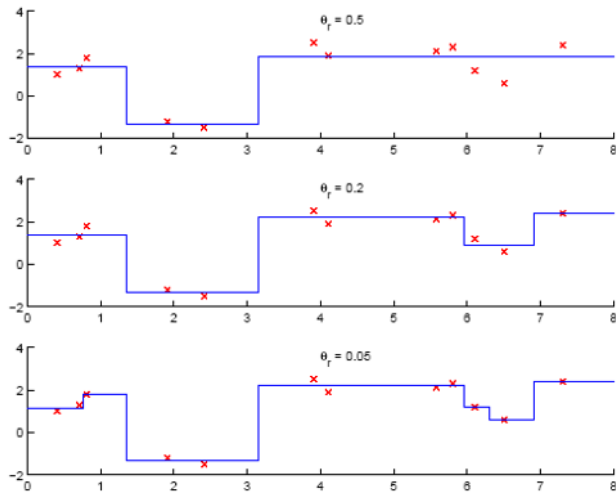
- ▶ Estimated value in branch j of node m :

$$g_{mj} = \frac{\sum_{\ell} b_{mj}(\mathbf{x}^{(\ell)}) r^{(\ell)}}{\sum_{\ell} b_{mj}(\mathbf{x}^{(\ell)})}$$

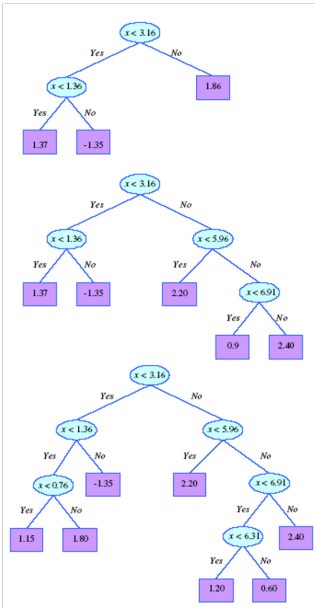
- ▶ Error after split:

$$E'_m = \frac{1}{N_m} \sum_j \sum_{\ell} (r^{(\ell)} - g_{mj})^2 b_{mj}(\mathbf{x}^{(\ell)})$$

Regression Trees for Different Values of θ_r



Univariate Trees



Best Split

- ▶ As in classification, we look for the split that results in the smallest error E'_m and then split the node to expand the tree.
- ▶ Besides the mean squared error, other error functions may also be used. E.g., **worst possible error:**

$$E_m = \max_j \max_{\ell} \left\{ |y^{(\ell)} - g_{mj}| b_m(\mathbf{x}^{(\ell)}) \right\}$$

which can guarantee that the error for any instance is never larger than a given threshold.

- ▶ Similar to going from running mean to running line in nonparametric regression, instead of taking an average at a leaf that implements a constant fit, we can also do a linear regression fit over the instances choosing the leaf:

$$g_m(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_{m0}$$

which makes the estimate dependent on \mathbf{x} and generates smaller trees, but introduces extra computation.

Outline

Introduction

Univariate Trees

Tree Pruning

Rule Extraction from Trees

Multivariate Trees

Pruning

- ▶ If very few training instances reach a node, decision based on the instances may give high generalization error.
- ▶ Solution 1 – **prepruning**:
 - Stop node split when the number of instances reaching a node is below a certain percentage of the training set regardless of the impurity or error.
- ▶ Solution 2 – **postpruning**:
 - Grow the tree full until all leaves are pure.
 - Find subtrees and do the following for each of them:
 - ▶ Replace the subtree by a leaf node set with an appropriate label (for either classification or regression) based on the training instances covered by the subtree.
 - ▶ If the leaf node does not perform worse than the subtree on the **pruning set** (a separate labeled data set), the subtree is pruned and replaced by the leaf node.
- ▶ Prepruning is **faster** but postpruning generally leads to **more accurate** trees.

Outline

Introduction

Univariate Trees

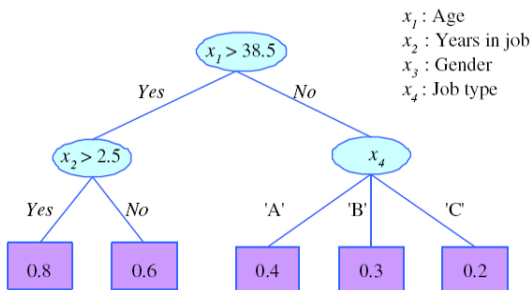
Tree Pruning

Rule Extraction from Trees

Multivariate Trees

Feature Extraction

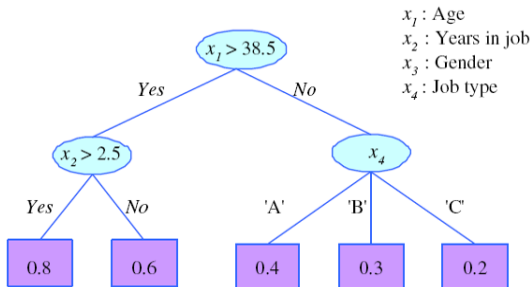
- ▶ The univariate tree only uses the variables that are necessary, so it is possible to use the tree for **feature extraction**.
- ▶ E.g., only features x_1, x_2, x_4 are extracted:



- ▶ In general, features closer to the root are more important **globally**.

Rule Extraction

- ▶ A main advantage of decision trees is **interpretability**.
- ▶ A set of **IF-THEN** rules (R), i.e., a rule base:



- R1: IF (age > 38.5) AND (years-in-job > 2.5) THEN $y = 0.8$
R2: IF (age > 38.5) AND (years-in-job ≤ 2.5) THEN $y = 0.6$
R3: IF (age ≤ 38.5) AND (job-type = 'A') THEN $y = 0.4$
R4: IF (age ≤ 38.5) AND (job-type = 'B') THEN $y = 0.3$
R5: IF (age ≤ 38.5) AND (job-type = 'C') THEN $y = 0.2$

More on Rules

- ▶ For classification, more than one leaf node may be labeled with the same class. So the condition part of the corresponding rule can be expressed as a **disjunction** (OR) of **conjunctions** (AND), e.g.

IF $(x_1 \leq w_{10})$ OR $((x_1 > w_{10}) \text{ AND } (x_2 \leq w_{20}))$ THEN C_1

- ▶ Pruning rules (i.e., pruning a term from one rule without touching other rules) is possible for simplification. But after the rules are pruned, it may not be possible to write them back as a tree anymore.
- ▶ Instead of extracting rules from a decision tree learned from data, it is also possible to learn the rules directly from data.
- ▶ **Rule induction** is similar to tree induction, but:
 - Tree induction is **breadth-first**.
 - Rule induction is **depth-first**; one rule at a time.

Outline

Introduction

Univariate Trees

Tree Pruning

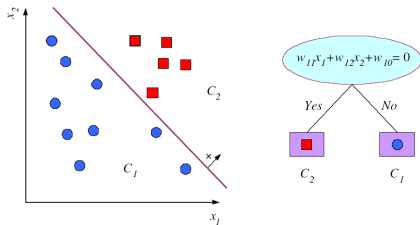
Rule Extraction from Trees

Multivariate Trees

Multivariate Trees

- ▶ At each decision node of a **multivariate tree**, all input attributes can be used to define a test function for the split.
- ▶ **Linear multivariate node** for numeric attributes:

$$f_m(\mathbf{x}) : \mathbf{w}_m^T \mathbf{x} + w_{m0} > 0$$



- ▶ A node $f_m(\mathbf{x})$ defines a hyperplane with arbitrary orientation and leaf nodes define polyhedra in the input space.
- ▶ A univariate tree can be seen as a special case.

Nonlinear Multivariate Nodes

- ▶ Nonlinear multivariate nodes provide even more flexibility, e.g.:
 - Quadratic multivariate node:

$$f_m(\mathbf{x}) : \mathbf{x}^T \mathbf{W}_m \mathbf{x} + \mathbf{w}_m^T \mathbf{x} + w_{m0} > 0$$

- Multilayer perceptron
- Sphere node:

$$f_m(\mathbf{x}) : \|\mathbf{x} - \mathbf{c}_m\| \leq \alpha_m$$

where \mathbf{c}_m is the center and α_m is the radius.

- ▶ Omnivariate decision tree: a hybrid tree architecture where the tree may have univariate, linear multivariate, or nonlinear multivariate nodes.
- ▶ While providing additional flexibility is good, multivariate decision trees also increase the computational requirement significantly. Univariate trees are still the more popular choice in practical applications.