

Ensemble Learning

Ziping Zhao

School of Information Science and Technology
ShanghaiTech University, Shanghai, China

CS182: Introduction to Machine Learning (Fall 2021)
<http://cs182.sist.shanghaitech.edu.cn>

Outline

Introduction

Voting

Boosting

Outline

Introduction

Voting

Boosting

Introduction

Rationale

- ▶ **No Free Lunch Theorem:** There is no single learning algorithm that in any domain always induces the most accurate learner.
- ▶ The usual approach is to try many and choose the one that performs the best on a separate validation set.
- ▶ Each learning algorithm dictates a certain model with a set of assumptions, i.e., the **inductive bias**, leading to error if the assumptions do not hold for the data.
- ▶ **Ensemble learning:**
 - We construct a group of **base-learners** which, when combined, has higher accuracy than the individual learners.
 - The base-learners are usually not chosen for their accuracy, but for their **simplicity**.
 - The base-learners should be **diverse**, that is, accurate on different instances, specializing in subdomains of the problem, so that they can **complement** each other.

Diverse Base-Learners

- ▶ **Different learning algorithms:** different algorithms make different assumptions about the data and lead to different classifiers.
- ▶ **Different hyperparameters of the same learning algorithm:** e.g., number of hidden units in a multilayer perceptron, k in k -nearest neighbor classifier, error threshold in a decision tree, initial state of an iterative procedure, etc.
- ▶ **Different representations of the same input object or event:** multiple sources of information are combined, e.g., both acoustic input and video sequence of lip movements for speech recognition.
- ▶ **Different training sets:** multiple base-learners are trained either in parallel or serially using different training sets.
- ▶ **Different subtasks:** the main task is defined in terms of a number of subtasks solved by different base-learners.

Combining Base-Learners

- ▶ There are different ways the multiple base-learners are combined to generate the final output.
- ▶ **Multiexpert combination methods:**
 - The base-learners work in parallel.
 - Given an instance, they all give their decisions which are then combined to give the final decision.
 - E.g., **voting**, **mixture of experts**, **stacked generalization**.
- ▶ **Multistage combination methods:**
 - The base-learners work in serial.
 - The base-learners are sorted in increasing complexity: a complex base learner is not used unless the preceding simpler base-learners are not confident.
 - E.g., **boosting**, **cascading**.
- ▶ Given L base-learners, we denote by $d_j(x)$ the prediction of the j th base-learner given the input x . The final prediction is calculated from

$$y = f(d_1, d_2, \dots, d_L \mid \Phi)$$

where $f(\cdot)$ is the combining/fusion function with Φ denoting its parameters.

Outline

Introduction

Voting

Boosting

Voting


Voting

- ▶ Voting takes a **convex combination** of the base-learners:

$$y = f(d_1, \dots, d_L \mid \Phi) = \sum_{j=1}^L w_j d_j$$

where y is the final prediction and $\Phi = (w_1, \dots, w_L)^T$ with w_j the weight satisfying

$$w_j \geq 0 \text{ and } \sum_{j=1}^L w_j = 1.$$



- ▶ When there are K outputs, for each learner there are $d_{ji}(x)$, $j = 1, \dots, L$, $i = 1, \dots, K$, and, combining them, we also generate K values, y_i , $i = 1, \dots, K$.
- ▶ For example in classification, we choose the class with the maximum y_i value:

$$\text{Choose } C_i \text{ if } y_i = \max_k y_k$$

Combination Rules

Rule	Fusion function $f(\cdot)$
Weighted sum	$y_i = \sum_j w_j d_{ji}, w_j \geq 0, \sum_j w_j = 1$
Median	$y_i = \text{median}_j d_{ji}$
Minimum	$y_i = \min_j d_{ji}$
Maximum	$y_i = \max_j d_{ji}$
Product	$y_i = \prod_j d_{ji}$

- ▶ Sum rule is the most intuitive and is the most widely used in practice.
- ▶ Median rule is more robust to outliers.
- ▶ Minimum and maximum rules are pessimistic and optimistic, respectively.
- ▶ With the product rule, each learner has veto power; regardless of the other ones, if one learner has an output of 0, the overall output goes to 0.

Mixture of Experts as A Voting


- ▶ In voting, the weights w_j are constant over the input space x .

- ▶ The **mixture of experts** architecture may be seen as a voting method based on the weighted sum in which the weights depend on the input and are in general different for different inputs.

- ▶ The weights of the experts are determined by a **gating network**:

$$y = \sum_{j=1}^L w_j(x) d_j$$

- ▶ The competitive learning algorithm used by the mixture of experts localizes the base-learners such that each of them becomes an expert in a different part of the input space and have its weight, $w_j(x)$, close to 1 in its region of expertise.



Voting for Classification

- ▶ For class C_i :

$$y_i = \sum_{j=1}^L w_j d_{ji}$$

where d_{ji} is the vote of learner j for C_i and w_j is the weight of its vote.

- ▶ Simple voting (a.k.a. plurality voting, or majority voting for 2-class problem):

$$w_j = \frac{1}{L}$$

- ▶ If the voters can also supply the additional information of how much they vote for each class (e.g., by the posterior probability).
- ▶ Another possible way to find w_j is to assess the accuracies of the learners (regressor or classifier) on a separate validation set and use that information to compute the weights, so that we give more weights to more accurate learners.

Voting Under A Bayesian Framework

- ▶ Voting schemes can be seen as approximations under a Bayesian framework.
 - weights w_j approximating prior model probabilities $P(\mathcal{M}_j)$, and
 - model decisions $d_{ji}(x)$ approximating model-conditional likelihoods $P(C_i | x, \mathcal{M}_j)$.
- ▶ For example, in classification **Bayesian model combination**:

$$P(C_i | x) = \sum_{\text{all models } \mathcal{M}_j} P(\mathcal{M}_j) P(C_i | x, \mathcal{M}_j)$$

- ▶ Simple voting corresponds to a uniform prior.
- ▶ If we have a prior distribution preferring simpler models, this would give larger weights to them.
- ▶ We cannot integrate over all models; we only choose a subset for which we believe $P(\mathcal{M}_j)$ is high, or we can have another Bayesian step and calculate $P(\mathcal{M}_j | \mathcal{X})$, the probability of a model given the sample, and sample high probable models from this density.

Analysis

- ▶ Let there be L independent two-class classifiers, where $E[d_j]$ and $\text{Var}(d_j)$ are the expected value and variance of d_j for classifier j .
- ▶ Expected value and variance of output for independent classifiers:

$$E[y] = E\left[\sum_j \frac{1}{L} d_j\right] \geq \frac{1}{L} L \min_j \{E[d_j]\} = \min_j \{E[d_j]\}$$

$$\text{Var}(y) = \text{Var}\left(\sum_j \frac{1}{L} d_j\right) = \frac{1}{L^2} \text{Var}\left(\sum_j d_j\right) \leq \frac{1}{L^2} L \max_j \{\text{Var}(d_j)\} = \frac{1}{L} \max_j \{\text{Var}(d_j)\}$$

As L increases, the expected value (and hence the **bias**) does not change but the variance decreases, and hence the **mean squared error** of the estimator y decreases, leading to an increase in accuracy.

- ▶ General case (non-independent classifiers):

$$\text{Var}(y) = \frac{1}{L^2} \text{Var}\left(\sum_j d_j\right) = \frac{1}{L^2} \left[\sum_j \text{Var}(d_j) + 2 \sum_j \sum_{i < j} \text{Cov}(d_j, d_i) \right]$$

Bagging

- ▶ **Bagging**, short for **bootstrap aggregating**, is a voting method whereby the base-learners are made different by training on slightly different training sets.
- ▶ The L different training sets are generated by **bootstrap**, which draws N instances randomly from a training set \mathcal{X} of size N **with replacement**.
- ▶ Bagging can be seen as a special case of **model averaging** which helps to reduce variance and hence improve accuracy.
- ▶ Unstable algorithms (e.g., decision trees and multilayer perceptrons) that cause large changes in the generated learner (i.e., **high variance**) with small changes in the training set can particularly benefit from bagging.

Outline

Introduction

Voting

Boosting


Boosting

- ▶ In bagging, generating complementary base-learners is left to chance and to the instability of the learning algorithm.
- ▶ In **boosting**, complementary base-learners are generated by training the next learner on the mistakes of the previous learners.
- ▶ Boosting combines **weak learners** (learners with accuracy just required to be better than random guessing, i.e., $> 1/K$ for K -class classification problems; weak but not too weak) to generate a **strong learner** (learners with arbitrarily small error probability).


AdaBoost

- ▶ **AdaBoost** (a short form for **adaptive boosting**) is an iterative procedure that generates a sequence of base-learners each focusing on the errors of previous ones.
- ▶ The original algorithm is **AdaBoost.M1**, but many variants of AdaBoost have also been proposed.
- ▶ AdaBoost modifies the probabilities of drawing instances for classifier training as a function of the error of the previous base-learner.
- ▶ Initially all N instances have the same probability of being drawn.
- ▶ Moving from one iteration to the next iteration, the probability of a correctly classified instance is **decreased** and that of a misclassified instance is **increased**.
- ▶ The success of AdaBoost is due to its property of increasing the **margin**, making the aim of AdaBoost similar to that of SVM.
- ▶ One statistical view of boosting is an **additive** form of **logistic regression**.

A Simple Example




A Simple Example: Base Learner 1




- ▶ All data points have equal weights.
- ▶ Base-learner D1, which is a simple decision tree with only one single level (a.k.a. **decision stump**), misclassifies three + (plus) data points as - (minus).

A Simple Example: Base Learner 2




- ▶ The three data points misclassified by D1 now have higher weights, making it more likely for base learner D2 to classify them correctly.
- ▶ D2 misclassifies three – (minus) data points as + (plus).

A Simple Example: Base Learner 3




- ▶ The three data points misclassified by D2 now have higher weights, making it more likely for base learner D3 to classify them correctly.
- ▶ D3 generates a horizontal decision boundary.

A Simple Example: Combining Base-Learners



- ▶ The three weak learners (D1, D2, D3) are combined to give a strong learner (D4) which can classify all data points correctly.
- ▶ Although the weak learners are simple linear classifiers, the strong learner is highly nonlinear.


A More Realistic Example




- ▶ The data distribution is radial with uniform angular distribution.
- ▶ Linear classifiers are used as weak learners.

A More Realistic Example: $t = 1$ and $t = 2$

► $t = 1$:




► $t = 2$:




A More Realistic Example: $t = 3$ and $t = 4$

► $t = 3$:




► $t = 4$:




A More Realistic Example: $t = 5$ and $t = 6$

► $t = 5$:




► $t = 6$:




A More Realistic Example: $t = 7$ and $t = 40$

► $t = 7$:



► $t = 40$:



AdaBoost Algorithm

Training:

For all $\{x^t, r^t\}_{t=1}^N \in \mathcal{X}$, initialize $p_1^t = 1/N$

For all base-learners $j = 1, \dots, L$

Randomly draw \mathcal{X}_j from \mathcal{X} with probabilities p_j^t

Train d_j using \mathcal{X}_j

For each (x^t, r^t) , calculate $y_j^t \leftarrow d_j(x^t)$

Calculate error rate: $\epsilon_j \leftarrow \sum_t p_j^t \cdot 1(y_j^t \neq r^t)$

If $\epsilon_j > 1/2$, then $L \leftarrow j - 1$; stop

$\beta_j \leftarrow \epsilon_j / (1 - \epsilon_j)$

For each (x^t, r^t) , decrease probabilities if correct:

If $y_j^t = r^t$, then $p_{j+1}^t \leftarrow \beta_j p_j^t$ Else $p_{j+1}^t \leftarrow p_j^t$

Normalize probabilities:

$Z_j \leftarrow \sum_t p_{j+1}^t$; $p_{j+1}^t \leftarrow p_{j+1}^t / Z_j$



Testing:

Given x , calculate $d_j(x), j = 1, \dots, L$

Calculate class outputs, $i = 1, \dots, K$:

$$y_i = \sum_{j=1}^L \left(\log \frac{1}{\beta_j} \right) d_{ji}(x)$$

Bagging vs. Boosting: Training and Test Error Curves



Methods Comparisons

