

Review of Linear Discrimination

Mengting Chen

October 27, 2022

1 Three Approaches to Decision Problems

- (a) First solve the inference problem of determining the class-conditional densities $p(\mathbf{x}|\mathcal{C}_k)$ for each class \mathcal{C}_k individually. Also separately infer the prior class probabilities $p(\mathcal{C}_k)$. Then use Bayes' theorem in the form

$$p(\mathcal{C}_k|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{p(\mathbf{x})} \quad (1.82)$$

to find the posterior class probabilities $p(\mathcal{C}_k|\mathbf{x})$. As usual, the denominator in Bayes' theorem can be found in terms of the quantities appearing in the numerator, because

$$p(\mathbf{x}) = \sum_k p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k). \quad (1.83)$$

Equivalently, we can model the joint distribution $p(\mathbf{x}, \mathcal{C}_k)$ directly and then normalize to obtain the posterior probabilities. Having found the posterior probabilities, we use decision theory to determine class membership for each new input \mathbf{x} . Approaches that explicitly or implicitly model the distribution of inputs as well as outputs are known as *generative models*, because by sampling from them it is possible to generate synthetic data points in the input space.

- (b) First solve the inference problem of determining the posterior class probabilities $p(\mathcal{C}_k|\mathbf{x})$, and then subsequently use decision theory to assign each new \mathbf{x} to one of the classes. Approaches that model the posterior probabilities directly are called *discriminative models*.
- (c) Find a function $f(\mathbf{x})$, called a discriminant function, which maps each input \mathbf{x} directly onto a class label. For instance, in the case of two-class problems, $f(\cdot)$ might be binary valued and such that $f = 0$ represents class \mathcal{C}_1 and $f = 1$ represents class \mathcal{C}_2 . In this case, probabilities play no role.

2 Approach (c)

2.1 Two classes

Two classes: $\{C_1, C_2\}$.

Linear discriminant function: $y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$.

Decision rules:
$$\begin{cases} y(\mathbf{x}) \geq 0 & \mathbf{x} \in C_1 \\ y(\mathbf{x}) < 0 & \mathbf{x} \in C_2 \end{cases}.$$

Decision boundary: $y(\mathbf{x}) = 0$.

2.2 Multiple classes

Multiple classes: $\{C_1, \dots, C_K\}$.

Linear discriminant function: $y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k,0}$.

Decision rules:
$$\begin{cases} y_k(\mathbf{x}) \geq y_j(\mathbf{x}) & \mathbf{x} \in C_k \\ y_k(\mathbf{x}) < y_j(\mathbf{x}) & \mathbf{x} \in C_j \end{cases}.$$

Decision boundary: $y_k(\mathbf{x}) = y_j(\mathbf{x})$ or $y_k(\mathbf{x}) - y_j(\mathbf{x}) = (\mathbf{w}_k - \mathbf{w}_j)^\top \mathbf{x} + (w_{k,0} - w_{j,0}) = 0$.

3 Probabilistic Generative Models

By **Bayes' theorem**: $P(C_k|\mathbf{x}) \propto P(\mathbf{x}|C_k)P(C_k)$.

Decision rules: Choose the class with largest posterior.

3.1 Two classes

The posterior probability for class C_1 can be written as

$$\begin{aligned} P(C_1|\mathbf{x}) &= \frac{P(C_1|\mathbf{x})}{P(C_1|\mathbf{x}) + P(C_2|\mathbf{x})} \\ &= \frac{P(\mathbf{x}|C_1)P(C_1)}{P(\mathbf{x}|C_1)P(C_1) + P(\mathbf{x}|C_2)P(C_2)} \\ &= \frac{1}{1 + \frac{P(\mathbf{x}|C_2)P(C_2)}{P(\mathbf{x}|C_1)P(C_1)}} \\ \sigma(a) &:= \frac{1}{1 + e^{-a}}, \end{aligned} \tag{1}$$

where

$$a(P(C_1|\mathbf{x})) := \ln \frac{P(\mathbf{x}|C_1)P(C_1)}{P(\mathbf{x}|C_2)P(C_2)}. \tag{2}$$

We call $\sigma(\cdot)$ as the **sigmoid** function, and $a(\cdot)$ as the **logit** function.




Figure 1: Sigmoid

Decision rules via posterior: Choose C_1 , if $P(C_1|\mathbf{x}) \geq 0.5$.

In order to classify an input \mathbf{x} , we need to compute posterior for each class, which is uniquely identified by the logit function a . So a is defined as **discriminant** here.

Decision rules via discriminant: Choose C_1 , if $a(P(C_1|\mathbf{x})) \geq 0$. ($\Leftrightarrow P(C_1|\mathbf{x}) \geq 0.5$)

3.1.1 Gaussian $p(\mathbf{x}|C_k)$ with the same Σ

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{N/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_k)^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_k) \right\}, k = 1, 2. \quad (3)$$

Plug (3) into (1), we have

$$\begin{aligned} a(P(C_1|\mathbf{x})) &= \ln \frac{P(\mathbf{x}|C_1)P(C_1)}{P(\mathbf{x}|C_2)P(C_2)} \\ &= \ln \frac{P(\mathbf{x}|C_1)}{P(\mathbf{x}|C_2)} + \ln \frac{P(C_1)}{P(C_2)} \\ &= \ln \frac{\exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_1)^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_1) \right\}}{\exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_2)^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_2) \right\}} + \ln \frac{P(C_1)}{P(C_2)} \\ &= \Sigma^{-1} \left[(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^\top \mathbf{x} + \left(-\frac{1}{2} \|\boldsymbol{\mu}_1\|^2 + \frac{1}{2} \|\boldsymbol{\mu}_2\|^2 \right) \right] + \ln \frac{P(C_1)}{P(C_2)} \\ &= \mathbf{w}^\top \mathbf{x} + w_0. \end{aligned}$$

So if class-conditional densities are Gaussian with the same covariance matrix, the discriminant is **linear** over \mathbf{x} .

Decision rules via discriminant: Choose C_1 , if $a(P(C_1|\mathbf{x})) = \mathbf{w}^\top \mathbf{x} + w_0 \geq 0$.

3.2 Multiple classes

The posterior probability for class C_k can be written as

$$\begin{aligned} P(C_k|\mathbf{x}) &= \frac{P(C_k|\mathbf{x})}{\sum_{j=1}^K P(C_j|\mathbf{x})} \\ &= \frac{P(\mathbf{x}|C_k)P(C_k)}{\sum_{j=1}^K P(\mathbf{x}|C_j)P(C_j)} \\ \text{softmax}(a_k) &:= \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}} \end{aligned} \tag{4}$$

where

$$a_k(\mathbf{x}) := \ln P(\mathbf{x}|C_k)P(C_k).$$

As posterior is identified by all a_k 's, they are defined as discriminants.

Decision rules via posterior: Choose C_k , if $P(C_k|\mathbf{x})$ is the largest.

Decision rules via discriminant: Choose C_k , if $a_k(\mathbf{x})$ is the largest.

Similarly, if class-conditional densities are Gaussian with the same covariance matrix, the discriminant is linear over \mathbf{x} .

Decision rules via discriminant: Choose C_k , if $a_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k,0} \geq a_j(\mathbf{x}) = \mathbf{w}_j^T \mathbf{x} + w_{j,0}$ for all $j \neq k$.




Figure 2: Three Gaussian models, the first two share the same covariance matrix

4 Probabilistic Discriminative Models

Directly **maximize a likelihood function** defined through the conditional distribution $p(C_k|\mathbf{x})$. Or equivalently, **minimize negative log-likelihood** (cross-entropy error).

4.1 Two classes

As in (1), posterior of class C_1 is in the form of sigmoid.

- 1) Build error function
- 2) Use gradient descent method to estimate parameters (the chain rule)
- 3) Compute $P(C_1|x) = \sigma(a)$, choose C_1 if $P(C_1|x) > 0.5$.

4.2 Multiple classes

As in (4), posterior of class C_k is in the form of softmax.

- 1) Build error function
- 2) Use gradient descent method to estimate parameters
- 3) Compute $P(C_k|x) = \text{sigmoid}(a_k)$, choose C_k if $P(C_k|x)$ is the largest.

Two-class Classification

Perceptron

Input vector $\mathbf{x} = [x_1, \dots, x_n]^T$, want to find

$$f(\mathbf{x}) = \text{sign} \left(\left(\sum_{j=1}^n w_j x_j \right) + w_0 \right)$$

The “bias weight” w_0 corresponds to the **threshold** when the neuron is triggered.

We have defined a Hypothesis set \mathcal{H} (dummy variable $x_0 \equiv 1$)

$$\mathcal{H} = \{f(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})\}$$

called the perceptron or linear separator

A perceptron fits the data by using a line to separate the +1 from -1 data

A simple learning model

- ▶ Input vector $\mathbf{x} = [x_1, \dots, x_d]^T$
- ▶ Given importance weights to the different inputs and compute a “Credit Score”
“Credit Score” = $\sum_{i=1}^d w_i x_i$.
- ▶ Approve credit if the “Credit Score” is acceptable
“Approve Score” = $\sum_{i=1}^d w_i x_i > \text{threshold}$. (“credit” is good)
“Deny Score” = $\sum_{i=1}^d w_i x_i < \text{threshold}$. (“credit” is bad)
- ▶ How to choose the importance weights w_i
 - input x_i is important \rightarrow large weight $|w_i|$
 - input x_i beneficial for credit \rightarrow positive weight $w_i > 0$
 - input x_i detrimental for credit \rightarrow negative weight $w_i < 0$

“Approve Score” = $\sum_{i=1}^d w_i x_i > \text{threshold}$. (“credit” is good)
 “Deny Score” = $\sum_{i=1}^d w_i x_i < \text{threshold}$. (“credit” is bad)

can be written formally as

$$f(\mathbf{x}) = \text{sign} \left(\left(\sum_{i=1}^d w_i x_i \right) + w_0 \right)$$

The “bias weight” w_0 corresponds to the threshold when the neuron is triggered

$$\mathbf{x} = [x_1, \dots, x_d], \mathbf{w}' = [w_1, \dots, w_d]$$

$$(1) \mathbf{w}'^T \mathbf{x} > \text{threshold}, Y;$$

$$(2) \mathbf{w}'^T \mathbf{x} \leq \text{threshold}, N;$$

(1) can be rewritten as $\mathbf{w}'^T \mathbf{x} - \text{threshold} = \mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T \mathbf{x} > 0$, where
 $w_0 = -\text{threshold}$, $\mathbf{w} = [w_1, \dots, w_d, w_0]$.

The perceptron learning algorithm (PLA)

The perceptron implements

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$$

Given the training set:

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

pick a **misclassified** point:

$$\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n$$

and update the weight vector:

$$\mathbf{w} \leftarrow \mathbf{w} + y_n \mathbf{x}_n$$



Why adding $\pm \mathbf{x}_i$ to \mathbf{w} ?

PLA implements our idea: start at some weights and try to improve it

“Incremental learning” on a single example at a time

$$\mathbf{a} \cdot \mathbf{b} = \cos \langle \mathbf{a}, \mathbf{b} \rangle \|\mathbf{a}\| \|\mathbf{b}\|$$

NOT REQUIRED!



Logistic Regression


Finding loss functions

A third linear prediction model

$$s = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \mathbf{x}$$


linear classification

$$h(\mathbf{x}) = \text{sign}(s)$$




linear regression

$$h(\mathbf{x}) = s$$




logistic regression

$$h(\mathbf{x}) = \theta(s)$$



The logistic function



$$\theta(s) = \frac{e^s}{1 + e^s} = \frac{1}{1 + e^{-s}}$$

Properties about θ :

$$\theta(-s) = 1 - \theta(s), \quad \theta'(s) = \frac{e^s}{(1 + e^s)^2} = \theta(s)(1 - \theta(s))$$

Error Measure: likelihood

$$P(y | \mathbf{x}) = \begin{cases} h(\mathbf{x}) & \text{for } y = +1; \\ 1 - h(\mathbf{x}) & \text{for } y = -1. \end{cases}$$

$$P(y | \mathbf{x}) = \theta(y \mathbf{w}^T \mathbf{x})$$

Properties about θ :

For an input \mathbf{x} , it has two possibilities: being labelled as $+1$, or -1 .

Compute $\text{score} = \mathbf{w}^T \mathbf{x}$,

1. if $\text{score} > 0$,

a. then it is more likely to be classified into $y = +1$,

$$P(y = +1 | \mathbf{x}) = h(\mathbf{x}) = \theta(\text{score}) \in (0.5, 1).$$

b. and is less likely to be classified into $y = -1$,

$$P(y = -1 | \mathbf{x}) = 1 - h(\mathbf{x}) = \theta(-\text{score}) \in (0, 0.5).$$

2. if $\text{score} < 0$,

a. then it is more likely to be classified into $y = -1$,

$$P(y = -1|\mathbf{x}) = h(\mathbf{x}) = \theta(\text{score}) \in (0, 0.5).$$

b. and is less likely to be classified into $y = +1$,

$$P(y = +1|\mathbf{x}) = 1 - h(\mathbf{x}) = \theta(-\text{score}) \in (0.5, 1).$$

So $P(y|\mathbf{x}) = \theta(y \cdot \text{score})$.

Likelihood of $\mathcal{D} = (\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ is

$$\prod_{n=1}^N P(y_n | \mathbf{x}_n) = \prod_{n=1}^N \theta(y_n \mathbf{w}^\top \mathbf{x}_n)$$

MLE

Maximize the likelihood, is to minimize:

$$\begin{aligned} & -\frac{1}{N} \ln \left(\prod_{n=1}^N \theta(y_n \mathbf{w}^\top \mathbf{x}_n) \right) \\ &= \frac{1}{N} \sum_{n=1}^N \ln \left(\frac{1}{\theta(y_n \mathbf{w}^\top \mathbf{x}_n)} \right) \end{aligned} \quad \left[\theta(s) = \frac{1}{1 + e^{-s}} \right]$$

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \underbrace{\ln \left(1 + e^{-y_n \mathbf{w}^\top \mathbf{x}_n} \right)}_{e(h(\mathbf{x}_n), y_n)} \quad \text{“cross-entropy” error}$$

Summary:

- $\text{score} = \mathbf{w}^T \mathbf{x}_i$, where $\mathbf{w} = [\mathbf{w}', w_0]$
- Perceptron: $\text{sign}(\text{score})$ for classification: $y_{\text{pred}} = \{+1, -1\}$
 - treat data locally, require the dataset to be linearly separable
 - if \mathbf{x}_i is misclassified, then update weights: $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$
- sigmoid(score) for probability: $P(y|\mathbf{x}) \in (0, 1)$
 - treat data globally, be tolerable to noise
 - Loss: negative log-likelihood