

## CS211 Computer Architecture II

### ISAs, Microprogramming and ROP

*Assigned 14/09/2020*

**Homework #1**

*Due 23:59:59 28/09/2020*

---

<https://toast-lab.gitee.io/courses/CS211@ShanghaiTech/Fall-2020/>

---

The homework is intended to help you learn the material, and we encourage you to collaborate with other students and to ask questions in discussion sections and office hours to understand the problems. However, each student must turn in their own solution to the problems.

The homework also provides essential background material for the mid-term and final exams. It will be graded primarily on an effort basis, but if you do not work through it, you are unlikely to succeed on the mid-term and final exams! We will distribute solutions to homework assignment. Note that each homework assignment is due at its respective due time, and all assignments are to be submitted **in English**. However, late submissions will **NOT** be accepted, except for extreme circumstances and/or with prior arrangement.

**Name:** \_\_\_\_\_

**ID:** \_\_\_\_\_

## Problem 1: CISC, RISC, and Stack: Comparing ISAs (2.5 points)

In this problem, your task is to compare three different ISAs: x86 (a CISC architecture with variable-length instructions), RISC-V (a load-store, RISC architecture with 32-bit instructions in its base form), and a stack-based ISA.

### Problem 1.A CISC (0.5 point)

---

Let us begin by considering the following C code:

```
int modulo(int x, int y) {
    for (; x >= y; x -= y);
    return x;
}
```

Using gcc and objdump on an x86 machine, we see that the above loop compiles to the following x86 instruction sequence. On entry to this code, register %edx contains x, and register %eax contains y. Throughout parts (1.A--1.C), we will ignore what happens in the done label and return statement.

```
start:
    cmp    %eax, %edx
    jl     done
loop:
    sub    %eax, %edx
    cmp    %edx, %eax
    jle    loop
done:
    ret
```

The meanings and instruction lengths of the instructions used above are given in the following table. Registers are denoted with  $R_{\text{SUBSCRIPT}}$ , register contents with  $\langle R_{\text{SUBSCRIPT}} \rangle$ .

Instruction	Operation	Length
sub $R_{\text{SRC}}, R_{\text{DEST}}$	$\langle R_{\text{DEST}} \rangle - \langle R_{\text{SRC}} \rangle \rightarrow R_{\text{DEST}}$	2 bytes
cmp $R_{\text{SRC1}}, R_{\text{SRC2}}$	$\langle R_{\text{SRC2}} \rangle - \langle R_{\text{SRC1}} \rangle \rightarrow \text{Temp}$	6 bytes
jl label	if ( $SF \neq OF$ ) jump to the address specified by label	2 bytes
jle label	if ( $(SF \neq OF) \text{ or } (ZF == 1)$ ) jump to the address specified by label	2 bytes
ret	return to the caller function	1 byte

Notice that the jump instructions jl and jle (jump if less, and jump if less or equal) depend on status flags, i.e., SF, ZF and OF in the table. Status flags are set by the instruction preceding a jump, based on the result of the computation. The meanings of the status flags are given in the following table:

Name	Purpose	Condition Reported
ZF	Zero flag	Result is zero ( <code>cmp</code> performs an actual subtraction)
SF	Sign flag	Set equal to high-order bit of result (0 if positive 1 if negative)
OF	Overflow flag	Set if result is too large a positive number or too small a negative number (excluding sign bit) to fit in destination operand; cleared otherwise.

How many bytes is the program? For the above x86 assembly code, how many bytes of instructions need to be fetched if  $x = 13$  and  $y = 3$ ? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored?

### Problem 1.B

### RISC (1 points)

Translate each of the x86 instructions in the following table into one or more RISC-V instructions. Place the `loop` label where appropriate. You should use the minimum number of instructions needed to translate each x86 instruction. Assume that upon entry, `x1` contains  $x$  and `x2` contains  $y$ . If needed, use `x3` as a condition register, and `x4`, `x5`, etc., for temporaries. You should not need to use any floating-point registers or instructions in your code. A description of the RISC-V instruction set architecture can be found in the class website alongside L03.

x86 instruction	label	RISC-V instruction sequence
	start:	
cmp     %eax, %edx		
j1       done		
sub     %eax, %edx		
cmp     %edx, %eax		
jle     loop		
	done:	
ret		

How many bytes is the RISC-V program using your direct translation? How many bytes of RISC-V instructions need to be fetched for  $x = 13$  and  $y = 3$  using your direct translation? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored?

**Problem 1.C****Stack (1 points)**

In a stack architecture, all operations occur on top of the stack. Only push and pop access memory, and all other instructions remove their operands from the stack and replace them with the result. The hardware implementation we will assume for this problem uses stack registers for the top two entries; accesses that involve other stack positions (e.g., pushing or popping something when the stack has more than two entries) use an extra memory reference. Assume each instruction occupies three bytes if it takes an address or label; other instructions occupy one byte.

Instruction	Definition
PUSH A	load value at M[A]; push value onto stack
POP A	pop stack; store value to M[A]
ADD	pop two values from the stack; ADD them; push result onto stack
SUB	pop two values from the stack; SUBtract top value from the 2nd; push result onto stack
INC	pop value from top of stack; increment value by one; push result onto stack
ZERO	zero out value at top of stack
DEC	pop value from top of stack; decrement value by one; push result onto stack
BEQZ <i>label</i>	pop value from stack; if it's zero, continue at <i>label</i> ; else, continue with next instruction
BNEZ <i>label</i>	pop value from stack; if it's not zero, continue at <i>label</i> ; else, continue with next instruction
BNEG <i>label</i>	pop value from stack; if it's negative (less than zero), continue at <i>label</i> ; else, continue with next instruction
JUMP <i>label</i>	continue execution at location <i>label</i>

Translate the `modulo` loop to the stack ISA. For uniformity, please use the same control flow as in parts (1.A) and (1.B). Assume that when we reach the loop, `x` and `y` are at the top of the stack. At the end of the loop, `x` should be at the top of the stack.

How many bytes is the stack program using your translation? How many bytes of instructions need to be fetched for `x = 13` and `y = 3` using your translation? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored? Would the number of bytes loaded and stored change if the stack could fit 8 entries in registers?

## Problem 2: ROP (0.5 point)

In this problem, we explore how to explore and leverage code snippets to launch a return-oriented programming (ROP) attack. In L03, we have studied that code snippets ended with `ret`, also known as gadgets, can be chained to achieve some purpose.

### Problem 2.A

### Modular multiplication (0.5 point)

Many cryptographic algorithms, such as RSA and Diffie-Hellman key exchange, are based on modular multiplication, e.g.,  $A * B \bmod N$ . In this problem, we would consider using ROP to make a modular multiplication. Some code snippets that might be useful for you are listed below.

```
/* For instructions with two operands, the left operand is R_SRC
(source register) or immediate value, and the right one is R_DEST (destination
register). */
```

```
start:                                /* address: 0x8000 */
    cmp    %eax, %edx
    jl     done
loop:                                     /* address: 0x8008 */
    sub    %eax, %edx
    cmp    %edx, %eax
    jle    loop
done:                                       /* address: 0x8011 */
    ret
...

G0:                                       /* address: 0x8080 */
    xor    %ebx, %ebx                    /* <R_SRC> xor <R_DEST> → R_DEST */
    mov    $0x08, %eax                  /* imm → R_DEST */
    ret
...

G1:                                       /* address: 0x9090 */
    add    %ecx, %ebx                    /* <R_SRC> + <R_DEST> → R_DEST */
    ret
...

G2:                                       /* address: 0xA0A0 */
    mov    %ebx, %edx                    /* <R_SRC> → <R_DEST> */
    mov    %ecx, %eax
    sub    %eax, %edx                    /* <R_DEST> - <R_SRC> → R_DEST */
    ret
...

G3:                                       /* address: 0xB0B0 */
    dec    %ebx                          /* <R_DEST> - 1 → <R_DEST> */
    ret
...
```

```

G4:                                     /* address: 0xC0C0 */
    pop    %ecx                       /* pop stack; store the value to R_DEST */
    ret
...

G5:                                     /* address: 0xD0D0 */
    shl    $0x08, %ebx               /* shift bits of <R_DEST> to the left for 0x08 times → <R_DEST> */
    ret
...

G6:                                     /* address: 0xE0E0 */
    and     $0x07, %ecx               /* <R_DEST> bitwise_and 0x07 → <R_DEST> */
    and     $0x07, %edx
    ret
...

G7:                                     /* address: 0xF0F0 */
    mul     %ebx, %ecx               /* <R_SRC> * <R_DEST> → R_DEST */
    mov     %edx, %ebx
    ret
...

```

Using all or some of these code snippets, could you build a ROP case to calculate  
 $(1689 * 255 \% 8)$

and store the result into `ebx` register?

If yes, please illustrate one case. Otherwise, give your reason. You can draw a diagram for illustration with descriptions. When you draw the memory space, you can use a simplified layout that is similar to ones shown in the lecture slides.

*Hint: You may find none, one, or multiple cases. Bonus points may be accredited to students that submit reasonable and interesting answer.*