



CS240 Algorithm Design and Analysis

Lecture 24

Approximation Algorithms

Fall 2021
2021.12.13



Approximation Algorithms



- Up to now, most of our algorithms have been exact, i.e., they find an optimal solution.
- But there are many problems for which we don't know how to find an optimal solution.
 - A key example is NP-complete problems. We don't know efficient algorithms for any NPC problem.
- Many such problems are important in practice. What do we do?
- If we can't get find the best answer, let's try for good enough.
- Approximation algorithms find an approximately optimal answer.



Approximation Ratio



- Let X be a maximization problem. Let A be an algorithm for X .
- Let $a > 1$ be a constant.
- A is an a -approximation algorithm for X if A always returns an answer that's at least $1/a$ times the optimal.
 - **Ex** If X is max-flow, A is a 2-approx algorithm if it always returns a flow that's at least $\frac{1}{2}$ the optimal.
 - The closer a is to 1, the better the approximation.
- If X is a minimization problem, A is an a -approximation algorithm for X if it always returns an answer that's at most a times larger than the optimal.
 - **Ex** If X is min-cut, A is a 2-approx algorithm if it always returns a cut that's at most 2 times the size of the optimal.
 - Again, the closer a is to 1, the better the approximation.







Coverings



- Suppose there's a set of teachers, and each can teach a certain set of classes.
 - Let S_i be the set of classes teacher i can teach.
- The entire set of classes is X .
- We want to pick the minimum set of teachers to teach all the classes.
 - Let T be set of teachers we pick.
 - We want $\bigcup_{i \in T} S_i = X$, and T to be the smallest possible.






Set Covering



- **Input** A collection F of sets. Each set has a cost. The union of all the sets is X .
- **Output** A subset G of F , whose union is X .
- **Goal** Minimize the total cost of the sets in G .



If all sets have same cost, S_3 , S_4 and S_5 is a min cost set cover of X .

- Minimum cost set cover is NP-complete.
- We'll see a $\ln(n)$ -approximation algorithm, where $n=|X|$.






A Greedy Approximation Algorithm



- A natural greedy heuristic is to choose sets which cover points most cheaply.
 - For each set, let c be its cost, and m be the number of points it covers.
 - We want to use the set with the smallest c/m value, because this is the cheapest way to cover some new points.
- After we pick this set, remove all the points it covers. Then we consider the per unit cost of the remaining sets and again choose the cheapest.





A Greedy Approximation Algorithm



- F is the entire collection of sets. The union of these sets is X .
- Each set S in F has a cost $\text{cost}(S)$.
- U is the set of elements of X we haven't covered yet.
- C is the set cover we eventually output.
- $U = X$
- $C = \emptyset$
- while $U \neq \emptyset$
 - choose $S \in F - C$ with $\min |\text{cost}(S)|/|S \cap U|$
 - $C = C \cup \{S\}$
 - $U = U - S$
- output C

Per unit cost to cover
new elements.






Proof of Correctness



- We always output a set cover, because the while loop continues till X is covered.
- We'll prove the approximation ratio is at most $1 + 1/2 + 1/3 + \dots + 1/n \approx \ln(n)$.
 - If the min cost of a set cover is V , our set cover costs at most $\ln(n) * V$.
- The basic plan is to bound the cost of the set cover the algorithm outputs using the "average cost" per element.

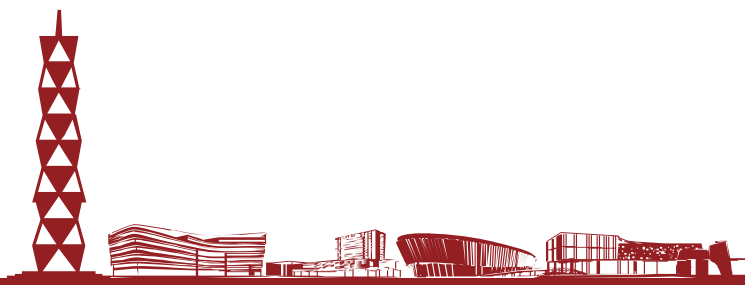




Proof of Correctness



- Order the sets in C by when they're added to C , earliest set first.
 - Let the order be S_1, S_2, \dots, S_m .
- Cost of the set cover is $L = \sum_i \text{cost}(S_i)$.
- Order the elements in X by when they're added, earliest element first.
 - Let the order be e_1, e_2, \dots, e_n .
 - So, the first few e 's are added by S_1 , the next few added by S_2 , etc.
 - Every element in X is in the list, because C covers X .





Proof of Correctness



- Let n_i be the number of new elements S_i covers.
 - So, n_i is the number of elements in S_i , but not in S_1, \dots, S_{i-1} .
 - Divide the cost of S_i evenly among the new elements it covers.
 - If e is newly covered by S_i , then $\text{cost}(e) = \text{cost}(S_i)/n_i$.
- $$\sum_k \text{cost}(e_k) = \sum_i n_i * \frac{\text{cost}(S_i)}{n_i} = \sum_i \text{cost}(S_i) = L$$
- Every element is covered by some S_i , and S_i covers n_i new elements.
- We'll prove $\text{cost}(e_k) \leq \text{OPT}/(n-k+1)$, for any k .
 - Suppose this is true, then
$$L = \sum_k \text{cost}(e_k) \leq \sum_k \text{OPT}/(n-k+1) \approx \ln(n) * \text{OPT}$$





The Per Element Cost



- Let's focus on some element e_k , and let S_j be the set which covers e_k for the first time.
- Let C_1, \dots, C_m be the sets in an optimal cover, each of which covers some elements of $U = \{e_k, e_{k+1}, e_{k+2}, \dots, e_n\}$.
 - Let n'_1, \dots, n'_m be the number of elements of U which C_1, \dots, C_m cover.
- **Obs 1: $\sum_i n'_i \geq n - k + 1$.**
 - Because C_1, \dots, C_m cover U .
- **Obs 2: $\sum_i \text{cost}(C_i) \leq \text{OPT}$.**
 - Because C_1, \dots, C_m are a subset of an optimal cover, which has cost OPT .





The Per Element Cost



- **Obs 3** None of C_1, \dots, C_m are among S_1, \dots, S_{j-1} .
 - If some C_i is among S_1, \dots, S_{j-1} , then since C_i covers some e in U , e would be covered by $\{S_1, \dots, S_{j-1}\}$. So, e would be among the first $k-1$ elements covered. Contradiction.
- **Obs 4** There exists some C_i among C_1, \dots, C_m with $\frac{\text{cost}(C_i)}{n'_i} \leq \text{OPT}/(n - k + 1)$.
 - If every C_i in C_1, \dots, C_m has $\frac{\text{cost}(C_i)}{n'_i} \geq \text{OPT}/(n - k + 1)$, then
$$\begin{aligned} & \text{OPT} \\ & \geq \sum_i \text{cost}(C_i) = \sum_i n'_i * \frac{\text{cost}(C_i)}{n'_i} > \sum_i n'_i * \frac{\text{OPT}}{n - k + 1} \geq \text{OPT}/(n - k + 1) \sum_i n'_i \geq \frac{\text{OPT}}{n - k + 1} * (n - k + 1) = \text{OPT} \end{aligned}$$
Contradiction.





Proof of Approximation Ratio




- Lemma $cost(S_j)/n_j \leq OPT/(n - k + 1)$
- **Proof** When choosing S_j , the only sets the algorithm is not allowed to choose are S_1, \dots, S_{j-1} .
 - By obs 3, C_1, \dots, C_m aren't in S_1, \dots, S_{j-1} .
 - By obs 4, there's some C_i in C_1, \dots, C_m , with $\frac{cost(C_j)}{n'_i} \leq OPT/(n - k + 1)$.
 - S_j was chosen so that $cost(S_j)/n_j$ is min among all sets not in S_1, \dots, S_{j-1} .
 - So $\frac{cost(S_j)}{n_j} \leq \frac{cost(C_i)}{n'_i} \leq OPT/(n - k + 1)$.
- Since $\frac{cost(S_j)}{n_j} = cost(e_k)$, we have $cost(e_k) \leq OPT/(n - k + 1)$.
- The approximation ratio follows because
$$L = \sum_k cost(e_k) = \sum_k \frac{OPT}{n-k+1} \approx \ln(n) * OPT$$





Scheduling







Parallel Computing and Scheduling



- Computers today are parallel.
 - Multiple processors in a system.
 - Multiple tasks for the processors to run.
- Multiprocessor scheduling is the problem of deciding which tasks to run on which processors at what time.
- Many possible objectives.
 - Throughput, fairness, energy usage.
 - Latency, i.e. finishing all jobs as fast as possible.






Makespan Scheduling



- n independent jobs.
 - Jobs have different sizes, i.e. time needed to perform job.
 - Jobs can be done in any order.
 - Any job can be done on any machine.
- m processors.
 - All have the same speed.
 - Each processors can do one job at a time.
- Assign the jobs to the processors.
- Makespan is when the last processor finishes all its jobs.
- Minimize the makespan.
 - I.e., finish all the jobs as fast as possible.

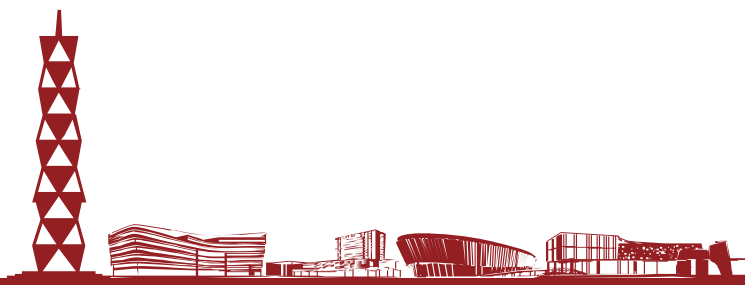




Minimizing makespan is NPC



- The decision version of scheduling is obviously in NP.
- SUBSET-SUM: given a set of numbers S and target t , is there a subset of S summing to t ?
 - **Ex** $S=\{1,3,8,9\}$. $t=9$, yes. $t=14$, no.
 - This is NP-complete. We reduce SUBSET-SUM to scheduling.
- Let (S,t) be an instance of SUBSET-SUM.
 - Let s be sum of all elements in S .
- Make a set of jobs $J = S \cup \{s-2t\}$, and schedule them on 2 processors.





Minimizing makespan is NPC



- **Claim** If some subset of S sums to t , then min makespan is $s-t$.
- **Proof** Say $S' \subseteq S$ sums to t . Schedule the jobs in S' and job $s-2t$ on processor 1. So proc 1 finishes at time $t+s-2t=s-t$. Proc 2 does the jobs in $S-S'$, so it finishes at time $s-t$ as well.
- **Claim** If the min makespan is $s-t$, there exists a subset of S that sums to t .
- **Proof** Suppose WLOG proc 1 does the $s-2t$ job. Since makespan is $s-t$, the other jobs proc 1 does must have total size $s-t-(s-2t)=t$.
- So (S,t) is yes instance of SUBSET-SUM iff makespan = $s-t$.
 - So SUBSET-SUM \leq_p scheduling, and scheduling is NP-complete.





Graham's List Scheduling



- Since scheduling is NPC, it's unlikely we can find the min makespan in polytime.
- List scheduling is a simple greedy algorithm.
 - Finds a schedule with makespan at most twice the minimum.
 - A 2-approximation.
- If there are n tasks and m processors, list scheduling only takes $O(n \log n)$ time.
 - Compare this to $n!$ $C(n+m-1, m-1)$ time to try all possible schedules and pick the best.






Graham's List Scheduling



- List the jobs in any order.
- As long as there are unfinished jobs.
 - If any processor doesn't have a job now, give it the next job in the list.
- Example
 - 3 processors. The jobs have length 2, 3, 3, 4, 5, 6, 8.
 - List them in any order. Say 4, 5, 3, 2, 6, 8, 3.
 - Initially, no proc has a job. Give first 3 jobs to the 3 procs.
 - At time 3, proc 3 is done. Give it next job in list, 2.
 - At time 4, proc 2 is done. Give it next job in list, 6.
 - At time 5, both 1, 3 are done. Give them next jobs in list, 8, 3.
 - Everybody finishes by time 13.
 - The makespan of this schedule is 13.







The Worst Case for LS



- How badly can list scheduling do compared to optimal?
- Say there are m^2 jobs with length 1, and one job with length m .
 - Suppose they're listed in the order $1, 1, 1, \dots, 1, m$.
 - LS has makespan $2m$. Optimal makespan is $m+1$.
 - $\text{makespan}(\text{LS}) / \text{makespan}(\text{opt}) = 2m/(m+1) \approx 2$.
- This is worst possible case for list scheduling.



$\text{makespan}(\text{LS}) = 2m$



$\text{makespan}(\text{opt}) = m+1$





LS is a 2-approximation




- Next, we prove LS always gives a schedule at most twice the optimal.
- Suppose LS gives makespan of M .
- Let the optimal schedule have makespan M^* .
- We prove that $M \leq 2M^*$.





LS is a 2-approximation




- The picture above is the schedule produced by list scheduling.
- Consider task X that finishes last.
 - Say X starts at time T, and has length t.
- **Claim 1** $M^* \geq t$.
 - In any schedule, X has to run on some process.
 - Since X takes t time, every schedule, including the opt, takes $\geq t$ time.





LS is a 2-approximation



■ Claim 2 $M^* \geq T$.

□ Up to time T , no processor is ever idle.

■ Up to T , there's always some unfinished job.

■ As soon as a processor finishes one job, it's assigned another one.

□ So at time T , each processor completed T units of work.

□ So total amount of work in all the jobs is $\geq mT$. Up to T : mT


□ In the opt schedule, m processors complete at most m units of work per time unit.

□ So length of opt schedule is $\geq (\text{total work})/m \geq mT/m = T$.





LS is a 2-approximation



- From Claims 1 and 2, we have $M^* \geq t$ and $M^* \geq T$.
- So $M^* \geq \max(T, t)$.
- $M = T + t$, because X is last job to finish.
- So $M/M^* \leq (T+t)/\max(T, t) \leq 2$.





LPT Scheduling



- Worst case for LS occurred when longest job was scheduled last.
 - Large jobs are “dangerous” at end.
- Let's try to schedule longest jobs first.
- Longest processing time (LPT) schedule is just like list scheduling, except it first sorts tasks by nonincreasing order of size.
- **Ex** For three processors and tasks with sizes 2, 3, 3, 4, 5, 6, 8, LPT first sorts the jobs as 8,6,5,4,3,3,2. Then it assigns p_1 tasks 8,3, p_2 tasks 6,3, p_3 tasks 5,4,2, for a makespan of 11.
- LPT has an approximation ratio of $4/3$.






LPT is a $4/3$ -approximation



- **Thm** Suppose the optimal makespan is M^* , and LPT produces a schedule with makespan M . Then $M \leq 4/3 M^*$.
- Let X be the last job to finish. Assume it starts at time T and has size t .
- Assume WLOG that X is the last job to start.
 - If not, then say Y starts after T .
 - Y finishes before $T+t$. So we can remove Y without increasing the makespan.
- **Cor 1** X is the smallest job.
 - X is the last job to start, so due to LPT scheduling it's the smallest.






LPT is a 4/3-approximation



- **Claim 1** LPT's makespan = $T + t \leq M^* + t$.
 - As in LS, no processor is idle up to time T , so $M^* \geq T$.
- **Case 1** $t \leq M^*/3$.
 - Then LPT's makespan $\leq M^* + t \leq M^* + M^*/3 = 4/3 M^*$.
- **Case 2** $t > M^*/3$.
 - Since X is the smallest task, all tasks have size $> M^*/3$.
 - So the optimal schedule has at most 2 tasks per processor. So $n \leq 2m$.
 - If $1 \leq n \leq m$, then LPT and optimal schedule both put one task per processor.
 - If $m < n \leq 2m$, then optimal schedule is to put tasks in nonincreasing order on processors $1, \dots, m$, then on $m, \dots, 1$.
 - LPT also schedules tasks this way, so it's optimal.

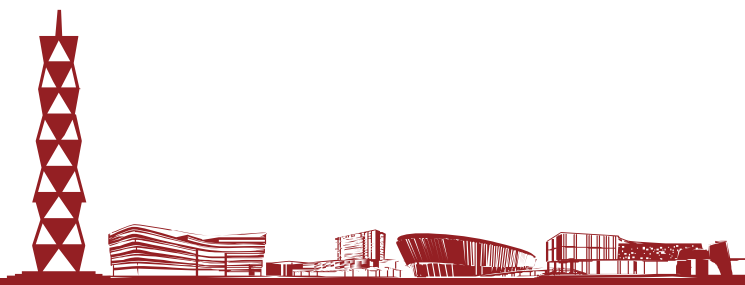




LS VS. LPT



- LPT gives better approximation ratio, has same running time. Why bother with LS?
- LS is online.
 - Imagine the jobs are coming one by one.
 - LS just puts them on any idle computer.
- LPT is offline
 - It needs to know all the jobs that will ever arrive, in order to sort them.
- In a realistic parallel computation, you get jobs on the fly.
 - Online is more realistic.
 - LS is usually more useful.





Next Time: Approximation algorithms (Cont.)

