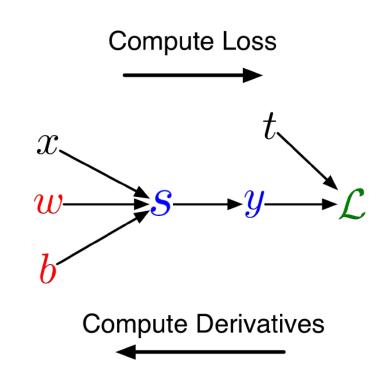
# Lecture 4: CNNs I - Architecture & Equivariance

Xuming He SIST, ShanghaiTech Fall, 2020



# Computation graph

- Represent the computations using a computation graph
  - □ Nodes: inputs & computed quantities
  - Edges: which nodes are computed directly as function of which other nodes





# **General Backpropagation**

#### Given a computation graph

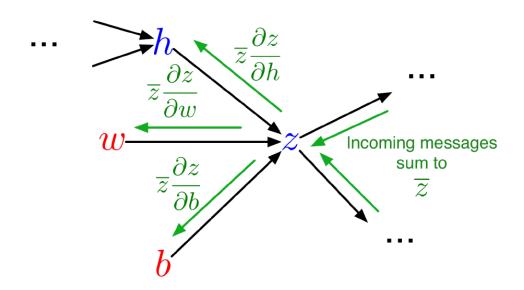
Let  $v_1, \ldots, v_N$  be a topological ordering of the computation graph (i.e. parents come before children.)

 $v_N$  denotes the variable we're trying to compute derivatives of (e.g. loss)



# **General Backpropagation**

Backprop as message passing:

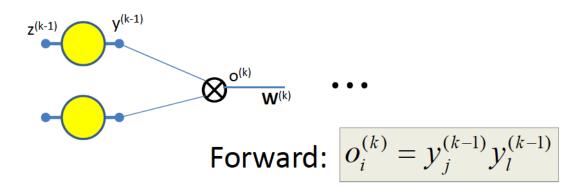


- Each node receives a set of messages from its children, which are aggregated into its error signal, then it passes messages to its parents
- Modularity: each node only has to know how to compute derivatives w.r.t. its arguments – local computation in the graph



#### Patterns in backward flow

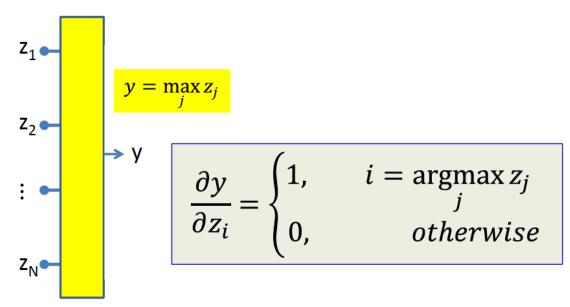
#### Multiplicative node



$$\frac{\partial L}{\partial y_i^{(k-1)}} = \frac{\partial L}{\partial o_i^{(k)}} \frac{\partial o_i^{(k)}}{\partial y_i^{(k-1)}} = y_l^{(k-1)} \frac{\partial L}{\partial o_i^{(k)}}$$

# Patterns in backward flow

#### Max node



- Vector equivalent of subgradient
  - 1 w.r.t. the largest incoming input
    - Incremental changes in this input will change the output
  - 0 for the rest
    - Incremental changes to these inputs will not change the output



# Vector form of BackProp

Review: Jacobian of vector functions

$$\mathbf{J} = \left[ egin{array}{cccc} rac{\partial \mathbf{f}}{\partial x_1} & \cdots & rac{\partial \mathbf{f}}{\partial x_n} \end{array} 
ight] = \left[ egin{array}{cccc} rac{\partial f_1}{\partial x_1} & \cdots & rac{\partial f_1}{\partial x_n} \ dots & \ddots & dots \ rac{\partial f_m}{\partial x_1} & \cdots & rac{\partial f_m}{\partial x_n} \end{array} 
ight].$$

Vectorized chain rule

$$\mathbf{x} \in \mathbb{R}^{m}, \mathbf{y} \in \mathbb{R}^{n} \qquad g : \mathbb{R}^{m} \to \mathbb{R}^{n}, \mathbf{y} = g(\mathbf{x})$$

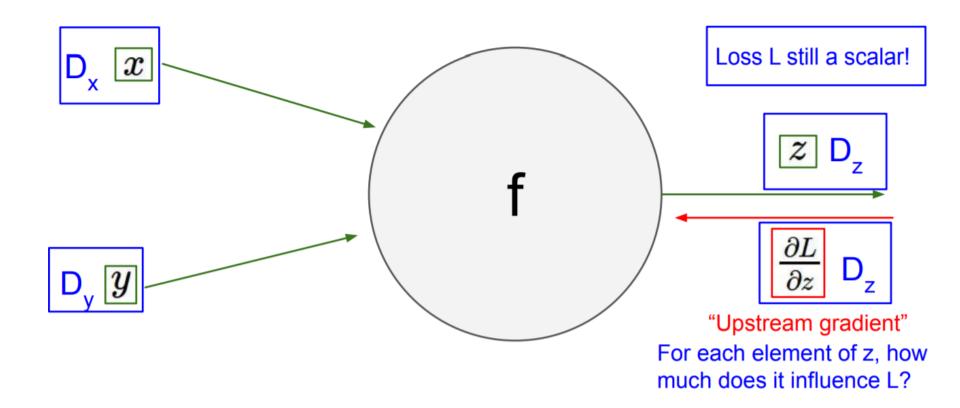
$$\frac{\partial z}{\partial \mathbf{x}_{i}} = \sum_{j=1}^{n} \frac{\partial z}{\partial \mathbf{y}_{j}} \frac{\partial \mathbf{y}_{j}}{\partial \mathbf{x}_{i}}$$

$$f : \mathbb{R}^{n} \to \mathbb{R}, z = f(\mathbf{y})$$

$$\nabla_{\mathbf{x}} z = \left[\frac{\partial \mathbf{y}_{j}}{\partial \mathbf{x}_{i}}\right] \nabla_{\mathbf{y}} z = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}}\right)^{T} \nabla_{\mathbf{y}} z$$

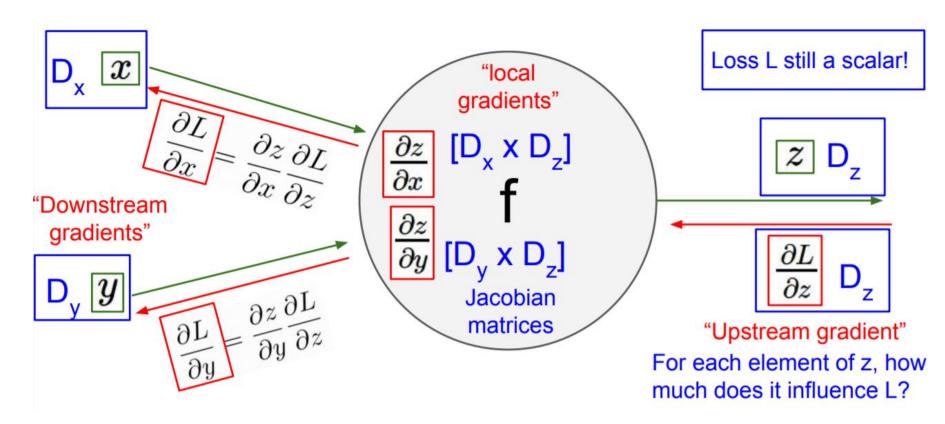
# Vector form of BackProp

Forward pass with vectors



# Vector form of BackProp

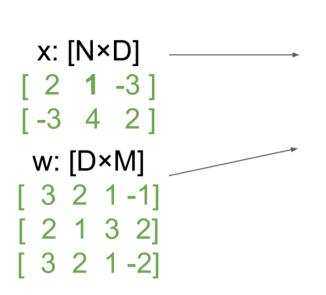
Note: here the Jacobian matrices are actually the transpose of the standard version.





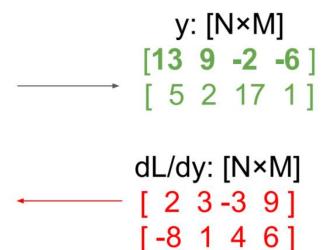
# Matrix form of BackProp

- Often used in mini-batches
  - □ N is the batch size, for instance.



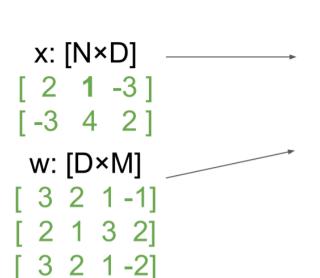
#### Matrix Multiply

$$y_{n,m} = \sum_{d} x_{n,d} w_{d,m}$$



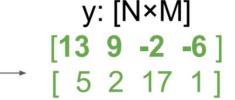


- Often used in mini-batches
  - □ N is the batch size, for instance.



#### Matrix Multiply

$$y_{n,m} = \sum_{d} x_{n,d} w_{d,m}$$



#### Jacobians:

dy/dx:  $[(N\times D)\times (N\times M)]$ dy/dw:  $[(D\times M)\times (N\times M)]$ 

For a neural net we may have
N=64, D=M=4096
Each Jacobian takes 256 GB of memory!
Must work with them implicitly!

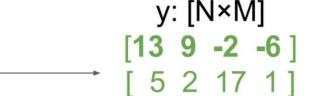


- Often used in mini-batches
  - □ N is the batch size, for instance.

# x: [N×D] [ 2 1 -3 ] [ -3 4 2 ] w: [D×M] [ 3 2 1 -1] [ 2 1 3 2] [ 3 2 1 -2]

#### Matrix Multiply

$$y_{n,m} = \sum_{d} x_{n,d} w_{d,m}$$



#### $[N\times D]$ $[N\times M]$ $[M\times D]$

$$\frac{\partial L}{\partial x} = \left(\frac{\partial L}{\partial y}\right) w^T$$

$$\frac{\partial L}{\partial x_{n,d}} = \sum_{m} \frac{\partial L}{\partial y_{n,m}} \frac{\partial y_{n,m}}{\partial x_{n,d}} = \sum_{m} \frac{\partial L}{\partial y_{n,m}} w_{d,m}$$



- Often used in mini-batches
  - N is the batch size, for instance.

#### **Matrix Multiply**

$$y_{n,m} = \sum_{d} x_{n,d} w_{d,m}$$

#### [N×D] [N×M] [M×D]

$$\frac{\partial L}{\partial x} = \left(\frac{\partial L}{\partial y}\right) w^T$$

#### [D×M] [D×N] [N×M]

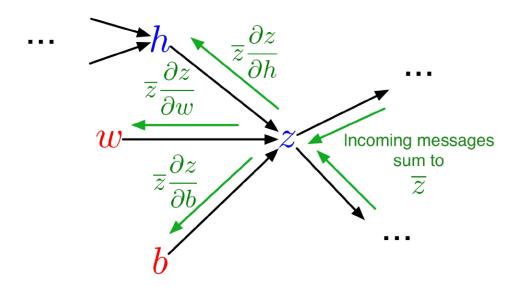
$$\frac{\partial L}{\partial w} = x^T \left( \frac{\partial L}{\partial y} \right)$$

These formulas are easy to remember: they are the only way to make shapes match up!



# Computation cost

- Forward pass: one add-multiply operation per weight
- Backward pass: two add-multiply operations per weight



 For a multilayer network, the cost is linear in the number of layers, quadratic in the number of units per layer



# Backpropagation

- Backprop is used to train the majority of neural nets
  - □ Even generative network learning, or advanced optimization algorithms (second-order) use backprop to compute the update of weights
- However, backprop seems biologically implausible
  - □ No evidence for biological signals analogous to error derivatives
  - All the existing biologically plausible alternatives learn much more slowly on computers.
  - □ So how on earth does the brain learn???



#### **Outline**

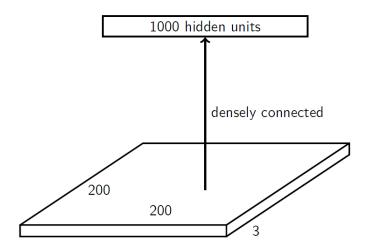
- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - Convolution layers & model complexity
  - Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes



### **Motivation**

- Visual recognition
  - Suppose we aim to train a network that takes a 200x200 RGB image as input



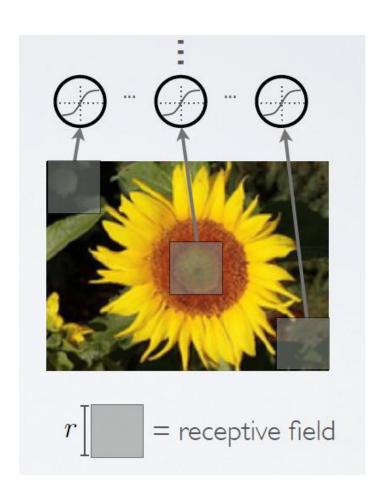
- □ What is the problem with have full connections in the first layer?
  - Too many parameters! 200x200x3x1000 = 120 million
  - What happens if the object in the image shifts a little?



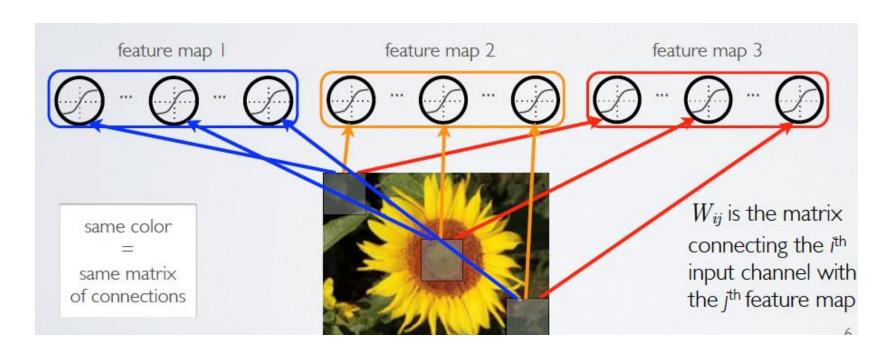
# Our goal

- Visual Recognition: Design a neural network that
  - □ Much deal with very high-dimensional inputs
  - □ Can exploit the 2D topology of pixels in images
  - Can build in invariance/equivariance to certain variations we can expect
    - Translation, small deformations, illumination, etc.
- Convolution networks leverage these ideas
  - Local connectivity
  - Parameter sharing
  - Pooling/subsampling hidden units

- First idea: Use a local connectivity of hidden units
  - Each hidden unit is connected only to a subregion (patch) of the input image
  - Usually it is connected to all channels
  - Each neuron has a local receptive field

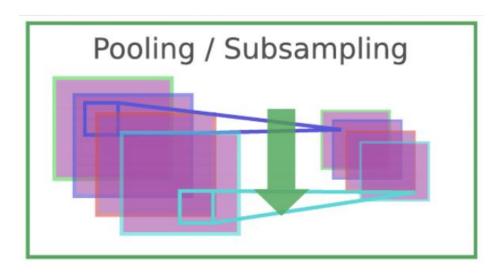


- Second idea: share weights across certain units
  - Units organized into the same "feature map" share weight parameters
  - Hidden units within a feature map cover different positions in the image

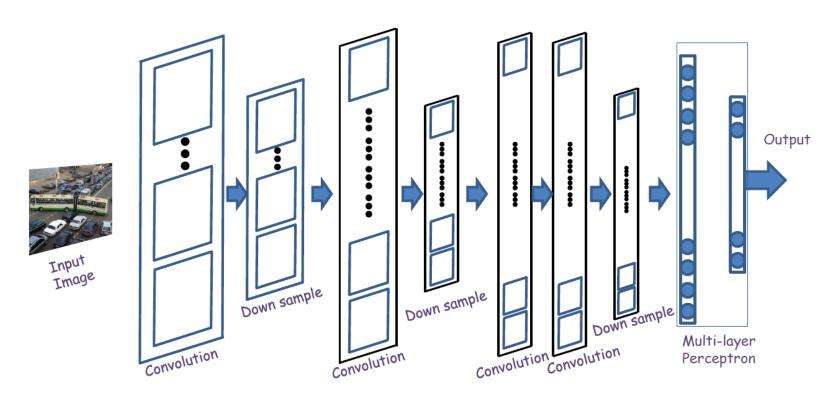




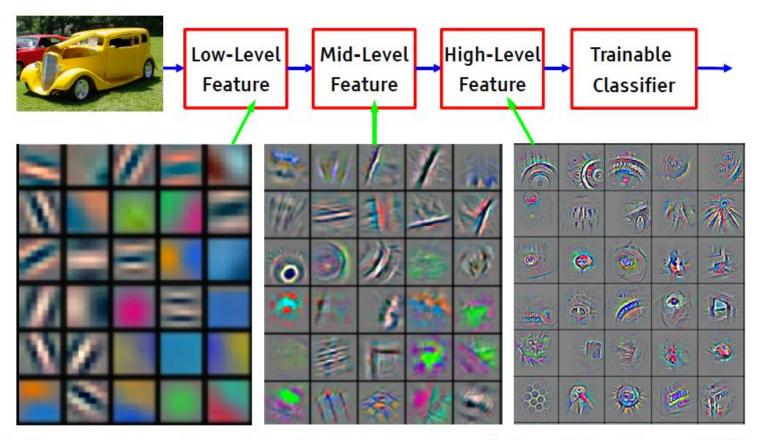
- Third idea: pool hidden units in the same neighborhood
  - □ Averaging or Discarding location information in a small region
  - Robust toward small deformations in object shapes by ignoring details.



- Fourth idea: Interleaving feature extraction and pooling operations
  - Extracting abstract, compositional features for representing semantic object classes



 Artificial visual pathway: from images to semantic concepts (Representation learning)



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



#### **Outline**

- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - Convolution layers & model complexity
  - Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes

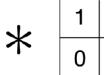


# 2D Convolution

If A and B are two 2-D arrays, then:

$$(A*B)_{ij} = \sum_{s} \sum_{t} A_{st} B_{i-s,j-t}.$$

1	3	1
0	-1	1
2	2	-1



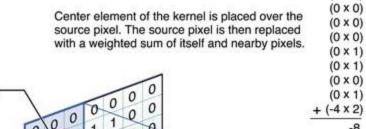
			1 0				
1	3	1	× 2 1	1	5	7	2
	4	4		0		_1	4
U	- 1	1			-2	-4	ı
2	2	-1		2	6	4	-3
				5	-2	-2	1
				U	-2		1

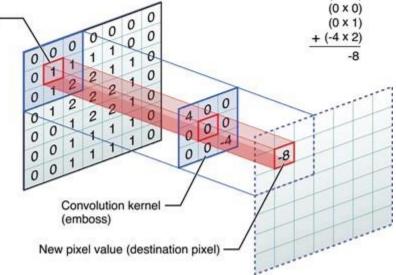
# 2D Convolution

 $(4 \times 0)$ 

If A and B are two 2-D arrays, then:

$$(A*B)_{ij} = \sum_{s} \sum_{t} A_{st} B_{i-s,j-t}.$$





1,	1,0	1,	0	0
0,0	1,	1,0	1	0
0,	0,0	1,	1	1
0	0	1	1	0
0	1	1	0	0



4	

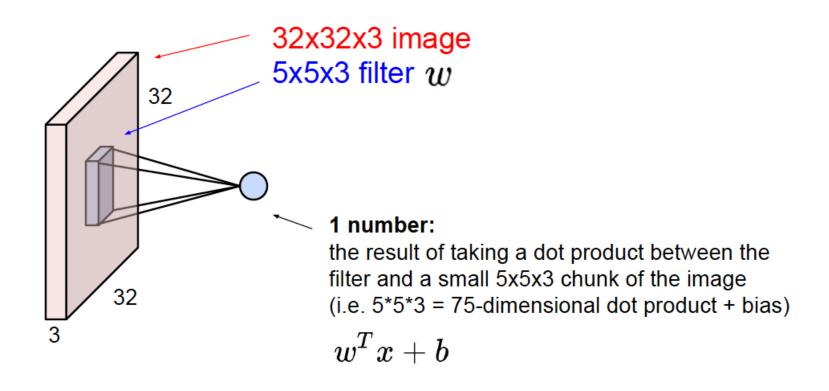
Convolved Feature

Picture Courtesy: developer.apple.com

Source pixel



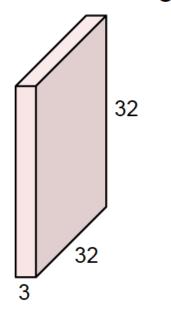
#### Formal definition





Define a neuron corresponding to a 5x5 filter

#### 32x32x3 image



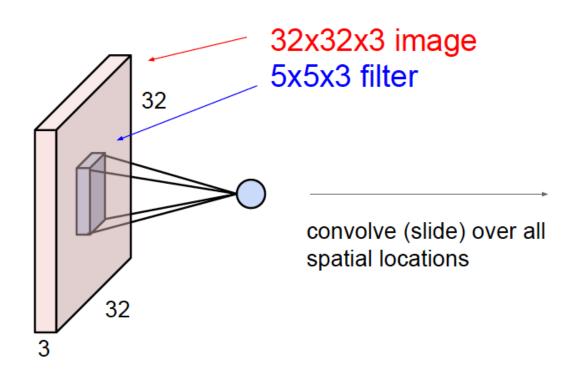
5x5x3 filter



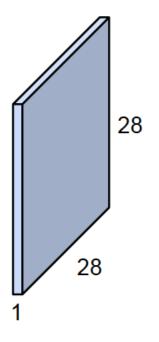
**Convolve** the filter with the image i.e. "slide over the image spatially, computing dot products"



- Convolution operation
  - Parameter sharing
  - Spatial information

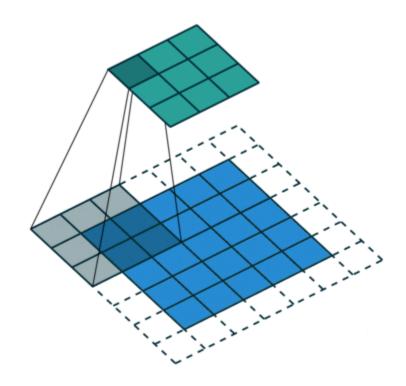


#### activation map





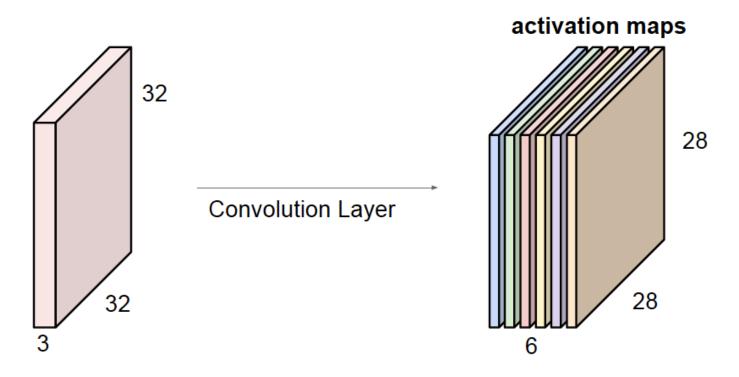
- Convolution operation
  - □ Parameter sharing
  - Spatial information





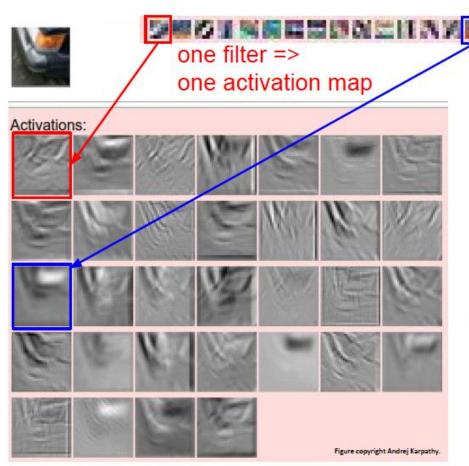
#### Multiple kernels/filters

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a "new image" of size 28x28x6!

Visualizing the filters and their outputs



example 5x5 filters (32 total)

We call the layer convolutional because it is related to convolution of two signals:

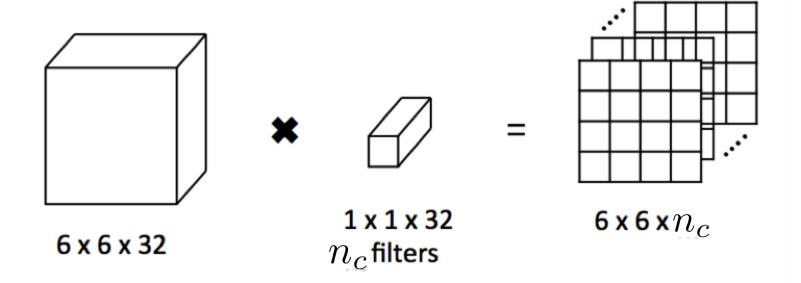
$$f[x,y] * g[x,y] = \sum_{n_1 = -\infty}^{\infty} \sum_{n_2 = -\infty}^{\infty} f[n_1, n_2] \cdot g[x - n_1, y - n_2]$$

elementwise multiplication and sum of a filter and the signal (image)



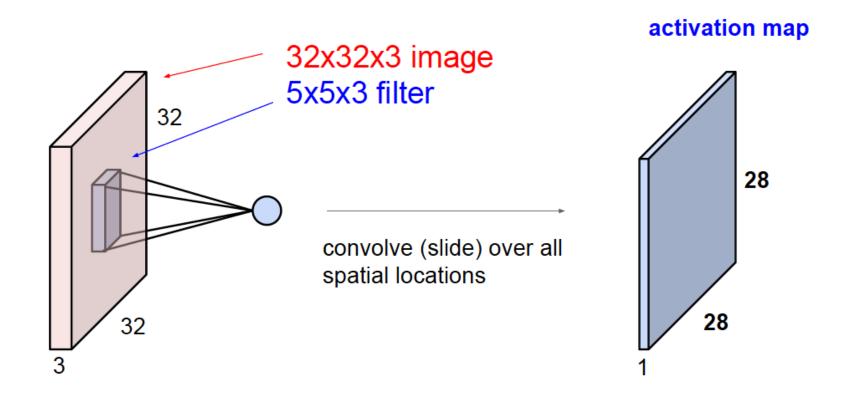
# **Special Convolutions**

- 1x1 convolutions
  - □ Used in Network-in-network, GoogleNet
  - Reduce or increase dimensionality
  - Can be considered as 'feature pooling"



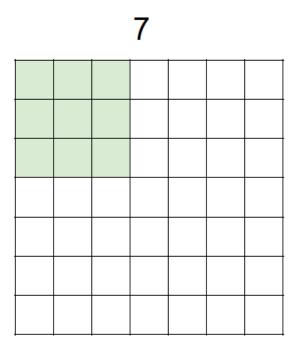
# Complexity of Convolution Layers

Sizes of activation maps and number of parameters



# Complexity of Convolution Layers

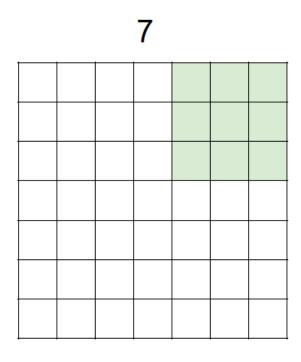
Size of activation maps



7x7 input (spatially) assume 3x3 filter

# Complexity of Convolution Layers

Size of activation maps

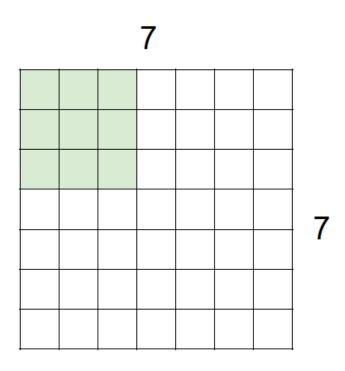


7x7 input (spatially) assume 3x3 filter

=> 5x5 output



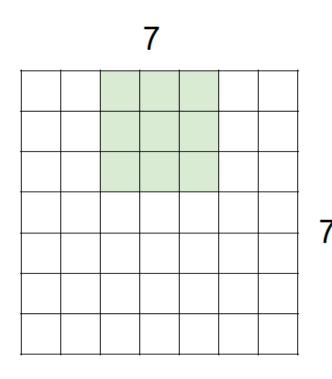
Case: Stride > 1



7x7 input (spatially) assume 3x3 filter applied with stride 2



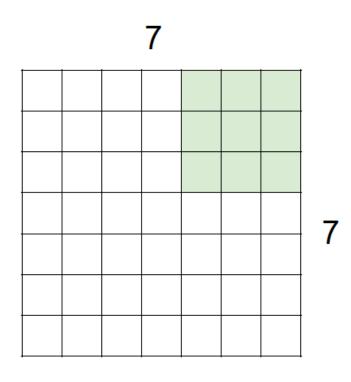
Case: Stride > 1



7x7 input (spatially) assume 3x3 filter applied with stride 2

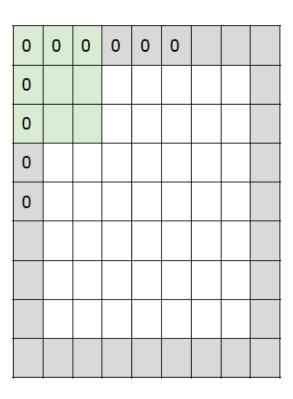


Case: Stride > 1



7x7 input (spatially) assume 3x3 filter applied with stride 2 => 3x3 output!

 Zero padding to handle non-integer cases or control the output sizes



e.g. input 7x7

3x3 filter, applied with stride 1

pad with 1 pixel border => what is the output?

7x7 output!

 Zero padding to handle non-integer cases or control the output sizes

0	0	0	0	0	0		
0							
0							
0							
0							

e.g. input 7x7

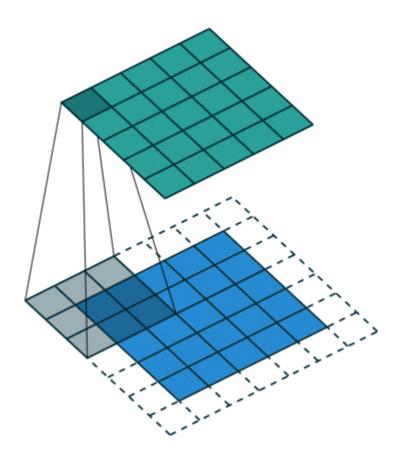
3x3 filter, applied with stride 1

pad with 1 pixel border => what is the output?

#### 7x7 output!

in general, common to see CONV layers with stride 1, filters of size FxF, and zero-padding with (F-1)/2. (will preserve size spatially)

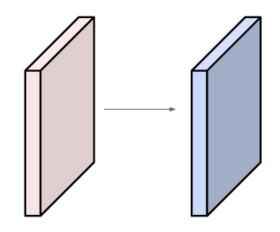
 Zero padding to handle non-integer cases or control the output sizes



### Examples time:

Input volume: 32x32x3

10 5x5 filters with stride 1, pad 2



Output volume size:

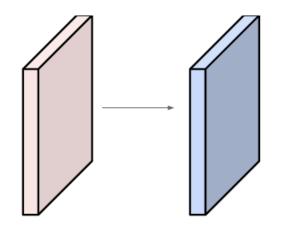
$$(32+2*2-5)/1+1 = 32$$
 spatially, so

32x32x10

### Examples time:

Input volume: 32x32x3

10 5x5 filters with stride 1, pad 2



Number of parameters in this layer? each filter has 5\*5\*3 + 1 = 76 params

(+1 for bias)



#### Summary

- Accepts a volume of size  $W_1 imes H_1 imes D_1$
- · Requires four hyperparameters:
  - Number of filters K.
  - their spatial extent F,
  - the stride S,
  - the amount of zero padding P.
- Produces a volume of size  $W_2 imes H_2 imes D_2$  where:
  - $W_2 = (W_1 F + 2P)/S + 1$
  - $\circ H_2 = (H_1 F + 2P)/S + 1$  (i.e. width and height are computed equally by symmetry)
  - $D_2 = K$
- With parameter sharing, it introduces  $F \cdot F \cdot D_1$  weights per filter, for a total of  $(F \cdot F \cdot D_1) \cdot K$  weights and K biases.
- In the output volume, the d-th depth slice (of size  $W_2 imes H_2$ ) is the result of performing a valid convolution of the d-th filter over the input volume with a stride of S, and then offset by d-th bias.



#### **Outline**

- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - □ Convolution layers & model complexity
  - Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

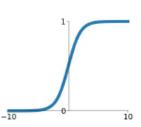
Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes

### Review: Activation Function

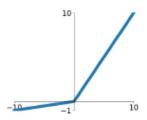
#### Zoo of Activation functions

#### **Sigmoid**

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

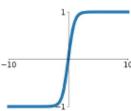


# Leaky ReLU $\max(0.1x, x)$



#### tanh

tanh(x)

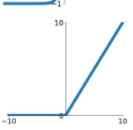


#### **Maxout**

 $\max(w_1^T x + b_1, w_2^T x + b_2)$ 

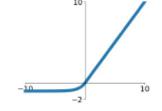
#### ReLU

 $\max(0, x)$ 

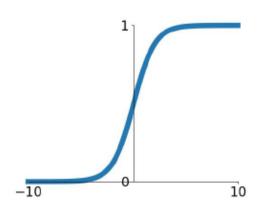


#### **ELU**

$$\begin{cases} x & x \ge 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Sigmoid function



**Sigmoid** 

$$\sigma(x)=1/(1+e^{-x})$$

- Squashes numbers to range [0,1]
- Historically popular since they have nice interpretation as a saturating "firing rate" of a neuron

#### 3 problems:

- Saturated neurons "kill" the gradients
- Sigmoid outputs are not zero-centered
- 3. exp() is a bit compute expensive

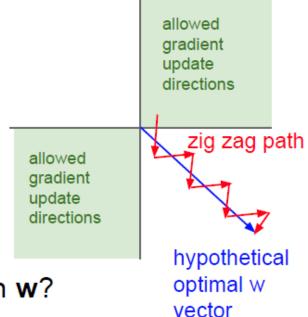


# Sigmoid function

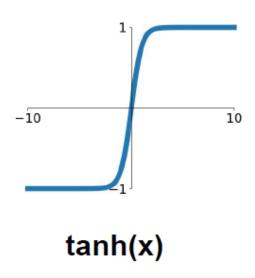
Consider what happens when the input to a neuron is always positive...

$$f\left(\sum_i w_i x_i + b
ight)$$

What can we say about the gradients on **w**? Always all positive or all negative :( (this is also why you want zero-mean data!)



### Tanh function



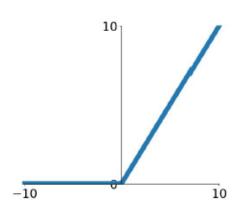
- Squashes numbers to range [-1,1]
- zero centered (nice)
- still kills gradients when saturated :(

[LeCun et al., 1991]

Recurrent neural networks: LSTM, GRU

# м

#### **Rectified Linear Unit**

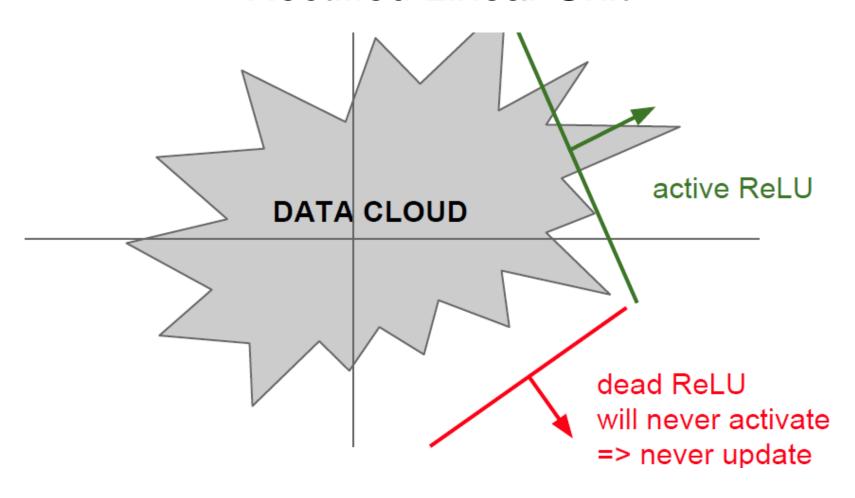


ReLU (Rectified Linear Unit)

- Computes f(x) = max(0,x)
- Does not saturate (in +region)
- Very computationally efficient
- Converges much faster than sigmoid/tanh in practice (e.g. 6x)
- Actually more biologically plausible than sigmoid
- Not zero-centered output
- An annoyance:

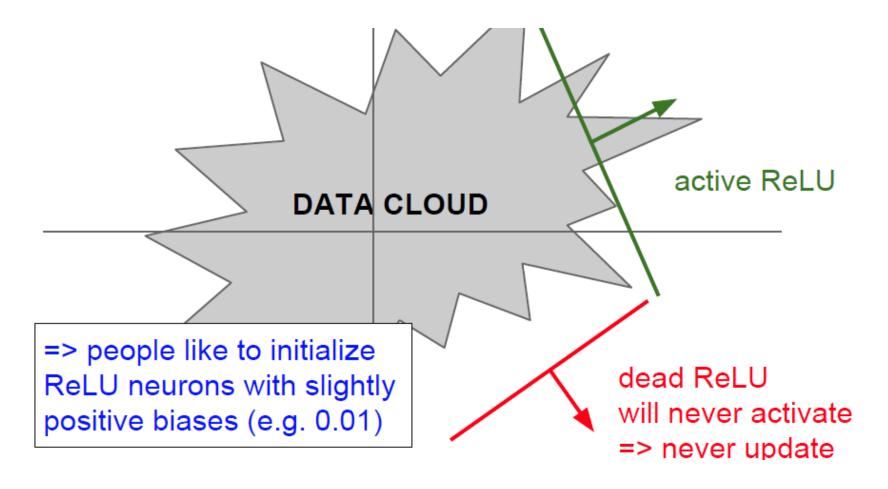
hint: what is the gradient when x < 0?

### **Rectified Linear Unit**





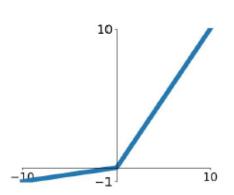
### **Rectified Linear Unit**





# Leaky ReLU

[Mass et al., 2013] [He et al., 2015]



- Does not saturate
- Computationally efficient
- Converges much faster than sigmoid/tanh in practice! (e.g. 6x)
- will not "die".

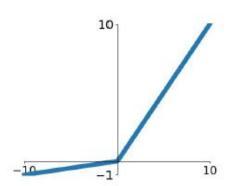
#### Leaky ReLU

$$f(x) = \max(0.01x, x)$$

# м

# Leaky ReLU





Leaky ReLU

$$f(x) = \max(0.01x, x)$$

- Does not saturate
- Computationally efficient
- Converges much faster than sigmoid/tanh in practice! (e.g. 6x)
- will not "die".

#### Parametric Rectifier (PReLU)

$$f(x) = \max(\alpha x, x)$$

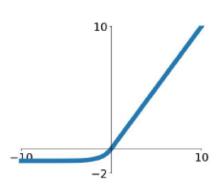
backprop into \alpha (parameter)



# Exponential Linear Units (ELU)

[Clevert et al., 2015]

#### Exponential Linear Units (ELU)

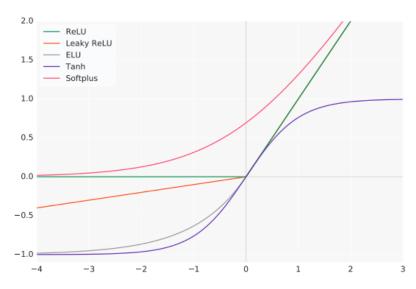


$$f(x) \, = \, \begin{cases} x & \text{if } x > 0 \\ \alpha \; (\exp(x) - 1) & \text{if } x \leq 0 \end{cases} \quad \text{- Computation requires exp()}$$

- All benefits of ReLU
- Closer to zero mean outputs
- Negative saturation regime compared with Leaky ReLU adds some robustness to noise

# Summary: Activation function

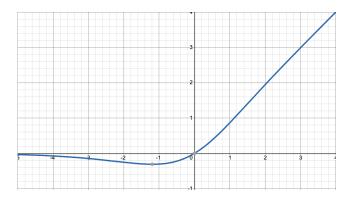
- For internal layers in CNNs
  - Use ReLU. Be careful with your learning rates
  - Try out Leaky ReLU / Maxout / ELU
  - Try out tanh but don't expect much
  - Don't use sigmoid
- For output layers
  - □ Task dependent
  - □ Related to your loss function



# Summary: Activation function

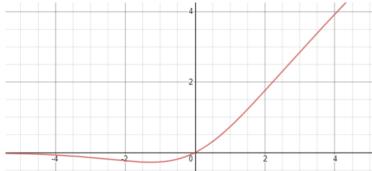
#### Recent progresses

$$\square$$
 Mish  $f(x) = x \cdot \tanh(\varsigma(x))$ ,  $\varsigma(x) = \ln(1 + e^x)$ .



□ Swish 
$$f(x) = x * (1 + \exp(-x))^{-1}$$

https://arxiv.org/abs/1908.08681



https://arxiv.org/abs/1710.05941



#### **Outline**

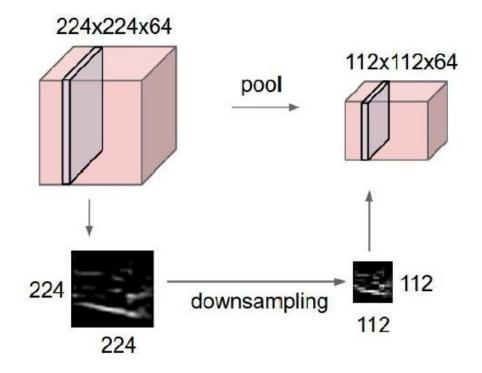
- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - □ Convolution layers & model complexity
  - □ Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes



# **Pooling Layers**

- Reducing the spatial size of the feature maps
  - □ Smaller representations
  - On each activation map independently
  - Low resolution means fewer details

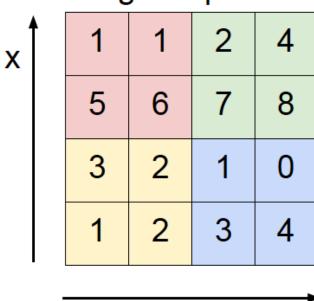




# **Pooling Layers**

Example: max pooling

#### Single depth slice



max pool with 2x2 filters and stride 2

6	8
3	4



# Complexity of Pooling Layers

#### Summary

- Accepts a volume of size  $W_1 imes H_1 imes D_1$
- Requires three hyperparameters:
  - their spatial extent F,
  - the stride S.
- Produces a volume of size  $W_2 imes H_2 imes D_2$  where:

$$W_2 = (W_1 - F)/S + 1$$

$$H_2 = (H_1 - F)/S + 1$$

$$O$$
  $D_2 = D_1$ 

- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

# Outline

- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - □ Convolution layers & model complexity
  - □ Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes



- What representations a CNN can capture in general?
- lacktriangle Consider a representation  $\phi$  as an abstract function

$$\phi: \mathbf{x} \to \phi(\mathbf{x}) \in \mathbb{R}^d$$

- We want to look at how the representation changes upon transformations of input image.
  - Transformations represent the potential variations in the natural images
  - □ Translation, scale change, rotation, local deformation etc.



- Two key properties of representations
  - □ Equivariance

A representation  $\phi$  is equivariant with a transformation g if the transformation can be transferred to the representation output.

$$\exists$$
 a map  $M_g : \mathbb{R}^d \to \mathbb{R}^d$  such that:  $\forall \mathbf{x} \in \mathcal{X} : \phi(g\mathbf{x}) \approx M_g \phi(\mathbf{x})$ 

□ Example: convolution w.r.t. translation



- Two key properties of representations
  - □ Invariance

A representation  $\phi$  is invariant with a transformation g if the transformation has no effect on the representation output.

$$\forall \mathbf{x} \in \mathcal{X} : \phi(g\mathbf{x}) \approx \phi(\mathbf{x})$$

Example: convolution+pooling+FC w.r.t. translation



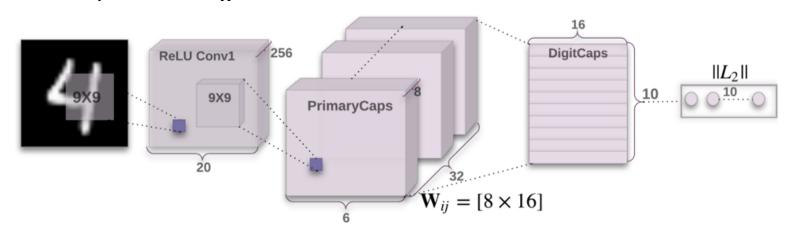


- Recent results on convolution layers
  - Convolutions are equivariant to translation
  - Convolutions are not equivariant to other isometries of the sampling lattice, e.g., rotation



- □ What if a CNN learns rotated copies of the same filter?
  - The stack of feature maps is equivariant to rotation.

- Recent results on convolution layers
  - □ Ordinary CNNs can be generalized to Group Equivariant
     Networks (Cohen and Welling ICML'16, Kondor and Trivedi ICML'18)
    - Redefining the convolution and pooling operations
    - Equivariant to more general transformation from some group G
  - Replacing pooling by other network designs
    - Capsule network (Sabour et al, 2017) https://arxiv.org/abs/1710.09829



# Outline

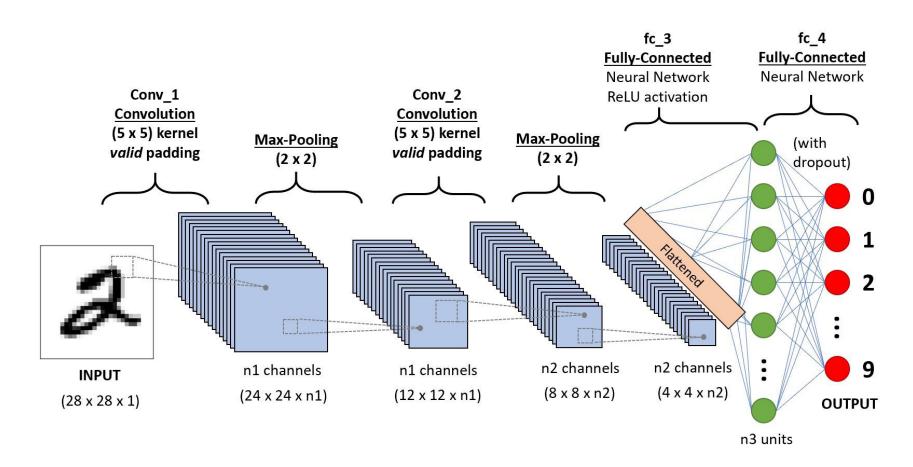
- Why Convolutional Neural Network (CNN)?
  - Motivation and overview
- What is the CNN?
  - □ Convolution layers & model complexity
  - □ Closer look at activation functions
  - □ Pooling layers & model complexity
  - Math properties
- Examples of CNNs

Acknowledgement: Roger Grosse @UofT & Feifei Li's cs231n notes



#### LeNet-5

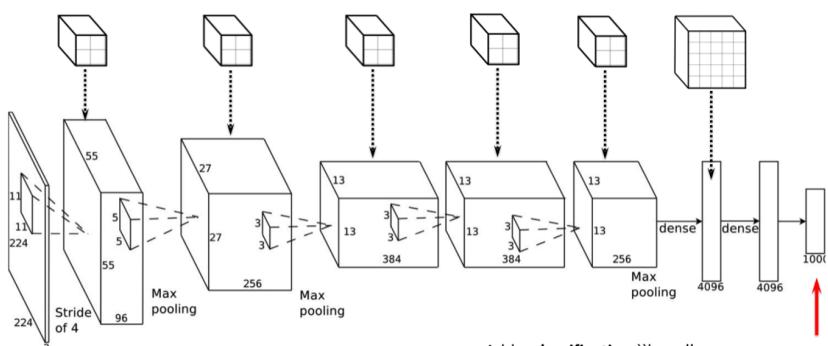
#### Handwritten digit recognition





#### **AlexNet**

#### Deeper network structure



Add a classification ``layer".

For an input image, the value in a particular dimension of this vector tells you the probability of the corresponding object class.



# Summary of CNNs

- CNN properties [Bronstein et al., 2018]
  - Convolutional (Translation invariance)
  - Scale Separation (Compositionality)
  - ☐ Filters localized in space (Deformation Stability)
  - □ O(1) parameters per filter (independent of input image size n)
  - □ O(n) complexity per layer (filtering done in the spatial domain)
  - □ O(log n) layers in classification tasks
- Next time ...
  - ☐ Structure design of Modern CNNs
- Reference
  - □ CS231n course notes <a href="http://cs231n.github.io/convolutional-networks/">http://cs231n.github.io/convolutional-networks/</a>
  - □ D2L Chapter 6 + DLBook Chapter 9