Lecture 15: Deep Generative Models I: Latent Variable Models & EM, AutoEncoder

Lan Xu

SIST, ShanghaiTech

Fall, 2021

Outline

- **Unsupervised learning**

    - Problem setup

- **Latent variable model**

    - EM algorithm and GMM

- **Representation learning**

    - AutoEncoder

Acknowledgement: Yingyu Liang@Princeton's & Feifei Li's cs231n notes

Unsupervised learning

■ Task formulation

Unsupervised Learning

Data: x
Just data, no labels!

Goal: Learn some underlying
hidden *structure* of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

# Unsupervised learning

- Task formulation

Unsupervised Learning

Data: x
Just data, no labels!

Goal: Learn some underlying
hidden *structure* of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

K-means clustering
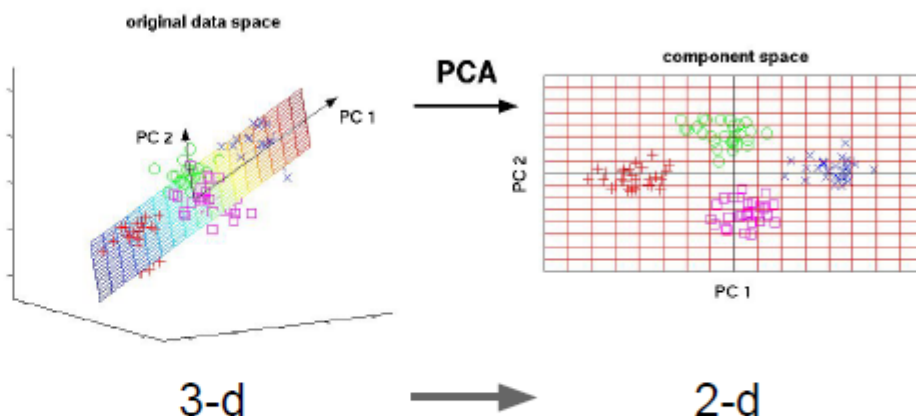
Unsupervised learning

- ■ Task formulation

Unsupervised Learning

Data: x
Just data, no labels!

Goal: Learn some underlying
hidden *structure* of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

original data space

PCA →

component space

3-d → 2-d

Principal Component Analysis
(Dimensionality reduction)

Unsupervised learning
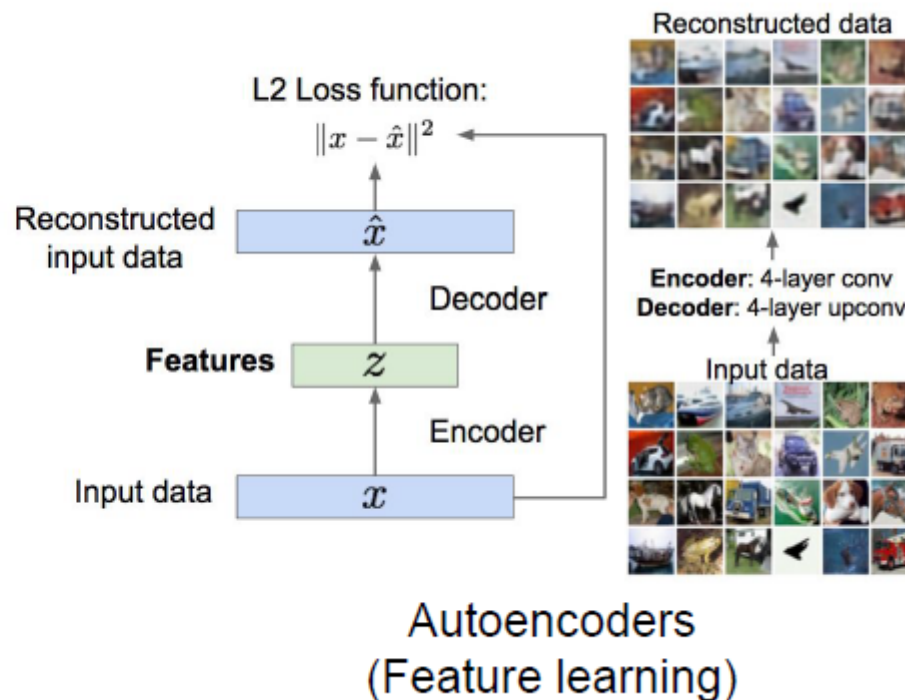
- ## Task formulation

Unsupervised Learning

Data: x
Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

L2 Loss function:
$$\|x - \hat{x}\|^2$$

Reconstructed input data — $\hat{x}$

Decoder

Features — z

Encoder

Input data — x

Reconstructed data

Encoder: 4-layer conv
Decoder: 4-layer upconv

Input data

Autoencoders
(Feature learning)

Unsupervised learning

- ■ Task formulation

Unsupervised Learning

Data: x
Just data, no labels!

Goal: Learn some underlying
hidden *structure* of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

Figure copyright Ian Goodfellow, 2016. Reproduced with permission.

1-d density estimation

2-d density estimation

Outline

- Unsupervised learning

    - Problem setup

- Latent variable model

    - EM algorithm and GMM

- Representation learning

    - AutoEncoder

Acknowledgement: Yingyu Liang@Princeton's & Feifei Li's cs231n notes

Latent variable model

- **Data generation process**
  - ☐ Latent variable z $p(z) = \text{something simple}$

  - ☐ A mapping from the latent space to observation x

$$p(x) = \int p(x, z) \; dz \quad \text{where} \quad p(x, z) = p(x \mid z)p(z)$$

For example, a Gaussian mixture model

$$p_\theta(x) = \sum_{k=1}^{K} p_\theta(z = k) p_\theta(x|z = k)$$

Review: Gaussian Mixture Model

- **Definition**

 - A Gaussian mixture model represents a **distribution** as

  $$p(\mathbf{x}) = \sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)$$

  with π_k the mixing coefficients, where:

  $$\sum_{k=1}^{K} \pi_k = 1 \quad \text{and} \quad \pi_k \geq 0 \quad \forall k$$

  - □ GMM is a density estimator

  - □ GMMs are universal approximators of densities (if you have enough Gaussians).

Review: 2D GMM example

Fitting GMMs: Maximum Likelihood

- Maximum likelihood maximizes

$$\ln p(\mathbf{X}|\pi, \mu, \Sigma) = \sum_{n=1}^{N} \ln \left(\sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k) \right)$$

w.r.t $\Theta = \{\pi_k, \mu_k, \Sigma_k\}$

- Problems:
  - ▶ Singularities: Arbitrarily large likelihood when a Gaussian explains a single point
  - ▶ Identifiability: Solution is up to permutations

- How would you optimize this?

- Can we have a closed form update?

- Don't forget to satisfy the constraints on π_k

Fitting GMMs as an LVM

■ **A latent variable model formulation**

- Introduce a hidden variable such that its knowledge would simplify the maximization

- We could introduce a hidden (latent) variable z which would represent which Gaussian generated our observation x, with some probability

- Let $z \sim \text{Categorical}(\pi)$ (where $\pi_k \geq 0$, $\sum_k \pi_k = 1$)

- Then:

$$p(\mathbf{x}) = \sum_{k=1}^{K} p(\mathbf{x}, z = k)$$

$$= \sum_{k=1}^{K} \underbrace{p(z = k)}_{\pi_k} \underbrace{p(\mathbf{x}|z = k)}_{\mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)}$$

Fitting GMMs

- A Gaussian mixture distribution:

$$p(\mathbf{x}) = \sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)$$

- We had: $z \sim \text{Categorical}(\pi)$ (where $\pi_k \geq 0$, $\sum_k \pi_k = 1$)
- Joint distribution: $p(\mathbf{x}, \mathbf{z}) = p(\mathbf{z})p(\mathbf{x}|\mathbf{z})$
- Log-likelihood:

$$\ell(\pi, \mu, \Sigma) = \ln p(\mathbf{X}|\pi, \mu, \Sigma) = \sum_{n=1}^{N} \ln p(\mathbf{x}^{(n)}|\pi, \mu, \Sigma)$$

$$= \sum_{n=1}^{N} \ln \sum_{z^{(n)}=1}^{K} p(\mathbf{x}^{(n)}|z^{(n)}; \mu, \Sigma)p(z^{(n)}|\pi)$$

- Note: We have a hidden variable $z^{(n)}$ for every observation
- General problem: sum inside the log
- How can we optimize this?

Expectation Maximization

- Optimization uses the Expectation Maximization algorithm, which alternates between two steps:

  1. E-step: Compute the posterior probability that each Gaussian generates each datapoint (as this is unknown to us)
  2. M-step: Assuming that the data really was generated this way, change the parameters of each Gaussian to maximize the probability that it would generate the data it is currently responsible for.

Expectation Maximization
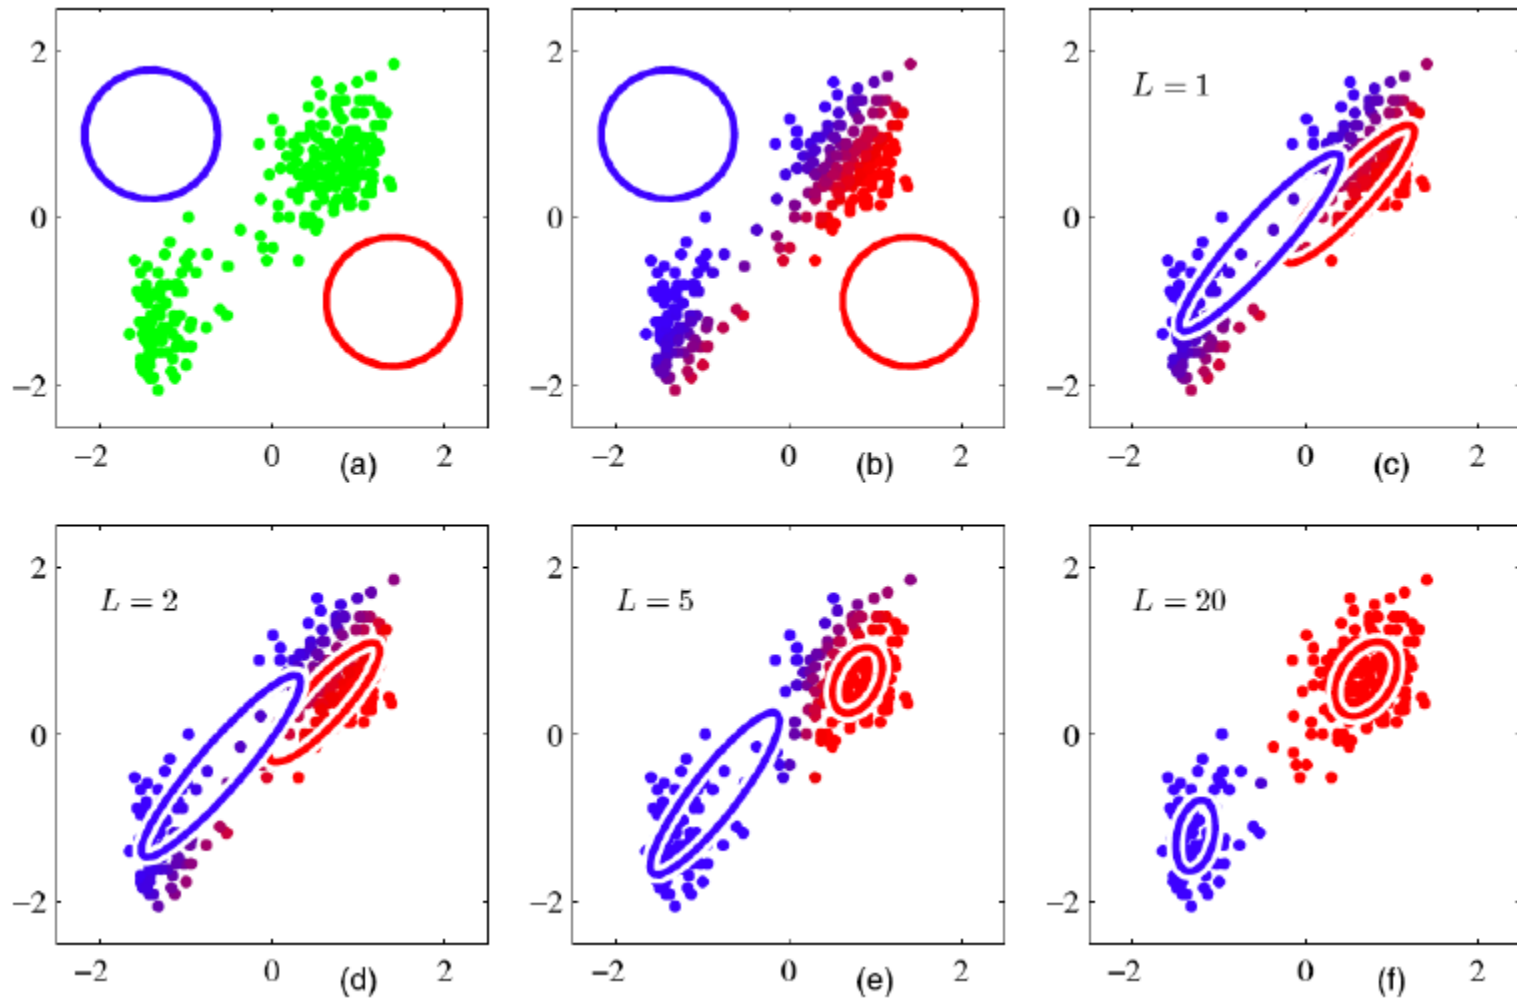
Expectation Maximization

- Elegant and powerful method for finding maximum likelihood solutions for models with latent variables

    1. E-step:
        - In order to adjust the parameters, we must first solve the inference problem: Which Gaussian generated each datapoint?
        - We cannot be sure, so it's a distribution over all possibilities.

        $$\gamma_k^{(n)} = p(z^{(n)} = k | \mathbf{x}^{(n)}; \pi, \mu, \Sigma)$$

    2. M-step:
        - Each Gaussian gets a certain amount of posterior probability for each datapoint.
        - At the optimum we shall satisfy

        $$\frac{\partial \ln p(\mathbf{X}|\pi, \mu, \Sigma)}{\partial \Theta} = 0$$

        - We can derive closed form updates for all parameters

GMM: E-Step

- Conditional probability (using Bayes rule) of **z** given **x**

$$
\gamma_k = p(z = k|\mathbf{x}) \; = \; \frac{p(z = k)p(\mathbf{x}|z = k)}{p(\mathbf{x})}
$$

$$
= \; \frac{p(z = k)p(\mathbf{x}|z = k)}{\sum_{j=1}^{K} p(z = j)p(\mathbf{x}|z = j)}
$$

$$
= \; \frac{\pi_k \mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)}{\sum_{j=1}^{K} \pi_j \mathcal{N}(\mathbf{x}|\mu_j, \Sigma_j)}
$$

- γ_k can be viewed as the responsibility

GMM: M-Step

- Log-likelihood:

$$\ln p(\mathbf{X}|\pi, \mu, \Sigma) = \sum_{n=1}^{N} \ln \left(\sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k) \right)$$

- Set derivatives to 0:

$$\frac{\partial \ln p(\mathbf{X}|\pi, \mu, \Sigma)}{\partial \mu_k} = 0 = \sum_{n=1}^{N} \frac{\pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k)}{\sum_{j=1}^{K} \pi_j \mathcal{N}(\mathbf{x}|\mu_j, \Sigma_j)} \Sigma_k(\mathbf{x}^{(n)} - \mu_k)$$

- We used:

$$\mathcal{N}(\mathbf{x}|\mu, \Sigma) = \frac{1}{\sqrt{(2\pi)^d |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu) \right)$$

and:

$$\frac{\partial(\mathbf{x}^T A \mathbf{x})}{\partial \mathbf{x}} = \mathbf{x}^T (A + A^T)$$

GMM: M-Step

$$\frac{\partial \ln p(\mathbf{X}|\pi, \mu, \Sigma)}{\partial \mu_k} = 0 = \sum_{n=1}^{N} \underbrace{\frac{\pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k)}{\sum_{j=1}^{K} \pi_j \mathcal{N}(\mathbf{x}|\mu_j, \Sigma_j)}}_{\gamma_k^{(n)}} \Sigma_k(\mathbf{x}^{(n)} - \mu_k)$$

- This gives

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)} \mathbf{x}^{(n)}$$

with N_k the effective number of points in cluster k

$$N_k = \sum_{n=1}^{N} \gamma_k^{(n)}$$

- We just take the center-of gravity of the data that the Gaussian is responsible for

- Just like in K-means, except the data is weighted by the posterior probability of the Gaussian.

- Guaranteed to lie in the convex hull of the data (Could be big initial jump)

GMM: M-Step

- We can get similarly expression for the variance

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)}(\mathbf{x}^{(n)} - \mu_k)(\mathbf{x}^{(n)} - \mu_k)^T$$

- We can also minimize w.r.t the mixing coefficients

$$\pi_k = \frac{N_k}{N}, \quad \text{with} \quad N_k = \sum_{n=1}^{N} \gamma_k^{(n)}$$

- The optimal mixing proportion to use (given these posterior probabilities) is just the fraction of the data that the Gaussian gets responsibility for.

- Note that this is not a closed form solution of the parameters, as they depend on the responsibilities $\gamma_k^{(n)}$, which are complex functions of the parameters

- But we have a simple iterative scheme to optimize

Summary: EM Algorithm for GMM

- Initialize the means μ_k, covariances Σ_k and mixing coefficients π_k

- Iterate until convergence:

  - ▶ E-step: Evaluate the responsibilities given current parameters

  $$\gamma_k^{(n)} = p(z^{(n)}|\mathbf{x}) = \frac{\pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k)}{\sum_{j=1}^{K} \pi_j \mathcal{N}(\mathbf{x}^{(n)}|\mu_j, \Sigma_j)}$$

  - ▶ M-step: Re-estimate the parameters given current responsibilities

  $$\mu_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)} \mathbf{x}^{(n)}$$

  $$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)} (\mathbf{x}^{(n)} - \mu_k)(\mathbf{x}^{(n)} - \mu_k)^T$$

  $$\pi_k = \frac{N_k}{N} \quad \text{with} \quad N_k = \sum_{n=1}^{N} \gamma_k^{(n)}$$

  - ▶ Evaluate log likelihood and check for convergence

  $$\ln p(\mathbf{X}|\pi, \mu, \Sigma) = \sum_{n=1}^{N} \ln \left(\sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k) \right)$$

An Alternative View of EM

- Hard to maximize (log-)likelihood of data directly

- General problem: sum inside the log

$$\ln p(\mathbf{x}|\Theta) = \ln \sum_z p(\mathbf{x}, \mathbf{z}|\Theta)$$

- Complete data $\{\mathbf{x}, \mathbf{z}\}$, and $\mathbf{x}$ is the incomplete data

- If we knew z, then easy to maximize (replace sum over k with just the k where $z = k$)

- Unfortunately we are not given the complete data, but only the incomplete.

An Alternative View of EM

- Our knowledge about the latent variables is $p(\mathbf{Z}|\mathbf{X}, \Theta^{old})$

- In the E-step we compute $p(\mathbf{Z}|\mathbf{X}, \Theta^{old})$

- In the M-step we maximize w.r.t Θ

$$Q(\Theta, \Theta^{old}) = \sum_{\mathbf{Z}} p(\mathbf{Z}|\mathbf{X}, \Theta^{old}) \ln p(\mathbf{X}, \mathbf{Z}|\Theta)$$

▶ E-step: Evaluate the responsibilities given current parameters

$$\gamma_k^{(n)} = p(z^{(n)}|\mathbf{x}) = \frac{\pi_k \mathcal{N}(\mathbf{x}^{(n)}|\mu_k, \Sigma_k)}{\sum_{j=1}^{K} \pi_j \mathcal{N}(\mathbf{x}^{(n)}|\mu_j, \Sigma_j)}$$

▶ M-step: Re-estimate the parameters given current responsibilities

$$\max \sum_{n=1}^{N} \sum_{z^{(n)}=1}^{K} p(z^{(n)}|x^{(n)}) \ln[p(x^{(n)}|z^{(n)}; \mu_k, \Sigma_k) p(z^{(n)}|\pi_k)]$$

$$\Leftrightarrow \quad \max \sum_{n=1}^{N} \sum_{z^{(n)}=1}^{K} \gamma_k^{(n)} \ln[\pi_k \mathcal{N}(x^{(n)}|\mu_k, \Sigma_k)]$$

$$\Leftrightarrow \quad \max_{\pi_k, \mu_k, \Sigma_k} \sum_{n=1}^{N} \gamma_k^{(n)} \ln[\pi_k \mathcal{N}(x^{(n)}|\mu_k, \Sigma_k)] \quad \Rightarrow$$

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)} \mathbf{x}^{(n)}$$

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^{N} \gamma_k^{(n)} (\mathbf{x}^{(n)} - \mu_k)(\mathbf{x}^{(n)} - \mu_k)^T$$

$$\pi_k = \frac{N_k}{N} \quad \text{with} \quad N_k = \sum_{n=1}^{N} \gamma_k^{(n)}$$

General EM Algorithm

1. Initialize Θ^{old}

2. E-step: Evaluate $p(\mathbf{Z}|\mathbf{X}, \Theta^{old})$

3. M-step:

$$\Theta^{new} = \arg\max_{\Theta} Q(\Theta, \Theta^{old})$$

where

$$Q(\Theta, \Theta^{old}) = \sum_{z} p(\mathbf{Z}|\mathbf{X}, \Theta^{old}) \ln p(\mathbf{X}, \mathbf{Z}|\Theta)$$

4. Evaluate log likelihood and check for convergence (or the parameters). If not converged, $\Theta^{old} = \Theta$, Go to step 2

Why it works?

- Updating each Gaussian definitely improves the probability of generating the data if we generate it from the same Gaussians after the parameter updates.

  ▶ But we know that the posterior will change after updating the parameters.

- A good way to show that this is OK is to show that there is a single function that is improved by both the E-step and the M-step.

  ▶ The function we need is called Free Energy.

Review: Jensen's Inequality

- For concave function f, we have

$$f(\mathbb{E}[X]) \geq \mathbb{E}[f(X)]$$

Figure: Jensen's Inequality

A simple latent variable model

- We assume that the data is generated i.i.d as

$$z \sim p(z) \qquad x \sim p(x|z)$$

- □ z is latent/hidden and x is observed
- Bounding the marginal likelihood

$$\log p(x) = \log \int_z p(x, z) \text{ (Multiply and divide by } q(z))$$

$$= \log \int_z \frac{q(z) p(x, z)}{q(z)} = \log \mathbb{E}_{z \sim q(z)} \left[\frac{p(x, z)}{q(z)} \right] \text{ (By Jensen's Inequality)}$$

$$\geq \int_z q(z) \log \frac{p(x, z)}{q(z)} = \mathcal{L}(x; \theta, \phi)$$

$$= \underbrace{\mathbb{E}_{q(z)}[\log p(x, z)]}_{\text{Expectation of Joint distribution}} + \underbrace{\mathrm{H}(q(z))}_{\text{Entropy}}$$

Evidence Lower Bound (ELBO)

- When is the lower bound tight?
- Let's look at: objective function – lower bound

$$\log p(x; \theta) - \mathcal{L}(x; \theta, \phi)$$

$$\log p(x) - \int_z q(z) \log \frac{p(x, z)}{q(z)}$$

$$= \int_z q(z) \log p(x) - \int_z q(z) \log \frac{p(x, z)}{q(z)}$$

$$= \int_z q(z) \log \frac{q(z) p(x)}{p(x, z)}$$

$$= \mathrm{KL}(q(z; \phi) \| p(z|x))$$

Visualization of ELBO

$$\mathcal{L}(q, \Theta) \leq \ln p(\mathbf{X}|\Theta)$$

$$\text{KL}(q||p)$$

$$\mathcal{L}(q, \boldsymbol{\theta}) \qquad \ln p(\mathbf{X}|\boldsymbol{\theta})$$

Key Point

The optimal $q(z; \phi)$ corresponds to the one that realizes
$$\text{KL}(q(z; \phi)||p(z|x)) = 0 \iff q(z; \phi) = p(z|x)$$

E-step and M-step

$$\ln p(\mathbf{X}|\Theta) = \mathcal{L}(q, \Theta) + KL(q||p(\mathbf{Z}|\mathbf{X}, \Theta))$$

- In the E-step we maximize w.r.t $q(\mathbf{Z})$ the lower bound $\mathcal{L}(q, \Theta)$
- Since $\ln p(\mathbf{X}|\theta)$ does not depend on $q(\mathbf{Z})$, the maximum $\mathcal{L}$ is obtained when the KL is 0
- This is achieved when $q(\mathbf{Z}) = p(\mathbf{Z}|\mathbf{X}, \Theta)$
- The lower bound $\mathcal{L}$ is then

$$
\begin{aligned}
\mathcal{L}(q, \Theta) &= \sum_{\mathbf{Z}} p(\mathbf{Z}|\mathbf{X}, \Theta^{old}) \ln p(\mathbf{X}, \mathbf{Z}|\Theta) - \sum_{\mathbf{Z}} p(\mathbf{Z}|\mathbf{X}, \Theta^{old}) \ln p(\mathbf{Z}|\mathbf{X}, \Theta^{old}) \\
&= Q(\Theta, \Theta^{old}) + \text{const}
\end{aligned}
$$

with the content the entropy of the q distribution, which is independent of Θ

- In the M-step the quantity to be maximized is the expectation of the complete data log-likelihood
- Note that Θ is only inside the logarithm and optimizing the complete data likelihood is easier

Visualization of E-step

$\text{KL}(q||p) = 0$

$\mathcal{L}(q, \boldsymbol{\theta}^{\text{old}})$ $\ln p(\mathbf{X}|\boldsymbol{\theta}^{\text{old}})$

- The q distribution equal to the posterior distribution for the current parameter values Θ^{old}, causing the lower bound to move up to the same value as the log likelihood function, with the KL divergence vanishing.

Visualization of M-step

- The distribution $q(\mathbf{Z})$ is held fixed and the lower bound $\mathcal{L}(q, \Theta)$ is maximized with respect to the parameter vector Θ to give a revised value Θ^{new}. Because the KL divergence is nonnegative, this causes the log likelihood $\ln p(\mathbf{X}|\Theta)$ to increase by at least as much as the lower bound does.

Visualization of the EM Algorithm

- The EM algorithm involves alternately computing a lower bound on the log likelihood for the current parameter values and then maximizing this bound to obtain the new parameter values.

Summary of EM algorithms

■ **EM is coordinate ascent in ELBO**

$$\int_z q(z) \log \frac{p(x,z)}{q(z)} = \mathcal{L}(x; \theta, \phi)$$

$$= \underbrace{\mathbb{E}_{q(z)}[\log p(x,z)]}_{\text{Expectation of Joint distribution}} + \underbrace{\mathrm{H}(q(z))}_{\text{Entropy}}$$

■ **Or coordinate descent in** Free Energy

$$\mathcal{F} = -\mathcal{L}(x; \theta, q) = \underbrace{E_q[-\log p(x,z)]}_{\text{Expected energy}} - \underbrace{H(q(z))}_{\text{Entropy}}$$

- The E-step minimizes F by finding the best distribution over hidden configurations for each data point.

- The M-step holds the distribution fixed and minimizes F by changing the parameters that determine the energy of a configuration.

Recall EM GMM

■ MLE: maximizing the log-likelihood

$$\ell(\theta) = \sum_{i=1}^{n} \log p(x^{(i)}; \theta)$$

■ ELBO: Evidence Lower Bound

$$
\begin{aligned}
\log p(\mathbf{x}) &= \log \int_z p(\mathbf{x}, z) \\
&= \log \int_z p(\mathbf{x}, z) \frac{q(\mathbf{z})}{q(\mathbf{z})} \\
&= \log(E_q[\frac{p(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})}]) \\
&\geq E_q[\log p(\mathbf{x}, \mathbf{z})] - E_q[\log q(\mathbf{z})] \\
&= ELBO
\end{aligned}
$$

$$
\begin{aligned}
KL(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) &= E_q[\log q(\mathbf{z})] - E_q[\log p(\mathbf{z}|\mathbf{x})] \\
&= E_q[\log q(\mathbf{z})] - E_q[\log p(\mathbf{z}, \mathbf{x})] + \log p(\mathbf{x}) \\
&= \log p(\mathbf{x}) - (E_q[\log p(\mathbf{z}, \mathbf{x})] - E_q[\log q(\mathbf{z})]) \\
&= \log p(\mathbf{x}) - ELBO
\end{aligned}
$$

$$
\begin{aligned}
EBLO &= E_q[\log p(\mathbf{x}, \mathbf{z})] - E_q[\log q(\mathbf{z})] \\
&= E_q[\log \frac{p(\mathbf{x}, \mathbf{z})}{p(\mathbf{z})}] - E_q[\log \frac{q(\mathbf{z})}{p(\mathbf{z})}] \\
&= E_q[\log p(\mathbf{x}|\mathbf{z})] - KL(q(\mathbf{z})||p(\mathbf{z}))
\end{aligned}
$$

Outline

- Unsupervised learning

  - Problem setup

- Latent variable model

  - EM algorithm and GMM

- **Representation learning**

  - AutoEncoder

Acknowledgement: Yingyu Liang@Princeton's & Feifei Li's cs231n notes

Autoencoder

- ## Feature representation learning

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

Autoencoder

- ## Feature representation learning

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

Originally: Linear + nonlinearity (sigmoid)
Later: Deep, fully-connected
Later: ReLU CNN

Features z

Input data x

Encoder

Autoencoder

- Feature representation learning

How to learn this feature representation?


Features z

Encoder

Input data x

Autoencoder

- **Feature representation learning**

How to learn this feature representation?
Train such that features can be used to reconstruct original data
"Autoencoding" - encoding itself

Reconstructed input data $\hat{x}$

Decoder

Features z

Encoder

Input data x

Reconstructed data

Encoder: 4-layer conv
Decoder: 4-layer upconv

Input data

Autoencoder

- Feature representation learning

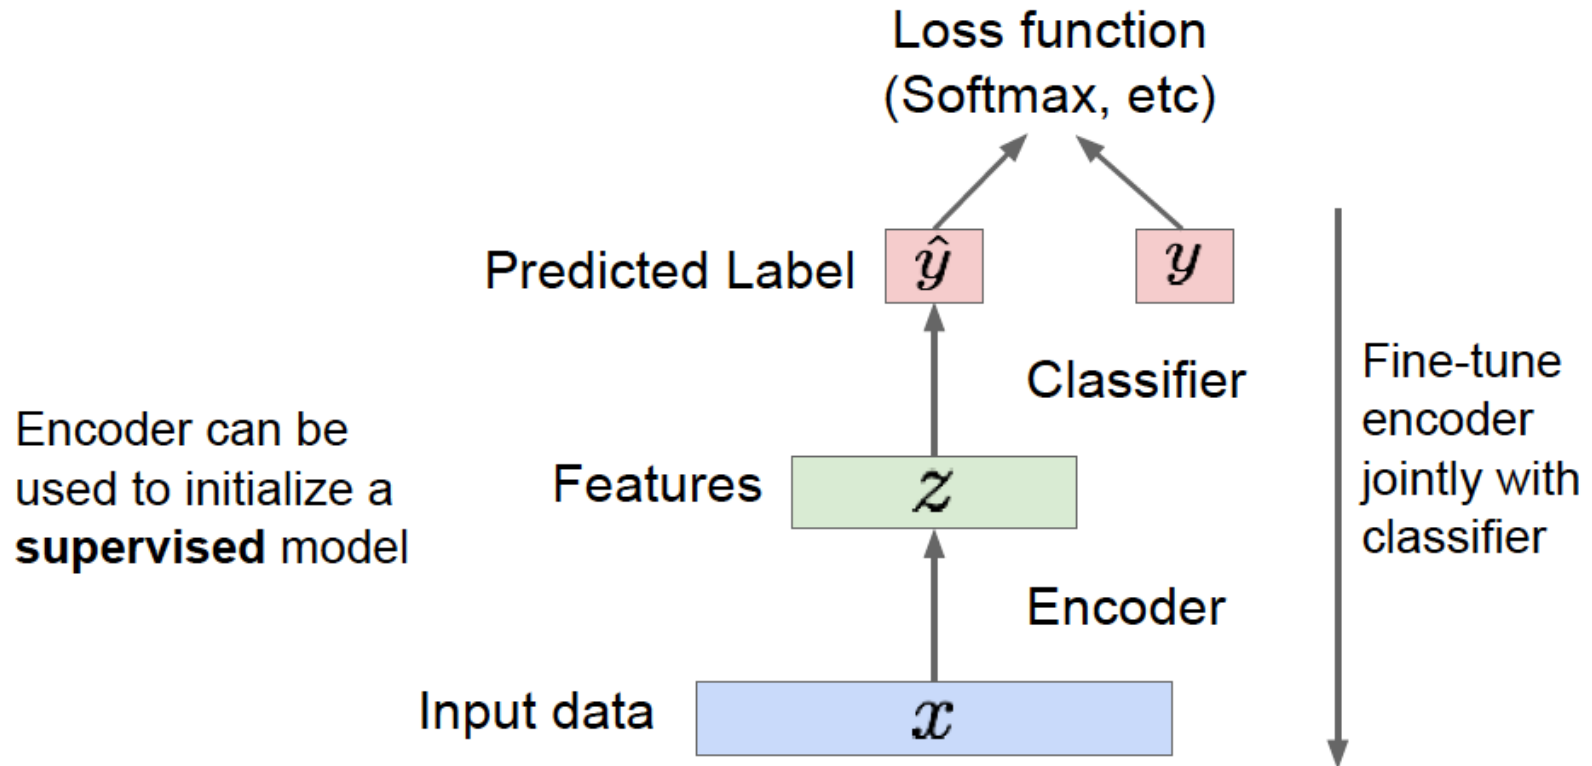Train such that features can be used to reconstruct original data

L2 Loss function:

$$\|x - \hat{x}\|^2$$

Reconstructed input data: $\hat{x}$

Decoder

Features: z

Encoder

Input data: x

Autoencoder

- Feature representation learning

Reconstructed input data — $\hat{x}$

Decoder

After training, throw away decoder

Features — z

Encoder

Input data — x

Autoencoder

- Feature representation learning

Loss function
(Softmax, etc)

Predicted Label $\hat{y}$ y

Classifier

Encoder can be used to initialize a **supervised** model

Features z

Fine-tune encoder jointly with classifier

Encoder

Input data x

Autoencoder

- Linear hidden layer example

$$\hat{X}$$

$$W^T$$

$$W$$

$$X$$

- A single hidden unit
- Hidden unit has *linear* activation
- What will this learn?

Lan Xu – CS 280 Deep Learning

Autoencoder

- **Linear hidden layer example**

Training: Learning W by minimizing L2 divergence

$$\hat{x} = w^T w x$$
$$div(\hat{x}, x) = \|x - \hat{x}\|^2 = \|x - w^T w x\|^2$$

$$\widehat{W} = \underset{W}{\operatorname{argmin}} \, E[div(\hat{x}, x)]$$
$$\widehat{W} = \underset{W}{\operatorname{argmin}} \, E[\|x - w^T w x\|^2]$$

- This is just PCA!

Autoencoder

■ Nonlinear hidden layer

- With non-linearity
  - "Non linear" PCA
  - Deeper networks can capture more complicated manifolds
    - "Deep" autoencoders

Autoencoder

- **Binary input example**

Decoder

$$\widehat{\mathbf{x}} = o(\widehat{\mathbf{a}}(\mathbf{x}))$$
$$= \underbrace{\text{sigm}}(\mathbf{c} + \mathbf{W}^*\mathbf{h}(\mathbf{x}))$$
for binary inputs

Encoder

$$\mathbf{h}(\mathbf{x}) = g(\mathbf{a}(\mathbf{x}))$$
$$= \text{sigm}(\mathbf{b} + \mathbf{W}\mathbf{x})$$

$$l(f(\mathbf{x})) = -\sum_k \left(x_k \log(\widehat{x}_k) + (1 - x_k)\log(1 - \widehat{x}_k) \right)$$

- cross-entropy (more precisely: sum of Bernoulli cross-entropies)

Regularization

■ Regularized autoencoders: add regularization term that encourages the model to have other properties

☐ Sparsity of the representation (sparse autoencoder)

☐ Robustness to noise or to the missing inputs (denoising autoencoder)

☐ Smallness of the derivative of the representation (contracitve autoencoder)

Regularization

- **Undercomplete representation**


- Hidden layer is undercomplete if smaller than the input layer
  - hidden layer "compresses" the input
  - will compress well only for the training distribution
- Hidden units will be
  - good features for the training distribution
  - but bad for other types of input

$\hat{\mathbf{x}}$

c_k

$\mathbf{W}^* = \mathbf{W}^\top$

(tied weights)

$\mathbf{h}(\mathbf{x})$

b_j

$\mathbf{W}$

$\mathbf{x}$

Regularization

$$L_R = L(x, g(f(x))) + R(h)$$

- Sparse autoencoder

  ☐ Constrain the code to have sparsity

  ☐ Training: minimize a loss function

$$L_R = L(x, g(f(x))) + \lambda|h|_1$$

Regularization

- ## Denoising autoencoder

- Idea: representation should be robust to introduction of noise:
  - random assignment of subset of inputs to 0, with probability ν
  - Gaussian additive noise

- Reconstruction $\widehat{\mathbf{x}}$ computed from the corrupted input $\widetilde{\mathbf{x}}$

- Loss function compares $\widehat{\mathbf{x}}$ reconstruction with the **noiseless input** $\mathbf{x}$

$\widehat{\mathbf{x}}$ ◯ c_k ◯◯◯◯◯

$\mathbf{W}^* = \mathbf{W}^{\top}$
(tied weights)

$\mathbf{h}(\widetilde{\mathbf{x}})$ ◯◯◯ b_j ◯◯◯◯

$\mathbf{W}$

$\widetilde{\mathbf{x}}$ ⓪◯◯⓪◯⓪

noise process
$p(\widetilde{\mathbf{x}}|\mathbf{x})$

$\mathbf{x}$ ◯◯◯◯◯◯

Regularization

- Denoising autoencoder

$$\widehat{\mathbf{x}} = \mathrm{sigm}(\mathbf{c} + \mathbf{W}^*\mathbf{h}(\widetilde{\mathbf{x}}))$$

$$p(\widetilde{\mathbf{x}}|\mathbf{x})$$
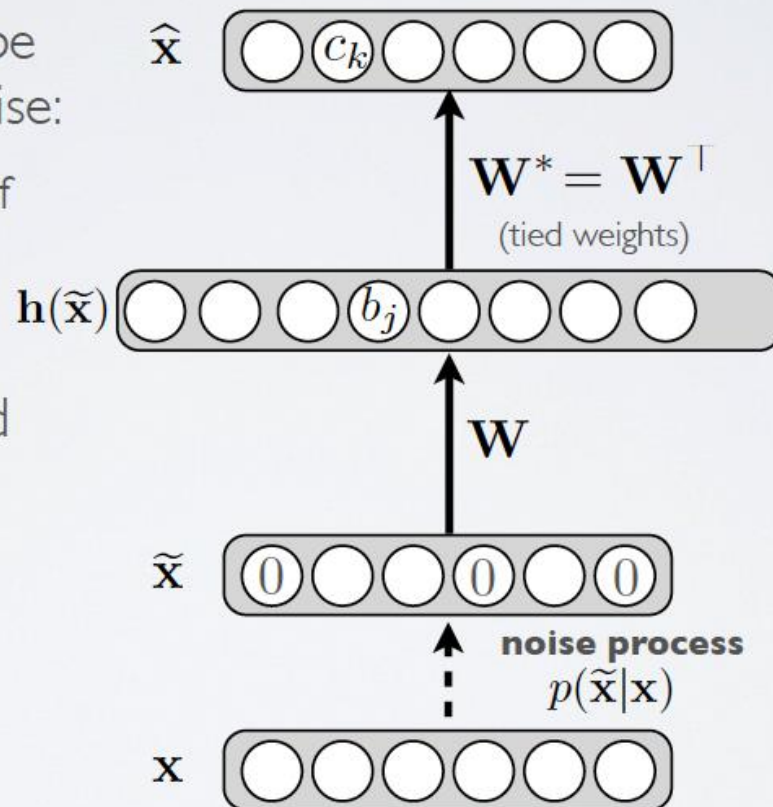
Summary

- **Four typical problems in unsupervised learning:**

  - Clustering, dimensionality reduction, representation learning, density estimation

- **Latent variable model and EM algorithm:**

  - ELBO and free energy

- **Representation learning with AutoEncoder:**

  - Reconstructing data with constraints

- **Reading material**

  - http://cs229.stanford.edu/notes2020spring/cs229-notes7a.pdf

  - http://cs229.stanford.edu/notes2020spring/cs229-notes7b.pdf

  - http://cs229.stanford.edu/notes2020spring/cs229-notes8.pdf

  Acknowledgement: Yingyu Liang@Princeton's & Feifei Li's cs231n notes