

Optimization and Machine Learning SI151

Lu Sun

School of Information Science and Technology
ShanghaiTech University

March 4, 2020

Today:

- Overview of supervised learning II
 - Statistical decision theory
 - Local methods in high dimensions
 - Statistical models
 - Model selection

Readings:

- The Element of Statistical Learning, Chapters 1 and 2
- Pattern Recognition and Machine Learning, Chapter 1

Overview of Supervised Learning II


--- Statistical Decision Theory

Statistical Decision Theory

- **Given:**
 - random input vector $X \in \mathbb{R}^p$,
 - random output variable $Y \in \mathbb{R}$,
 - joint distribution $\Pr(X, Y)$,
- **Goal:** we seek a function $f(X)$ for predicting Y given values of X .
- To penalize prediction errors, we introduce the *loss function* $L(Y, f(X))$.
- **Squared error loss:**
$$L(Y, f(X)) = (Y - f(X))^2.$$
- **Expected prediction error (EPE):**
$$\begin{aligned} \text{EPE}(f) &= E(Y - f(X))^2 \\ &= \int (y - f(x))^2 \Pr(dx, dy). \end{aligned}$$
- Since $\Pr(X, Y) = \Pr(Y|X) \Pr(X)$, EPE can also be written as
$$\text{EPE}(f) = E_X E_{Y|X}([Y - f(X)]^2 | X).$$
- Thus, it suffices to minimize EPE *pointwise*:
$$f(x) = \operatorname{argmin}_c E_{Y|X}([Y - c]^2 | X = x)$$

Regression function: $f(x) = E(Y|X = x)$.

Statistical Decision Theory



Regression function: $f(x) = E(Y|X = x)$.

Statistical Decision Theory

- Nearest neighbor methods try to directly implement this recipe
$$\hat{f}(x) = \text{Ave}(y_i | x_i \in N_k(x)).$$
- Two approximations:
 - expectation is approximated by averaging over sample data;
 - conditioning at a point is relaxed to conditioning on neighborhood.
- As $N, k \rightarrow \infty$ and $\frac{k}{N} \rightarrow 0$, we have
$$\hat{f}(x) \rightarrow E(Y | X = x).$$
- But usually we do not have very large samples.
- By making assumptions (**linearity**), we can reduce the required number of observations greatly.
- As increasing the number p of dimensions, the number N of observations required in the training data set **increases exponentially**.
- Thus the *rate of convergence* to the true estimator (with increasing k) decreases.

Regression function: $f(x) = E(Y | X = x).$

Statistical Decision Theory

- Linear regression assumes that the regression function is approximately linear

$$f(x) \approx x^T \beta.$$

- This is a model-based approach.
- Plugging this $f(x)$ into EPE,

$$\begin{aligned} \text{EPE}(f) &= E(Y - f(X))^2 \\ &= E((Y - X^T \beta)^T (Y - X^T \beta)) \end{aligned}$$

- Differentiating w.r.t. β , leads to

$$\beta = [E(XX^T)]^{-1} E(XY)$$

- Again, linear regression replaces the theoretical expectation by averaging over the observed data

$$\text{RSS}(\beta) = \sum_{i=1}^N (y_i - x_i^T \beta)^2$$

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- Summary – approximation of $f(X)$
 - Least squares:
globally linear function
 - Nearest neighbors:
locally constant function.

Regression function: $f(x) = E(Y|X = x)$.

Statistical Decision Theory

- Additional methods in our course are often model-based but more flexible than the linear model.
- For example, additive models

$$f(X) = \sum_{j=1}^p f_j(X_j)$$

- Coordinate function f_j is arbitrary.
- Approximate *univariate* conditional expectations *simultaneously* for each f_j .
- Model assumption: **additivity**.

- What happens if we use another loss function?

$$L_1(Y, f(X)) = E|Y - f(X)|$$

- In this case,
$$\hat{f}(x) = \text{median}(Y|X = x)$$
- More **robust** than the conditional mean.
- Summary:
 - L_1 criterion **not differentiable**.
 - Squared error is the most popular.

Statistical Decision Theory

- Procedure for **categorical output variable** G with values from \mathcal{G} .
- **Loss function** is $K \times K$ matrix \mathbf{L} , where $K = \text{card}(\mathcal{G})$
- $\mathbf{L}(k, l)$ is the price paid for misclassifying an observation belonging to class \mathcal{G}_k as class \mathcal{G}_l
- \mathbf{L} is zero on the diagonal
- We often use the **zero-one loss** function

$$\mathbf{L}(k, l) = 1 - \delta_{kl}$$

where $\delta_{kl} = 1$ if $k = l$, otherwise $\delta_{kl} = 0$

- **Expected prediction error (EPE)**

$$\text{EPE} = \mathbb{E}[L(G, \hat{G}(X))]$$

where expectation taken w.r.t. $\Pr(G, X)$

- Conditioning on X yields

$$\text{EPE} = \mathbb{E}_X \sum_{k=1}^K L[\mathcal{G}_k, \hat{G}(X)] \Pr(\mathcal{G}_k | X)$$

- Again, it suffices to pointwise minimization


$$\hat{G}(x) = \operatorname{argmin}_{g \in \mathcal{G}} \sum_{k=1}^K L(\mathcal{G}_k, g) \Pr(\mathcal{G}_k | X = x)$$

- Or simply

$$\hat{G}(x) = \max_{g \in \mathcal{G}} \Pr(g | X = x)$$

Bayes classifier

Statistical Decision Theory



Since the generating density is known for each class, this boundary can be calculated exactly.

- Expected prediction error (EPE)

$$\text{EPE} = E[L(G, \hat{G}(X))]$$

where expectation taken w.r.t. $\Pr(G, X)$

- Conditioning on X yields

$$\text{EPE} = E_X \sum_{k=1}^K L[\mathcal{G}_k, \hat{G}(X)] \Pr(\mathcal{G}_k | X)$$

- Again, it suffices to pointwise minimization


$$\hat{G}(x) = \operatorname{argmin}_{g \in G} \sum_{k=1}^K L(\mathcal{G}_k, g) \Pr(\mathcal{G}_k | X = x)$$

- Or simply

$$\hat{G}(x) = \max_{g \in G} \Pr(g | X = x)$$

Bayes classifier

Statistical Decision Theory



- Expected prediction error (EPE)

$$\text{EPE} = E[L(G, \hat{G}(X))]$$

where expectation taken w.r.t. $\Pr(G, X)$

- Conditioning on X yields

$$\text{EPE} = E_X \sum_{k=1}^K L[\mathcal{G}_k, \hat{G}(X)] \Pr(\mathcal{G}_k | X)$$

- Again, it suffices to pointwise minimization

$$\hat{G}(x) = \operatorname{argmin}_{g \in G} \sum_{k=1}^K L(\mathcal{G}_k, g) \Pr(\mathcal{G}_k | X = x)$$

- Or simply

$$\hat{G}(x) = \max_{g \in G} \Pr(g | X = x)$$


Bayes classifier

Overview of Supervised Learning II

--- Local Methods in High Dimensions


Local Models in High Dimensions

- **Curse of Dimensionality:**
Local neighborhoods become increasingly global, as the number of dimension increases
- **Example:**
Points uniformly distributed in a p -dimensional unit hypercube.
- Hypercubical neighborhood in p dimensions that captures a fraction r of the data
 - edge length: $e_p(r) = r^{\frac{1}{p}}$
 - $e_{10}(0.01) = 0.63$
 - $e_{10}(0.1) = 0.80$




In **ten** dimensions we need to cover **63%** (**80%**) of the range of each coordinate to capture **1%** (**10%**) of the data.

Local Models in High Dimensions



Reducing r reduces the number of observations and thus the stability.



In **ten** dimensions we need to cover **63%** (**80%**) of the range of each coordinate to capture **1%** (**10%**) of the data.

Local Models in High Dimensions

- In high dimensions, all sample points are close to the edge of the sample
- N data points uniformly distributed in a p -dimensional unit ball centered at the origin
- **Median distance** from the closest point to the origin

$$d(p, N) = \left(1 - \frac{1}{2}\right)^{1/p}$$

- $d(10,500) \approx 0.52$: more than **half** the way to the boundary

$$(1) \quad \prod_{i=1}^N \Pr(\|x_i\| > r) = \frac{1}{2}$$

$$(2) \quad \Pr(\|x_i\| > r) = 1 - \Pr(\|x_i\| \leq r) = 1 - r^p$$

$$(3) \quad (1 - r^p)^N = \frac{1}{2}$$

$$\text{Volume of a } p\text{-ball: } V_p(r) = \frac{\pi^{\frac{p}{2}}}{\Gamma(\frac{p}{2}+1)} r^p$$

Local Models in High Dimensions


- In high dimensions, all sample points are close to the edge of the sample
- N data points uniformly distributed in a p -dimensional unit ball centered at the origin
- **Median distance** from the closest point to the origin
- Sampling density is proportional to $N^{1/p}$
- If $N_1 = 100$ is a dense sample for one input, then $N_{10} = 100^{10}$ is an equally dense sample for 10 inputs.
- Thus in high dimensions all feasible training samples **sparsely populate** the input space.

$$d(p, N) = \left(1 - \frac{1^{1/N}}{2}\right)^{1/p}$$

- $d(10, 500) \approx 0.52$: more than **half** the way to the boundary

Local Models in High Dimensions

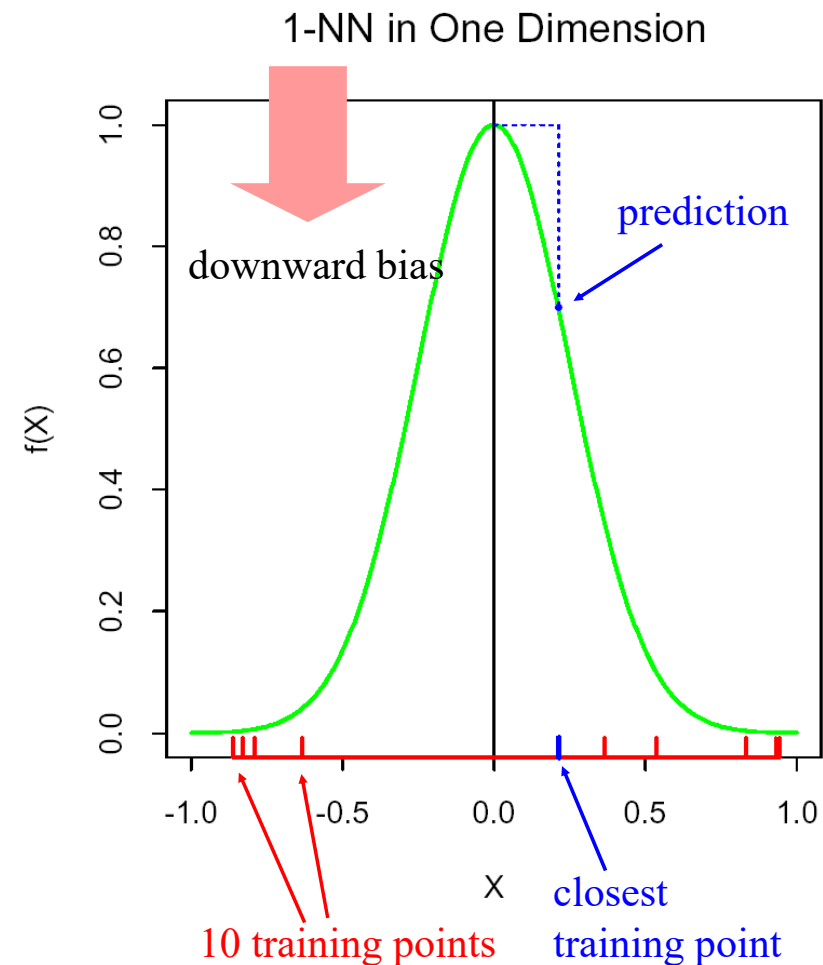
- Another example
- \mathcal{T} : set of training points x_i generated uniformly in $[-1,1]^p$ (red)
- Functional relationship between X and Y (green)
$$Y = f(X) = e^{-8\|X\|^2}$$
- No measurement error
- Error of a 1-nearest neighbor classifier in estimating $f(0)$ (blue)



Local Models in High Dimensions

- Another example
- Problem deterministic:
Prediction error is the **mean-squared error** for estimating $f(0)$

$$\begin{aligned}\text{MSE}(x_0) &= \mathbb{E}_{\mathcal{T}}[f(x_0) - \hat{y}_0]^2 \\ &= \mathbb{E}_{\mathcal{T}}[\hat{y}_0 - \mathbb{E}_{\mathcal{T}}(\hat{y}_0)]^2 \\ &\quad + [\mathbb{E}_{\mathcal{T}}(\hat{y}_0) - f(x_0)]^2 \\ &= \text{Var}_{\mathcal{T}}(\hat{y}_0) + \text{Bias}^2(\hat{y}_0)\end{aligned}$$



Local Models in High Dimensions

$$\begin{aligned}\text{MSE}(x_0) &= E_{\mathcal{T}}[f(x_0) - \hat{y}_0]^2 \\ &= E_{\mathcal{T}}[\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0) + E_{\mathcal{T}}(\hat{y}_0) - f(x_0)]^2 \\ &= E_{\mathcal{T}} \left[(\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0))^2 + \underbrace{2(\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0))(E_{\mathcal{T}}(\hat{y}_0) - f(x_0))}_{E_{\mathcal{T}}(\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0))(E_{\mathcal{T}}(\hat{y}_0) - f(x_0)) = 0} + \underbrace{(E_{\mathcal{T}}(\hat{y}_0) - f(x_0))^2}_{\text{Constant}} \right] \\ &= E_{\mathcal{T}} \left[(\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0))^2 \right] + (E_{\mathcal{T}}(\hat{y}_0) - f(x_0))^2 \\ &= \text{Var}_{\mathcal{T}}(\hat{y}_0) + \text{Bias}^2(\hat{y}_0)\end{aligned}$$

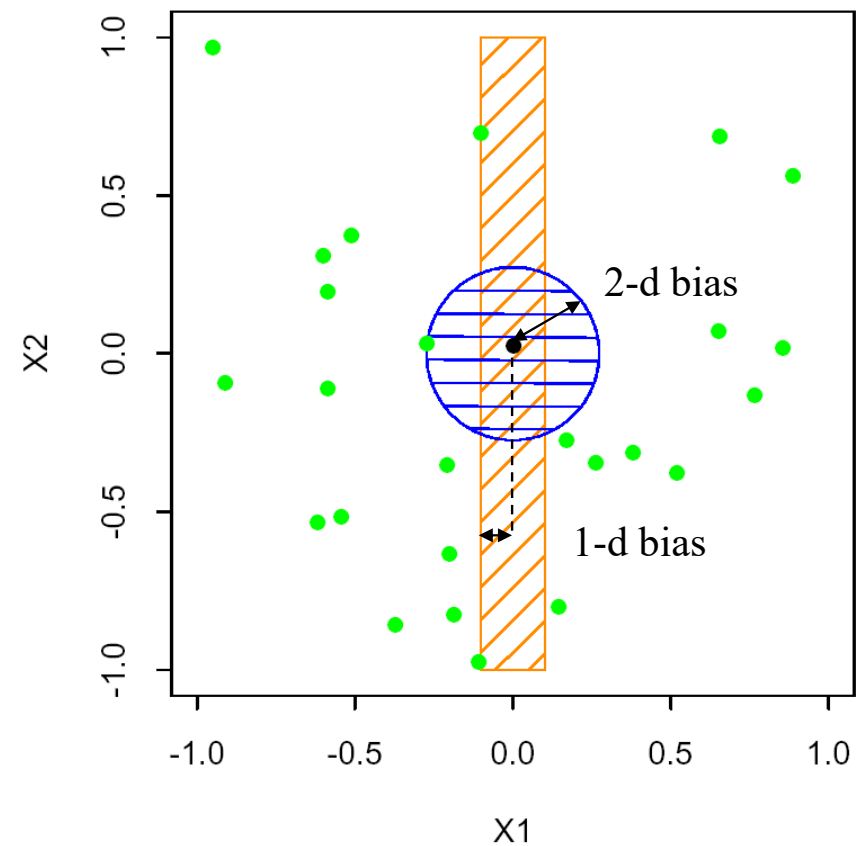
This is known as **the bias-variance decomposition**.

Local Models in High Dimensions

- Another example
- 1-d (red) vs 2-d (blue)
- As p increases, the bias increases

$$\begin{aligned}\text{MSE}(x_0) &= E_{\mathcal{T}}[f(x_0) - \hat{y}_0]^2 \\ &= E_{\mathcal{T}}[\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0)]^2 \\ &\quad + [E_{\mathcal{T}}(\hat{y}_0) - f(x_0)]^2 \\ &= \text{Var}_{\mathcal{T}}(\hat{y}_0) + \text{Bias}^2(\hat{y}_0)\end{aligned}$$

1-NN in One vs. Two Dimensions




Local Models in High Dimensions

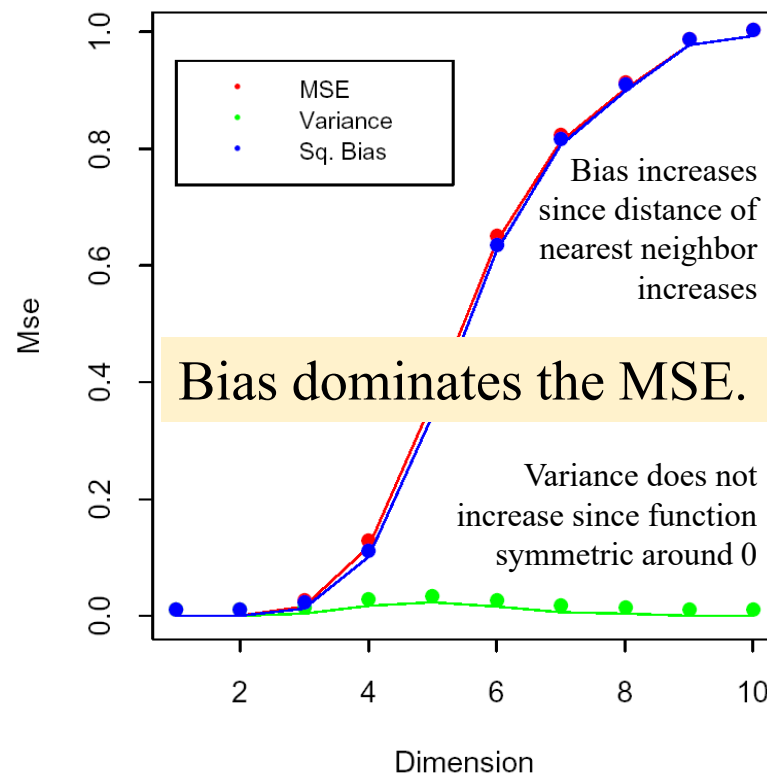
$$Y = f(X) = e^{-8\|X\|^2}$$

- The case on $N=1000$ training points

Distance to 1-NN vs. Dimension



MSE vs. Dimension




Local Models in High Dimensions


- Yet another example

$$Y = f(X) = \frac{1}{2}(X_1 + 1)^3$$

1-NN in One Dimension



MSE vs. Dimension



Local Models in High Dimensions

- Suppose a linear relationship with measurement error

$$Y = X^T \beta + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We fit the model by **least squares**.

- For an arbitrary test point x_0 ,

$$\hat{y}_0 = x_0^T \hat{\beta} = x_0^T \beta + \sum_{i=1}^N \ell_i(x_0) \varepsilon_i$$

where $\ell_i(x_0)$ is the i -th element of $\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} x_0$

Q: How can we get it?

Hint: solution of least squares:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Local Models in High Dimensions

- Suppose a linear relationship with measurement error

$$Y = X^T \beta + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We fit the model by **least squares**
- For an arbitrary test point x_0 ,

$$\hat{y}_0 = x_0^T \hat{\beta} = x_0^T \beta + \sum_{i=1}^N \ell_i(x_0) \varepsilon_i$$

where $\ell_i(x_0)$ is the i -th element of $\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} x_0$

- We find that

$$\text{EPE}(x_0) = \sigma^2 + E_{\mathcal{T}}[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} x_0] \sigma^2$$

$$\begin{aligned} \text{EPE}(x_0) &= E_{y_0|x_0} E_{\mathcal{T}}(y_0 - \hat{y}_0)^2 \\ &= \text{Var}(y_0|x_0) + E_{\mathcal{T}}[y_0 - E_{\mathcal{T}}(\hat{y}_0)]^2 + [E_{\mathcal{T}}(\hat{y}_0) - E_{\mathcal{T}}(y_0)]^2 \\ &= \text{Var}(y_0|x_0) + \text{Var}_{\mathcal{T}}(\hat{y}_0) + \text{Bias}^2(\hat{y}_0) \\ &= \sigma^2 + E_{\mathcal{T}}[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} x_0] \sigma^2 + 0^2 \end{aligned}$$

Additional variance

Local Models in High Dimensions

- Suppose a linear relationship with measurement error

$$Y = X^T \beta + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We fit the model by **least squares**.
- For an arbitrary test point x_0 ,

$$\hat{y}_0 = x_0^T \hat{\beta} = x_0^T \beta + \sum_{i=1}^N \ell_i(x_0) \varepsilon_i$$

where $\ell_i(x_0)$ is the i -th element of $\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} x_0$

- We find that

$$\text{EPE}(x_0) = \sigma^2 + E_{\mathcal{T}}[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} x_0] \sigma^2$$

- Additional variance σ^2 originates from the **nondeterministic** output
- Variance depends on x_0
- If N is large, we get

$$E_{x_0} \text{EPE}(x_0) \sim \frac{\sigma^2}{N} p + \sigma^2$$

$$\begin{aligned} E_{x_0} \text{EPE}(x_0) &\sim \frac{\sigma^2}{N} E_{x_0} x_0^T \text{Cov}(X)^{-1} x_0 + \sigma^2 \\ &= \frac{\sigma^2}{N} \text{trace}(\text{Cov}(X)^{-1} \text{Cov}(x_0)) + \sigma^2 \\ &= \frac{\sigma^2}{N} p + \sigma^2 \end{aligned}$$

Assume $E(X) = 0$, we have $\mathbf{X}^T \mathbf{X} \rightarrow N \text{Cov}(\mathbf{X})$

Local Models in High Dimensions

- Suppose a linear relationship with measurement error

$$Y = X^T \beta + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We fit the model by **least squares**.
- For an arbitrary test point x_0 ,

$$\hat{y}_0 = x_0^T \hat{\beta} = x_0^T \beta + \sum_{i=1}^N \ell_i(x_0) \varepsilon_i$$

where $\ell_i(x_0)$ is the i -th element of $\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} x_0$

- We find that

$$\text{EPE}(x_0) = \sigma^2 + E_{\mathcal{T}}[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} x_0] \sigma^2$$

- Additional variance σ^2 originates from the **nondeterministic** output
- Variance depends on x_0
- If N is large, we get

$$E_{x_0} \text{EPE}(x_0) \sim \frac{\sigma^2}{N} p + \sigma^2$$

- As p increases, variance grows negligible for large N or small σ^2
- No bias
- Curse of dimensionality controlled

Local Models in High Dimensions


- More generally

$$Y = f(X) + \varepsilon,$$

X uniform, $\varepsilon \sim \mathcal{N}(0,1)$

- Sample size: $N = 500$
- Linear case
 - EPE (Least Squares) is slightly above 1 no bias
 - EPE (1-NN) always above 2, grows slowly as nearest training point strays from target

$$\text{EPE ratio} = \frac{\text{EPE (1-NN)}}{\text{EPE (least squares)}}$$



Local Models in High Dimensions


- More generally

$$Y = f(X) + \varepsilon,$$

X uniform, $\varepsilon \sim \mathcal{N}(0,1)$

- Sample size: $N = 500$
- Cubic case
 - EPE (Least Squares)
is biased, thus ratio is smaller

$$\text{EPE ratio} = \frac{\text{EPE (1 - NN)}}{\text{EPE (least squares)}}$$



Overview of Supervised Learning II

--- Statistical Models

Statistical Models

- NN methods are the direct implementation of

$$f(x) = E(Y|X = x)$$

- But it can fail in two ways
 - With high dimensions NN need not be close to the target point
 - If special structure exists in the problem, this can be used to reduce variance and bias
- We anticipate using other classes of models for $f(x)$, and discuss **a framework for incorporating them into the prediction problem.**

Statistical Models – A Statistical Model for $\Pr(X, Y)$

- Assume **additive error model**

$$Y = f(X) + \varepsilon$$

$$E(\varepsilon) = 0$$

ε independent of X

- Then $\Pr(Y|X)$ depends only on the conditional mean of $f(x)$
- This model is a good approximation in many cases
- In many cases, $f(x)$ is deterministic and error enters through uncertainty in the **input**. This can often be mapped on uncertainty in the **output** with deterministic input.

Statistical Models – Function Approximation


- **Data:** pairs (x_i, y_i) that are points in $(p + 1)$ -dimensional Euclidean space, we fit $f: \mathbb{R}^p \rightarrow \mathbb{R}$ by
- **Linear basis expansions** have the more general form

$$y_i = f(x_i) + \varepsilon_i$$

$$f_{\theta}(x) = \sum_{k=1}^K h_k(x) \theta_k$$

- More general input spaces are possible
- **Goal:** a good approximation of $f(x)$ in some region of input space, given the training set \mathcal{T}
- Many models have certain parameters θ
 - E.g. for the linear model $f(x) = x^T \beta$ and $\theta = \beta$
- h_k : a suitable set of functions or transformations of the input vector x .
- Examples:
 - Polynomial expansions: $h_k(x) = x_1 x_2^2$
 - Trigonometric expansions: $h_k(x) = \cos(x_1)$
 - Sigmoid expansion:

$$h_k(x) = \frac{1}{1 + \exp(-x^T \beta_k)}$$



Statistical Models – Function Approximation

- Approximating f_θ by minimizing the **residual sum of squares**


$$\text{RSS}(\theta) = \sum_{i=1}^N (y_i - f_\theta(x_i))^2$$

- Linear basis expansions** have the more general form

$$f_\theta(x) = \sum_{k=1}^K h_k(x) \theta_k$$

- h_k : a suitable set of functions or transformations of the input vector x .
- Examples:
 - Polynomial expansions: $h_k(x) = x_1 x_2^2$
 - Trigonometric expansions: $h_k(x) = \cos(x_1)$
 - Sigmoid expansion:

$$h_k(x) = \frac{1}{1 + \exp(-x^T \beta_k)}$$




Statistical Models – Function Approximation

- Approximating f_θ by minimizing the **residual sum of squares**

$$\text{RSS}(\theta) = \sum_{i=1}^N (y_i - f_\theta(x_i))^2$$

- Intuition**
 - f surface in $(p + 1)$ –space
 - Observe noisy realizations
 - Want **fitted surface as close to the observed points as possible**
 - Distance measured by RSS
- Methods**
 - Closed form**: if basis function have no hidden parameters
 - Iterative**: otherwise




Statistical Models – Function Approximation

- Approximating f_θ by **maximum likelihood estimation (MLE)**
- Assume an independently drawn random sample $y_i, i = 1, \dots, N$ from a probability density $\text{Pr}_\theta(y)$.
- The log-probability of observing the sample is

$$L(\theta) = \sum_{i=1}^N \log \text{Pr}_\theta(y_i)$$

$$\begin{aligned} L(\theta) &= \log \text{Pr}_\theta(y_1, y_2, \dots, y_N) \\ &= \log \prod_{i=1}^N \text{Pr}_\theta(y_i) \end{aligned}$$

independent and identically distributed (i.i.d.)



Statistical Models – Function Approximation


- Approximating f_θ by **maximum likelihood estimation (MLE)**
- Assume an independently drawn random sample $y_i, i = 1, \dots, N$ from a probability density $\text{Pr}_\theta(y)$.
- The log-probability of observing the sample is

$$L(\theta) = \sum_{i=1}^N \log \text{Pr}_\theta(y_i)$$

- **Set θ to maximize $L(\theta)$**

Intuition:

Under the assumed statistical model, the observed data is most probable.



Statistical Models – Function Approximation

- Approximating f_θ by **maximum likelihood estimation (MLE)**
 - Assume an independently drawn random sample $y_i, i = 1, \dots, N$ from a probability density $\text{Pr}_\theta(y)$.
 - The log-probability of observing the sample is
- Least squares with additive error model**
$$Y = f_\theta(X) + \varepsilon$$
$$\varepsilon \sim \mathcal{N}(0, \sigma^2)$$

is **equivalent to maximum likelihood** with the conditional likelihood

$$L(\theta) = \sum_{i=1}^N \log \text{Pr}_\theta(y_i)$$

- Set θ to maximize $L(\theta)$**

$$\text{Pr}_\theta(y|X=x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{y - f_\theta(x)}{\sigma}\right)^2\right)$$

$$\text{Pr}_\theta(Y|X) = \mathcal{N}(f_\theta(X), \sigma^2)$$

- This is, because in this case the *log-likelihood* of data is

$$L(\theta) = -\frac{N}{2} \log(2\pi) - N \log \sigma - \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - f_\theta(x_i))^2$$

Statistical Models – Function Approximation

- Approximating f_θ by **maximum likelihood estimation (MLE)**
 - Assume an independently drawn random sample $y_i, i = 1, \dots, N$ from a probability density $\text{Pr}_\theta(y)$.
 - The log-probability of observing the sample is
- Least squares with additive error model**
$$Y = f_\theta(X) + \varepsilon$$
$$\varepsilon \sim \mathcal{N}(0, \sigma^2)$$

is **equivalent to maximum likelihood** with the conditional likelihood

$$\text{Pr}_\theta(Y|X) = \mathcal{N}(f_\theta(X), \sigma^2)$$

- This is, because in this case the *log-likelihood* of data is

$$L(\theta) = -\frac{N}{2} \log(2\pi) - N \log \sigma$$

- Set θ to maximize $L(\theta)$**

Proportional to RSS

$$\text{argmax}_\theta L(\theta) = \text{argmin}_\theta \text{RSS}(\theta) = \text{argmin}_\theta \sum_{i=1}^N (y_i - f_\theta(x_i))^2$$

$$-\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - f_\theta(x_i))^2$$

Overview of Supervised Learning II

--- Classes of Restricted Estimators

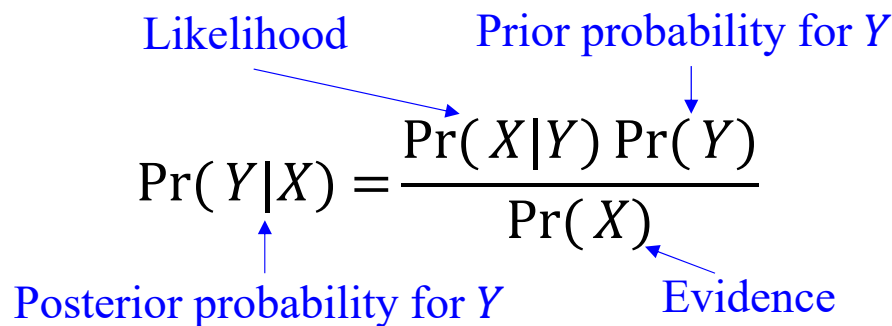
- Bayesian Methods and Roughness Penalty
- Kernel Methods and Local Regression
- Basis Functions

Bayesian Methods and Roughness Penalty

- Bayesian methods
- Formula for joint probabilities

$$\begin{aligned}\Pr(X, Y) &= \Pr(Y|X) \Pr(X) \\ &= \Pr(X|Y) \Pr(Y)\end{aligned}$$

- Bayes's theorem



The diagram shows the formula for Bayes's theorem: $\Pr(Y|X) = \frac{\Pr(X|Y) \Pr(Y)}{\Pr(X)}$. Blue arrows point from labels to parts of the formula: 'Likelihood' points to $\Pr(X|Y)$, 'Prior probability for Y' points to $\Pr(Y)$, 'Posterior probability for Y' points to $\Pr(Y|X)$, and 'Evidence' points to $\Pr(X)$.

$$\Pr(Y|X) = \frac{\Pr(X|Y) \Pr(Y)}{\Pr(X)}$$

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$

- RSS is penalized with a roughness penalty

$$\text{PRSS}(f; \lambda) = \text{RSS}(f) + \lambda J(f)$$

- $J(f)$ is large for ragged functions
 - E.g. cubic smoothing spline is the solution for the least-squares problem

$$\begin{aligned}\text{PRSS}(f; \lambda) &= \sum_{i=1}^N (y_i - f(x_i))^2 \\ &\quad + \lambda \int [f''(x)]^2 dx\end{aligned}$$

- Large second derivative is penalized

Bayesian Methods and Roughness Penalty

- Introducing penalty functions is a type of **regularization**
 - It works against **overfitting**
 - It implements beliefs about unseen parts of the problem
- In a **Bayesian** framework
 - Penalty J is the **log-prior** (probability distribution)
 - PRSS is the **log-posterior** (probability distribution)
- RSS is penalized with a **roughness penalty**
$$\text{PRSS}(f ; \lambda) = \text{RSS}(f) + \lambda J(f)$$
- $J(f)$ is large for ragged functions
 - E.g. **cubic smoothing spline** is the solution for the least-squares problem
$$\text{PRSS}(f ; \lambda) = \sum_{i=1}^N (y_i - f(x_i))^2 + \lambda \int [f''(x)]^2 dx$$
 - Large second derivative is penalized

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$

Kernel Methods and Local Regression

- Kernel functions
 - model the local neighborhoods in NN methods

- Gaussian kernel

$$K_{\lambda}(x_0, x) = \frac{1}{\lambda} \exp \left[-\frac{\|x - x_0\|^2}{2\lambda} \right]$$

- assigns weights to points that die exponentially with the square of the distance from the point x_0
- λ controls the variance of the Gaussian density (neighborhood width)

- Simplest kernel estimate:
Nadaraya-Watson weighted average

$$\hat{f}(x_0) = \frac{\sum_{i=1}^N K_{\lambda}(x_0, x_i) y_i}{\sum_{i=1}^N K_{\lambda}(x_0, x_i)}$$

- General local regression estimate of $f(x_0)$ is $f_{\hat{\theta}}(x_0)$, where $\hat{\theta}$ minimizes

$$\text{RSS}(f_{\theta}, x_0) = \sum_{i=1}^N K_{\lambda}(x_0, x_i) (y_i - f_{\theta}(x_i))^2$$

Kernel Methods and Local Regression

- f_θ is a parameterized function, such as a low-order polynomial
 - $f_\theta = \theta_0$
Nadaraya-Watson estimate
 - $f_\theta = \theta_0 + \theta^T x$
local linear regression model
 - NN methods can be regarded as kernel methods with a special metric
- Simplest Kernel estimate:
Nadaraya-Watson weighted average
 - General local regression estimate of $f(x_0)$ is $f_{\hat{\theta}}(x_0)$, where $\hat{\theta}$ minimizes

$$\hat{f}(x_0) = \frac{\sum_{i=1}^N K_\lambda(x_0, x_i) y_i}{\sum_{i=1}^N K_\lambda(x_0, x_i)}$$

$$\text{RSS}(f_\theta, x_0) = \sum_{i=1}^N K_\lambda(x_0, x_i) (y_i - f_\theta(x_i))^2$$

$x_{(k)}$ training sample ranked k in distance from x_0
 I indicator function

Basis Functions

- Include linear and polynomial expansions and more

- General form

$$f_{\theta}(x) = \sum_{m=1}^M \theta_m h_m(x)$$

- The term linear refers to the action of the parameters θ
- Cover a wide range of methods

- Radial Basis Functions

$$f_{\theta}(x) = \sum_{m=1}^M K_{\lambda_m}(\mu_m, x) \theta_m$$

- Parameters are

- ♦ Centroids μ_m
- ♦ Scales λ_m

- E.g. the Gaussian kernel

$$K_{\lambda}(\mu, x) = \exp\left(-\frac{\|x - \mu\|^2}{2\lambda}\right)$$

- Change the regression problem from a linear problem to a combinatorially nonlinear problem.

Basis Functions

- Include linear and polynomial expansions and more

- **General form**

$$f_{\theta}(x) = \sum_{m=1}^M \theta_m h_m(x)$$

- The term linear refers to the action of the parameters θ
- Cover a wide range of methods


- **Neural Networks**

- Single-layer feed-forward model

$$f_{\theta}(x) = \sum_{m=1}^M \beta_m \sigma(\alpha_m^T x + b_m)$$

Neuron weight

Neuron output



Basis Functions

- Include linear and polynomial expansions and more

- General form

$$f_{\theta}(x) = \sum_{m=1}^M \theta_m h_m(x)$$

- The term linear refers to the action of the parameters θ
- Cover a wide range of methods


- Neural Networks

- Single-layer feed-forward model

$$f_{\theta}(x) = \sum_{m=1}^M \beta_m \sigma(\alpha_m^T x + b_m)$$

- $\sigma(x)$ is known as *activation function*

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Overview of Supervised Learning II

--- Model Selection

Model Selection

- Smoothing and complexity parameters
 - Coefficient of the penalty term
 - Width of the kernel
 - Number of basis functions
- The setting of the parameters implements a trade-off between bias and variance

- Example: k -NN methods

$$Y = f(X) + \varepsilon$$

$$E(\varepsilon) = 0$$

$$\text{Var}(\varepsilon) = \sigma^2$$

- Generalization error

$$\begin{aligned} \text{EPE}_k(x_0) &= E[Y - \hat{f}_k(x_0) | X = x_0] \\ &= \sigma^2 + [\text{Bias}^2(\hat{f}_k(x_0)) + \text{Var}_{\mathcal{T}}(\hat{f}_k(x_0))] \\ &= \sigma^2 + \underbrace{\left[f(x_0) - \frac{1}{k} \sum_{\ell=1}^k f(x_{(\ell)}) \right]^2}_{\text{mean-square error}} + \frac{\sigma^2}{k} \end{aligned}$$

irreducible error

mean-square error

