


Neural Network & Reinforcement Learning

田鹏超

tianpch@shanghaitech.edu.cn

Neural Network



$\sigma(x)$ is the sigmoid function

$$\frac{1}{1 + e^{-x}}$$

Nice property: $\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$

MLE and MSE (Lecture21, p8)

$$y = f(x) + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

$$\Rightarrow y \sim \mathcal{N}(f(x), \sigma^2)$$

$$\begin{aligned} l(\mathbf{w}, \sigma) &= \ln \prod_{i=1}^N p(t|x_i, \mathbf{w}, \sigma) \\ &= \sum_{i=1}^N \ln p(t|x_i, \mathbf{w}, \sigma) \\ &= \sum_{i=1}^N \ln \left[\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(t_i - f(x_i, \mathbf{w}))^2}{2\sigma^2}\right) \right] \\ &= -\frac{1}{2\sigma^2} \sum_{i=1}^N (t_i - f(x_i, \mathbf{w}))^2 - \sum_{i=1}^N \ln(\sqrt{2\pi}\sigma) \end{aligned}$$

Thus, to maximize $l(\mathbf{w}, \sigma)$ is equal to minimize $\sum_{i=1}^N (t_i - f(x_i, \mathbf{w}))^2$, which is the sum-of-squares error (or we can convert into MSE).

MAP and regularized MSE (Lecture21, p9)

More, if we assume that the polynomial coefficients \mathbf{w} is distributed as the Gaussian distribution of the form

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha \mathbf{I})$$

We can write the poster probability function as

$$p(\theta|\mathcal{D}) = \prod_{i=1}^N p(\mathbf{w}|t_i) = \prod_{i=1}^N \frac{p(t_i|x_i, \mathbf{w}, \sigma)p(\mathbf{w}|\alpha)}{p(t_i|x_i, \sigma)}$$

Then


$$\text{maximize } \ln p(\theta|\mathcal{D}) \equiv \text{maximize } \ln \left[\prod_{i=1}^N p(t_i|x_i, \mathbf{w}, \sigma)p(\mathbf{w}|\alpha) \right]$$

Denote D is the dimension of \mathbf{w} ,

$$\begin{aligned} \ln \left[\prod_{i=1}^N p(t_i|x_i, \mathbf{w}, \sigma)p(\mathbf{w}|\alpha) \right] &= \ln \prod_{i=1}^N p(t_i|x_i, \mathbf{w}, \sigma) + \ln p(\mathbf{w}|\alpha) \\ &= -\frac{1}{2\sigma^2} \sum_{i=1}^N (t_i - f(x_i, \mathbf{w}))^2 - \sum_{i=1}^N \ln(\sqrt{2\pi}\sigma) \\ &\quad + \ln \left[\frac{1}{(2\pi)^{D/2}} \frac{1}{|\alpha \mathbf{I}|^{1/2}} \exp \left(-\frac{1}{2} \mathbf{w}^T (\alpha \mathbf{I})^{-1} \mathbf{w} \right) \right] \\ &= -\frac{1}{2\sigma^2} \sum_{i=1}^N (t_i - f(x_i, \mathbf{w}))^2 - \frac{1}{2\alpha} \mathbf{w}^T \mathbf{w} + \text{const} \end{aligned}$$

Backpropagation (MLE v.s. MAP)

Error Gradient for a Sigmoid Unit



$$\begin{aligned}\frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \left(\frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 \right) \\ &= \frac{1}{2} \sum_d \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_d 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\ &= \sum_d (t_d - o_d) \left(-\frac{\partial o_d}{\partial w_i} \right) \\ &= - \sum_d (t_d - o_d) \frac{\partial o_d}{\partial net_d} \frac{\partial net_d}{\partial w_i}\end{aligned}$$

x_d = input

t_d = target output

o_d = observed unit output

w_i = weight i

But we know:

$$\frac{\partial o_d}{\partial net_d} = \frac{\partial \sigma(net_d)}{\partial net_d} = o_d(1 - o_d)$$

$$\frac{\partial net_d}{\partial w_i} = \frac{\partial (\vec{w} \cdot \vec{x}_d)}{\partial w_i} = x_{i,d}$$

So:

$$\frac{\partial E}{\partial w_i} = - \sum_{d \in D} (t_d - o_d) o_d (1 - o_d) x_{i,d}$$

$$\left(\frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 + \lambda \|\mathbf{w}\|_2^2 \right)$$

Batch GD, SGD, Mini-batch GD

Batch mode Gradient Descent:

Do until satisfied

1. Compute the gradient $\nabla E_D[\vec{w}]$
2. $\vec{w} \leftarrow \vec{w} - \eta \nabla E_D[\vec{w}]$

Incremental mode Gradient Descent:

Do until satisfied

- For each training example d in D

1. Compute the gradient $\nabla E_d[\vec{w}]$
2. $\vec{w} \leftarrow \vec{w} - \eta \nabla E_d[\vec{w}]$

$$E_D[\vec{w}] \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

$$E_d[\vec{w}] \equiv \frac{1}{2} (t_d - o_d)^2$$

Incremental Gradient Descent can approximate
Batch Gradient Descent arbitrarily closely if η
made small enough

with size b

$$E_b[\vec{w}] = \frac{1}{2} \sum_{\substack{d \in B \\ B \subset D}} (t_d - o_d)^2$$

size $|D|$

only one observation

Reinforcement Learning

Value function & Q-Learning

Define new function, closely related to V^*

$$V^*(s) = E[r(s, \pi^*(s))] + \gamma E_{s'|\pi^*(s)}[V^*(s')]$$

$V^*(s)$ is the expected discounted reward of following the optimal policy from time 0 onward.

$$Q(s, a) = E[r(s, a)] + \gamma E_{s'|a}[V^*(s')]$$

$Q(s, a)$ is the expected discounted reward of first doing action a and then following the optimal policy from the next step onward.

If agent knows $Q(s, a)$, it can choose optimal action without knowing $P(s_{t+1}|s_t, a)$!

$$\pi^*(s) = \arg \max_a Q(s, a) \quad V^*(s) = \max_a Q(s, a)$$

Just chose the action that maximizes the Q value

And, it can learn Q without knowing $P(s_{t+1}|s_t, a)$

using something very much like the dynamic programming algorithm we used to compute V^* .

基于此假设
MDP

$$P(s_{t+1} | s_t, s_{t-1}, \dots) \\ = P(s_{t+1} | s_t)$$