

The Vapnik-Chervonenkis dimension; Sauer's Lemma

Let C be a concept class over an instance space X , i.e. a set of functions from X to $\{0, 1\}$ (where both C and X may be infinite). For any $S \subseteq X$, let's denote by $C(S)$ the set of all labelings or dichotomies on S that are induced or realized by C , i.e. if $S = \{x_1, \dots, x_m\}$, then $C(S) \subseteq \{0, 1\}^m$ and

$$C(S) = \{(c(x_1), \dots, c(x_m)); c \in C\}.$$

Also, for any natural number m , we consider $C[m]$ to be the maximum number of ways to split m points using concepts in C , that is

$$C[m] = \max \{|C(S)|; |S| = m, S \subseteq X\}.$$

Definition 1 If $|C(S)| = 2^{|S|}$ then S is **shattered** by C .

Definition 2 The **Vapnik-Chervonenkis dimension** of C , denoted as $VCdim(C)$, is the cardinality of the largest set S shattered by C . If arbitrarily large finite sets can be shattered by C , then $VCdim(C) = \infty$.

Note 1 In order to show that the VC dimension of a class is at least d we must simply find some shattered set of size d . In order to show that the VC dimension is at most d we must show that no set of size $d + 1$ is shattered.

Examples

1. Let C be the concept class of thresholds on the real number line. Clearly samples of size 1 can be shattered by this class. However, no sample of size 2 can be shattered since it is impossible to choose threshold such that x_1 is labeled positive and x_2 is labeled negative. Hence the $VCdim(C) = 1$.
2. Let C be the concept class intervals on the real line. Here a sample of size 2 is shattered, but no sample of size 3 is shattered, since no concept can satisfy a sample whose middle point is negative and outer points are positive. Hence, $VCdim(C) = 2$.
3. Let C be the concept class of k non-intersecting intervals on the real line. A sample of size $2k$ shatters (just treat each pair of points as a separate case of example 2) but no sample of size $2k + 1$ shatters, since if the sample points are alternated positive/negative, starting with a positive point, the positive points can't be covered by only k intervals. Hence $VCdim(C) = 2k$.

4. Let C the class of linear separators in \mathbf{R}^2 . Three points can be shattered, but four cannot; hence $VCdim(C) = 3$. To see why four points can never be shattered, consider two cases. The trivial case is when one point can be placed within a triangle formed by the other three; then if the middle point is positive and the others are negative, no half space can contain only the positive points. If however the points cannot be arranged in that pattern, then label two points diagonally across from each other as positive, and the other two as negative. In general, one can show that the VCdimension of the class of linear separators in \mathbf{R}^n is $n + 1$.
5. The class of axis-aligned rectangles in the plane has $VC_{DIM} = 4$. The trick here is to note that for any collection of five points, at least one of them must be interior to or on the boundary of any rectangle bounded by the other four; hence if the bounding points are positive, the interior point cannot be made negative.



Sauer's Lemma

Lemma 1 If $d = VCdim(C)$, then for all m , $C[m] \leq \Phi_d(m)$, where $\Phi_d(m) = \sum_{i=0}^d \binom{m}{i}$.

Proof: The proof proceeds by induction on both d and m . We have two base cases: when $m = 0$ and d is arbitrary, and when $d = 0$ and m is arbitrary. When $m = 0$, there can only be one subset, hence $C[0] \leq 1 = \Phi_d(0)$. When $d = VCdim(C) = 0$, no set of points can be shattered, hence all points can be labeled only one way. From this we conclude that $C[m] = 1 \leq \Phi_0(m)$. So the lemma holds for the base case.

We assume for induction that for all m', d' such that $m' \leq m$ and $d' \leq d$ and at least one of these inequalities is strict, we have $C[m'] \leq \Phi_{d'}(m')$.

Now suppose we have a set $S = \{x_1, x_2, \dots, x_m\}$ of cardinality m . Let H be a class of functions defined only over $\{x_1, x_2, \dots, x_m\}$ such that $C(S) = H(S) = H$. Since any $\tilde{S} \subseteq S$ that is shattered by H is also shattered by C , we have $VCdim(H) \leq VCdim(C)$.

We now construct H_1 and H_2 on which we apply our induction hypothesis as follows: for each possible labeling of $\{x_1, x_2, \dots, x_{m-1}\}$ induced by a function in H , we add a representative function from H to H_1 ; we let $H_2 = H \setminus H_1$. So for each $h \in H_2$, $\exists \tilde{h} \in H_1$ such that $h(x_i) = \tilde{h}(x_i)$ for $i \in \{1, \dots, m-1\}$ and $h(x_m) \neq \tilde{h}(x_m)$. For convenience, let's choose the representatives such that $h(x_m) = 1$ and $\tilde{h}(x_m) = 0$, so all $h \in H_2$ label x_m as positive.

By construction we have

$$|C(S)| = |H(S)| = |H_1(S)| + |H_2(S)|.$$

Since $H_1 \subseteq H$ we have $VCdim(H_1) \leq VCdim(H) \leq d$. Moreover, we can show

$$|H_1(S)| = |H_1(S \setminus \{x_m\})|.$$

In one direction it is clear that $|H_1(S)| \geq |H_1(S \setminus \{x_m\})|$. In the other direction we have $|H_1(S)| \leq |H_1(S \setminus \{x_m\})|$ since there is no labeling h of $\{x_1, x_2, \dots, x_{m-1}\}$ such that both $(h(x_1), h(x_2), \dots, h(x_{m-1}), 0)$ and $(h(x_1), h(x_2), \dots, h(x_{m-1}), 1)$ are in H_1 .

By induction we have:

$$|H_1(S)| \leq \Phi_d(m-1).$$

Now note that if T is shattered by H_2 , then $T \cup \{x_m\}$ is shattered by H . If T is shattered by H_2 then (a) $x_m \notin T$ (because all $h \in H_2$ label x_m as positive), and (b) $T \cup \{x_m\}$ is shattered by H (because each $h \in H_2$ has a twin $\tilde{h} \in H_1$ that is identical except on x_m) So

$$VCdim(H_2) \leq VCdim(H) - 1 \leq d - 1.$$

We can also show

$$|H_2(S)| = |H_2(\{x_1, x_2, \dots, x_{m-1}\})|,$$

and by induction we get:

$$|H_2(S)| \leq \Phi_{d-1}(m-1).$$

Combining all these we get

$$|C(S)| \leq \Phi_d(m-1) + \Phi_{d-1}(m-1).$$

Since

$$\sum_{i=0}^d \binom{m-1}{i} + \sum_{i=0}^{d-1} \binom{m-1}{i} = \binom{m}{0} + \sum_{i=1}^d \binom{m-1}{i} + \sum_{i=1}^d \binom{m-1}{i-1} = \sum_{i=0}^d \binom{m}{i},$$

we get $|C(S)| \leq \Phi_d(m)$, as desired.

■

Note that for C the class of intervals we achieve $C[m] = \Phi_d(m)$, where $d = VCdim(C)$.

Lemma 2 For $m > d$ we have:

$$\Phi_d(m) \leq \left(\frac{em}{d}\right)^d.$$

Proof: Since $m > d$, we have $0 \leq \frac{d}{m} < 1$. Therefore:

$$\left(\frac{d}{m}\right)^d \sum_{i=0}^d \binom{m}{i} \leq \sum_{i=0}^d \left(\frac{d}{m}\right)^i \binom{m}{i} \leq \sum_{i=0}^m \left(\frac{d}{m}\right)^i \binom{m}{i} = \left(1 + \frac{d}{m}\right)^m \leq e^d.$$

■

General Sample Complexity Results

We can now use Sauer's lemma to get a nice closed form expression on sample complexity. For example, we have shown last time that:

Theorem 1 Let C be an arbitrary hypothesis space. Let D be an arbitrary, fixed unknown probability distribution over X and let c^* be an arbitrary unknown target function. For any $\epsilon, \delta > 0$, if we draw a sample S from D of size m satisfying

$$2C[2m]2^{-\epsilon m/2} \leq \delta$$

then with probability at least $1 - \delta$, all the hypotheses in C with $err_D(h) > \epsilon$ are inconsistent with the data, i.e., $err_S(h) \neq 0$.

Using Sauer's lemma we can show that:

Theorem 2 *Let C be an arbitrary hypothesis space. Let D be an arbitrary, fixed unknown probability distribution over X and let c^* be an arbitrary unknown target function. For any $\epsilon, \delta > 0$, if we draw a sample S from D of size m satisfying*

$$m \geq \frac{8}{\epsilon} \left[d \ln \left(\frac{16}{\epsilon} \right) + \ln \left(\frac{2}{\delta} \right) \right].$$

then with probability at least $1 - \delta$, all the hypotheses in C with $\text{err}_D(h) > \epsilon$ are inconsistent with the data, i.e., $\text{err}_S(h) \neq 0$.



Proof: Note that it suffices to set $2^{\epsilon m/2} \geq \frac{2C[2m]}{\delta}$. To do so it suffices to set $m \geq \frac{4}{\epsilon} \ln \left(\frac{2C[2m]}{\delta} \right)$. We can now use Sauer's lemma to show that it is enough to set

$$m \geq \frac{4}{\epsilon} \left[d \ln \left(\frac{2me}{d} \right) + \ln \left(\frac{2}{\delta} \right) \right]$$

or

$$m \geq \frac{4}{\epsilon} \left[d \ln m + d \ln \left(\frac{2e}{d} \right) + \ln \left(\frac{2}{\delta} \right) \right]$$

We now use the inequality $\ln x \leq \alpha x - \ln \alpha - 1$ for $\alpha, x > 0$ to show

$$\frac{4d}{\epsilon} \ln m \leq \frac{4d}{\epsilon} \left[\frac{\epsilon}{8d} m + \ln \left(\frac{8d}{\epsilon} \right) - 1 \right] = \frac{m}{2} + \frac{4d}{\epsilon} \ln \left(\frac{8d}{e\epsilon} \right).$$

So it suffices to set

$$m \geq \frac{m}{2} + \frac{4d}{\epsilon} \ln \left(\frac{8d}{e\epsilon} \right) + \frac{4d}{\epsilon} \ln \left(\frac{2e}{d} \right) + \frac{4}{\epsilon} \ln \left(\frac{2}{\delta} \right).$$

Simplifying we get:

$$m \geq \frac{8}{\epsilon} \left[d \ln \left(\frac{16}{\epsilon} \right) + \ln \left(\frac{2}{\delta} \right) \right].$$

■