

Clustering.

Unsupervised Learning

Maria-Florina Balcan

04/06/2015

Reading:

- Chapter 14.3: Hastie, Tibshirani, Friedman.

Additional resources:

- Center Based Clustering: A Foundational Perspective.
Awasthi, Balcan. Handbook of Clustering Analysis. 2015.

Clustering, Informal Goals

Goal: Automatically partition **unlabeled** data into groups of similar datapoints.

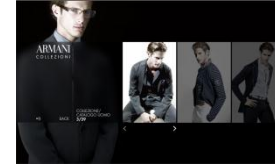
Question: When and why would we want to do this?

Useful for:

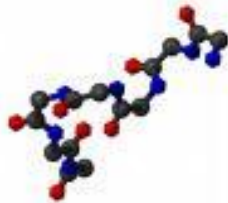
- Automatically organizing data.
- Understanding hidden structure in data.
- Preprocessing for further analysis.
 - Representing high-dimensional data in a low-dimensional space (e.g., for visualization purposes).

Applications (Clustering comes up everywhere...)

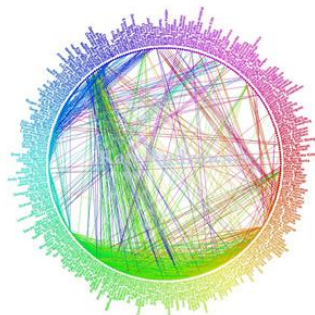
- Cluster news articles or web pages or search results by topic.



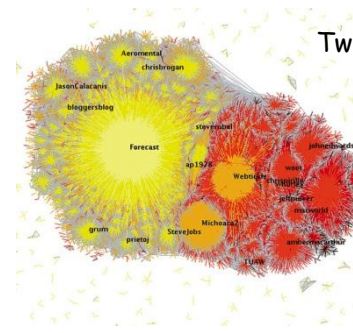
- Cluster protein sequences by function or genes according to expression profile.



- Cluster users of social networks by interest (community detection).



Facebook network



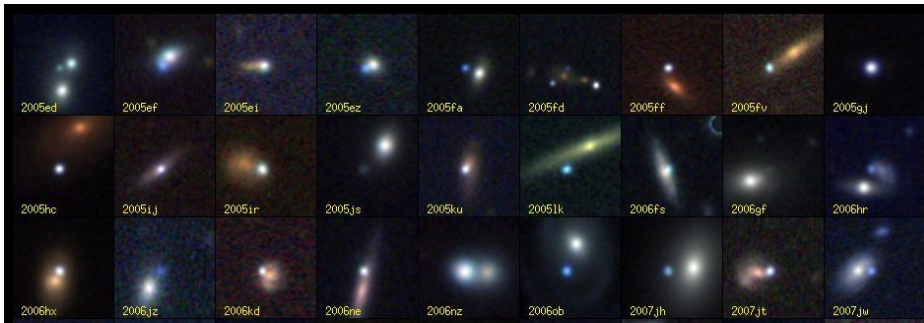
Twitter Network

Applications (clustering comes up everywhere...)

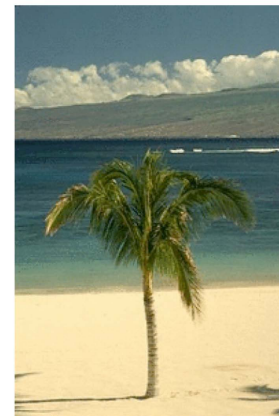
- Cluster customers according to purchase history.



- Cluster galaxies or nearby stars (e.g. Sloan Digital Sky Survey)



- Image segmentation (clustering pixels)



Clustering

Today:

- Objective based clustering *Flat*
- hard* • Hierarchical clustering *Tree*
- soft* • Mention overlapping clusters

[March 4th: EM-style algorithm for clustering for mixture of Gaussians (specific probabilistic model).]

distance metrics learning

$$d(x_1, x_2) = (x_1 - x_2)^T A (x_1 - x_2) \\ = \|L(x_1 - x_2)\|_2^2$$

$$A = L^T L$$

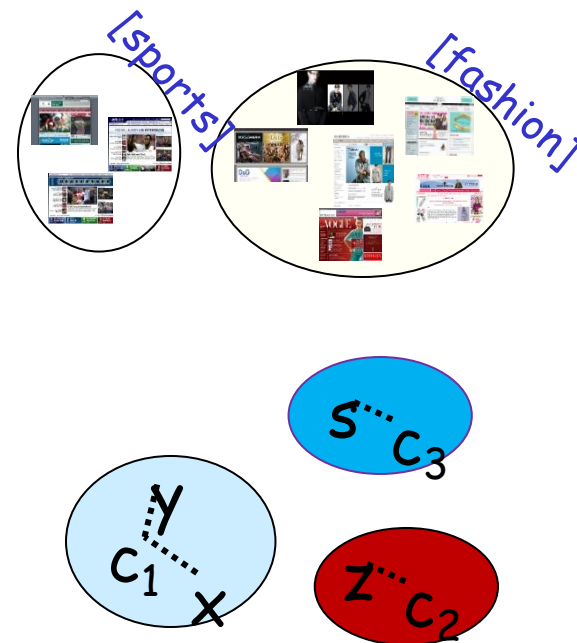
Objective Based Clustering

Input: A set S of n points, also a **distance/dissimilarity** measure specifying the distance $d(x,y)$ between pairs (x,y) .

E.g., # keywords in common, edit distance, wavelets coef., etc.

Goal: output a **partition** of the data.

- **k-means:** find center pts c_1, c_2, \dots, c_k to minimize $\sum_{i=1}^n \min_{j \in \{1, \dots, k\}} d^2(x^i, c_j)$
- **k-median:** find center pts c_1, c_2, \dots, c_k to minimize $\sum_{i=1}^n \min_{j \in \{1, \dots, k\}} d(x^i, c_j)$
- **K-center:** find partition to minimize the maxim radius



Euclidean k-means Clustering

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d
target #clusters k

Output: k representatives $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$


Objective: choose $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ to minimize

$$\min_{\mathbf{c}} \sum_{i=1}^n \min_{j \in \{1, \dots, k\}} \underbrace{\|\mathbf{x}^i - \mathbf{c}_j\|^2}_{W.}$$

within-cluster cov. (W)
between-cluster cov. (B)
total cov. (T)

$$T = W + B$$

$$\downarrow W = T - B \uparrow$$



Euclidean k-means Clustering

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d
target #clusters k


Output: k representatives $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$

Objective: choose $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ to minimize

non-convex

$$\sum_{i=1}^n \min_{j \in \{1, \dots, k\}} \|\mathbf{x}^i - \mathbf{c}_j\|^2$$

Natural assignment: each point assigned to its closest center, leads to a Voronoi partition.



Euclidean k-means Clustering

$$\min_C \sum_{i=1}^n \min_{j=1,2,\dots,k} \|x_i - c_j\|_2^2$$

$$\|a\|_2^2 = \langle a, a \rangle = a^T a$$

$$\|A\|_F^2 = \langle A, A \rangle = \text{Tr}(A^T A) = \text{Tr}(A A^T)$$


$$\min_{R, C} \sum_{i=1}^n \sum_{j=1}^k r_{ij} \|x_i - c_j\|_2^2$$

s.t. $r_{ij} \in \{0, 1\}, R \mathbb{I}_k = \mathbb{I}_n$

$$r_{ij} = \begin{cases} 1, & x_i \in C_j \\ 0, & \text{otherwise} \end{cases}$$

$$\min_{R, C} \|X - RC\|_F^2$$

s.t. $R \in \{0, 1\}^{n \times k}, R \mathbb{I}_k = \mathbb{I}_n$



R : indicator matrix

$$R^T R: \begin{matrix} k \\ \text{diagonal} \end{matrix}$$

$$R R^T: \begin{matrix} n \\ (R R^T)_{ij} \end{matrix}$$

Euclidean k-means Clustering

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d
target #clusters k

Output: k representatives $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$

Objective: choose $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ to minimize

$$\sum_{i=1}^n \min_{j \in \{1, \dots, k\}} \left\| \mathbf{x}^i - \mathbf{c}_j \right\|^2$$

Computational complexity:

NP hard: even for $k = 2$ [Dagupta'08] or
 $d = 2$ [Mahajan-Nimbhorkar-Varadarajan09]

There are a couple of easy cases...



An Easy Case for k-means: $k=1$

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d

Output: $\mathbf{c} \in \mathbb{R}^d$ to minimize $\sum_{i=1}^n \left\| \mathbf{x}^i - \mathbf{c} \right\|^2$

Solution: The optimal choice is $\boldsymbol{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}^i$

Idea: bias/variance like decomposition

$$\frac{1}{n} \sum_{i=1}^n \left\| \mathbf{x}^i - \mathbf{c} \right\|^2 = \left\| \boldsymbol{\mu} - \mathbf{c} \right\|^2 + \frac{1}{n} \sum_{i=1}^n \left\| \mathbf{x}^i - \boldsymbol{\mu} \right\|^2$$

Avg k-means cost wrt \mathbf{c}

Avg k-means cost wrt $\boldsymbol{\mu}$

Quiz 

So, the optimal choice for \mathbf{c} is $\boldsymbol{\mu}$.

Common Heuristic in Practice: The Lloyd's method

[Least squares quantization in PCM, Lloyd, IEEE Transactions on Information Theory, 1982]

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d

Initialize centers $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ and
clusters C_1, C_2, \dots, C_k in any way.

Repeat until there is no further change in the cost.

E-step • For each j : $C_j \leftarrow \{x \in S \text{ whose closest center is } \mathbf{c}_j\}$

M-step • For each j : $\mathbf{c}_j \leftarrow \text{mean of } C_j$

$\|x - \mathbf{c}_j\|_F^2$

R

C

Common Heuristic in Practice: The Lloyd's method

[Least squares quantization in PCM, Lloyd, IEEE Transactions on Information Theory, 1982]

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d

Initialize centers $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ and
clusters C_1, C_2, \dots, C_k in any way.

$$\|\mathbf{x} - \mathbf{c}_j\|_F^2$$

Repeat until there is no further change in the cost.

- For each j : $C_j \leftarrow \{x \in S \text{ whose closest center is } \mathbf{c}_j\}$

- For each j : $\mathbf{c}_j \leftarrow \text{mean of } C_j \leftarrow \mathbf{c}_j = \frac{1}{|C_j|} \sum_{x_i \in C_j} x_i$

$$C_j = \{x_i \mid j = \arg \min_{k=1, \dots, k} \|\mathbf{x}_i - \mathbf{c}_k\|_2^2, i=1, 2, \dots, n\}$$

Holding $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k$ fixed,
pick optimal C_1, C_2, \dots, C_k

Holding C_1, C_2, \dots, C_k fixed,
pick optimal $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k$

Common Heuristic: The Lloyd's method

Input: A set of n datapoints x^1, x^2, \dots, x^n in \mathbb{R}^d

Initialize centers $c_1, c_2, \dots, c_k \in \mathbb{R}^d$ and
clusters C_1, C_2, \dots, C_k in any way.

Repeat until there is no further change in the cost.

1. R • For each j : $C_j \leftarrow \{x \in S \text{ whose closest center is } c_j\}$
2. C • For each j : $c_j \leftarrow \text{mean of } C_j$

Note: it always converges.

- the cost always drops and
- there is only a finite #s of Voronoi partitions
(so a finite # of values the cost could take)

Common Heuristic: The Lloyd's method

$$\min_{R, C} \|X - RC\|_F^2 = Q(R, C).$$

$$\text{s.t. } R \in \{0, 1\}^{n \times k}, R D_k = I_n$$

$$I_k = \left\{ \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} \right\}_k, \quad I_n = \left\{ \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} \right\}_n.$$

t -th iteration.

$$R^{(t-1)}, C^{(t-1)}$$

$$1. R^{(t)} = \arg \min_R Q(R, C^{(t-1)})$$

$$Q(R^{(t)}, C^{(t-1)}) \leq Q(R^{(t-1)}, C^{(t-1)})$$

$$2. C^{(t)} = \arg \min_C Q(R^{(t)}, C)$$

$$Q(R^{(t)}, C^{(t)}) \leq Q(R^{(t)}, C^{(t-1)})$$

$$\Rightarrow Q(R^{(t)}, C^{(t)}) \leq Q(R^{(t-1)}, C^{(t-1)})$$

Initialization for the Lloyd's method

Input: A set of n datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ in \mathbb{R}^d

Initialize centers $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k \in \mathbb{R}^d$ and
clusters C_1, C_2, \dots, C_k in any way.

Repeat until there is no further change in the cost.


- For each j : $C_j \leftarrow \{x \in S \text{ whose closest center is } \mathbf{c}_j\}$
- For each j : $\mathbf{c}_j \leftarrow \text{mean of } C_j$

- **Initialization is crucial** (how fast it converges, quality of solution output)
- Discuss techniques commonly used in practice
 - Random centers from the datapoints (repeat a few times)
 - Furthest traversal
 - K-means ++ (works well and has provable guarantees)

Lloyd's method: Random Initialization


Lloyd's method: Random Initialization

Example: Given a set of datapoints




Lloyd's method: Random Initialization

Select initial centers at random




Lloyd's method: Random Initialization

Assign each point to its nearest center




Lloyd's method: Random Initialization

Recompute optimal centers given a fixed clustering




Lloyd's method: Random Initialization

Assign each point to its nearest center




Lloyd's method: Random Initialization

Recompute optimal centers given a fixed clustering



Lloyd's method: Random Initialization




Assign each point to its nearest center



Lloyd's method: Random Initialization


Recompute optimal centers given a fixed clustering

Step 2: $C^{(3)}$



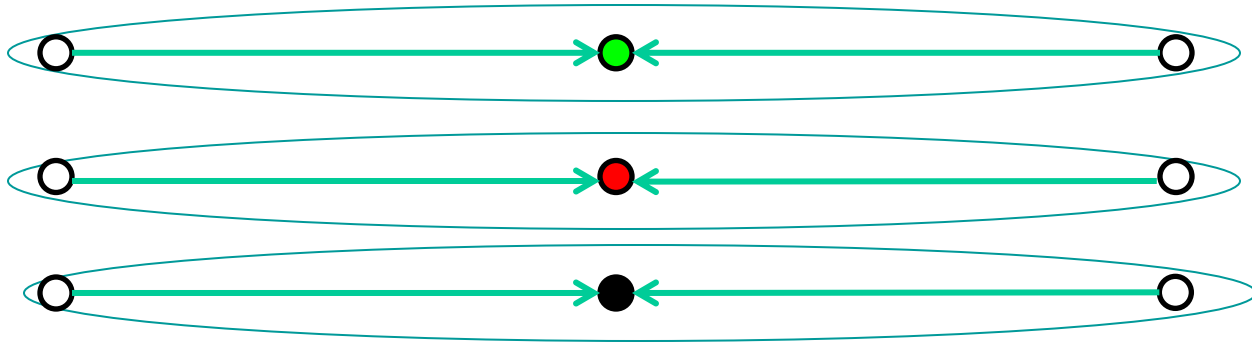
Get a good quality solution in this example.

Lloyd's method: Performance



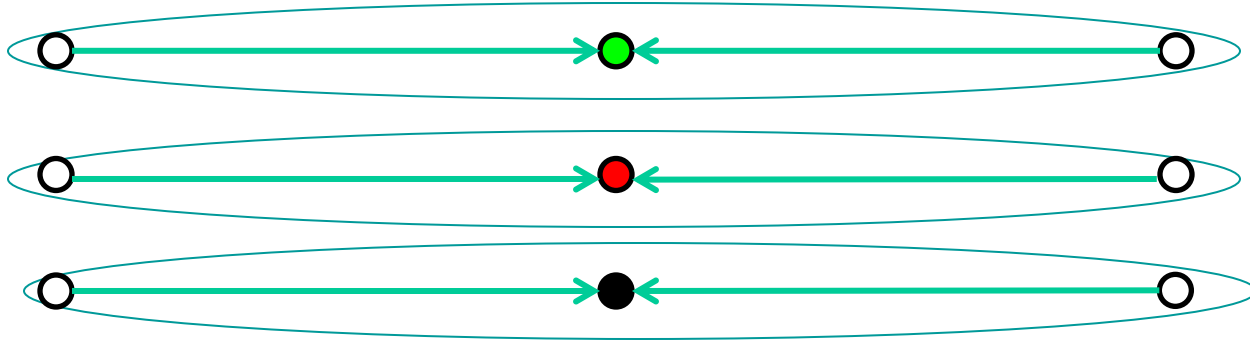
It always converges, but it may converge at a local optimum that is different from the global optimum, and in fact could be arbitrarily worse in terms of its score.

Lloyd's method: Performance

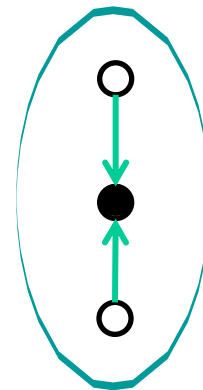
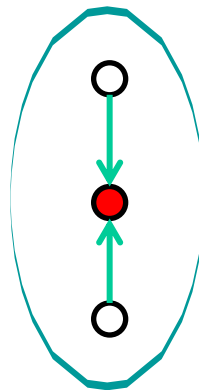
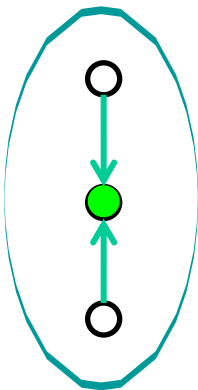


Local optimum: every point is assigned to its nearest center and every center is the mean value of its points.




Lloyd's method: Performance



.It is arbitrarily worse than optimum solution.... 🤔

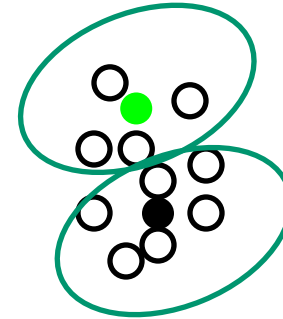
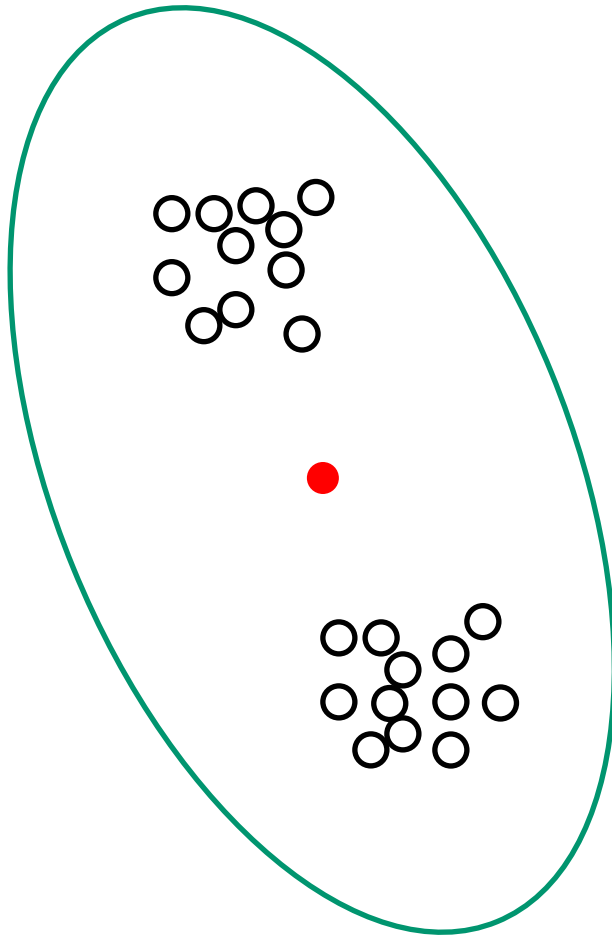


Lloyd's method: Performance



This bad performance, can happen even with well separated Gaussian clusters.

Lloyd's method: Performance



This bad performance, can happen even with well separated Gaussian clusters.

Some Gaussian are combined.....



Lloyd's method: Performance

- If we do random initialization, as k increases, it becomes more likely we won't have perfectly picked one center per Gaussian in our initialization (so Lloyd's method will output a bad solution).
- For k equal-sized Gaussians, $\Pr[\text{each initial center is in a different Gaussian}] \approx \frac{k!}{k^k} \approx \frac{1}{e^k}$
 - $k \cdot (k-1) \cdot (k-2) \cdot \dots \cdot 2 \cdot 1 = k!$
 - $\binom{k+k-1}{k-1} = O(k^k)$
- Becomes unlikely as k gets large.




Another Initialization Idea: Furthest Point Heuristic

Choose \mathbf{c}_1 arbitrarily (or at random).

- For $j = 2, \dots, k$
 - Pick \mathbf{c}_j among datapoints $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^d$ that is farthest from previously chosen $\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_{j-1}$


Fixes the Gaussian problem. But it can be thrown off by outliers....

- o Furthest point heuristic does well on previous example



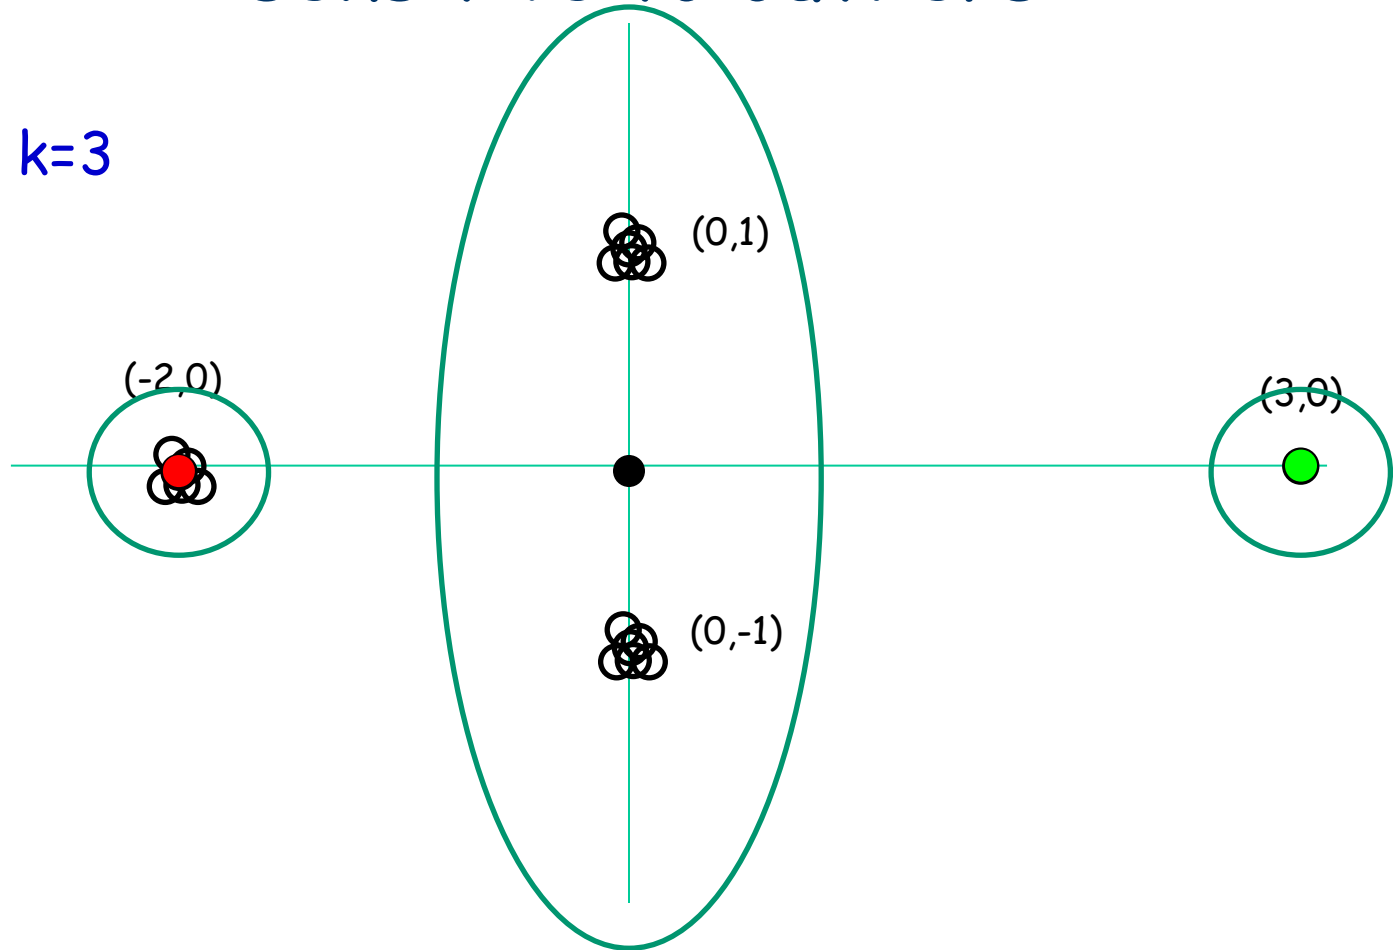
Furthest point initialization heuristic sensitive to outliers

Assume $k=3$



Furthest point initialization heuristic sensitive to outliers

Assume $k=3$



K-means++ Initialization: D^2 sampling [AV07]

- Interpolate between random and furthest point initialization
- Let $D(x)$ be the distance between a point x and its nearest center. Chose the next center proportional to $D^2(x)$.

- Choose c_1 at random.
- For $j = 2, \dots, k$
 - Pick c_j among x^1, x^2, \dots, x^n according to the distribution

$$\Pr(c_j = x^i) \propto \min_{j' < j} \|x^i - c_{j'}\|^2 D^2(x^i)$$

Theorem: K-means++ always attains an $O(\log k)$ approximation to optimal k-means solution in expectation.

Running Lloyd's can only further improve the cost.

K-means++ Idea: D^2 sampling

- Interpolate between random and furthest point initialization
- Let $D(x)$ be the distance between a point x and its nearest center. Chose the next center proportional to $D^\alpha(x)$. ℓ_α -norm

- $\alpha = 0$, random sampling

- $\alpha = \infty$, furthest point (Side note: it actually works well for k-center)

- $\alpha = 2$, k-means++

$$\|a\|_\alpha = \left(\sum_{j=1}^d |a_j|^\alpha \right)^{\frac{1}{\alpha}}$$


$$\|a\|_0 = \sum_{j=1}^d \mathbb{I}(a_j \neq 0)$$

$$\|a\|_1 = \sum_{j=1}^d |a_j|$$

$$\|a\|_\infty = \max_{j=1, \dots, d} |a_j|$$

Side note: $\alpha = 1$, works well for k-median

K-means ++ Fix



K-means++/ Lloyd's Running Time

- K-means ++ initialization: $O(nd)$ and one pass over data to select next center. So $O(nkd)$ time in total.

- Lloyd's method

Repeat until there is no change in the cost.

Step 1. For each j : $C_j \leftarrow \{x \in S \text{ whose closest center is } c_j\}$

Step 2. For each j : $c_j \leftarrow \text{mean of } C_j$

$$Q(R, c) = \|x - R\|_F^2$$

Each round takes

time $O(nkd)$. (T , #iter.)

$\rightarrow O(Tnkd)$

In total: $O(nkd) + O(Tnkd)$

- Exponential # of rounds in the worst case [AV07].
- Expected polynomial time in the smoothed analysis model!

K-means++/ Lloyd's Summary

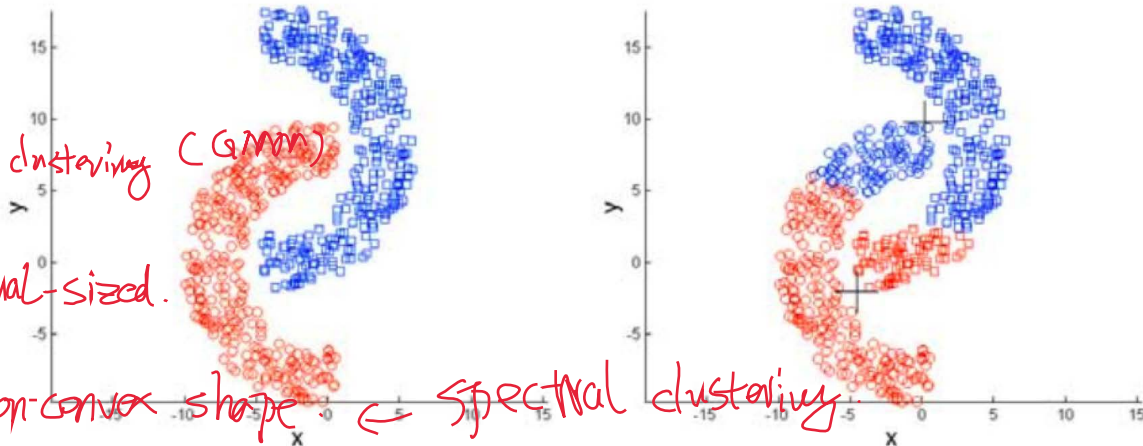
- K-means++ always attains an $O(\log k)$ approximation to optimal k-means solution in expectation.
- Running Lloyd's can only further improve the cost.
- Exponential # of rounds in the worst case [AV07].
- Expected polynomial time in the smoothed analysis model!
- Does well in practice.

Limitations:

① hard clustering \leftarrow soft \leftarrow prob. clustering (GMM)

② work well only clusters are equal-sized.

③ work badly if clusters have non-convex shape \leftarrow spectral clustering



Kernel K-means

kernel trick

$$k(\langle x_i, x_j \rangle) = \langle \phi(x_i), \phi(x_j) \rangle$$


$$\min_{R, C} \|X - RC\|_F^2 = Q(R, C).$$

$$\text{s.t. } R \in \{0, 1\}^{n \times k}, R \mathbb{I}_k = \mathbb{I}_n.$$

$$\begin{aligned} \|A\|_F^2 &= \langle A, A \rangle \\ &= \text{Tr}(A^T A) \\ &= \text{Tr}(A A^T) \end{aligned}$$

$$\begin{aligned} \text{Tr}(ABC) &= \text{Tr}(CAB) \\ &= \text{Tr}(BCA) \end{aligned}$$

$$\begin{aligned} C &= (R^T R)^{-1} R^T X \\ &= D R^T X \end{aligned}$$



$$R D R^T R D R^T = R D R^T$$

$$\langle X, R D R^T X \rangle = \text{Tr}(F^T X X^T F)$$

$$= \text{Tr}(X^T R D R^T X) = \text{Tr}(X^T F F^T X)$$

$$(F = R (R^T R)^{-\frac{1}{2}} : \text{weighted indicator})$$

$$\min_Q(R) = \|X - R D R^T X\|_F^2$$

R

$$= \langle X - R D R^T X, X - R D R^T X \rangle$$

$$= \langle X, X \rangle - 2 \langle X, R D R^T X \rangle + \langle R D R^T X, R D R^T X \rangle$$

$$= \langle X, X \rangle - \max_R \langle X, R D R^T X \rangle$$

$$= \text{Tr}(X^T \underbrace{R D R^T R D R^T}_{} X)$$

$$R D R^T$$

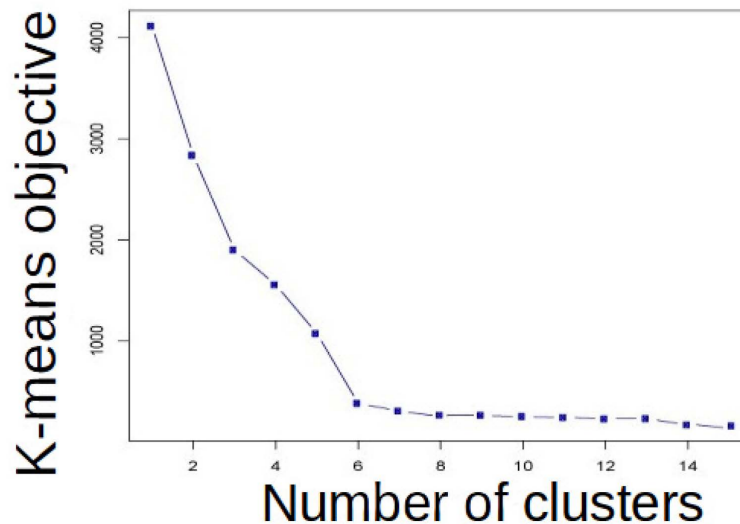
Gram.

$$\begin{aligned} & \max_R \langle X, R D R^T X \rangle \\ & \text{s.t. } R \in \{0, 1\}^{n \times k}, R \mathbb{I}_k = \mathbb{I}_n \end{aligned}$$

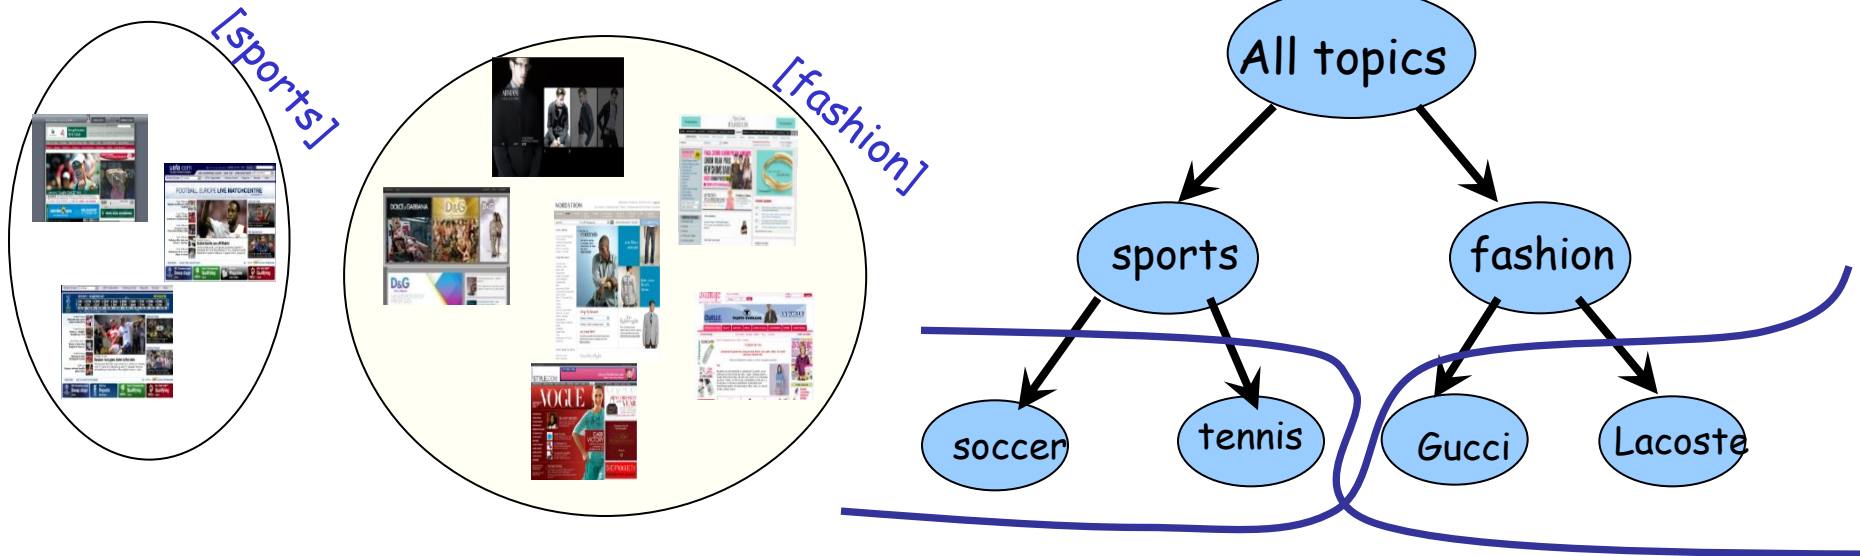
$$\begin{aligned} & \max_F \text{Tr}(F^T X X^T F) \quad \text{s.t. } F^T F = I \end{aligned}$$

What value of k???

- Heuristic: Find large gap between $k-1$ -means cost and k -means cost.
- According to information criteria (AIC, BIC, etc.)
- Try hierarchical clustering.



Hierarchical Clustering




- A hierarchy might be more natural.
- Different users might care about different levels of granularity or even prunings.

Hierarchical Clustering


Top-down (divisive)

- Partition data into 2-groups (e.g., 2-means)
- Recursively cluster each group.



Bottom-Up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the "closest" two clusters.
- Different defs of "closest" give different algorithms.

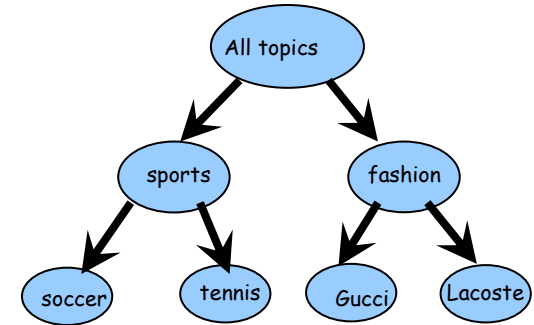


Bottom-Up (agglomerative)

Have a **distance** measure on pairs of objects.

$d(x,y)$ - distance between x and y

E.g., # keywords in common, edit distance, etc

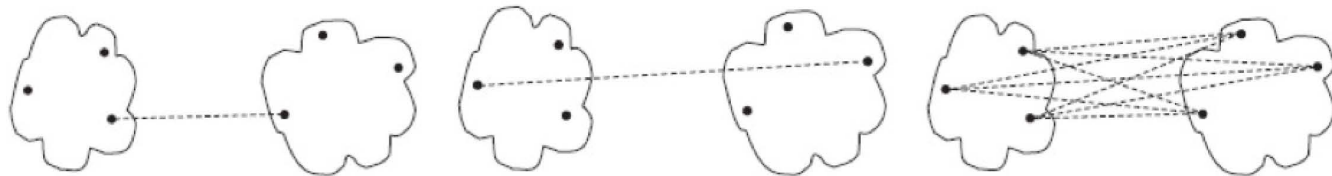


- Single linkage: $\text{dist}(A, B) = \min_{x \in A, x' \in B'} \text{dist}(x, x')$

- Complete linkage: $\text{dist}(A, B) = \max_{x \in A, x' \in B'} \text{dist}(x, x')$

- Average linkage: $\text{dist}(A, B) = \text{avg}_{x \in A, x' \in B'} \text{dist}(x, x')$

- Ward's method



(a) MIN (single link.)

(b) MAX (complete link.)

(c) Group average.

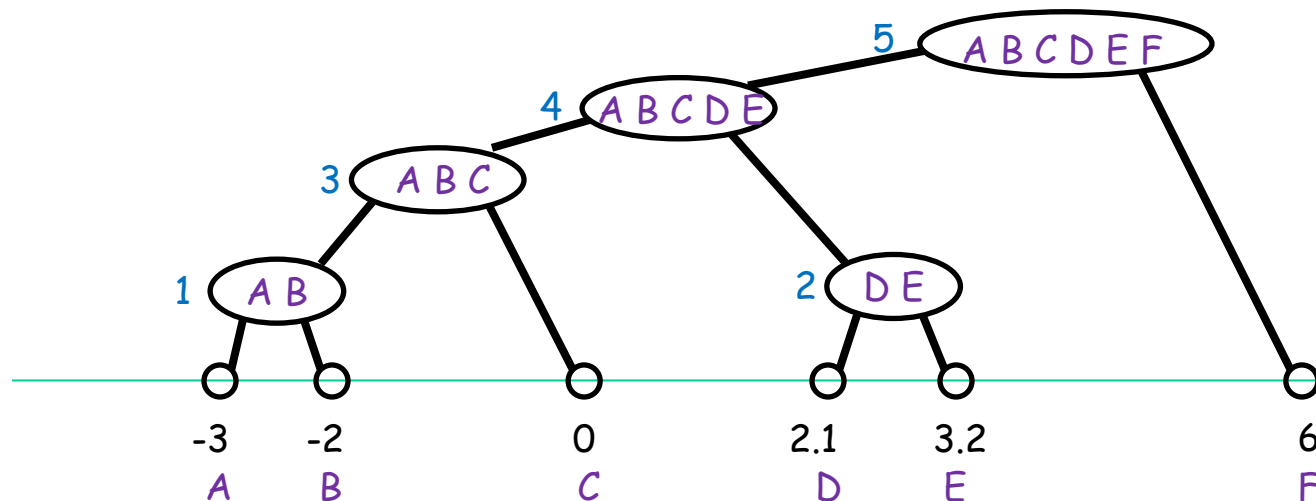
Single Linkage

Bottom-up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the "closest" two clusters.

Single linkage: $\text{dist}(A, B) = \min_{x \in A, x' \in B} \text{dist}(x, x')$

Dendrogram



Single Linkage

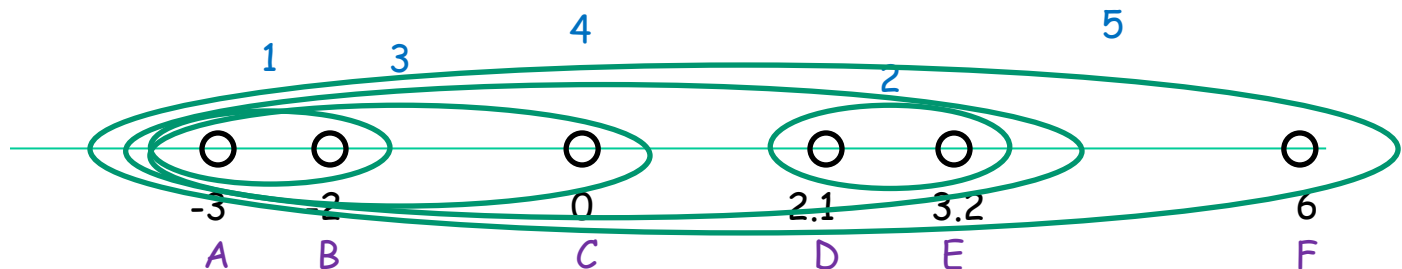
Bottom-up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the "closest" two clusters.

Single linkage: $\text{dist}(A, B) = \min_{x \in A, x' \in B} \text{dist}(x, x')$

One way to think of it: at any moment, we see connected components of the graph where connect any two pts of distance $< r$.

Watch as r grows (only $n-1$ relevant values because we only merge at value of r corresponding to values of r in different clusters).



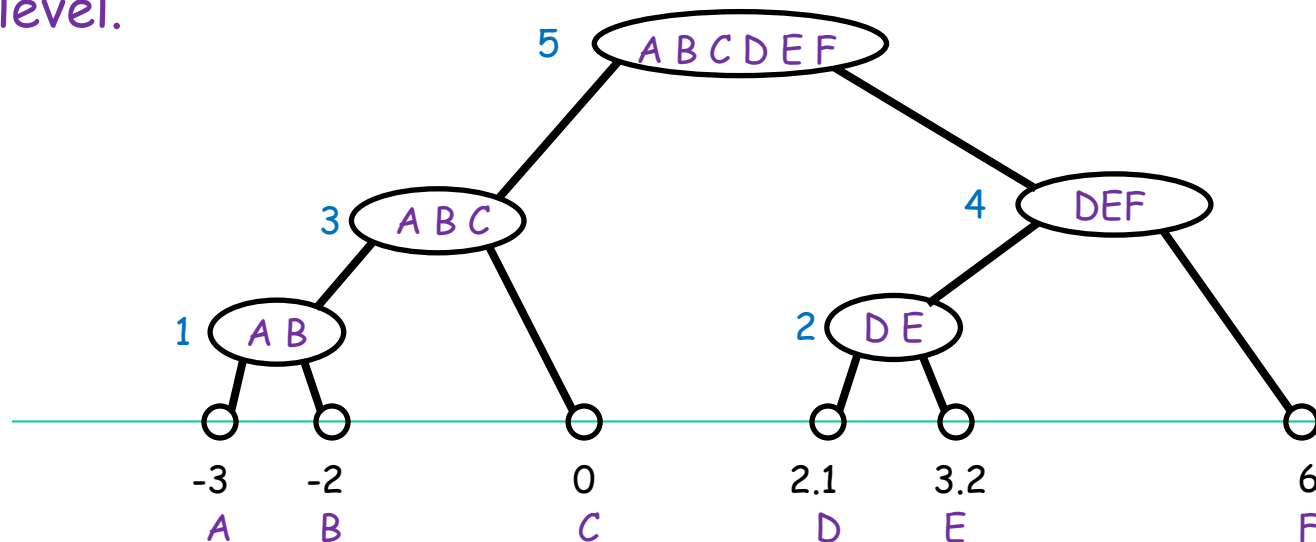
Complete Linkage

Bottom-up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the "closest" two clusters.

Complete linkage: $\text{dist}(A, B) = \max_{x \in A, x' \in B} \text{dist}(x, x')$

One way to think of it: keep max diameter as small as possible at any level.



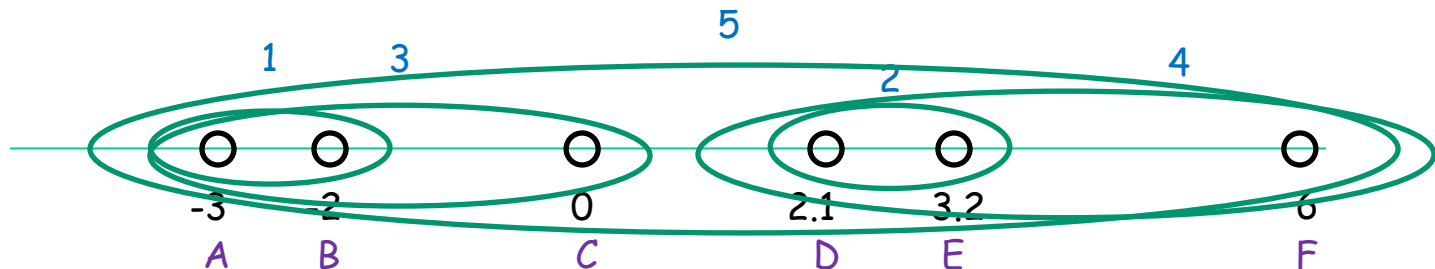
Complete Linkage

Bottom-up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the "closest" two clusters.

Complete linkage: $\text{dist}(A, B) = \max_{x \in A, x' \in B} \text{dist}(x, x')$

One way to think of it: keep max diameter as small as possible.



Ward's Method


Bottom-up (agglomerative)

- Start with every point in its own cluster.
- Repeatedly merge the “closest” two clusters.

Ward's method: $\text{dist}(C, C') = \frac{|C| \cdot |C'|}{|C| + |C'|} \|\text{mean}(C) - \text{mean}(C')\|^2$

Merge the two clusters such that the **increase in k-means cost** is as small as possible.

Works well in practice.



Running time

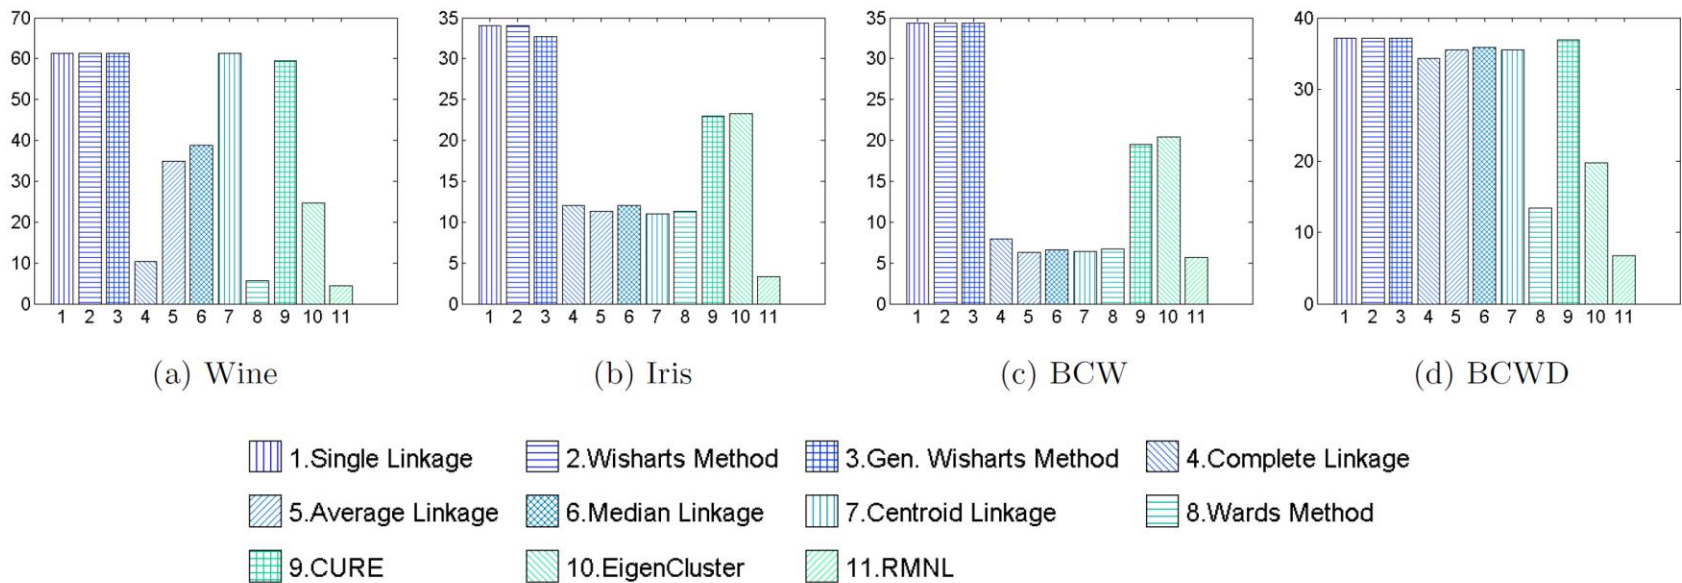
- Each algorithm starts with N clusters, and performs $N-1$ merges.
- For each algorithm, computing $\text{dist}(C, C')$ can be done in time $O(|C| \cdot |C'|)$. (e.g., examining $\text{dist}(x, x')$ for all $x \in C, x' \in C'$)
- Time to compute all pairwise distances and take smallest is $O(N^2)$.
- Overall time is $O(N^3)$.

In fact, can run all these algorithms in time $O(N^2 \log N)$.

See: Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008. <http://www-nlp.stanford.edu/IR-book/>

Hierarchical Clustering Experiments

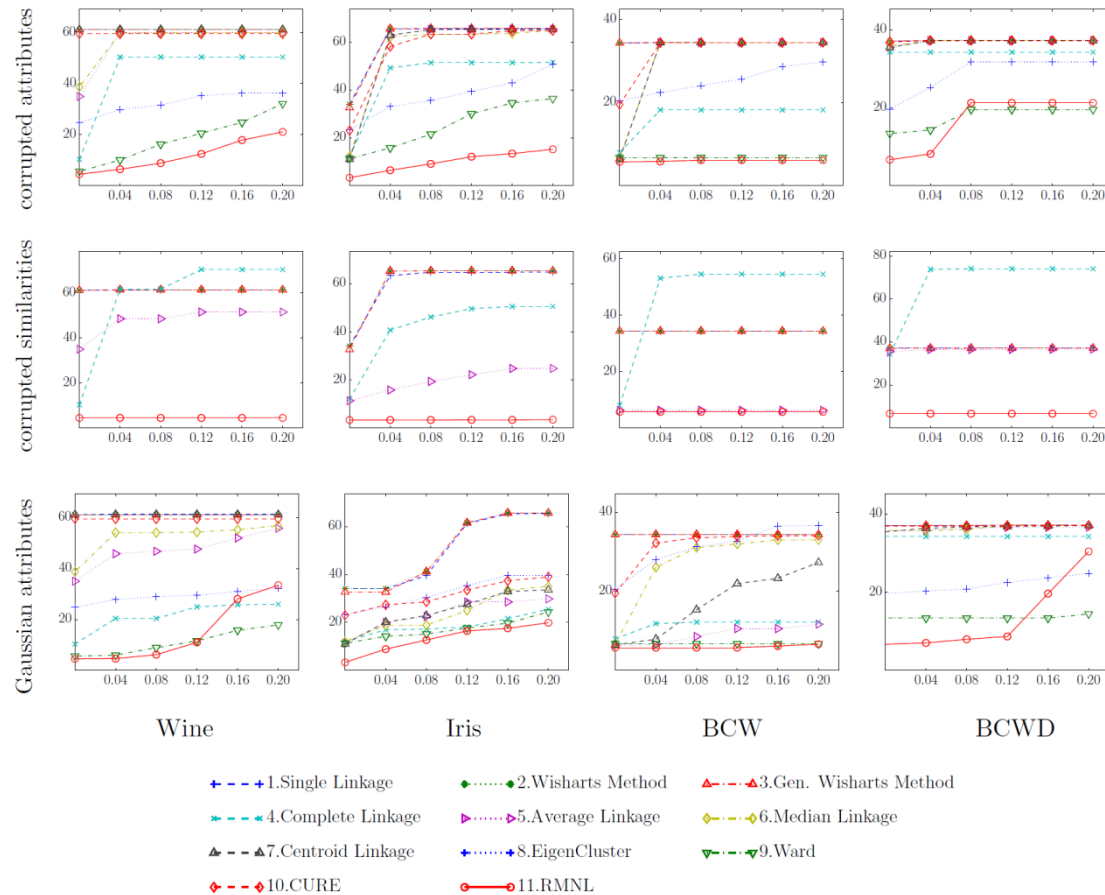
[BLG, JMLR'15]



Ward's method does the best among classic techniques.

Hierarchical Clustering Experiments

[BLG, JMLR'15]



Ward's method does the best among classic techniques.

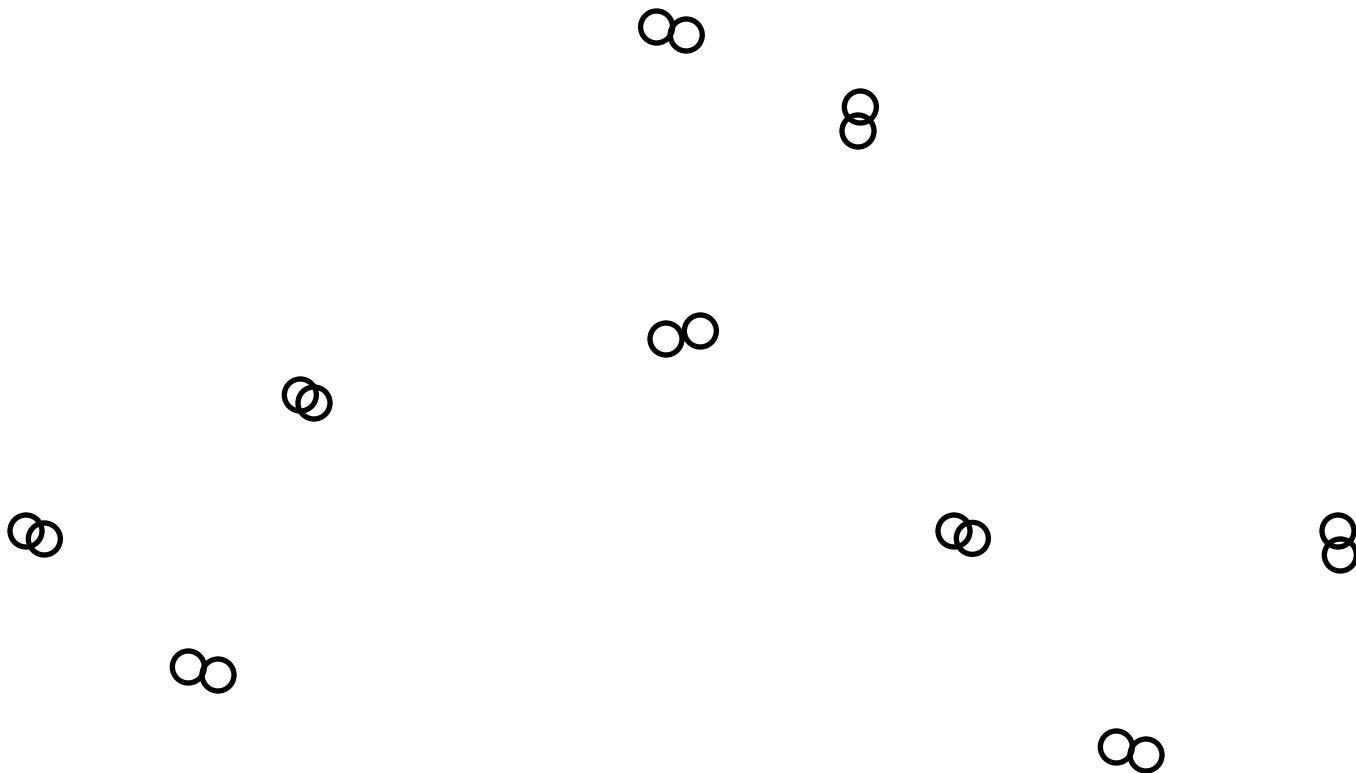
What You Should Know

- Partitional Clustering. k-means and k-means ++
 - Lloyd's method
 - Initialization techniques (random, furthest traversal, k-means++)
- Hierarchical Clustering.
 - Single linkage, Complete Linkage, Ward's method

Additional Slides

Smoothed analysis model

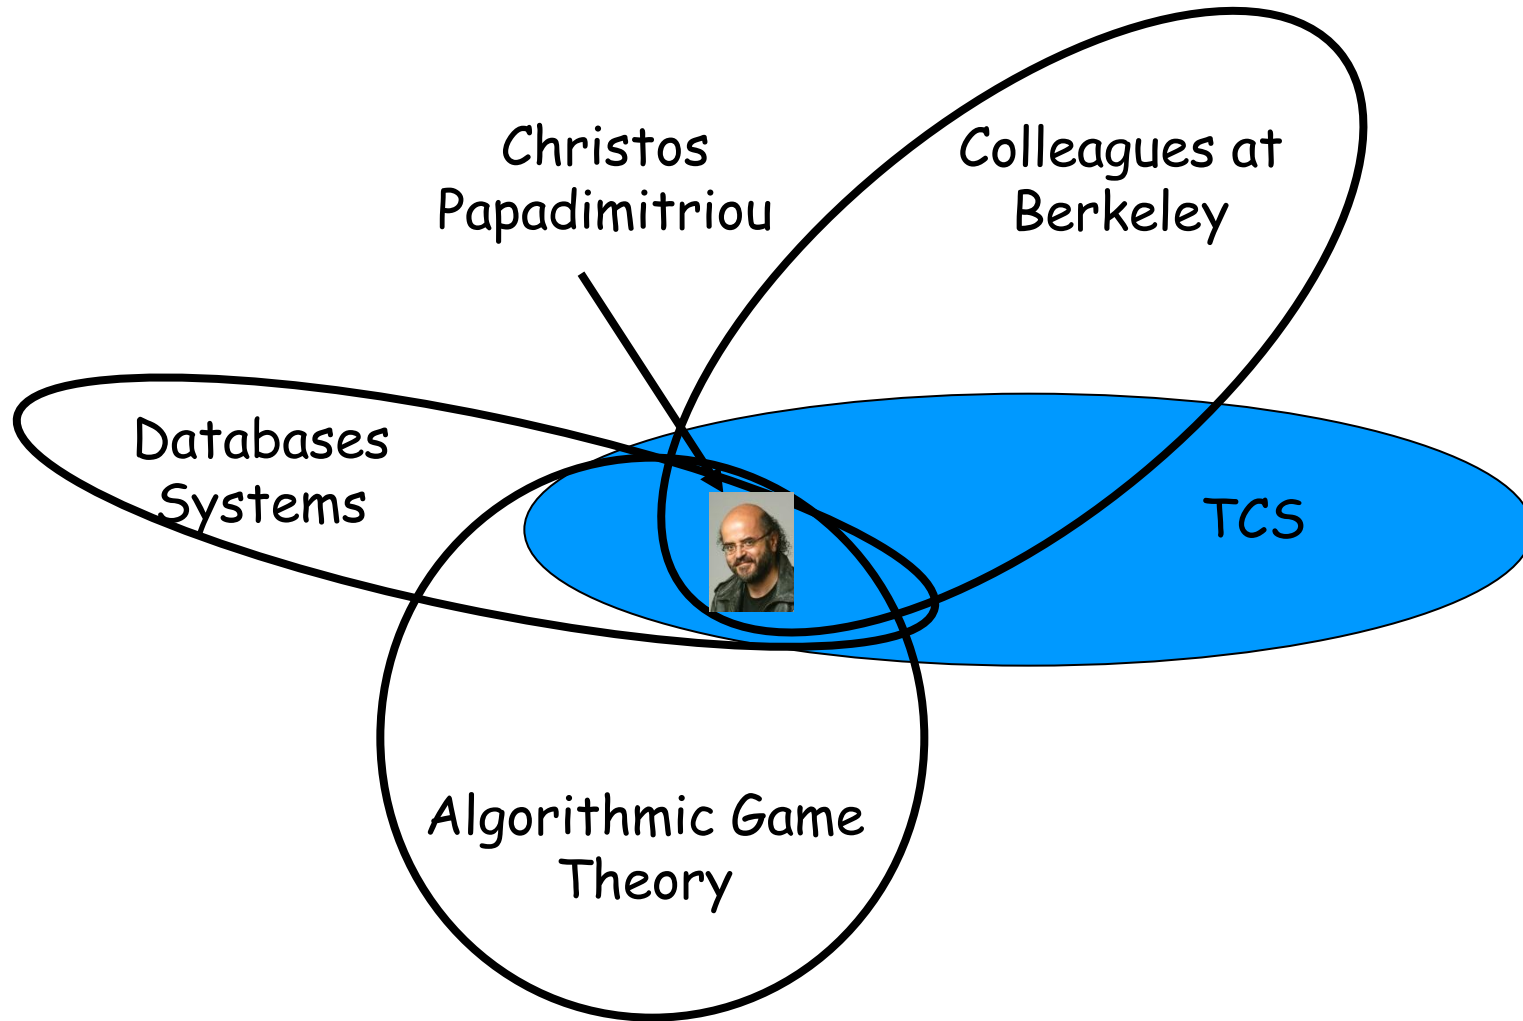
- Imagine a worst-case input.
- But then add small *Gaussian* perturbation to each data point.



Smoothed analysis model

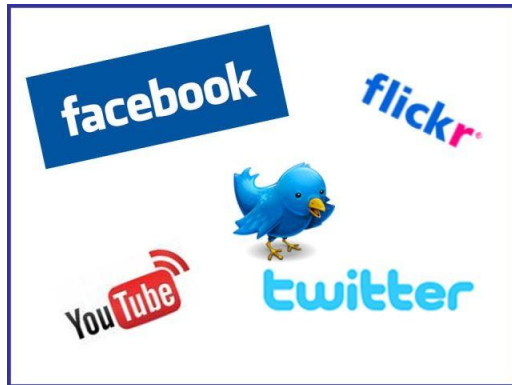
- Imagine a worst-case input.
- But then add small Gaussian perturbation to each data point.
- Theorem [Arthur-Manthey-Roglin 2009]:
 - $E[\text{number of rounds until Lloyd's converges}]$ if add Gaussian perturbation with variance σ^2 is polynomial in $n, 1/\sigma$.
 - The actual bound is : $O\left(\frac{n^{34}k^{34}d^8}{\sigma^6}\right)$
- Might still find local opt that is far from global opt.

Overlapping Clusters: Communities



Overlapping Clusters: Communities

- Social networks
- Professional networks



- Product Purchasing Networks, Citation Networks, Biological Networks, etc

Overlapping Clusters: Communities

